

Molecular hydrogen in EXHALE

Abstract

EXHALE carries H_2 as a single species. The state vector holds one number density for it, the network reacts it as one reagent, and no rovibrational level is transported. Every quantity that depends on how the molecules are distributed over (v, J) therefore has that distribution supplied by assumption — and the assumption is not the same in every place. The infrared line cooling is explicitly LTE and says so in its own function name; the thermal dissociation and its three-body reverse rest on vibrational equilibrium and are now built from one another through the equilibrium constant, so that axis is consistent by construction; the Lyman–Werner photodissociation rate no longer imports a fit at all — it reads a table of a line-by-line calculation, with the lines absorbing each other’s beam, run at the temperature and density of this layer, and states where that calculation itself stops being a value; and the H_3^+ cooling, alone among the four, carries an explicit departure factor. This memo records each assumption, the published evidence behind it, the range over which it holds, and where the layer EXHALE actually computes leaves that range.

1. What is and is not carried

H_2 enters the state vector as species `isp_H2` of `src/modules/init/species_table.f90`, one of the four molecular columns (`H2`, `H2p`, `H3p`, `HeHp`) that the molecular option adds. It is written to `output/Ion_species.txt` and `output/Ion_species_adv.txt` as the number density of the molecule; the four molecular columns are the last four of the file, but their absolute position depends on which metal elements the run carries, so they must be located from the `#` columns schema header rather than counted.

What the state vector does *not* carry is any resolution of the rovibrational ladder. There is one H_2 continuity equation, not one per level, and no state-to-state collisional network. Nothing in the code transports a vibrational temperature, an ortho/para ratio, or a level population of any kind. The rovibrational ladder appears twice as *tabulated data* — the partition function of the equilibrium constant (`mol_rates.f90`) and the line list of the infrared cooling (`molecular_infrared_data.f90`) — but in both places it is evaluated at the local kinetic temperature and never solved for.

The consequence is the subject of this memo. A rate coefficient measured on one distribution of levels does not transfer to another, and EXHALE uses coefficients measured under at least three different distributions in the same cell. Each of them is defensible where it is used, and the point of collecting them is to say exactly where that is.

1.1 Code map

Table 1 lists every file that touches H_2 and what it does with it. Three groups are worth separating in advance, because the level-distribution assumption enters each of them differently: the files that supply a *rate*, those that supply an *emissivity*, and those that supply a *composition*.

1.2 The layer the assumptions are evaluated in, and what the numbers below are read from

Every claim about whether a fit is used inside or outside its range needs a statement of the conditions EXHALE actually reaches. The reference case here is `mol_lyman_werner` of the regression matrix — the hot-Uranus molecular base with

File	What it does with H ₂
init/	
species_table.f90	Declares isp_H2 and the nucleus weights (2 H nuclei per molecule) the element budget counts with
set_IC.f90	Seeds the H ₂ column from the chemical-equilibrium fit at the local pressure and T_0 , capped by the available neutral H
lower_atmosphere/	
mol_rates.f90	The Koskinen et al. (2022) Table-1 network, plus the H ₂ thermochemistry: partition function, equilibrium constant, three-body recombination
lower_column.f90	q_h2_equilibrium, the chemical-equilibrium mixing ratio, and the mean molecular weight of the analytic column built from it
lyman_werner.f90, h2_self_shielding_table.f90	Solomon-process photodissociation; the absolute dissociation cross section and its two branching ratios tabulated from a line-by-line calculation with overlapping lines, and the Draine & Bertoldi (1996) equivalent width that still sets the band share
molecular_infrared_cooling.f90, . . . _data.f90	The H ₂ quadrupole and magnetic-dipole line cooling, and the H ₂ O and CO bands beside it
h3p_cooling.f90	H ₃ ⁺ infrared cooling; H ₂ enters as the collider density of the non-LTE departure factor
nonlinear_system_solver/	
System_HeH_mol.f90, System_HeH_mol_metals.f90	The coupled ionization system: unknown $x(4) = 2n_{\text{H}_2}/n_{\text{H}}$, the H nuclei bound in H ₂
constrained_chemical_equilibrium.f90	The same network in logarithmic densities, with a continuation in the radiation field seeded from the H ₂ dissociation equilibrium
functions/	
composition.f90	h2_mixing_ratio_base, h2_bound_fraction, comp_ntot_bc: the H nuclei bound into H ₂ leave the base particle count
cross_sec.f90	sigma_H2, the total photoionization cross section above 15.4 eV, and frac_H2_dissociative_ionization, the share of it that leaves H + H ⁺ instead of H ₂ ⁺
binary_element_diffusion.f90	H ₂ as a diffusion carrier of the hydrogen component, with mass 2 and its own polarizability
radiation/ and post_process/	
util_ion_eq.f90	fuv_lw_photon_field: the shared 912–1110 Å beam, the H ₂ column, and the band fraction the pumping lines remove
util_ion_eq.f90	he_rec_coupling: H ₂ competes with H I (and, above 24.6 eV, He I) for the He II recombination photons, and wins essentially all of them in the molecular layer
Cool_coeff.f90	ioniz_HeI23S_H2, the He(2 ³ S) + H ₂ Penning and associative ionization rate
post_process_adv.f90	Does <i>not</i> carry H ₂ : the _adv reconstruction is molecule-free by construction, and says so in its header

Table 1: Where H₂ appears. Paths are relative to src/modules/.

Stellar LW flux [erg/cm²/s]: 343.0 — because it is the only default case that carries the Lyman–Werner channel. Measured on its output/Lyman_Werner.txt and output/Ion_species.txt:

Layer definition	cells	r/R_p	T [K]	n_H [cm ⁻³]
$x_{H_2} > 0.5$	227	0.9998–1.127	1149–1307	1.67×10^{12} – 9.27×10^{13}
$x_{H_2} > 0.01$	251	0.9998–1.172	945–1307	6.24×10^{11} – 9.27×10^{13}
$x_{H_2} > 10^{-6}$	—	0.9998–3.853	945–2430	1.98×10^7 – 9.27×10^{13}

with $x_{H_2} = 2n_{H_2}/n_H$; the H₂ density itself peaks at 4.63×10^{13} cm⁻³ in the base cell and falls to 4.42×10^9 cm⁻³ at the 1% edge. The three rows matter because they give three different answers. The often-quoted range 945–2430 K is the extent of gas that carries *any* H₂ at all, and its upper half is gas where H₂ is a 10⁻⁵ trace; the layer where H₂ is actually a constituent spans 945–1307 K. Statements below use the narrow definition unless they say otherwise.

These are relaxation snapshots, not converged solutions. Six of the eight default regression cases — every molecular one, plus `lower_profile` — are pinned at 12,000 steps by a `maxsteps` file and stop there rather than at a convergence criterion. `mol_lyman_werner` ends at $du = 2.24$, i.e. a 224% radial spread in the mass flux $\rho v r^2$, three orders of magnitude from the 10⁻³ threshold a converged run meets; only the two wasp cases converge. Every profile number in this memo therefore describes a relaxation snapshot. That is adequate for the purpose — the question is which decade of density and temperature a fit is being asked about, not the fourth digit — but it is not adequate for quoting a mass-loss rate, and none is quoted here.

And the snapshots do not currently match the goldens. At the time of writing, all eight default cases differ from `backup/regression/golden/` in `Hydro_ioniz.txt` and `Ion_species.txt`. The working tree carries uncommitted source changes, including the one described in §2.2, and the goldens predate them. The numbers above are from the case output directories as they stand, not from the goldens.

2. The level-distribution assumptions

2.1 Infrared line cooling: LTE, stated explicitly

`molecular_infrared_cooling.f90` evaluates the H₂ line emission as a Boltzmann sum,

$$E_{H_2}(T) = \sum_l f_u(T) A_l \Delta E_l, \quad f_u(T) = \frac{g_u e^{-T_u/T}}{Z(T)}, \quad Z(T) = \sum_{lev} g e^{-T_{lev}/T}, \quad (1)$$

and names the routine that returns it `h2_line_emission_lte`. There is no ambiguity about the assumption here; the question is only where it holds.

The line list is Roueff et al. (2019), A&A 630, A58, table 2 (electric quadrupole plus magnetic dipole), and the level ladder comes from the same table. As carried in `molecular_infrared_data.f90` it is `n_h2_lev = 302` levels and `n_h2_line = 1833` lines.

Rotational levels are not the constraint: they thermalize at densities far below anything in this layer. The vibrational levels are. The code records the comparison at its own site. The critical density of $v = 1$ is $n_{crit} = A_{10}/\gamma_{10}$ (Hollenbach

& McKee 1979, ApJS 41, 555, p. 576), and with their $A_{10} = 8.3 \times 10^{-7} \text{ s}^{-1}$ (p. 583, J -averaged, from Turner, Kirby-Docken & Dalgarno 1977) and the rate coefficients of their eq. (6.29) it is 2×10^6 – $6 \times 10^6 \text{ cm}^{-3}$ over 900–1350 K with H_2 as the collider and 5×10^4 – $8 \times 10^4 \text{ cm}^{-3}$ with atomic H. The measured base density of the molecular gate case, $n(\text{H}_2) = 4.63 \times 10^{13} \text{ cm}^{-3}$, is seven to nine decades above either.

Two caveats on that comparison. Neither Hollenbach & McKee nor Burton, Hollenbach & Tielens (1990) gives a *helium* collision rate for $\text{H}_2 v = 1$ — their colliders are H, H_2 and, in the latter, electrons — so a helium-dominated gas is outside both sources; omitting helium from the collider sum understates the de-excitation rate and therefore errs against claiming LTE, not for it. And the figure quoted here previously, 10^9 – 10^{11} cm^{-3} , carried no citation and is three to five decades too high; correcting it widens the margin rather than narrowing it.

Above the base the density falls, but so does the temperature, and the vibrational bands then carry a small share of the emission. That is the quantitative statement that makes the LTE overestimate harmless where LTE is least justified, and it can be evaluated directly from the line list: sum the emission of the lines whose upper level has $v \geq 1$ and divide by the total. Table 2 does so.

T [K]	$E_{\text{H}_2}(T)$ [erg s ⁻¹ molecule ⁻¹]	share from $v_u \geq 1$
400	3.034×10^{-23}	0.72%
500	7.717×10^{-23}	5.70%
600	1.789×10^{-22}	18.2%
1000	2.726×10^{-21}	65.7%

Table 2: Share of the LTE H_2 line emission carried by transitions out of a vibrationally excited upper level, evaluated on the implemented line list. Outside the tabulated range the share is 0.016% at 300 K and 81.5% at 2000 K. The level table carries no (v, J) index, so the vibrational assignment was recovered from the Roueff et al. (2019) source table and matched back onto the 1833 selected lines; classifying instead by $v_u \neq v_l$ (a band transition rather than a pure rotational one) changes the shares by less than 0.1 percentage point up to 1000 K.

The reading is that the LTE assumption is safe at the base, where the density supports it, and that where the density no longer supports it the vibrational bands — the levels the density fails to thermalize — have already stopped carrying the emission. Below about 500 K the error LTE can make is bounded by a few per cent of the total, because that is all the vibrational levels contribute.

2.2 Thermal dissociation and three-body recombination: vibrational equilibrium

The pair R12 ($\text{H}_2 + \text{M} \rightarrow \text{H} + \text{H} + \text{M}$) and R15 ($\text{H} + \text{H} + \text{M} \rightarrow \text{H}_2 + \text{M}$) is one channel run in two directions. Since the change recorded below, `mol_rates.f90` builds R12 from R15 through the equilibrium constant,

$$k_{\text{diss}}(T) = \frac{k_{\text{rec}}(T)}{K_{\text{eq}}(T)}, \quad K_{\text{eq}}(T) = \frac{n_{\text{H}_2}}{n_{\text{H}}^2} [\text{cm}^3], \quad (2)$$

rather than transcribing an independent dissociation fit.

That construction is what makes this axis consistent. K_{eq} is built in `h2_thermochemistry_init` from the translational partition functions and the *rovibrational* internal partition function `q_rovib_H2`, summed over the Morse ladder of H_2

$X^1\Sigma_g^+$ with the ortho/para nuclear-spin weights. It therefore contains vibrational excitation explicitly: both sides of the equilibrium refer to a vibrationally equilibrated gas, and the reverse rate built from it inherits that and nothing else.

The measured coefficient it is anchored on refers to the same state, and the source says so. Breshears & Bird (1973), 14th Symposium (International) on Combustion, 211, read their initial rate off the minimum of the postshock density-gradient dip, which they identify as “the rate corresponding to conditions of vibrational equilibrium in the thermal sense but with negligible extent of dissociation”. The same paper bounds where its rate law applies: the “observation of induction periods for diatomic dissociation clearly indicates that the assumptions implicit in Eqs. (2)–(4) are not valid during the very early course of the reaction; nevertheless, they may be applicable during the ‘steady state’ which follows the induction period.” A hydrodynamic model that resolves nothing faster than a vibrational relaxation time is in that steady state throughout. Their Eq. (3), $k_d/k_r = K_{eq}$, states the detailed-balance relation outright, so deriving one direction from the other is the relation the measured coefficient already carries rather than an assumption added here. The critical density for vibrational thermalization, 10^4 – 10^8 cm⁻³ depending on the collider, sits many orders below the 10^{13} – 10^{15} cm⁻³ of the base, which is the independent check that the premise holds where the rate is evaluated.

Status of this change. R12 as detailed balance is present in the working tree and *not* in the last commit: `git show HEAD:src/modules/lower_atmosphere/mol_rates.f90` still carries the independent Baulch et al. (1992) fit $1.5 \times 10^{-9} e^{-48350/T}$, and `docs/Update_EXHALE.md` ends at section 114 with no entry for the change. The changelog is therefore behind the code on this point.

2.3 Lyman–Werner self-shielding: no longer an imported fit

This was the weakest of the four assumptions, and it is the one that has been replaced by a calculation rather than better-argued. `lyman_werner.f90` took the self-shielding factor from Draine & Bertoldi (1996), ApJ 468, 269 (hereafter DB96), their eq. (37); for a time it took it from the temperature-dependent fit of Richings, Schaye & Oppenheimer (2014) instead. Both are closed forms of the same two-term shape, both were built for interstellar gas, and neither had ever been tested at the temperature of a planetary molecular base.

What was done instead. Both fits were measured against a level-resolved CLOUDY c25.00 calculation run on our own grid — $T = 700$ – 3200 K at $n_H = 10^{12}$ – 10^{14} cm⁻³, columns to 5×10^{21} cm⁻², pure H/He, no dust, thermal Doppler width, and our own band flux — and neither survived it. Over $10^{18} \leq N_{H_2} \leq 5 \times 10^{20}$ cm⁻², the columns where our H₂ sits, DB96 runs from $1.85 \times$ high at 900 K to $2.3 \times$ low at 2700 K, and Richings is low by 3.3 – $4.3 \times$ at every temperature of the layer. The failure is one of *shape* and not of level: the calculated $f_{\text{shield}}(N_{H_2})$ has a trough near 10^{16} – 10^{17} and a shoulder near 10^{19} cm⁻² that a two-term algebraic form of this kind has no freedom to place, so no rescaling of either fit could have worked.

And then the CLOUDY calculation was replaced as well (2026-09-03). It carried two errors of opposite sign that crossed near $N_{H_2} = 2 \times 10^{20}$ cm⁻² and cancelled below it. First, the quantity tabulated was CLOUDY’s `Shield(H2)` column, which is the dissociation rate of H2g alone — the sub-reservoir $v = 0$, $J \leq 8$ that CLOUDY splits off at `ENERGY_H2_STAR = 4100` cm⁻¹ — divided by the H2g density, while EXHALE carries one H₂ species and applied

it to all of it. The vibrationally excited part starts from levels 10^2 to 10^3 times less populated, saturates far later, and keeps dissociating after the H2g lines have gone black; weighting the two reservoirs by their own densities raises the rate by up to a factor 11 at $N_{\text{H}_2} = 10^{19} \text{ cm}^{-2}$. Second, CLOUDY applies the Federman, Glassgold & Kwan (1979) shielding function to each transition using that transition's own optical depth, so no line absorbs any other line's beam. Sec. 135 of Update_EXHALE and docs/p38_line_overlap_shielding.md are the measurement.

What is tabulated now is a line-by-line calculation of the same three quantities on the same $39 \times 9 \times 3$ grid: a plane-parallel slab, a normally incident beam, flat F_λ normalized to $F = 343 \text{ erg cm}^{-2} \text{ s}^{-1}$ over 912–1110 Å, and $\tau(\nu, N) = N \sum_i x_i \sigma_i(\nu)$ summed over every Lyman and Werner transition between 911.75 and 1200 Å on a 4×10^5 -point frequency grid — so every line is shielded by every other line. The profiles are Voigt with $b = (2kT/m_{\text{H}_2})^{1/2}$, the populations LTE at the imposed T , and the line data are Abgrall, Roueff & Drira (2000), the same measurements CLOUDY's `dissprob_*` and `transprob_*` files and the Meudon PDR code carry, so no line-list difference enters the comparison. `h2_self_shielding_table.f90` holds $\log_{10} \sigma_{\text{diss}}$, $\log_{10} p_{\text{single}}$ and $\log_{10} p_{\text{eff}}$, trilinear in $(\log_{10} T, \log_{10} n_{\text{H}}, \log_{10} N_{\text{H}_2})$, with every axis clamped rather than extrapolated. The generator is `src/utls/h2_shielding_table_line_by_line.py`; `docs/h2_self_shielding_cloudy.md` and `src/utls/h2_shielding_table_from_cloudy.py` record the CLOUDY runs, which still supply the fluorescent trapping.

The one piece that is not line-by-line. The trapping ratio $p_{\text{eff}}/p_{\text{single}}$ — the suppression of the branching by re-absorption of the fluorescent decay photons — is read node by node from the same 27 CLOUDY runs and multiplied in, rather than reimplemented as a second and unvalidated version of the same thing. Its geometry is plane-parallel, illuminated on one face and closed on the other. Our layer has the same thick-inward side, which is the side that matters, but it is spherical and open outward. This is the only ingredient of the table that carries a geometry our layer does not have, and it is an open item.

What the level distribution now is. It is an LTE population at the imposed T , not one imported from a 10^2 K interstellar front, and the assumption is checked rather than asserted: at $n_{\text{H}} \geq 10^{12} \text{ cm}^{-3}$ the collisional rates exceed the pumping rate by orders of magnitude, and the CLOUDY runs on the same grid, which solve the populations, return departure coefficients of 1.000 at every level. That is what our base is, and it is precisely the population no published fit was built on. The open item this subsection used to carry, “a calculation of f_{shield} on a 10^3 K collisionally equilibrated population would settle whether the fit transfers”, has been done, and the answer was that it does not.

Three axes, now measured rather than assumed.

- **Density.** The line-by-line pumping has no density in it at all: the populations are LTE, so the tabulated $\sigma_{\text{pump}} = \sigma_{\text{diss}}/p_{\text{eff}}$ is identical to every stored digit across $n_{\text{H}} = 10^{12}$ – 10^{14} cm^{-3} . The axis survives only because the trapping ratio came from CLOUDY, and through it $f_{\text{shield}} = \sigma_{\text{diss}}(N)/\sigma_{\text{diss}}(N_{\text{min}})$ moves by at most 2.5% below $N_{\text{H}_2} = 10^{19} \text{ cm}^{-2}$ and by up to 23% at 10^{21} — measured on the installed table. The absence of a density argument in either published fit was never what was wrong with them.

- **Radiation field.** One incident flux only, $F = 343 \text{ erg cm}^{-2} \text{ s}^{-1}$ over the band. With LTE populations the line-by-line shielding cannot depend on it; in the CLOUDY runs, where the populations were solved, dropping the band flux by four decades (G_0 from 1.25×10^6 to 125) moved f_{shield} by at most 9% anywhere and by under 5% below $N_{\text{H}_2} = 10^{20}$. There is no radiation-field axis to carry.
- **Doppler parameter.** Thermal only, $b = (2kT/m_{\text{H}_2})^{1/2}$, 2.7–4.7 km s⁻¹ over the grid — the same expression the code’s own `h2_doppler_parameter` evaluates, so the table and the code agree on it by construction. The table needs no b argument: b is a function of T , which is already an axis.

Where the table stops being a value, and why that is the honest statement. There is no longer a column above which the tabulated factor is an upper bound. Line overlap — what DB96 §5.2 say produces “the rapid falloff ... for $N_2 > 10^{20} \text{ cm}^{-2}$ ” — is inside the calculation, and at the columns a planetary molecular base reaches it is not a correction but the dominant term. Two calculations differing only in whether the lines absorb each other’s beam give, at 1300 K, ratios of 0.62 at $N_{\text{H}_2} = 10^{19}$, 0.33 at 10^{20} , 0.042 at 10^{21} and 0.0016 at $4.4 \times 10^{21} \text{ cm}^{-2}$.

`h2_shielding_overlap_column`, which returned the column at which the summed equivalent width of the pumping lines reached the whole band ($1.4\text{--}1.8 \times 10^{21} \text{ cm}^{-2}$ at 700 K falling to 2.5×10^{20} at 3200 K), is gone from the module together with the plane it interpolated. The boundary it drew was about a decade optimistic: at the column it returned, the one-line-at-a-time table already stood 7 to 10 times above the overlapping-line answer.

What replaces it is a grid boundary and nothing more. `h2_shield_max_column()` returns the top of the tabulated column axis, $5.0 \times 10^{21} \text{ cm}^{-2}$, above which the edge value is used, and the two run-time warnings that quoted the overlap column now report that extrapolation instead. Our base sits at $N_{\text{H}_2} = 4.4 \times 10^{21}$, inside the axis; only the transient excursion to 5.3×10^{21} crosses it.

What the calculation was checked against. At 100 K it reproduces the DB96 self-shielding function to 2–9 per cent over $10^{16} \leq N_{\text{H}_2} \leq 10^{21} \text{ cm}^{-2}$. That fit rests on a calculation which includes overlap (their §5.2), and it is the one published overlap-including curve there is to check against; the same calculation run one line at a time departs from DB96 exactly where DB96 say overlap sets in, by $3.3 \times$ at 10^{21} and $12 \times$ at 4.4×10^{21} . Against a second code that carries overlap — the Meudon PDR code (Le Petit et al. 2006) with its exact UV line transfer, 30 H₂ levels and 20 global iterations on pure H/He at $n_{\text{H}} = 10^{13} \text{ cm}^{-3}$ — the two agree to 5–22 per cent over $10^{19} \leq N_{\text{H}_2} \leq 10^{21} \text{ cm}^{-2}$ at 900, 1300 and 1800 K. At the illuminated face, the lowest tabulated column $N_{\text{H}_2} = 10^{12} \text{ cm}^{-2}$, the absolute quantities are $\sigma_{\text{pump}} = 2.59 \times 10^{-17}$ (700 K) to $2.88 \times 10^{-17} \text{ cm}^2$ (3200 K) against DB96’s own $3.4520 \times 10^{-18} / 0.135 = 2.557 \times 10^{-17} \text{ cm}^2$, and $p_{\text{single}} = 0.143$ to 0.159, which at 1300 K is 0.1475 against 0.1465 from the CLOUDY runs. The Fortran table reproduces the Python calculation at all 1053 grid points to 3.1×10^{-9} relative on σ_{diss} .

The rate is the tabulated cross section itself (2026-09-02). The module no longer multiplies a constant cross section by this factor. What is tabulated is $\sigma_{\text{diss}}(N_{\text{H}_2}, T, n_{\text{H}})$, the dissociation cross section per *incident* band photon, with the self-shielding and the trapping of the fluorescent decay photons both inside it; f_{shield} survives as the ratio $\sigma_{\text{diss}}(N)/\sigma_{\text{diss}}(N_{\text{min}})$, a diagnostic. The earlier reading of the same runs — that σ_{LW} was 23–60% low and $\langle p_{\text{diss}} \rangle$ 31% high — was a comparison of two different quantities: the line data agree to 0.8 per cent, and the gap is line trapping. One

absorption has *two* branching ratios and the code now carries both from the same table: p_{single} , the untrapped $D/(D + \sum A_{ul})$, which fixes how many fluorescent decays accompany each dissociation, and p_{eff} , which fixes how many band photons the beam loses per dissociation. Sec. 132 of Update_EXHALE and docs/p39_lw_cross_section_sources.md carry the derivation; Sec. 135 and docs/p38_line_overlap_shielding.md carry the rebuild of the table these three are read from.

What DB96 still owns. The band share A that H_2 removes from the shared 912–1110 Å beam is DB96 eq. (39), the closed-form integral of their own eq. (37), and it stays there. Its normalization is the one quantity in this module that no line transfer touches — the *pump* cross section — and there the two sources agree: $2.640 \times 10^{-17} \text{ cm}^2$ at 1300 K at the illuminated face of the line-by-line calculation against 2.557×10^{-17} from DB96’s own Table 1–2, 3 per cent.

The energy the fluorescence leaves behind is computed, not adopted. Burton, Hollenbach & Tielens (1990) adopt 2.0 eV per pump in one sentence of their Appendix A, without a derivation. The same quantity follows from the Abgrall, Roueff & Drira (2000) transition probabilities: summing every decay of every in-band pump, weighted by the $X(v'', J'')$ energy it lands on, gives a mean landing energy of 2.060 eV at 700 K rising to 2.129 eV at 3200 K per *bound* fluorescence — 3 to 6 per cent above the adopted value. `h2_energy_per_bound_fluorescence_eV(T)` replaces the constant, and the photoelectron channel of `electron_energy_degradation.f90` shares it with its own state-resolved branching. The average is over unshielded line weights; with shielded weights it is 2.17–2.25 eV, so inside the layer the term is a lower bound by about as much as the correction it applies.

2.4 H_3^+ cooling: a departure factor, not LTE

The one place where EXHALE does not assume LTE is the H_3^+ infrared cooling, and it is worth stating because it runs the other way from §2.1. `h3p_cooling.f90` multiplies the Miller et al. (2013) LTE emission fit by the departure factor $s(T, n_{\text{H}_2})$ of their Table 6, bilinearly interpolated in $(T, \log_{10} n_{\text{H}_2})$, clamped at the table edges and set to 1 for $n_{\text{H}_2} \geq 10^{14} \text{ cm}^{-3}$. Here H_2 enters as the collider density that sets how far the H_3^+ levels depart from a Boltzmann distribution.

The net exchange is written in Kirchhoff form,

$$\Lambda_{\text{net}} = n_{\text{H}_3^+} 4\pi s(T, n_{\text{H}_2}) [E(T) - W E(T_{\text{rad}})], \quad (3)$$

with W the dilution of the field the unresolved lower atmosphere presents, switched on by the input key `Base IR field` (default `False`). The departure factor multiplies both terms, which it must: $s < 1$ applied to the emission alone would heat a gas hotter than the field, and the form above is exactly zero at $T = T_{\text{rad}}$, $W = 1$.

The factor is doing real work rather than sitting at its LTE value: evaluated over the $x_{\text{H}_2} > 0.01$ layer of `mol_lyman_werner`, s runs from 0.996 in the base cell ($n_{\text{H}_2} = 4.63 \times 10^{13} \text{ cm}^{-3}$) down to 0.438 at the top of that layer ($4.42 \times 10^9 \text{ cm}^{-3}$, $r = 1.172 R_p$). Assuming LTE there would overstate the H_3^+ cooling by a factor 2.3.

The table is a vacuum departure factor — collisions against spontaneous decay, no incident field — so a field strong enough to matter pumps the levels back toward LTE and the true departure is smaller than s . The treatment therefore underestimates the magnitude of the exchange, not its sign or its fixed point.

For comparison, Frelikh & Murray-Clay (2026), ApJ 996, 96 treat the same coolant in pure LTE. On this one channel EXHALE carries the more complete physics.

2.5 The chemical-equilibrium mixing ratio

`q_h2_equilibrium` (`lower_column.f90`) is the Koskinen et al. (2022) Eq. (11) fit to the Visscher et al. (2006) chemical-equilibrium H_2/H partition,

$$q_{\text{H}_2} = \frac{1.9845 + 10^u - \sqrt{10^u(3.9690 + 10^u)}}{2.3670}, \quad u = -\frac{23672}{T} - \log_{10} p[\text{bar}] + 6.2645, \quad (4)$$

with limits $1.9845/2.3670 = 0.83840$ as $u \rightarrow -\infty$ (fully molecular, the fit's approximation of the solar-composition volume mixing ratio $0.5/(0.5 + \text{He}/\text{H}) = 0.863$) and 0 as $u \rightarrow +\infty$.

It is used in four places: the mean molecular weight of the analytic lower column, the base H_2 mixing ratio when no `base.inp` supplies a photochemical one, the initial-condition seed of the H_2 column, and the $\lambda \rightarrow 0$ end of the continuation path in `constrained_chemical_equilibrium.f90`.

The assumption here is not about level populations but about the chemistry: this is *thermochemical* equilibrium, and it knows nothing about photodissociation. It is therefore not a statement about an irradiated wind, and the code says so at its own site — in chemical equilibrium $q_{\text{H}_2}(1 \mu\text{bar})$ stays large up to $T \sim 2000$ K and the gas is fully atomic only above ~ 2400 K, while photochemistry plus the hotter actual base temperature dissociate H_2 well below that. Using it as a *seed* for a solve that then applies the full radiation field is exactly the use it supports; using it as a discriminant for whether a cell should be molecular is not. Section 114 of `docs/Update_EXHALE.md` records an attempt in that direction being withdrawn, together with the repair of an asymptotic guard in the same function that had been inverted since transcription and handed every cell below about 529 K a fully atomic state labelled “the dense molecular limit”.

2.6 Three sites that inherit a ground-state assumption silently

Three further quantities are measured on ground-state H_2 and used without a level-distribution correction. None is flagged in the code as such, and none is likely to matter at the level the three above do, but the list should be complete.

- **Photoionization.** `sigma_H2` in `functions/cross_sec.f90` is the Yan, Sadeghpour & Dalgarno (1998), ApJ 496, 1044 fit (their Eqs. 17–19), threshold 15.4 eV, with the near-threshold resonance structure smoothed as in the source fit. It is a cross section for H_2 in its ground vibrational state; a vibrationally hot molecule has a lower threshold and different near-edge structure. The branching between its two product channels (`frac_H2_dissociative_ionization`, Sec. 3.) carries the same assumption: it is measured on a room-temperature target.
- **Diffusion.** `binary_element_diffusion.f90` gives H_2 the static dipole polarizability $\alpha = 5.315 a_0^3$ (Schwerdtfeger & Nagle 2019) for the ion-neutral coefficient. This is the $v = 0$ value; polarizability rises with vibrational excitation.
- **Metastable-helium Penning ionization.** `ioniz_HeI23S_H2` in `radiation/Cool_coeff.f90` is the Maxwell–Boltzmann average of the Cohen & Lane (1977) cross sections, taken from the Garcia Munoz (2025) network

file, valid over 500–10⁴ K following its tabulation. It is a thermal average over relative translational velocity with the internal state of the H₂ fixed.

3. The two product channels of the H₂ photoionization

A photon absorbed by H₂ above threshold leaves one of two things behind, and the network treats them as different reactions because they are:

$$(a) \quad \text{H}_2 + h\nu \rightarrow \text{H}_2^+ + e^-, \quad E_{\text{th}} = 15.4 \text{ eV}, \quad (5)$$

$$(b) \quad \text{H}_2 + h\nu \rightarrow \text{H} + \text{H}^+ + e^-, \quad E_{\text{th}} = 18.076 \text{ eV}. \quad (6)$$

Channel (b), *dissociative* photoionization, destroys the molecule without making H₂⁺, so it does not feed R8 (H₂⁺ + H₂ → H₃⁺ + H) and does not reach the H₃⁺ cooling at all; what it produces is a proton and an H atom. In the shielded molecular base, where the H I photoionization that normally supplies protons is starved, it is the dominant photon-driven proton source: measured on the converged He/H = 0.0793 hot Uranus, $P_{\text{H}_2, \text{di}} n_{\text{H}_2}$ is 21 times $P_{\text{HI}} n_{\text{HI}}$ at 1.004 R_p and falls through unity near 1.04 R_p .

sigma_H2 is the *total* photoionization cross section, so (b) is a branching of it and not an addition: the H₂ destruction rate is unchanged, and only the products move. The branching is

$$f_{\text{di}}(E) = \frac{\sigma(\text{H}^+)}{\sigma(\text{H}^+) + \sigma(\text{H}_2^+)},$$

tabulated at the 69 energies of Chung, Lee, Masuoka & Samson (1993), J. Chem. Phys. 99, 885, Table II, and interpolated linearly in E . It rises within 0.03 eV of threshold to a plateau near 0.02, carries the Q_1 Rydberg autoionization resonance (0.108 at 35 eV, back to 0.091 at 37.5 eV), peaks at 0.222 at 74 eV, and is held at its 124 eV value 0.2048 above the top of the measured range. The rows of the coupled system this drives are (4) H₂, which sees the whole rate; (5) H₂⁺, which sees $P_{\text{H}_2} - P_{\text{H}_2, \text{di}}$; and (1) H⁺, which gains $P_{\text{H}_2, \text{di}} n_{\text{H}_2}$. The H atom of the pair is supplied by the H-nucleus closure, exactly as the two Lyman–Werner fragments are.

The secondary electrons do the same thing: Dalgarno, Yan & Liu (1999) give the dissociative channel a mean energy per ion of 22 $W_m(\text{H}_2^+)$, so the fast electrons make one such proton for every 22 H₂⁺ ions, an additional yield rather than a share of the H₂⁺ one.

Yan et al. (1998) misdescribe this branching, and the code does not follow them. Their Sec. 4 gives the dissociative “ratio to the total photoionization cross section” as 0.284 at 76 eV and 0.258 at 124 eV. Those are Chung et al.’s H⁺/H₂⁺ *column*, a ratio to the H₂⁺ partial cross section; the ratio to the total is $r/(1+r)$, i.e. 0.221 and 0.205. Using the quoted numbers as a branching of sigma_H2 overstates the channel by a factor 1.26–1.28 above 70 eV and understates it by up to 20× below 20 eV.

4. The third body

R12, R13 and R15 are three-body reactions written as a two-body coefficient times a density. Koskinen et al. (2022) print that density as “ n ” and do not say which particles it counts; their models are H₂/H-dominated Jovian and

Neptunian envelopes, where M is H₂ and H. EXHALE passes the total gas-particle density with electrons excluded (`ion_cell_state%ntot`), so every heavy particle is treated as an equally efficient third body.

The coefficient this assumes uniform is the M = H₂ one: R15 carries Cohen & Westberg's (1983) recommended $k_1(\text{H}_2) = 2.8 \times 10^{-31} T^{-0.6} \text{ cm}^6 \text{ s}^{-1}$ over 50–5000 K, labelled M = H₂ on their data sheet and separate from their $k_1(\text{H})$.

Breshears & Bird measured all three of their colliders in one experiment with one analysis, so their dissociation coefficients are the clean statement of how much the third body matters:

$$\begin{aligned} k_d(\text{M} = \text{Ar}) &= 1.55 \times 10^{-10} e^{-44736/T}, \\ k_d(\text{M} = \text{H}_2) &= 5.48 \times 10^{-9} e^{-52989/T}, \quad [\text{cm}^3 \text{ s}^{-1}] \\ k_d(\text{M} = \text{H}) &= 3.52 \times 10^{-9} e^{-43881/T}. \end{aligned} \quad (7)$$

Evaluated here, $k_d(\text{H})/k_d(\text{H}_2)$ is 8.67 at 3500 K, 6.26 at 4000 K, 3.97 at 5000 K and 2.01 at 8000 K — a factor of a few across their measured 3500–8000 K range. Cohen & Westberg's independent *recombination* pair, $k_1(\text{H}) = 8.8 \times 10^{-33}$ against $k_1(\text{H}_2) = 2.8 \times 10^{-31} T^{-0.6}$, gives 4.21 and 5.21 at 3500 and 5000 K, which is the agreement detailed balance requires of two independent sources inside their ranges.

Below the measured range the two Arrhenius forms diverge, because their activation temperatures differ by 9108 K: the ratio reaches 13.4 at 3000 K, 61.0 at 2000 K and 5.80×10^3 at 1000 K, against 1.98 from Cohen & Westberg at 1000 K. Take a factor of 2 to 5 as the collider spread where H₂ exists, and read the 5.8×10^3 as the cost of extrapolating an Arrhenius fit rather than as physics.

Helium is the collider nobody measured. The equal-efficiency assumption is weakest in two places, and only one of them is quantified above. Above the H₂ front atomic H dominates the collider population, and the factor of 2–5 is the size of the error there — carried, at least, by a vanishing H₂ density. In a helium-rich envelope M is helium, and there is no measurement to compare against at all. Breshears & Bird shocked H₂ in argon and in xenon (5.0, 10.0, 20.0% H₂ in Ar; 5.0, 10.0, 20.2% in Xe) and extracted M = Ar, Xe, H₂ and H; helium appears in their paper only as a shock-driver gas and in “helium–neon laser”. None of the sources behind R12, R13 or R15 supplies a helium efficiency, and none is applied. On general grounds a monatomic third body has no internal modes to take up the released energy and should be *less* efficient than H₂, so in a helium-dominated base these three rates are upper bounds; docs/molecular_chemistry_audit_he_rich.md §4.1 puts the size of a plausible factor at order 0.5 and records it as a guess rather than a measurement. Nothing in this repository supports a helium efficiency near unity, and no such number should be quoted from it.

5. The caloric equation of state of the molecular layer

The thermal equation of state of the code is composition exact: the pressure is $p = (n_{\text{tot}} + n_e)kT$ with the species count of the cell. The caloric half, the relation between the internal energy density and the pressure, was not. It used the monatomic $e = p/(\gamma - 1)$ with $\gamma = 5/3$ everywhere, which asserts that every particle stores $\frac{3}{2}kT$.

Molecular hydrogen does not. Its rotational constant gives $\theta_{\text{rot}} = 85.35 \text{ K}$, so at the 400–1300 K the molecular layer runs the rotational ladder is fully excited and an H₂ molecule stores about $\frac{5}{2}kT$; the vibrational ladder adds to that above

~ 800 K. In a gas that is 80 per cent H_2 the specific heat is therefore wrong by a factor 1.67.

5.1 The mixture

For a mixture in which every atom, ion and free electron carries $\frac{3}{2}k$ and H_2 carries its own rovibrational heat capacity, the internal energy density is

$$e = \frac{3}{2}(n_{\text{tot}} + n_e)kT + n(\text{H}_2)ku_{\text{rv}}(T), \quad (8)$$

with u_{rv} the mean rovibrational energy of one H_2 molecule in units of k (so u_{rv} has the dimension of a temperature). Because $p = (n_{\text{tot}} + n_e)kT$ is unchanged, the exponent that makes $e = p/(\gamma_{\text{eff}} - 1)$ exact is

$$\gamma_{\text{eff}} = 1 + \frac{n_{\text{tot}} + n_e}{\frac{3}{2}(n_{\text{tot}} + n_e) + n(\text{H}_2)c_{\text{rv}}(T)}, \quad c_{\text{rv}} = \frac{du_{\text{rv}}}{dT}. \quad (9)$$

γ_{eff} is a function of the composition *and* of the temperature, and it is not a free parameter: both quantities in (8) are already carried by the composition module.

5.2 Where the level data comes from

u_{rv} and c_{rv} are evaluated from the complete bound rovibrational ladder of $\text{H}_2 X^1\Sigma_g^+$ — 302 levels reaching 51 966 K, the 4.478 eV dissociation limit — which the code already carries for its infrared line cooling (Roueff et al. 2019). The level term energies and the weights $g_l(2J+1)$ come from that same table, so there is no second copy of the H_2 spectroscopic constants in the code to drift out of step with the first.

Using the observed ladder is not a refinement for its own sake. A rigid rotor plus a harmonic oscillator, with $\theta_{\text{rot}} = 85.35$ K and $\theta_{\text{vib}} = 5987$ K, gives $c_v(\text{H}_2)/k$ smaller by 1.4 per cent at 200–900 K, 2.3 per cent at 1300 K and 7.3 per cent at 3000 K.

Because the sum runs over a single partition function with the nuclear-spin weights inside it, (8) is the ortho/para *equilibrium* mixture. A frozen 3:1 mixture gives the same c_{rv} to four decimals at and above 300 K and differs only below ~ 250 K, which is colder than the layer reaches, so the distinction does not arise here.

H_2^+ , H_3^+ and HeH^+ are given $\frac{3}{2}k$ each, as if monatomic. Measured over every cell in which H_2 exceeds 10^{-3} of the particle count, the largest ratios to H_2 are 1.7×10^{-7} (H_2^+), 1.2×10^{-7} (HeH^+) and 3.2×10^{-8} (H_3^+), so giving all three the full H_2 heat capacity would move C_V by less than 10^{-6} .

5.3 Inverting it

With the vibrational ladder in c_v , the map from internal energy to temperature stops being a multiplication. It is inverted exactly, by a bisection-safeguarded Newton iteration on the bracket $0 < T \leq u/1.5$, which holds because u_{rv} is non-negative and increasing in T . It is *not* lagged at a previous temperature: a residual that depended on the previous state would blur both the fixed point of the steady solver and the finite-difference Jacobian it is found with.

To keep the cost of that iteration bounded, u_{rv} is tabulated once on a 4096-point logarithmic grid from 1 K to 5×10^4 K and interpolated by a cubic Hermite in $\ln T$ built on exact nodal values of both u_{rv} and $du_{\text{rv}}/d\ln T$. The interpolant is therefore C^1 and its own derivative — the quantity the Newton iteration differentiates — is exact at the nodes. Against the direct 302-level sum over 200–5000 K the largest relative error is 10^{-12} in u_{rv} and 9×10^{-10} in c_{rv} .

5.4 What carries the local index, and what does not

The physical content of γ sits entirely in the $E \leftrightarrow p$ conversion, and that is where the local value is used: the primitive/conservative conversions, the geometric source, the total energies and the physical flux of the Riemann solver, the low-Mach dissipation pressure, the semi-implicit energy update, the heat capacity of the conduction operator, the state scales of the steady solver, and the advected energy equation of the post-processor. The sound speeds likewise use each cell's own γ_{eff} .

The reconstruction needs no index at all: it works in primitive variables. Each face state's energy and sound speed are built with the heat capacity of the cell it was reconstructed from, which is well defined because the two stored composition quantities are ratios of number densities to ρ — functions of the mass fractions alone — and so are unchanged by the density difference between a cell average and its face value.

One place still needs a single index for both sides: the Roe average $a_{\text{avg}} = \sqrt{(\gamma-1)(H_{\text{avg}} - v_{\text{avg}}^2/2)}$ and the adaptive p_* estimate of Toro's hybrid scheme, both derived for one ideal gas. They are given the arithmetic mean of the two face values. This is a wave-speed estimate, not a definition of the steady state: the converged solution is the root of a residual assembled from the conservative fluxes, and those carry each side's own equation of state. The HLLC branch needs no such average, its Davis–Einfeldt speeds taking one sound speed per side already.

5.5 Reduction to the atomic case

An atomic run must reproduce the constant- γ arithmetic to the bit, and the formula cannot be relied on to do it: with $n(\text{H}_2) = 0$, (9) evaluates $1 + 1/1.5$, which is not the same double as $5/3$. The reduction is therefore structural. A run with molecular chemistry off leaves every map on the legacy expression verbatim, and inside a molecular run a cell with exactly zero H_2 takes the same expression. Every non-molecular case of the regression matrix is byte-identical across the change.

6. Validity ranges

Table 3 collects every range in one place: what each treatment stands on, and what EXHALE asks of it.

One of those rows is about the closure and not about a coefficient, and it is the one that fails. Everywhere in this document H_2 is a root of a local algebraic balance, recomputed from scratch in each cell by `ioniz_eq`. That is a statement about time scales, not a property of the network: it holds where the chemistry is fast against the flow and nowhere else. Evaluated on our own converged solution it does not hold across the H_2 front — the term the closure drops is two decades larger than the whole net chemical rate there (docs/supersonic_molecular_base.md section 13.6) — so the front position and the H_3^+ extent that follow from it are properties of the closure as much as of the chemistry. The transported alternative exists and is described in `Update_EXHALE.md` sections 128 and 134; it is off by default while the question of whether the transported problem has a steady state on that planet is still open.

The transported operator has a validity range of its own, and it is numerical rather than chemical. Its advection was first-order donor-cell until 2026-09-03, carrying a numerical diffusivity $|v|\Delta r/2$ that on the 500-cell grid of the hot Uranus is 8 to 22 times the molecular diffusivity of H_2 over $1.15\text{--}1.50R_p$, so the front it produced was set by the scheme

as much as by the flow. It is now second order — van Leer limited slope added as a deferred correction — which on the coarse grid reproduces what doubling the grid gave. The residual front motion after that fix is grid-converged, so it is the state relaxing and not the discretization; but the state is not steady, the steady solver stalls on the energy row inside the front, and no number from a transported run is a converged one yet (TO_BE_DONE.md item (Y)).

And the way the transported problem was being closed against the wind was itself wrong, in sign. Until 2026-09-03 the steady route alternated a wind solve at a frozen H_2 partition with a carrier relaxation at a frozen wind. That alternation does not converge and cannot: its acceptance gate stood below the smallest move the transport operator can make, the fixed-wind steady state it relaxed to sat on the hydrogen element ceiling over a quarter of the grid, and its convergence measure — the carrier row divided by $n_{\text{H}}(|v| + c_s)/\Delta r$ — understated the row's own imbalance by a factor $10^{2.7}$ to $10^{6.4}$, so a state carrying 0.60 of its own terms at the front was reported at 1.6×10^{-4} . Worse, the splitting moved the H_2 front OUTWARD at 83% of the local gas speed, whereas solving the carriers and the wind together moves it INWARD — from $x_{\text{H}_2} = 0.5$ at 1.142 to $1.062 R_p$ — because H_2 costs particles, fewer particles at the wind's own pressure means a hotter cell, and a hotter cell destroys H_2 . A fixed wind cannot see that feedback. The carrier row now carries the same convergence statement as every other row (its own largest terms), and Coupled carrier solve puts $n(\text{H}_2)$ among the Newton unknowns so that the two halves are solved at once (Update_EXHALE.md section 139). The coupled solve improves the residual by a factor three and the front's own carrier imbalance by four, and then stagnates with the molecular layer 0.6 to 4% out of carrier balance, so the closure question above stands: the transported problem still has no converged steady state on that planet, and the option stays off by default.

Everything else in this section is about how well a rate is known; this row is about whether the equation is the right one. The evaluated column is the $x_{\text{H}_2} > 0.01$ layer of §1.2 unless stated, with the caveats recorded there about relaxation snapshots.

Table 3: What each treatment stands on and what it is asked to do.

Treatment	Range of the data	Range evaluated	Verdict
H_2 LTE line emission	Vibrational thermalization needs $n \gg n_{\text{crit}}(v=1)$; A_{10}/γ_{10} of Hollenbach & McKee (1979) is $2\text{--}6 \times 10^6 \text{ cm}^{-3}$ (H_2 collider) or $5\text{--}8 \times 10^4 \text{ cm}^{-3}$ (H collider) at 900–1350 K	$n(\text{H}_2) = 4.4 \times 10^9\text{--}4.6 \times 10^{13} \text{ cm}^{-3}$	Inside everywhere, by three to nine decades. No helium collision rate is published in either source, so a helium-dominated gas is an extrapolation — one that understates the de-excitation rate
H_2O and CO band cross sections	Tabulated 50–2000 K, clamped outside	945–1307 K in the molecular layer	Inside. Both molecules are largely dissociated above 2000 K in any case

Table 3 (continued)

Treatment	Range of the data	Range evaluated	Verdict
R15 three-body recombination, $k_1(\text{H}_2) = 2.8 \times 10^{-31} T^{-0.6}$	Cohen & Westberg (1983), recommended 50–5000 K; ± 0.2 in $\log k$ at 300 K rising to ± 0.4 at 5000 K	945–1307 K	Inside
R12 thermal dissociation, $k_{\text{rec}}/K_{\text{eq}}$	Inherits R15's range; K_{eq} tabulated 100–20000 K and clamped	945–1307 K	Inside. Against the Baulch et al. (1992) fit the ratio is 0.97–1.18 over that fit's own 2500–8000 K and departs only below it
Third body M = everything	Measured for M = H ₂ , H, Ar, Xe (Breshears & Bird 1973); no helium coefficient exists	Total heavy-particle density	H ₂ -dominated base: inside. Above the front, M is atomic H, a factor 2–5. Helium: unquantified
Lyman–Werner σ_{diss} , column axis	Tabulated 10^{12} – $5 \times 10^{21} \text{ cm}^{-2}$, overlap included; cross-checked against the Meudon PDR code to 10^{21}	up to 4.36×10^{21} , and 5.31×10^{21} in transient	Inside the axis. The transient excursion is 6% past its top and is clamped there; the second-code cross-check stops at 10^{21}
Lyman–Werner σ_{diss} , Doppler axis	Not an axis: the profiles are Voigt with thermal $b = (2kT/m_{\text{H}_2})^{1/2}$ at the tabulated T	thermal $b = 2.79$ – 4.48 km s^{-1}	Inside. h2_doppler_parameter evaluates the same expression, so table and code agree by construction
Lyman–Werner σ_{diss} , level distribution	LTE at the tabulated $T = 700$ – 3200 K ; the CLOUDY runs on the same grid return departure coefficients of 1.000	945–1307 K, n_{H} up to $9.3 \times 10^{13} \text{ cm}^{-3}$, collisional	Inside. Turbulent broadening and any non-thermal population are not carried

Table 3 (continued)

Treatment	Range of the data	Range evaluated	Verdict
Lyman–Werner trapping ratio $p_{\text{eff}}/p_{\text{single}}$	CLOUDY escape probabilities in a plane-parallel slab, illuminated on one face and closed on the other	a spherical layer, thick inward and open outward	Outside on the geometry. The thick-inward side is the one that matters, but the open outward side is not what was computed; the size of the mismatch is unquantified
Band fraction A from DB96 eq. (39)	Their own ceiling, §4.3 eq. (29), is $W_{\text{max}} \approx \ln(1110/912) \approx 0.2$, i.e. $A = 1$	asymptote 1.0843 at $b = 2.79$, 1.1050 at 4.48 km s^{-1}	Above the authors' own cap by 8.4–10.5%; clamped at 1 in the code
H_3^+ emission fit $E(T)$	Miller et al. (2013) Table 5, piecewise 30–5000 K; frozen above	945–1307 K	Inside. The published segments do not join; each is used on its own range
H_3^+ departure factor $s(T, n_{\text{H}_2})$	Table 6 grid $T = 300$ –5000 K, $\log_{10} n_{\text{H}_2} = 6$ –14; clamped, $s = 1$ above 10^{14}	$T = 945$ –1307 K, $n_{\text{H}_2} = 4.4 \times 10^9$ – 4.6×10^{13}	Inside; $s = 0.44$ –0.996 across the layer
q_h2_equilibrium	Koskinen et al. (2022) Eq. (11), validated against NASA CEA at 1000–2500 K; thermochemical equilibrium only	Seed and base-composition use only	Inside as a seed. Not applicable to an irradiated wind: it has no photodissociation

Table 3 (continued)

Treatment	Range of the data	Range evaluated	Verdict
Photoelectron energy degradation in H/H ₂ /He (Dalgarno et al. 1999, electron_energy_degradation.f90)	Fits stated for ionization fraction $x \leq 0.1$; tables at $E_0 = 30\text{--}1000$ eV, clamped outside; helium fixed at $n(\text{He}) = 0.1[n(\text{H}) + 2n(\text{H}_2)]$ and no metal at all; H ₂ in its ground vibrational level, so $T \lesssim 2000$ K	$x = 10^{-16}\text{--}5 \times 10^{-5}$ in the layer and $\rightarrow 1$ in the wind above it; $E_0 > 30$ eV (E_sec_ion); 945–1307 K	Inside on x in the layer, on E_0 and on T . Outside on x in the ionized wind , where eq. (14) is extrapolated above 0.1 with a neutral-fraction closure that forces the $x = 1$ endpoint; that extrapolation is accepted on a measurement against Shull & van Steenberg’s Table 1 rather than on the closure (§118.6 of Update_EXHALE). Outside on helium for any arm at $\text{He}/\text{H} \sim 1\text{--}10$, by one to two orders in the direction of more helium: the branching between targets is renormalized onto the cell’s own densities, but the total ionization budget and the heating fraction still carry the sources’ own 0.1. Outside on metals , whose electrons enter the ionized fraction although neither source has a metal in its gas

Table 3 (continued)

Treatment	Range of the data	Range evaluated	Verdict
H ₂ vibrational energy: where it goes (Hollenbach & McKee 1979; Burton et al. 1990, h2_vibrational_relaxation.f90)	Rate coefficients of HM79 eq. (6.29) fitted “good to 20% for $T \gtrsim 500$ K”; colliders H and H ₂ only (Burton et al. add electrons); no helium rate in either source; $E_* = 2.0$ eV is an adopted value for <i>photon</i> pumping	$T = 945\text{--}1307$ K; colliders are H ₂ , H and He with $n(\text{H}_2) = 4.4 \times 10^9\text{--}4.6 \times 10^{13} \text{ cm}^{-3}$ against $n_{\text{cr}}(v=1) = 5 \times 10^4\text{--}1 \times 10^6 \text{ cm}^{-3}$	Inside on T and, by three to nine decades, on density: the quench fraction is 1.0000 throughout the layer. Outside on helium , which is left out of the collider sum for want of a rate — an omission that lowers the quench fraction, so it errs against claiming the energy as heat. Applying E_* to electron-impact excitation of B and C is an assumption: same upper states and cascade, different distribution over upper vibrational levels
sigma_H2 photoionization	Yan et al. (1998) Eqs. 17–19, $E \geq 15.4$ eV, ground vibrational state, resonances smoothed	the full stellar spectrum above threshold	Inside in energy; the vibrational-state assumption is unquantified
frac_H2_dissociative_ionization	Chen et al. (1993) Table II, 18.076–124 eV, room-temperature target, stated uncertainty $\pm 4\text{--}5\%$	the full stellar spectrum above 18.076 eV	Inside up to 124 eV; held flat above it , where the source has no data. Above ~ 47 eV the measured $\sigma(\text{H}^+)$ counts the protons of the double-ionization channel as well ($\sim 20\%$ of it by 80 eV), so the proton source is right and the co-product assignment is not
He(2^3S) + H ₂ ionization	Cohen & Lane (1977) via Garcia Munoz (2025), tabulated 500– 10^4 K	layer temperatures, evaluated unclamped outside	Inside over the molecular layer

Table 3 (continued)

Treatment	Range of the data	Range evaluated	Verdict
H ₂ solved as a LOCAL chemical equilibrium (the closure itself)	Valid where the chemistry is fast against the flow, i.e. where the Damköhler number $\tau_{\text{adv}}/\tau_{\text{chem}} \gg 1$	The whole domain, in every run that does not set Molecular carrier transport: True	OUTSIDE across the H₂ front. Measured on the converged He/H = 0.0793 hot Uranus, the omitted advection term $ r^{-2}\partial_r(r^2vn_{\text{H}_2}) $ is 44 and 117 times the whole net chemical rate at 1.05 and 1.15 R_p , and only at 1.25 R_p does the chemistry win (0.6). Inside beyond about 1.25 R_p and at a base where $v = 0$. The transport operator that removes the restriction is diffusive_photochemistry.f90; it is default off for a molecular-only run

7. Open items

- **No (v, J) resolution, and why it has not been added.** Carrying the ladder as transported species would mean 302 further continuity equations, the 1833-line radiative network as a level-coupling matrix instead of a summed emissivity, and a state-to-state collisional set that the code does not have and that none of its current sources supply. Against that cost, the place it would most obviously buy something — the self-shielding — is now supplied from outside by a line-by-line calculation that carries the ladder for us (§2.3), so the transported ladder would buy consistency rather than a missing number. The trade is recorded here rather than settled.
- **Self-shielding at our temperature: DONE, and it did not transfer.** This item asked for f_{shield} on a 10^3 K collisionally equilibrated population at the columns of Table 3. It was computed (§2.3), both published fits failed against it, and the rate now reads the calculation.
- **Line overlap: DONE; the geometry of the trapping is what is left.** The successor item — that above $N_{\text{overlap}}(T)$ the tabulated calculation was itself an upper bound, because it shielded one line at a time — is closed. The table is now built with every Lyman and Werner line absorbing every other line's beam, h2_shielding_overlap_column is gone, and the only column boundary left is the top of the grid (§2.3).

What replaces it is narrower: the trapping ratio $p_{\text{eff}}/p_{\text{single}}$ is still CLOUDY's, computed in a plane-parallel slab illuminated on one face and closed on the other, while our layer is spherical and open outward. The thick-inward side is the side that dominates, but the size of the mismatch has not been quantified.

- **An internal inconsistency in the published DB96 fit.** DB96 cap the overlap-corrected equivalent width of the pumping lines at $W_{\text{max}} \approx \ln(1110/912) \approx 0.2$ (their §4.3, eqs. 29–30), i.e. at a band fraction of exactly 1. Their closed-form integral, eq. (39), has the asymptote $1.05(1+0.0117b_5)$ as $N_{\text{H}_2} \rightarrow \infty$, which at our Doppler parameters is 1.0843 ($b = 2.79 \text{ km s}^{-1}$) to 1.1050 (4.48 km s^{-1}): 8.4–10.5% above their own ceiling. Their remark below eq. (39) is only that it “corresponds to an equivalent width $W \approx \Delta \ln \nu \approx 0.2$ in the limit $N_2 \rightarrow \infty$ ”. The overshoot is a property of the published fit, not of our use of it, and is the same whether the band fraction is taken from eq. (39) or from integrating f_{shield} on the grid. `fuv_lw_photon_field` clamps it at 1 so the transmission handed to the continuum absorbers of the same interval cannot go negative.
- **The `mol_ir_bands` regression case does not test what its name says.** It sets `Base IR field: True` and `Molecular IR bands: True` but leaves `Oxygen chemistry` off, as its `EXHALE_resolved.out` confirms, so there are no H_2O or CO carriers. Measured on `output/Cooling_breakdown.txt`, columns 13 and 14 are exactly zero in all 504 rows, while column 12, the H_2 line channel, is nonzero everywhere and peaks at $2.934 \times 10^{-7} \text{ erg cm}^{-3} \text{ s}^{-1}$ at $r = 1.0025 R_p$, $T = 1316 \text{ K}$. The zeros are structural rather than numerical: with $n(\text{H}_2\text{O}) = n(\text{CO}) = 0$ the two net-rate functions return 0 on their own guard. The case gates the H_2 line channel and nothing else, and a case that exercises the H_2O and CO paths would have to add `Oxygen chemistry: True`. *Note also* that the golden directory for this case carries only `Hydro_ioniz.txt`, `Ion_species.txt` and `run.log`; `Cooling_breakdown.txt` is not snapshotted, so the infrared columns are not under regression control at all.

References

- [Abgrall, Roueff & Drira (2000)] H. Abgrall, E. Roueff & I. Drira, *Total transition probability and spontaneous radiative dissociation of B, C, B' and D states of molecular hydrogen*, A&AS 141, 297 (2000).
- [Black & Dalgarno (1977)] J. H. Black & A. Dalgarno, *Models of interstellar clouds. I. The Zeta Ophiuchi cloud*, ApJS 34, 405 (1977).
- [Breshears & Bird (1973)] W. D. Breshears & P. F. Bird, *Density gradient measurements of H_2 dissociation in shock waves*, 14th Symposium (International) on Combustion, 211 (1973).
- [Cohen & Westberg (1983)] N. Cohen & K. R. Westberg, *Chemical kinetic data sheets for high-temperature chemical reactions*, J. Phys. Chem. Ref. Data 12, 531 (1983), p. 559.
- [Draine & Bertoldi (1996)] B. T. Draine & F. Bertoldi, *Structure of stationary photodissociation fronts*, ApJ 468, 269 (1996).
- [Federman, Glassgold & Kwan (1979)] S. R. Federman, A. E. Glassgold & J. Kwan, *Atomic to molecular hydrogen transition in interstellar clouds*, ApJ 227, 466 (1979).

- [Frelikh & Murray-Clay (2026)] R. Frelikh & R. A. Murray-Clay, *Efficiency of hydrodynamic atmospheric escape in hot Jupiters and super-Earths*, ApJ 996, 96 (2026).
- [Koskinen et al. (2022)] T. T. Koskinen et al., *Mass loss by atmospheric escape from extremely close-in planets*, ApJ 929, 52 (2022).
- [Le Petit et al. (2006)] F. Le Petit, C. Nehmé, J. Le Bourlot & E. Roueff, *A model for atomic and molecular interstellar gas: the Meudon PDR code*, ApJS 164, 506 (2006).
- [Miller et al. (2013)] S. Miller, T. Stallard, J. Tennyson & H. Melin, *Cooling by H_3^+ emission*, J. Phys. Chem. A 117, 9770 (2013).
- [Roueff et al. (2019)] E. Roueff et al., *The full infrared spectrum of molecular hydrogen*, A&A 630, A58 (2019).
- [Yan et al. (1998)] M. Yan, H. R. Sadeghpour & A. Dalgarno, *Photoionization cross sections of He and H_2* , ApJ 496, 1044 (1998).