

EXHALE: Update Log

***Purpose.** This document is a running log of incremental updates to EXHALE as the code is extended toward reproducing Huang et al. (2023) WASP-121b (see Huang_update_plan.md). Each update is a self-contained section: what changed, why, which files were touched, and the quantitative check that was run. New updates are appended as new top-level sections.*

Contents

1. Runtime domain-extent option: Roche vs Spherical	5
1.1 Overview	5
1.2 The two domain modes	5
1.2.1 Gravitational potential	5
1.2.2 Outer boundary	6
1.2.3 Why spherical mode must drop the tidal terms	6
1.3 Input interface and backward compatibility	6
1.4 Implementation	7
1.5 Validation on WASP-121b	7
1.6 Status and relation to the Huang plan	8
2. Ionization network: add magnesium (Mg I / II / III)	8
2.1 Overview	8
2.2 Atomic data and rates	8
2.3 Touched files	9
2.4 Validation on WASP-121b	10
2.5 Status and relation to the Huang plan	11
3. Phase 1a: interface / data-structure refactor (no new physics)	11
3.1 Overview	11
3.2 What changed (Steps A–F)	11
3.3 Gate: bit-for-bit reproduction of the C/N/O + Mg run	12
3.4 What is deliberately <i>not</i> generalized yet	12
3.5 Touched files	12
3.6 Status and relation to the Huang plan	13
4. Phase 1b: correct C/N/O cross sections and add Si, Ca, Na, K, S	13
4.1 Overview	13
4.2 Variable-stage solver (2- and 3-stage elements in one packing)	13
4.3 Atomic data and rates (sources, not fabricated)	14
4.4 Rename: System_HeHCO → System_HeH_metals	14

4.5	Validation on WASP-121b	14
4.6	Caveats and gaps	15
4.7	Regression reference for Phase 1c	15
4.8	Status and relation to the Huang plan	16
5.	Phase 1c: add iron (Fe I / II / III)	16
5.1	Overview	16
5.2	Atomic data and rates (sources, not fabricated)	16
5.3	Touched files	17
5.4	Iron abundance	17
5.5	Validation on WASP-121b	17
5.6	Caveats and gaps	18
5.7	Regression reference for Phase 1d / Phase 2	19
5.8	Status and relation to the Huang plan	19
6.	Phase 1d: charge exchange with hydrogen (Huang Table 4)	19
6.1	Overview	19
6.2	Atomic data and rates (sources, not fabricated)	19
6.3	Touched files	20
6.4	Validation on WASP-121b	20
6.5	Caveats and gaps	21
6.6	Regression reference for Phase 2	22
6.7	Status and relation to the Huang plan	22
7.	Phase 2: metal line cooling — validation gate closed	22
7.1	Active line coolants	22
7.2	Channel-by-channel cooling diagnostic (this phase)	22
7.3	Validation on WASP-121b (the gate)	23
7.4	Touched files	23
7.5	Caveats and deferred	23
7.6	Status and relation to the Huang plan	24
8.	Phase 3a: in-code excited H(n=2) + Balmer proton source and heating	24
8.1	Physics and atomic data (sources, not fabricated)	24
8.2	Coupling into the solver	24
8.3	Runtime interface	25
8.4	Diagnostics	25
8.5	Touched files	25
8.6	Validation on WASP-121b (coupling ON vs OFF)	26
8.7	Caveats and gaps	26

8.8	Two analysis gotchas (recorded so they are not repeated)	27
8.9	Status and relation to the Huang plan	27
9.	Phase 2 refinement: Fe I line cooling (NIST f-values + Van Regemorter)	27
9.1	Atomic data and method (sources, not fabricated)	27
9.2	Touched files	28
9.3	Validation on WASP-121b	28
9.4	Status and relation to the Huang plan	29
10.	Phase 2 refinement: density-dependent Fe II line cooling (SE table)	29
10.1	Why the coronal table is wrong at the base	29
10.2	The density-dependent coefficient (multilevel statistical equilibrium)	29
10.3	Step 1 — the 2-D table generator (done)	30
10.4	Step 2 — the 2-D table and bilinear interpolator in Cool_coeff.f90 (done)	30
10.5	Step 3 — wiring the density-dependent Fe II into the cooling assembly (done)	31
10.6	Step 4 — WASP-121b run and the cooling budget (numbers for the reader)	31
11.	Post-process refinement: re-solve metal ionization in the advected wind (pp_metals 2)	32
11.1	Implementation — reuse the coupled residual with H/He pinned	33
11.2	The cell-by-cell sweep in post_process_adv	34
11.3	Clean A/B and what changes	34
12.	Phase 3b: in-code Lyα radiative transfer ($\bar{J}_{\text{Ly}\alpha}$ via escape probability)	35
12.1	The escape-probability field (lya_rt.f90 : jlya_escape_prob)	36
12.2	WASP-121b validation setup (spherical Case A, matched to Huang)	37
12.3	Trend vs Huang Fig. 11 (qualitative — exact match not expected)	37
13.	Phase 4: RLOF / tidal potential (Case B validated; Case D blocker = escape radius, transonic IC added, geometry-limited)	37
13.1	Base gravity for each case (Huang Table 3)	38
13.2	Case B — validated	38
13.3	Case D — real blocker found (escape radius), transonic IC added	38
14.	Phase 5: velocity-broadened transmission spectrum (in progress)	40
14.1	5a — metal resonance lines (Mg II, Ca II, Na I D), spherical — done	40
14.2	5b — Roche-equipotential reconstruction + broadening — implemented	41
14.3	5b — direct Case A comparison: hydro validated, metals over-extended	41
15.	Post-process advection guard at a “breathing” base (2026-06-07)	42
16.	He 2³S triplet + metals simultaneous solve, Lyα-flux guard, and convergence workflow (2026-06-08)	43

17. Metals in the bulk-gas mass/charge/particle budget (eos_metals, 2026-06-13)	46
18. Automatic initial-condition selection (IC mode: auto, 2026-06-13)	47
19. Base boundary conditions and the flux-based convergence decision (2026-06-15)	48
20. Tidal (Roche) potential with an extended outer boundary (2026-06-16)	50
21. Photoionization cross-section benchmarks and a TOPbase high-energy extension of He 2³S (2026-07-01)	50
22. He I (1¹S) photoionization → Verner+1996 by default, with ATES-rate switches (2026-07-01)	52
23. He 2³S photoionization → two VFKY96 (Verner) forms + power-law bridge (2026-07-01)	53
24. Temperature-dependent Penning ionization of He 2³S (2026-07-01)	53
25. He/H diffusive separation — Phase 1 (2026-07-01, experimental, default OFF)	54
26. He/H diffusion — Phase 2 (2026-07-02): ambipolar, thermal diffusion, metals	55
27. Diffusive separation: detailed physics and numerics	56
27.1 Governing equation	56
27.2 Conservative implicit discretization	56
27.3 The $n_{\text{He}}-n_{\text{H}}$ mass coupling: why a self-consistent n_{H} is <i>not</i> needed	57
27.4 The near-base pile-up and the physical cap	58
27.5 Settling discretization: a Péclet hybrid	58
27.6 Ambipolar field, thermal diffusion, and metals	58
27.7 The metal-rescale ratchet bug (found in review, fixed 2026-07-02)	59
27.8 Boundary conditions and validation	59
28. Lower-atmosphere connection: Tier 1–3 (2026-07-02)	59
29. Tier-2 molecular chemistry core: H₂/H₂⁺/H₃⁺/HeH⁺ (2026-07-03)	60
30. Newton–diffusion co-convergence, solver hygiene, and the Lyα lower-BC audit (2026-07-03)	60
31. Tier-3 first light: VULCAN → base.inp (2026-07-03)	61
32. VULCAN lower-atmosphere pre-step: obtained, not included in the tree (2026-07-06)	61
33. TPM.py renamed to EXHALE_transit.py (2026-07-06)	62
34. Sub-level-resolved Balmer opacity and LaRT 2s inclusion (2026-07-06)	62
35. Molecular equilibrium seeding fix, molecular-base coupling, Tier-2 gate refresh (2026-07-16)	62

36. Default H/He rates → Badnell+Mao / Voronov, van Hoof free-free Gaunt factor, free-free charge fix, and secondary ionization (2026-07-17)	63
36.1 What changed and why	63
36.2 Files touched	64
36.3 Quantitative check	65
36.4 Effect on existing results	65
37. He recombination radiation → HI ionization coupling (use_he_rec_coupling, 2026-07-17)	65
37.1 What changed and why	66
37.2 Files touched	66
37.3 Quantitative check (HD 209458 b)	67
A Code-size summary: EXHALE vs. the original ATES (2026-06-07; updated 2026-07-02)	70
B Fortran source inherited unchanged from ATES	71

1. Runtime domain-extent option: Roche vs Spherical

Added 2026-06-04.

1.1 Overview

ATES historically hardcodes the full Roche potential (planetary gravity plus the stellar-tidal and centrifugal terms) and truncates the radial grid at the Hill/L1 radius. For close-in planets this Roche lobe is small; for WASP-121b the outer boundary lands at $r_{\max} \approx 2.01 R_p$. Huang et al. (2023) Case A, by contrast, is a *spherical* model (no tidal terms) extending to $\sim 19 R_p$.

To support both geometries from a single executable, a runtime option was added with two modes:

- **Roche** (default): full Roche potential, domain truncated at the Hill/L1 radius. Identical to the previous ATES behavior.
- **Spherical**: pure planetary potential $\phi = -b_0/r$ (tidal and centrifugal terms dropped), domain extended to a user-specified outer radius.

The mode is selected by two optional lines appended to `input.inp`; when they are absent the code reverts to Roche, so pre-existing input files are unaffected.

1.2 The two domain modes

1.2.1 Gravitational potential

All quantities are in code-normalized units: radius r in units of the planetary radius $R_0 = R_p$, with $b_0 = GM_p\mu/(k_B T_0 R_0)$ the surface Jeans parameter, $M_r = M_*/M_p$ the star-to-planet mass ratio, and $\tilde{a} = a_{\text{orb}}/R_0$ the orbital distance in planetary

radii.

Roche mode. The dimensionless potential and its gradient are

$$\phi(r) = -\frac{b_0}{r} - \frac{b_0 M_r}{\tilde{a} - r} - \frac{b_0(1+M_r)}{2\tilde{a}^3} \left(\frac{\tilde{a} M_r}{1+M_r} - r \right)^2, \quad (1)$$

$$\frac{d\phi}{dr} = \frac{b_0}{r^2} - \frac{b_0 M_r}{(\tilde{a} - r)^2} + \frac{b_0(1+M_r)}{\tilde{a}^3} \left(\frac{\tilde{a} M_r}{1+M_r} - r \right). \quad (2)$$

The three terms are planetary gravity, stellar tidal attraction, and the centrifugal potential of the co-rotating frame.

Spherical mode. Only the planetary term is kept:

$$\phi(r) = -\frac{b_0}{r}, \quad \frac{d\phi}{dr} = \frac{b_0}{r^2}. \quad (3)$$

1.2.2 Outer boundary

- **Roche:** $r_{\max} = (3M_r)^{-1/3} \tilde{a}$, the Hill-radius approximation to the L1 distance (unchanged from before).
- **Spherical:** $r_{\max}$ is set directly by the user-supplied “Outer radius [R_p]” value (in R_p , since the grid is normalized to $R_0 = R_p$).

1.2.3 Why spherical mode must drop the tidal terms

The outer boundary cannot simply be pushed past L1 while keeping the full Roche potential. Beyond L1 the Roche potential *turns over*: the effective gravity changes sign (points toward the star), so a monotonic, steadily accelerating outward wind is no longer defined there and the sonic-point/steady-state treatment breaks down. Removing the stellar-tidal and centrifugal terms restores a monotonic potential ($\phi = -b_0/r$ increases monotonically toward zero as $r \rightarrow \infty$), which is what permits the domain to extend arbitrarily far. This is exactly the spherical geometry of Huang Case A. A guard enforces a sensible configuration: selecting Spherical mode without an “Outer radius [R_p]” > 1 stops the run with an explicit error rather than silently falling back.

1.3 Input interface and backward compatibility

Two optional lines are appended at the end of `input.inp`, after the “Force start:” line:

Domain mode: Spherical

Outer radius [R_p]: 19.0

- “Domain mode:” accepts `Spherical`; any other value (or a missing line) means Roche.
- “Outer radius [R_p]:” is read only in Spherical mode; it is ignored in Roche mode, where the Hill radius is computed automatically.

The reader scans the trailing lines by *keyword* (using `index(line, . . .)`), not by fixed line position, so blank trailing lines and line ordering do not matter, and `input.inp` files written before this change (which lack both lines) keep the default Roche behavior. The active mode is echoed in the setup report `EXHALE_setup.out`.

1.4 Implementation

Touched files.

- `src/modules/init/parameters.f90`: new globals `spherical_domain` (logical, default `.false.`) and `r_out_user` (the outer radius in R_p).
- `src/modules/functions/grav_field.f90`: `phi` and `Dphi` branch on `spherical_domain`.
- `src/modules/files_IO/input_read.f90`: keyword scan of the trailing lines, and the $r_{\max}$ branch with the validity guard.
- `src/modules/files_IO/write_setup_report.f90`: prints the active mode to `EXHALE_setup.out`.

Gravity branch.

```
double precision function phi(r_in)
real*8, intent(in) :: r_in
if (spherical_domain) then
  phi = -b0/r_in
else
  phi = -b0/r_in &
        -b0*Mrapp/(atilde-r_in) &
        -b0*(1.0+Mrapp)/(2.0*atilde**3.0)* &
        (atilde*Mrapp/(1.0+Mrapp)-r_in)**2.0
endif
end function phi
```

1.5 Validation on WASP-121b

Two runs were made that differ *only* in the domain-extent option; `input.inp` was otherwise identical (power-law SED, $\text{He/H} = 0.0851$, trace metals C/N/O on, “Rate/2 + Mdot/2” 2-D flux approximation). Both reached steady state via the proper criterion (momentum constant, $du < du_{\text{th}}$), not a stall or a false convergence.

mode	$r_{\max} [R_p]$	$\log_{10} \dot{M}$	sonic $[R_p]$	$r(50\% \text{ HII})$	$T_{\text{peak}} [\text{K}]$
Roche (default)	2.01	12.68	1.825	1.178	9662
Spherical	19.00	11.81	5.270	1.152	10752
Huang Case A (ref.)	~ 19	12.57	—	—	~ 12000

$\dot{M}$ in $\log_{10} \text{ g s}^{-1}$. Convergence took $\approx 5,392$ iterations (Roche) and $\approx 154,183$ (Spherical, the larger domain).

Observations (for evaluation, not pass/fail).

- The Roche run reproduces the prior Phase-0 Case A result exactly ($\log_{10} \dot{M} = 12.68$), confirming the new optional block does not perturb default behavior.

- Spherical T_{peak} (10,752 K) is higher than Roche (9,662 K) and closer to Huang's $\sim 12,000$ K — the extended domain lets the profile develop past the $2-R_p$ L1 clip.
- Spherical $\dot{M}$ (11.81) is lower than both Roche (12.68) and Huang (12.57), consistent with the loss of tidal enhancement, but ~ 0.76 dex below Huang.
- The sonic point moves outward from 1.83 to $5.27 R_p$ (gentler acceleration without tidal forces).

Caveats. Two confounders remain before the spherical $\dot{M}/T$ can be compared to Huang cleanly: (a) only C/N/O metals are included here, not Huang's full Mg/Fe/Si/Ca/Na/K/S set (Phase 1); and (b) the same $N = 500$ cells now span $1 \rightarrow 19 R_p$, so the inner resolution is much coarser than in the Roche run — a denser-grid spherical re-run may be needed before drawing conclusions.

1.6 Status and relation to the Huang plan

This option closes the Phase-0 geometry gap: the Roche mode is the validated baseline, and the Spherical mode provides the Case A geometry needed for later comparison. The tidal RLOF physics of Huang Phase 4 is a separate, richer treatment (substellar streamline, L1 location, terminator mapping); the present Roche mode is the existing single-streamline truncation, not the full RLOF model. The Spherical mode is the immediate tool for Case A reproduction.

2. Ionization network: add magnesium (Mg I / II / III)

Added 2026-06-05.

2.1 Overview

The first trace metal beyond C/N/O, following the Huang et al. (2023) WASP-121b species set (Mg, Fe, Si, O, C, N, S, Ca, Na, K). Per the Phase-1 plan, elements are added one at a time (order Mg $\rightarrow$ Fe $\rightarrow$ Si $\rightarrow$ Ca $\rightarrow$ Na $\rightarrow$ K $\rightarrow$ S); this section covers **Mg only**. Magnesium is solved to second ionization (Mg I / Mg II / Mg III) inside the existing coupled MINPACK ionization system, raising the species count to 18 and the ratio-unknown count to 11. The change is confined to the ionization/radiation plumbing: there are no hydrodynamic or mass-density changes, and `calc_ne/calc_rho` still count only H/He — metal electrons enter solely inside the MINPACK system.

2.2 Atomic data and rates

Sourced, not fabricated:

- **Photoionization:** Verner et al. (1996) outer-shell fits, with Mg I threshold 7.646 eV and Mg II threshold 15.035 eV (`cross_sec.f90: sigma_MgI, sigma_MgII`).
- **Recombination:** Badnell RR+DR — the same source already used for C/N/O — via `lookup_rec` cases and `rec_MgII` ($\rightarrow$ Mg I) / `rec_MgIII` ($\rightarrow$ Mg II). This is a *deliberate deviation* from the validation-rate prescription

in `Huang_update_plan.md` §1.3 (Shull & van Steenberg 1982 for Mg); it corresponds to the plan's Phase-6 *production* rates, applied early to match the existing C/N/O treatment. The Mg ion/neutral balance below should be read as the Badnell result, not a Huang-faithful one.

- **Collisional ionization:** Voronov (1997) fits, `ion_coeff_MgI` and `ion_coeff_MgII`.
- **Line cooling:** *not yet ported* — `cool_MgI` = `cool_MgII` = 0. No published Mg two-level/CHIANTI cooling fit is implemented (that is Phase 2). Magnesium therefore acts only as an electron donor and a photoelectric-heating channel here, not as a coolant.
- **Charge exchange Mg + H⁺:** omitted in this first cut.

Grid extension below 13.6 eV. Because Mg I ionizes at 7.646 eV, turning Mg on triggers the existing low-IP-metal branch (`thereis_lowIP_metal`), which extends the energy grid and the stellar SED down to 7.646 eV (`set_energy_vectors.f90`). This is required to photoionize Mg I, but has a side effect noted under Caveats.

2.3 Touched files

- `init/parameters.f90`: `n_species` = 18; `e_th_MgI/e_th_MgII`; `X_Mg`; `s_mgi/s_mgii`; `thereis_lowIP_metal`.
- `functions/cross_sec.f90`: `sigma_MgI`, `sigma_MgII`.
- `radiation/Cool_coeff.f90`: Badnell fits, `lookup_rec`, `rec_MgII/rec_MgIII`, Voronov `ion_coeff_MgI/MgII`, `cool_MgI/MgII` (= 0, flagged).
- `radiation/util_ion_eq.f90`: `PH_heat_HHe` and `eval_cool` carry Mg.
- `radiation/ionization_equilibrium.f90`: Mg densities, params slots 49–55, MINPACK unknowns 10/11 (Mg II/Mg III ratios), solution unpack.
- `init/set_energy_vectors.f90`: 7.646 eV sub-grid; allocate/fill `s_mgi/s_mgii`.
- `files_IO/input_read.f90`: `X_Mg` default; `N_eq` = 11; `metals/low-IP` flags. `files_IO/metals_input_read.f90`: Mg I abundance key. `files_IO/write_output.f90`: Mg I/II/III as `Ion_species.txt` columns 17–19.
- Threading through arrays, I/O and post-processing: `init/set_IC.f90`, `files_IO/load_IC.f90`, `init/init.f90`, `EXHALE_main.f90`, `time_step/energy_semi_implicit.f90`, `functions/utilities.f90`, `post_process/post_process_adv.f90`.

`post_process_adv.f90` also received a latent-bug fix: its advection systems fill only the H/He equations, so they now use a local `Neq_adv` rather than the global `N_eq` = 11, which would otherwise feed `hybrd1` uninitialized metal entries.

2.4 Validation on WASP-121b

Four runs with identical input.inp (power-law SED, He/H = 0.0851, Roche domain), differing only in which metals are present. All reached steady state via $du < du_{\text{th}}$. Mg/H is conserved at the input $X_{\text{Mg}} = 3.98 \times 10^{-5}$ across the domain.

run	metals	grid floor	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$
no metals	H/He only	13.6 eV	12.92	13722	1.703
Mg only	Mg	7.646 eV	12.92	13682	1.697
C/N/O	C, N, O	13.6 eV	12.68	9662	1.199
C/N/O + Mg	C, N, O, Mg	7.646 eV	12.76	9828	1.224
Huang Case A	full set	—	12.57	~12000	—

$\dot{M}$ in $\log_{10} \text{ g s}^{-1}$; r in R_p .

Observations (for evaluation, not pass/fail).

- **Clean Mg isolation** (Mg only vs no metals): $\Delta \log_{10} \dot{M} = +0.000$ dex, $\Delta T_{\text{peak}} = -40$ K. Adding magnesium alone leaves the wind essentially unchanged, as expected for a 4×10^{-5} trace species with no line cooling: it neither heats nor cools appreciably, and the soft-UV band it adds is optically thin without other absorbers.
- **Mg ionization structure**: predominantly Mg II (Mg^+) through the lower thermosphere, predominantly Mg III (Mg^{2+}) in the upper thermosphere (Mg III fraction ~ 0.47 at $1.4 R_p$ rising to ~ 0.81 at $1.9 R_p$); neutral Mg I stays subdominant. Same qualitative ordering as Huang+2023 Fig. 12.
- **MINPACK convergence**: the 11-unknown system converges without NaNs across the full temperature range.
- Mg-only $\dot{M}$ (12.92) is $\sim 2.2 \times$ the Huang Case A spherical value (12.57), in the same range as the prior C/N/O Roche run; the metal-set and geometry confounds from Phase 0 still apply.

Caveats and gaps.

- **C I grid-extension confound** (C/N/O + Mg vs C/N/O). The $\dot{M}/T$ offset between these two (+0.080 dex, +165 K) is *not* a magnesium effect: extending the grid to 7.646 eV exposes **C I** (threshold 11.26 eV) to the new 11.26–13.6 eV photons the 13.6 eV-floor C/N/O run never sees, adding C I photoionization heating. This is why the clean Mg measurement uses Mg-only vs no-metals, where C is absent.
- **No Mg line cooling** ($\text{cool_MgI} = \text{cool_MgII} = 0$). Huang’s headline Mg result — Mg II h&k emission dominating the cooling at $1.15\text{--}1.4 R_p$ — cannot appear here by construction. That is Phase 2.
- **No Mg + H^+ charge exchange** yet.
- **Recombination is Badnell**, not the §1.3 validation rate.

2.5 Status and relation to the Huang plan

This completes the **Mg** step of Phase 1 (network expansion). The ionization structure satisfies the qualitative Fig. 12 ordering and the “ $\dot{M}$ within a factor ~ 2 of Case A” part of the gate; the thermal-balance part of Phase 1 is deferred to Phase 2 (Mg cooling). Next element in the planned order: **Fe**. Scratch runs and the comparison notebook live in `mg_validation/`.

3. Phase 1a: interface / data-structure refactor (no new physics)

3.1 Overview

Each metal added so far (C/N/O, then Mg in §2) widened the named-argument lists of `PH_heat_HHe`, `eval_cool`, `calc_column_dens_metals`, `write_output`, the MINPACK system `ion_system_HeH_metals`, and the driver `ionization_equilibrium.f90` by several arguments for each ion, and threaded the MINPACK coefficients for each cell through a hand-indexed flat `params(60)`. This step converts that plumbing to a **species-metadata table + array-indexed (2D) data flow**, so that adding an element becomes “add rows to a table” instead of “thread N new arguments through M routines.” **No physics changes** — the network is still C/N/O + Mg, and the result is gated bit-for-bit against §2 (below).

3.2 What changed (Steps A–F)

- **A — Metadata module** (`init/species_table.f90`, new). One canonical table holds metadata for each ion (length `n_mion` = 12): `f_sp` column (`mion_fsp`), parent element (`mion_elem`), stage / charge² (`mion_stage`, `mion_z2`), photo-ionizable flag and photo-table column (`mion_isphot`, `mion_iphot`), threshold (`mion_ethr`), cooling flag (`mion_iscool`); and metadata for each element (length `n_melem` = 4): nuclear charge (`melem_Z`), neutral-ion index (`melem_i0`), top stage (`melem_top`). The canonical ion order (C I..Mg III = `f_sp` columns 7–18) is defined here once; element indices `iel_C/O/N/Mg`.
- **B — write_output 2D**. Metal ion densities pass as `nm(:, 1:n_mion)` and are written with an implied-do over the table (`((nm(j,i)*n0, i = 1,n_mion))`), not a named write for each ion.
- **C — PH_heat_HHe + column density 2D**. Optical depth, photoheating, and the photoionization rates loop over photo-ionizable ions via `sigma_tab(:,mion_iphot(i))` and `Nm_col`, returning a 2D rate array `P_m(:, 1:n_mion)` (inert top stages stay zero) instead of separate `P_CI..P_MgII` arguments, one for each ion.
- **D — eval_cool metal channels 2D**. Recombination and collisional-ionization coefficients return as 2D arrays (`rec_m`, `aion_m`); bremsstrahlung sums `mion_z2(i)*GF_elem(:,mion_elem(i))*nm(:,i)` over ions. No separate coefficient arguments for each ion.
- **E — Generalized MINPACK system**. `ion_system_HeH_metals` keeps `hybrid1`’s fixed (`N_eq,x,fvec,iflag,params`) signature but (i) receives the metal coefficients for each cell through a

module-level block in `System_HeH_metals` (`set_metal_coeffs` stores `met_ntot/g0/g1/b0/b1/a1/a2`), set by the driver before each `hybrd1` call — safe because the ionization-equilibrium cell loop is serial — and (ii) assembles `fvec` by **looping over elements** (two balance equations each, force-zeroing absent elements). `params` now carries only H/He (1–11), HeITR (12–18), and the C/N/O charge-transfer rates (40–48); the old slots 49–55, one for each element, are gone. The driver builds metal densities, initial guesses, and the solution unpack by looping over `n_melem/n_mion`. The system size is now the formula $N_{eq} = 3 + 2 \times n_{melem}$ (= 11 for the current four elements), set in `input_read.f90`.

- **F — Generalized grid floor.** `set_energy_vectors.f90` no longer hardwires the sub-13.6 eV floor to Mg’s 7.646 eV; it computes `e_sub_low = min over active elements of the neutral threshold mion_ethr(melem_i0(e))` that lies below `e_th_HI`. The trigger flag `thereis_lowIP_metal` is set generically in `input_read.f90` (any active element with a sub-13.6 eV neutral threshold), driven by a new canonical-order abundance array `melem_ab` (declared in `parameters.f90`, populated in `input_read.f90`). For the current set only CI (11.26 eV) and Mg I (7.646 eV) lie below 13.6 eV and the minimum is Mg I, so the floor is unchanged at 7.646 eV when Mg is present. (Consequence: a future C/N/O-without-Mg run would now extend the grid to 11.26 eV, which it did not before — and once Phase 1b adds Na/K/Ca, whose neutral thresholds 5.14/4.34/6.11 eV sit below Mg’s, the floor drops further.)

3.3 Gate: bit-for-bit reproduction of the C/N/O + Mg run

With C/N/O + Mg active and the same `input.inp/metals.inp` as §2, the refactored code reproduces the `mg_validation/cno_mg` reference to round-off: $\log_{10} \dot{M} = 12.76$; `Hydro_ioniz.txt`, `Hydro_ioniz_adv.txt`, and `Ion_species.txt` are byte-identical to the frozen reference. `Ion_species_adv.txt` differs at exactly one number — line 1 (innermost ghost cell), column 7 (the HeITR fraction): the reference holds 4.0048×10^{-5} left in uninitialized stack memory while the deterministic build writes 0.0; every other column of that line and every other line match. This is a pre-existing reference artifact (a HeITR ghost-cell slot never used in the H/He + metals run), not an effect of the refactor.

3.4 What is deliberately *not* generalized yet

- **Charge exchange with H** is still hardcoded for C/N/O only (the `fvec(1)` C,N,O sum and the `cxlo/cxup` mapping via `iel_C/O/N` in `System_HeH_metals`). Generalizing it is Phase 1d.
- **Atomic data for each ion and abundance for each element remain manual entries** — the `sigma_tab(:, 1..8)` assignments and the `sigma_XX` cross-section functions in `set_energy_vectors.f90`, the abundance scalars `X_C/X_O/X_N/X_Mg`, the `metals.inp` keys, and the `melem_ab` population block. These are the atomic-data / abundance rows expected to grow per element in Phase 1b/1c; they are data, not interface plumbing.

3.5 Touched files

- **New:** `init/species_table.f90` (Step A).
- `init/parameters.f90`: `melem_ab(:)` abundance array; generalized `thereis_lowIP_metal` comment.

- `files_IO/input_read.f90`: $N_{\text{eq}} = 3 + 2 \times n_{\text{melem}}$; `melem_ab` population; generalized low-IP-metal trigger.
- `files_IO/write_output.f90`: 2D metal-column writer (B).
- `radiation/util_ion_eq.f90`: `PH_heat_HHe` + column density 2D (C); `eval_cool` metal channels 2D (D).
- `nonlinear_system_solver/System_HeH_metals.f90`: module-level metal block for each cell + element-loop residual (E).
- `radiation/ionization_equilibrium.f90`: metadata-driven density build, params packing (H/He + charge transfer only), initial guesses, solution unpack (E).
- `init/set_energy_vectors.f90`: min-over-active-metals grid floor (F).

3.6 Status and relation to the Huang plan

Phase 1a is complete: the element-by-element plumbing of §2 is now table-driven. Adding the Phase-1b metals should require only metadata rows plus atomic/abundance data, with **no changes to `System_HeH_metals.f90` or `ionization_equilibrium.f90`**. Next: **Phase 1b** — Si, Ca, Na, K, S added together as a uniform-template batch. Gate scratch run: `/tmp/ates_gate/cno_mg/` vs frozen reference `mg_validation/cno_mg/output/`.

4. Phase 1b: correct C/N/O cross sections and add Si, Ca, Na, K, S

4.1 Overview

Two changes are made together here, both enabled by the Phase 1a (§3) table-driven plumbing:

1. **Corrected the C/N/O photoionization cross sections** to the full Verner, Ferland, Korista & Yakovlev (1996; VFKY96) Eq. (1) form. The previous code used the simpler Verner & Yakovlev (1995) inner-shell form (no y_0/y_1 offset/asymptote parameters) for C/N/O; that form is exact only when $y_0 = y_1 = 0$, which does not hold for the C/N/O outer shells. This is the dominant *physics* change.
2. **Added the five uniform-template metals Si, Ca, Na, K, S** in one batch. With Mg (§2) the in-code element set is now C, N, O, Mg, Si, Ca, Na, K, S — the full Huang+2023 §2.4 set *except Fe* (Phase 1c). Counts: `n_species = 30`, `n_mion = 24`, `n_melem = 9`.

The metal MINPACK solver was also **renamed** (below).

4.2 Variable-stage solver (2- and 3-stage elements in one packing)

Huang+2023 carries Si, Ca (like Mg) to second ionization but Na, K, S only to first. The solver keeps a uniform packing of **two unknowns for each element** ($N_{\text{eq}} = 3 + 2 \times n_{\text{melem}} = 21$ with all nine elements) and a `melem_top` flag for each element:

- `melem_top = 2` (C, O, N, Mg, Si, Ca): solve both $X^0 \leftrightarrow X^+$ and $X^+ \leftrightarrow X^{2+}$.

- `melem_top = 1` (Na, K, S): solve only $X^0 \leftrightarrow X^+$ and **pin the unused upper unknown to zero** (`fvec(ix+1) = x(ix+1)`), keeping the Jacobian non-singular. Every `ix+1` access is guarded by `melem_top` so a two-stage element never touches its neighbor's column.

Zero-abundance elements are still force-zeroed, so the same 21-unknown system solves C-only, C/N/O+Mg, or all nine without renumbering.

4.3 Atomic data and rates (sources, not fabricated)

- **Photoionization cross sections** (`functions/cross_sec.f90`): full VFKY96 Eq. (1) form (`sigma_VFKY96`) for C I/II, N I/II, O I/II, Si I/II, Ca I/II, S I; simpler Verner & Yakovlev 1995 form (`sigma_Verner96`, `y_0=y_1=0`) for Mg I/II, Na I, K I (their published fits have `y_0=y_1=0`, so the forms agree exactly). Neutral thresholds [eV]: Si I 8.152, Ca I 6.113, Na I 5.139, K I 4.341, S I 10.36.
- **Collisional ionization**: Voronov (1997) for all new ions.
- **Recombination**: Badnell RR+DR via the swappable dispatcher, as for C/N/O and Mg — the *same deliberate deviation* from the `Huang_update_plan.md` §1.3 validation rates already accepted for Mg (the Phase-6 production rate applied early). Read the balances as the Badnell result.
- **Line cooling**: *not* ported (`mion_iscool = .false.`), flagged To Be Checked; the five act only as electron donors and photoelectric-heating channels (Phase 2).
- **Charge exchange with H**: *not* added; still C/N/O only (Phase 1d).

4.4 Rename: `System_HeHCO` → `System_HeH_metals`

The metal solver was named after the AIOLOS heritage (AIOLOS tracked only C and O as metals $\Rightarrow$ “HCO” = H, C, O). It now carries H + He + nine metals, so the C/O-specific name was misleading. Renamed to parallel the existing `System_HeH`/`System_HeH_TR` family: module `System_HeHCO` → `System_HeH_metals`; `ion_system_HeHCO` → `ion_system_HeH_metals`; `set_hco_metals` → `set_metal_coeffs`; `hco_*` → `met_*`. Pure rename, no behavior change (§3 above is updated to the new names).

4.5 Validation on WASP-121b

Three converged runs, identical planet / SED / grid (power-law SED, He/H=0.0851, Roche domain, “Rate/2 + $\dot{M}$ dot/2”); they differ only in code version and metal set.

key	code	metals	C/N/O σ	floor [eV]	$\log_{10} \dot{M}$	T_{peak} [K]	steps
ref	old	C,N,O,Mg	Verner96	7.646	12.76	9828	5003
A	new	C,N,O,Mg	VFKY96	7.646	12.85	12197	15382
B	new	+Si,Ca,Na,K,S	VFKY96	4.341	12.85	12135	15505

$\dot{M}$ in $\log_{10}$ g/s; Huang+2023 Case A reference $\log_{10}\dot{M} = 12.57$. ref and A share the 7.646 eV grid floor (Mg sets the minimum in the C/N/O+Mg set), so ref vs A isolates the C/N/O cross-section change with no grid-floor confound; A vs B changes both the element set and the floor (K I pulls it to 4.341 eV).

Observations (for evaluation, not pass/fail):

- **C/N/O correction (ref→A):** $\Delta \log_{10}\dot{M} = +0.090$ dex; T_{peak} rises 9828 $\rightarrow$ 12197 K and the peak moves outward from 1.22 to $\sim 1.90 R_p$. The corrected VFY96 opacity deposits more photoheating; the hotter peak now sits near Huang's ~ 12000 K. A genuine model change, not numerical.
- **Adding Si/Ca/Na/K/S (A→B):** $\Delta \log_{10}\dot{M} = +0.000$ dex, $\Delta T_{\text{peak}} = -62$ K, despite the floor dropping to 4.341 eV — at solar abundance and with no line cooling these five are radiatively near-passive donors.
- **Mg moves indirectly.** Mg's own cross sections are byte-identical between ref and A, yet its ionization fractions shift by up to ~ 0.1 : an indirect response to the hotter, more extended profile (changed n_e , optical depth, photoionizing flux), not a change in Mg atomic data. Mg is coupled to the global state, hence *not* an isolated control.
- **Numerics.** Zero NaN/Inf; the 21-unknown system converges over the full T range; X/H for each element is conserved to $\sim 10^{-13}$ for all nine elements including the two-stage Na/K/S — direct evidence the variable-stage packing does not leak between element columns.
- **New-metal ionization (run B).** Si, Ca mostly *doubly* ionized in the upper thermosphere; Na, K, S mostly *singly* ionized (no second stage carried) — the qualitative Huang+2023 Fig. 12 ordering.

4.6 Caveats and gaps

- **No metal line cooling** for Si/Ca/Na/K/S (Phase 2); the thermal structure will not match Huang until Fe II/Mg II/Ca II cooling is in.
- **No charge exchange with H** for the new elements or Mg (Phase 1d).
- **Recombination is Badnell, not the §1.3 validation rate** (deliberate, as Mg).
- $\dot{M}$ is $\sim 1.9 \times$ Huang Case A (12.85 vs 12.57); the metal-set (Fe still missing) and geometry confounds from Phase 0 still apply.
- **Not done here:** Huang Fig. 12 is not digitized/overlaid (only the qualitative ordering is checked); the atomic-data values are not re-derived beyond conservation and smoothness.

4.7 Regression reference for Phase 1c

The accepted run A output is frozen at phase1b_validation/.gate_ref_cno_mg_vfky96/output/ as the post-Phase-1b C/N/O+Mg baseline. When Phase 1c adds Fe, running with Fe *off* must reproduce it; because adding Fe bumps `n_species/n_mion` (and the `Ion_species.txt` column count), the Phase-1c check is a **round-off match on the**

shared physical columns, not a byte-for-byte compare (the `mg_validation/cno_mg` byte-gate applied only to the physics-neutral Phase 1a refactor, where `n_species` was unchanged).

4.8 Status and relation to the Huang plan

This completes **Phase 1b**: the network carries C/N/O/Mg/Si/Ca/Na/K/S, the C/N/O cross sections are corrected to VFKY96, and the variable-stage solver handles 1st- and 2nd-ionization elements uniformly. The gate’s “converges/conserves/Fig. 12 ordering/ $\dot{M}$ within $\sim 2\times$ Case A” items are addressed by the numbers above; the thermal-balance part is deferred to Phase 2 (metal line cooling). Next: **Phase 1c** — iron (Fe I/II/III), which needs non-uniform cross-section (Zatsarinny+2019) and recombination data. Scratch runs and notebook: `phase1b_validation/Phase1b_validation.ipynb`.

5. Phase 1c: add iron (Fe I / II / III)

5.1 Overview

Iron is added as the third three-stage trace metal, completing the Huang+2023 §2.4 element set (C, N, O, Mg, Si, Ca, Na, K, S, **Fe**). No solver, cooling-loop, heating, output, or post-processing code changed: the variable-stage MINPACK packing and every metal loop are generic over `n_melem`, so adding iron is purely new metadata rows plus four new rate functions. Counts: `n_species = 33`, `n_mion = 27`, `n_melem = 10`, `n_mphot = 17`; the system grows to `N_eq = 3 + 2*n_melem = 23`. Iron carries `melem_top = 2` (three stages, like C/O/N/Mg/Si/Ca).

5.2 Atomic data and rates (sources, not fabricated)

Iron is the one element whose data are *not* the uniform-template choice. Each quantity was checked against its primary source before encoding; the Voronov fit in particular was cross-checked because an earlier abandoned attempt had its P/X/K columns transposed.

- **Photoionization cross sections** (`functions/cross_sec.f90`): single-shell VFKY96 form (`sigma_VFKY96`) for Fe I and Fe II, parameters from the Verner+1996 `phfit2.f` PH2 block. Neutral thresholds [eV]: Fe I **7.902** (NIST ionization potential), Fe II **16.199**. *To Be Checked / known limitation*: Huang+2023 use the multi-shell Zatsarinny+2019 Fe I cross section; the single-shell fit here under-represents inner-shell absorption at high XUV energy, so the iron photoionization rate is a lower bound at the hardest photons (flagged inline). A multi-shell upgrade is left for later.
- **Collisional ionization**: Voronov (1997) `cfit.dat`. Fe I (`dE=7.9`, `A=2.52e-7`, `P=0`, `X=0.701`, `K=0.25`), Fe II (`dE=16.2`, `A=2.21e-8`, `P=1`, `X=0.033`, `K=0.45`).
- **Recombination**: the Huang+2023 Eqs. (5)–(6) total fits (a dielectronic Arrhenius term $A T^{-1.5} e^{-T_0/T} (1 + B e^{-T_1/T})$ plus a radiative power law $C (T/10^4)^{-\eta}$) used *directly*, not the Badnell dispatcher used for the lighter metals. Deliberate: the Huang iron rates are the validation target, so the model

runs on exactly those rates; the Badnell-swap for iron is deferred to Phase 6. $\text{Fe II}+e \rightarrow \text{Fe I}$: $A=2.833\text{e-}8$, $T0=5.731\text{e}4$, $B=1.383\text{e}4$, $T1=120.4$; $C=1.248\text{e-}12$, $\text{eta}=0.485$. $\text{Fe III}+e \rightarrow \text{Fe II}$: $A=1.094\text{e-}5$, $T0=1.490\text{e}4$, $B=36.74$, $T1=1.153\text{e}5$; $C=1.728\text{e-}12$, $\text{eta}=0.618$.

- **Line cooling:** *not* ported (`mion_iscool = .false.` for all Fe stages) — Phase 2. This matters more for iron than for the lighter metals: Fe II is the *dominant* low-altitude line coolant in Huang+2023, so a metals-on thermal profile is *expected to run hotter than Huang* until Phase 2.
- **Charge exchange with H:** *not* added; still C/N/O only (Phase 1d). Iron's $\text{Fe}^+ + \text{H}^+ \leftrightarrow \text{Fe}^{2+} + \text{H}$ reactions (Huang Table 4) come in Phase 1d.

5.3 Touched files

Metadata rows + four rate functions + abundance/IC wiring; no solver, cooling, heating, output, or post-process changes:

- `init/species_table.f90` — `n_mion=27`, `n_melem=10`, `n_mphot=17`; Fe I/II/III rows (`f_sp` cols 31–33); `iel_Fe=10`; `melem_top` appends 2.
- `functions/cross_sec.f90` — `sigma_FeI`, `sigma_FeII`.
- `init/set_energy_vectors.f90` — `sigma_tab(:,16)=sigma_FeI`, `(:,17)=sigma_FeII`.
- `radiation/Cool_coeff.f90` — Huang recombination `alpha_rec_FeI/FeII_Huang` with `rec_FeII/rec_FeIII` wrappers; Voronov `ion_coeff_FeI/FeII`; dispatcher cases for FeII/FeIII recombination and FeI/FeII ionization. No `cool_coeff_metal` case for Fe (Phase 2).
- `init/parameters.f90`, `files_IO/input_read.f90`, `files_IO/metals_input_read.f90` — `n_species=33`; `X_Fe`; `metals.inp` keys `FeI/Fe`; `melem_ab(iel_Fe)`; `thereis_metals` trigger.
- `init/set_IC.f90`, `files_IO/load_IC.f90` — neutral-iron IC for `f_sp` cols 31–33.

5.4 Iron abundance

Asplund et al. (2009) solar $\text{Fe}/\text{H} = 3.16\text{e-}5$ ($\log \text{eps} = 7.50$). Huang+2023 adopt $4\times$ solar metallicity for WASP-121b, so the validation run uses $X_{\text{Fe}} = 1.26\text{e-}4$ by number relative to H; C/N/O/Mg are held at solar to isolate the iron change.

5.5 Validation on WASP-121b

Two converged runs (identical planet / SED / grid; power-law SED, $\text{He}/\text{H}=0.0851$, Roche domain, “Rate/2 + Mdot/2”) plus the frozen Phase-1b C/N/O+Mg baseline:

key	metals	iron	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$	steps
ref	C,N,O,Mg	frozen 1b run A	12.85	12197	1.899	15382
off	C,N,O,Mg	off ($X_{\text{Fe}}=0$)	12.85	12197	1.899	15382
on	+Fe	on, $4\times$ solar	12.87	12223	1.915	16566

$\dot{M}$ in $\log_{10}$ g/s; Huang+2023 Case A reference $\log_{10} \dot{M} = 12.57$.

Observations (for evaluation, not pass/fail):

- **Regression, force-zeroed iron (off vs ref).** The Phase-1c binary carries three extra species and two extra MINPACK unknowns, yet run off reproduces the frozen baseline on every shared column to the round-off floor: worst relative difference 2.0×10^{-13} across the seven hydro columns and 3.7×10^{-13} (NI) across the 30 shared species. The three new Fe columns are *identically zero, no NaNs*. With $X_{\text{Fe}}=0$ the iron block reduces to identity rows in the residual ($f_{\text{vec}}=x$, rooted at zero) and contributes exactly zero electrons, so the 23-unknown system returns the 21-unknown answer — adding iron is non-perturbative when off.
- **Iron ionization structure (on).** Fe II dominates the base (fraction 0.89 at $r = 1$, peaking ~ 0.96 near $r = 1.02$); Fe III dominates the escaping flow (0.97 at the outer edge), with the Fe II/Fe III crossover near $r \simeq 1.25$; Fe I is minor everywhere (max ~ 0.11 at the base). This is the qualitative Huang+2023 Fig. 12 arrangement — Fe I confined to the base, Fe II at the base, Fe III through the outflow.
- **Iron as an electron donor.** At $4\times$ solar, iron supplies up to $\sim 26\%$ of the free-electron density, so it feeds back on the whole coupled solution, not just its own columns.
- **Switching iron on (off $\rightarrow$ on).** $\Delta \log_{10} \dot{M} = +0.02$ dex, $\Delta T_{\text{peak}} = +26$ K ($12197 \rightarrow 12223$). Small and in the *heating* direction: with no Fe line cooling yet, iron acts only as an added electron donor / photoelectric-heating channel, so the thermosphere runs slightly hotter, not cooler.
- **Numerics.** Run on converged with zero NaNs; X/H for each element is conserved to 8.3×10^{-13} across all five active elements including iron.

5.6 Caveats and gaps

- **No iron line cooling** (`mion_iscool=.false.`) — Phase 2. Because Fe II is the dominant low-altitude coolant in Huang+2023, the run-on *temperature profile is not yet comparable to Huang* and is expected to sit above it. The Phase-1c physics target is the iron *ionization structure*, not $T(r)$.
- **No iron charge exchange with H** — Phase 1d.
- **Single-shell VFKY96 Fe I/II cross sections** under-represent inner-shell XUV vs Huang’s Zatsarinny+2019 Fe I (flagged in `cross_sec.f90`).
- **Recombination is the Huang fit**, not Badnell (deliberate; the Badnell-swap is Phase 6).
- $\dot{M}$ is still $\sim 2\times$ Huang Case A (12.87 vs 12.57); the missing metal cooling and the Phase-0 geometry confounds still apply.

- **Huang Fig. 12 is not digitized/overlaid** — only the qualitative Fe II-base / Fe III-outflow ordering is checked.

5.7 Regression reference for Phase 1d / Phase 2

The run-on output (C/N/O/Mg + 4×-solar Fe) is frozen at `phase1c_validation/.gate_ref_cno_mg_fe/` (input, `metals.inp`, `run.log`, `output/`) as the post-Phase-1c iron-on baseline. Phase 1d (iron charge exchange) and Phase 2 (iron line cooling) both *change* the iron ionization and the thermal structure deliberately, so this is the “before” snapshot for those comparisons, not a bit-reproduction target.

5.8 Status and relation to the Huang plan

Phase 1c **implementation is complete** and the iron ionization structure is presented for validation against Huang+2023 Fig. 12 (numbers and figures in the notebook; no pass/fail asserted here — that judgement is the reader’s). The ionization-network half of Phase 1 now carries the full Huang element set with verified iron atomic data; the thermal-balance half remains deferred (until Fe II cooling lands in Phase 2 the metals-on temperature profile will not match Huang). Next: **Phase 1d** — charge exchange with H for all metals (Huang Table 4), including $\text{Fe}^+ + \text{H}^+ \leftrightarrow \text{Fe}^{2+} + \text{H}$. Scratch runs and notebook: `phase1c_validation/Phase1c_validation.ipynb`.

6. Phase 1d: charge exchange with hydrogen (Huang Table 4)

6.1 Overview

Phase 1d adds the charge-exchange (charge-transfer, CT) reactions of Huang+2023 Table 4 as a single table-driven module. The default active set is the 23 **Group A** metal + H / H^+ reactions, which couple each metal’s ionization balance to the H^+ fraction. A runtime switch `cx_full 1` in `metals.inp` enables the rest of the table — **Group B** (He + H, 2 rows), **Group C** (metal + He, 6 rows), **Group D** (metal + metal, 32 rows). This *replaces* the previous hard-coded Kingdon & Ferland (1996) C/N/O + H charge-exchange terms in the solver with a generic implementation covering all ten metals plus He.

Re-sourcing C/N/O to Huang Table 4 has one deliberate consequence: Table 4 tabulates only the neutral $\leftrightarrow$ singly-ionized CT for C, N, O (rows A13–A18), so the switch *drops* the doubly-ionized $\text{C}^{2+}/\text{N}^{2+}/\text{O}^{2+} + \text{H}$ recombination terms the old KF96 implementation carried. This is an intentional behavior change to match the paper; it raises the C/N/O double-ionization fractions and reshapes the thermal profile (below). No solver structure, cooling loop, heating, output, or post-processing code changed beyond the CT wiring.

6.2 Atomic data and rates (sources, not fabricated)

All 63 rate rows are transcribed verbatim from `docs/charge_exchange_table4.md`, the image-verified transcription of Huang+2023 Table 4 (23 Group A + 2 Group B + 6 Group C + 32 Group D). The paper text says “65 reactions”; the 2-row discrepancy is immaterial and noted in the reference doc.

- **Functional forms** (radiation/charge_exchange.f90): KF96 bracket $a T_4^b (1 + c e^{dT_4}) e^{-E_k/T_4}$ (kf); power law $a (T/300)^b e^{-E_k/T_4}$ (gp); 4th-order $\ln T$ polynomial (p4) for the radiative-CT rows of N / S / Na / K (Lin/Zhao/-Dutta/Watanabe); and the Satta (2013) $f_{\text{Si}}(T)$ and $D_{\text{CSi}}(T)$ helpers for the Si + He and C + Si rows.
- **Spot checks at $T = 10^4$ K** (against the table): $\text{Fe} + \text{H}^+ = 4.000 \times 10^{-9}$ (RV72), $\text{Fe}^+ + \text{H} = 2.011 \times 10^{-12}$, $\text{Fe}^{2+} + \text{H} = 1.260 \times 10^{-9}$ (ND87), $\text{O} + \text{H}^+ = 2.091 \times 10^{-9}$ (near-resonant, S99) — all match.
- **Rate clamp** $[0, 10^{-6}] \text{ cm}^3 \text{ s}^{-1}$: the upper bound guards the $\ln T$ polynomials, which diverge as $T \rightarrow 1$ K (never reached); the lower bound zeroes a few KF96 bracket fits that dip slightly negative below their validity floor (e.g. $\text{Mg}^+ + \text{H}^+$ at $T \lesssim 1500$ K), where the true rate is negligible.
- **Reference-doc notes**: Ca has *no* direct H charge exchange (only Group D); K has only the forward $\text{K} + \text{H}^+$ row. Reverse rates for pairs lacking direct data are Huang’s detailed-balance estimates, carried verbatim.

6.3 Touched files

A new module plus minimal wiring; the solver’s old element-by-element C/N/O CT block is deleted, not extended:

- radiation/charge_exchange.f90 (**new**) — the full Table-4 descriptor arrays, the cx_rate dispatcher, and the generic residual assembly. Entry points: cx_init (build the active list from cx_full), cx_set_cell(T) (rates in each cell), cx_add_to_fvec (add CT terms to the MINPACK residual via the cx_fvidx element/stage $\rightarrow$ row map), and the cx_full switch. The zero-abundance guard is automatic: each rate is $k n_{\text{donor}} n_{\text{acceptor}}$.
- nonlinear_system_solver/System_HeH_metals.f90 — the hard-coded C/N/O + H charge-exchange terms are *removed*; a single call cx_add_to_fvec(...) now feeds CT contributions into the H, He, and every metal-stage residual row.
- radiation/ionization_equilibrium.f90 — call cx_set_cell(T_K(j)) before each cell’s hybrd1 solve.
- files_IO/metals_input_read.f90 — parses cx_full <0|1> from metals.inp.
- files_IO/input_read.f90 — if (thereis_metals) call cx_init after the input read, before the ionization sweep.

6.4 Validation on WASP-121b

Two new converged runs (identical planet / SED / grid as Phase 1c) compared against the frozen Phase-1c iron-on output as the regression reference:

key	charge exchange	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$	steps
ref	KF96 C/N/O + H (frozen 1c on)	12.87	12223	1.915	16566
cx	Huang Table 4 Group A (default)	12.89	12955	1.778	13942
full	Huang Table 4 full, cx_full 1	12.88	12951	1.765	11661

$\dot{M}$ in $\log_{10} \text{ g/s}$; Huang+2023 Case A reference $\log_{10} \dot{M} = 12.57$. All three runs converged with **zero NaNs**.

Observations (for evaluation, not pass/fail):

- **H⁺ sink magnitude.** In the ionized bulk ($x_{\text{HII}} > 0.5$) the net metal charge exchange is at most 3.7×10^{-3} of radiative recombination (median 1.0×10^{-3}) — negligible as an H⁺ sink, consistent with Huang §4.2 (“Fe–H CX lower than H photoionization and recombination”). At the near-neutral base ($x_{\text{HII}} \sim 7.8 \times 10^{-8}$) the *ratio* is large (~ 118) only because H⁺ is itself a trace species there.
- **Near-resonant balance.** $\text{O} + \text{H}^+ \leftrightarrow \text{O}^+ + \text{H}$ is near-resonant and near-equilibrium at the base (gross destroy 3.66×10^5 vs produce 3.69×10^5 , net $+3.0 \times 10^3$). The net base H⁺ sink is instead driven by the fast, one-sided $\text{Fe} + \text{H}^+ \rightarrow \text{Fe}^+ + \text{H}$ channel (net -2.9×10^4).
- **C/N/O double ionization rises** (the dropped 2+ CT). Column-max fractions ref $\rightarrow$ cx: O III 0.871 $\rightarrow$ 0.946, N III 0.897 $\rightarrow$ 0.944, C III 0.851 $\rightarrow$ 0.856.
- **Iron base shift.** Fe II base fraction 0.889 $\rightarrow$ 0.896; Fe III essentially vanishes at the base (0.002 $\rightarrow$ 0.000): $\text{Fe} + \text{H}^+ \rightarrow \text{Fe}^+ + \text{H}$ pushes iron toward Fe⁺.
- **Thermal structure.** T_{peak} rises 12223 $\rightarrow$ 12955 K and moves inward (1.915 $\rightarrow$ 1.778 R_p). This is the *combined* effect of (i) the C/N/O re-sourcing that drops the 2+ recombination terms and (ii) the new metal + H CT. Charge exchange adds *no* radiative-loss channel, so it cannot cool the gas; Fe II / metal line cooling remains deferred to Phase 2, and the metals-on profile therefore still runs *hotter* than Huang.
- **Group A vs full Table 4 (cx vs full).** Turning on Groups B/C/D changes the bulk temperature by at most 1.5% (worst relative difference), with max $|\Delta T| = 556$ K at $r = 1.008$; at $r = 1.02$ the full run is 357 K cooler with more base H⁺ and less He II. $\dot{M}$ moves -0.01 dex and the outflow is nearly unchanged.
- **Dynamics.** v_{out} and sonic radius: ref 18.83 km/s @ 1.884 R_p ; cx 20.28 km/s @ 1.826 R_p ; full 20.40 km/s @ 1.826 R_p . The hotter, more inward thermal peak drives a slightly faster wind with the sonic point $\sim 0.06 R_p$ further in.

6.5 Caveats and gaps

- **No metal line cooling yet** (Phase 2). Charge exchange adds no loss term, so $T(r)$ is not comparable to Huang; the C/N/O re-sourcing makes the base hotter, not cooler.
- **C/N/O 2+ + H terms intentionally dropped** to match Table 4 — a real change from the Phase-1c (KF96) reference, and the dominant driver of the O III/N III rise rather than the new metal CT alone.
- $\dot{M}$ still $\sim 2 \times$ Huang Case A (12.88–12.89 vs 12.57); the missing metal cooling and the Phase-0 geometry confounds still apply.
- **Paper lists 65 reactions, 63 transcribed** — immaterial, documented in charge_exchange_table4.md.
- **No Huang figure digitized/overlaid** for the H⁺ budget; only the relative-magnitude statement (net CT $\ll$ recombination in the bulk) is checked.

6.6 Regression reference for Phase 2

The default Group-A run is frozen at `phase1d_validation/.gate_ref_cno_mg_fe_cx/` (input, `metals.inp`, `run.log`, output/) as the post-Phase-1d baseline. Because the C/N/O re-sourcing deliberately shifted both the ionization and the thermal structure away from the Phase-1c frozen reference, *this* snapshot — not the Phase-1c one — is the “before” reference for Phase 2 (metal line cooling, including Fe II).

6.7 Status and relation to the Huang plan

Phase 1d **implementation is complete**: the 23 Group-A metal + H reactions are active by default and the full Table 4 sits behind `cx_full`. With this, the **ionization-network half of Phase 1 is physically complete** relative to Huang+2023 §2.4 + Table 4 — the full element set, verified atomic data, and the metal–H charge-exchange coupling are all in place. The thermal-balance half remains deferred: until metal line cooling (Fe II in particular) lands in **Phase 2**, the metals-on temperature profile stays above Huang. Next: **Phase 2** — metal line cooling. Scratch runs and notebook: `phase1d_validation/Phase1d_validation.ipynb`.

7. Phase 2: metal line cooling — validation gate closed

Phase 2 brings the CHIANTI-based metal line cooling (Mg I/II, Ca II, Na I, Fe II, on top of the C/N/O coolants) into the coupled energy balance and **closes the Phase 2 validation gate** against Huang et al. (2023) Fig. 10. The atomic data for each coolant and the CHIANTI effective-cooling tables are documented in `Update_EXHALE_early_phase` (Parts II–III); this section records the integrated WASP-121b gate and the diagnostic used to close it.

7.1 Active line coolants

The line-cooling ions are flagged by `mion_iscool` in `species_table.f90` (canonical order): C I, C II, O I, O II, N I, N II, Mg I, Mg II, Ca II, Na I, Fe II. Each contributes $\beta_{\text{esc}} n_e n_{\text{ion}} \Lambda_{\text{ion}}(T)$ with the effective cooling coefficient Λ_{ion} interpolated from the hardcoded `cool_logL_*` tables in `Cool_coeff.f90` ($N_{\text{COOLT}} = 41$, $\log_{10} T = 3.0\text{--}5.0$, $\Delta \log_{10} T = 0.05$). In the optically thin limit $\beta_{\text{esc}} = 1$ (see `Update_EXHALE_early_phase`, Part II, “Adopted assumption: 100% escape”). N I/II carry the `iscool` flag but have no tabulated line cooling ($\Lambda = 0$), so they contribute exactly zero.

7.2 Channel-by-channel cooling diagnostic (this phase)

To close the gate quantitatively, `eval_cool` (`util_ion_eq.f90`) gained an optional `cool_chan` output — an **exact** decomposition of the total cooling rate into H/He recombination, collisional ionization, collisional excitation, bremsstrahlung, and one column per metal ion. It is read straight from the arrays `eval_cool` already builds, so the channel sum reproduces the cool total to machine precision (no offline re-derivation, no divergence risk). A companion writer `write_cool_breakdown_eq` (same module, called once from `EXHALE_main.f90` after the final equilibrium write) dumps output/`Cooling_breakdown.txt`: columns r/R_p , T , n_e , `cool_total`, the four H/He channels, then 27 metal-ion columns. The run prints `max |sum(channels)/cool - 1|` as an internal consistency check.

7.3 Validation on WASP-121b (the gate)

Converged equilibrium run (full metals, the curated `metals.inp`; same SED/grid as Phase 1d):

- Internal consistency: $\max|\sum \text{channels/cool} - 1| = 3.7 \times 10^{-16}$ (machine eps) — the breakdown is a bit-exact decomposition of the `Hydro_ioniz.txt` cool column.
- In the $1.15 < r/R_p < 1.4$ band the dominant radiative coolants are **Fe II** (~30%) and **Mg II** (~25%), jointly ~55%; metals are **57–85%** of all radiative cooling through the band. Fe II leads the inner edge (top coolant at $1.1\text{--}1.2 R_p$, 53% at $1.1 R_p$), Mg II takes over by $1.3\text{--}1.4 R_p$, and H recombination only overtakes above $\sim 1.5 R_p$ — the Fe II $\rightarrow$ Mg II handoff Huang report in their Fig. 10.
- $\dot{M}$ unchanged ($\log_{10} \dot{M} = 12.51$ g/s); zero NaNs.

This matches the Phase 2 gate (“a metal-driven cooling feature appears at $1.15\text{--}1.4 R_p$, Fig. 10”) and `Huang_update_expected_results` (Mg II + Fe II become the most important coolant in $1.15 < r/R_p < 1.4$).

7.4 Touched files

- `src/modules/radiation/util_ion_eq.f90` — optional `cool_chan` out-arg on `eval_cool` (exact channel-by-channel split); new `write_cool_breakdown_eq` writer; `mion_fsp` added to the `species_table` use-list.
- `EXHALE_main.f90` — use `utils_ion_eq`, only: `write_cool_breakdown_eq`; call it after the final equilibrium `write_output`.
- Cooling coefficients themselves live in `Cool_coeff.f90` (provenance in `Update_EXHALE_early_phase`, Parts II–III); unchanged here.

7.5 Caveats and deferred

- The ion-by-ion metal decomposition is exact in the default branch (`use_2lev_cool = .false.`, the WASP-121b setting). In the optional two-level fine-structure branch the terms for each ion are the resonance-line approximation and need not sum to the two-level `cool_M`.
- **Fe II** uses Boltzmann/coronal lower-level populations ($E_{\text{cut}} = 38459 \text{ cm}^{-1}$), validated vs Huang Fig. 6 coronal curve within $\sim 1.5\times$ — sufficient for the gate. The Hollenbach & McKee IR fine-structure + density-dependent two-level refinement (Huang “Sum” curve) is deferred.
- **Fe I** line cooling (NIST f -values + Van Regemorter, Huang Fig. 5) is not yet added; Fe I is a minor coolant in the gate band. (*Now added in §9; confirmed negligible — peak 3×10^{-5} of total cooling.*)
- The **CHIANTI Ly α swap** is intentionally skipped in Phase 2; the $0.35\times$ Ly α is a Phase-3 radiative-transfer suppression, not a cooling-rate change.
- Carrying metal cooling into the post-processed `_adv` profile is handled separately (runtime `pp_metals`, default frozen); see the post-process notes.

7.6 Status and relation to the Huang plan

Phase 2 is complete and the gate is closed (user-validated 2026-06-05). The metals-on thermal balance now reproduces the Huang Fig. 10 Fe II/Mg II cooling feature, closing the thermal-balance half of the network that Phase 1 left open. Next: **Phase 3** — in-code excited $H(n=2) + Ly\alpha$ radiative transfer.

8. Phase 3a: in-code excited $H(n=2) + \text{Balmer}$ proton source and heating

Added 2026-06-05.

Phase 3a is the first half of Phase 3: an **in-code $H(n=2)$ population model** whose output feeds back into both the hydrogen ionization balance (an extra proton source) and the energy equation (Balmer photoelectric heating), fully coupled. The second half (a $Ly\alpha$ radiative-transfer solve to supply a self-consistent mean intensity $\bar{J}_{Ly\alpha}(r)$) is deferred; this phase uses a **parameterized $\bar{J}_{Ly\alpha}$** with a switch to read an external RT profile instead. Everything is **default-off**: with no stellar parameters in `input.inp` the build is byte-identical to the Phase-2 result (`use_excited_H = .false.`).

8.1 Physics and atomic data (sources, not fabricated)

- **$H(2s)/H(2p)$ populations**: a 2×2 statistical-equilibrium system following Christie, Arras & Li (2013) — recombination cascade into $n=2$, collisional $2s \leftrightarrow 2p$ mixing, two-photon $2s \rightarrow 1s$ decay, $Ly\alpha$ $2p \rightarrow 1s$ with a $2p$ radiative-pumping term, and $2s/2p$ photoionization. Solved per cell (`n2_populations`).
- **$Ly\alpha$ mean intensity** (the $2p$ pump): parameterized as $\bar{J}_{Ly\alpha} \approx 0.1 F_{Ly\alpha} / \Delta\nu_D$ (Huang et al. 2017, Eq. 6) with a **top-down $1/(1 + \tau_{Ly\alpha})$ line-center attenuation** so the optically thick base ($\tau_{Ly\alpha} \sim 1.1 \times 10^8$ here) is not over-pumped. Oscillator strength $f_{Ly\alpha} = 0.4162$.
- **$n=2$ Balmer-continuum photoionization rate Γ_2** : from a **diluted stellar blackbody** above the $n=2$ threshold $E_2 = 3.40$ eV (`gamma_n2_balmer`), with $\sigma_{LyC} = 6.3 \times 10^{-18} \text{ cm}^2$ and threshold cross-section $s_2 = 1.4 \times 10^{-17} \text{ cm}^2$. Γ_2 carries the **same ξ dayside-dilution factor as ground-state EUV ionization** (`Rate/2` $\rightarrow 0.5$, `Rate/4` $\rightarrow 0.25$); for WASP-121b's “`Rate/2 + Mdot/2`” ($\xi = 0.5$) this gives $\Gamma_2 \approx 613 \text{ s}^{-1}$, comparable to Huang's $\approx 576 \text{ s}^{-1}$.
- **Heating**: Balmer photoelectric heating $H_{pe} = n_2 \cdot (\text{photoelectron energy} \times \Gamma_2)$ per cell (`heat_n2_balmer`); optional collisional de-excitation heating behind `incl_deexc_heat` (off by default — it overlaps the existing HI collisional-excitation cooling).

8.2 Coupling into the solver

The model runs **decoupled (lagged one outer step)** to keep it out of the cell-by-cell Newton solve: `excited_H_update` fills two global arrays (one entry per cell) from the previous converged state, and `ioniz_eq` injects them:

- `gph_balmer_HI [s-1]` is added to the HI photoionization rate `P_HI` — an extra **proton source** ($S_{\text{proton}} = \Gamma_2 n_2$).
- `heat_balmer [erg cm-3 s-1]` is added to the photoheating rate.

excited_H_update is called once per outer iteration in EXHALE_main.f90 before ioniz_eq, and once more after convergence to refresh the diagnostics. The post-loop refresh deliberately does *not* re-solve ioniz_eq (the Balmer heating is $\sim 10^{-4}$ of the total, not worth desyncing T/f_{sp} from the converged hydro state).

8.3 Runtime interface

Append to input.inp (omit for the default Phase-2-identical behavior):

Stellar Teff [K]: 6459.0

Stellar radius [R_sun]: 1.458

These set use_excited_H (both must be > 0). Optional further tail lines: Jlya RT file: <path> switches $\bar{J}_{\text{Ly}\alpha}$ to **mode 1** (read an external RT profile and use it directly, instead of the mode-0 parameterized field); Deexc heat: True enables collisional de-excitation heating. As with all trailing lines the scan is by keyword, so order and blank lines do not matter.

8.4 Diagnostics

write_excited_H dumps output/Excited_H.txt (14 columns) from the converged state: r/R_p , T , n_{HI} , n_e , $\bar{J}_{\text{Ly}\alpha}$, n_{2s} , n_{2p} , $S_{\text{proton}}(n=2)$, H_{pe} , H_{dx} , $\tau_{\text{Ly}\alpha}$, then the three **proton-budget** columns added in this phase — **col12 ground-state photoionization** ($P_{\text{HI}} n_{\text{HI}}$), **col13 collisional ionization** ($C n_e n_{\text{HI}}$), **col14 recombination** ($\alpha_B n_e n_{\text{HII}}$), all volumetric [$\text{cm}^{-3} \text{s}^{-1}$]. The three rate coefficients are captured inside ioniz_eq on the converged pass (globals gph_ground_HI, cion_HI, arec_HII), so the budget is self-consistent with the solver rather than re-derived offline.

8.5 Touched files

- init/parameters.f90 — Phase-3a globals (use_excited_H, T_star_eff, R_star, gamma2_bal, hpe2_bal, jlya_mode, jlya_rt_file); feedback arrays gph_balmer_HI, heat_balmer; diagnostics Jlya_arr, n2s_arr, n2p_arr, Sproton_arr, Hpe_arr, Hdx_arr; and the proton-budget capture arrays gph_ground_HI, cion_HI, arec_HII.
- radiation/excited_hydrogen.f90 (**new**) — excited_H_update, n2_populations, gamma_n2_balmer, heat_n2_balmer, load_jlya_rt, write_excited_H.
- radiation/ionization_equilibrium.f90 — inject gph_balmer_HI/heat_balmer under use_excited_H; capture the ground-state / collisional / recombination coefficients for the proton-budget dump.
- files_IO/input_read.f90 — scan the stellar / Jlya RT file / Deexc heat tail lines; set use_excited_H, jlya_mode.
- EXHALE_main.f90 — in-loop excited_H_update before ioniz_eq; post-loop refresh + write_excited_H.
- Makefile — excited_hydrogen.f90 in SRC.

8.6 Validation on WASP-121b (coupling ON vs OFF)

Two converged runs, same build, differing only in the `use_excited_H` toggle (stellar tail lines present/absent). Both reach steady state via $du < du_{th}$ with $\log_{10} \dot{M} = 12.51$ (unchanged from Phase 2 — the coupling does not destabilize the energy solver). Diagnostics in WASP-121b/{`analyze_excited_H.py`, `analyze_proton_budget.py`, `compare_on_off.py`}.

Proton budget (cf. Huang Figs. 11/27), volumetric rates [$\text{cm}^{-3} \text{s}^{-1}$]:

r/R_p	S_{n2}	S_{ground}	S_{recomb}	S_{n2}/S_{ground}	S_{n2}/S_{source}
1.10	5.1e4	6.6e6	6.6e6	0.008	0.008
1.30	1.2e4	4.3e5	4.4e5	0.029	0.028
1.50	4.9e3	7.6e4	8.1e4	0.064	0.060
1.80	1.9e3	1.5e4	1.7e4	0.125	0.111
2.00	2.0e3	5.9e3	7.9e3	0.34	0.25
2.01	2.9e3	5.1e3	7.9e3	0.57	0.36

Observations (for evaluation, not pass/fail):

- **Self-consistency:** at the base $S_{\text{ground}} \approx S_{\text{recomb}}$ (6.6×10^6 vs 6.6×10^6 at $1.1 R_p$), i.e. photoionization equilibrium holds; collisional ionization is negligible (T too low).
- **n=2 share:** the n=2 Balmer channel is a minor proton source below $\sim 1.5 R_p$ ($\sim 1\text{--}6\%$ of ground-state photoionization) but climbs to **$\sim 34\text{--}57\%$ of the ground-state rate ($\sim 25\text{--}36\%$ of the total proton source) near $2 R_p$** , where the gas is already ionized and EUV is attenuated (ground-state photoionization depleted) while the thin-Ly α -fed n=2 channel is not. This is the “significant proton source approaching $2 R_p$ ” behavior Huang describe, here concentrated near $2 R_p$ rather than flat across $1\text{--}2 R_p$.

Heating budget (cf. Huang Figs. 10/26): Balmer photoelectric heating is present in the budget, peaking at **$\sim 1.8\%$ of total heating** at the Ly α photosphere ($\tau_{\text{Ly}\alpha} \sim 1$ near $r \sim 2 R_p$) and $\sim 10^{-4}$ globally — a small but nonzero entry, dominating nowhere, as expected.

ON vs OFF temperature and ionization: the coupling makes the gas slightly **more ionized** (up to $\sim 24\%$ less neutral H near $2 R_p$, Δx_{HII} up to $\sim 1.1 \times 10^{-2}$) and slightly **cooler** (up to ~ 179 K at $r \sim 2 R_p$, mean ~ 42 K). The cooling is **physical, not numerical:** the extra n=2 photoionization converts $\text{HI} \rightarrow \text{HII}$, lowering n_{HI} ; the dominant EUV photoelectric heating scales with n_{HI} , so $\sim 24\%$ less n_{HI} at $r \sim 2 R_p$ means correspondingly less local EUV heating there, which the small Balmer H_{pe} does not offset.

8.7 Caveats and gaps

- **Parameterized $\bar{J}_{\text{Ly}\alpha}$ (mode 0).** The $0.1 F_{\text{LyC}}/\Delta\nu_D$ field with $1/(1 + \tau_{\text{Ly}\alpha})$ attenuation is a placeholder for the Phase-3b Ly α RT solve. Because our n=2 contribution is peaked near $2 R_p$ rather than flat across $1\text{--}2 R_p$, if Huang’s Figs. 11/27 are flatter/larger across the band the parameterized field is under-pumping n=2 in $1.1\text{--}1.8 R_p$; mode 1 (external RT $\bar{J}_{\text{Ly}\alpha}$) is the path to test this once a Ly α RT result exists.

- **Convergence stop is momentum-uniformity** ($du < du_{th}$), not a thermal/ionization steady-state criterion, so a small residual convergence-path sensitivity rides on top of the physical ON–OFF difference.
- **De-excitation heating off by default** (overlaps the existing H I coex cooling); left behind `incl_deexc_heat`.

8.8 Two analysis gotchas (recorded so they are not repeated)

- `Hydro_ioniz.txt` columns 6/7 (heat/cool) are **already physical** (`write_output` multiplies by q_0); do *not* multiply by q_0 again in analysis.
- `Ion_species.txt` column 1 is r/R_p , so $n_{HI} = \text{col2}$, $n_{HII} = \text{col3}$ — the recombination sink uses col3, not col2.

8.9 Status and relation to the Huang plan

Phase 3a is complete and the proton-source comparison is user-accepted (2026-06-05). The H(n=2) model is fully coupled into the ionization balance and the energy equation; the n=2 photoionization proton source reproduces the qualitative Huang Figs. 11/27 behavior and the Balmer photoelectric heating appears in the budget without destabilizing the solver. Next: **Phase 3b** — a $Ly\alpha$ radiative-transfer solve to replace the parameterized $\bar{J}_{Ly\alpha}$ with a self-consistent $\bar{J}_{Ly\alpha}(r)$ (feeding mode 1).

9. Phase 2 refinement: Fe I line cooling (NIST f -values + Van Regemorter)

Added 2026-06-05.

Fe I line cooling was one of the two coolants left deferred when the Phase 2 gate closed (§7, “Caveats and deferred”). This update adds it. CHIANTI v11 has **no Fe I model atom**, so — following Huang et al. (2023) §2.5 — Fe I cooling is built from NIST permitted-line oscillator strengths plus the Van Regemorter collision-strength approximation, with the lower levels populated in Boltzmann equilibrium over the metastable manifold. The new coolant flows through the same $\beta_{esc} n_e n_{FeI} \Lambda_{FeI}(T)$ channel as the other Phase-2 metals.

9.1 Atomic data and method (sources, not fabricated)

The cooling coefficient per $(n_e n_{FeI, total})$ in the optically thin limit [$\text{erg cm}^3 \text{ s}^{-1}$] is

$$\Lambda(T) = \sum_{\text{lines}} f_l(T) q_{lu}(T) \Delta E_{lu},$$

$$f_l(T) = \frac{g_l e^{-E_l/kT}}{U(T)}, \quad U(T) = \sum_{\substack{e3en \text{ le3els} \\ E < E_{cut}}} g e^{-E/kT}, \quad q_{lu} = 2.16 a^{-1.68} e^{-a} T^{-3/2} f_{lu},$$

with $a = \Delta E/kT$ and the absorption oscillator strength f_{lu} from the NIST A_{ki} , statistical weights, and wavelength (Huang Eq. 10).

- **Raw data:** the auditable NIST dumps `fe1_nist_lines.tsv` (permitted E1 lines: A_{ki} , g_i , g_k , E_i , E_k) and `fe1_nist_levels.tsv` (every level g , E for the partition function), fetched by `fetch_fe1_nist.py`.

- **Metastable manifold = EVEN-parity levels below the lowest odd level** (z^7D° , $19351\text{ cm}^{-1} = 2.40\text{ eV}$). Below that energy there is no lower level for an E1 decay, so the even levels are genuinely metastable and hold the Fe I population. This is *not* a plain energy cut: the Fe II treatment can use an energy cut because its cut conveniently lands at the first odd level, but for Fe I a pure energy cut would wrongly include the radiatively-depopulated z^7D° odd levels. The default cut yields **17 even levels and 828 permitted cooling lines**.
- Upper levels of the cooling lines are odd (E1) and assumed to decay radiatively. The builder is `cooling_data/fe1_cooling.py`; it emits the 41-point `cool_logL_FeI(NCOOLT)` Fortran block on the ATES grid ($\log_{10} T = 3.0\text{--}5.0$, $\Delta \log T = 0.05$).

Coefficient values (per $n_e n_{\text{FeI}}$, $\text{erg cm}^3\text{ s}^{-1}$): $\Lambda(5 \times 10^3\text{ K}) = 1.27 \times 10^{-22}$, $\Lambda(10^4\text{ K}) = 1.23 \times 10^{-20}$, $\Lambda(2 \times 10^4\text{ K}) = 1.31 \times 10^{-19}$ — the rising, then flattening shape of Huang Fig. 5.

9.2 Touched files

- `cooling_data/fetch_fe1_nist.py`, `cooling_data/fe1_cooling.py` (**new**) — raw NIST fetch and the Boltzmann-manifold + Van Regemorter builder.
- `radiation/Cool_coeff.f90` — `cool_logL_FeI(NCOOLT)` table constant; new `cool_FeI(T, out)` (interpolates the table via `interp_cool_table`); dispatch cases FeI in `cool_coeff_metal` and `cool_coeff_metal_scalar`. The table-header comment block now documents Fe I alongside the other coolants.
- `init/species_table.f90` — `mion_iscool(FeI) = .true.` (canonical ion index 25), so the `eval_cool` metal-cooling loop and the channel-by-channel breakdown both pick up Fe I automatically.

9.3 Validation on WASP-121b

Re-ran the converged full-metals WASP-121b equilibrium (same SED/grid as the Phase 2 gate) with Fe I cooling on; the channel-by-channel `Cooling_breakdown.txt` then gives the Fe I share directly. Diagnostics: `WASP-121b/analyze_FeI_cooling.py` (left panel = the ported $\Lambda_{\text{FeI}}(T)$ coefficient read straight from `Cool_coeff.f90` vs the Fe II coronal curve, Huang Fig. 5 style; right panel = cooling in each channel vs radius, Huang Fig. 10 style), figure `WASP-121b/FeI_cooling_validation.png`. Observations (for evaluation, not pass/fail):

- **Fe I is a negligible coolant** along the whole profile: its peak share of the total cooling is 3.0×10^{-5} at $r = 1.06 R_p$ ($T = 8303\text{ K}$), and it is far smaller elsewhere ($\sim 3 \times 10^{-8}$ of total at the base, where Fe II dominates). Adding Fe I does *not* perturb the gate-band Fe II $\rightarrow$ Mg II handoff.
- **Mdot unchanged** at $\log_{10} \dot{M} = 12.51$ (Fe I adds no measurable cooling), zero NaNs.
- This is the expected outcome: at WASP-121b base temperatures iron is almost entirely Fe II (Fe I fraction < 0.11 , §5), and Fe I's permitted lines have high excitation energies, so its Boltzmann-suppressed cooling cannot compete with the Fe II / Mg II forbidden + resonance lines. Fe I is carried for completeness and to mirror Huang's coolant list, not because it shapes the WASP-121b energy budget.

9.4 Status and relation to the Huang plan

Fe I line cooling is **implemented and its coefficient matches the Huang Fig. 5 shape/magnitude**; the WASP-121b run confirms it is radiatively negligible there, surfaced as a figure for the reader's judgment (no pass/fail asserted). One of the two Phase-2 deferred cooling items is now closed. The remaining one — the Fe II density-dependent (IR fine-structure) low-temperature refinement — is §10.

10. Phase 2 refinement: density-dependent Fe II line cooling (SE table)

Added 2026-06-05. (Implemented step-by-step; this section is updated as each step lands — currently: the 2-D table generator is built and verified.)

This is the second deferred Phase-2 cooling item (§7): replacing the **coronal** Fe II line-cooling coefficient with a **density-dependent** one that follows the Fe II level populations from the coronal (low- n_e) limit into the collisionally saturated (high- n_e , LTE) limit. It directly affects the cool, dense base of the WASP-121b thermosphere, where the coronal table overestimates Fe II cooling by $\sim 10^4 \times$.

10.1 Why the coronal table is wrong at the base

The Phase-2 Fe II table is the optically-thin **coronal** limit: only the ground level is populated and every collisional excitation is assumed to be followed by a radiative decay, so the cooling is strictly $\propto n_e$ (per $n_e n_{\text{FeII}}$ it is a function of T alone). That assumption fails at the cool, **dense** base of a hot Jupiter thermosphere, because the low- T Fe II cooling there is carried almost entirely by **forbidden** metastable / a⁶D fine-structure lines ($A \lesssim 5 \times 10^{-3} \text{ s}^{-1}$, far-IR 26–87 μm), whose critical densities

$$n_{\text{crit}} = A_{ul} / \sum_c n_c k_{ul}^c \sim 10^4\text{--}10^7 \text{ cm}^{-3} \quad (\text{electron collider})$$

sit far *below* the base $n_e \sim 10^8 \text{ cm}^{-3}$. Those lines are therefore collisionally **saturated** (thermalized to their LTE populations); the coronal coefficient overestimates their cooling by $\sim n_e / n_{\text{crit}}$.

10.2 The density-dependent coefficient (multilevel statistical equilibrium)

cooling_data/fe2_cooling.py solves the full CHIANTI multilevel statistical-equilibrium (SE) problem for the Fe II level populations at each (T, n_e) — 175 levels, electron-collisional excitation/de-excitation (Burgess–Tully descaled Υ from the .scups file) plus radiative decay (.wgfa A -values), normalized to $\sum_i n_i = 1$ — so the cooling smoothly interpolates between the two limits:

$$\frac{\Lambda_{3\text{ol}}}{n_{\text{FeII}}} = \sum_{u>l} n_u A_{ul} \Delta E_{ul} \quad [\text{erg s}^{-1} \text{ per ion}], \quad \Lambda_{\text{eff}}(T, n_e) = \frac{\Lambda_{3\text{ol}} / n_{\text{FeII}}}{n_e} \quad [\text{erg cm}^3 \text{ s}^{-1}].$$

At low n_e , $\Lambda_{\text{eff}} \rightarrow \Lambda_{\text{coronal}}(T)$; at high n_e it falls off $\sim 1/n_e$ as the numerator saturates. This Λ_{eff} is exactly the per- $(n_e n_{\text{FeII}})$ coefficient EXHALE already tabulates, so the existing assembly $\text{cool_M} = \beta_{\text{esc}} n_e \sum_i n_{\text{ion}}(i) \Lambda(i)$ carries it unchanged — the n_e in the prefactor times Λ_{eff} gives $\beta_{\text{esc}} n_{\text{FeII}} \times (\text{cooling per ion})$, i.e. the correct saturated base cooling.

Why electron-only SE is sufficient at the (mostly neutral) base. In the saturated limit the upper levels reach their **Boltzmann (LTE) populations regardless of which collider thermalizes them**, and the cooling of each ion becomes collider- and n_e -independent. At the base $n_e \sim 10^8$ already exceeds n_{crit} for the dominant forbidden lines, so the electron-only SE already reaches that LTE plateau; adding the (more abundant) neutral-H collider would only push the few still-unsaturated higher- n_{crit} lines further toward the same LTE limit — a second-order correction. An explicit H I / proton collider channel is flagged as a future refinement in the module.

10.3 Step 1 — the 2-D table generator (done)

`fe2_cooling.py` gained an **optimized** table builder `lambda_eff_table(Tgrid, negrid)` and Fortran emitters (`fortran_grid_1d`, `fortran_block_2d`, `emit_fortran_table`; run `python3 fe2_cooling.py --table`). The optimization is required because the naive cell-by-cell solve rebuilds a cubic spline for every one of the ~ 4300 transitions at every grid cell: instead the descaled Υ is evaluated **once per transition over the whole T-grid** (`upsilon()` accepts an array T), and the rate matrix is split $M = n_e Q_{\text{coll}}(T) + A_{\text{rad}}$, so the n_e -loop only reassembles and solves. The full 41×29 table builds in ~ 3 s and matches the naive cell-by-cell solve to 1.5×10^{-15} (machine precision).

The emitted table is on the ATES temperature grid (`NCOOLT` = 41, $\log_{10} T = 3.0\text{--}5.0$, $\Delta \log T = 0.05$) $\times$ an electron-density axis `cool_logne` (`NCOOLNE` = 29), $\log_{10} n_e = 0\text{--}14$, $\Delta \log = 0.5$ (brackets the WASP-121b run's $\log_{10} n_e \approx 7.9\text{--}9.7$ with the coronal plateau below and LTE saturation above). It is stored as `cool_logL_FeII_ne(NCOOLT, NCOOLNE)` = $\log_{10} \Lambda_{\text{eff}}$.

Two built-in sanity checks (printed to `stderr` by `--table`):

- **Low- n_e edge reproduces the coronal table.** At $n_e = 1 \text{ cm}^{-3}$ the SE coefficient equals `cooling_effective(pop='coronal')` to **0.1%** across $T = 10^3\text{--}10^5$ K — so swapping the 1-D coronal table for the $n_e = 1$ column of the 2-D table is a no-op, and the Phase-2 gate's coronal Fe II validation (vs Huang Fig. 6, within $\sim 1.5\times$) carries over unchanged.
- **Base suppression.** At the WASP-121b base ($T \approx 2239$ K, $n_e = 10^8$) the coronal coefficient 2.80×10^{-20} drops to the SE value 1.79×10^{-24} — a **15,663 $\times$ suppression**. The suppression is strongest at low T (forbidden lines) and weak at high T ($T \gtrsim 3 \times 10^4$ K, where the resonance lines with large A dominate and stay near-coronal): e.g. only $\sim 1.4\times$ at $T = 3 \times 10^4$, $n_e = 10^8$.

10.4 Step 2 — the 2-D table and bilinear interpolator in `Cool_coeff.f90` (done)

The generated table and its interpolators are now in `src/modules/radiation/Cool_coeff.f90`:

- **Table.** `integer, parameter :: NCOOLNE = 29`, the `cool_logne(NCOOLNE)` axis ($\log_{10} n_e = 0\text{--}14$, $d\log = 0.5$), and the data block `cool_logL_FeII_ne(NCOOLT, NCOOLNE)` (a `reshape([...], [NCOOLT, NCOOLNE])` with T on the inner index and n_e on the outer), pasted straight from `fe2_cooling.py --table`. The $n_e = 1$ column matches the existing 1-D `cool_logL_FeII` to the fifth decimal (first node -19.78002 vs -19.77975).
- **Bilinear interpolator.** `interp_cool_table_2d(logL2d, T, ne, out)` — log-log bilinear in $(\log_{10} T, \log_{10} n_e)$. The T axis reuses the uniform `cool_dlogT` scheme; the n_e axis uses the uniform `dlogne = 0.5`. Both axes clamp to the nearest edge outside the table (so n_e below the grid gives the coronal plateau, above it the LTE

branch). A scalar twin `interp_cool_table_2d_scalar(logL2d, Ts, nes)` mirrors it bit-for-bit for the cell-by-cell post-process solve.

- **Coolant wrapper.** `cool_FeII_ne(T, ne, out)` calls the 2-D interpolator on `cool_logL_FeII_ne`. The 1-D `cool_FeII` is kept as the low- n_e edge / scalar fallback; its header `!To Be Checked` note was rewritten to record that the density-dependent table now supersedes the coronal-only rate in the cooling assembly.

Builds clean (gfortran, full make); the new routines are additive and not yet wired, so behavior is unchanged until Step 3.

10.5 Step 3 — wiring the density-dependent Fe II into the cooling assembly (done)

The 2-D coefficient is now wired into both cooling paths so the local n_e selects the coronal $\rightarrow$ LTE-saturated branch automatically:

- **Main solver (`util_ion_eq.f90`, `eval_cool`).** n_e is already computed at the top of `eval_cool` (call `calc_ne`). Immediately after the canonical `c_metal(:, i)` fill loop, a short loop locates Fe II by name and overrides its column with `call cool_FeII_ne(T_K, ne, ...)`. Nothing else changes: the existing assembly $\text{cool_M} = \beta_{\text{esc}} n_e \sum_i n_{m,i} c_{\text{metal},i}$ then multiplies by n_e , which *cancels* the $1/n_e$ in $\Lambda_{\text{eff}} = (\sum_u n_u A_{ul} \Delta E_{ul}) / n_e$, leaving the correct cooling for each Fe II ion. The saturated (LTE) populations are collider-independent, so the electron-only SE solve is exact in that limit; at low n_e , Λ_{eff} reduces to the coronal rate, so optically-thin upper-atmosphere cells are unaffected. The override flows automatically into the `use_2lev_cool` branch (which already reads `c_metal(:, 26)`) and into the channel-by-channel `cool_chan` diagnostic (which reads `c_metal(:, i)`), so the breakdown stays an exact decomposition of `cool`.
- **Post-process cell-by-cell solve (`T_equation.f90`).** This path computes its own local n_e (from H, He, and the frozen metal stages) and sums the metal coolants with the scalar dispatcher. The Fe II term is special-cased to call the new `cool_FeII_ne_scalar(TT, ne)` (a thin wrapper over `interp_cool_table_2d_scalar`), so the post-process temperature balances the same density-dependent Fe II cooling.
- **species_table.f90.** The `!To Be Checked` Fe II comment was rewritten to state that `eval_cool` now applies the density-dependent SE coefficient (no longer “deferred”).

Builds clean (gfortran, full make, all four touched modules relink).

10.6 Step 4 — WASP-121b run and the cooling budget (numbers for the reader)

A fresh WASP-121b run (solar metals, `pp_metals 1`, otherwise the same `input.inp` as the FeI-cooling run) was compared against the snapshot of the previous (coronal Fe II) run kept in `output_coronalFeII/`. Both share the same grid, so `Cooling_breakdown.txt` columns are differenced directly (`analyze_FeII_density_correction.py` $\rightarrow$ `FeII_density_correction_compare.png`).

At the dense base ($r = 1.0$, $T = 2358\text{ K}$ fixed by the lower BC). The Fe II line cooling drops from the coronal 2.23×10^{-4} to the SE $1.71 \times 10^{-8} \text{ erg cm}^{-3} \text{ s}^{-1}$ — a $\approx 1.3 \times 10^4 \times$ suppression, consistent with the offline fixed- (T, n_e) check in Step 1 ($15,663 \times$ at $T = 2239\text{ K}$, $n_e = 10^8$). Because coronal Fe II carried **99.1%** of the base cooling, the *total* radiative cooling at the base falls $\approx 122 \times$ ($2.25 \times 10^{-4} \rightarrow 1.84 \times 10^{-6}$); Fe II’s share there is now **0.9%**. The previous run’s post-process printed “35 of 502 cells fell back to eq T (stiff base band)”; with the spurious base overcooling removed that message is **gone** — the cell-by-cell T solve no longer stiffens at the base.

Transition region 1.15–1.4 R_p (Huang Fig. 10). The band-integrated metal-line cooling share in the new run is

coolant	share
Mg II	44.0%
O II	18.9%
Fe II	13.4%
O I	7.3%
C II	6.3%
Ca II	5.1%

so **Mg II is now unambiguously the dominant metal coolant** across the transition region, with Fe II demoted to a $\sim 13\%$ contributor — the qualitative ordering Huang’s Fig. 10 shows. (In the old coronal run Fe II still topped the budget out to $\sim 1.2 R_p$.)

Global response. Removing the base overcooling raises the heating efficiency, so the steady-state mass-loss rate rises from $\log_{10} \dot{M} = 12.51$ to **12.80 g/s** ($\approx 2 \times$). The self-consistent wind is correspondingly denser and, above the base, slightly *cooler* (e.g. $T(1.05)$: $7810 \rightarrow 5744\text{ K}$) — the freed energy goes into mechanical escape and adiabatic expansion rather than being radiated, the expected energy-limited behavior. Whether this $\dot{M}$ and the transition-region budget match Huang’s WASP-121b quantitatively is left to the reader’s judgment against the paper; the figure and table above are surfaced for that comparison, not asserted as a pass.

Caveat on the figure’s panel (b). The “suppression factor” there is the *ratio of Fe II cooling between the two runs*, not a fixed-condition coefficient ratio. At the base (BC-fixed T) it equals the coefficient suppression ($\sim 1.3 \times 10^4$); above $\sim 1.2 R_p$ the coronal and SE coefficients converge (the hot resonance lines dominate and stay near-coronal), so the residual ratio there reflects the changed — denser — wind profile, not the coefficient.

11. Post-process refinement: re-solve metal ionization in the advected wind (pp_metals 2)

Added 2026-06-05.

The `_adv` profiles that `EXHALE_transit.py` reads are built by `post_process_adv.f90`, which re-solves the *temperature in each cell* at the advection-corrected structure (the wind is advected, then its $T(r)$ is re-converged against heat-

ing/cooling). Until now the metal *ion* densities carried into those profiles were handled in one of two ways, selected at runtime by `pp_metals` in `metals.inp` ($\rightarrow$ `pp_metal_mode`):

- `pp_metals 0` — **metal-free legacy** `_adv` (the pre-metals behavior).
- `pp_metals 1` — **frozen equilibrium** metals (the default): the converged equilibrium metal stages are carried into `_adv` unchanged, so they track the *equilibrium* structure, not the (slightly hotter, more ionized) advected wind.

This step implements the deferred third option:

- `pp_metals 2` — **re-solve**: the metal ionization balance is recomputed in each cell at the advection-corrected H/He and the post-process temperature, so the metal stage split tracks the advected wind self-consistently.

11.1 Implementation — reuse the coupled residual with H/He pinned

The metals are pure equilibrium species (no advection equations of their own): in the main loop they are solved *coupled* with H/He through `ion_system_HeH_metals` + `hybrd1` (photo + collisional ionization + recombination + charge exchange, no time derivative). “Re-solve” therefore means re-evaluating that same equilibrium balance at the post-process local T and the advected H/He — *not* a new physics path.

Rather than duplicate the residual, mode 2 reuses it verbatim through a thin wrapper, `ion_system_metals_pp` (new, in `post_process_adv.f90`):

```
subroutine ion_system_metals_pp(N_in,x,fvec,iflag,params)
...
call ion_system_HeH_metals(N_in,x,fvec,iflag,params) ! full coupled residual
fvec(1) = x(1) - pp_xHII_fix      ! pin HII fraction
fvec(2) = x(2) - pp_xHeII_fix     ! pin HeII fraction
fvec(3) = x(3) - pp_xHeIII_fix    ! pin HeIII fraction
end subroutine
```

The wrapper overwrites only the three H/He rows with identity equations that pin H/He to their advection-corrected fractions (stored in module-level `pp_xHII_fix` / `pp_xHeII_fix` / `pp_xHeIII_fix`), while leaving the $2n_{\text{melem}} = 20$ metal-stage rows untouched. So `hybrd1` moves only the metal unknowns; H/He stay fixed at the advected wind. Two properties fall out for free:

- **Self-consistent n_e .** The reused residual computes the electron density *internally* from the pinned H/He plus the metals it is currently solving — exactly as the equilibrium solver does — so the re-solve sees the same n_e (metals) feedback, not a frozen n_e .
- **Charge exchange with the advected H.** The Kingdon & Ferland $\text{H} \leftrightarrow \text{metal}$ couplings inside `ion_system_HeH_metals` see the *pinned* H^+/He^+ densities, so the metal balance reacts to the advected hydrogen ionization, not the equilibrium one.

11.2 The cell-by-cell sweep in post_process_adv

Mode 2 adds one sweep over cells, placed after the post-advection calc_ne and before the photoheating refresh, inside the same k -loop that re-converges T , so metals, H/He and T are driven to a *joint* fixed point. Per cell j (skipping $n_h(j) \leq 0$):

1. `cx_set_cell(T_K(j))` — charge-exchange rate coefficients at the local T .
2. Build the coefficients for each element (canonical order) from the existing post-process rate arrays: `meg_g0 = P_m` (photoionization), `meg_b0 = aion_m_pp` (collisional ionization), `meg_a1 = rec_m_pp` (recombination), and the second-stage $g1/b1/a2$ where `melem_top  $\geq$  2`; hand them to `set_metal_coeffs`.
3. Pin `pp_xHII_fix = nHII/nh`, `pp_xHeII_fix = nHeII/nhe`, `pp_xHeIII_fix = nHeIII/nhe`; set `params(7) = nh`, `params(8) = nhe`.
4. Seed `sys_x` with the pinned H/He fractions plus the current metal split and call `hybrd1(ion_system_metals_pp, . . .)`.
5. **Conserve element totals:** the pre-solve totals `nm_tot_pp(:,im)` (summed over each element's stages once, before the k -loop) are multiplied by the re-solved fractions to write back `nm_w`. Only the *stage split* changes; the element abundance is untouched.

Supporting edits: the use block pulls in `ion_system_HeH_metals`, `set_metal_coeffs`, and `cx_set_cell`; `params` was widened from `dimension(25)` to `dimension(60)` to match the residual's dummy length (only entries 1–11 are consumed by the metal solve); module-level `pp_x*_fix` and locals `nm_tot_pp / meg_* / i0`, `top`, `im` were added. Builds clean (gfortran, full make).

11.3 Clean A/B and what changes

Because the **equilibrium** solve in the main RK loop is mode-independent and deterministic, a full `pp_metals 2` run reproduces `Hydro_ioniz.txt`, `Ion_species.txt`, and `Cooling_breakdown.txt` **byte-identically** to the `pp_metals 1` run (verified with `diff -q`); *only* the `*_adv.txt` profiles differ. So the A/B isolates exactly the post-process metal treatment. The mode-1 default is preserved in `output_ppm1_dd/`, the mode-2 run in `output_ppm2_resolve/`, and the canonical `output/` is left at the mode-1 default (`metals.inp` reset to `pp_metals 1`); `analyze_metal_resolve.py` differences the two and writes `metal_resolve_compare.png`.

Singly-ionized fraction X^+/X_{tot} in the _adv wind (eq frozen $\rightarrow$ re-solve). Only the elements that move appreciably are tabulated (Mg and Na are unchanged to $\leq 2\%$); the full output for each element is in the analysis script.

r/R_p	T_{ad3} [K]	O	N	Ca	Fe
1.05	5798	.00 $\rightarrow$.00	.01 $\rightarrow$.02	.64 $\rightarrow$.59	1.00 $\rightarrow$ 1.00
1.10	7984	.02 $\rightarrow$.01	.12 $\rightarrow$.13	.65 $\rightarrow$.60	.99 $\rightarrow$.99
1.15	9044	.10 $\rightarrow$.06	.38 $\rightarrow$.40	.66 $\rightarrow$.59	.99 $\rightarrow$.99
1.20	10035	.37 $\rightarrow$.26	.51 $\rightarrow$.47	.61 $\rightarrow$.54	.93 $\rightarrow$.96
1.30	10274	.34 $\rightarrow$.34	.35 $\rightarrow$.34	.38 $\rightarrow$.38	.39 $\rightarrow$.40
1.40	10680	.23 $\rightarrow$.23	.23 $\rightarrow$.23	.25 $\rightarrow$.25	.18 $\rightarrow$.19

Largest density change. The element-by-element scan of $\max |\Delta n(X^+)/n(X^+)|$ over the wind 1.0–1.48 R_p :

element	max rel. change	at r/R_p	$n(X^+)$ frozen $\rightarrow$ resolve
Fe	+9.5%	1.232	$1.19 \times 10^5 \rightarrow 1.30 \times 10^5$
Ca	−11.9%	1.191	$1.24 \times 10^4 \rightarrow 1.09 \times 10^4$
Na	−6.2%	1.000	$3.50 \times 10^6 \rightarrow 3.28 \times 10^6$
C	−2.1%	1.216	$1.22 \times 10^6 \rightarrow 1.19 \times 10^6$
Mg	−2.0%	1.000	$9.82 \times 10^7 \rightarrow 9.62 \times 10^7$

The corrections are **modest and concentrated in the 1.1–1.25 R_p ionization transition**, where T_{ad3} runs a few hundred K hotter than T_{eq} : **Fe II** is *higher* across the 1.18–1.27 R_p shoulder (blue above red in the figure), **Ca II** is *lower* by ~ 5 –7 points through 1.05–1.25 R_p , and **every element converges back to the equilibrium split** in the outer wind ($\gtrsim 1.27 R_p$) where T_{ad3} rejoins T_{eq} . The post-process printed **no** stiff-base T -fallback message in mode 2 (the density-dependent Fe II cooling of the previous section already removed the base overcooling).

Tiny-denominator artifacts (not physical). The same scan reports **O** +6538% at $r = 1.000$ and **N** +17.9% at $r = 1.013$. These are at the dense base where O and N are essentially *fully neutral* (O II fraction ≈ 0.00 in the table), so $n(\text{O II})$ is a vanishingly small number; a tiny absolute shift ($n^+ \approx 9.7 \times 10^2 \rightarrow 6.4 \times 10^4 \text{ cm}^{-3}$, against an O total $\sim 10^8$) reads as a huge *relative* change but is negligible in absolute and in any observable. They are denominator artifacts, not a physical re-solve effect.

Reader’s judgment, not a pass. The re-solve makes the metal stage split in the `_adv` profiles consistent with the advected wind rather than the equilibrium structure, with the physically expected sign (more ionization where the advected wind is hotter) and magnitude (single-digit-percent, transition-localized). The default remains `pp_metals 1` (frozen): the deliberate, documented choice (the equilibrium split is the converged one, and EXHALE_transit.py’s lines are dominated by H and He). Mode 2 is provided for studies that need the metal ionization to track the advected wind. The table and `metal_resolve_compare.png` are surfaced for that comparison, not asserted as a validation pass.

12. Phase 3b: in-code $\text{Ly}\alpha$ radiative transfer ($\bar{J}_{\text{Ly}\alpha}$ via escape probability)

Phase 3a drove the $\text{H}(n=2)$ population with a *parameterized* $\text{Ly}\alpha$ mean intensity $\bar{J}_{\text{Ly}\alpha} = 0.1 F_{\text{LyC}} / \Delta\nu_D / (1 + \tau_{\text{Ly}\alpha})$ (mode 0). That estimate peaks the $n=2$ population near $2R_p$ and *under-pumps* it across 1.1–1.8 R_p , leaving a `jlya_mode=1` hook (`load_jlya_rt`) to ingest a real RT field. Phase 3b supplies the real $\bar{J}_{\text{Ly}\alpha}(r)$. Three findings reshaped the original plan (a Monte Carlo with an offline two-pass and a β -on-cooling term):

1. **Method $\rightarrow$ escape probability, not Monte Carlo.** The WASP-121b column is extremely thick: $\tau_0(\text{Ly}\alpha) \approx 5.5 \times 10^7$ at the base (Voigt $a \approx 5 \times 10^{-4}$, so $a\tau_0 \approx 10^4 \gg 1$). A faithful acceleration-free MC needs $\sim \tau_0 \approx 10^7$ – 10^8 scatterings/photon (intractable), and core-skipping is both forbidden and incompatible with the core-dominated pumping field. For $a\tau_0 \gg 1$ the Neufeld (1990)/Harrington analytic wing-escape solution is accurate, so we use a local escape-probability source function.

2. **In-code, not offline.** Being a cheap formula, $\bar{J}_{\text{Ly}\alpha}$ is evaluated every timestep inside the run (`jlya_mode=2`, the same lagged-explicit decoupling as the Phase-3a Balmer feedback); a single run is self-consistent. The `jlya_mode=1` option to read an external (e.g. real MC) profile is kept.
3. **No β -on-cooling term.** From Huang §2.5/Fig. 7 directly: the “Ly α cooling 10 $\times$ lower” is the Black (1981) vs CHIANTI *collisional-excitation rate-coefficient* difference, not photon trapping. EXHALE already uses the CHIANTI-consistent Cen `coex_rate_HI` (§7), so there is no 10 $\times$ to apply. The trapping feedback is instead the Phase-3a H($n=2$) heating (photoelectric Hpe + collisional de-excitation Hdx), which strengthens once $\bar{J}_{\text{Ly}\alpha}$ is correct. The scaffolding β -on-coex hook was reverted.

12.1 The escape-probability field (`lya_rt.f90` : `jlya_escape_prob`)

Per cell, from the converged T , $n_{\text{HI}} = n_{1s}$, n_{HII} , n_e :

- $\Delta\nu_D = \nu_{\text{Ly}\alpha} \sqrt{2kT/m_{\text{H}}}/c$; Voigt $a = (A_{2p1s}/4\pi)/\Delta\nu_D$; optical depth $\tau(r) = \int_r^\infty n_{1s}\sigma_0 dr'$, $\sigma_0 = C_{\text{Ly}\alpha}/\Delta\nu_D$.
- Wing escape probability $\beta = \pi^{-1/4} \sqrt{a/\tau}$ (≤ 1): escape once the photon diffuses to $x_* = (a\tau/\sqrt{\pi})^{1/2}$ where the wing optical depth is unity.
- Internal sources $P = \alpha_B n_e n_{\text{HII}} + q_{1s2p} n_e n_{\text{HI}}$ (recombination cascade + electron-impact $1s \rightarrow 2p$). The escape-probability closure $\bar{J} = S(1 - \beta)$ with the trapped pile-up $n_{2p} = P/(A_{2p1s}\beta)$ gives $J_{\text{int}} = (2h\nu^3/c^2)(g_{1s}/g_{2p})P(1 - \beta)/(A_{2p1s}\beta n_{1s})$. The $(1 - \beta)$ factor makes $J_{\text{int}} \rightarrow S$ when thick and $\rightarrow 0$ when thin ($\beta \rightarrow 1$), instead of $S \propto 1/n_{1s}$ diverging as the gas ionizes in the outer wind (which spuriously raised $\bar{J}$ outward).
- Stellar beam (trapped): Ly α is *scattered*, not destroyed, so it is not killed by $e^{-\tau}$; the broad stellar line (half-width $\Delta v_s \approx \pm 70 \text{ km s}^{-1}$) penetrates through its wings. A photon at offset x reaches r if $|x| > x_1 = \sqrt{a\tau/\sqrt{\pi}}$; for a Gaussian stellar profile of width $X_s = \Delta v_s/v_{\text{th}}$, $T_*(r) = \text{erfc}(x_1/(\sqrt{2}X_s))$. The profile-averaged mean intensity uses the *stellar* line width $\Delta\nu_* = \nu_0 \Delta v_s/c$ (constant), NOT the planetary $\Delta\nu_D$ ($\bar{J} = \int J_\nu \phi d\nu \approx J_\nu(\nu_0)$ for a stellar field flat across the narrow planetary line; dividing by $\Delta\nu_D$ makes J_{star} rise as the wind cools). A bounded buildup $E = 1 + (\text{boost} - 1)(1 - \beta)T_*$ enhances the penetrating beam in the thick wind and $\rightarrow 1$ when thin or at the deep edge, avoiding both the $1/\beta$ over-count (which blew the $n=2$ source to 6300 $\times$ ground-state) and a flat cap. Hence $J_{\text{star}} = \xi F_* T_* E / (4\pi \Delta\nu_*)$, ξ the dayside dilution. $\bar{J}_{\text{Ly}\alpha} = J_{\text{int}} + J_{\text{star}}$.

New `input.inp` lines (default-off $\Rightarrow$ byte-identical to 3a; a 3b-0 scaffolding step verified this, frozen reference `output_pre3b_baseline/`):

```
Jlya escape-prob: True           # jlya_mode = 2 (in-line escape-prob RT)
Stellar Lya flux [erg/cm2/s]: 1.0e5 # incident stellar Lya at the planet (x0.35)
Lya stellar halfwidth [km/s]: 70   # broad stellar line Dv_s (penetration + Dnu_star)
Lya stellar boost [-]: 5           # bounded trapping-buildup cap (left un-tuned)
Deexc heat: True                  # Phase-3a collisional de-excitation heating
```

12.2 WASP-121b validation setup (spherical Case A, matched to Huang)

Huang’s Fig. 11 is **Case A** (spherical, no RLOF), so the comparison must match his geometry *and* radial frame — which exposed a setup error: the run had been Roche-truncated at $\sim 2R_p$, and EXHALE’s base ($r=1$) sits at **1 μ bar** ($n_0 = 3.1 \times 10^{12}$, $T = 2358$ K), but Huang defines R_p at the **4 mbar transit radius** and places his 1 μ bar hydro bottom at $\sim 1.25R_p$ (Fig. 9 top-pressure axis / Fig. 12 green band; the earlier “1.46 R_p ” note was wrong). The setup was corrected to spherical Case A with R_0 set to Huang’s 1 μ bar radius ($R_0 = 1.25 \times 1.766 = 2.2075 R_J$, planet mass unchanged $\Rightarrow$ correct base gravity) and outer radius 15.2 ($= 19R_p^{\text{Huang}}$); EXHALE r maps as $R_{\text{Huang}} = 1.25r$. The Case-A $\dot{M} = 0.077 M_p/\text{Gyr}$ is within $\sim 1.5\times$ of Huang’s 0.052 (Table 3).

12.3 Trend vs Huang Fig. 11 (qualitative — exact match not expected)

Huang’s $\bar{J}_{\text{Ly}\alpha}$ is a **plane-parallel Monte Carlo** RT; ours is an analytic escape-probability formula, so the two cannot agree quantitatively — a quantitative comparison is only sensible once *all* Phase-3b features (incl. the velocity term) are in. Qualitatively (lya_caseA_huang11.py), both show a significant inner peak (EXHALE ~ 0.6 at 1.35, Huang ~ 0.5 at 1.6 R_p) and comparable magnitude (0.05–0.6). EXHALE has a dip at $\sim 1.8R_p$ and rises gently outward where Huang gently declines; the outward rise traces to the $n2s$ recombination cascade in the Phase-3a H($2s/2p$) model ($n=2$ stays recombination-populated while ground-state H vanishes in the ionized wind), *not* $\bar{J}$, which after the two fixes below is physically sound (J_{int} peaks then declines; $J_{\text{star}} \approx \text{const}$ outward). H($n=2$) heating stays $\sim 1\%$ of total with $\text{Hdx} \gtrsim \text{Hpe}$ near the n_2 peak (cf. Fig. 10).

Two $\bar{J}$ bugs fixed during the trend check. (1) J_{star} used the planetary $\Delta\nu_D$ instead of the stellar width $\Delta\nu_*$, making $\bar{J}$ rise as the wind cooled; (2) J_{int} lacked the $(1-\beta)$ closure, so S diverged ($\propto 1/n_{1s}$) in the ionized wind.

Wind-velocity (Sobolev) escape — implemented, negligible here. $\beta_{\text{tot}} = 1 - (1 - \beta_{\text{static}})(1 - \beta_{\text{Sob}})$ with $\tau_S = C_{\text{sob}} n_{1s} / |dv/dr|$ is now in jlya_escape_prob (the outflow $v(r)$ threaded through excited_H_update). For WASP-121b it changes $\bar{J}$ by $< 1\%$: where it could help (thick inner wind) $\tau_S \sim \tau_{\text{static}} \sim 10^7$ so $\beta_{\text{Sob}} \sim 10^{-7}$, and the thin outer wind already has $\beta_{\text{static}} \approx 1$ — the velocity gradient and the Ly α opacity scale together ($\propto n_{1s}$), opening no new escape route. **Phase 3b is thus feature-complete.** *Deferred per the trend-first guidance:* the post-all-features quantitative comparison and the $n2s$ outward-rise revisit (a Phase-3a 2s effect, not the RT); boost, Δv_s , and the $\pi^{-1/4}$ prefactor are left un-tuned.

13. Phase 4: RLOF / tidal potential (Case B validated; Case D blocker = escape radius, transonic IC added, geometry-limited)

Phase 4 adds the Roche-lobe-overflow (RLOF) tidal physics that drives Huang’s enhanced mass loss — **Case B** $\dot{M} = 0.32$, **Case D** $= 1.03 M_p/\text{Gyr}$ — the regime that matches the observed transit depths (spherical Case A underestimates them). EXHALE already has the **1D Roche potential** (grav_field.f90, the spherical_domain=.false.

branch: planet + stellar-tidal + centrifugal), the L1-truncated domain, and the tidal momentum source (Source.f90) — Caldiroli’s ATES-v2 Roche mode. So the substellar tidal hydrodynamics is already in place.

13.1 Base gravity for each case (Huang Table 3)

The key to each case is the $1\ \mu\text{bar}$ base radius R_0 . Huang’s Table 3 lists $\log g @ 1\ \mu\text{bar} = GM_p/r_{1\mu\text{bar}}^2$ per case — exactly the R_0 -setting quantity ($R_0 = \sqrt{GM_p/g}$):

Case	tidal	$M_p(M_J)$	dT	XUV/Ly α /Fe	$\log g$	$R_0(R_J)$
A	N	1.1824	0	1/1/1	2.84	2.058
B	Y	1.1824	0	1/1/1	2.70	2.418
D	Y	1.1204	350	0.74/0.35/4	2.62	2.581

R_0 grows A $\rightarrow$ B $\rightarrow$ D as the tidal force + the dT+350 lower-atmosphere heating puff the $1\ \mu\text{bar}$ level outward (weaker base gravity), driving the RLOF.

13.2 Case B — validated

WASP-121b in Roche mode (tidal on) + the corrected base + metals + Ly α mode-2 gives $\dot{M} = 0.337 M_p/\text{Gyr} \approx$ Huang’s Case B (0.32), within $\sim 5\%$. With the base-radius fix both cases agree (Case A $0.077 \approx 1.5 \times$ Huang’s 0.052; Case B $0.337 \approx 1.05 \times$ Huang’s 0.32). EXHALE’s existing Roche potential thus captures the RLOF enhancement — no geometric-reconstruction code was needed for the Case B $\dot{M}$ gate.

13.3 Case D — real blocker found (escape radius), transonic IC added

Case D’s base ($R_0 = 2.581 R_J$, $\log g = 2.62$) sits at $\sim 76\%$ of the L1 distance: the L1-truncated Roche domain reaches only $r_{\text{max}} = 1.353 R_p$. Re-investigating the earlier “won’t launch” report uncovered the *actual* blocker, and it is *not* the initial condition.

Root cause — the escape radius lies outside the L1-truncated domain. `define_grid` sets the constant-momentum convergence window $[j_{\text{min}} : N]$ from the **escape radius** (Escape radius `[R_p] :`, default 1.50). It scans for the first cell with $r(j) \geq r_{\text{esc}}$; if *no* cell qualifies — which happens precisely when $r_{\text{esc}} > r_{\text{max}}$ — the DO loop falls through leaving $j = N + N_g + 1$, so $j_{\text{min}} > N$ and the window is **empty**. Fortran `maxval/minval` over a zero-size array return $\mp \text{HUGE}$, so $\text{mom}_{\text{max}} = -\text{HUGE}$, $\text{mom}_{\text{min}} = +\text{HUGE} \Rightarrow du = \text{Inf}$, and the time-derivative norm $dtu = -\text{HUGE}$. Since $dtu = -\text{HUGE} < dtu_{\text{th}}$, the loop reports **“steady state reached” at step 0** — a *false* convergence unrelated to the wind. The earlier bracket (“works $R_0 \leq 2.30$, fails $R_0 \geq 2.40$ ”) is exactly r_{max} crossing the 1.50 escape radius at $R_0 \approx 2.37$, confirming the escape radius, not the IC, as the discriminant. (The earlier “isothermal IC cannot launch the deep-RLOF wind” diagnosis was therefore wrong.)

Fixes. (1) A guard in `define_grid.f90`: if $j_{\text{min}} > N$, re-anchor the window to mid-domain ($r \geq 1 + \frac{1}{2}(r_{\text{max}} - 1)$) and print a loud WARNING — so this can never again masquerade as convergence. (2) For Case D, set Escape

radius [R_p]: 1.20 (inside the 1.353 domain). With a valid window the diagnostics become finite and meaningful, and **Case D launches and runs** — from *both* the hydrostatic and the transonic IC. (The earlier $du = |(\text{mom}_{\text{max}} - \text{mom}_{\text{min}}) / \max(\text{mom}_{\text{min}}, 10^{-30})|$ floor at EXHALE_main.f90:308 is retained as defense-in-depth against a genuine transient zero-momentum cell, but it was *not* what unblocked Case D — with an empty window $\text{mom}_{\text{min}} = +\text{HUGE}$, not zero, so the floor never engaged.)

Transonic-wind IC (implemented, Transonic IC: True). set_IC.f90 gained an optional transonic isothermal-wind profile, solved from the EXHALE potential via the algebraic (Bernoulli) integral of the steady isothermal-wind equation,

$$\frac{1}{2}v^2 - c^2 \ln v - 2c^2 \ln r + \phi(r) = B, \quad c^2 = (1 + \delta p_{bc}) / \rho_{bc},$$

with B fixed by the sonic point ($v = c_s$ where $D\phi(r_c) = 2c^2/r_c$, located by bisection). At each radius the relation has a subsonic ($v < c_s$) and a supersonic ($v > c_s$) root; the wind takes the subsonic root for $r < r_c$ and the supersonic root for $r > r_c$ (1-D bisection on each monotone branch), and ρ then follows from mass conservation $\rho v r^2 = \dot{M}$ pinned to the base BC. Because that IC already satisfies $\rho v r^2 = \text{const}$, $du \approx 0$ at step 0 would itself trip the momentum-constant exit before the cold wind heats; the flag therefore **auto-enables force_start**. If no interior sonic point exists it **falls back to the hydrostatic IC**. Default **off** $\Rightarrow$ existing cases byte-identical (Case B re-verified at $\dot{M} = 0.337$).

Case D outcome — launches, but the 1-D L1-truncated domain undershoots. With the escape-radius fix the wind launches and relaxes, but it **never reaches a sonic point inside the domain**: the heated metal wind is **deeply subsonic everywhere** (Mach $\lesssim 0.05$ — rising to ~ 0.050 at $r \approx 1.25$, then *falling* to ~ 0.040 at L1). With a subsonic outer boundary the zero-gradient outflow BC is only marginally well-posed and $\rho v r^2$ stays non-constant (peaks $\sim 1.7\times$ the L1 value), so the run settles into a **breathing limit cycle**: after the $\sim 10^2$ heating transient damps, du keeps oscillating in $\sim 0.04\text{--}0.8$ (touching but never crossing the strict $du_{\text{th}} = 0.02$) for as long as it runs. The mass flux is correspondingly time-variable — the quasi-steady $\dot{M}$ oscillates over $\sim 0.3\text{--}0.6 M_p/\text{Gyr}$, well *below* Huang's 1.03, the deficit set by the L1 truncation cutting off the still-accelerating subsonic wind before it reaches the sonic point. The transonic IC behaves identically to the hydrostatic IC here: both launch the same subsonic flow. (Warm-starting via load_IC reintroduces a base NaN — the documented grid re-read inconsistency — so it is not a workaround.)

Conclusion. Case D's deep RLOF ($\dot{M} = 1.03$) is **not reachable in EXHALE's 1-D L1-truncated Roche mode**: the wind is subsonic at L1, so the truncated domain structurally undershoots. Capturing it needs the 3-D geometric treatment (Phase 5) or a domain extended past L1 — *not* a different IC. The transonic IC and the escape-radius guard are correct, necessary infrastructure (they remove the false convergence and let the deep-RLOF wind launch and run with meaningful diagnostics) but are not sufficient on their own for Case D's full mass-loss rate. The geometric reconstruction (substellar $\rightarrow$ terminator mapping, triaxial radii Table 2, the 4/9 mass-loss conversion) feeds Phase 5 (§14), and also supplies the geometry this subsonic L1-truncated wind needs.

14. Phase 5: velocity-broadened transmission spectrum (in progress)

Phase 5 is the headline deliverable — the NUV/optical **transit-depth spectrum** the whole upgrade exists to match (Huang Table 3 / Figs. 17–24). It extends the existing spherical post-processor EXHALE_transit.py in two stages: **5a** adds the metal resonance lines (spherical, gated on Case A); **5b** replaces the spherical geometry with the 3-D Roche equipotential reconstruction + velocity broadening (gated on Case D). Pure Python post-processing — no Fortran change.

Line	Case A	Case D	Observed
Mg II $\lambda 2796$ (4 Å)	0.182	0.302	0.309 ± 0.036
Ca II K $\lambda 3934$	0.199	0.278	0.281 ± 0.009
H α $\lambda 6563$	0.201	0.185	0.186 ± 0.003
H β $\lambda 4861$	0.174	0.135	0.143 ± 0.005
Na D2 $\lambda 5890$	0.152	0.147	0.147 ± 0.002

These are the **effective transit radius** R_{eff}/R_* , not an absorption percent (the expected-results doc states the unit is R_p/R_*). Our code returns the line-center absorption fraction h , mapped to the same quantity via $R_{\text{eff}}/R_* = \sqrt{(R_p/R_*)^2 + h}$.

14.1 5a — metal resonance lines (Mg II, Ca II, Na I D), spherical — done

EXHALE_transit.py gained Mg II $\lambda 2796$, Ca II K $\lambda 3934$, Na I D2 $\lambda 5890$ as an **isolated appended block** (the validated He 10830 / Ly α / H α / H β pipeline above it is untouched). These are resonance lines whose lower level is the ion ground state; at $\sim 10^4$ K the excited fine-structure levels are Boltzmann-negligible, so the lower-level density is the ion density itself ($n_{\text{lower}} \approx n_{\text{ion}}$). The ion densities live past the H/He block of Ion_species_adv.txt, in mion_fsp order (species_table.f90): numpy columns **17 = Mg II**, **23 = Ca II**, **25 = Na I**. The block reuses the existing spherical chord / Voigt (wofz) / trapezoid- τ / disk-average / instrument convolution path (resonance_depth()).

Case A result (line-center % / 4 Å-band %) and diagnosis. Mg II: 26.1 / 2.3 ($R_{\text{eff}} = 4.2 R_p$); Ca II K: 5.5 / 0.33 ($1.9 R_p$); H α : 3.9 (1.6 R_p); Na D2: 0.8 / 0.06 (0.7 R_p). Two real gaps: (1) **Mg II already reaches Huang's extent** ($R_{\text{eff}} = 4.2 R_p$ vs his clustered $\sim 3.6 R_p$) — the metal treatment is sound — but **Ca II / Na I / H α are too optically thin** ($R_{\text{eff}} = 0.7\text{--}1.9 R_p$), a line-specific ionization / abundance / $n = 2$ question, not geometry. (2) The lines are **too narrow**: the Case A wind reaches only ~ 29 km/s, so a 26%-deep core washes out to $\sim 2\%$ over a 4 Å (± 215 km/s) bin, whereas Huang's lines stay deep *across* the bin — they are velocity-broadened to hundreds of km/s by the fast Roche outflow + rotation, i.e. **Phase-5b physics**. (The existing H α tool shows the same shortfall, so it is systematic.) *Units — resolved*: Huang's table is R_{eff}/R_* (radius ratio; the expected-results doc header states the unit is R_p/R_*), not absorption percent. The Case A spherical run over-extends (15 R_p domain $\Rightarrow$ Mg II reaches $4.2 R_p$, $R_{\text{eff}}/R_* = 0.51$ vs Huang 0.182); the proper test is the L1-confined Roche run in R_{eff}/R_* units (§5b).

14.2 5b — Roche-equipotential reconstruction + broadening — implemented

New module `roche_recon.py`. Assuming a uniform state on Roche-equipotential surfaces, the field at any (x, y, z) equals the 1-D substellar state at the equivalent radius r_{eff} where $\phi_{\text{sub}}(r_{\text{eff}}) = \phi_{3\text{D}}(x, y, z)$. The dimensionless Roche potential (R_p units) mirrors `grav_field.f90`:

$$\phi(x, y, z) = -\frac{1}{r} - \frac{q}{r_s} - \frac{1+q}{2a^3} [(x-x_{\text{cm}})^2 + y^2], \quad q = M_*/M_p, \quad a = \text{orbit}[R_p].$$

Tidal terms $\rightarrow 0$ recovers $\phi = -1/r$ (spherical), so the triaxial reconstruction reduces to the 5a spherical case for Case A. The module gives `roche_phi`, the L1 root of $d\phi_{\text{sub}}/dr = 0$, a monotone $\phi_{\text{sub}} \rightarrow r_{\text{eff}}$ inverse (`ReconMap`), the triaxial radii, and the Eggleton lobe radius. `EXHALE_transit.py` gained a `geometry='triaxial'` switch (`triaxial_depth()`): the transit LOS runs along $+x$; the state at each 3-D point is the substellar state at r_{eff} ; and the **LOS velocity in each sector** $v_{\text{sub}}(r_{\text{eff}})x/r - \Omega y$ (wind + tidally-locked rotation, Huang Eq. 16) is integrated over **20 angular sectors**.

Demonstrated on the converged Case B (Roche) run. Geometry self-consistent: $L1 = 1.574 R_p$; $R_{\text{RL}}/L1 = 0.714 \approx 2/3$; tidally triaxial $R_{px} = 1.05 > R_{py} = 0.919 > R_{pz} = 0.893$ (substellar bulge, polar compression).

Result — the model reproduces Huang’s transit radii. Converting the Case B triaxial depths to R_{eff}/R_* (the table unit):

Line	Huang A	Huang D	Case B (this work)
Mg II	0.182	0.302	0.234
Ca II	0.199	0.278	0.222
Na D2	0.152	0.147	0.163
H α	0.201	0.185	0.239
H β	0.174	0.135	0.232

Case B is the *intermediate* model (between Case A and Case D), and the **metals (Mg II, Ca II, Na) land right between Huang’s Case A and Case D** — the lobe-confined reconstruction correctly reproduces his transit radii in magnitude and trend. The Balmer lines run $\sim 25\%$ high (H α 0.239 vs the expected ~ 0.19), consistent with the 5a finding that $n = 2/T$ is slightly over-predicted.

Correction. An earlier draft claimed the lobe-confined model “undershoots Huang because it omits the beyond-L1 escaping tail.” That was wrong twice over: (1) Huang’s reconstruction is *also* lobe-confined (equipotential mapping to the Roche-lobe boundary, §1.8) — neither model has a beyond-L1 tail, so it cannot be the difference; and (2) the apparent shortfall was a **units error** (our absorption *percent* vs Huang’s R_{eff}/R_*). In the correct units the two agree (table above). The small triaxial-vs-spherical reduction (terminator compression $R_{py} < R_{px}$) is minor; Huang’s area-preserving $R_{py}R_{pz} = R_p^2$ normalization would offset it.

14.3 5b — direct Case A comparison: hydro validated, metals over-extended

A clean spherical Case A at the *correct* 1 μbar base $R_0 = 2.058 R_J$ (Table 3 $\log g = 2.84$; the old `output_caseA` used $R_0 = 1.766$, the 4 mbar transit radius) was run to convergence.

- **Hydro / base validated:** $\dot{M} = 0.056 M_p / \text{Gyr} \approx \text{Huang's Case A } 0.052$ (the wrong base gave a different $\dot{M}$) — base and wind are right.
- **Metal transit radii — consistent once compared like-for-like** (Huang quotes Mg II in a $4 \text{ \AA}$ NUV bin; Ca II / Na are optical line-center): Mg II ($4 \text{ \AA}$) **0.220** vs Huang 0.182 ($\times 1.2$); Na D2 (LC) **0.175** vs 0.152 ($\times 1.15$); Ca II K (LC) **0.306** vs 0.199 ($\times 1.5$, high).

The “over-extension” was a comparison-definition artifact, not physics. An earlier draft reported Mg II $R_{\text{eff}}/R_* = 0.60$ ($3.3\times$ Huang) — that compared the *line-center* R_{eff} against Huang’s $4 \text{ \AA}$ NUV-bin value. In the same $4 \text{ \AA}$ bin Mg II $= 0.22 \approx \text{Huang's } 0.182$; Na matches too; only Ca II is genuinely $\sim 1.5\times$ high.

The metal ionization is correct and is NOT the lever. All channels are wired and physical (Verner+1996 σ , Badnell RR+DR, Voronov, Huang+2023 charge exchange; $a_\tau = 0$). The Mg II/III balance is a close competition between a weak photoionization ($\Gamma \approx 2 \times 10^{-3} \text{ s}^{-1}$, weak because $\sigma_{\text{MgII}} \approx 0.24 \text{ Mb}$ at threshold) and RR- n_e ; DR ($\sim 5 \times 10^5 \text{ K}$ resonances) and collisional ionization ($e^{-15 \text{ eV}/kT}$ at 10^4 K) are negligible. A test showed $\times 30$ more Mg II photoionization moves R_{eff} only $0.57 \rightarrow 0.32$ — the resonance line is so optically thick that ionizing it harder cannot confine it, so the ionization is a dead end as a lever (and is not the cause once the bin is matched).

Status: the full Phase-5 pipeline is implemented and runs end-to-end; the hydro/ $\dot{M}$ reproduces Huang (Case A 0.056, Case B 0.337); and the metal transit radii match Huang once compared in his own bins (Mg II $4 \text{ \AA}$ 0.22 vs 0.182, Na 0.175 vs 0.152). Remaining items are minor: Ca II $\sim 1.5\times$ high (line-center; plausibly Ca II $\rightarrow$ Ca III or the Ca abundance), the slightly-high Balmer $n = 2$, the $R_{py}R_{pz} = R_p^2$ normalization (methodological), a terminating Case D run, and phase5_transmission.ipynb. Comparisons must use Huang’s line-by-line definition ($4 \text{ \AA}$ bin for the NUV metals, line-center for the optical lines) — the two differ by $\sim 3\times$ for the optically-thick Mg II.

15. Post-process advection guard at a “breathing” base (2026-06-07)

Symptom. The spherical-vs-Roche example figure (user manual Fig. 4, WASP-121b) showed a sharp temperature spike at $r \approx 1.1 R_p$ in the *advection-corrected* (`_adv`) Roche profile, with a sawtooth band just below it. The converged equilibrium (eq) profile is smooth there, so the spike is a post-processing artifact, not a feature of the hydro solution.

Diagnosis (two compounding causes).

1. *A breathing, inflowing base (1-D limitation).* WASP-121b nearly fills its Roche lobe, so the 1-D wind is sub-sonic at L_1 and the dense lower atmosphere settles into a slowly recirculating state with small *negative* (inflow) velocities and a stagnation point ($v = 0$) near the wind base ($r \approx 1.1$). This is intrinsic to spherical symmetry, not a bug.
2. *A spurious hot root in the metal-cooled post-processor.* The advection post-processor re-solves the energy/ionization balance in each cell assuming an outflow (each cell upwinds from the next-inner cell). In the dense base the non-monotonic metal line-cooling curve gives the energy equation a second, spurious *hot* root that the Newton (MINPACK) solver can land on; the upwind coupling then cascades one bad cell outward, producing the sawtooth and the overshoot spike. The pre-existing $2\times$ band guard around eq T did not catch the $\sim 1.4\times$ spike.

Fix (option c: eq fallback where the flow is not a clean outflow). In `post_process_adv.f90`, wherever $v \leq 0$ (the breathing inflow base) the post-processor now *skips* the advection correction and keeps the converged eq ionization and temperature. Three guards were added, all gated on `pp_metal_on` so the metals-off mode stays byte-identical: the no-He and He H/He advection loops (set the H/He fractions to eq and cycle when $v(j) \leq 0$ or $v(j-1) \leq 0$) and the cell-by-cell temperature solve (set $T_{\text{out}} = T_{\text{in}}$ and cycle when $v \leq 0$). Keeping the inner cells at eq removes the spurious root *and* breaks the upwind cascade at its source, while the genuine `_adv` correction is retained in the outflow above the stagnation point.

Validation (PP-only re-run on the converged Case B state). Rebuilt and re-ran the post-processor (Do only PP on the converged IC):

- eq reproduced: $\max |\Delta T|/T = 2.6 \times 10^{-3}$, $\max |\Delta v| = 0.015 \text{ km s}^{-1}$; $\log_{10} \dot{M} = 13.38$ unchanged ($\approx 0.34 M_p \text{ Gyr}^{-1}$); Mg II reproduced to 1.6% in the line-forming region.
- `_adv` base oscillation eliminated: temperature sign changes in $[1.02, 1.10] R_p$ dropped from **36 to 0**; the spike is gone; the eq $\rightarrow$ `_adv` handoff at $v = 0$ is continuous (0.7% step) and the few-percent wind correction is preserved.
- Spherical Case A ($v > 0$ everywhere) is untouched — the guard never triggers, so its `_adv` is byte-equivalent.

Cause 1 (the breathing base) remains a 1-D limitation but is now represented there by the smooth eq solution; resolving the inflow itself would require a 2-D/3-D model.

Documentation updates (this session). The user manual (`EXHALE_user_manual.tex`) gained: §1.3 *The gravitational potential* (spherical, Roche substellar $\phi(r)$ and $d\phi/dr$, and the 3-D off-axis form, matching `grav_field.f90`); §3.5 *Spectral coverage for metal photoionization* (neutral-metal thresholds down to K I at 4.34 eV/2856 Å, and how the power-law sub-grid extends below 13.6 eV vs. a loaded SED that must itself reach those wavelengths); §8 *Atomic and molecular data* (a by-process source table — Asplund09 abundances, Verner96 photoionization, Voronov97 collisional ionization, Hui & Gnedin97 H/He + Badnell06 RR/ADF48 DR metals + Shull & Van Steenberg82 Ca I + Huang23 Fe, Huang23 charge exchange, AIOLOS+CHIANTI cooling, Christie13/Draine11/Huang17 excited-H, NIST line data; and the explicit note that the code is purely atomic/ionic with *no* molecular data); and the §4 *breathing-base caveat* documenting the issue and fix above. The Fig. 4 example was regenerated from the corrected `_adv` output.

16. He 2³S triplet + metals simultaneous solve, Ly α -flux guard, and convergence workflow (2026-06-08)

Motivation. Until now the He metastable triplet (He 2³S, “HeITR”) and the trace-metal ionization network could not be solved in the same run: both placed an extra unknown at `x(4)` of the coupled ionization residual, and `input_read.f90` allocated the metal unknowns only if `(thereis_metals .and. .not. thereis_HeITR)`. Modelling HD189733b with the triplet *and* solar metals requires both at once.

Merged solver (factual). New module `nonlinear_system_solver/System_HeH_TR_metals.f90` (ion_system_HeH_TR_metals, 84 SLOC). The H/He/triplet rows ($x_1 = \text{HII}$, $x_2 = \text{HeII}$, $x_3 = \text{HeIII}$, $x_4 = \text{He } 2^3\text{S}$) are taken verbatim from `System_HeH_TR`; the metal rows are taken verbatim from `System_HeH_metals` and shifted to $x(5+2(e-1))$. The two blocks share one electron density $n_e = n_{\text{HII}} + n_{\text{HeII}} + 2n_{\text{HeIII}} + \sum_e (n_{X^+} + 2n_{X^{++}})$ (the neutral triplet is excluded), and charge exchange is added through `cx_add_to_fvec`. Solved with `hybrd1`. Supporting edits: `charge_exchange.f90` gained a settable base index `cx_metal_base` (default 4; the merged dispatch sets 5) so the metal charge-exchange rows land at the right offset; `input_read.f90` sets $N_{\text{eq}} = 4 + 2n_{\text{melem}}$ when both are active; `ionization_equilibrium.f90` uses a generalized metal base index (`mbase = 4`, or 5 with `HeITR`) in the initial guess / solution scatter and a both-on dispatch branch; the new source was added to the Makefile. The radiation grid, cross sections, photo-heating, IC setup and post-processing already handled both species; only the equilibrium *solver* needed merging.

Regression & physical check. With metals on / triplet off the run is byte-identical to the pre-merge binary over 8 048 steps (the new branch is not taken). With both on, the coupled solve is clean (no NaN) and was exercised end-to-end on WASP-121b with the full stack (triplet + metals + excited-H + $\text{Ly}\alpha$), reaching a flat converged solution (peak $\text{He } 2^3\text{S}$ fraction $\sim 5 \times 10^{-6}$).

$\text{Ly}\alpha$ escape-probability input guard (factual). `input_read.f90`: if `Jlya escape-prob: True` (`jlya_mode=2`) but `Stellar Lya flux [erg/cm2/s]`: is missing or ≤ 0 , the run now stops with an explanatory message rather than silently dropping the stellar beam (in mode 2 $\bar{J}_{\text{Ly}\alpha} = \bar{J}_{\text{int}} + \bar{J}_*$ with $\bar{J}_* \propto F_{\text{Ly}\alpha,*}$, so $F_{\text{Ly}\alpha,*} = 0$ would make the halfwidth/boost settings silent no-ops).

Convergence threshold and Reconstruction scheme:-controlled two-stage reconstruction (factual).

- `du_th` is now a runtime variable (was a compile-time parameter) with default $1.0\text{e-}3$ — the original ATES-Code-main value. EXHALE had carried $2.0\text{e-}2$, at which a run flagged “converged” can still show $\sim 2\%$ spatial spread in the supersonic mass flux $\rho v r^2$.
- The single-stage vs. two-stage choice is now set by the `Reconstruction scheme: line`. `Reconstruction scheme: PLM+WENO3` selects a two-stage run: with the input line `du_th [PLM,WENO3]: <du_plm> <du_final>` supplying both thresholds (`du_th_plm` and `du_th`), `EXHALE_main.f90` starts in PLM and switches `rec_method` to WENO3 once $du < du_{\text{plm}}$ (or PLM stalls), then converges at $du < du_{\text{final}}$. This automates the manual `PLM → Load IC+WENO3` workflow recommended in the ATES README. `Reconstruction scheme: PLM` or `WENO3` is single-stage and uses only the first `du_th` value (the second is ignored).
- CFL is now runtime and settable via the input line `CFL: <value>`.

Other source-term changes (kept, with caveats). `energy_semi_implicit.f90` now damps the implicit energy update with $\partial F / \partial T = 1 + c |dC/dT|$ (was $\max(0, dC/dT)$), adding damping on the falling cooling branch; it removed one oscillation in testing but did *not* by itself fix the convergence behaviour below. `Cool_coeff.f90` now interpolates the cooling tables with monotone PCHIP (C^1) rather than linear (C^0); this was implemented to test an interpolation-kink hypothesis (not supported, below) and is retained as a smoothness improvement.

Convergence — provisional observations (not conclusions). *The following is a research log, deliberately tentative; several intermediate “non-convergence” readings in earlier sessions were later traced to methodology, not physics.* Two artifacts are recorded so they are not repeated: an over-short step cap (weak-wind cases judged “floored” near 25k steps while du was still descending — HD209458b later reached $du < 2 \times 10^{-2}$ near 50k steps), and a run-script bug that killed the backgrounding subshell rather than the EXHALE.x child, so runs continued orphaned and the reported du was a premature snapshot (use `kill -x EXHALE.x`). Because of these, the earlier framing that “metal line cooling destabilizes the wind” should be treated as *not established*. With the two-stage strict- du_{th} pipeline we currently see, across planets:

- **WASP-121b** (inflated, strongly irradiated): converges quickly and flat ($\sim 0.16\%$ supersonic mass-flux spread, vs $\sim 2\%$ for single-stage PLM). This case looks solid.
- **HD209458b**: reaches a small $du \sim 0.02\text{--}0.025$ that then oscillates rather than dropping to 10^{-3} . *Tentatively* usable as a quasi-steady state: a He 2^3S+7 -metal snapshot from the two-stage pipeline gives a smooth transonic wind with $T_{\max} \approx 6.8 \times 10^3$ K and $\log_{10} \dot{M} \approx 10.1$ (median over the supersonic region), consistent with the literature $\sim 10^{10} \text{ g s}^{-1}$; see `heitr_metals_dev/HD209458b_quasisteady.ipynb`.
- **HD189733b** (strongly bound, $\beta_0 \approx 192$): du oscillates around ~ 2 in a sustained limit cycle and does not reach the threshold — the problematic case.

A plausible (but unproven) reading is that the difficulty tracks how weakly driven / strongly bound the wind is, rather than the metal content per se.

Base-breathing hypothesis and the fixes that did *not* work. The HD189733b oscillation resembles the known base-breathing behaviour, and the lower boundary (`Apply_BC.f90`, `BC_component_constrho`) hard-pins the ghost density and pressure to cold reservoir values with a one-way velocity valve $\max(v_1, 0)$ — a configuration that can reflect acoustic waves. Three variants were tried and reverted:

Test	Outcome
Lower CFL ($0.6 \rightarrow 0.2$)	du floor barely moved ($\sim 2.8 \rightarrow \sim 1.9$)
Zero-gradient base velocity (drop the valve)	du still ~ 2.3
Riemann-invariant base velocity $v = v_1 + \frac{2(c_p - c_l)}{\gamma - 1}$	did not help HD189733b and broke WASP-121b

So the base-breathing reading is *consistent with* the observations but is **not demonstrated**, and these simple BC variants are not solutions. Candidate routes if pursued (both substantial, uncertain payoff): a characteristic / non-reflecting (NSCBC-style) base BC; a steady-state Newton / BVP solver (which would not orbit a time-marching limit cycle if a steady solution exists); or accepting HD189733b as a marginal, time-averaged quasi-steady case. A separate, cheaper idea for *speed* (not the oscillation) is local cell-by-cell time-stepping, since the global dt is currently set by the smallest base cell (`eval_dt.f90`). Full working notes: `docs/heitr_metals_and_convergence_notes.md`; solver-option discussion: `docs/numerical_methods.md`.

Code-size note. The merge adds one Fortran module (System_HeH_TR_metals.f90, 84 SLOC), so the 2026-06-07 file tally of 57 becomes 58; the remaining changes are small edits to charge_exchange.f90, input_read.f90, ionization_equilibrium.f90, parameters.f90, EXHALE_main.f90, energy_semi_implicit.f90 and Cool_coeff.f90.

17. Metals in the bulk-gas mass/charge/particle budget (eos_metals, 2026-06-13)

Motivation. Until this update the trace metals were strictly *passive* species: they contributed to heating/cooling and were solved in the coupled ionization equilibrium, but they were excluded from the mass density, the base density ρ_{bc} , and the electron/particle densities used everywhere *outside* the ionization solver (EOS temperature, opacities, cooling calls, energy update). The only self-consistent place was the charge balance *inside* the MINPACK residuals (System_HeH_metals.f90). At solar abundances the missing metal mass is $\sim 1.3\%$ of the H/He mass — small but systematic, and inconsistent with solving the metal ionization to high precision. Recorded in TO_BE_DONE.md on 2026-06-12; implemented 2026-06-13.

Runtime interface. A new key eos_metals 0|1 in metals.inp (documented in inputdata/metals.inp.example). 1 (the new default) includes the metals in the bulk-gas budget; 0 restores the legacy trace approximation exactly. Metals-off runs (no metals.inp) are identical either way.

What changed (factual).

- **Composition constants** (input_read.f90): the mass per hydrogen nucleus is now $\mu_H \equiv \text{mass_per_H} = 1 + 4Y + \sum_X a_X A_X$ (with $Y = \text{HeH}$ and a_X, A_X the metal abundance and atomic weight), and the base density becomes $\rho_{bc} = \mu_H/(1+Y)$ in units of $m_H n_0$. Atomic weights live in species_table.f90:melem_A. The normalization n_0 keeps its H/He-nuclei meaning, so $n_H = n_0/(1+Y)$ is unchanged. (Correction to the original TO_BE_DONE entry: v_0, p_0, b_0 are m_H -based normalizations and never contained a mean molecular weight — they were never affected.)
- **Mass density:** calc_rho takes an optional nm argument and adds $\sum A_X n_{X,i}$ over all stages; passed from ionization_equilibrium.
- **Electron density:** calc_ne takes an optional nm and adds $\sum q_i n_{X,i}$; passed at every site where nm is available — composition (main-loop EOS), ionization_equilibrium (opacity/cooling), util_ion_eq (eval_cool + cooling dump), energy_semi_implicit, excited_hydrogen, post_process_adv. (The scalar T_equation already carried metal electrons via pp_nm_cell.)
- **Particle density:** calc_ntot likewise adds the metal nuclei, so the EOS $T = p/(n_{\text{tot}} + n_e)$ sees them.
- **Ghost pressure / dp_bc:** the lower-BC ghost pressure is now $n_{\text{tot_bc}} + \text{dp_bc}$ with $n_{\text{tot_bc}} = (1+Y + \sum_X a_X)/(1+Y)$ — so $T = T_0$ holds *exactly* at the base — and dp_bc includes the metal electrons $\sum q_i n_{X,i}/n_0$ at the first cell. set_IC/load_IC use mass_per_H/ntot_bc consistently. Full formulas: EXHALE_BC_and_IC.tex.

Validation (WASP-121b regression matrix).

- `eos_metals 0` is **byte-identical** to the pre-change goldens for both regression cases (`wasp_full 7296` steps, `wasp_he23off 7300`) — the legacy path is untouched.
- `eos_metals 1` (new default) converges cleanly (7145/7142 steps, $\sim 2\%$ fewer); profile shifts are at the expected $\sim 1\%$ mass scale: median $|\Delta n| \sim 2.4\%$, $|\Delta p| \sim 2.3\%$, $|\Delta T| \sim 0.9\%$ (max $|\Delta T| = 132$ K in the breathing-base region $r \sim 1.09$). The *number* flux nvr^2 drops 1.1% ($\log_{10} \dot{M}$: 13.31 $\rightarrow$ 13.30), which offsets the +1.2% mass per particle — the *mass*-loss rate is essentially unchanged.
- Goldens re-snapshotted under the new default; a repeat run reproduces them byte-identically (single-thread determinism intact).

Still open (small). The `use_2lev_cool` legacy cooling branch remains AIOLOS-hard-coded (default off), and the cell-by-cell scalar `T_equation` uses `pp_nm_cell` only in the post-process path (the time-marching cell-by-cell solve receives n_e via the now-metal-aware `energy_semi_implicit`).

18. Automatic initial-condition selection (IC mode: auto, 2026-06-13)

Motivation. The code now carries three IC families — cold hydrostatic (default), transonic isothermal wind (added for the Phase-4 deep-RLOF Case D, where the cold IC false-converges), and the hot-Parker warm seed (benchmarked: no speedup) — and the choice was manual. Kubyshkina+2018 build an IC for each planet automatically; the design study is `docs/auto_ic_design.md` and the literature comparison `docs/code_comparison.tex`.

The selector (one exact probe, no tunable threshold). `select_IC_auto` (`set_IC.f90`) evaluates the *cold* base sound speed in code units, $c_{\text{cold}}^2 = (n_{\text{tot,bc}} + \Delta p_{\text{bc}}) / \rho_{\text{bc}}$, and calls the existing `find_sonic` bisection for a critical point $\phi'(r_c) = 2c^2/r_c$ in the *actual* potential (Roche or spherical). If an interior sonic point exists, the cold isothermal atmosphere already wants to blow a transonic wind, so the transonic IC is selected; otherwise the cold hydrostatic IC is kept. One criterion catches both regimes that need the transonic start: deep RLOF (through the Roche/L1 topology) and low-gravity boil-off ($r_c \simeq b_0/(2c_0^2)$ in a point-mass potential). The Jeans parameter b_0 (Kubyshkina’s Λ) is logged as a diagnostic but is *not* a decision variable; the hot-Parker seed is never auto-selected.

Runtime interface. New input line `IC mode: <cold|transonic|hot_parker|auto>`. The named modes simply set the legacy flags; `auto` defers the decision to `set_IC`. The legacy keys (`Transonic IC:`, `Hot Parker IC:`) take precedence over `auto`, and the setup report (`EXHALE_setup.out`) records the family actually in effect, with a “(chosen automatically: IC mode = auto)” tag. Touched files: `parameters.f90` (`ic_mode`), `input_read.f90` (key parsing), `set_IC.f90` (`select_IC_auto`), `write_setup_report.f90`.

Validation gates (all run 2026-06-13).

Gate	Result
A-1 no IC mode key	byte-identical regression (defaults untouched)
A-2 IC mode: cold	byte-identical to A-1
A-3 spherical tutorial, auto	$b_0 = 83.2$, no interior $r_c \rightarrow$ cold; IC_dump byte-identical
B-1 Case D (deep RLOF), auto	$r_c = 1.297$ (just inside $L1 = 1.353$) $\rightarrow$ transonic, same as the manual key
B-2 boil-off planet ($0.0189M_J$, $0.446R_J$, 1100 K)	$b_0 = 8.5$, $r_c = 5.2$ (as predicted) $\rightarrow$ transonic; wind launches (initial residual 0.19 vs 14 cold)
B-3 wasp_full (Roche), auto	$\rightarrow$ transonic; du -stop in 5876 vs 7145 steps (−18%)

Honest speed verdict. The −18% step count is *Roche/deep-RLOF-specific* (the steady wind there is close to the cold isothermal transonic solution). For the mild spherical boil-off case B-2 the cold control *also* launched and converged first ($\log_{10} \dot{M} = 11.20$), while the auto run relaxed monotonically toward the same flux from $\sim 2\times$ above — the cold- c^2 Parker profile *overestimates* the flux of a weakly heated wind. So auto is not harmful (same attractor) but not faster there.

Key byproduct (a trap for all $\dot{M}$ work). While adjudicating gate B-3 it was found that the $du < 10^{-3}$ stopping criterion is **path-dependent at the $\sim 5\%$ level**: the cold-start and transonic-start wasp_full runs both “converge” by du but sit 4.8% apart in mass flux, each only 1–2% flux-flat. Newton-finishing *both* states (Load IC + Solver: Newton) collapses them onto the *same* fixed point (flux ratio 0.9996, flatness 0.005%). Practical rule, now standing: **for quantitative $\dot{M}$ always finish with Solver: Newton**; never compare bare du -stops at the percent level.

Deferred. A Phase-C retry ladder (escalate cold $\rightarrow$ warm seed on NaN/stall) is designed in auto_ic_design.md §6 but *not* implemented — the B-2 result suggests failure-triggered escalation is the right shape if a real case ever needs it. IC formulas and the selection logic are also documented in EXHALE_BC_and_IC.tex §5 and the user manual.

19. Base boundary conditions and the flux-based convergence decision (2026-06-15)

Context. A systematic study of how the lower boundary and the stabilizers interact with convergence on HD 209458b is collected in docs/EXHALE_BC_and_IC.tex (test matrix, two stabilizer tables, and a convergence-criterion history). This entry records the resulting default changes to the code; the companion document holds the full evidence.

Shapiro filter: off by default. The optional Shapiro (1970) spatial filter — ported from CETIMB (Koskinen et al. 2013a), whose use of it was found through Huang et al. (2023), to damp the same gravity-unbalanced base sound-wave instability — damps the transient “breathing” of the base cell, but on a clean cold-start it also drives the wind into a slow infall and prevents convergence. Isolated in a 2×2 (filter $\times$ base-valve) sweep, the filter — not the valve — was the cause. It is therefore *off* by default (shapiro_eps < 0) and opt-in only via Shapiro filter: <eps> [every].

Base boundary condition: density vs. pressure. The base can now be anchored either on density (the legacy fixed n_0) or on a fixed base pressure, selected by Base BC: density|pressure [p_ubar] (default density; pressure fixes p_{base} in μbar and derives n_0). In practice the converged wind is nearly insensitive to which anchor is used — the real lever on $\dot{M}$ is the base *density* itself, not the BC form.

The flux-based convergence decision (history). EXHALE measures two distinct quantities at the base: the flux metric du (the radial spread of $\rho v r^2$, which is exactly the constant-mass-flux criterion of ATES, $\Delta\dot{M}/\dot{M} < 10^{-3}$, and of CETIMB, $F_c = \rho v r^2$ altitude-flat), and the steady residual $\|R\| = \|\partial_t u\|$ over mass/momentum/energy, a *stricter* EXHALE-specific check. The convergence rule evolved: before the Newton finish it was purely flux-based; the Newton work then added the stronger $\|R\|$ criterion; but the reference codes converge on flux alone. The standing decision, adopted here, is therefore **flux-based** — $du < du_{\text{th}}$ is the acceptance test, and $\|R\|$ is computed and reported *for reference only* (it gates the stop only if Resid tol: is set).

Volume-weighted residual norm. When $\|R\|$ is reported, the default is now a cell-volume-weighted relative norm, $\sum_j |R_j| r_j^2 \Delta r_j / \sum_j |u_j| r_j^2 \Delta r_j$ (resid_vol, settable Resid norm: vol|Linf), rather than the bare L_∞ maximum. On the non-uniform radial grid the L_∞ max is dominated by the few tiny near-base cells and overstates the residual; the volume weighting gives a physically representative number. The infall diagnostic used in the test matrix was likewise switched to a *radius-weighted* negative-velocity fraction (weighted by Δr over the linear radius range), so that sub- 10 m s^{-1} near-base noise is not counted as bulk infall.

du-keyed Newton/JFNK hand-off. The production Newton finish (Solver: Newton) now hands off from the marching warm-up to the JFNK steady solve when the *flux* metric drops below newton_du_switch (default 10^{-2}), consistent with the flux-based decision and adjustable via an optional third token, Solver: Newton [du_switch]. This replaces the earlier residual-keyed hand-off ($\|R\| < 5 \times 10^{-2}$, variable newton_R_switch, now removed); the two are equivalent in practice, but keying on *du* avoids tying the hand-off to the stricter residual. The JFNK solve still *targets* $\|R\| < \text{resid_th}$ (10^{-3}); on failure it keeps its best iterate and falls back to marching with the *du*-based stops re-armed.

The full-physics residual floor (base $j = 1$). With He 2^3S + metals, the *du*-keyed hand-off plus the volume-weighted norm drove the JFNK residual from ~ 0.17 down to $\sim 2 \times 10^{-3}$ — about two orders of magnitude better — but it then *stalls* (info $\neq 0$) on a localized momentum imbalance at the base cell $j = 1$ (the “breathing base”), short of 10^{-3} . The wind is nonetheless flux-converged (e.g. HD 209458b full physics: $\log_{10} \dot{M}$ steady, $\rho v r^2$ spread $< 1\%$), which is a converged run by the reference standard. Closing the last factor of ~ 2 in $\|R\|$ requires physically damping the base momentum — the explicit-viscosity task at the top of TO_BE_DONE.md.

Touched files. parameters.f90 (shapiro_eps, base_bc_mode/base_p_ubar, base_v_massflux, resid_vol, newton_du_switch; newton_R_switch removed); input_read.f90 (Base BC:, Resid norm:, Solver: Newton [du_switch] parsing); steady_residual.f90 (residual_norms_vol); steady_newton.f90 (resid_relnorm vol/ L_∞ branch); EXHALE_main.f90 (*du*-keyed JFNK trigger, always-on flux/residual diagnostic). Full detail and the test matrix: docs/EXHALE_BC_and_IC.{tex,pdf}.

20. Tidal (Roche) potential with an extended outer boundary (2026-06-16)

Motivation. The two domain modes were coupled to two different gravitational potentials. Domain mode: Spherical extends the grid to a user Outer radius but *drops* the stellar-tidal and centrifugal terms (pure $-b_0/r$); the default Roche mode keeps the full tidal potential but truncates the domain at the inner Lagrange point L_1 . Neither matched the escape models of Yan et al. (2022) / Huang et al. (2023), which retain the tidal force *and* integrate out to $\sim 10R_p$ (well past L_1). For WASP-52b, for example, $L_1 \simeq 2.53R_p$, so Roche mode truncated the domain far short of the $10R_p$ used by Yan, cutting off the outer He 10830 / H α absorbing layers.

Change. The outer-boundary extent is now decoupled from the choice of potential. In Roche mode (no Domain mode line, so the full tidal potential of grav_field.f90 is retained), supplying an explicit Outer radius [R_p]: > 1 sets r_max to that value instead of the L_1 radius:

```
else ! Roche mode: full tidal potential retained
    r_max = (3.0*Mrapp)**(-1.0/3.0)*atilde ! L1 (default)
    if (r_out_user .gt. 1.0d0) r_max = r_out_user ! NEW: extend past L1
endif
```

This reproduces the Yan/Huang “tidal force + extended domain” setup with a two-line change and no new keyword (the existing Outer radius is simply honored in Roche mode too).

Caveats and usage. Beyond L_1 the Roche potential is past its maximum (net outward force), so the 1-D radial flow there is a spherical approximation to what is physically an L_1 funnel — the same approximation Yan/Huang make. The cold hydrostatic IC is invalid past L_1 , so pair the extended domain with IC mode: auto (validated: the auto selector finds an interior cold sonic point and seeds a transonic wind). The tidal singularity of the $-b_0 q/(\tilde{a}-r)$ term sits at the star ($r = \tilde{a} \gg 10$), so a $10R_p$ domain is numerically safe. Domain mode: Spherical (no tidal) remains available for Huang Case-A-like comparisons.

Touched files. input_read.f90 (honor r_out_user in the Roche branch of the r_max assignment). Documented in docs/EXHALE_user_manual.tex (§ potential and the input-options table).

21. Photoionization cross-section benchmarks and a TOPbase high-energy extension of He 2³S (2026-07-01)

Motivation. The He 2³S (metastable triplet) photoionization cross section sigma_HeI23S — the drain term that sets the He 10830 lower-level population — was a four-segment broken power-law fit to Norcross (1971), hard-cut to zero above 59.2 eV. Two questions: (i) are EXHALE’s H/He/He 2³S photoionization cross sections consistent with modern atomic data, and (ii) does the unphysical 59 eV cutoff matter? Both are settled in the new memo docs/photoion_cross_sections.{tex,pdf} (reproducible via docs/compare_photoion_cross_sections.py).

Benchmarks. H I and He II use the exact hydrogenic form; He I uses the two-term ATES fit, which agrees with the Verner+1996 (VFKY96) single-shell fit to within 14% over 24.6–500 eV. The He 2³S broken PL reproduces its parent Norcross/p-winds table to ~10% and — checked against the **TOPbase / Opacity Project** R-matrix cross section (extracted from the CDS server; raw 922-point table saved to `data/topbase_HeI_2_3S_photo.dat`) — agrees to ~8% in the smooth near-threshold region (5–20 eV). TOPbase additionally resolves a 40–55 eV autoionizing-resonance forest (peaks to 3223 Mb); their **net effect on the 2³S rate is only** ~1–2% once the few numerically under-resolved spikes are excluded (the raw table must *not* be fed into a broadband rate integral — a quadrature artifact). Extending the cross section above 59 eV raises the 2³S rate by only 0.1–1.2%.

Change. `sigma_HeI23S` is extended with a minimal, fully-connected edit (`src/modules/functions/cross_sec.f90`): the broken PL is kept unchanged up to the 45.6 eV bump; the **last power-law segment is re-aimed** so it runs continuously from the bump peak to a power-law fit of the TOPbase smooth background at the junction $E_J = 70$ eV,

$$\sigma_{\text{tail}}(E) = 10^{5.4623} E^{-2.9964} \text{ Mb} \quad (E > 70 \text{ eV}),$$

(the last-segment slope changes only $-3.039 \rightarrow -2.658$; the intercept is derived in-code for exact continuity). The old 59 eV cutoff is dropped, so 59–1240 eV now carries the TOPbase tail instead of zero. The result is C^0 -continuous at *both* joins (45.6 eV and 70 eV); the original pre-modification formula is kept in comments. A unified full-range PCHIP node table and a tail-only power law are documented as alternatives, but this is the recommended drop-in (smallest rate change, retains the validated threshold-to-bump shape exactly).

Validation. A controlled before/after on WASP-52b (`fxuv0p25_he98_L1`, Include He23S? True, ϵ Eri SED with X-rays), same IC, only `cross_sec.f90` differing: the bulk wind is unchanged ($\log \dot{M} = 11.49$ both), and the He 10830 line changes by $< 0.001\%$ in every metric (line-center depth $15.2675\% \rightarrow 15.2675\%$; EW $0.0877 \rightarrow 0.0877 \text{ \AA}$) — below the Newton convergence noise. The extension is more physically correct yet observationally harmless for He 10830, as predicted by the +0.65% rate estimate (even smaller here because the ϵ Eri SED is soft).

Recombination cross-check. A companion memo `docs/recombination_coefficients.{tex,pdf}` (`compare_recombination.py`) documents the H/He/He 2³S recombination: H II, He II, He III use the Hui & Gnedin (1997) case-B fits ($\approx$ Draine 2011 / Verner & Ferland to ~2%); He II $\rightarrow$ He I splits into Benjamin+1999 singlet/triplet (Oklopčić & Hirata 2018) when the 2³S network is on. The He 2³S triplet rate is identical to the latest escape models (Benjamin+1999). Newer nebular-community data (Porter+2012/2013, Del Zanna & Storey 2022) differ ~5–15% but are not yet adopted by escape codes; H/He case-B needs no update.

Touched files. `src/modules/functions/cross_sec.f90` (`sigma_HeI23S`: re-aimed last segment + TOPbase tail, 59 eV cutoff removed, original kept in comments). New docs: `photoion_cross_sections.{tex,pdf}`, `compare_photoion_cross_sections.py`, `data/topbase_HeI_2_3S_photo.dat`, `recombination_coefficients.{tex,pdf}`, `compare_recombination.py`, and figures under `docs/figures/`.

22. He I (1^1S) photoionization $\rightarrow$ Verner+1996 by default, with ATES-rate switches (2026-07-01)

Motivation. The He I *ground-state* (1^1S , singlet) photoionization used the legacy ATES two-term fit, which runs $\sim 12\text{--}14\%$ below the Verner+1996 (VFKY96) single-shell fit in the 50–100 eV band (`photoion_cross_sections.tex`, §4). We make Verner the default, keep the old behavior behind a switch, and add a parallel switch for the recombination data.

Singlet/triplet separation. He photoionization in EXHALE is already split by spin: `sigma_HeI` is the 1^1S **ground (singlet)** cross section (updated here to Verner — Verner’s He I *is* the 1^1S ground state), while `sigma_HeI23S` is the 2^3S **metastable (triplet)** cross section (Norcross/TOPbase, §21). The 2^1S singlet metastable is not tracked (only 2^3S , the 10830 lower level), so the Verner update correctly touches only the singlet ground state.

Change. A single backward-compatible switch in `input.inp`: `ATES_photoionization_rate: True` reverts He I (1^1S) photoionization to the legacy two-term fit (default = Verner+1996, via `sigma_VFKY96` with He I parameters $[E_{\text{th}}, E_0, \sigma_0, y_a, P, y_w, y_0, y_1] = 24.59, 13.61, 949.2, 1.469, 3.188, 2.039, 0.4434, 2.136$). No recombination switch is needed (see below).

Comparison. WASP-52b (`fxuv0p25_he98_L1`, He23S active, same converged IC). Switching He I photoionization legacy $\rightarrow$ Verner changes the He ionization balance and hence the metastable population (Table 1): unlike the $> 59\text{ eV}$ tail (§21, $< 0.001\%$), the ground-state photoionization choice is a **real $\sim 3\%$ effect on He 10830** — it matters for the absolute depth and the inferred He abundance. `ATES_photoionization_rate: True` reproduces the old result exactly.

Table 1: He I photoionization legacy (ATES two-term) vs Verner+1996 (new default), WASP-52b He 10830.

metric	legacy (ATES)	Verner (new)	change
$\log \dot{M}$ [g/s]	11.49	11.48	-2%
He 10830 line-center depth	14.516 %	14.933 %	$+2.9\%$
He 10830 equivalent width	$0.0877 \text{ \AA}$	$0.0910 \text{ \AA}$	$+3.8\%$

Recombination — no update needed (EXHALE already current). The companion update (He recombination $\rightarrow$ Porter+2012/2013 / Del Zanna & Storey 2022) was investigated and found to be a *no-op*: those are *emissivity* papers — they update the He I line cascade / line ratios, *not* the total (effective) recombination to 2^3S , and provide no drop-in $\alpha(2^3\text{S})/\alpha(1^1\text{S})$. The escape-relevant quantity — total recombination into the triplet ladder, all funnelling to the metastable 2^3S — is set by (case-B total) $\times$ (triplet spin fraction $3/4$) and is robust across modern treatments: EXHALE (Benjamin+1999) has $\alpha_3 = 2.10 \times 10^{-13} T_4^{-0.778}$, $\alpha_1 = 1.54 \times 10^{-13} T_4^{-0.486}$, while the 2025 thermosphere paper (arXiv:2509.14499) uses $\alpha_3 = (3/4)(2.72 \times 10^{-13} T_4^{-0.789}) = 2.04 \times 10^{-13} T_4^{-0.789}$ and the *identical* singlet $\alpha_1 = 1.54 \times 10^{-13} T_4^{-0.486}$ — so α_3 agrees to $\sim 2\text{--}3\%$ and α_1 is identical. EXHALE’s He recombination already matches the latest practice; no change is warranted. (A trial direct sum of Cloudy’s `he_iso_recomb.dat` level-by-level radiative

recombination gave a total $\sim 2\times$ the case-A value owing to its unmapped 1642-level indexing, confirming the raw file is not a usable shortcut.) No recombination switch was added.

Touched files. `src/modules/init/parameters.f90` (ates_photoion_rate switch), `src/modules/files_IO/input_read.f90` (parse the optional key), `src/modules/functions/cross_sec.f90` (sigma_HeI branches Verner default / legacy two-term). Verified: Verner is the default (He I σ at 30/50/100 eV = 5.36/2.02/0.394 Mb); ATES_photoionization_rate: True recovers the legacy values (5.24/1.78/0.347 Mb).

23. He 2³S photoionization → two VFKY96 (Verner) forms + power-law bridge (2026-07-01)

Motivation. `sigma_HeI23S` was a bespoke five-segment broken power law (Norcross-1971 fit + TOPbase-tail extension, §21). Since EXHALE already uses the Verner+1996 (VFKY96) `sigma_VFKY96` routine for He I and all metals, the metastable triplet is now expressed in the same form, so every photoionization cross section shares one analytic kernel.

Change. (`src/modules/functions/cross_sec.f90`.) `sigma_HeI23S` is rebuilt from two VFKY96 single-shell forms joined by a log-linear (power-law) bridge, used piecewise, with $[E_{\text{th}}, E_0, \sigma_0, y_a, P, y_w, y_0, y_1]$:

- **wing A** ($E \leq x_3 \simeq 34.70$ eV, threshold → Cooper minimum): 4.78, 2.645, 20.8, 1.0×10^{12} , 3.42, 2.681, 1.956, 2.603;
- **bridge** ($34.70 < E < 45.59$ eV): straight log–log line joining wing A(x_3) to wing B(x_4) (the “power-law part in the middle”);
- **wing B** ($E \geq x_4 \simeq 45.59$ eV, bump → tail): 45.59, 49.68, 1052, 0.04393, 2.941, 1.717, 5.488×10^{-5} , 1.118.

Each wing is a $\sim 4\%$ fit to the resonance-averaged cross section; the construction is **continuous at both** x_3 **and** x_4 , and threshold ($E < 4.78$ eV) and the smooth high- E tail ($\sim E^{-7/2}$) are supplied by the wings (no hard cutoff). The former broken power law is retained, commented out, in the routine.

Verification. Rebuilt EXHALE.x; `sigma_HeI23S` reproduces the former broken power law to $\sim 4\%$ (12 eV: 1.466 vs 1.465; 30: 0.289; 45.6: 2.57 vs 2.68 — bump 4% low; 70: 0.855 vs 0.858; 100: 0.294; 300: 0.0110), continuous across the bridge (34.70 eV: 0.222 → 38 : 0.502 → 40 : 0.795 → 45.59 : 2.570 → 46 : 2.521 Mb).

Touched files. `src/modules/functions/cross_sec.f90` (`sigma_HeI23S`: two VFKY96 wings + power-law bridge; former broken PL kept commented out).

24. Temperature-dependent Penning ionization of He 2³S (2026-07-01)

Motivation. The metastable He(2³S) destruction by neutral hydrogen, He(2³S)+H, was carried as a *temperature-independent* constant $Q_{31} = 5 \times 10^{-10} \text{ cm}^3 \text{ s}^{-1}$ — the ATES/Oklopčić & Hirata (2018) value, itself the Roberge & Dalgarno

(1982) *sum* of Penning and associative ionization, as also used by Lampón et al. (2020). Taylor et al. (2025, ApJ 989:68) replaced this with a Maxwell–Boltzmann-averaged, **temperature-dependent** fit to the cross sections of Morgner & Niehaus (1979) and Cohen & Lane (1971). Because Penning ionization is often the *dominant* 2^3S loss near the base (where neutral H is abundant), this directly affects the He 10830 prediction EXHALE is built to compute.

New default rate (Taylor et al. 2025, Table 2). A two-branch power law ($\text{cm}^3 \text{s}^{-1}$, T in K):

$$Q_{31}(T) = \begin{cases} 1.9 \times 10^{-9} (300/T)^{0.07}, & T \leq 4000 \text{ K}, \\ 9.1 \times 10^{-9} (300/T)^{0.50}, & T > 4000 \text{ K}. \end{cases} \quad (4)$$

This is $\sim 3\times$ the legacy constant across the relevant 2000–9000 K range. Note the new rate is the Penning channel only ($\text{He}(2^3\text{S})+\text{H} \rightarrow \text{He}(1^1\text{S})+\text{H}^++\text{e}^-$); the associative-ionization channel ($\text{He}(2^3\text{S})+\text{H} \rightarrow \text{HeH}^++\text{e}^-$) is not represented separately in EXHALE (no HeH^+), so the single Q_{31} continues to carry the dominant $2^3\text{S}+\text{H}$ destruction term $-n_{\text{HeITR}} n_{\text{HI}} Q_{31}$.

Unconditional default. The temperature-dependent rate is used *unconditionally*: the legacy constant 5×10^{-10} is no longer selectable. (The `ATES_photoionization_rate` flag continues to revert only the He I (1^1S) photoionization cross section to the legacy ATES 2-term fit; it does *not* affect the Penning rate.)

Implementation. Q_{31} became an array over the cells (it was a scalar). A new `penning_HeI_23S(T,coeff)` in `Cool_coeff.f90` evaluates the piecewise fit over the whole grid; `HeITR_coeffs` calls it unconditionally. The two callers (`ionization_equilibrium.f90`, `post_process_adv.f90`) pass $Q_{31}(j)$ per cell into `params`; the `System_HeH_TR*` residuals are unchanged (they already read Q_{31} from `params`).

Verification. Rebuilt EXHALE.x; ran HD 189733b (metals + He23S, step-capped) with the new rate and, for comparison, with the former 5×10^{-10} constant. The temperature dependence is active and physically sensible: relative to the old constant, $\text{He}(2^3\text{S})$ is reduced near the base (ratio ≈ 0.28 at $1.0R_p$, 0.92 at $1.1R_p$; column-integrated over $1\text{--}2R_p$ drops to $0.59\times$), and the profiles converge to unity above $\sim 1.3R_p$ where H is ionized and Penning switches off — the expected neutral-H-gated signature.

Touched files. `src/modules/radiation/Cool_coeff.f90` (new `penning_HeI_23S`); `src/modules/radiation/util_ion_eq.f90` (`HeITR_coeffs`: Q_{31} array + default/legacy branch); `src/modules/radiation/ionization_equilibrium.f90` and `src/modules/post_process/post_process_adv.f90` (Q_{31} array, `params` assignment in each cell).

25. He/H diffusive separation — Phase 1 (2026-07-01, experimental, default OFF)

Motivation. EXHALE is single-fluid and re-solves composition by *local* ionization equilibrium (conserving the element ratio), so He/H was frozen at the input HeH at all radii — it could not represent the diffusive He/H separation of Taylor et al. (2025) (8%→2.5%) and Xing et al. (2023). Top recommendation of `docs/methodology_aiolos_taylor_xing`.

Method. New module `species_diffusion.f90` transports the He element ratio $f = n_{\text{He}}/n_{\text{H}}$ with bulk advection and a molecular-diffusion drift of He relative to H, $\Phi = -Dn_{\text{H}}(\partial f/\partial r + fG)$, $G = (m_{\text{He}} - m_{\text{H}})g/(kT)$, $D = 1.52 \times 10^{18}(1/m_{\text{H}} + 1/m_{\text{He}})^{1/2}T^{1/2}/n$ (Banks & Kockarts 1973), plus eddy diffusion K_{zz} on the gradient term. Solved as a **fully conservative** implicit (backward-Euler tridiagonal) update of n_{He} ; the new split is written back into `f_sp` conserving mass ($\sum_s m_s f_{\text{sp},s} = 1$, metals frozen to H). Gated on `He_diffusion` (default off; `He_Kzz` sets K_{zz}). Derivation: `docs/design_hehe_diffusion.md`.

Root-cause bug (design doc §7). A local time scale named `t0` ($= R_0/v_0$) shadowed the *global temperature* `T0` (Fortran is case-insensitive) in `TK = Tcode*T0`, zeroing `TK` (floored to 1 K) and inflating the settling coefficient $G \propto 1/TK$ by $\sim 5800\times$, causing total over-settling. The two constants are numerically near-equal ($R_0/v_0 \approx 2.83 \times 10^4$ s, $T_0 \approx 2.83 \times 10^4$ K), which hid it. Fix: rename the local to `tscale`.

Result (HD 209458b, He23S+metals). Physical mild separation — $(\text{He}/\text{H})/\text{HeH} \approx 0.6\text{--}0.9$ in the inner thermosphere, falling to $\sim 0.17\text{--}0.24$ aloft — consistent in magnitude with Taylor ($\rightarrow 0.3\times$) and Xing. Stable (no NaN); `He_diffusion OFF` is byte-identical.

Phase-1 limitations. A near-base pile-up from the lagged- n_{H} nonlinearity is held by a physical limiter $f \leq \text{HeH}$ (aloft observable unaffected); proper fix is a self-consistent n_{H} inner iteration. Ambipolar field, thermal diffusion, metal diffusion, and a Newton-finished quantitative $\dot{M}$ are Phase-2. Flag stays default OFF.

Touched files. `new src/modules/functions/species_diffusion.f90`; `parameters.f90` (`he_diffusion`, `he_kzz`); `input_read.f90` (`He_diffusion`, `He_Kzz`); `EXHALE_main.f90`; `Makefile`.

26. He/H diffusion — Phase 2 (2026-07-02): ambipolar, thermal diffusion, metals

Extends §25 (design doc §7e). The Phase-1 kernel was factored into a shared subroutine `solve_1elem` (one conservative implicit advection–diffusion–settling solve for any element vs the H background), reused by He and each metal.

- **P2b — ambipolar-corrected settling** (default ON, `He_ambipolar`). The polarization field lifts ions, so the He-vs-H effective settling mass is $\Delta m_{\text{eff}} = 3 - 0.5(\bar{Z}_{\text{He}} - \bar{Z}_{\text{H}})$ (mean charges from the local ionization state): 3 at the neutral base, 2.5 in the fully ionized wind. Modest effect ($\sim 2\%$ less depletion aloft), correct direction.
- **P2c — thermal diffusion** (default `He_alphaT=0`). Adds $\alpha_T \partial \ln T / \partial r$ to the settling coefficient; off by default.
- **P2d — element-by-element metal diffusion** (default OFF, `He_metal_diffusion`). Each trace metal element diffuses independently vs n_{H} with its own mass (`melem_A`), binary D , and ambipolar correction; ion stages rescaled to the diffused total. Metals trace $\Rightarrow$ no n_{H} feedback. **Validated (HD 209458b C/N/O):** heavier elements deplete more — at $3 R_{\text{p}}$, `element/base` $\approx$ He 0.171, C 0.158, N 0.136, O 0.105 (monotonic in mass; matches Xing et al. 2023). Regression: OFF $\Rightarrow$ He/H unchanged, C/H frozen at 1.0.

- **P2a — self-consistent n_{H} : attempted, NOT adopted.** A Picard $n_{\text{He}} \leftrightarrow n_{\text{H}}$ iteration diverges (near-base settling genuinely concentrates He, driving $n_{\text{H}} \rightarrow 0$); the physical cap $f \leq \text{HeH}$ remains the correct limiter.

Touched files. `species_diffusion.f90` (solve_1elem kernel; ambipolar dmeff; thermal term; metal-diffusion loop); `parameters.f90` (he_ambipolar, he_alphaT, he_metal_diffusion); `input_read.f90` (He_ambipolar, He_alphaT, He_metal_diffusion).

27. Diffusive separation: detailed physics and numerics

This section consolidates §25–§26 into a single, more complete treatment of the He/H (and metal) diffusion, with emphasis on the two design decisions that are easy to get wrong: the $n_{\text{He}}-n_{\text{H}}$ mass coupling (§27.3) and the settling discretization (§27.5).

27.1 Governing equation

By default EXHALE is single-fluid: the hydro evolves $(\rho, \rho v, E)$ with one bulk velocity, and the composition is re-solved every step by *local* ionization equilibrium, which conserves the element ratios. Consequently the helium-to-hydrogen ratio $f \equiv n_{\text{He}}/n_{\text{H}}$ is frozen at the input value at all radii — EXHALE cannot represent diffusive separation without an added transport equation.

We transport the He element ratio f with advection at the bulk velocity plus a molecular-diffusion *settling* drift of He relative to H,

$$\frac{\partial f}{\partial t} + v \frac{\partial f}{\partial r} = \frac{1}{n_{\text{H}}} \frac{1}{r^2} \frac{\partial}{\partial r} \left[r^2 n_{\text{H}} (D + K_{\text{zz}}) \frac{\partial f}{\partial r} + r^2 n_{\text{H}} D f G \right], \quad (5)$$

with the diffusive number flux (relative to the bulk)

$$\Phi_{\text{He}} = -n_{\text{H}} \left[(D + K_{\text{zz}}) \frac{\partial f}{\partial r} + D f G \right], \quad G = \frac{\Delta m g}{kT} + \alpha_T \frac{\partial \ln T}{\partial r}. \quad (6)$$

Here $D = 1.52 \times 10^{18} (1/m_{\text{H}} + 1/m_{\text{He}})^{1/2} T^{1/2} / n$ is the Banks & Kockarts (1973) binary He-in-H diffusion coefficient ($\text{cm}^2 \text{s}^{-1}$), $g = GM_{\text{p}}/r^2$, and $\Delta m = m_{\text{He}} - m_{\text{H}}$ ($= 3$ amu in the neutral limit; §27.6). K_{zz} is the eddy diffusion coefficient (He_Kzz, default $10^9 \text{ cm}^2 \text{s}^{-1}$): it enters only the *gradient* term and drives f toward uniform (mixing), whereas the *settling* term $D f G$ is molecular only. Their competition sets the homopause: below it ($K_{\text{zz}} \gg D$) the atmosphere is mixed, above it ($D \gg K_{\text{zz}}$) helium settles. In diffusive equilibrium ($\Phi_{\text{He}} = 0$, no eddy) $\partial \ln f / \partial r = -G$, i.e. f decreases outward with the (small) He scale height $1/G$.

27.2 Conservative implicit discretization

An early *non-conservative* form (material-derivative advection $-v \partial f / \partial r$ combined with a conservative diffusive flux) proved flux-inconsistent and *drained* helium spuriously. The production scheme is instead **fully conservative**: it evolves the He number density n_{He} with a single divergence of the total face flux

$$J_{j+1/2} = \underbrace{n_{\text{He}} v}_{\text{advection}} + \Phi_{\text{He}}, \quad \frac{n_{\text{He}}^{n+1} - n_{\text{He}}^n}{\Delta t} = -\frac{1}{r^2 \Delta r} \left[r_{j+1/2}^2 J_{j+1/2} - r_{j-1/2}^2 J_{j-1/2} \right]. \quad (7)$$

Because the cold, dense base has a tiny He scale height, the diffusion is stiff and an explicit scheme is unstable; the update is therefore **implicit** (backward Euler), assembled as a tridiagonal system in n_{He}^{n+1} and solved by the Thomas algorithm. Implicit backward Euler is unconditionally stable, and its fixed point ($n_{\text{He}}^{n+1} = n_{\text{He}}^n$) is independent of Δt , so local cell-by-cell time steps are used for steady-state acceleration. The new element split is written back into `f_sp` conserving mass ($\sum_s m_s f_{\text{sp},s} = 1$), and the same kernel (`solve_1elem`) is reused for each trace metal (§27.6).

27.3 The $n_{\text{He}}-n_{\text{H}}$ mass coupling: why a self-consistent n_{H} is *not* needed

Diffusion redistributes species at *fixed* total mass density ρ (owned by the hydro), so in every cell

$$\rho = m_{\text{H}}n_{\text{H}} + m_{\text{He}}n_{\text{He}} + (\text{metals}) \implies n_{\text{H}} = \frac{\rho(1 - c_{\text{met}}) - m_{\text{He}}n_{\text{He}}}{m_{\text{H}}}. \quad (8)$$

Thus n_{H} is itself a function of n_{He} : where helium accumulates, hydrogen must be displaced. But the flux (6) *uses* n_{H} (in the coefficient and in $f = n_{\text{He}}/n_{\text{H}}$), so the solve is genuinely nonlinear in n_{He} . Two options:

- **Lagged n_{H} (adopted).** Hold n_{H} at its start-of-step value during the solve; update it afterward, in the mass-conserving rescale, for the next step.
- **Self-consistent n_{H} (“P2a”, tried, rejected).** Iterate within the step (Picard): solve n_{He} , reset n_{H} from (8), re-solve, until n_{He} and n_{H} are mutually consistent.

Why the Picard iteration diverges. Write the He mass fraction $Y = m_{\text{He}}n_{\text{He}}/\rho \in [0, 1]$. Near the base, settling genuinely concentrates helium ($n_{\text{He}} \uparrow$); enforcing (8) exactly then drives

$$n_{\text{H}} = \frac{\rho(1 - c_{\text{met}}) - m_{\text{He}}n_{\text{He}}}{m_{\text{H}}} \xrightarrow{Y \rightarrow 1} 0,$$

so $f = n_{\text{He}}/n_{\text{H}}$ and the $1/n_{\text{H}}$ flux coefficients diverge — a genuine singularity as helium approaches a pure-He layer. Because the concentration is *real* (settling against a fixed reservoir base), the iteration has a positive feedback and blows up (NaN). The lagged scheme avoids this: with n_{H} held at a finite value, $1/n_{\text{H}}$ stays bounded, and the physical cap $f \leq \text{HeH}$ (§27.4) prevents $Y \rightarrow 1$.

Why lagging is nonetheless harmless. Four independent arguments:

1. *The lag vanishes at the fixed point.* The diffusion runs inside a relaxation loop. At steady state $n_{\text{He}}^{n+1} = n_{\text{He}}^n$ and $n_{\text{H}}^{n+1} = n_{\text{H}}^n$, so the lagged $n_{\text{H}} = n_{\text{H}}^n$ used in the solve *equals* the final consistent n_{H}^{n+1} . The lag affects only the transient path, not the converged state: where the scheme is stable, lagged and self-consistent reach the *same* steady solution, and mass consistency is achieved automatically.
2. *The region where it matters is the (unphysical) pile-up.* Lagged and self-consistent differ appreciably only where $Y \rightarrow 1$, a thin layer just above the fixed reservoir base — below the He 10830 line-forming region, hence irrelevant to the observable, and in any case capped (§27.4).
3. *Helium is a minor mass fraction where the observable lives.* In the upper thermosphere, where He/H falls to $\sim 0.2-0.5 \times$ (the part that weakens 10830), $m_{\text{He}}n_{\text{He}} \ll m_{\text{H}}n_{\text{H}}$, so $n_{\text{H}} \approx \rho/m_{\text{H}}$ is nearly insensitive to n_{He} and the lag is negligible.

4. *Trace metals have no back-reaction.* Metals are $\sim 10^{-3}$ of the mass, so their diffusion does not change n_{H} measurably; for them the lagged n_{H} is effectively exact.

Hence the divergent Picard iteration is not merely inconvenient but *unnecessary*: the lagged- n_{H} solve plus the physical cap is stable and converges to the correct steady state.

27.4 The near-base pile-up and the physical cap

Settling transports helium down toward the reservoir base. Against a fixed-composition Dirichlet base it tends to pile up in the first cells above the base (numerically unbounded with a lagged n_{H}). In a settling-plus-escape column, however, He/H *cannot* exceed the reservoir value: there is no mechanism to enrich helium above the base abundance. The code therefore applies the physical limiter

$$f \leq \text{HeH} \quad (9)$$

(and, for each metal, $f_X \leq f_{X,\text{base}}$). This encodes the expectation that the reservoir *absorbs* settled helium rather than letting it build a pure-He layer — so the cap is a physical boundary condition, not merely a numerical band-aid. The aloft separation (the observable) sits well below HeH and is unaffected by the cap.

27.5 Settling discretization: a Péclet hybrid

The settling term DfG is an advective drift (velocity $w_s = DG$, directed inward). Central differencing of an advective term is stable and accurate only at low Péclet number, $|DG| \Delta r < 2(D + K_{zz})$. For light helium this always holds (small G), and central differencing is used (least numerical diffusion). For *heavy metals* (e.g. Fe, mass 56) G is $\sim 18\times$ larger; central differencing then loses diagonal dominance and the tridiagonal solve NaNs. The kernel therefore uses a **Péclet-based central/upwind hybrid**:

$$\text{settling drift: } \begin{cases} \text{central} & |DG| \Delta r < 2(D + K_{zz}) \quad (\text{light species}), \\ \text{first-order upwind} & \text{otherwise} \quad (\text{stiff heavy-metal settling}). \end{cases} \quad (10)$$

Upwind keeps the matrix diagonally dominant for any settling strength; the metal-stage rescale additionally skips cells where the element is negligible and clamps the scale factor to avoid $0 \times \infty$. These changes were required to run WASP-121b (Fe) stably; on HD 209458b (light C/N/O), the hybrid selects central and reproduces the He-only result.

27.6 Ambipolar field, thermal diffusion, and metals

In the ionized wind the polarization field $eE = -(1/n_e)dp_e/dr$ (Koskinen et al. 2013) lifts ions; for an H^+ background it supports $\sim$ half a proton weight per unit charge, so a species of charge Z has effective mass $m - Zm_{\text{H}}/2$. The He-vs-H relative settling mass becomes

$$\Delta m_{\text{eff}} = 3 - \frac{1}{2}(\bar{Z}_{\text{He}} - \bar{Z}_{\text{H}}), \quad (11)$$

with $\bar{Z}$ the mean charges from the local ionization state — $\Delta m_{\text{eff}} = 3$ at the neutral base, 2.5 in the fully ionized $\text{He}^{++}/\text{H}^+$ wind (He_ambipolar, default on). Thermal diffusion adds $\alpha_T \partial \ln T / \partial r$ to G (He_alphaT, default 0). Each trace metal (He_metal_diffusion, default off) is diffused independently against the post-rescale n_{H} with its own mass

(`melem_A`), binary D_X , and ambipolar-corrected Δm_{eff} ; being heavier than helium, metals deplete faster with altitude where diffusion competes with the wind.

27.7 The metal-rescale ratchet bug (found in review, fixed 2026-07-02)

The first version of the metal-stage rescale contained two bugs that together made metal depletion *irreversible*: (i) an $r_X \leq 1$ clamp — redundant for the metal/H $\leq$ reservoir cap (already enforced by $\min(n_X, f_{X,\text{base}} n_{\text{H}})$) but forbidding any *replenishment* of a previously depleted cell (a one-way ratchet); and (ii) a skip of cells where the element was negligible, which made exhausted cells permanent holes. During early relaxation (wind undeveloped) settling transiently depletes metals; the ratchet locked that in. **Fix:** r_X is no longer clamped above (the cap alone bounds it), and an exhausted cell that the solve replenishes is re-seeded through the neutral stage (the ionization equilibrium re-partitions the stages next step). **Retraction:** the previously reported WASP-121b “metal homopause” ($\text{Fe} \rightarrow 0$ by $\sim 1.25 R_p$, with a ~ 2500 K hotter thermosphere) was an artifact of this ratchet — with the fix, WASP-121b metals track He (even Fe is advection-dominated there, $w_s/v \sim 10^{-3}$) and the v1/v2 temperature structures agree. The HD 209458b mass ordering ($\text{He} > \text{C} > \text{N} > \text{O}$ aloft) is unaffected.

27.8 Boundary conditions and validation

The lower boundary is Dirichlet at the reservoir composition ($f = \text{HeH}$); the outer boundary carries a zero diffusive flux plus advective outflow (helium escapes with the wind). The two-planet before/after comparison (`docs/version_compare.{md,tex,pdf}`) shows the expected regimes: on HD 209458b helium itself separates ($\text{He} 10830 \text{ } 70.5\% \rightarrow 27.1\%$, $-2.6\times$) and heavier C/N/O deplete more than He aloft, whereas on WASP-121b ($\log \dot{M} \approx 13.3$) advection dominates for every species — helium *and* metals are dragged out essentially unfractiated (10830 unchanged; metal/H tracks He/H). Both are stable (no NaN) and leave $\dot{M}$ essentially unchanged (metals are trace). All diffusion flags are **default off**, so standard runs are byte-identical.

28. Lower-atmosphere connection: Tier 1–3 (2026-07-02)

Survey + proposal + implementation: `docs/lower_atmosphere_coupling.{md,tex,pdf}`. New folder `src/modules/lower_atmosphere` everything **opt-in / default off** (regression: the standard HD 209458 b 3000-step run reproduces v1.0’s $\log \dot{M} = 8.97$ exactly).

- **Tier 1** `lower_column.f90` + key `Lower column::` the Koskinen+2022 analytic isothermal- T_{eq} hypsometric column with Visscher chemical-equilibrium $\text{H}_2/\text{H}/\text{He}$; reports the derived 1- μbar base radius (equilibrium + fully-atomic bracket), base $q_{\text{H}_2}/q_{\text{H}}/q_{\text{He}}$ and μ next to the input “Planet radius”. Gate passed: Model A $r_0/R_{1\text{bar}} = 1.343$ (paper 1.34), $q_{\text{H}_2} = 0.838$ (0.84), $q_{\text{H}} = 0.027$ (0.026).
- **Tier 2 foundation** `h3p_cooling.f90` (Miller+2013 Table-5 LTE emission fits + Table-6 non-LTE factors $s(T, n_{\text{H}_2})$; anchors reproduced to $<0.5\%$) and `mol_rates.f90` (Koskinen+2022 Table-1 rates R1–R23, verified against the PDF). Remaining: the coupled `System_HeH_mol` solver + EOS.

- **Tier 2a** key Molecular base: True: EOS-only base correction removing the H₂-bound particles from `ntot_bc` via the equilibrium fit (chemistry stays atomic; crude, documented).
- **Tier 3** optional `base.inp` (`T_base/r_base/HeH_base/Kzz_base`; echoed, no-op when absent) consumed by `input_read`, + the driver `src/utls/run_lower.py` (isothermal or Guillot 2010 semi-grey $T(p)$); VULCAN photochemistry is the documented upgrade path.

Physics note from the Tier-1 gate: the chemical-equilibrium base stays strongly molecular up to $T \sim 2000$ K, so the atomic base of $T_{\text{eq}} \sim 1000\text{--}2000$ K hot Jupiters rests on *photochemical* dissociation (Moses 2011; Koskinen 2013a) — quantified motivation for the VULCAN tier.

29. Tier-2 molecular chemistry core: H₂/H₂⁺/H₃⁺/HeH⁺ (2026-07-03)

Opt-in key Molecular chemistry: True (requires He; v1 excludes trace metals). New coupled equilibrium system `System_HeH_mol.f90` (7–8 unknowns; atomic rows = EXHALE’s own rates, so the molecule-free limit reproduces the atomic systems; molecular channels = Koskinen+2022 Table-1 via `mol_rates`). σ_{H_2} (Yan+1998 Eqs. 17–19; verified against their Table 7) added to `cross_sec.f90` and wired into the opacity/photoionization/heating integrals; H₃⁺ IR cooling (Miller+2013 + non-LTE factor) enters `cool`; the EOS is molecule-aware (particle count, electrons, masses). `f_sp` grows 33→37 (output schema + Load-IC extended; columns zero when off). **Gates:** (0) molecules-off regression $\log \dot{M}=8.97$ unchanged; (1) HD 209458 b molecules-on: fully molecular base, sharp H₂ → H front at $r=1.019 R_p$, wind above it essentially atomic ($\log \dot{M}$ 8.91 vs 8.97, He 2³S peak +1.4%); (2) hot-Uranus-like: front at $1.148 R_p$, H₃⁺ active below. Caveats: local equilibrium only (no molecular advection), no Lyman–Werner photodissociation, P4/P5 and the 4.48 eV sink omitted, $\gamma=5/3$ retained. Details: `docs/lower_atmosphere_coupling.pdf` §4.

30. Newton–diffusion co-convergence, solver hygiene, and the Ly α lower-BC audit (2026-07-03)

Stall-based Newton hand-off. `He_diffusion` runs plateau the flux metric just above the hand-off threshold (observed: du frozen at 1.084×10^{-2} vs the 10^{-2} switch for 10^6 steps). The warm-up now tracks the du plateau and hands over to JFNK once stalled within $5 \times$ the switch. *Validated:* the hand-off engaged at $du = 2.2 \times 10^{-2}$; the JFNK then hit the known base-momentum holdout (worst cell $j=1$, $\|R\| \sim 7.7 \times 10^{-3}$ while the volume-weighted reference residual was 2.2×10^{-4}) and fell back gracefully — the base-cell viscosity work remains the enabler. **Newton + He_diffusion co-convergence.** Solver: Newton previously froze the diffused He/H field at the hand-off state (the JFNK residual carries no operator-split diffusion). An outer iteration (the excited-H pattern) now alternates the JFNK solve with 500 diffusion-relaxation steps at the converged wind until the He/H field drift $< 10^{-3}$ (max 5 passes, drift reported). **Riemann floor:** `p_min` floored at $10^{-30} p_{\text{max}}$ in `speed_estimate_HLLC/ROE` (no-op for healthy states). **PH_heat_HHe OMP CRITICAL removed** (distinct- j writes are race-free; values unchanged). **regression/ run_fcheck.sh:** periodic `-fcheck` regression per the 2026-07-02 review. **lya_rt audit** (vs. Huang et al. 2017): the escape-probability closure $\bar{J} = S(1 - \beta)$ is local and lacks Huang’s *absorbing bottom* (H₂ accidental-resonance sink at $N_{\text{H}_2} \sim 10^{14} \text{ cm}^{-2}$),

so $\bar{J}$ — and the $H(n=2)$ population — is likely overestimated in the bottom few scale heights; the $H\alpha$ -forming region lies above, so the TPM impact is expected limited. Documented; a bottom-loss β channel is the follow-up.

31. Tier-3 first light: VULCAN $\rightarrow$ base.inp (2026-07-03)

The public VULCAN photochemical kinetics code was cloned (`./VULCAN`; FastChem compiled) and run to steady state for HD 189733 b (SNCHO-2025 network, `use_photo=True`; 2356 steps, ~ 6 h wall). The new converter `src/utls/vulcan_to_base.py` turns the `.vul` output into an EXHALE `base.inp`: photochemical $q_{H_2}/q_H/q_{He}$ at $1\ \mu\text{bar}$, a hypsometric base radius using VULCAN's own $\mu(p)$ and $T(p)$, and the molecular mixing ratios as comments (*no* metal release — metal/alkali chemistry is outside VULCAN's scope and remains Lavvas-only). **Result** (HD 189733 b, $1\ \mu\text{bar}$): $q_{H_2} = 0.63$, $q_H = 0.23$ — photochemistry dissociates $\sim 11\times$ more H than the chemical-equilibrium column ($q_H = 0.020$), directly quantifying the Tier-1 caveat (Moses 2011; Koskinen 2013a); $r_{\text{base}} = 1.168$ vs the analytic $1.174 R_J$; $T_{\text{base}} = 863\ \text{K}$ (the Moses11 $T(p)$, cooler than the isothermal $T_{\text{eq}} = 1183\ \text{K}$). End-to-end chain demonstrated: VULCAN $\rightarrow$ converter $\rightarrow$ base.inp $\rightarrow$ EXHALE startup override (echoed).

32. VULCAN lower-atmosphere pre-step: obtained, not included in the tree (2026-07-06)

VULCAN (with FastChem inside it, as shipped) is used as a lower-atmosphere pre-step invoked by EXHALE itself. VULCAN is a third-party code and is *not* committed to the EXHALE repository (VULCAN/ is git-ignored, avoiding the embedded- repository problem); it is fetched into EXHALE/VULCAN/ by the new `src/utls/setup_vulcan.sh` (clone + two patches + FastChem build), which the driver also runs automatically on first use if the tree is missing. The key

```
Lower atmosphere: vulcan <R_1bar>
```

in `input.inp` runs the new driver `src/utls/vulcan_driver.py` at startup whenever no `base.inp` exists: it copies the VULCAN tree into `<run_dir>/vulcan_work/`, builds the planet's Guillot $T(p)/K_{zz}$ atmosphere and picks a stellar UV spectrum by host T_{eff} , compiles FastChem once, runs VULCAN to steady state (hours on the first run, cached `.vul` afterwards; `--force` to redo), converts the result to `base.inp`, and then proceeds with the wind solve — i.e. VULCAN acts as a subroutine of EXHALE. `Lower atmosphere: analytic <R_1bar>` invokes the fast chemical-equilibrium column (`run_lower.py`) instead; omitting the key disables the pre-step (fully optional). EXHALE_ROOT overrides the code-root path for relocated installs, and an existing `base.inp` always wins. Implementation: `parameters.f90` (`lower_atm_mode/lower_atm_r1bar`), `input_read.f90` (`run_lower_atm_prestep` via `execute_command_line` before `read_base_inp`), and the added `vulcan_driver.py`. Tested: analytic branch end-to-end, VULCAN branch with a cached `.vul` (HD 189733 b photochemical base $T = 863\ \text{K}$ applied), and the opt-out (no key) path. The download URLs, citation requirements (VULCAN: Tsai et al. 2017, 2021; FastChem: Stock et al. 2018, 2022) and the exact list of modifications `setup_vulcan.sh` applies are documented in the README “Obtaining VULCAN and FastChem” section and the user manual (§“Obtaining VULCAN and FastChem”).

33. TPM.py renamed to EXHALE_transit.py (2026-07-06)

The transmission-spectrum post-processor TPM.py (the acronym clashed with “Trusted Platform Module” and did not describe the tool) is renamed EXHALE_transit.py, matching the EXHALE_* convention. Run-time overrides are now EXHALE_TRANSIT_<NAME> (PATH, SAVE_PREFIX, HE_LMIN/LMAX/N, RSTAR_RSUN, ROTP, TSTAR); the old TPM_<NAME> names remain as backward-compatible fallbacks (a `_tenv()` helper checks the new name first, then the old). Output filenames are unchanged (<prefix>tpm_He10830.txt etc.) so downstream notebooks and figure scripts keep working. All in-repo references were updated; VULCAN/ untouched.

34. Sub-level-resolved Balmer opacity and LaRT 2s inclusion (2026-07-06)

The $H\alpha/H\beta$ transmission was computed as $f_{23}n_2$ with the multiplet oscillator strength ($f_{23} = 0.6407$ for $H\alpha$), which implicitly assumes the $2s:2p$ populations sit in the 1:3 statistical ratio. They do not—the metastable $2s$ ($A_{2s1s} = 8.26 \text{ s}^{-1}$) and the $\text{Ly}\alpha$ -pumped $2p$ depart strongly from it—and the two sub-levels have *different* Balmer cross sections (only $2s \rightarrow 3p$ is allowed from $2s$; $2p \rightarrow 3s, 3d$ from $2p$). Fixed to sum per sub-level, $\tau \sim f_{2s}n_{2s} + f_{2p}n_{2p}$, with $f_{2s} = 0.4349$, $f_{2p} = 0.70941$ ($H\alpha$) and 0.1028, 0.125886 ($H\beta$); the components share the Doppler-dominated profile.

EXHALE_transit.py (escape-prob path) now carries n_{2s}/n_{2p} separately; the effect is small ($\lesssim$ few %) since the in-code $J_{\text{Ly}\alpha}$ keeps $2s:2p$ near statistical in the forming region. examples/tpm_halpha_lart2d.py (LaRT path) previously used the $2p$ -only population, *discarding the 2s absorption entirely*; it now returns and sums both (the LaRT scattering rate pumps only $1s \rightarrow 2p$, entering the S_{2p} source term, while $2s$ is recombination/collisionally fed throughout the atmosphere). Effect: +21% (HD 189733 b) to +67% (WASP-121 b). The $\text{Ly}\alpha$ absorption line itself is unchanged and already correct— $1s \rightarrow 2p$ only ($f = 0.41641$ on the ground-state n_{HI} ; $1s \rightarrow 2s$ is a forbidden two-photon transition).

benchmarks/hd209/input.inp was missing Stellar Lya flux, Stellar Teff, and Stellar radius; without the $\text{Ly}\alpha$ flux the LaRT converter read $F_{\text{Ly}\alpha} = 0 \Rightarrow L_{\text{Ly}\alpha} = 0 \Rightarrow n_{2p} \simeq 0$, a spurious $H\alpha = 0.000\%$. Added $T_{\text{eff}} = 6065 \text{ K}$, $R_* = 1.155 R_{\odot}$, and $F_{\text{Ly}\alpha} = 3.5 \times 10^3 \text{ erg cm}^{-2} \text{ s}^{-1}$ (quiet-solar-analog, R_*^2 scaled; provisional, since the absolute LaRT depth scales with it). With both paths on the same VULCAN-BC winds and the sub-level opacity, the escape-prob/LaRT $H\alpha$ ratios become HD 209458 b 0.84/0.43 (2.0), HD 189733 b 1.23/0.91 (1.4), WASP-121 b 4.41/2.33 (1.9), WASP-52 b 5.69/2.20 (2.6)—so the inline closure overestimates $H\alpha$ by only $\sim 1.4\text{--}2.6\times$, far less than the earlier $\sim 4\text{--}30\times$ (pre-VULCAN winds with the $2p$ -only LaRT). Paper abstract, methods, results table, summary, and figs/make_lart_fig.py were updated accordingly.

35. Molecular equilibrium seeding fix, molecular-base coupling, Tier-2 gate refresh (2026-07-16)

Frozen-atomic-base bug (diagnosed and fixed). The 8-unknown molecular equilibrium (ion_system_HeH_mol, hybrd1) is bistable in its initial guess: from a zero-molecular warm start it fails to converge at the dense, optically-thick base, and the `thereis_mol` branch never checked info—the failed state was silently adopted and then frozen indefinitely ($x_{\text{H}_2} = 0.0122$ at the base, flat across 130k steps and independent of T ; numerically *not* a root of the

coded network—it would require an H_2 photoionization rate $15\times$ the unattenuated stellar bound). The true dark-base equilibrium is strongly molecular (3-body formation R15 vs. thermal dissociation R12), which the earlier gate runs had reached only because the then-uninitialized molecular `f_sp` columns acted as an accidental nonzero seed; the composition-audit zero-init exposed the latent fragility. Fixes: (i) `info` is checked and a failed solve retries once from a physically-informed guess seeded by the chemical-equilibrium H_2 fit `q_h2_equilibrium(p,T)` (still-failing cells keep the previous state and are counted in a one-line summary warning); (ii) `set_IC` seeds the molecular H_2 column from the same fit (H nuclei conserved) instead of zero.

Molecular base coupled on. Solving molecular chemistry now implies the molecular base particle count (Molecular base: True behavior): with an atomic `ntot_bc` the base temperature is inflated by the particle-count ratio ($\sim 1.8\times$ for a fully molecular base), thermally dissociating the H_2 layer the chemistry just built. A note is printed when the coupling engages.

Tier-2 gates re-established (deterministic; inputs pinned at `docs/lower_atmosphere_figs/data_g*/input.inp`). Gate 0 (HD 209458 b atomic, 3000 steps): $\log \dot{M} = 8.97$ unchanged. Gate 1 (HD 209458 b molecular): fully molecular base ($x_{\text{H}_2} = 0.997$, base T unchanged at ~ 1464 K), sharp $\text{H}_2 \rightarrow \text{H}$ front at $r = 1.019 R_p$, $\log \dot{M} = 8.87$ vs. 8.97 atomic, He 2^3S peak +1.3%; stable over 12000 steps with zero solve failures. Gate 2 (hot-Uranus-like: $0.0457 M_J$, $R_p = 0.49 R_J$, $T_{\text{eq}} = 1140$ K, HD 209458 b orbit/spectrum—input recovered by matching the archived domain fingerprint $r_{\text{max}} = 4.814118$ to 6 decimals): fully molecular base, front at $r = 1.166 R_p$, H_3^+ active in the molecular layer. Atomic runs byte-identical throughout (make check + OMP identity).

36. Default H/He rates $\rightarrow$ Badnell+Mao / Voronov, van Hoof free-free Gaunt factor, free-free charge fix, and secondary ionization (2026-07-17)

Six coupled changes to the H/He microphysics. Items 1–2 are gated by the new `legacy_hhe_rates` flag (default False = new rates); items 3–4 (free-free) apply always; item 5 (secondary ionization) is gated by `use_sec_ion` (default True); item 6 rides along with the heating rework. Full cross-code context and the benchmark table: `docs/atomic_data_EXHALE_vs_MoCHII`.

36.1 What changed and why

1. H/He recombination default. The pre-update H/He case-B coefficients were unattributed fits inherited from ATEs; a numerical check showed them to be exactly the Hui & Gnedin (1997) case-B fits. The new default is

$$\alpha^B(T) = \alpha_{\text{Badnell}}^A(T) - \alpha_{\text{I}}^{\text{Mao}}(T), \quad (12)$$

the Badnell (2023) radiative total minus the Mao & Kaastra (2016, A&A 587, A84) $n = 1$ coefficient, with the Badnell three-term dielectronic recombination added to α^A for He II. This matches the construction MoCHII already uses, so the two codes now share the H/He recombination coefficients. Setting `legacy_hhe_rates = True` restores the Hui & Gnedin fits exactly.

2. H/He collisional ionization default. The legacy HI / He I rates were Abel et al. (1997) (reproducing Janev et al. 1987) and He II a Hui & Gnedin (1997) fit, none source-annotated. The Voronov (1997, ADNDT 65, 1) fits,

$$k(T) = A \frac{(1 + P\sqrt{U})U^K e^{-U}}{X + U}, \quad U = \Delta E[\text{eV}]/k_B T, \quad (13)$$

are now the default, matching the metal rates (already Voronov in both codes) and MoCHII. Reverted together with item 1 by `legacy_hhe_rates`.

3. Free-free Gaunt factor. The frequency-averaged non-relativistic $\langle g_{\text{ff}} \rangle(\gamma^2)$ (with $\gamma^2 = Z_{\text{ion}}^2 \text{Ry}/k_B T$) of van Hoof et al. (2014, MNRAS 444, 420) is tabulated at 161 points in $\log_{10} \gamma^2$ from -6 to 10 in 0.1 dex steps (generator `cooling_data/gauntff_thermal_avg.py`), replacing the earlier unattributed two-branch logarithmic fit. Applies always, independent of `legacy_hhe_rates`.

4. Free-free charge-weighting correction (physical bug). The previous code weighted He II (net charge $+1$) by $Z^2 = 4$ and evaluated the metal Gaunt factor at the nuclear atomic number ($Z = 8$ for O, 26 for Fe). Free-free emission scales with the ion *net* charge, not the nuclear number: $Z_{\text{ion}} = 1$ for H II, He II, and singly-ionized metals; $Z_{\text{ion}} = 2$ for He III and doubly-ionized metals. The most visible consequence is that the free-free cooling of a He II-dominated layer drops by a factor of four. Applies always.

5. Secondary ionization (Shull & van Steenberg 1985) in the main loop. A photoelectron ejected with $E_0 = e_\nu - E_{\text{th}}$ above ~ 40 eV does not thermalize immediately: it deposits only $f_{\text{heat}}(x)$ of its excess as heat and drives secondary HI and He I ionizations, where x is the ionized fraction of the H+He nuclei,

$$f_{\text{heat}}(x) = 0.9971 (1 - (1 - x^{0.2663})^{1.3163}), \quad (14)$$

$$f_{\text{ion,HI}}(x) = 0.3908 (1 - x^{0.4092})^{1.7592}, \quad (15)$$

$$f_{\text{ion,HeI}}(x) = 0.0554 (1 - x^{0.4614})^{1.6660}. \quad (16)$$

Below 40 eV the photoelectron thermalizes fully. Gated by `use_sec_ion` (default `True`); `use_sec_ion = False` is bit-identical to the legacy full-thermalization path. This differs from `wind_ae`, which combines the same SvS85 partition with Dere (2007) tables tied to its fixed spectral grid; the main loop runs on a runtime SED and applies the SvS85 species-resolved fits directly.

6. He I 2^3S triplet photoheating channel added. The metastable He I 2^3S (threshold 4.8 eV) previously contributed only opacity and its photoionization rate but no photoheating. Its photoelectron energy now enters the heating and absorbed-energy budgets and the secondary-ionization source consistently with every other absorber.

36.2 Files touched

`Cool_coeff.f90` (`rr_badnell`, `dr_HeII_badnell`, `rr_mao`, `alphaB*_new`, `voronov_ci/ci*_new`, the 161-point `gff_avg` table); `util_ion_eq.f90` (free-free net-charge weighting in `eval_cool`; secondary ionization and He 2^3S heating in `PH_heat_H/PH_heat_HHe`); `T_equation.f90` (free-free net-charge weighting); `ionization_equilibrium.f90`

and `post_process_adv.f90` (callers wired for the new signatures); `input_read.f90`, `parameters.f90`, `write_setup_report.f90` (legacy_hhe_rates and use_sec_ion declared, defaulted, parsed, echoed); `cooling_data/gauntff_thermal_avg.py` (van Hoof table generator); `docs/input_schema.md` (keys K14b, K14c).

36.3 Quantitative check

New default rates (α_B , collisional ionization ci ; $\text{cm}^3 \text{s}^{-1}$), with the legacy value at 10^4 K for comparison. The legacy $\alpha_B(\text{HII})$ at 10^4 K is the Hui & Gnedin (1997) fit, 2.592×10^{-13} .

Coefficient	10^4 K	10^5 K	10^6 K	legacy (10^4 K)
$\alpha_B(\text{HII})$	2.606e-13	3.070e-14	2.163e-15	2.592e-13
$\alpha_B(\text{He II})$	2.807e-13	5.006e-13	1.015e-12	2.616e-13
$\alpha_B(\text{He III})$	1.536e-12	2.370e-13	2.244e-14	—
$ci(\text{HI})$	7.448e-16	3.963e-9	3.103e-8	7.247e-16

The Fortran reproduces an independent Python recomputation to all printed digits, and the legacy branch reproduces the pre-update values exactly. The new $\alpha_B(\text{HII})$ at 10^4 K differs from the legacy Hui & Gnedin value by $\sim 1\%$; He II and He III differ more, mainly because the new He II value carries the Badnell dielectronic contribution explicitly. The van Hoof thermal Gaunt factor at $T = 10^4 \text{ K}$ and $Z = 1$ is $\langle g_{\text{ff}} \rangle = 1.266$. A smoke run confirms the expected factor-of-four drop in free-free cooling across a He II-dominated layer, and `use_sec_ion = False` reproduces the legacy heating and photoionization rates bit-identically.

36.4 Effect on existing results

Because the default rates change, results computed with the new default differ from the pre-update golden and gate outputs. `legacy_hhe_rates = True` recovers the old H/He ionization balance, with two always-on exceptions: the free-free charge-weighting correction (item 4) and the van Hoof Gaunt table (item 3), both physical corrections intentionally not tied to the switch. Secondary ionization (item 5) is a further default change; `use_sec_ion = False` restores the full-thermalization rates bit-identically.

37. He recombination radiation $\rightarrow$ H I ionization coupling (use_he_rec_coupling, 2026-07-17)

EXHALE carries no diffuse radiation field, so pure case B implicitly assumes every $\text{He II} \rightarrow \text{He I}$ recombination photon is reabsorbed locally by helium ($\alpha_B = \alpha_A - \alpha_1$, ground capture dropped). In H-dominated gas that is wrong: part of the He recombination radiation — the $\geq 24.6 \text{ eV}$ ground-capture continuum, the $584 \text{ \AA}$ resonance line, the $19.8 \text{ eV } 2^3\text{S}$ line, and the $> 13.6 \text{ eV}$ part of the He I two-photon continuum — ionizes H instead. The physics and the choice of recipe were worked out in `docs/QUESTIONS_2026-07-17.md`. This section documents an opt-in coupling using Draine’s (2011) on-the-spot y/z parametrization. New input key `He_rec_coupling` (flag `use_he_rec_coupling`), default `False` $\rightarrow$ pure case B, bit-identical to the legacy path.

37.1 What changed and why

When on, the He recombination photons that ionize H are added as an extra H I photoionization rate with its photoelectron heating, and the He II recombination coefficient becomes $\alpha_B + y\alpha_1$ rather than pure case B. Two local parameters set the split:

- $y = 1/(1 + R n_{\text{HeI}}/n_{\text{HI}})$ (Draine Eq. 14.16), the fraction of the ≥ 24.6 eV ground-capture continuum absorbed by H, with the cross-section ratio $R = \sigma_{\text{He}}/\sigma_{\text{H}}$ at the 24.6 eV He I ground edge. Evaluated from the code cross sections, $R = 6.004$.
- $z = 0.67 + 0.29/(1 + n_e/n_{\text{crit}})$ (Draine Sec. 15.5), the density-dependent fraction of case-B (excited-state) cascade photons that ionize H, with the 2^3S critical density $n_{\text{crit}} = 1100 e^{1.2/T_4} T_4^{1/2} \text{ cm}^{-3}$; it interpolates between Draine's two limits (~ 0.96 low density $\rightarrow \sim 0.67$ high density).

Atomic (case-B) mode. No explicit 2^3S : α_1 is the Mao & Kaastra (2016) ground ($1s^2$) capture and α_B the active He II case B. The He II recombination coefficient is $\alpha_{\text{eff}} = \alpha_B + y\alpha_1$; the extra H I photoionization rate is $n_{\text{HeII}}n_e[z\alpha_B + y\alpha_1]/n_{\text{HI}}$; the photoelectron heating uses the ground-channel energy $E_{\text{gnd}} = 24.6 - 13.6 = 11.0$ eV and a cascade-averaged $E_{\text{casc}} \approx 6.3$ eV.

He I 2^3S (triplet) mode. The mode that computes 10830: the decay channels are explicit, so z is not used — a 2^3S atom destroyed by photoionization or Penning emits no 19.8 eV photon, so the cascade must be summed at the actual channel rates (α_1 is the 1^1S channel). The 1^1S channel coefficient is $y\alpha_1 + 0.25\alpha_B$: the net ground capture plus the singlet-excited capture channel ($0.25\alpha_B$) the current network omits. The H-ionizing photon production is summed over channels: ground $y\alpha_1 n_{\text{HeII}}n_e$; singlet-excited $0.85 \times 0.25\alpha_B n_{\text{HeII}}n_e$ ($0.85 = \frac{2}{3}$ from the 584 Å line + $\frac{1}{3} \times 0.56$ from the 2^1S two-photon continuum above 13.6 eV); 2^3S radiative decay $A_{31}n(2^3\text{S})$ (19.8 eV line, always ionizes H); and 2^3S collisionally converted to the singlets $n_en(2^3\text{S})(0.56q_{31a} + q_{31b})$. The channel deposit energies are 11.0 (ground), 5.6 (singlet-excited), 6.2 (19.8 eV line), and 3.0/7.6 eV (collisionally converted two-photon / 584 Å).

The corrections are applied at the lagged (pre-solve) densities in the main ionization-equilibrium loop and re-evaluated at the advection-corrected densities in the post-processor (rate/coefficient before the advection solve, heating after). The He II recombination *cooling* is left at $kT\alpha_B$; the mismatch is at most $y\alpha_1 kT$, negligible.

37.2 Files touched

util_ion_eq.f90 (new he_rec_coupling subroutine returning the corrected He II recombination coefficient, the extra H I photoionization rate, and the photoelectron heating; all zero when off); Cool_coeff.f90 (new alpha1_HeII_mao, the Mao & Kaastra 2016 ground $1s^2$ capture, i.e. the α_1 subtracted in alphaB_HeII_new); ionization_equilibrium.f90 and post_process_adv.f90 (callers wired to rcheiIB/P_HI/heat in the main loop and to the advection ODE and theat in post-process); input_read.f90, parameters.f90, write_setup_report.f90 (use_he_rec_coupling declared, defaulted .false., parsed, echoed); docs/input_schema.md (key K14d).

37.3 Quantitative check (HD 209458 b)

Protocol: HD 209458 b `input.inp+metals.inp`, the converged profiles loaded as IC, then four runs $\{\text{atomic}, \text{He } 2^3\text{S}\} \times \{\text{off}, \text{on}\}$ each relaxed for 15000 steps from the same IC. The IC predates the 2026-07-17 rate update, so both branches share the same partial re-relaxation and **only the on/off difference is meaningful** (not the absolute values). Across the wind $y \sim 0.65\text{--}0.69$ and $R = \sigma_{\text{He}}/\sigma_{\text{H}}(24.6\text{eV}) = 6.004$. Script `docs/plot_he_rec_coupling.py`; figures `docs/figures/he_rec_coupling_{atomic,heitr}.pdf`.

Quantity (on/off)	Atomic mode	He I 2^3S mode
He I	$\times 1.04$ ($1.2 R_p$) $\rightarrow \times 1.25$ plateau ($1.5\text{--}3.5 R_p$) $\rightarrow \times 1.32$ ($3 R_p$)	$\times 1.01\text{--}1.16$
He II (min.)	$\times 0.72\text{--}0.74$ ($1.1\text{--}1.2 R_p$) $\rightarrow \sim 1.0$ ($2 R_p$)	$\times 0.885$
$n(2^3\text{S})$	— (not computed)	$\times 0.89$ (1.05), $\times 0.92$ (1.2), $\times 0.98\text{--}1.01$ ($1.5\text{--}2 R_p$)
Temperature	$\leq +1\%$ (outer $+3\%$)	$\leq +0.3\%$
$\log \dot{M}$	$9.58 \rightarrow 9.65$ ($+0.07$ dex, $\sim +17\%$)	$9.57 \rightarrow 9.61$ ($+0.04$ dex, $\sim +10\%$)

The atomic mode carries the larger He ionization-structure change (case B alone, $y = 0$, under-recombines He), and its extra H photoionization raises the mass-loss rate by $\sim 17\%$. In the 2^3S mode the 10830-forming region ($1.5\text{--}2 R_p$) sees $n(2^3\text{S})$ within $\sim 1\text{--}2\%$ of the off case, so the He I 10830 line is essentially unchanged, while the base 2^3S drops $\sim 11\%$ (the TR-mode 1^1S channel coefficient change feeds directly into n_{HeII}). The temperature response is small in both modes. Default off is bit-identical to the pre-update path; existing goldens and gates are unaffected.




Figure 1: He recombination radiation $\rightarrow$ HI coupling, **atomic mode** (HD 209458 b). Clockwise from top left: He ionization fractions off vs. on; the on/off ratio for He I and He II; the photoheating rate; and the temperature. Turning the coupling on raises He I and lowers the He II minimum, with the $\sim 17\%$ mass-loss increase reflected in the extra photoheating.




Figure 2: Same comparison in the **He I 2^3S mode**. The extra panel and ratio curve track the metastable 2^3S density: it drops $\sim 11\%$ at the base but stays within $\sim 1\text{--}2\%$ of the off case through the 10830-forming region ($1.5\text{--}2 R_p$), so the He I 10830 line is essentially unchanged.

A Code-size summary: EXHALE vs. the original ATES (2026-06-07; updated 2026-07-02)

A line-by-line comparison of the Fortran source against the pristine upstream code (ATES/ATES-Code-main) quantifies the extent of the in-house extension. *Pure code lines (SLOC)* exclude blank lines and comments (the comment stripper is string-aware, so a `!` inside a quoted string is not treated as a comment); the line-by-line diff additionally normalizes whitespace and case, so reformatting and comment-only edits are *not* counted as changes. Updated 2026-07-02 after the Newton steady solver, the He 2^3S +metals merge, the composition module, He/H+metal diffusive separation, and the Wind-AE IC port; the *core* rows exclude the ported Wind-AE solver (listed separately).

	Fortran files	code lines (SLOC)
Original ATES (ATES-Code-main)	47	3 514
EXHALE core (excl. Wind-AE port)	62	9 245
net change (core)	+15	+5 731 ($\approx 2.6\times$)
Wind-AE IC-generator port (wind_ae/)	+36	+4 746
EXHALE total	98	13 991 ($\approx 4.0\times$)

Line-by-line diff, core (comments + blanks excluded; whitespace/case normalized): **5 970** code lines added in EXHALE (new modules + the new side of modified lines) and **239** removed from the original (deleted + the old side of modified lines), for **6 209** total changed lines. At the file level, of the 47 original files **27 were modified** and **20 are unchanged** (none were removed — the unchanged set is chiefly the 11 included MINPACK routines, the HLLC/ROE Riemann solvers with the flux/reconstruction pieces, and small utility modules), and **15 new core modules** were added:

New module	Purpose	SLOC
time_step/steady_newton.f90	JFNK steady-state solve (Newton finish)	540
radiation/charge_exchange.f90	charge-exchange network (Huang Table 4)	314
radiation/excited_hydrogen.f90	H($n=2$) model + Balmer feedback	305
functions/species_diffusion.f90	He/H (+ metal) diffusive separation	216
radiation/opacity_models.f90	opacity dispatcher (A/C/P/T)	180
nonlinear_system_solver/System_HeH_metals.f90	coupled H/He+metals ionization (+analytic Jacobian)	171
nonlinear_system_solver/newton_solver.f90	analytic-Jacobian Newton (+hybrd1 fallback)	124
files_IO/metals_input_read.f90	parse metals.inp	111
time_step/energy_semi_implicit.f90	semi-implicit energy solver	106
init/species_table.f90	metal-species metadata table	89
files_IO/opacity_input_read.f90	parse opacity.inp	85
nonlinear_system_solver/System_HeH_TR_metals.f90	merged He 2^3S +metals system	84
radiation/lya_rt.f90	Ly α escape-probability RT	68
functions/composition.f90	single-point species/EOS policy	46
time_step/steady_residual.f90	steady-residual assembly/norms	45
new-module total		2 484

So the 15 new core modules contribute 2 484 of the 5 970 added code lines; the remaining $\sim 3\,486$ added (and 239 removed) lines are edits spread across the 27 modified files — dominated by the cooling module (`Cool_coeff.f90`, 442 $\rightarrow$ 2 230 raw lines: CHIANTI metal-line fits, Fe tables, temperature-dependent Penning rate), followed by `EXHALE_transit.py`-side physics and the main driver, parameters/input parsing, IC construction, cross sections, ionization equilibrium, and the post-processor, which wire the metals, opacity dispatcher, Roche potential, transonic/auto IC, excited-H/Ly α physics, diffusion hooks, and the Brent/Newton solver upgrades into the existing solver. In total the EXHALE core is now $\approx 2.6\times$ the original code base (3 514 $\rightarrow$ 9 245 SLOC), touching a bit over half of the original files (27 of 47 modified, none deleted, plus 15 new); with the ported Wind-AE IC generator the full Fortran tree is 98 files / 13 991 SLOC ($\approx 4\times$). Beyond Fortran, the tree also carries the Python tooling (GUI, the `EXHALE_transit.py` transmission post-processor, plotting/analysis) and $\approx 20\,000$ lines of documentation under `docs/`.

B Fortran source inherited unchanged from ATES

Of the 47 upstream Fortran files (`ATES/ATES-Code-main`), **18 are byte-for-byte identical** in EXHALE — the numerics EXHALE reuses without touching. They fall into three groups: the included MINPACK nonlinear solver (the hybrid Powell hybrid method and its linear-algebra helpers, called by the ionization-equilibrium solver), the finite-volume hydrodynamics core (HLLC fluxes, reconstruction, source terms, Runge–Kutta right-hand side), and two radiation/ionization pieces. Table 2 lists them with their role. Two further files — the HLLC and ROE wave-speed estimators (`flux/speed_estimate_HLLC.f90`, `speed_estimate_ROE.f90`) — are *effectively* untouched: their original logic is preserved verbatim and only a one-line defensive floor `p_min=max(p_min, 1e-30 p_max)` (plus its comment) was appended in the 2026-07-02 code review, so no original line was altered.

Table 2: The 18 Fortran files byte-identical between the original ATES (ATES-Code-main) and EXHALE, grouped by role.

File	Purpose
<i>MINPACK nonlinear-system solver</i> (Argonne; dir nonlinear_system_solver/)	
hybrd.f90	Powell hybrid method: find a zero of n nonlinear equations in n unknowns
hybrd1.f90	Simplified user-facing driver for hybrd
fdjac1.f90	Forward-difference approximation of the Jacobian
qrfac.f90	QR factorization via Householder transformations
qform.f90	Accumulate the orthogonal factor Q from the QR factorization
r1updt.f90	Rank-one update of the QR factorization
r1mpyq.f90	Apply the Givens-rotation sequence of the rank-one update
dogleg.f90	Dogleg (trust-region) step combining Gauss–Newton and steepest descent
enorm.f90	Euclidean norm of a vector (overflow/underflow safe)
dpmpar.f90	Machine floating-point precision / range constants
<i>Hydrodynamics & grid core</i> (finite-volume Godunov)	
flux/Num_Fluxes.f90	Numerical interface flux assembly (HLLC Riemann fluxes)
functions/UW_conversions.f90	Conserved $\leftrightarrow$ primitive variable conversions
states/PLM_rec.f90	Piecewise-linear (2nd-order) cell reconstruction
states/Source.f90	Gravitational + geometric source terms
time_step/RK_rhs.f90	Runge–Kutta right-hand side (convection + source)
init/set_gravity_grid.f90	Pre-evaluate gravity at cell centers and interfaces
<i>Radiation & ionization</i>	
radiation/J_inc.f90	Energy-dependent incident stellar spectrum for photoionization
nonlinear_system_solver/System_implicit_adv_H.f90	H-only advection-coupled implicit ionization system (He/metal work went into the HeH/HeH_TR variants, leaving this untouched)