

EXHALE: Update Log

***Purpose.** This document is a running log of incremental updates to EXHALE as the code is extended toward reproducing Huang et al. (2023) WASP-121b (see Huang_update_plan.md). Each update is a self-contained section: what changed, why, which files were touched, and the quantitative check that was run. New updates are appended as new top-level sections.*

Contents

1. Runtime domain-extent option: Roche vs Spherical	6
1.1 Overview	6
1.2 The two domain modes	6
1.2.1 Gravitational potential	6
1.2.2 Outer boundary	7
1.2.3 Why spherical mode must drop the tidal terms	7
1.3 Input interface and backward compatibility	7
1.4 Implementation	8
1.5 Validation on WASP-121b	8
1.6 Status and relation to the Huang plan	9
2. Ionization network: add magnesium (Mg I / II / III)	9
2.1 Overview	9
2.2 Atomic data and rates	9
2.3 Touched files	10
2.4 Validation on WASP-121b	11
2.5 Status and relation to the Huang plan	12
3. Phase 1a: interface / data-structure refactor (no new physics)	12
3.1 Overview	12
3.2 What changed (Steps A–F)	12
3.3 Gate: bit-for-bit reproduction of the C/N/O + Mg run	13
3.4 What is deliberately <i>not</i> generalized yet	13
3.5 Touched files	13
3.6 Status and relation to the Huang plan	14
4. Phase 1b: correct C/N/O cross sections and add Si, Ca, Na, K, S	14
4.1 Overview	14
4.2 Variable-stage solver (2- and 3-stage elements in one packing)	14
4.3 Atomic data and rates (sources, not fabricated)	15
4.4 Rename: System_HeHCO → System_HeH_metals	15

4.5	Validation on WASP-121b	15
4.6	Caveats and gaps	16
4.7	Regression reference for Phase 1c	16
4.8	Status and relation to the Huang plan	17
5.	Phase 1c: add iron (Fe I / II / III)	17
5.1	Overview	17
5.2	Atomic data and rates (sources, not fabricated)	17
5.3	Touched files	18
5.4	Iron abundance	18
5.5	Validation on WASP-121b	18
5.6	Caveats and gaps	19
5.7	Regression reference for Phase 1d / Phase 2	20
5.8	Status and relation to the Huang plan	20
6.	Phase 1d: charge exchange with hydrogen (Huang Table 4)	20
6.1	Overview	20
6.2	Atomic data and rates (sources, not fabricated)	20
6.3	Touched files	21
6.4	Validation on WASP-121b	21
6.5	Caveats and gaps	22
6.6	Regression reference for Phase 2	23
6.7	Status and relation to the Huang plan	23
7.	Phase 2: metal line cooling — validation gate closed	23
7.1	Active line coolants	23
7.2	Channel-by-channel cooling diagnostic (this phase)	23
7.3	Validation on WASP-121b (the gate)	24
7.4	Touched files	24
7.5	Caveats and deferred	24
7.6	Status and relation to the Huang plan	25
8.	Phase 3a: in-code excited H(n=2) + Balmer proton source and heating	25
8.1	Physics and atomic data (sources, not fabricated)	25
8.2	Coupling into the solver	25
8.3	Runtime interface	26
8.4	Diagnostics	26
8.5	Touched files	26
8.6	Validation on WASP-121b (coupling ON vs OFF)	27
8.7	Caveats and gaps	27

8.8	Two analysis gotchas (recorded so they are not repeated)	28
8.9	Status and relation to the Huang plan	28
9.	Phase 2 refinement: Fe I line cooling (NIST f-values + Van Regemorter)	28
9.1	Atomic data and method (sources, not fabricated)	28
9.2	Touched files	29
9.3	Validation on WASP-121b	29
9.4	Status and relation to the Huang plan	30
10.	Phase 2 refinement: density-dependent Fe II line cooling (SE table)	30
10.1	Why the coronal table is wrong at the base	30
10.2	The density-dependent coefficient (multilevel statistical equilibrium)	30
10.3	Step 1 — the 2-D table generator (done)	31
10.4	Step 2 — the 2-D table and bilinear interpolator in Cool_coeff.f90 (done)	31
10.5	Step 3 — wiring the density-dependent Fe II into the cooling assembly (done)	32
10.6	Step 4 — WASP-121b run and the cooling budget (numbers for the reader)	32
11.	Post-process refinement: re-solve metal ionization in the advected wind (pp_metals 2)	33
11.1	Implementation — reuse the coupled residual with H/He pinned	34
11.2	The cell-by-cell sweep in post_process_adv	35
11.3	Clean A/B and what changes	35
12.	Phase 3b: in-code Lyα radiative transfer ($\bar{J}_{\text{Ly}\alpha}$ via escape probability)	36
12.1	The escape-probability field (lya_rt.f90 : jlya_escape_prob)	37
12.2	WASP-121b validation setup (spherical Case A, matched to Huang)	38
12.3	Trend vs Huang Fig. 11 (qualitative — exact match not expected)	38
13.	Phase 4: RLOF / tidal potential (Case B validated; Case D blocker = escape radius, transonic IC added, geometry-limited)	38
13.1	Base gravity for each case (Huang Table 3)	39
13.2	Case B — validated	39
13.3	Case D — real blocker found (escape radius), transonic IC added	39
14.	Phase 5: velocity-broadened transmission spectrum (in progress)	41
14.1	5a — metal resonance lines (Mg II, Ca II, Na I D), spherical — done	41
14.2	5b — Roche-equipotential reconstruction + broadening — implemented	42
14.3	5b — direct Case A comparison: hydro validated, metals over-extended	42
15.	Post-process advection guard at a “breathing” base (2026-06-07)	43
16.	He 2³S triplet + metals simultaneous solve, Lyα-flux guard, and convergence workflow (2026-06-08)	44

17. Metals in the bulk-gas mass/charge/particle budget (eos_metals, 2026-06-13)	47
18. Automatic initial-condition selection (IC mode: auto, 2026-06-13)	48
19. Base boundary conditions and the flux-based convergence decision (2026-06-15)	49
20. Tidal (Roche) potential with an extended outer boundary (2026-06-16)	51
21. Photoionization cross-section benchmarks and a TOPbase high-energy extension of He 2 ³ S (2026-07-01)	51
22. He I (1 ¹ S) photoionization → Verner+1996 by default, with ATES-rate switches (2026-07-01)	53
23. He 2 ³ S photoionization → two VFXY96 (Verner) forms + power-law bridge (2026-07-01)	54
24. Temperature-dependent Penning ionization of He 2 ³ S (2026-07-01)	54
25. He/H diffusive separation — Phase 1 (2026-07-01, experimental, default OFF)	55
26. He/H diffusion — Phase 2 (2026-07-02): ambipolar, thermal diffusion, metals	56
27. Diffusive separation: detailed physics and numerics	57
27.1 Governing equation	57
27.2 Conservative implicit discretization	57
27.3 The $n_{\text{He}}-n_{\text{H}}$ mass coupling: why a self-consistent n_{H} is <i>not</i> needed	58
27.4 The near-base pile-up and the physical cap	59
27.5 Settling discretization: a Péclet hybrid	59
27.6 Ambipolar field, thermal diffusion, and metals	59
27.7 The metal-rescale ratchet bug (found in review, fixed 2026-07-02)	60
27.8 Boundary conditions and validation	60
28. Lower-atmosphere connection: Tier 1–3 (2026-07-02)	60
29. Tier-2 molecular chemistry core: H ₂ /H ₂ ⁺ /H ₃ ⁺ /HeH ⁺ (2026-07-03)	61
30. Newton–diffusion co-convergence, solver hygiene, and the Ly α lower-BC audit (2026-07-03)	61
31. Tier-3 first light: VULCAN → base.inp (2026-07-03)	62
32. VULCAN lower-atmosphere pre-step: obtained, not included in the tree (2026-07-06)	62
33. TPM.py renamed to EXHALE_transit.py (2026-07-06)	63
34. Sub-level-resolved Balmer opacity and LaRT 2s inclusion (2026-07-06)	63
35. Molecular equilibrium seeding fix, molecular-base coupling, Tier-2 gate refresh (2026-07-16)	63

36. Default H/He rates → Badnell+Mao / Voronov, van Hoof free-free Gaunt factor, free-free charge fix, and secondary ionization (2026-07-17)	64
36.1 What changed and why	65
36.2 Files touched	66
36.3 Quantitative check	66
36.4 Effect on existing results	66
37. He recombination radiation → H I ionization coupling (use_he_rec_coupling, 2026-07-17)	67
37.1 What changed and why	67
37.2 Files touched	68
37.3 Quantitative check (HD 209458 b)	68
38. Staged activation of secondary ionization (2026-07-23)	71
38.1 The regression	71
38.2 The fix: converge first without the coupling, then switch it on	71
38.3 The Immediate override	71
38.4 Files touched	71
38.5 Effect on existing results (wasp_full validation)	72
39. He-triplet physics series from the Falorca & Vidotto (2026) review (2026-07-23)	72
40. du triggers require a descending crossing (2026-08-11)	73
41. JFNK line search decides on the true residual (2026-08-11)	73
42. Metal-line cooling at a cold base: line trapping and the coronal-fit validity floor (2026-08-11)	74
42.1 (a) The escape probability was hardwired to 1	74
42.2 (b) The CHIANTI coronal fits were evaluated far below their validity floor	74
42.3 Measured	75
42.4 Not fixed by this	75
43. Base grid resolution as an input key (Base_grid [dr,cells], 2026-08-11)	76
44. Base ghost temperature, runtime step cap, coronal-cutoff width as input keys (2026-08-11)	77
44.1 Base ghost temperature: isothermal continuous	77
44.2 Max steps: <N>	78
44.3 Coronal cutoff width: <w>	79
45. Ionization-equilibrium roots validated against the physical simplex (2026-08-11)	79
46. Line trapping in the thick metal resonance lines: measured, no change (2026-08-11)	80
47. Metal electrons and a validity range for the post-process advection correction (2026-08-12)	81

48. Ground-term fine-structure statistical equilibrium for C I, C II, N II and O I (2026-08-12)	83
49. Base ghost built from the composition of the state it bounds (2026-08-12)	85
50. $H(n=2)$ rate coefficients, the trapped- $2p$ balance, and de-excitation heating (2026-08-12)	85
51. Molecular chemistry and trace metals solved in one system (2026-08-13)	87
52. One electron density for the cooling and for the ionization solver (2026-08-13)	90
53. A fourth regression case for the molecular + metals system (2026-08-13)	92
A Code-size summary: EXHALE vs. the original ATES (2026-06-07; updated 2026-07-02)	93
B Fortran source inherited unchanged from ATES	94

1. Runtime domain-extent option: Roche vs Spherical

Added 2026-06-04.

1.1 Overview

ATES historically hardcodes the full Roche potential (planetary gravity plus the stellar-tidal and centrifugal terms) and truncates the radial grid at the Hill/L1 radius. For close-in planets this Roche lobe is small; for WASP-121b the outer boundary lands at $r_{\max} \approx 2.01 R_p$. Huang et al. (2023) Case A, by contrast, is a *spherical* model (no tidal terms) extending to $\sim 19 R_p$.

To support both geometries from a single executable, a runtime option was added with two modes:

- **Roche** (default): full Roche potential, domain truncated at the Hill/L1 radius. Identical to the previous ATES behavior.
- **Spherical**: pure planetary potential $\phi = -b_0/r$ (tidal and centrifugal terms dropped), domain extended to a user-specified outer radius.

The mode is selected by two optional lines appended to `input.inp`; when they are absent the code reverts to Roche, so pre-existing input files are unaffected.

1.2 The two domain modes

1.2.1 Gravitational potential

All quantities are in code-normalized units: radius r in units of the planetary radius $R_0 = R_p$, with $b_0 = GM_p\mu/(k_B T_0 R_0)$ the surface Jeans parameter, $M_r = M_*/M_p$ the star-to-planet mass ratio, and $\tilde{a} = a_{\text{orb}}/R_0$ the orbital distance in planetary radii.

Roche mode. The dimensionless potential and its gradient are

$$\phi(r) = -\frac{b_0}{r} - \frac{b_0 M_r}{\tilde{a} - r} - \frac{b_0(1+M_r)}{2\tilde{a}^3} \left(\frac{\tilde{a} M_r}{1+M_r} - r \right)^2, \quad (1)$$

$$\frac{d\phi}{dr} = \frac{b_0}{r^2} - \frac{b_0 M_r}{(\tilde{a} - r)^2} + \frac{b_0(1+M_r)}{\tilde{a}^3} \left(\frac{\tilde{a} M_r}{1+M_r} - r \right). \quad (2)$$

The three terms are planetary gravity, stellar tidal attraction, and the centrifugal potential of the co-rotating frame.

Spherical mode. Only the planetary term is kept:

$$\phi(r) = -\frac{b_0}{r}, \quad \frac{d\phi}{dr} = \frac{b_0}{r^2}. \quad (3)$$

1.2.2 Outer boundary

- **Roche:** $r_{\max} = (3M_r)^{-1/3} \tilde{a}$, the Hill-radius approximation to the L1 distance (unchanged from before).
- **Spherical:** $r_{\max}$ is set directly by the user-supplied “Outer radius [R_p]” value (in R_p , since the grid is normalized to $R_0 = R_p$).

1.2.3 Why spherical mode must drop the tidal terms

The outer boundary cannot simply be pushed past L1 while keeping the full Roche potential. Beyond L1 the Roche potential *turns over*: the effective gravity changes sign (points toward the star), so a monotonic, steadily accelerating outward wind is no longer defined there and the sonic-point/steady-state treatment breaks down. Removing the stellar-tidal and centrifugal terms restores a monotonic potential ($\phi = -b_0/r$ increases monotonically toward zero as $r \rightarrow \infty$), which is what permits the domain to extend arbitrarily far. This is exactly the spherical geometry of Huang Case A. A guard enforces a sensible configuration: selecting Spherical mode without an “Outer radius [R_p]” > 1 stops the run with an explicit error rather than silently falling back.

1.3 Input interface and backward compatibility

Two optional lines are appended at the end of `input.inp`, after the “Force start:” line:

Domain mode: Spherical

Outer radius [R_p]: 19.0

- “Domain mode:” accepts `Spherical`; any other value (or a missing line) means Roche.
- “Outer radius [R_p]:” is read only in Spherical mode; it is ignored in Roche mode, where the Hill radius is computed automatically.

The reader scans the trailing lines by *keyword* (using `index(line, . . .)`), not by fixed line position, so blank trailing lines and line ordering do not matter, and `input.inp` files written before this change (which lack both lines) keep the default Roche behavior. The active mode is echoed in the setup report `EXHALE_setup.out`.

1.4 Implementation

Touched files.

- `src/modules/init/parameters.f90`: new globals `spherical_domain` (logical, default `.false.`) and `r_out_user` (the outer radius in R_p).
- `src/modules/functions/grav_field.f90`: `phi` and `Dphi` branch on `spherical_domain`.
- `src/modules/files_IO/input_read.f90`: keyword scan of the trailing lines, and the $r_{\max}$ branch with the validity guard.
- `src/modules/files_IO/write_setup_report.f90`: prints the active mode to `EXHALE_setup.out`.

Gravity branch.

```
double precision function phi(r_in)
real*8, intent(in) :: r_in
if (spherical_domain) then
  phi = -b0/r_in
else
  phi = -b0/r_in &
        -b0*Mrapp/(atilde-r_in) &
        -b0*(1.0+Mrapp)/(2.0*atilde**3.0)* &
        (atilde*Mrapp/(1.0+Mrapp)-r_in)**2.0
endif
end function phi
```

1.5 Validation on WASP-121b

Two runs were made that differ *only* in the domain-extent option; `input.inp` was otherwise identical (power-law SED, $\text{He/H} = 0.0851$, trace metals C/N/O on, “Rate/2 + Mdot/2” 2-D flux approximation). Both reached steady state via the proper criterion (momentum constant, $du < du_{\text{th}}$), not a stall or a false convergence.

mode	$r_{\max} [R_p]$	$\log_{10} \dot{M}$	sonic $[R_p]$	$r(50\% \text{ HII})$	$T_{\text{peak}} [\text{K}]$
Roche (default)	2.01	12.68	1.825	1.178	9662
Spherical	19.00	11.81	5.270	1.152	10752
Huang Case A (ref.)	~ 19	12.57	—	—	~ 12000

$\dot{M}$ in $\log_{10} \text{ g s}^{-1}$. Convergence took $\approx 5,392$ iterations (Roche) and $\approx 154,183$ (Spherical, the larger domain).

Observations (for evaluation, not pass/fail).

- The Roche run reproduces the prior Phase-0 Case A result exactly ($\log_{10} \dot{M} = 12.68$), confirming the new optional block does not perturb default behavior.

- Spherical T_{peak} (10,752 K) is higher than Roche (9,662 K) and closer to Huang's $\sim 12,000$ K — the extended domain lets the profile develop past the $2-R_p$ L1 clip.
- Spherical $\dot{M}$ (11.81) is lower than both Roche (12.68) and Huang (12.57), consistent with the loss of tidal enhancement, but ~ 0.76 dex below Huang.
- The sonic point moves outward from 1.83 to $5.27 R_p$ (gentler acceleration without tidal forces).

Caveats. Two confounders remain before the spherical $\dot{M}/T$ can be compared to Huang cleanly: (a) only C/N/O metals are included here, not Huang's full Mg/Fe/Si/Ca/Na/K/S set (Phase 1); and (b) the same $N = 500$ cells now span $1 \rightarrow 19 R_p$, so the inner resolution is much coarser than in the Roche run — a denser-grid spherical re-run may be needed before drawing conclusions.

1.6 Status and relation to the Huang plan

This option closes the Phase-0 geometry gap: the Roche mode is the validated baseline, and the Spherical mode provides the Case A geometry needed for later comparison. The tidal RLOF physics of Huang Phase 4 is a separate, richer treatment (substellar streamline, L1 location, terminator mapping); the present Roche mode is the existing single-streamline truncation, not the full RLOF model. The Spherical mode is the immediate tool for Case A reproduction.

2. Ionization network: add magnesium (Mg I / II / III)

Added 2026-06-05.

2.1 Overview

The first trace metal beyond C/N/O, following the Huang et al. (2023) WASP-121b species set (Mg, Fe, Si, O, C, N, S, Ca, Na, K). Per the Phase-1 plan, elements are added one at a time (order Mg $\rightarrow$ Fe $\rightarrow$ Si $\rightarrow$ Ca $\rightarrow$ Na $\rightarrow$ K $\rightarrow$ S); this section covers **Mg only**. Magnesium is solved to second ionization (Mg I / Mg II / Mg III) inside the existing coupled MINPACK ionization system, raising the species count to 18 and the ratio-unknown count to 11. The change is confined to the ionization/radiation plumbing: there are no hydrodynamic or mass-density changes, and `calc_ne/calc_rho` still count only H/He — metal electrons enter solely inside the MINPACK system.

2.2 Atomic data and rates

Sourced, not fabricated:

- **Photoionization:** Verner et al. (1996) outer-shell fits, with Mg I threshold 7.646 eV and Mg II threshold 15.035 eV (`cross_sec.f90: sigma_MgI, sigma_MgII`).
- **Recombination:** Badnell RR+DR — the same source already used for C/N/O — via `lookup_rec` cases and `rec_MgII` ($\rightarrow$ Mg I) / `rec_MgIII` ($\rightarrow$ Mg II). This is a *deliberate deviation* from the validation-rate prescription

in `Huang_update_plan.md` §1.3 (Shull & van Steenberg 1982 for Mg); it corresponds to the plan's Phase-6 *production* rates, applied early to match the existing C/N/O treatment. The Mg ion/neutral balance below should be read as the Badnell result, not a Huang-faithful one.

- **Collisional ionization:** Voronov (1997) fits, `ion_coeff_MgI` and `ion_coeff_MgII`.
- **Line cooling:** *not yet ported* — `cool_MgI` = `cool_MgII` = 0. No published Mg two-level/CHIANTI cooling fit is implemented (that is Phase 2). Magnesium therefore acts only as an electron donor and a photoelectric-heating channel here, not as a coolant.
- **Charge exchange Mg + H⁺:** omitted in this first cut.

Grid extension below 13.6 eV. Because Mg I ionizes at 7.646 eV, turning Mg on triggers the existing low-IP-metal branch (`thereis_lowIP_metal`), which extends the energy grid and the stellar SED down to 7.646 eV (`set_energy_vectors.f90`). This is required to photoionize Mg I, but has a side effect noted under Caveats.

2.3 Touched files

- `init/parameters.f90`: `n_species` = 18; `e_th_MgI/e_th_MgII`; `X_Mg`; `s_mgi/s_mgii`; `thereis_lowIP_metal`.
- `functions/cross_sec.f90`: `sigma_MgI`, `sigma_MgII`.
- `radiation/Cool_coeff.f90`: Badnell fits, `lookup_rec`, `rec_MgII/rec_MgIII`, Voronov `ion_coeff_MgI/MgII`, `cool_MgI/MgII` (= 0, flagged).
- `radiation/util_ion_eq.f90`: `PH_heat_HHe` and `eval_cool` carry Mg.
- `radiation/ionization_equilibrium.f90`: Mg densities, params slots 49–55, MINPACK unknowns 10/11 (Mg II/Mg III ratios), solution unpack.
- `init/set_energy_vectors.f90`: 7.646 eV sub-grid; allocate/fill `s_mgi/s_mgii`.
- `files_IO/input_read.f90`: `X_Mg` default; `N_eq` = 11; `metals/low-IP` flags. `files_IO/metals_input_read.f90`: Mg I abundance key. `files_IO/write_output.f90`: Mg I/II/III as `Ion_species.txt` columns 17–19.
- Threading through arrays, I/O and post-processing: `init/set_IC.f90`, `files_IO/load_IC.f90`, `init/init.f90`, `EXHALE_main.f90`, `time_step/energy_semi_implicit.f90`, `functions/utilities.f90`, `post_process/post_process_adv.f90`.

`post_process_adv.f90` also received a latent-bug fix: its advection systems fill only the H/He equations, so they now use a local `Neq_adv` rather than the global `N_eq` = 11, which would otherwise feed `hybrd1` uninitialized metal entries.

2.4 Validation on WASP-121b

Four runs with identical input.inp (power-law SED, He/H = 0.0851, Roche domain), differing only in which metals are present. All reached steady state via $du < du_{\text{th}}$. Mg/H is conserved at the input $X_{\text{Mg}} = 3.98 \times 10^{-5}$ across the domain.

run	metals	grid floor	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$
no metals	H/He only	13.6 eV	12.92	13722	1.703
Mg only	Mg	7.646 eV	12.92	13682	1.697
C/N/O	C, N, O	13.6 eV	12.68	9662	1.199
C/N/O + Mg	C, N, O, Mg	7.646 eV	12.76	9828	1.224
Huang Case A	full set	—	12.57	~12000	—

$\dot{M}$ in $\log_{10} \text{ g s}^{-1}$; r in R_p .

Observations (for evaluation, not pass/fail).

- **Clean Mg isolation** (Mg only vs no metals): $\Delta \log_{10} \dot{M} = +0.000$ dex, $\Delta T_{\text{peak}} = -40$ K. Adding magnesium alone leaves the wind essentially unchanged, as expected for a 4×10^{-5} trace species with no line cooling: it neither heats nor cools appreciably, and the soft-UV band it adds is optically thin without other absorbers.
- **Mg ionization structure**: predominantly Mg II (Mg^+) through the lower thermosphere, predominantly Mg III (Mg^{2+}) in the upper thermosphere (Mg III fraction ~ 0.47 at $1.4 R_p$ rising to ~ 0.81 at $1.9 R_p$); neutral Mg I stays subdominant. Same qualitative ordering as Huang+2023 Fig. 12.
- **MINPACK convergence**: the 11-unknown system converges without NaNs across the full temperature range.
- Mg-only $\dot{M}$ (12.92) is $\sim 2.2 \times$ the Huang Case A spherical value (12.57), in the same range as the prior C/N/O Roche run; the metal-set and geometry confounds from Phase 0 still apply.

Caveats and gaps.

- **C I grid-extension confound** (C/N/O + Mg vs C/N/O). The $\dot{M}/T$ offset between these two (+0.080 dex, +165 K) is *not* a magnesium effect: extending the grid to 7.646 eV exposes **C I** (threshold 11.26 eV) to the new 11.26–13.6 eV photons the 13.6 eV-floor C/N/O run never sees, adding C I photoionization heating. This is why the clean Mg measurement uses Mg-only vs no-metals, where C is absent.
- **No Mg line cooling** ($\text{cool_MgI} = \text{cool_MgII} = 0$). Huang’s headline Mg result — Mg II h&k emission dominating the cooling at 1.15 – $1.4 R_p$ — cannot appear here by construction. That is Phase 2.
- **No Mg + H^+ charge exchange** yet.
- **Recombination is Badnell**, not the §1.3 validation rate.

2.5 Status and relation to the Huang plan

This completes the **Mg** step of Phase 1 (network expansion). The ionization structure satisfies the qualitative Fig. 12 ordering and the “ $\dot{M}$ within a factor ~ 2 of Case A” part of the gate; the thermal-balance part of Phase 1 is deferred to Phase 2 (Mg cooling). Next element in the planned order: **Fe**. Scratch runs and the comparison notebook live in `mg_validation/`.

3. Phase 1a: interface / data-structure refactor (no new physics)

3.1 Overview

Each metal added so far (C/N/O, then Mg in §2) widened the named-argument lists of `PH_heat_HHe`, `eval_cool`, `calc_column_dens_metals`, `write_output`, the MINPACK system `ion_system_HeH_metals`, and the driver `ionization_equilibrium.f90` by several arguments for each ion, and threaded the MINPACK coefficients for each cell through a hand-indexed flat `params(60)`. This step converts that plumbing to a **species-metadata table + array-indexed (2D) data flow**, so that adding an element becomes “add rows to a table” instead of “thread N new arguments through M routines.” **No physics changes** — the network is still C/N/O + Mg, and the result is gated bit-for-bit against §2 (below).

3.2 What changed (Steps A–F)

- **A — Metadata module** (`init/species_table.f90`, new). One canonical table holds metadata for each ion (length `n_mion` = 12): `f_sp` column (`mion_fsp`), parent element (`mion_elem`), stage / charge² (`mion_stage`, `mion_z2`), photo-ionizable flag and photo-table column (`mion_isphot`, `mion_iphot`), threshold (`mion_ethr`), cooling flag (`mion_iscool`); and metadata for each element (length `n_melem` = 4): nuclear charge (`melem_Z`), neutral-ion index (`melem_i0`), top stage (`melem_top`). The canonical ion order (C I..Mg III = `f_sp` columns 7–18) is defined here once; element indices `iel_C/O/N/Mg`.
- **B — write_output 2D**. Metal ion densities pass as `nm(:, 1:n_mion)` and are written with an implied-do over the table (`((nm(j,i)*n0, i = 1,n_mion))`), not a named write for each ion.
- **C — PH_heat_HHe + column density 2D**. Optical depth, photoheating, and the photoionization rates loop over photo-ionizable ions via `sigma_tab(:,mion_iphot(i))` and `Nm_col`, returning a 2D rate array `P_m(:, 1:n_mion)` (inert top stages stay zero) instead of separate `P_CI..P_MgII` arguments, one for each ion.
- **D — eval_cool metal channels 2D**. Recombination and collisional-ionization coefficients return as 2D arrays (`rec_m`, `aion_m`); bremsstrahlung sums `mion_z2(i)*GF_elem(:,mion_elem(i))*nm(:,i)` over ions. No separate coefficient arguments for each ion.
- **E — Generalized MINPACK system**. `ion_system_HeH_metals` keeps `hybrid1`’s fixed (`N_eq,x,fvec,iflag,params`) signature but (i) receives the metal coefficients for each cell through a

module-level block in `System_HeH_metals` (`set_metal_coeffs` stores `met_ntot/g0/g1/b0/b1/a1/a2`), set by the driver before each `hybrd1` call — safe because the ionization-equilibrium cell loop is serial — and (ii) assembles `fvec` by **looping over elements** (two balance equations each, force-zeroing absent elements). `params` now carries only H/He (1–11), HeITR (12–18), and the C/N/O charge-transfer rates (40–48); the old slots 49–55, one for each element, are gone. The driver builds metal densities, initial guesses, and the solution unpack by looping over `n_melem/n_mion`. The system size is now the formula $N_{eq} = 3 + 2 \times n_{melem}$ (= 11 for the current four elements), set in `input_read.f90`.

- **F — Generalized grid floor.** `set_energy_vectors.f90` no longer hardwires the sub-13.6 eV floor to Mg’s 7.646 eV; it computes `e_sub_low = min over active elements of the neutral threshold mion_ethr(melem_i0(e))` that lies below `e_th_HI`. The trigger flag `thereis_lowIP_metal` is set generically in `input_read.f90` (any active element with a sub-13.6 eV neutral threshold), driven by a new canonical-order abundance array `melem_ab` (declared in `parameters.f90`, populated in `input_read.f90`). For the current set only CI (11.26 eV) and Mg I (7.646 eV) lie below 13.6 eV and the minimum is Mg I, so the floor is unchanged at 7.646 eV when Mg is present. (Consequence: a future C/N/O-without-Mg run would now extend the grid to 11.26 eV, which it did not before — and once Phase 1b adds Na/K/Ca, whose neutral thresholds 5.14/4.34/6.11 eV sit below Mg’s, the floor drops further.)

3.3 Gate: bit-for-bit reproduction of the C/N/O + Mg run

With C/N/O + Mg active and the same `input.inp/metals.inp` as §2, the refactored code reproduces the `mg_validation/cno_mg` reference to round-off: $\log_{10} \dot{M} = 12.76$; `Hydro_ioniz.txt`, `Hydro_ioniz_adv.txt`, and `Ion_species.txt` are byte-identical to the frozen reference. `Ion_species_adv.txt` differs at exactly one number — line 1 (innermost ghost cell), column 7 (the HeITR fraction): the reference holds 4.0048×10^{-5} left in uninitialized stack memory while the deterministic build writes 0.0; every other column of that line and every other line match. This is a pre-existing reference artifact (a HeITR ghost-cell slot never used in the H/He + metals run), not an effect of the refactor.

3.4 What is deliberately *not* generalized yet

- **Charge exchange with H** is still hardcoded for C/N/O only (the `fvec(1)` C,N,O sum and the `cxlo/cxup` mapping via `iel_C/O/N` in `System_HeH_metals`). Generalizing it is Phase 1d.
- **Atomic data for each ion and abundance for each element remain manual entries** — the `sigma_tab(:, 1..8)` assignments and the `sigma_XX` cross-section functions in `set_energy_vectors.f90`, the abundance scalars `X_C/X_O/X_N/X_Mg`, the `metals.inp` keys, and the `melem_ab` population block. These are the atomic-data / abundance rows expected to grow per element in Phase 1b/1c; they are data, not interface plumbing.

3.5 Touched files

- **New:** `init/species_table.f90` (Step A).
- `init/parameters.f90`: `melem_ab(:)` abundance array; generalized `thereis_lowIP_metal` comment.

- `files_IO/input_read.f90`: $N_{\text{eq}} = 3 + 2 \times n_{\text{melem}}$; `melem_ab` population; generalized low-IP-metal trigger.
- `files_IO/write_output.f90`: 2D metal-column writer (B).
- `radiation/util_ion_eq.f90`: `PH_heat_HHe` + column density 2D (C); `eval_cool` metal channels 2D (D).
- `nonlinear_system_solver/System_HeH_metals.f90`: module-level metal block for each cell + element-loop residual (E).
- `radiation/ionization_equilibrium.f90`: metadata-driven density build, params packing (H/He + charge transfer only), initial guesses, solution unpack (E).
- `init/set_energy_vectors.f90`: min-over-active-metals grid floor (F).

3.6 Status and relation to the Huang plan

Phase 1a is complete: the element-by-element plumbing of §2 is now table-driven. Adding the Phase-1b metals should require only metadata rows plus atomic/abundance data, with **no changes to `System_HeH_metals.f90` or `ionization_equilibrium.f90`**. Next: **Phase 1b** — Si, Ca, Na, K, S added together as a uniform-template batch. Gate scratch run: `/tmp/ates_gate/cno_mg/` vs frozen reference `mg_validation/cno_mg/output/`.

4. Phase 1b: correct C/N/O cross sections and add Si, Ca, Na, K, S

4.1 Overview

Two changes are made together here, both enabled by the Phase 1a (§3) table-driven plumbing:

1. **Corrected the C/N/O photoionization cross sections** to the full Verner, Ferland, Korista & Yakovlev (1996; VFKY96) Eq. (1) form. The previous code used the simpler Verner & Yakovlev (1995) inner-shell form (no y_0/y_1 offset/asymptote parameters) for C/N/O; that form is exact only when $y_0 = y_1 = 0$, which does not hold for the C/N/O outer shells. This is the dominant *physics* change.
2. **Added the five uniform-template metals Si, Ca, Na, K, S** in one batch. With Mg (§2) the in-code element set is now C, N, O, Mg, Si, Ca, Na, K, S — the full Huang+2023 §2.4 set *except Fe* (Phase 1c). Counts: `n_species = 30`, `n_mion = 24`, `n_melem = 9`.

The metal MINPACK solver was also **renamed** (below).

4.2 Variable-stage solver (2- and 3-stage elements in one packing)

Huang+2023 carries Si, Ca (like Mg) to second ionization but Na, K, S only to first. The solver keeps a uniform packing of **two unknowns for each element** ($N_{\text{eq}} = 3 + 2 \times n_{\text{melem}} = 21$ with all nine elements) and a `melem_top` flag for each element:

- `melem_top = 2` (C, O, N, Mg, Si, Ca): solve both $X^0 \leftrightarrow X^+$ and $X^+ \leftrightarrow X^{2+}$.

- `melem_top = 1` (Na, K, S): solve only $X^0 \leftrightarrow X^+$ and **pin the unused upper unknown to zero** (`fvec(ix+1) = x(ix+1)`), keeping the Jacobian non-singular. Every `ix+1` access is guarded by `melem_top` so a two-stage element never touches its neighbor's column.

Zero-abundance elements are still force-zeroed, so the same 21-unknown system solves C-only, C/N/O+Mg, or all nine without renumbering.

4.3 Atomic data and rates (sources, not fabricated)

- **Photoionization cross sections** (`functions/cross_sec.f90`): full VFKY96 Eq. (1) form (`sigma_VFKY96`) for C I/II, N I/II, O I/II, Si I/II, Ca I/II, S I; simpler Verner & Yakovlev 1995 form (`sigma_Verner96`, `y_0=y_1=0`) for Mg I/II, Na I, K I (their published fits have `y_0=y_1=0`, so the forms agree exactly). Neutral thresholds [eV]: Si I 8.152, Ca I 6.113, Na I 5.139, K I 4.341, S I 10.36.
- **Collisional ionization**: Voronov (1997) for all new ions.
- **Recombination**: Badnell RR+DR via the swappable dispatcher, as for C/N/O and Mg — the *same deliberate deviation* from the `Huang_update_plan.md` §1.3 validation rates already accepted for Mg (the Phase-6 production rate applied early). Read the balances as the Badnell result.
- **Line cooling**: *not* ported (`mion_iscool = .false.`), flagged To Be Checked; the five act only as electron donors and photoelectric-heating channels (Phase 2).
- **Charge exchange with H**: *not* added; still C/N/O only (Phase 1d).

4.4 Rename: `System_HeHCO` → `System_HeH_metals`

The metal solver was named after the AIOLOS heritage (AIOLOS tracked only C and O as metals $\Rightarrow$ “HCO” = H, C, O). It now carries H + He + nine metals, so the C/O-specific name was misleading. Renamed to parallel the existing `System_HeH`/`System_HeH_TR` family: module `System_HeHCO` → `System_HeH_metals`; `ion_system_HeHCO` → `ion_system_HeH_metals`; `set_hco_metals` → `set_metal_coeffs`; `hco_*` → `met_*`. Pure rename, no behavior change (§3 above is updated to the new names).

4.5 Validation on WASP-121b

Three converged runs, identical planet / SED / grid (power-law SED, He/H=0.0851, Roche domain, “Rate/2 + $\dot{M}$ dot/2”); they differ only in code version and metal set.

key	code	metals	C/N/O σ	floor [eV]	$\log_{10} \dot{M}$	T_{peak} [K]	steps
ref	old	C,N,O,Mg	Verner96	7.646	12.76	9828	5003
A	new	C,N,O,Mg	VFKY96	7.646	12.85	12197	15382
B	new	+Si,Ca,Na,K,S	VFKY96	4.341	12.85	12135	15505

$\dot{M}$ in $\log_{10}$ g/s; Huang+2023 Case A reference $\log_{10}\dot{M} = 12.57$. ref and A share the 7.646 eV grid floor (Mg sets the minimum in the C/N/O+Mg set), so ref vs A isolates the C/N/O cross-section change with no grid-floor confound; A vs B changes both the element set and the floor (K I pulls it to 4.341 eV).

Observations (for evaluation, not pass/fail):

- **C/N/O correction (ref→A):** $\Delta \log_{10}\dot{M} = +0.090$ dex; T_{peak} rises 9828 $\rightarrow$ 12197 K and the peak moves outward from 1.22 to $\sim 1.90 R_p$. The corrected VFky96 opacity deposits more photoheating; the hotter peak now sits near Huang's ~ 12000 K. A genuine model change, not numerical.
- **Adding Si/Ca/Na/K/S (A→B):** $\Delta \log_{10}\dot{M} = +0.000$ dex, $\Delta T_{\text{peak}} = -62$ K, despite the floor dropping to 4.341 eV — at solar abundance and with no line cooling these five are radiatively near-passive donors.
- **Mg moves indirectly.** Mg's own cross sections are byte-identical between ref and A, yet its ionization fractions shift by up to ~ 0.1 : an indirect response to the hotter, more extended profile (changed n_e , optical depth, photoionizing flux), not a change in Mg atomic data. Mg is coupled to the global state, hence *not* an isolated control.
- **Numerics.** Zero NaN/Inf; the 21-unknown system converges over the full T range; X/H for each element is conserved to $\sim 10^{-13}$ for all nine elements including the two-stage Na/K/S — direct evidence the variable-stage packing does not leak between element columns.
- **New-metal ionization (run B).** Si, Ca mostly *doubly* ionized in the upper thermosphere; Na, K, S mostly *singly* ionized (no second stage carried) — the qualitative Huang+2023 Fig. 12 ordering.

4.6 Caveats and gaps

- **No metal line cooling** for Si/Ca/Na/K/S (Phase 2); the thermal structure will not match Huang until Fe II/Mg II/Ca II cooling is in.
- **No charge exchange with H** for the new elements or Mg (Phase 1d).
- **Recombination is Badnell, not the §1.3 validation rate** (deliberate, as Mg).
- $\dot{M}$ is $\sim 1.9 \times$ Huang Case A (12.85 vs 12.57); the metal-set (Fe still missing) and geometry confounds from Phase 0 still apply.
- **Not done here:** Huang Fig. 12 is not digitized/overlaid (only the qualitative ordering is checked); the atomic-data values are not re-derived beyond conservation and smoothness.

4.7 Regression reference for Phase 1c

The accepted run A output is frozen at phase1b_validation/.gate_ref_cno_mg_vfky96/output/ as the post-Phase-1b C/N/O+Mg baseline. When Phase 1c adds Fe, running with Fe *off* must reproduce it; because adding Fe bumps `n_species/n_mion` (and the `Ion_species.txt` column count), the Phase-1c check is a **round-off match on the**

shared physical columns, not a byte-for-byte compare (the `mg_validation/cno_mg` byte-gate applied only to the physics-neutral Phase 1a refactor, where `n_species` was unchanged).

4.8 Status and relation to the Huang plan

This completes **Phase 1b**: the network carries C/N/O/Mg/Si/Ca/Na/K/S, the C/N/O cross sections are corrected to VFKY96, and the variable-stage solver handles 1st- and 2nd-ionization elements uniformly. The gate’s “converges/conserves/Fig. 12 ordering/ $\dot{M}$ within $\sim 2\times$ Case A” items are addressed by the numbers above; the thermal-balance part is deferred to Phase 2 (metal line cooling). Next: **Phase 1c** — iron (Fe I/II/III), which needs non-uniform cross-section (Zatsarinny+2019) and recombination data. Scratch runs and notebook: `phase1b_validation/Phase1b_validation.ipynb`.

5. Phase 1c: add iron (Fe I / II / III)

5.1 Overview

Iron is added as the third three-stage trace metal, completing the Huang+2023 §2.4 element set (C, N, O, Mg, Si, Ca, Na, K, S, **Fe**). No solver, cooling-loop, heating, output, or post-processing code changed: the variable-stage MINPACK packing and every metal loop are generic over `n_melem`, so adding iron is purely new metadata rows plus four new rate functions. Counts: `n_species` = 33, `n_mion` = 27, `n_melem` = 10, `n_mphot` = 17; the system grows to `N_eq` = 3 + 2*n_melem = 23. Iron carries `melem_top` = 2 (three stages, like C/O/N/Mg/Si/Ca).

5.2 Atomic data and rates (sources, not fabricated)

Iron is the one element whose data are *not* the uniform-template choice. Each quantity was checked against its primary source before encoding; the Voronov fit in particular was cross-checked because an earlier abandoned attempt had its P/X/K columns transposed.

- **Photoionization cross sections** (`functions/cross_sec.f90`): single-shell VFKY96 form (`sigma_VFKY96`) for Fe I and Fe II, parameters from the Verner+1996 `phfit2.f` PH2 block. Neutral thresholds [eV]: Fe I **7.902** (NIST ionization potential), Fe II **16.199**. *To Be Checked / known limitation*: Huang+2023 use the multi-shell Zatsarinny+2019 Fe I cross section; the single-shell fit here under-represents inner-shell absorption at high XUV energy, so the iron photoionization rate is a lower bound at the hardest photons (flagged inline). A multi-shell upgrade is left for later.
- **Collisional ionization**: Voronov (1997) `cfit.dat`. Fe I (`dE=7.9`, `A=2.52e-7`, `P=0`, `X=0.701`, `K=0.25`), Fe II (`dE=16.2`, `A=2.21e-8`, `P=1`, `X=0.033`, `K=0.45`).
- **Recombination**: the Huang+2023 Eqs. (5)–(6) total fits (a dielectronic Arrhenius term $A T^{-1.5} e^{-T_0/T} (1 + B e^{-T_1/T})$ plus a radiative power law $C (T/10^4)^{-\eta}$) used *directly*, not the Badnell dispatcher used for the lighter metals. Deliberate: the Huang iron rates are the validation target, so the model

runs on exactly those rates; the Badnell-swap for iron is deferred to Phase 6. $\text{Fe II}+e \rightarrow \text{Fe I}$: $A=2.833\text{e-}8$, $T0=5.731\text{e}4$, $B=1.383\text{e}4$, $T1=120.4$; $C=1.248\text{e-}12$, $\text{eta}=0.485$. $\text{Fe III}+e \rightarrow \text{Fe II}$: $A=1.094\text{e-}5$, $T0=1.490\text{e}4$, $B=36.74$, $T1=1.153\text{e}5$; $C=1.728\text{e-}12$, $\text{eta}=0.618$.

- **Line cooling:** *not* ported (`mion_iscool = .false.` for all Fe stages) — Phase 2. This matters more for iron than for the lighter metals: Fe II is the *dominant* low-altitude line coolant in Huang+2023, so a metals-on thermal profile is *expected to run hotter than Huang* until Phase 2.
- **Charge exchange with H:** *not* added; still C/N/O only (Phase 1d). Iron’s $\text{Fe}^+ + \text{H}^+ \leftrightarrow \text{Fe}^{2+} + \text{H}$ reactions (Huang Table 4) come in Phase 1d.

5.3 Touched files

Metadata rows + four rate functions + abundance/IC wiring; no solver, cooling, heating, output, or post-process changes:

- `init/species_table.f90` — `n_mion=27`, `n_melem=10`, `n_mphot=17`; Fe I/II/III rows (`f_sp` cols 31–33); `iel_Fe=10`; `melem_top` appends 2.
- `functions/cross_sec.f90` — `sigma_FeI`, `sigma_FeII`.
- `init/set_energy_vectors.f90` — `sigma_tab(:,16)=sigma_FeI`, `(:,17)=sigma_FeII`.
- `radiation/Cool_coeff.f90` — Huang recombination `alpha_rec_FeI/FeII_Huang` with `rec_FeII/rec_FeIII` wrappers; Voronov `ion_coeff_FeI/FeII`; dispatcher cases for FeII/FeIII recombination and FeI/FeII ionization. No `cool_coeff_metal` case for Fe (Phase 2).
- `init/parameters.f90`, `files_IO/input_read.f90`, `files_IO/metals_input_read.f90` — `n_species=33`; `X_Fe`; `metals.inp` keys `FeI/Fe`; `melem_ab(iel_Fe)`; `thereis_metals` trigger.
- `init/set_IC.f90`, `files_IO/load_IC.f90` — neutral-iron IC for `f_sp` cols 31–33.

5.4 Iron abundance

Asplund et al. (2009) solar $\text{Fe/H} = 3.16\text{e-}5$ ($\log \epsilon_{\text{Fe}} = 7.50$). Huang+2023 adopt $4\times$ solar metallicity for WASP-121b, so the validation run uses $X_{\text{Fe}} = 1.26\text{e-}4$ by number relative to H; C/N/O/Mg are held at solar to isolate the iron change.

5.5 Validation on WASP-121b

Two converged runs (identical planet / SED / grid; power-law SED, $\text{He/H}=0.0851$, Roche domain, “Rate/2 + Mdot/2”) plus the frozen Phase-1b C/N/O+Mg baseline:

key	metals	iron	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$	steps
ref	C,N,O,Mg	frozen 1b run A	12.85	12197	1.899	15382
off	C,N,O,Mg	off ($X_{\text{Fe}}=0$)	12.85	12197	1.899	15382
on	+Fe	on, $4\times$ solar	12.87	12223	1.915	16566

$\dot{M}$ in $\log_{10}$ g/s; Huang+2023 Case A reference $\log_{10} \dot{M} = 12.57$.

Observations (for evaluation, not pass/fail):

- **Regression, force-zeroed iron (off vs ref).** The Phase-1c binary carries three extra species and two extra MINPACK unknowns, yet run off reproduces the frozen baseline on every shared column to the round-off floor: worst relative difference 2.0×10^{-13} across the seven hydro columns and 3.7×10^{-13} (NI) across the 30 shared species. The three new Fe columns are *identically zero, no NaNs*. With $X_{\text{Fe}}=0$ the iron block reduces to identity rows in the residual ($f_{\text{vec}}=x$, rooted at zero) and contributes exactly zero electrons, so the 23-unknown system returns the 21-unknown answer — adding iron is non-perturbative when off.
- **Iron ionization structure (on).** Fe II dominates the base (fraction 0.89 at $r = 1$, peaking ~ 0.96 near $r = 1.02$); Fe III dominates the escaping flow (0.97 at the outer edge), with the Fe II/Fe III crossover near $r \simeq 1.25$; Fe I is minor everywhere (max ~ 0.11 at the base). This is the qualitative Huang+2023 Fig. 12 arrangement — Fe I confined to the base, Fe II at the base, Fe III through the outflow.
- **Iron as an electron donor.** At $4\times$ solar, iron supplies up to $\sim 26\%$ of the free-electron density, so it feeds back on the whole coupled solution, not just its own columns.
- **Switching iron on (off $\rightarrow$ on).** $\Delta \log_{10} \dot{M} = +0.02$ dex, $\Delta T_{\text{peak}} = +26$ K ($12197 \rightarrow 12223$). Small and in the *heating* direction: with no Fe line cooling yet, iron acts only as an added electron donor / photoelectric-heating channel, so the thermosphere runs slightly hotter, not cooler.
- **Numerics.** Run on converged with zero NaNs; X/H for each element is conserved to 8.3×10^{-13} across all five active elements including iron.

5.6 Caveats and gaps

- **No iron line cooling** (`mion_iscool=.false.`) — Phase 2. Because Fe II is the dominant low-altitude coolant in Huang+2023, the run-on *temperature profile is not yet comparable to Huang* and is expected to sit above it. The Phase-1c physics target is the iron *ionization structure*, not $T(r)$.
- **No iron charge exchange with H** — Phase 1d.
- **Single-shell VFKY96 Fe I/II cross sections** under-represent inner-shell XUV vs Huang’s Zatsarinny+2019 Fe I (flagged in `cross_sec.f90`).
- **Recombination is the Huang fit**, not Badnell (deliberate; the Badnell-swap is Phase 6).
- $\dot{M}$ is still $\sim 2\times$ Huang Case A (12.87 vs 12.57); the missing metal cooling and the Phase-0 geometry confounds still apply.

- **Huang Fig. 12 is not digitized/overlaid** — only the qualitative Fe II-base / Fe III-outflow ordering is checked.

5.7 Regression reference for Phase 1d / Phase 2

The run-on output (C/N/O/Mg + 4×-solar Fe) is frozen at `phase1c_validation/.gate_ref_cno_mg_fe/` (input, `metals.inp`, `run.log`, `output/`) as the post-Phase-1c iron-on baseline. Phase 1d (iron charge exchange) and Phase 2 (iron line cooling) both *change* the iron ionization and the thermal structure deliberately, so this is the “before” snapshot for those comparisons, not a bit-reproduction target.

5.8 Status and relation to the Huang plan

Phase 1c **implementation is complete** and the iron ionization structure is presented for validation against Huang+2023 Fig. 12 (numbers and figures in the notebook; no pass/fail asserted here — that judgement is the reader’s). The ionization-network half of Phase 1 now carries the full Huang element set with verified iron atomic data; the thermal-balance half remains deferred (until Fe II cooling lands in Phase 2 the metals-on temperature profile will not match Huang). Next: **Phase 1d** — charge exchange with H for all metals (Huang Table 4), including $\text{Fe}^+ + \text{H}^+ \leftrightarrow \text{Fe}^{2+} + \text{H}$. Scratch runs and notebook: `phase1c_validation/Phase1c_validation.ipynb`.

6. Phase 1d: charge exchange with hydrogen (Huang Table 4)

6.1 Overview

Phase 1d adds the charge-exchange (charge-transfer, CT) reactions of Huang+2023 Table 4 as a single table-driven module. The default active set is the 23 **Group A** metal + H / H^+ reactions, which couple each metal’s ionization balance to the H^+ fraction. A runtime switch `cx_full 1` in `metals.inp` enables the rest of the table — **Group B** (He + H, 2 rows), **Group C** (metal + He, 6 rows), **Group D** (metal + metal, 32 rows). This *replaces* the previous hard-coded Kingdon & Ferland (1996) C/N/O + H charge-exchange terms in the solver with a generic implementation covering all ten metals plus He.

Re-sourcing C/N/O to Huang Table 4 has one deliberate consequence: Table 4 tabulates only the neutral $\leftrightarrow$ singly-ionized CT for C, N, O (rows A13–A18), so the switch *drops* the doubly-ionized $\text{C}^{2+}/\text{N}^{2+}/\text{O}^{2+} + \text{H}$ recombination terms the old KF96 implementation carried. This is an intentional behavior change to match the paper; it raises the C/N/O double-ionization fractions and reshapes the thermal profile (below). No solver structure, cooling loop, heating, output, or post-processing code changed beyond the CT wiring.

6.2 Atomic data and rates (sources, not fabricated)

All 63 rate rows are transcribed verbatim from `docs/charge_exchange_table4.md`, the image-verified transcription of Huang+2023 Table 4 (23 Group A + 2 Group B + 6 Group C + 32 Group D). The paper text says “65 reactions”; the 2-row discrepancy is immaterial and noted in the reference doc.

- **Functional forms** (radiation/charge_exchange.f90): KF96 bracket $a T_4^b (1 + c e^{dT_4}) e^{-E_k/T_4}$ (kf); power law $a (T/300)^b e^{-E_k/T_4}$ (gp); 4th-order $\ln T$ polynomial (p4) for the radiative-CT rows of N / S / Na / K (Lin/Zhao/-Dutta/Watanabe); and the Satta (2013) $f_{\text{Si}}(T)$ and $D_{\text{CSi}}(T)$ helpers for the Si + He and C + Si rows.
- **Spot checks at $T = 10^4$ K** (against the table): $\text{Fe} + \text{H}^+ = 4.000 \times 10^{-9}$ (RV72), $\text{Fe}^+ + \text{H} = 2.011 \times 10^{-12}$, $\text{Fe}^{2+} + \text{H} = 1.260 \times 10^{-9}$ (ND87), $\text{O} + \text{H}^+ = 2.091 \times 10^{-9}$ (near-resonant, S99) — all match.
- **Rate clamp** $[0, 10^{-6}] \text{ cm}^3 \text{ s}^{-1}$: the upper bound guards the $\ln T$ polynomials, which diverge as $T \rightarrow 1$ K (never reached); the lower bound zeroes a few KF96 bracket fits that dip slightly negative below their validity floor (e.g. $\text{Mg}^+ + \text{H}^+$ at $T \lesssim 1500$ K), where the true rate is negligible.
- **Reference-doc notes**: Ca has *no* direct H charge exchange (only Group D); K has only the forward $\text{K} + \text{H}^+$ row. Reverse rates for pairs lacking direct data are Huang’s detailed-balance estimates, carried verbatim.

6.3 Touched files

A new module plus minimal wiring; the solver’s old element-by-element C/N/O CT block is deleted, not extended:

- radiation/charge_exchange.f90 (**new**) — the full Table-4 descriptor arrays, the cx_rate dispatcher, and the generic residual assembly. Entry points: cx_init (build the active list from cx_full), cx_set_cell(T) (rates in each cell), cx_add_to_fvec (add CT terms to the MINPACK residual via the cx_fvidx element/stage $\rightarrow$ row map), and the cx_full switch. The zero-abundance guard is automatic: each rate is $k n_{\text{donor}} n_{\text{acceptor}}$.
- nonlinear_system_solver/System_HeH_metals.f90 — the hard-coded C/N/O + H charge-exchange terms are *removed*; a single call cx_add_to_fvec(...) now feeds CT contributions into the H, He, and every metal-stage residual row.
- radiation/ionization_equilibrium.f90 — call cx_set_cell(T_K(j)) before each cell’s hybrd1 solve.
- files_IO/metals_input_read.f90 — parses cx_full <0|1> from metals.inp.
- files_IO/input_read.f90 — if (thereis_metals) call cx_init after the input read, before the ionization sweep.

6.4 Validation on WASP-121b

Two new converged runs (identical planet / SED / grid as Phase 1c) compared against the frozen Phase-1c iron-on output as the regression reference:

key	charge exchange	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$	steps
ref	KF96 C/N/O + H (frozen 1c on)	12.87	12223	1.915	16566
cx	Huang Table 4 Group A (default)	12.89	12955	1.778	13942
full	Huang Table 4 full, cx_full 1	12.88	12951	1.765	11661

$\dot{M}$ in $\log_{10} \text{ g/s}$; Huang+2023 Case A reference $\log_{10} \dot{M} = 12.57$. All three runs converged with **zero NaNs**.

Observations (for evaluation, not pass/fail):

- **H⁺ sink magnitude.** In the ionized bulk ($x_{\text{HII}} > 0.5$) the net metal charge exchange is at most 3.7×10^{-3} of radiative recombination (median 1.0×10^{-3}) — negligible as an H⁺ sink, consistent with Huang §4.2 (“Fe–H CX lower than H photoionization and recombination”). At the near-neutral base ($x_{\text{HII}} \sim 7.8 \times 10^{-8}$) the *ratio* is large (~ 118) only because H⁺ is itself a trace species there.
- **Near-resonant balance.** $\text{O} + \text{H}^+ \leftrightarrow \text{O}^+ + \text{H}$ is near-resonant and near-equilibrium at the base (gross destroy 3.66×10^5 vs produce 3.69×10^5 , net $+3.0 \times 10^3$). The net base H⁺ sink is instead driven by the fast, one-sided $\text{Fe} + \text{H}^+ \rightarrow \text{Fe}^+ + \text{H}$ channel (net -2.9×10^4).
- **C/N/O double ionization rises** (the dropped 2+ CT). Column-max fractions ref $\rightarrow$ cx: O III 0.871 $\rightarrow$ 0.946, N III 0.897 $\rightarrow$ 0.944, C III 0.851 $\rightarrow$ 0.856.
- **Iron base shift.** Fe II base fraction 0.889 $\rightarrow$ 0.896; Fe III essentially vanishes at the base (0.002 $\rightarrow$ 0.000): $\text{Fe} + \text{H}^+ \rightarrow \text{Fe}^+ + \text{H}$ pushes iron toward Fe⁺.
- **Thermal structure.** T_{peak} rises 12223 $\rightarrow$ 12955 K and moves inward (1.915 $\rightarrow$ 1.778 R_p). This is the *combined* effect of (i) the C/N/O re-sourcing that drops the 2+ recombination terms and (ii) the new metal + H CT. Charge exchange adds *no* radiative-loss channel, so it cannot cool the gas; Fe II / metal line cooling remains deferred to Phase 2, and the metals-on profile therefore still runs *hotter* than Huang.
- **Group A vs full Table 4 (cx vs full).** Turning on Groups B/C/D changes the bulk temperature by at most 1.5% (worst relative difference), with max $|\Delta T| = 556$ K at $r = 1.008$; at $r = 1.02$ the full run is 357 K cooler with more base H⁺ and less He II. $\dot{M}$ moves -0.01 dex and the outflow is nearly unchanged.
- **Dynamics.** v_{out} and sonic radius: ref 18.83 km/s @ 1.884 R_p ; cx 20.28 km/s @ 1.826 R_p ; full 20.40 km/s @ 1.826 R_p . The hotter, more inward thermal peak drives a slightly faster wind with the sonic point $\sim 0.06 R_p$ further in.

6.5 Caveats and gaps

- **No metal line cooling yet** (Phase 2). Charge exchange adds no loss term, so $T(r)$ is not comparable to Huang; the C/N/O re-sourcing makes the base hotter, not cooler.
- **C/N/O 2+ + H terms intentionally dropped** to match Table 4 — a real change from the Phase-1c (KF96) reference, and the dominant driver of the O III/N III rise rather than the new metal CT alone.
- $\dot{M}$ still $\sim 2 \times$ Huang Case A (12.88–12.89 vs 12.57); the missing metal cooling and the Phase-0 geometry confounds still apply.
- **Paper lists 65 reactions, 63 transcribed** — immaterial, documented in charge_exchange_table4.md.
- **No Huang figure digitized/overlaid** for the H⁺ budget; only the relative-magnitude statement (net CT $\ll$ recombination in the bulk) is checked.

6.6 Regression reference for Phase 2

The default Group-A run is frozen at `phase1d_validation/.gate_ref_cno_mg_fe_cx/` (input, `metals.inp`, `run.log`, output/) as the post-Phase-1d baseline. Because the C/N/O re-sourcing deliberately shifted both the ionization and the thermal structure away from the Phase-1c frozen reference, *this* snapshot — not the Phase-1c one — is the “before” reference for Phase 2 (metal line cooling, including Fe II).

6.7 Status and relation to the Huang plan

Phase 1d **implementation is complete**: the 23 Group-A metal + H reactions are active by default and the full Table 4 sits behind `cx_full`. With this, the **ionization-network half of Phase 1 is physically complete** relative to Huang+2023 §2.4 + Table 4 — the full element set, verified atomic data, and the metal–H charge-exchange coupling are all in place. The thermal-balance half remains deferred: until metal line cooling (Fe II in particular) lands in **Phase 2**, the metals-on temperature profile stays above Huang. Next: **Phase 2** — metal line cooling. Scratch runs and notebook: `phase1d_validation/Phase1d_validation.ipynb`.

7. Phase 2: metal line cooling — validation gate closed

Phase 2 brings the CHIANTI-based metal line cooling (Mg I/II, Ca II, Na I, Fe II, on top of the C/N/O coolants) into the coupled energy balance and **closes the Phase 2 validation gate** against Huang et al. (2023) Fig. 10. The atomic data for each coolant and the CHIANTI effective-cooling tables are documented in `Update_EXHALE_early_phase` (Parts II–III); this section records the integrated WASP-121b gate and the diagnostic used to close it.

7.1 Active line coolants

The line-cooling ions are flagged by `mion_iscool` in `species_table.f90` (canonical order): C I, C II, O I, O II, N I, N II, Mg I, Mg II, Ca II, Na I, Fe II. Each contributes $\beta_{\text{esc}} n_e n_{\text{ion}} \Lambda_{\text{ion}}(T)$ with the effective cooling coefficient Λ_{ion} interpolated from the hardcoded `cool_logL_*` tables in `Cool_coeff.f90` ($N_{\text{COOLT}} = 41$, $\log_{10} T = 3.0\text{--}5.0$, $\Delta \log_{10} T = 0.05$). In the optically thin limit $\beta_{\text{esc}} = 1$ (see `Update_EXHALE_early_phase`, Part II, “Adopted assumption: 100% escape”). N I/II carry the `iscool` flag but have no tabulated line cooling ($\Lambda = 0$), so they contribute exactly zero.

7.2 Channel-by-channel cooling diagnostic (this phase)

To close the gate quantitatively, `eval_cool` (`util_ion_eq.f90`) gained an optional `cool_chan` output — an **exact** decomposition of the total cooling rate into H/He recombination, collisional ionization, collisional excitation, bremsstrahlung, and one column per metal ion. It is read straight from the arrays `eval_cool` already builds, so the channel sum reproduces the cool total to machine precision (no offline re-derivation, no divergence risk). A companion writer `write_cool_breakdown_eq` (same module, called once from `EXHALE_main.f90` after the final equilibrium write) dumps output/`Cooling_breakdown.txt`: columns r/R_p , T , n_e , `cool_total`, the four H/He channels, then 27 metal-ion columns. The run prints `max |sum(channels)/cool - 1|` as an internal consistency check.

7.3 Validation on WASP-121b (the gate)

Converged equilibrium run (full metals, the curated `metals.inp`; same SED/grid as Phase 1d):

- Internal consistency: $\max|\sum \text{channels/cool} - 1| = 3.7 \times 10^{-16}$ (machine eps) — the breakdown is a bit-exact decomposition of the `Hydro_ioniz.txt` cool column.
- In the $1.15 < r/R_p < 1.4$ band the dominant radiative coolants are **Fe II** (~30%) and **Mg II** (~25%), jointly ~55%; metals are **57–85%** of all radiative cooling through the band. Fe II leads the inner edge (top coolant at $1.1\text{--}1.2 R_p$, 53% at $1.1 R_p$), Mg II takes over by $1.3\text{--}1.4 R_p$, and H recombination only overtakes above $\sim 1.5 R_p$ — the Fe II $\rightarrow$ Mg II handoff Huang report in their Fig. 10.
- $\dot{M}$ unchanged ($\log_{10} \dot{M} = 12.51$ g/s); zero NaNs.

This matches the Phase 2 gate (“a metal-driven cooling feature appears at $1.15\text{--}1.4 R_p$, Fig. 10”) and `Huang_update_expected_results` (Mg II + Fe II become the most important coolant in $1.15 < r/R_p < 1.4$).

7.4 Touched files

- `src/modules/radiation/util_ion_eq.f90` — optional `cool_chan` out-arg on `eval_cool` (exact channel-by-channel split); new `write_cool_breakdown_eq` writer; `mion_fsp` added to the `species_table` use-list.
- `EXHALE_main.f90` — use `utils_ion_eq`, only: `write_cool_breakdown_eq`; call it after the final equilibrium `write_output`.
- Cooling coefficients themselves live in `Cool_coeff.f90` (provenance in `Update_EXHALE_early_phase`, Parts II–III); unchanged here.

7.5 Caveats and deferred

- The ion-by-ion metal decomposition is exact in the default branch (`use_2lev_cool = .false.`, the WASP-121b setting). In the optional two-level fine-structure branch the terms for each ion are the resonance-line approximation and need not sum to the two-level `cool_M`.
- **Fe II** uses Boltzmann/coronal lower-level populations ($E_{\text{cut}} = 38459 \text{ cm}^{-1}$), validated vs Huang Fig. 6 coronal curve within $\sim 1.5\times$ — sufficient for the gate. The Hollenbach & McKee IR fine-structure + density-dependent two-level refinement (Huang “Sum” curve) is deferred.
- **Fe I** line cooling (NIST f -values + Van Regemorter, Huang Fig. 5) is not yet added; Fe I is a minor coolant in the gate band. (*Now added in §9; confirmed negligible — peak 3×10^{-5} of total cooling.*)
- The **CHIANTI Ly α swap** is intentionally skipped in Phase 2; the $0.35\times$ Ly α is a Phase-3 radiative-transfer suppression, not a cooling-rate change.
- Carrying metal cooling into the post-processed `_adv` profile is handled separately (runtime `pp_metals`, default frozen); see the post-process notes.

7.6 Status and relation to the Huang plan

Phase 2 is complete and the gate is closed (user-validated 2026-06-05). The metals-on thermal balance now reproduces the Huang Fig. 10 Fe II/Mg II cooling feature, closing the thermal-balance half of the network that Phase 1 left open. Next: **Phase 3** — in-code excited $H(n=2) + Ly\alpha$ radiative transfer.

8. Phase 3a: in-code excited $H(n=2) + Balmer$ proton source and heating

Added 2026-06-05.

Phase 3a is the first half of Phase 3: an **in-code $H(n=2)$ population model** whose output feeds back into both the hydrogen ionization balance (an extra proton source) and the energy equation (Balmer photoelectric heating), fully coupled. The second half (a $Ly\alpha$ radiative-transfer solve to supply a self-consistent mean intensity $\bar{J}_{Ly\alpha}(r)$) is deferred; this phase uses a **parameterized $\bar{J}_{Ly\alpha}$** with a switch to read an external RT profile instead. Everything is **default-off**: with no stellar parameters in `input.inp` the build is byte-identical to the Phase-2 result (`use_excited_H = .false.`).

8.1 Physics and atomic data (sources, not fabricated)

- **$H(2s)/H(2p)$ populations**: a 2×2 statistical-equilibrium system following Christie, Arras & Li (2013) — recombination cascade into $n=2$, collisional $2s \leftrightarrow 2p$ mixing, two-photon $2s \rightarrow 1s$ decay, $Ly\alpha$ $2p \rightarrow 1s$ with a $2p$ radiative-pumping term, and $2s/2p$ photoionization. Solved per cell (`n2_populations`).
- **$Ly\alpha$ mean intensity** (the $2p$ pump): parameterized as $\bar{J}_{Ly\alpha} \approx 0.1 F_{Ly\alpha} / \Delta\nu_D$ (Huang et al. 2017, Eq. 6) with a **top-down $1/(1 + \tau_{Ly\alpha})$ line-center attenuation** so the optically thick base ($\tau_{Ly\alpha} \sim 1.1 \times 10^8$ here) is not over-pumped. Oscillator strength $f_{Ly\alpha} = 0.4162$.
- **$n=2$ Balmer-continuum photoionization rate Γ_2** : from a **diluted stellar blackbody** above the $n=2$ threshold $E_2 = 3.40$ eV (`gamma_n2_balmer`), with $\sigma_{LyC} = 6.3 \times 10^{-18} \text{ cm}^2$ and threshold cross-section $s_2 = 1.4 \times 10^{-17} \text{ cm}^2$. Γ_2 carries the **same ξ dayside-dilution factor as ground-state EUV ionization** (`Rate/2` $\rightarrow 0.5$, `Rate/4` $\rightarrow 0.25$); for WASP-121b's “`Rate/2 + Mdot/2`” ($\xi = 0.5$) this gives $\Gamma_2 \approx 613 \text{ s}^{-1}$, comparable to Huang's $\approx 576 \text{ s}^{-1}$.
- **Heating**: Balmer photoelectric heating $H_{pe} = n_2 \cdot (\text{photoelectron energy} \times \Gamma_2)$ per cell (`heat_n2_balmer`); optional collisional de-excitation heating behind `incl_deexc_heat` (off by default — it overlaps the existing HI collisional-excitation cooling).

8.2 Coupling into the solver

The model runs **decoupled (lagged one outer step)** to keep it out of the cell-by-cell Newton solve: `excited_H_update` fills two global arrays (one entry per cell) from the previous converged state, and `ioniz_eq` injects them:

- `gph_balmer_HI [s-1]` is added to the HI photoionization rate `P_HI` — an extra **proton source** ($S_{\text{proton}} = \Gamma_2 n_2$).
- `heat_balmer [erg cm-3 s-1]` is added to the photoheating rate.

excited_H_update is called once per outer iteration in EXHALE_main.f90 before ioniz_eq, and once more after convergence to refresh the diagnostics. The post-loop refresh deliberately does *not* re-solve ioniz_eq (the Balmer heating is $\sim 10^{-4}$ of the total, not worth desyncing T/f_{sp} from the converged hydro state).

8.3 Runtime interface

Append to input.inp (omit for the default Phase-2-identical behavior):

Stellar Teff [K]: 6459.0

Stellar radius [R_sun]: 1.458

These set use_excited_H (both must be > 0). Optional further tail lines: Jlya RT file: <path> switches $\bar{J}_{\text{Ly}\alpha}$ to **mode 1** (read an external RT profile and use it directly, instead of the mode-0 parameterized field); Deexc heat: True enables collisional de-excitation heating. As with all trailing lines the scan is by keyword, so order and blank lines do not matter.

8.4 Diagnostics

write_excited_H dumps output/Excited_H.txt (14 columns) from the converged state: r/R_p , T , n_{HI} , n_e , $\bar{J}_{\text{Ly}\alpha}$, n_{2s} , n_{2p} , $S_{\text{proton}}(n=2)$, H_{pe} , H_{dx} , $\tau_{\text{Ly}\alpha}$, then the three **proton-budget** columns added in this phase — **col12 ground-state photoionization** ($P_{\text{HI}} n_{\text{HI}}$), **col13 collisional ionization** ($C n_e n_{\text{HI}}$), **col14 recombination** ($\alpha_B n_e n_{\text{HII}}$), all volumetric [$\text{cm}^{-3} \text{s}^{-1}$]. The three rate coefficients are captured inside ioniz_eq on the converged pass (globals gph_ground_HI, cion_HI, arec_HII), so the budget is self-consistent with the solver rather than re-derived offline.

8.5 Touched files

- init/parameters.f90 — Phase-3a globals (use_excited_H, T_star_eff, R_star, gamma2_bal, hpe2_bal, jlya_mode, jlya_rt_file); feedback arrays gph_balmer_HI, heat_balmer; diagnostics Jlya_arr, n2s_arr, n2p_arr, Sproton_arr, Hpe_arr, Hdx_arr; and the proton-budget capture arrays gph_ground_HI, cion_HI, arec_HII.
- radiation/excited_hydrogen.f90 (**new**) — excited_H_update, n2_populations, gamma_n2_balmer, heat_n2_balmer, load_jlya_rt, write_excited_H.
- radiation/ionization_equilibrium.f90 — inject gph_balmer_HI/heat_balmer under use_excited_H; capture the ground-state / collisional / recombination coefficients for the proton-budget dump.
- files_IO/input_read.f90 — scan the stellar / Jlya RT file / Deexc heat tail lines; set use_excited_H, jlya_mode.
- EXHALE_main.f90 — in-loop excited_H_update before ioniz_eq; post-loop refresh + write_excited_H.
- Makefile — excited_hydrogen.f90 in SRC.

8.6 Validation on WASP-121b (coupling ON vs OFF)

Two converged runs, same build, differing only in the `use_excited_H` toggle (stellar tail lines present/absent). Both reach steady state via $du < du_{th}$ with $\log_{10} \dot{M} = 12.51$ (unchanged from Phase 2 — the coupling does not destabilize the energy solver). Diagnostics in WASP-121b/{`analyze_excited_H.py`, `analyze_proton_budget.py`, `compare_on_off.py`}.

Proton budget (cf. Huang Figs. 11/27), volumetric rates [$\text{cm}^{-3} \text{s}^{-1}$]:

r/R_p	S_{n2}	S_{ground}	S_{recomb}	S_{n2}/S_{ground}	S_{n2}/S_{source}
1.10	5.1e4	6.6e6	6.6e6	0.008	0.008
1.30	1.2e4	4.3e5	4.4e5	0.029	0.028
1.50	4.9e3	7.6e4	8.1e4	0.064	0.060
1.80	1.9e3	1.5e4	1.7e4	0.125	0.111
2.00	2.0e3	5.9e3	7.9e3	0.34	0.25
2.01	2.9e3	5.1e3	7.9e3	0.57	0.36

Observations (for evaluation, not pass/fail):

- **Self-consistency:** at the base $S_{\text{ground}} \approx S_{\text{recomb}}$ (6.6×10^6 vs 6.6×10^6 at $1.1 R_p$), i.e. photoionization equilibrium holds; collisional ionization is negligible (T too low).
- **n=2 share:** the n=2 Balmer channel is a minor proton source below $\sim 1.5 R_p$ ($\sim 1\text{--}6\%$ of ground-state photoionization) but climbs to **$\sim 34\text{--}57\%$ of the ground-state rate ($\sim 25\text{--}36\%$ of the total proton source) near $2 R_p$** , where the gas is already ionized and EUV is attenuated (ground-state photoionization depleted) while the thin-Ly α -fed n=2 channel is not. This is the “significant proton source approaching $2 R_p$ ” behavior Huang describe, here concentrated near $2 R_p$ rather than flat across $1\text{--}2 R_p$.

Heating budget (cf. Huang Figs. 10/26): Balmer photoelectric heating is present in the budget, peaking at **$\sim 1.8\%$ of total heating** at the Ly α photosphere ($\tau_{\text{Ly}\alpha} \sim 1$ near $r \sim 2 R_p$) and $\sim 10^{-4}$ globally — a small but nonzero entry, dominating nowhere, as expected.

ON vs OFF temperature and ionization: the coupling makes the gas slightly **more ionized** (up to $\sim 24\%$ less neutral H near $2 R_p$, Δx_{HII} up to $\sim 1.1 \times 10^{-2}$) and slightly **cooler** (up to ~ 179 K at $r \sim 2 R_p$, mean ~ 42 K). The cooling is **physical, not numerical:** the extra n=2 photoionization converts $\text{HI} \rightarrow \text{HII}$, lowering n_{HI} ; the dominant EUV photoelectric heating scales with n_{HI} , so $\sim 24\%$ less n_{HI} at $r \sim 2 R_p$ means correspondingly less local EUV heating there, which the small Balmer H_{pe} does not offset.

8.7 Caveats and gaps

- **Parameterized $\bar{J}_{\text{Ly}\alpha}$ (mode 0).** The $0.1 F_{\text{LyC}}/\Delta\nu_D$ field with $1/(1 + \tau_{\text{Ly}\alpha})$ attenuation is a placeholder for the Phase-3b Ly α RT solve. Because our n=2 contribution is peaked near $2 R_p$ rather than flat across $1\text{--}2 R_p$, if Huang’s Figs. 11/27 are flatter/larger across the band the parameterized field is under-pumping n=2 in $1.1\text{--}1.8 R_p$; mode 1 (external RT $\bar{J}_{\text{Ly}\alpha}$) is the path to test this once a Ly α RT result exists.

- **Convergence stop is momentum-uniformity** ($du < du_{th}$), not a thermal/ionization steady-state criterion, so a small residual convergence-path sensitivity rides on top of the physical ON–OFF difference.
- **De-excitation heating off by default** (overlaps the existing H I coex cooling); left behind `incl_deexc_heat`.

8.8 Two analysis gotchas (recorded so they are not repeated)

- `Hydro_ioniz.txt` columns 6/7 (heat/cool) are **already physical** (`write_output` multiplies by q_0); do *not* multiply by q_0 again in analysis.
- `Ion_species.txt` column 1 is r/R_p , so $n_{HI} = \text{col2}$, $n_{HII} = \text{col3}$ — the recombination sink uses col3, not col2.

8.9 Status and relation to the Huang plan

Phase 3a is complete and the proton-source comparison is user-accepted (2026-06-05). The H(n=2) model is fully coupled into the ionization balance and the energy equation; the n=2 photoionization proton source reproduces the qualitative Huang Figs. 11/27 behavior and the Balmer photoelectric heating appears in the budget without destabilizing the solver. Next: **Phase 3b** — a $Ly\alpha$ radiative-transfer solve to replace the parameterized $\bar{J}_{Ly\alpha}$ with a self-consistent $\bar{J}_{Ly\alpha}(r)$ (feeding mode 1).

9. Phase 2 refinement: Fe I line cooling (NIST f -values + Van Regemorter)

Added 2026-06-05.

Fe I line cooling was one of the two coolants left deferred when the Phase 2 gate closed (§7, “Caveats and deferred”). This update adds it. CHIANTI v11 has **no Fe I model atom**, so — following Huang et al. (2023) §2.5 — Fe I cooling is built from NIST permitted-line oscillator strengths plus the Van Regemorter collision-strength approximation, with the lower levels populated in Boltzmann equilibrium over the metastable manifold. The new coolant flows through the same $\beta_{esc} n_e n_{FeI} \Lambda_{FeI}(T)$ channel as the other Phase-2 metals.

9.1 Atomic data and method (sources, not fabricated)

The cooling coefficient per $(n_e n_{FeI, total})$ in the optically thin limit [$\text{erg cm}^3 \text{ s}^{-1}$] is

$$\Lambda(T) = \sum_{\text{lines}} f_l(T) q_{lu}(T) \Delta E_{lu},$$

$$f_l(T) = \frac{g_l e^{-E_l/kT}}{U(T)}, \quad U(T) = \sum_{\substack{\text{even levels} \\ E < E_{\text{cut}}}} g e^{-E/kT}, \quad q_{lu} = 2.16 a^{-1.68} e^{-a} T^{-3/2} f_{lu},$$

with $a = \Delta E/kT$ and the absorption oscillator strength f_{lu} from the NIST A_{ki} , statistical weights, and wavelength (Huang Eq. 10).

- **Raw data:** the auditable NIST dumps `fe1_nist_lines.tsv` (permitted E1 lines: A_{ki} , g_i , g_k , E_i , E_k) and `fe1_nist_levels.tsv` (every level g , E for the partition function), fetched by `fetch_fe1_nist.py`.

- **Metastable manifold = EVEN-parity levels below the lowest odd level** (z^7D° , $19351\text{ cm}^{-1} = 2.40\text{ eV}$). Below that energy there is no lower level for an E1 decay, so the even levels are genuinely metastable and hold the Fe I population. This is *not* a plain energy cut: the Fe II treatment can use an energy cut because its cut conveniently lands at the first odd level, but for Fe I a pure energy cut would wrongly include the radiatively-depopulated z^7D° odd levels. The default cut yields **17 even levels and 828 permitted cooling lines**.
- Upper levels of the cooling lines are odd (E1) and assumed to decay radiatively. The builder is `cooling_data/fe1_cooling.py`; it emits the 41-point `cool_logL_FeI(NC00LT)` Fortran block on the ATES grid ($\log_{10} T = 3.0\text{--}5.0$, $\Delta \log T = 0.05$).

Coefficient values (per $n_e n_{\text{FeI}}$, $\text{erg cm}^3\text{ s}^{-1}$): $\Lambda(5 \times 10^3\text{ K}) = 1.27 \times 10^{-22}$, $\Lambda(10^4\text{ K}) = 1.23 \times 10^{-20}$, $\Lambda(2 \times 10^4\text{ K}) = 1.31 \times 10^{-19}$ — the rising, then flattening shape of Huang Fig. 5.

9.2 Touched files

- `cooling_data/fetch_fe1_nist.py`, `cooling_data/fe1_cooling.py` (**new**) — raw NIST fetch and the Boltzmann-manifold + Van Regemorter builder.
- `radiation/Cool_coeff.f90` — `cool_logL_FeI(NC00LT)` table constant; new `cool_FeI(T, out)` (interpolates the table via `interp_cool_table`); dispatch cases FeI in `cool_coeff_metal` and `cool_coeff_metal_scalar`. The table-header comment block now documents Fe I alongside the other coolants.
- `init/species_table.f90` — `mion_iscool(FeI) = .true.` (canonical ion index 25), so the `eval_cool` metal-cooling loop and the channel-by-channel breakdown both pick up Fe I automatically.

9.3 Validation on WASP-121b

Re-ran the converged full-metals WASP-121b equilibrium (same SED/grid as the Phase 2 gate) with Fe I cooling on; the channel-by-channel `Cooling_breakdown.txt` then gives the Fe I share directly. Diagnostics: `WASP-121b/analyze_FeI_cooling.py` (left panel = the ported $\Lambda_{\text{FeI}}(T)$ coefficient read straight from `Cool_coeff.f90` vs the Fe II coronal curve, Huang Fig. 5 style; right panel = cooling in each channel vs radius, Huang Fig. 10 style), figure `WASP-121b/FeI_cooling_validation.png`. Observations (for evaluation, not pass/fail):

- **Fe I is a negligible coolant** along the whole profile: its peak share of the total cooling is 3.0×10^{-5} at $r = 1.06 R_p$ ($T = 8303\text{ K}$), and it is far smaller elsewhere ($\sim 3 \times 10^{-8}$ of total at the base, where Fe II dominates). Adding Fe I does *not* perturb the gate-band Fe II $\rightarrow$ Mg II handoff.
- **Mdot unchanged** at $\log_{10} \dot{M} = 12.51$ (Fe I adds no measurable cooling), zero NaNs.
- This is the expected outcome: at WASP-121b base temperatures iron is almost entirely Fe II (Fe I fraction < 0.11 , §5), and Fe I's permitted lines have high excitation energies, so its Boltzmann-suppressed cooling cannot compete with the Fe II / Mg II forbidden + resonance lines. Fe I is carried for completeness and to mirror Huang's coolant list, not because it shapes the WASP-121b energy budget.

9.4 Status and relation to the Huang plan

Fe I line cooling is **implemented and its coefficient matches the Huang Fig. 5 shape/magnitude**; the WASP-121b run confirms it is radiatively negligible there, surfaced as a figure for the reader's judgment (no pass/fail asserted). One of the two Phase-2 deferred cooling items is now closed. The remaining one — the Fe II density-dependent (IR fine-structure) low-temperature refinement — is §10.

10. Phase 2 refinement: density-dependent Fe II line cooling (SE table)

Added 2026-06-05. (Implemented step-by-step; this section is updated as each step lands — currently: the 2-D table generator is built and verified.)

This is the second deferred Phase-2 cooling item (§7): replacing the **coronal** Fe II line-cooling coefficient with a **density-dependent** one that follows the Fe II level populations from the coronal (low- n_e) limit into the collisionally saturated (high- n_e , LTE) limit. It directly affects the cool, dense base of the WASP-121b thermosphere, where the coronal table overestimates Fe II cooling by $\sim 10^4 \times$.

10.1 Why the coronal table is wrong at the base

The Phase-2 Fe II table is the optically-thin **coronal** limit: only the ground level is populated and every collisional excitation is assumed to be followed by a radiative decay, so the cooling is strictly $\propto n_e$ (per $n_e n_{\text{FeII}}$ it is a function of T alone). That assumption fails at the cool, **dense** base of a hot Jupiter thermosphere, because the low- T Fe II cooling there is carried almost entirely by **forbidden** metastable / a⁶D fine-structure lines ($A \lesssim 5 \times 10^{-3} \text{ s}^{-1}$, far-IR 26–87 μm), whose critical densities

$$n_{\text{crit}} = A_{ul} / \sum_c n_c k_{ul}^c \sim 10^4\text{--}10^7 \text{ cm}^{-3} \quad (\text{electron collider})$$

sit far *below* the base $n_e \sim 10^8 \text{ cm}^{-3}$. Those lines are therefore collisionally **saturated** (thermalized to their LTE populations); the coronal coefficient overestimates their cooling by $\sim n_e / n_{\text{crit}}$.

10.2 The density-dependent coefficient (multilevel statistical equilibrium)

cooling_data/fe2_cooling.py solves the full CHIANTI multilevel statistical-equilibrium (SE) problem for the Fe II level populations at each (T, n_e) — 175 levels, electron-collisional excitation/de-excitation (Burgess–Tully descaled Υ from the .scups file) plus radiative decay (.wgfa A -values), normalized to $\sum_i n_i = 1$ — so the cooling smoothly interpolates between the two limits:

$$\frac{\Lambda_{\text{vol}}}{n_{\text{FeII}}} = \sum_{u>l} n_u A_{ul} \Delta E_{ul} \quad [\text{erg s}^{-1} \text{ per ion}], \quad \Lambda_{\text{eff}}(T, n_e) = \frac{\Lambda_{\text{vol}} / n_{\text{FeII}}}{n_e} \quad [\text{erg cm}^3 \text{ s}^{-1}].$$

At low n_e , $\Lambda_{\text{eff}} \rightarrow \Lambda_{\text{coronal}}(T)$; at high n_e it falls off $\sim 1/n_e$ as the numerator saturates. This Λ_{eff} is exactly the per- $(n_e n_{\text{FeII}})$ coefficient EXHALE already tabulates, so the existing assembly $\text{cool_M} = \beta_{\text{esc}} n_e \sum_i n_{\text{ion}}(i) \Lambda(i)$ carries it unchanged — the n_e in the prefactor times Λ_{eff} gives $\beta_{\text{esc}} n_{\text{FeII}} \times (\text{cooling per ion})$, i.e. the correct saturated base cooling.

Why electron-only SE is sufficient at the (mostly neutral) base. In the saturated limit the upper levels reach their **Boltzmann (LTE) populations regardless of which collider thermalizes them**, and the cooling of each ion becomes collider- and n_e -independent. At the base $n_e \sim 10^8$ already exceeds n_{crit} for the dominant forbidden lines, so the electron-only SE already reaches that LTE plateau; adding the (more abundant) neutral-H collider would only push the few still-unsaturated higher- n_{crit} lines further toward the same LTE limit — a second-order correction. An explicit H I / proton collider channel is flagged as a future refinement in the module.

10.3 Step 1 — the 2-D table generator (done)

`fe2_cooling.py` gained an **optimized** table builder `lambda_eff_table(Tgrid, negrid)` and Fortran emitters (`fortran_grid_1d`, `fortran_block_2d`, `emit_fortran_table`; run `python3 fe2_cooling.py --table`). The optimization is required because the naive cell-by-cell solve rebuilds a cubic spline for every one of the ~ 4300 transitions at every grid cell: instead the descaled Υ is evaluated **once per transition over the whole T-grid** (`upsilon()` accepts an array T), and the rate matrix is split $M = n_e Q_{\text{coll}}(T) + A_{\text{rad}}$, so the n_e -loop only reassembles and solves. The full 41×29 table builds in ~ 3 s and matches the naive cell-by-cell solve to 1.5×10^{-15} (machine precision).

The emitted table is on the ATES temperature grid (`NCOOLT` = 41, $\log_{10} T = 3.0\text{--}5.0$, $\Delta \log T = 0.05$) $\times$ an electron-density axis `cool_logne` (`NCOOLNE` = 29), $\log_{10} n_e = 0\text{--}14$, $\Delta \log = 0.5$ (brackets the WASP-121b run's $\log_{10} n_e \approx 7.9\text{--}9.7$ with the coronal plateau below and LTE saturation above). It is stored as `cool_logL_FeII_ne(NCOOLT, NCOOLNE)` = $\log_{10} \Lambda_{\text{eff}}$.

Two built-in sanity checks (printed to `stderr` by `--table`):

- **Low- n_e edge reproduces the coronal table.** At $n_e = 1 \text{ cm}^{-3}$ the SE coefficient equals `cooling_effective(pop='coronal')` to **0.1%** across $T = 10^3\text{--}10^5$ K — so swapping the 1-D coronal table for the $n_e = 1$ column of the 2-D table is a no-op, and the Phase-2 gate's coronal Fe II validation (vs Huang Fig. 6, within $\sim 1.5\times$) carries over unchanged.
- **Base suppression.** At the WASP-121b base ($T \approx 2239$ K, $n_e = 10^8$) the coronal coefficient 2.80×10^{-20} drops to the SE value 1.79×10^{-24} — a **15,663 $\times$ suppression**. The suppression is strongest at low T (forbidden lines) and weak at high T ($T \gtrsim 3 \times 10^4$ K, where the resonance lines with large A dominate and stay near-coronal): e.g. only $\sim 1.4\times$ at $T = 3 \times 10^4$, $n_e = 10^8$.

10.4 Step 2 — the 2-D table and bilinear interpolator in `Cool_coeff.f90` (done)

The generated table and its interpolators are now in `src/modules/radiation/Cool_coeff.f90`:

- **Table.** `integer, parameter :: NCOOLNE = 29`, the `cool_logne(NCOOLNE)` axis ($\log_{10} n_e = 0\text{--}14$, $d\log = 0.5$), and the data block `cool_logL_FeII_ne(NCOOLT, NCOOLNE)` (a `reshape([...], [NCOOLT, NCOOLNE])` with T on the inner index and n_e on the outer), pasted straight from `fe2_cooling.py --table`. The $n_e = 1$ column matches the existing 1-D `cool_logL_FeII` to the fifth decimal (first node -19.78002 vs -19.77975).
- **Bilinear interpolator.** `interp_cool_table_2d(logL2d, T, ne, out)` — log-log bilinear in $(\log_{10} T, \log_{10} n_e)$. The T axis reuses the uniform `cool_dlogT` scheme; the n_e axis uses the uniform `dlogne = 0.5`. Both axes clamp to the nearest edge outside the table (so n_e below the grid gives the coronal plateau, above it the LTE

branch). A scalar twin `interp_cool_table_2d_scalar(logL2d, Ts, nes)` mirrors it bit-for-bit for the cell-by-cell post-process solve.

- **Coolant wrapper.** `cool_FeII_ne(T, ne, out)` calls the 2-D interpolator on `cool_logL_FeII_ne`. The 1-D `cool_FeII` is kept as the low- n_e edge / scalar fallback; its header `!To Be Checked` note was rewritten to record that the density-dependent table now supersedes the coronal-only rate in the cooling assembly.

Builds clean (gfortran, full make); the new routines are additive and not yet wired, so behavior is unchanged until Step 3.

10.5 Step 3 — wiring the density-dependent Fe II into the cooling assembly (done)

The 2-D coefficient is now wired into both cooling paths so the local n_e selects the coronal $\rightarrow$ LTE-saturated branch automatically:

- **Main solver (`util_ion_eq.f90`, `eval_cool`).** n_e is already computed at the top of `eval_cool` (call `calc_ne`). Immediately after the canonical `c_metal(:, i)` fill loop, a short loop locates Fe II by name and overrides its column with `call cool_FeII_ne(T_K, ne, ...)`. Nothing else changes: the existing assembly $\text{cool_M} = \beta_{\text{esc}} n_e \sum_i n_{m,i} c_{\text{metal},i}$ then multiplies by n_e , which *cancels* the $1/n_e$ in $\Lambda_{\text{eff}} = (\sum_u n_u A_{ul} \Delta E_{ul}) / n_e$, leaving the correct cooling for each Fe II ion. The saturated (LTE) populations are collider-independent, so the electron-only SE solve is exact in that limit; at low n_e , Λ_{eff} reduces to the coronal rate, so optically-thin upper-atmosphere cells are unaffected. The override flows automatically into the `use_2lev_cool` branch (which already reads `c_metal(:, 26)`) and into the channel-by-channel `cool_chan` diagnostic (which reads `c_metal(:, i)`), so the breakdown stays an exact decomposition of `cool`.
- **Post-process cell-by-cell solve (`T_equation.f90`).** This path computes its own local n_e (from H, He, and the frozen metal stages) and sums the metal coolants with the scalar dispatcher. The Fe II term is special-cased to call the new `cool_FeII_ne_scalar(TT, ne)` (a thin wrapper over `interp_cool_table_2d_scalar`), so the post-process temperature balances the same density-dependent Fe II cooling.
- **`species_table.f90`.** The `!To Be Checked` Fe II comment was rewritten to state that `eval_cool` now applies the density-dependent SE coefficient (no longer “deferred”).

Builds clean (gfortran, full make, all four touched modules relink).

10.6 Step 4 — WASP-121b run and the cooling budget (numbers for the reader)

A fresh WASP-121b run (solar metals, `pp_metals 1`, otherwise the same `input.inp` as the Fe I-cooling run) was compared against the snapshot of the previous (coronal Fe II) run kept in `output_coronalFeII/`. Both share the same grid, so `Cooling_breakdown.txt` columns are differenced directly (`analyze_FeII_density_correction.py  $\rightarrow$  FeII_density_correction_compare.png`).

At the dense base ($r = 1.0$, $T = 2358\text{ K}$ fixed by the lower BC). The Fe II line cooling drops from the coronal 2.23×10^{-4} to the SE $1.71 \times 10^{-8} \text{ erg cm}^{-3} \text{ s}^{-1}$ — a $\approx 1.3 \times 10^4 \times$ suppression, consistent with the offline fixed- (T, n_e) check in Step 1 ($15,663 \times$ at $T = 2239\text{ K}$, $n_e = 10^8$). Because coronal Fe II carried **99.1%** of the base cooling, the *total* radiative cooling at the base falls $\approx 122 \times$ ($2.25 \times 10^{-4} \rightarrow 1.84 \times 10^{-6}$); Fe II’s share there is now **0.9%**. The previous run’s post-process printed “35 of 502 cells fell back to eq T (stiff base band)”; with the spurious base overcooling removed that message is **gone** — the cell-by-cell T solve no longer stiffens at the base.

Transition region 1.15–1.4 R_p (Huang Fig. 10). The band-integrated metal-line cooling share in the new run is

coolant	share
Mg II	44.0%
O II	18.9%
Fe II	13.4%
O I	7.3%
C II	6.3%
Ca II	5.1%

so **Mg II is now unambiguously the dominant metal coolant** across the transition region, with Fe II demoted to a $\sim 13\%$ contributor — the qualitative ordering Huang’s Fig. 10 shows. (In the old coronal run Fe II still topped the budget out to $\sim 1.2 R_p$.)

Global response. Removing the base overcooling raises the heating efficiency, so the steady-state mass-loss rate rises from $\log_{10} \dot{M} = 12.51$ to **12.80 g/s** ($\approx 2 \times$). The self-consistent wind is correspondingly denser and, above the base, slightly *cooler* (e.g. $T(1.05)$: $7810 \rightarrow 5744\text{ K}$) — the freed energy goes into mechanical escape and adiabatic expansion rather than being radiated, the expected energy-limited behavior. Whether this $\dot{M}$ and the transition-region budget match Huang’s WASP-121b quantitatively is left to the reader’s judgment against the paper; the figure and table above are surfaced for that comparison, not asserted as a pass.

Caveat on the figure’s panel (b). The “suppression factor” there is the *ratio of Fe II cooling between the two runs*, not a fixed-condition coefficient ratio. At the base (BC-fixed T) it equals the coefficient suppression ($\sim 1.3 \times 10^4$); above $\sim 1.2 R_p$ the coronal and SE coefficients converge (the hot resonance lines dominate and stay near-coronal), so the residual ratio there reflects the changed — denser — wind profile, not the coefficient.

11. Post-process refinement: re-solve metal ionization in the advected wind (pp_metals 2)

Added 2026-06-05.

The `_adv` profiles that `EXHALE_transit.py` reads are built by `post_process_adv.f90`, which re-solves the *temperature in each cell* at the advection-corrected structure (the wind is advected, then its $T(r)$ is re-converged against heat-

ing/cooling). Until now the metal *ion* densities carried into those profiles were handled in one of two ways, selected at runtime by `pp_metals` in `metals.inp` ($\rightarrow$ `pp_metal_mode`):

- `pp_metals 0` — **metal-free legacy** `_adv` (the pre-metals behavior).
- `pp_metals 1` — **frozen equilibrium** metals (the default): the converged equilibrium metal stages are carried into `_adv` unchanged, so they track the *equilibrium* structure, not the (slightly hotter, more ionized) advected wind.

This step implements the deferred third option:

- `pp_metals 2` — **re-solve**: the metal ionization balance is recomputed in each cell at the advection-corrected H/He and the post-process temperature, so the metal stage split tracks the advected wind self-consistently.

11.1 Implementation — reuse the coupled residual with H/He pinned

The metals are pure equilibrium species (no advection equations of their own): in the main loop they are solved *coupled* with H/He through `ion_system_HeH_metals + hybrd1` (photo + collisional ionization + recombination + charge exchange, no time derivative). “Re-solve” therefore means re-evaluating that same equilibrium balance at the post-process local T and the advected H/He — *not* a new physics path.

Rather than duplicate the residual, mode 2 reuses it verbatim through a thin wrapper, `ion_system_metals_pp` (new, in `post_process_adv.f90`):

```
subroutine ion_system_metals_pp(N_in,x,fvec,iflag,params)
...
call ion_system_HeH_metals(N_in,x,fvec,iflag,params) ! full coupled residual
fvec(1) = x(1) - pp_xHII_fix      ! pin HII fraction
fvec(2) = x(2) - pp_xHeII_fix     ! pin HeII fraction
fvec(3) = x(3) - pp_xHeIII_fix    ! pin HeIII fraction
end subroutine
```

The wrapper overwrites only the three H/He rows with identity equations that pin H/He to their advection-corrected fractions (stored in module-level `pp_xHII_fix` / `pp_xHeII_fix` / `pp_xHeIII_fix`), while leaving the $2n_{\text{elem}} = 20$ metal-stage rows untouched. So `hybrd1` moves only the metal unknowns; H/He stay fixed at the advected wind. Two properties fall out for free:

- **Self-consistent n_e .** The reused residual computes the electron density *internally* from the pinned H/He plus the metals it is currently solving — exactly as the equilibrium solver does — so the re-solve sees the same n_e (metals) feedback, not a frozen n_e .
- **Charge exchange with the advected H.** The Kingdon & Ferland $\text{H} \leftrightarrow \text{metal}$ couplings inside `ion_system_HeH_metals` see the *pinned* H^+/He^+ densities, so the metal balance reacts to the advected hydrogen ionization, not the equilibrium one.

11.2 The cell-by-cell sweep in post_process_adv

Mode 2 adds one sweep over cells, placed after the post-advection calc_ne and before the photoheating refresh, inside the same k -loop that re-converges T , so metals, H/He and T are driven to a *joint* fixed point. Per cell j (skipping $n_h(j) \leq 0$):

1. `cx_set_cell(T_K(j))` — charge-exchange rate coefficients at the local T .
2. Build the coefficients for each element (canonical order) from the existing post-process rate arrays: `meg_g0 = P_m` (photoionization), `meg_b0 = aion_m_pp` (collisional ionization), `meg_a1 = rec_m_pp` (recombination), and the second-stage $g1/b1/a2$ where `melem_top  $\geq$  2`; hand them to `set_metal_coeffs`.
3. Pin `pp_xHII_fix = nHII/nh`, `pp_xHeII_fix = nHeII/nhe`, `pp_xHeIII_fix = nHeIII/nhe`; set `params(7) = nh`, `params(8) = nhe`.
4. Seed `sys_x` with the pinned H/He fractions plus the current metal split and call `hybrd1(ion_system_metals_pp, . . .)`.
5. **Conserve element totals:** the pre-solve totals `nm_tot_pp(:,im)` (summed over each element's stages once, before the k -loop) are multiplied by the re-solved fractions to write back `nm_w`. Only the *stage split* changes; the element abundance is untouched.

Supporting edits: the use block pulls in `ion_system_HeH_metals`, `set_metal_coeffs`, and `cx_set_cell`; `params` was widened from `dimension(25)` to `dimension(60)` to match the residual's dummy length (only entries 1–11 are consumed by the metal solve); module-level `pp_x*_fix` and locals `nm_tot_pp / meg_* / i0`, `top`, `im` were added. Builds clean (gfortran, full make).

11.3 Clean A/B and what changes

Because the **equilibrium** solve in the main RK loop is mode-independent and deterministic, a full `pp_metals 2` run reproduces `Hydro_ioniz.txt`, `Ion_species.txt`, and `Cooling_breakdown.txt` **byte-identically** to the `pp_metals 1` run (verified with `diff -q`); *only* the `*_adv.txt` profiles differ. So the A/B isolates exactly the post-process metal treatment. The mode-1 default is preserved in `output_ppm1_dd/`, the mode-2 run in `output_ppm2_resolve/`, and the canonical `output/` is left at the mode-1 default (`metals.inp` reset to `pp_metals 1`); `analyze_metal_resolve.py` differences the two and writes `metal_resolve_compare.png`.

Singly-ionized fraction X^+/X_{tot} in the _adv wind (eq frozen $\rightarrow$ re-solve). Only the elements that move appreciably are tabulated (Mg and Na are unchanged to $\leq 2\%$); the full output for each element is in the analysis script.

r/R_p	T_{adv} [K]	O	N	Ca	Fe
1.05	5798	.00 $\rightarrow$.00	.01 $\rightarrow$.02	.64 $\rightarrow$.59	1.00 $\rightarrow$ 1.00
1.10	7984	.02 $\rightarrow$.01	.12 $\rightarrow$.13	.65 $\rightarrow$.60	.99 $\rightarrow$.99
1.15	9044	.10 $\rightarrow$.06	.38 $\rightarrow$.40	.66 $\rightarrow$.59	.99 $\rightarrow$.99
1.20	10035	.37 $\rightarrow$.26	.51 $\rightarrow$.47	.61 $\rightarrow$.54	.93 $\rightarrow$.96
1.30	10274	.34 $\rightarrow$.34	.35 $\rightarrow$.34	.38 $\rightarrow$.38	.39 $\rightarrow$.40
1.40	10680	.23 $\rightarrow$.23	.23 $\rightarrow$.23	.25 $\rightarrow$.25	.18 $\rightarrow$.19

Largest density change. The element-by-element scan of $\max |\Delta n(X^+)/n(X^+)|$ over the wind 1.0–1.48 R_p :

element	max rel. change	at r/R_p	$n(X^+)$ frozen $\rightarrow$ resolve
Fe	+9.5%	1.232	$1.19 \times 10^5 \rightarrow 1.30 \times 10^5$
Ca	−11.9%	1.191	$1.24 \times 10^4 \rightarrow 1.09 \times 10^4$
Na	−6.2%	1.000	$3.50 \times 10^6 \rightarrow 3.28 \times 10^6$
C	−2.1%	1.216	$1.22 \times 10^6 \rightarrow 1.19 \times 10^6$
Mg	−2.0%	1.000	$9.82 \times 10^7 \rightarrow 9.62 \times 10^7$

The corrections are **modest and concentrated in the 1.1–1.25 R_p ionization transition**, where T_{adv} runs a few hundred K hotter than T_{eq} : **Fe II** is *higher* across the 1.18–1.27 R_p shoulder (blue above red in the figure), **Ca II** is *lower* by ~ 5 –7 points through 1.05–1.25 R_p , and **every element converges back to the equilibrium split** in the outer wind ($\gtrsim 1.27 R_p$) where T_{adv} rejoins T_{eq} . The post-process printed **no** stiff-base T -fallback message in mode 2 (the density-dependent Fe II cooling of the previous section already removed the base overcooling).

Tiny-denominator artifacts (not physical). The same scan reports **O** +6538% at $r = 1.000$ and **N** +17.9% at $r = 1.013$. These are at the dense base where O and N are essentially *fully neutral* (O II fraction ≈ 0.00 in the table), so $n(\text{O II})$ is a vanishingly small number; a tiny absolute shift ($n^+ \approx 9.7 \times 10^2 \rightarrow 6.4 \times 10^4 \text{ cm}^{-3}$, against an O total $\sim 10^8$) reads as a huge *relative* change but is negligible in absolute and in any observable. They are denominator artifacts, not a physical re-solve effect.

Reader’s judgment, not a pass. The re-solve makes the metal stage split in the `_adv` profiles consistent with the advected wind rather than the equilibrium structure, with the physically expected sign (more ionization where the advected wind is hotter) and magnitude (single-digit-percent, transition-localized). The default remains `pp_metals 1` (frozen): the deliberate, documented choice (the equilibrium split is the converged one, and EXHALE_transit.py’s lines are dominated by H and He). Mode 2 is provided for studies that need the metal ionization to track the advected wind. The table and `metal_resolve_compare.png` are surfaced for that comparison, not asserted as a validation pass.

12. Phase 3b: in-code $\text{Ly}\alpha$ radiative transfer ($\bar{J}_{\text{Ly}\alpha}$ via escape probability)

Phase 3a drove the $\text{H}(n=2)$ population with a *parameterized* $\text{Ly}\alpha$ mean intensity $\bar{J}_{\text{Ly}\alpha} = 0.1 F_{\text{LyC}} / \Delta\nu_D / (1 + \tau_{\text{Ly}\alpha})$ (mode 0). That estimate peaks the $n=2$ population near $2R_p$ and *under-pumps* it across 1.1–1.8 R_p , leaving a `jlya_mode=1` hook (`load_jlya_rt`) to ingest a real RT field. Phase 3b supplies the real $\bar{J}_{\text{Ly}\alpha}(r)$. Three findings reshaped the original plan (a Monte Carlo with an offline two-pass and a β -on-cooling term):

1. **Method $\rightarrow$ escape probability, not Monte Carlo.** The WASP-121b column is extremely thick: $\tau_0(\text{Ly}\alpha) \approx 5.5 \times 10^7$ at the base (Voigt $a \approx 5 \times 10^{-4}$, so $a\tau_0 \approx 10^4 \gg 1$). A faithful acceleration-free MC needs $\sim \tau_0 \approx 10^7$ – 10^8 scatterings/photon (intractable), and core-skipping is both forbidden and incompatible with the core-dominated pumping field. For $a\tau_0 \gg 1$ the Neufeld (1990)/Harrington analytic wing-escape solution is accurate, so we use a local escape-probability source function.

2. **In-code, not offline.** Being a cheap formula, $\bar{J}_{\text{Ly}\alpha}$ is evaluated every timestep inside the run (`jlya_mode=2`, the same lagged-explicit decoupling as the Phase-3a Balmer feedback); a single run is self-consistent. The `jlya_mode=1` option to read an external (e.g. real MC) profile is kept.
3. **No β -on-cooling term.** From Huang §2.5/Fig. 7 directly: the “Ly α cooling 10 $\times$ lower” is the Black (1981) vs CHIANTI *collisional-excitation rate-coefficient* difference, not photon trapping. EXHALE already uses the CHIANTI-consistent Cen `coex_rate_HI` (§7), so there is no 10 $\times$ to apply. The trapping feedback is instead the Phase-3a H($n=2$) heating (photoelectric Hpe + collisional de-excitation Hdx), which strengthens once $\bar{J}_{\text{Ly}\alpha}$ is correct. The scaffolding β -on-coex hook was reverted.

12.1 The escape-probability field (`lya_rt.f90` : `jlya_escape_prob`)

Per cell, from the converged T , $n_{\text{HI}} = n_{1s}$, n_{HII} , n_e :

- $\Delta\nu_D = \nu_{\text{Ly}\alpha} \sqrt{2kT/m_{\text{H}}}/c$; Voigt $a = (A_{2p1s}/4\pi)/\Delta\nu_D$; optical depth $\tau(r) = \int_r^\infty n_{1s}\sigma_0 dr'$, $\sigma_0 = C_{\text{Ly}\alpha}/\Delta\nu_D$.
- Wing escape probability $\beta = \pi^{-1/4} \sqrt{a/\tau}$ (≤ 1): escape once the photon diffuses to $x_* = (a\tau/\sqrt{\pi})^{1/2}$ where the wing optical depth is unity.
- Internal sources $P = \alpha_B n_e n_{\text{HII}} + q_{1s2p} n_e n_{\text{HI}}$ (recombination cascade + electron-impact $1s \rightarrow 2p$). The escape-probability closure $\bar{J} = S(1 - \beta)$ with the trapped pile-up $n_{2p} = P/(A_{2p1s}\beta)$ gives $J_{\text{int}} = (2h\nu^3/c^2)(g_{1s}/g_{2p})P(1 - \beta)/(A_{2p1s}\beta n_{1s})$. The $(1 - \beta)$ factor makes $J_{\text{int}} \rightarrow S$ when thick and $\rightarrow 0$ when thin ($\beta \rightarrow 1$), instead of $S \propto 1/n_{1s}$ diverging as the gas ionizes in the outer wind (which spuriously raised $\bar{J}$ outward).
- Stellar beam (trapped): Ly α is *scattered*, not destroyed, so it is not killed by $e^{-\tau}$; the broad stellar line (half-width $\Delta v_s \approx \pm 70 \text{ km s}^{-1}$) penetrates through its wings. A photon at offset x reaches r if $|x| > x_1 = \sqrt{a\tau/\sqrt{\pi}}$; for a Gaussian stellar profile of width $X_s = \Delta v_s/v_{\text{th}}$, $T_*(r) = \text{erfc}(x_1/(\sqrt{2}X_s))$. The profile-averaged mean intensity uses the *stellar* line width $\Delta\nu_* = \nu_0 \Delta v_s/c$ (constant), NOT the planetary $\Delta\nu_D$ ($\bar{J} = \int J_\nu \phi d\nu \approx J_\nu(\nu_0)$ for a stellar field flat across the narrow planetary line; dividing by $\Delta\nu_D$ makes J_{star} rise as the wind cools). A bounded buildup $E = 1 + (\text{boost} - 1)(1 - \beta)T_*$ enhances the penetrating beam in the thick wind and $\rightarrow 1$ when thin or at the deep edge, avoiding both the $1/\beta$ over-count (which blew the $n=2$ source to 6300 $\times$ ground-state) and a flat cap. Hence $J_{\text{star}} = \xi F_* T_* E / (4\pi \Delta\nu_*)$, ξ the dayside dilution. $\bar{J}_{\text{Ly}\alpha} = J_{\text{int}} + J_{\text{star}}$.

New `input.inp` lines (default-off $\Rightarrow$ byte-identical to 3a; a 3b-0 scaffolding step verified this, frozen reference `output_pre3b_baseline/`):

```
Jlya escape-prob: True           # jlya_mode = 2 (in-line escape-prob RT)
Stellar Lya flux [erg/cm2/s]: 1.0e5 # incident stellar Lya at the planet (x0.35)
Lya stellar halfwidth [km/s]: 70   # broad stellar line Dv_s (penetration + Dnu_star)
Lya stellar boost [-]: 5           # bounded trapping-buildup cap (left un-tuned)
Deexc heat: True                  # Phase-3a collisional de-excitation heating
```

12.2 WASP-121b validation setup (spherical Case A, matched to Huang)

Huang’s Fig. 11 is **Case A** (spherical, no RLOF), so the comparison must match his geometry *and* radial frame — which exposed a setup error: the run had been Roche-truncated at $\sim 2R_p$, and EXHALE’s base ($r=1$) sits at $1\mu\text{bar}$ ($n_0 = 3.1 \times 10^{12}$, $T = 2358\text{ K}$), but Huang defines R_p at the **4 mbar transit radius** and places his $1\mu\text{bar}$ hydro bottom at $\sim 1.25R_p$ (Fig. 9 top-pressure axis / Fig. 12 green band; the earlier “ $1.46R_p$ ” note was wrong). The setup was corrected to spherical Case A with R_0 set to Huang’s $1\mu\text{bar}$ radius ($R_0 = 1.25 \times 1.766 = 2.2075R_J$, planet mass unchanged $\Rightarrow$ correct base gravity) and outer radius 15.2 ($= 19R_p^{\text{Huang}}$); EXHALE r maps as $R_{\text{Huang}} = 1.25r$. The Case-A $\dot{M} = 0.077M_p/\text{Gyr}$ is within $\sim 1.5\times$ of Huang’s 0.052 (Table 3).

12.3 Trend vs Huang Fig. 11 (qualitative — exact match not expected)

Huang’s $\bar{J}_{\text{Ly}\alpha}$ is a **plane-parallel Monte Carlo** RT; ours is an analytic escape-probability formula, so the two cannot agree quantitatively — a quantitative comparison is only sensible once *all* Phase-3b features (incl. the velocity term) are in. Qualitatively (`lya_caseA_huang11.py`), both show a significant inner peak (EXHALE ~ 0.6 at 1.35, Huang ~ 0.5 at $1.6R_p$) and comparable magnitude (0.05–0.6). EXHALE has a dip at $\sim 1.8R_p$ and rises gently outward where Huang gently declines; the outward rise traces to the $n2s$ recombination cascade in the Phase-3a $\text{H}(2s/2p)$ model ($n=2$ stays recombination-populated while ground-state H vanishes in the ionized wind), *not* $\bar{J}$, which after the two fixes below is physically sound (J_{int} peaks then declines; $J_{\text{star}} \approx \text{const}$ outward). $\text{H}(n=2)$ heating stays $\sim 1\%$ of total with $\text{Hdx} \gtrsim \text{Hpe}$ near the n_2 peak (cf. Fig. 10).

Two $\bar{J}$ bugs fixed during the trend check. (1) J_{star} used the planetary $\Delta\nu_D$ instead of the stellar width $\Delta\nu_*$, making $\bar{J}$ rise as the wind cooled; (2) J_{int} lacked the $(1-\beta)$ closure, so S diverged ($\propto 1/n_{1s}$) in the ionized wind.

Wind-velocity (Sobolev) escape — implemented, negligible here. $\beta_{\text{tot}} = 1 - (1 - \beta_{\text{static}})(1 - \beta_{\text{Sob}})$ with $\tau_S = C_{\text{sob}}n_{1s}/|dv/dr|$ is now in `jlya_escape_prob` (the outflow $v(r)$ threaded through `excited_H_update`). For WASP-121b it changes $\bar{J}$ by $< 1\%$: where it could help (thick inner wind) $\tau_S \sim \tau_{\text{static}} \sim 10^7$ so $\beta_{\text{Sob}} \sim 10^{-7}$, and the thin outer wind already has $\beta_{\text{static}} \approx 1$ — the velocity gradient and the $\text{Ly}\alpha$ opacity scale together ($\propto n_{1s}$), opening no new escape route. **Phase 3b is thus feature-complete.** *Deferred per the trend-first guidance:* the post-all-features quantitative comparison and the $n2s$ outward-rise revisit (a Phase-3a $2s$ effect, not the RT); boost, Δv_s , and the $\pi^{-1/4}$ prefactor are left un-tuned.

13. Phase 4: RLOF / tidal potential (Case B validated; Case D blocker = escape radius, transonic IC added, geometry-limited)

Phase 4 adds the Roche-lobe-overflow (RLOF) tidal physics that drives Huang’s enhanced mass loss — **Case B** $\dot{M} = 0.32$, **Case D** $= 1.03M_p/\text{Gyr}$ — the regime that matches the observed transit depths (spherical Case A underestimates them). EXHALE already has the **1D Roche potential** (`grav_field.f90`, the `spherical_domain=.false.`

branch: planet + stellar-tidal + centrifugal), the L1-truncated domain, and the tidal momentum source (Source.f90) — Caldiroli’s ATES-v2 Roche mode. So the substellar tidal hydrodynamics is already in place.

13.1 Base gravity for each case (Huang Table 3)

The key to each case is the $1\ \mu\text{bar}$ base radius R_0 . Huang’s Table 3 lists $\log g @ 1\ \mu\text{bar} = GM_p/r_{1\mu\text{bar}}^2$ per case — exactly the R_0 -setting quantity ($R_0 = \sqrt{GM_p/g}$):

Case	tidal	$M_p(M_J)$	dT	XUV/Ly α /Fe	$\log g$	$R_0(R_J)$
A	N	1.1824	0	1/1/1	2.84	2.058
B	Y	1.1824	0	1/1/1	2.70	2.418
D	Y	1.1204	350	0.74/0.35/4	2.62	2.581

R_0 grows A $\rightarrow$ B $\rightarrow$ D as the tidal force + the dT+350 lower-atmosphere heating puff the $1\ \mu\text{bar}$ level outward (weaker base gravity), driving the RLOF.

13.2 Case B — validated

WASP-121b in Roche mode (tidal on) + the corrected base + metals + Ly α mode-2 gives $\dot{M} = 0.337 M_p/\text{Gyr} \approx$ Huang’s Case B (0.32), within $\sim 5\%$. With the base-radius fix both cases agree (Case A $0.077 \approx 1.5 \times$ Huang’s 0.052; Case B $0.337 \approx 1.05 \times$ Huang’s 0.32). EXHALE’s existing Roche potential thus captures the RLOF enhancement — no geometric-reconstruction code was needed for the Case B $\dot{M}$ gate.

13.3 Case D — real blocker found (escape radius), transonic IC added

Case D’s base ($R_0 = 2.581 R_J$, $\log g = 2.62$) sits at $\sim 76\%$ of the L1 distance: the L1-truncated Roche domain reaches only $r_{\text{max}} = 1.353 R_p$. Re-investigating the earlier “won’t launch” report uncovered the *actual* blocker, and it is *not* the initial condition.

Root cause — the escape radius lies outside the L1-truncated domain. `define_grid` sets the constant-momentum convergence window $[j_{\text{min}} : N]$ from the **escape radius** (Escape radius `[R_p] :`, default 1.50). It scans for the first cell with $r(j) \geq r_{\text{esc}}$; if *no* cell qualifies — which happens precisely when $r_{\text{esc}} > r_{\text{max}}$ — the DO loop falls through leaving $j = N + N_g + 1$, so $j_{\text{min}} > N$ and the window is **empty**. Fortran `maxval/minval` over a zero-size array return $\mp\text{HUGE}$, so $\text{mom}_{\text{max}} = -\text{HUGE}$, $\text{mom}_{\text{min}} = +\text{HUGE} \Rightarrow du = \text{Inf}$, and the time-derivative norm $dtu = -\text{HUGE}$. Since $dtu = -\text{HUGE} < dtu_{\text{th}}$, the loop reports **“steady state reached” at step 0** — a *false* convergence unrelated to the wind. The earlier bracket (“works $R_0 \leq 2.30$, fails $R_0 \geq 2.40$ ”) is exactly r_{max} crossing the 1.50 escape radius at $R_0 \approx 2.37$, confirming the escape radius, not the IC, as the discriminant. (The earlier “isothermal IC cannot launch the deep-RLOF wind” diagnosis was therefore wrong.)

Fixes. (1) A guard in `define_grid.f90`: if $j_{\text{min}} > N$, re-anchor the window to mid-domain ($r \geq 1 + \frac{1}{2}(r_{\text{max}} - 1)$) and print a loud WARNING — so this can never again masquerade as convergence. (2) For Case D, set Escape

radius [R_p]: 1.20 (inside the 1.353 domain). With a valid window the diagnostics become finite and meaningful, and **Case D launches and runs** — from *both* the hydrostatic and the transonic IC. (The earlier $du = |(\text{mom}_{\text{max}} - \text{mom}_{\text{min}}) / \max(\text{mom}_{\text{min}}, 10^{-30})|$ floor at EXHALE_main.f90:308 is retained as defense-in-depth against a genuine transient zero-momentum cell, but it was *not* what unblocked Case D — with an empty window $\text{mom}_{\text{min}} = +\text{HUGE}$, not zero, so the floor never engaged.)

Transonic-wind IC (implemented, Transonic IC: True). set_IC.f90 gained an optional transonic isothermal-wind profile, solved from the EXHALE potential via the algebraic (Bernoulli) integral of the steady isothermal-wind equation,

$$\frac{1}{2}v^2 - c^2 \ln v - 2c^2 \ln r + \phi(r) = B, \quad c^2 = (1 + \delta p_{bc}) / \rho_{bc},$$

with B fixed by the sonic point ($v = c_s$ where $D\phi(r_c) = 2c^2/r_c$, located by bisection). At each radius the relation has a subsonic ($v < c_s$) and a supersonic ($v > c_s$) root; the wind takes the subsonic root for $r < r_c$ and the supersonic root for $r > r_c$ (1-D bisection on each monotone branch), and ρ then follows from mass conservation $\rho v r^2 = \dot{M}$ pinned to the base BC. Because that IC already satisfies $\rho v r^2 = \text{const}$, $du \approx 0$ at step 0 would itself trip the momentum-constant exit before the cold wind heats; the flag therefore **auto-enables force_start**. If no interior sonic point exists it **falls back to the hydrostatic IC**. Default **off** $\Rightarrow$ existing cases byte-identical (Case B re-verified at $\dot{M} = 0.337$).

Case D outcome — launches, but the 1-D L1-truncated domain undershoots. With the escape-radius fix the wind launches and relaxes, but it **never reaches a sonic point inside the domain**: the heated metal wind is **deeply subsonic everywhere** (Mach $\lesssim 0.05$ — rising to ~ 0.050 at $r \approx 1.25$, then *falling* to ~ 0.040 at L1). With a subsonic outer boundary the zero-gradient outflow BC is only marginally well-posed and $\rho v r^2$ stays non-constant (peaks $\sim 1.7\times$ the L1 value), so the run settles into a **breathing limit cycle**: after the $\sim 10^2$ heating transient damps, du keeps oscillating in $\sim 0.04\text{--}0.8$ (touching but never crossing the strict $du_{\text{th}} = 0.02$) for as long as it runs. The mass flux is correspondingly time-variable — the quasi-steady $\dot{M}$ oscillates over $\sim 0.3\text{--}0.6 M_p/\text{Gyr}$, well *below* Huang's 1.03, the deficit set by the L1 truncation cutting off the still-accelerating subsonic wind before it reaches the sonic point. The transonic IC behaves identically to the hydrostatic IC here: both launch the same subsonic flow. (Warm-starting via load_IC reintroduces a base NaN — the documented grid re-read inconsistency — so it is not a workaround.)

Conclusion. Case D's deep RLOF ($\dot{M} = 1.03$) is **not reachable in EXHALE's 1-D L1-truncated Roche mode**: the wind is subsonic at L1, so the truncated domain structurally undershoots. Capturing it needs the 3-D geometric treatment (Phase 5) or a domain extended past L1 — *not* a different IC. The transonic IC and the escape-radius guard are correct, necessary infrastructure (they remove the false convergence and let the deep-RLOF wind launch and run with meaningful diagnostics) but are not sufficient on their own for Case D's full mass-loss rate. The geometric reconstruction (substellar $\rightarrow$ terminator mapping, triaxial radii Table 2, the 4/9 mass-loss conversion) feeds Phase 5 (§14), and also supplies the geometry this subsonic L1-truncated wind needs.

14. Phase 5: velocity-broadened transmission spectrum (in progress)

Phase 5 is the headline deliverable — the NUV/optical **transit-depth spectrum** the whole upgrade exists to match (Huang Table 3 / Figs. 17–24). It extends the existing spherical post-processor EXHALE_transit.py in two stages: **5a** adds the metal resonance lines (spherical, gated on Case A); **5b** replaces the spherical geometry with the 3-D Roche equipotential reconstruction + velocity broadening (gated on Case D). Pure Python post-processing — no Fortran change.

Line	Case A	Case D	Observed
Mg II $\lambda 2796$ (4 Å)	0.182	0.302	0.309 ± 0.036
Ca II K $\lambda 3934$	0.199	0.278	0.281 ± 0.009
H α $\lambda 6563$	0.201	0.185	0.186 ± 0.003
H β $\lambda 4861$	0.174	0.135	0.143 ± 0.005
Na D2 $\lambda 5890$	0.152	0.147	0.147 ± 0.002

These are the **effective transit radius** R_{eff}/R_* , not an absorption percent (the expected-results doc states the unit is R_p/R_*). Our code returns the line-center absorption fraction h , mapped to the same quantity via $R_{\text{eff}}/R_* = \sqrt{(R_p/R_*)^2 + h}$.

14.1 5a — metal resonance lines (Mg II, Ca II, Na I D), spherical — done

EXHALE_transit.py gained Mg II $\lambda 2796$, Ca II K $\lambda 3934$, Na I D2 $\lambda 5890$ as an **isolated appended block** (the validated He 10830 / Ly α / H α / H β pipeline above it is untouched). These are resonance lines whose lower level is the ion ground state; at $\sim 10^4$ K the excited fine-structure levels are Boltzmann-negligible, so the lower-level density is the ion density itself ($n_{\text{lower}} \approx n_{\text{ion}}$). The ion densities live past the H/He block of Ion_species_adv.txt, in mion_fsp order (species_table.f90): numpy columns **17 = Mg II**, **23 = Ca II**, **25 = Na I**. The block reuses the existing spherical chord / Voigt (wofz) / trapezoid- τ / disk-average / instrument convolution path (resonance_depth()).

Case A result (line-center % / 4 Å-band %) and diagnosis. Mg II: 26.1 / 2.3 ($R_{\text{eff}} = 4.2 R_p$); Ca II K: 5.5 / 0.33 ($1.9 R_p$); H α : 3.9 ($1.6 R_p$); Na D2: 0.8 / 0.06 ($0.7 R_p$). Two real gaps: (1) **Mg II already reaches Huang's extent** ($R_{\text{eff}} = 4.2 R_p$ vs his clustered $\sim 3.6 R_p$) — the metal treatment is sound — but **Ca II / Na I / H α are too optically thin** ($R_{\text{eff}} = 0.7\text{--}1.9 R_p$), a line-specific ionization / abundance / $n = 2$ question, not geometry. (2) The lines are **too narrow**: the Case A wind reaches only ~ 29 km/s, so a 26%-deep core washes out to $\sim 2\%$ over a 4 Å (± 215 km/s) bin, whereas Huang's lines stay deep *across* the bin — they are velocity-broadened to hundreds of km/s by the fast Roche outflow + rotation, i.e. **Phase-5b physics**. (The existing H α tool shows the same shortfall, so it is systematic.) *Units — resolved*: Huang's table is R_{eff}/R_* (radius ratio; the expected-results doc header states the unit is R_p/R_*), not absorption percent. The Case A spherical run over-extends ($15 R_p$ domain $\Rightarrow$ Mg II reaches $4.2 R_p$, $R_{\text{eff}}/R_* = 0.51$ vs Huang 0.182); the proper test is the L1-confined Roche run in R_{eff}/R_* units (§5b).

14.2 5b — Roche-equipotential reconstruction + broadening — implemented

New module `roche_recon.py`. Assuming a uniform state on Roche-equipotential surfaces, the field at any (x, y, z) equals the 1-D substellar state at the equivalent radius r_{eff} where $\phi_{\text{sub}}(r_{\text{eff}}) = \phi_{3\text{D}}(x, y, z)$. The dimensionless Roche potential (R_p units) mirrors `grav_field.f90`:

$$\phi(x, y, z) = -\frac{1}{r} - \frac{q}{r_s} - \frac{1+q}{2a^3} [(x-x_{\text{cm}})^2 + y^2], \quad q = M_*/M_p, \quad a = \text{orbit}[R_p].$$

Tidal terms $\rightarrow 0$ recovers $\phi = -1/r$ (spherical), so the triaxial reconstruction reduces to the 5a spherical case for Case A. The module gives `roche_phi`, the L1 root of $d\phi_{\text{sub}}/dr = 0$, a monotone $\phi_{\text{sub}} \rightarrow r_{\text{eff}}$ inverse (`ReconMap`), the triaxial radii, and the Eggleton lobe radius. `EXHALE_transit.py` gained a `geometry='triaxial'` switch (`triaxial_depth()`): the transit LOS runs along $+x$; the state at each 3-D point is the substellar state at r_{eff} ; and the **LOS velocity in each sector** $v_{\text{sub}}(r_{\text{eff}})x/r - \Omega y$ (wind + tidally-locked rotation, Huang Eq. 16) is integrated over **20 angular sectors**.

Demonstrated on the converged Case B (Roche) run. Geometry self-consistent: $L1 = 1.574 R_p$; $R_{\text{RL}}/L1 = 0.714 \approx 2/3$; tidally triaxial $R_{px} = 1.05 > R_{py} = 0.919 > R_{pz} = 0.893$ (substellar bulge, polar compression).

Result — the model reproduces Huang’s transit radii. Converting the Case B triaxial depths to R_{eff}/R_* (the table unit):

Line	Huang A	Huang D	Case B (this work)
Mg II	0.182	0.302	0.234
Ca II	0.199	0.278	0.222
Na D2	0.152	0.147	0.163
H α	0.201	0.185	0.239
H β	0.174	0.135	0.232

Case B is the *intermediate* model (between Case A and Case D), and the **metals (Mg II, Ca II, Na) land right between Huang’s Case A and Case D** — the lobe-confined reconstruction correctly reproduces his transit radii in magnitude and trend. The Balmer lines run $\sim 25\%$ high (H α 0.239 vs the expected ~ 0.19), consistent with the 5a finding that $n = 2/T$ is slightly over-predicted.

Correction. An earlier draft claimed the lobe-confined model “undershoots Huang because it omits the beyond-L1 escaping tail.” That was wrong twice over: (1) Huang’s reconstruction is *also* lobe-confined (equipotential mapping to the Roche-lobe boundary, §1.8) — neither model has a beyond-L1 tail, so it cannot be the difference; and (2) the apparent shortfall was a **units error** (our absorption *percent* vs Huang’s R_{eff}/R_*). In the correct units the two agree (table above). The small triaxial-vs-spherical reduction (terminator compression $R_{py} < R_{px}$) is minor; Huang’s area-preserving $R_{py}R_{pz} = R_p^2$ normalization would offset it.

14.3 5b — direct Case A comparison: hydro validated, metals over-extended

A clean spherical Case A at the *correct* 1 μbar base $R_0 = 2.058 R_J$ (Table 3 $\log g = 2.84$; the old `output_caseA` used $R_0 = 1.766$, the 4 mbar transit radius) was run to convergence.

- **Hydro / base validated:** $\dot{M} = 0.056 M_p / \text{Gyr} \approx \text{Huang's Case A } 0.052$ (the wrong base gave a different $\dot{M}$) — base and wind are right.
- **Metal transit radii — consistent once compared like-for-like** (Huang quotes Mg II in a $4 \text{ \AA}$ NUV bin; Ca II / Na are optical line-center): Mg II ($4 \text{ \AA}$) **0.220** vs Huang 0.182 ($\times 1.2$); Na D2 (LC) **0.175** vs 0.152 ($\times 1.15$); Ca II K (LC) **0.306** vs 0.199 ($\times 1.5$, high).

The “over-extension” was a comparison-definition artifact, not physics. An earlier draft reported Mg II $R_{\text{eff}}/R_* = 0.60$ ($3.3\times$ Huang) — that compared the *line-center* R_{eff} against Huang’s $4 \text{ \AA}$ NUV-bin value. In the same $4 \text{ \AA}$ bin Mg II $= 0.22 \approx \text{Huang's } 0.182$; Na matches too; only Ca II is genuinely $\sim 1.5\times$ high.

The metal ionization is correct and is NOT the lever. All channels are wired and physical (Verner+1996 σ , Badnell RR+DR, Voronov, Huang+2023 charge exchange; $a_\tau = 0$). The Mg II/III balance is a close competition between a weak photoionization ($\Gamma \approx 2 \times 10^{-3} \text{ s}^{-1}$, weak because $\sigma_{\text{MgII}} \approx 0.24 \text{ Mb}$ at threshold) and RR- n_e ; DR ($\sim 5 \times 10^5 \text{ K}$ resonances) and collisional ionization ($e^{-15 \text{ eV}/kT}$ at 10^4 K) are negligible. A test showed $\times 30$ more Mg II photoionization moves R_{eff} only $0.57 \rightarrow 0.32$ — the resonance line is so optically thick that ionizing it harder cannot confine it, so the ionization is a dead end as a lever (and is not the cause once the bin is matched).

Status: the full Phase-5 pipeline is implemented and runs end-to-end; the hydro/ $\dot{M}$ reproduces Huang (Case A 0.056, Case B 0.337); and the metal transit radii match Huang once compared in his own bins (Mg II $4 \text{ \AA}$ 0.22 vs 0.182, Na 0.175 vs 0.152). Remaining items are minor: Ca II $\sim 1.5\times$ high (line-center; plausibly Ca II $\rightarrow$ Ca III or the Ca abundance), the slightly-high Balmer $n = 2$, the $R_{py}R_{pz} = R_p^2$ normalization (methodological), a terminating Case D run, and phase5_transmission.ipynb. Comparisons must use Huang’s line-by-line definition ($4 \text{ \AA}$ bin for the NUV metals, line-center for the optical lines) — the two differ by $\sim 3\times$ for the optically-thick Mg II.

15. Post-process advection guard at a “breathing” base (2026-06-07)

Symptom. The spherical-vs-Roche example figure (user manual Fig. 4, WASP-121b) showed a sharp temperature spike at $r \approx 1.1 R_p$ in the *advection-corrected* (`_adv`) Roche profile, with a sawtooth band just below it. The converged equilibrium (eq) profile is smooth there, so the spike is a post-processing artifact, not a feature of the hydro solution.

Diagnosis (two compounding causes).

1. *A breathing, inflowing base (1-D limitation).* WASP-121b nearly fills its Roche lobe, so the 1-D wind is sub-sonic at L_1 and the dense lower atmosphere settles into a slowly recirculating state with small *negative* (inflow) velocities and a stagnation point ($v = 0$) near the wind base ($r \approx 1.1$). This is intrinsic to spherical symmetry, not a bug.
2. *A spurious hot root in the metal-cooled post-processor.* The advection post-processor re-solves the energy/ionization balance in each cell assuming an outflow (each cell upwinds from the next-inner cell). In the dense base the non-monotonic metal line-cooling curve gives the energy equation a second, spurious *hot* root that the Newton (MINPACK) solver can land on; the upwind coupling then cascades one bad cell outward, producing the sawtooth and the overshoot spike. The pre-existing $2\times$ band guard around eq T did not catch the $\sim 1.4\times$ spike.

Fix (option c: eq fallback where the flow is not a clean outflow). In `post_process_adv.f90`, wherever $v \leq 0$ (the breathing inflow base) the post-processor now *skips* the advection correction and keeps the converged eq ionization and temperature. Three guards were added, all gated on `pp_metal_on` so the metals-off mode stays byte-identical: the no-He and He H/He advection loops (set the H/He fractions to eq and cycle when $v(j) \leq 0$ or $v(j-1) \leq 0$) and the cell-by-cell temperature solve (set $T_{\text{out}} = T_{\text{in}}$ and cycle when $v \leq 0$). Keeping the inner cells at eq removes the spurious root *and* breaks the upwind cascade at its source, while the genuine `_adv` correction is retained in the outflow above the stagnation point.

Validation (PP-only re-run on the converged Case B state). Rebuilt and re-ran the post-processor (Do only PP on the converged IC):

- eq reproduced: $\max |\Delta T|/T = 2.6 \times 10^{-3}$, $\max |\Delta v| = 0.015 \text{ km s}^{-1}$; $\log_{10} \dot{M} = 13.38$ unchanged ($\approx 0.34 M_p \text{ Gyr}^{-1}$); Mg II reproduced to 1.6% in the line-forming region.
- `_adv` base oscillation eliminated: temperature sign changes in $[1.02, 1.10] R_p$ dropped from **36 to 0**; the spike is gone; the eq $\rightarrow$ `_adv` handoff at $v = 0$ is continuous (0.7% step) and the few-percent wind correction is preserved.
- Spherical Case A ($v > 0$ everywhere) is untouched — the guard never triggers, so its `_adv` is byte-equivalent.

Cause 1 (the breathing base) remains a 1-D limitation but is now represented there by the smooth eq solution; resolving the inflow itself would require a 2-D/3-D model.

Documentation updates (this session). The user manual (`EXHALE_user_manual.tex`) gained: §1.3 *The gravitational potential* (spherical, Roche substellar $\phi(r)$ and $d\phi/dr$, and the 3-D off-axis form, matching `grav_field.f90`); §3.5 *Spectral coverage for metal photoionization* (neutral-metal thresholds down to K I at 4.34 eV/2856 Å, and how the power-law sub-grid extends below 13.6 eV vs. a loaded SED that must itself reach those wavelengths); §8 *Atomic and molecular data* (a by-process source table — Asplund09 abundances, Verner96 photoionization, Voronov97 collisional ionization, Hui & Gnedin97 H/He + Badnell06 RR/ADF48 DR metals + Shull & Van Steenberg82 Ca I + Huang23 Fe, Huang23 charge exchange, AIOLOS+CHIANTI cooling, Christie13/Draine11/Huang17 excited-H, NIST line data; and the explicit note that the code is purely atomic/ionic with *no* molecular data); and the §4 *breathing-base caveat* documenting the issue and fix above. The Fig. 4 example was regenerated from the corrected `_adv` output.

16. He 2³S triplet + metals simultaneous solve, Ly α -flux guard, and convergence workflow (2026-06-08)

Motivation. Until now the He metastable triplet (He 2³S, “HeITR”) and the trace-metal ionization network could not be solved in the same run: both placed an extra unknown at `x(4)` of the coupled ionization residual, and `input_read.f90` allocated the metal unknowns only if `(thereis_metals .and. .not. thereis_HeITR)`. Modeling HD189733b with the triplet *and* solar metals requires both at once.

Merged solver (factual). New module `nonlinear_system_solver/System_HeH_TR_metals.f90` (ion_system_HeH_TR_metals, 84 SLOC). The H/He/triplet rows ($x_1 = \text{HII}$, $x_2 = \text{HeII}$, $x_3 = \text{HeIII}$, $x_4 = \text{He } 2^3\text{S}$) are taken verbatim from `System_HeH_TR`; the metal rows are taken verbatim from `System_HeH_metals` and shifted to $x(5+2(e-1))$. The two blocks share one electron density $n_e = n_{\text{HII}} + n_{\text{HeII}} + 2n_{\text{HeIII}} + \sum_e (n_{X^+} + 2n_{X^{++}})$ (the neutral triplet is excluded), and charge exchange is added through `cx_add_to_fvec`. Solved with `hybrd1`. Supporting edits: `charge_exchange.f90` gained a settable base index `cx_metal_base` (default 4; the merged dispatch sets 5) so the metal charge-exchange rows land at the right offset; `input_read.f90` sets $N_{\text{eq}} = 4 + 2n_{\text{melem}}$ when both are active; `ionization_equilibrium.f90` uses a generalized metal base index (`mbase = 4`, or 5 with HeITR) in the initial guess / solution scatter and a both-on dispatch branch; the new source was added to the Makefile. The radiation grid, cross sections, photo-heating, IC setup and post-processing already handled both species; only the equilibrium *solver* needed merging.

Regression & physical check. With metals on / triplet off the run is byte-identical to the pre-merge binary over 8 048 steps (the new branch is not taken). With both on, the coupled solve is clean (no NaN) and was exercised end-to-end on WASP-121b with the full stack (triplet + metals + excited-H + $\text{Ly}\alpha$), reaching a flat converged solution (peak He 2^3S fraction $\sim 5 \times 10^{-6}$).

$\text{Ly}\alpha$ escape-probability input guard (factual). `input_read.f90`: if `Jlya escape-prob: True` (`jlya_mode=2`) but `Stellar Lya flux [erg/cm2/s]`: is missing or ≤ 0 , the run now stops with an explanatory message rather than silently dropping the stellar beam (in mode 2 $\bar{J}_{\text{Ly}\alpha} = \bar{J}_{\text{int}} + \bar{J}_*$ with $\bar{J}_* \propto F_{\text{Ly}\alpha,*}$, so $F_{\text{Ly}\alpha,*} = 0$ would make the halfwidth/boost settings silent no-ops).

Convergence threshold and Reconstruction scheme:-controlled two-stage reconstruction (factual).

- `du_th` is now a runtime variable (was a compile-time parameter) with default $1.0\text{e-}3$ — the original ATES-Code-main value. EXHALE had carried $2.0\text{e-}2$, at which a run flagged “converged” can still show $\sim 2\%$ spatial spread in the supersonic mass flux $\rho v r^2$.
- The single-stage vs. two-stage choice is now set by the `Reconstruction scheme: line`. `Reconstruction scheme: PLM+WENO3` selects a two-stage run: with the input line `du_th [PLM,WENO3]: <du_plm> <du_final>` supplying both thresholds (`du_th_plm` and `du_th`), `EXHALE_main.f90` starts in PLM and switches `rec_method` to WENO3 once $du < du_{\text{plm}}$ (or PLM stalls), then converges at $du < du_{\text{final}}$. This automates the manual `PLM → Load IC+WENO3` workflow recommended in the ATES README. `Reconstruction scheme: PLM` or `WENO3` is single-stage and uses only the first `du_th` value (the second is ignored).
- CFL is now runtime and settable via the input line `CFL: <value>`.

Other source-term changes (kept, with caveats). `energy_semi_implicit.f90` now damps the implicit energy update with $\partial F / \partial T = 1 + c |dC/dT|$ (was $\max(0, dC/dT)$), adding damping on the falling cooling branch; it removed one oscillation in testing but did *not* by itself fix the convergence behavior below. `Cool_coeff.f90` now interpolates the cooling tables with monotone PCHIP (C^1) rather than linear (C^0); this was implemented to test an interpolation-kink hypothesis (not supported, below) and is retained as a smoothness improvement.

Convergence — provisional observations (not conclusions). *The following is a research log, deliberately tentative; several intermediate “non-convergence” readings in earlier sessions were later traced to methodology, not physics.* Two artifacts are recorded so they are not repeated: an over-short step cap (weak-wind cases judged “floored” near 25k steps while du was still descending — HD209458b later reached $du < 2 \times 10^{-2}$ near 50k steps), and a run-script bug that killed the backgrounding subshell rather than the EXHALE.x child, so runs continued orphaned and the reported du was a premature snapshot (use `kill -x EXHALE.x`). Because of these, the earlier framing that “metal line cooling destabilizes the wind” should be treated as *not established*. With the two-stage strict- du_{th} pipeline we currently see, across planets:

- **WASP-121b** (inflated, strongly irradiated): converges quickly and flat ($\sim 0.16\%$ supersonic mass-flux spread, vs $\sim 2\%$ for single-stage PLM). This case looks solid.
- **HD209458b**: reaches a small $du \sim 0.02\text{--}0.025$ that then oscillates rather than dropping to 10^{-3} . *Tentatively* usable as a quasi-steady state: a He 2^3S+7 -metal snapshot from the two-stage pipeline gives a smooth transonic wind with $T_{\max} \approx 6.8 \times 10^3$ K and $\log_{10} \dot{M} \approx 10.1$ (median over the supersonic region), consistent with the literature $\sim 10^{10} \text{ g s}^{-1}$; see `heitr_metals_dev/HD209458b_quasisteady.ipynb`.
- **HD189733b** (strongly bound, $\beta_0 \approx 192$): du oscillates around ~ 2 in a sustained limit cycle and does not reach the threshold — the problematic case.

A plausible (but unproven) reading is that the difficulty tracks how weakly driven / strongly bound the wind is, rather than the metal content per se.

Base-breathing hypothesis and the fixes that did *not* work. The HD189733b oscillation resembles the known base-breathing behavior, and the lower boundary (`Apply_BC.f90`, `BC_component_constrho`) hard-pins the ghost density and pressure to cold reservoir values with a one-way velocity valve $\max(v_1, 0)$ — a configuration that can reflect acoustic waves. Three variants were tried and reverted:

Test	Outcome
Lower CFL ($0.6 \rightarrow 0.2$)	du floor barely moved ($\sim 2.8 \rightarrow \sim 1.9$)
Zero-gradient base velocity (drop the valve)	du still ~ 2.3
Riemann-invariant base velocity $v = v_1 + \frac{2(c_p - c_1)}{\gamma - 1}$	did not help HD189733b and broke WASP-121b

So the base-breathing reading is *consistent with* the observations but is **not demonstrated**, and these simple BC variants are not solutions. Candidate routes if pursued (both substantial, uncertain payoff): a characteristic / non-reflecting (NSCBC-style) base BC; a steady-state Newton / BVP solver (which would not orbit a time-marching limit cycle if a steady solution exists); or accepting HD189733b as a marginal, time-averaged quasi-steady case. A separate, cheaper idea for *speed* (not the oscillation) is local cell-by-cell time-stepping, since the global dt is currently set by the smallest base cell (`eval_dt.f90`). Full working notes: `docs/heitr_metals_and_convergence_notes.md`; solver-option discussion: `docs/numerical_methods.md`.

Code-size note. The merge adds one Fortran module (System_HeH_TR_metals.f90, 84 SLOC), so the 2026-06-07 file tally of 57 becomes 58; the remaining changes are small edits to charge_exchange.f90, input_read.f90, ionization_equilibrium.f90, parameters.f90, EXHALE_main.f90, energy_semi_implicit.f90 and Cool_coeff.f90.

17. Metals in the bulk-gas mass/charge/particle budget (eos_metals, 2026-06-13)

Motivation. Until this update the trace metals were strictly *passive* species: they contributed to heating/cooling and were solved in the coupled ionization equilibrium, but they were excluded from the mass density, the base density ρ_{bc} , and the electron/particle densities used everywhere *outside* the ionization solver (EOS temperature, opacities, cooling calls, energy update). The only self-consistent place was the charge balance *inside* the MINPACK residuals (System_HeH_metals.f90). At solar abundances the missing metal mass is $\sim 1.3\%$ of the H/He mass — small but systematic, and inconsistent with solving the metal ionization to high precision. Recorded in TO_BE_DONE.md on 2026-06-12; implemented 2026-06-13.

Runtime interface. A new key eos_metals 0|1 in metals.inp (documented in inputdata/metals.inp.example). 1 (the new default) includes the metals in the bulk-gas budget; 0 restores the legacy trace approximation exactly. Metals-off runs (no metals.inp) are identical either way.

What changed (factual).

- **Composition constants** (input_read.f90): the mass per hydrogen nucleus is now $\mu_H \equiv \text{mass_per_H} = 1 + 4Y + \sum_X a_X A_X$ (with $Y = \text{HeH}$ and a_X, A_X the metal abundance and atomic weight), and the base density becomes $\rho_{bc} = \mu_H / (1 + Y)$ in units of $m_H n_0$. Atomic weights live in species_table.f90:melem_A. The normalization n_0 keeps its H/He-nuclei meaning, so $n_H = n_0 / (1 + Y)$ is unchanged. (Correction to the original TO_BE_DONE entry: v_0, p_0, b_0 are m_H -based normalizations and never contained a mean molecular weight — they were never affected.)
- **Mass density:** calc_rho takes an optional nm argument and adds $\sum A_X n_{X,i}$ over all stages; passed from ionization_equilibrium.
- **Electron density:** calc_ne takes an optional nm and adds $\sum q_i n_{X,i}$; passed at every site where nm is available — composition (main-loop EOS), ionization_equilibrium (opacity/cooling), util_ion_eq (eval_cool + cooling dump), energy_semi_implicit, excited_hydrogen, post_process_adv. (The scalar T_equation already carried metal electrons via pp_nm_cell.)
- **Particle density:** calc_ntot likewise adds the metal nuclei, so the EOS $T = p / (n_{\text{tot}} + n_e)$ sees them.
- **Ghost pressure / dp_bc:** the lower-BC ghost pressure is now $n_{\text{tot_bc}} + \text{dp_bc}$ with $n_{\text{tot_bc}} = (1 + Y + \sum_X a_X) / (1 + Y)$ — so $T = T_0$ holds *exactly* at the base — and dp_bc includes the metal electrons $\sum q_i n_{X,i} / n_0$ at the first cell. set_IC/load_IC use mass_per_H/ntot_bc consistently. Full formulas: EXHALE_BC_and_IC.tex.

Validation (WASP-121b regression matrix).

- `eos_metals 0` is **byte-identical** to the pre-change goldens for both regression cases (`wasp_full 7296` steps, `wasp_he23off 7300`) — the legacy path is untouched.
- `eos_metals 1` (new default) converges cleanly (7145/7142 steps, $\sim 2\%$ fewer); profile shifts are at the expected $\sim 1\%$ mass scale: median $|\Delta n| \sim 2.4\%$, $|\Delta p| \sim 2.3\%$, $|\Delta T| \sim 0.9\%$ (max $|\Delta T| = 132$ K in the breathing-base region $r \sim 1.09$). The *number* flux nvr^2 drops 1.1% ($\log_{10} \dot{M}$: 13.31 $\rightarrow$ 13.30), which offsets the +1.2% mass per particle — the *mass*-loss rate is essentially unchanged.
- Goldens re-snapshotted under the new default; a repeat run reproduces them byte-identically (single-thread determinism intact).

Still open (small). The `use_2lev_cool` legacy cooling branch remains AIOLOS-hard-coded (default off), and the cell-by-cell scalar `T_equation` uses `pp_nm_cell` only in the post-process path (the time-marching cell-by-cell solve receives n_e via the now-metal-aware `energy_semi_implicit`).

18. Automatic initial-condition selection (IC mode: auto, 2026-06-13)

Motivation. The code now carries three IC families — cold hydrostatic (default), transonic isothermal wind (added for the Phase-4 deep-RLOF Case D, where the cold IC false-converges), and the hot-Parker warm seed (benchmarked: no speedup) — and the choice was manual. Kubyshkina+2018 build an IC for each planet automatically; the design study is `docs/auto_ic_design.md` and the literature comparison `docs/code_comparison.tex`.

The selector (one exact probe, no tunable threshold). `select_IC_auto` (`set_IC.f90`) evaluates the *cold* base sound speed in code units, $c_{\text{cold}}^2 = (n_{\text{tot,bc}} + \Delta p_{\text{bc}}) / \rho_{\text{bc}}$, and calls the existing `find_sonic` bisection for a critical point $\phi'(r_c) = 2c^2/r_c$ in the *actual* potential (Roche or spherical). If an interior sonic point exists, the cold isothermal atmosphere already wants to blow a transonic wind, so the transonic IC is selected; otherwise the cold hydrostatic IC is kept. One criterion catches both regimes that need the transonic start: deep RLOF (through the Roche/L1 topology) and low-gravity boil-off ($r_c \simeq b_0/(2c_0^2)$ in a point-mass potential). The Jeans parameter b_0 (Kubyshkina’s Λ) is logged as a diagnostic but is *not* a decision variable; the hot-Parker seed is never auto-selected.

Runtime interface. New input line `IC mode: <cold|transonic|hot_parker|auto>`. The named modes simply set the legacy flags; `auto` defers the decision to `set_IC`. The legacy keys (`Transonic IC:`, `Hot Parker IC:`) take precedence over `auto`, and the setup report (`EXHALE_setup.out`) records the family actually in effect, with a “(chosen automatically: IC mode = auto)” tag. Touched files: `parameters.f90` (`ic_mode`), `input_read.f90` (key parsing), `set_IC.f90` (`select_IC_auto`), `write_setup_report.f90`.

Validation gates (all run 2026-06-13).

Gate	Result
A-1 no IC mode key	byte-identical regression (defaults untouched)
A-2 IC mode: cold	byte-identical to A-1
A-3 spherical tutorial, auto	$b_0 = 83.2$, no interior $r_c \rightarrow$ cold; IC_dump byte-identical
B-1 Case D (deep RLOF), auto	$r_c = 1.297$ (just inside $L1 = 1.353$) $\rightarrow$ transonic, same as the manual key
B-2 boil-off planet ($0.0189M_J$, $0.446R_J$, 1100 K)	$b_0 = 8.5$, $r_c = 5.2$ (as predicted) $\rightarrow$ transonic; wind launches (initial residual 0.19 vs 14 cold)
B-3 wasp_full (Roche), auto	$\rightarrow$ transonic; du -stop in 5876 vs 7145 steps (−18%)

Honest speed verdict. The −18% step count is *Roche/deep-RLOF-specific* (the steady wind there is close to the cold isothermal transonic solution). For the mild spherical boil-off case B-2 the cold control *also* launched and converged first ($\log_{10} \dot{M} = 11.20$), while the auto run relaxed monotonically toward the same flux from $\sim 2\times$ above — the cold- c^2 Parker profile *overestimates* the flux of a weakly heated wind. So auto is not harmful (same attractor) but not faster there.

Key byproduct (a trap for all $\dot{M}$ work). While adjudicating gate B-3 it was found that the $du < 10^{-3}$ stopping criterion is **path-dependent at the $\sim 5\%$ level**: the cold-start and transonic-start wasp_full runs both “converge” by du but sit 4.8% apart in mass flux, each only 1–2% flux-flat. Newton-finishing *both* states (Load IC + Solver: Newton) collapses them onto the *same* fixed point (flux ratio 0.9996, flatness 0.005%). Practical rule, now standing: **for quantitative $\dot{M}$ always finish with Solver: Newton**; never compare bare du -stops at the percent level.

Deferred. A Phase-C retry ladder (escalate cold $\rightarrow$ warm seed on NaN/stall) is designed in auto_ic_design.md §6 but *not* implemented — the B-2 result suggests failure-triggered escalation is the right shape if a real case ever needs it. IC formulas and the selection logic are also documented in EXHALE_BC_and_IC.tex §5 and the user manual.

19. Base boundary conditions and the flux-based convergence decision (2026-06-15)

Context. A systematic study of how the lower boundary and the stabilizers interact with convergence on HD 209458b is collected in docs/EXHALE_BC_and_IC.tex (test matrix, two stabilizer tables, and a convergence-criterion history). This entry records the resulting default changes to the code; the companion document holds the full evidence.

Shapiro filter: off by default. The optional Shapiro (1970) spatial filter — ported from CETIMB (Koskinen et al. 2013a), whose use of it was found through Huang et al. (2023), to damp the same gravity-unbalanced base sound-wave instability — damps the transient “breathing” of the base cell, but on a clean cold-start it also drives the wind into a slow infall and prevents convergence. Isolated in a 2×2 (filter $\times$ base-valve) sweep, the filter — not the valve — was the cause. It is therefore *off* by default (shapiro_eps < 0) and opt-in only via Shapiro filter: <eps> [every].

Base boundary condition: density vs. pressure. The base can now be anchored either on density (the legacy fixed n_0) or on a fixed base pressure, selected by Base BC: `density|pressure [p_ubar]` (default density; pressure fixes p_{base} in μbar and derives n_0). In practice the converged wind is nearly insensitive to which anchor is used — the real lever on $\dot{M}$ is the base *density* itself, not the BC form.

The flux-based convergence decision (history). EXHALE measures two distinct quantities at the base: the flux metric du (the radial spread of $\rho v r^2$, which is exactly the constant-mass-flux criterion of ATES, $\Delta\dot{M}/\dot{M} < 10^{-3}$, and of CETIMB, $F_c = \rho v r^2$ altitude-flat), and the steady residual $\|R\| = \|\partial_t u\|$ over mass/momentum/energy, a *stricter* EXHALE-specific check. The convergence rule evolved: before the Newton finish it was purely flux-based; the Newton work then added the stronger $\|R\|$ criterion; but the reference codes converge on flux alone. The standing decision, adopted here, is therefore **flux-based** — $du < du_{\text{th}}$ is the acceptance test, and $\|R\|$ is computed and reported *for reference only* (it gates the stop only if `Resid tol:` is set).

Volume-weighted residual norm. When $\|R\|$ is reported, the default is now a cell-volume-weighted relative norm, $\sum_j |R_j| r_j^2 \Delta r_j / \sum_j |u_j| r_j^2 \Delta r_j$ (`resid_vol`, settable `Resid norm: vol|Linf`), rather than the bare L_∞ maximum. On the non-uniform radial grid the L_∞ max is dominated by the few tiny near-base cells and overstates the residual; the volume weighting gives a physically representative number. The infall diagnostic used in the test matrix was likewise switched to a *radius-weighted* negative-velocity fraction (weighted by Δr over the linear radius range), so that sub- 10 m s^{-1} near-base noise is not counted as bulk infall.

du -keyed Newton/JFNK hand-off. The production Newton finish (`Solver: Newton`) now hands off from the marching warm-up to the JFNK steady solve when the *flux* metric drops below `newton_du_switch` (default 10^{-2}), consistent with the flux-based decision and adjustable via an optional third token, `Solver: Newton [du_switch]`. This replaces the earlier residual-keyed hand-off ($\|R\| < 5 \times 10^{-2}$, variable `newton_R_switch`, now removed); the two are equivalent in practice, but keying on du avoids tying the hand-off to the stricter residual. The JFNK solve still *targets* $\|R\| < \text{resid_th}$ (10^{-3}); on failure it keeps its best iterate and falls back to marching with the du -based stops re-armed.

The full-physics residual floor (base $j = 1$). With $\text{He } 2^3\text{S} + \text{metals}$, the du -keyed hand-off plus the volume-weighted norm drove the JFNK residual from ~ 0.17 down to $\sim 2 \times 10^{-3}$ — about two orders of magnitude better — but it then *stalls* (`info` $\neq 0$) on a localized momentum imbalance at the base cell $j = 1$ (the “breathing base”), short of 10^{-3} . The wind is nonetheless flux-converged (e.g. HD 209458b full physics: $\log_{10} \dot{M}$ steady, $\rho v r^2$ spread $< 1\%$), which is a converged run by the reference standard. Closing the last factor of ~ 2 in $\|R\|$ requires physically damping the base momentum — the explicit-viscosity task at the top of `TO_BE_DONE.md`.

Touched files. `parameters.f90` (`shapiro_eps`, `base_bc_mode/base_p_ubar`, `base_v_massflux`, `resid_vol`, `newton_du_switch`; `newton_R_switch` removed); `input_read.f90` (`Base BC:`, `Resid norm:`, `Solver: Newton [du_switch]` parsing); `steady_residual.f90` (`residual_norms_vol`); `steady_newton.f90` (`resid_relnorm vol/L $\infty$` branch); `EXHALE_main.f90` (du -keyed JFNK trigger, always-on flux/residual diagnostic). Full detail and the test matrix: `docs/EXHALE_BC_and_IC.{tex,pdf}`.

20. Tidal (Roche) potential with an extended outer boundary (2026-06-16)

Motivation. The two domain modes were coupled to two different gravitational potentials. Domain mode: Spherical extends the grid to a user Outer radius but *drops* the stellar-tidal and centrifugal terms (pure $-b_0/r$); the default Roche mode keeps the full tidal potential but truncates the domain at the inner Lagrange point L_1 . Neither matched the escape models of Yan et al. (2022) / Huang et al. (2023), which retain the tidal force *and* integrate out to $\sim 10R_p$ (well past L_1). For WASP-52b, for example, $L_1 \simeq 2.53R_p$, so Roche mode truncated the domain far short of the $10R_p$ used by Yan, cutting off the outer He 10830 / H α absorbing layers.

Change. The outer-boundary extent is now decoupled from the choice of potential. In Roche mode (no Domain mode line, so the full tidal potential of `grav_field.f90` is retained), supplying an explicit Outer radius `[R_p]: > 1` sets `r_max` to that value instead of the L_1 radius:

```
else ! Roche mode: full tidal potential retained
  r_max = (3.0*Mrapp)**(-1.0/3.0)*atilde ! L1 (default)
  if (r_out_user .gt. 1.0d0) r_max = r_out_user ! NEW: extend past L1
endif
```

This reproduces the Yan/Huang “tidal force + extended domain” setup with a two-line change and no new keyword (the existing Outer radius is simply honored in Roche mode too).

Caveats and usage. Beyond L_1 the Roche potential is past its maximum (net outward force), so the 1-D radial flow there is a spherical approximation to what is physically an L_1 funnel — the same approximation Yan/Huang make. The cold hydrostatic IC is invalid past L_1 , so pair the extended domain with IC mode: `auto` (validated: the auto selector finds an interior cold sonic point and seeds a transonic wind). The tidal singularity of the $-b_0 q/(\tilde{a}-r)$ term sits at the star ($r = \tilde{a} \gg 10$), so a $10R_p$ domain is numerically safe. Domain mode: Spherical (no tidal) remains available for Huang Case-A-like comparisons.

Touched files. `input_read.f90` (honor `r_out_user` in the Roche branch of the `r_max` assignment). Documented in `docs/EXHALE_user_manual.tex` (§ potential and the input-options table).

21. Photoionization cross-section benchmarks and a TOPbase high-energy extension of He 2³S (2026-07-01)

Motivation. The He 2³S (metastable triplet) photoionization cross section `sigma_HeI23S` — the drain term that sets the He 10830 lower-level population — was a four-segment broken power-law fit to Norcross (1971), hard-cut to zero above 59.2 eV. Two questions: (i) are EXHALE’s H/He/He 2³S photoionization cross sections consistent with modern atomic data, and (ii) does the unphysical 59 eV cutoff matter? Both are settled in the new memo `docs/photoion_cross_sections.{tex,pdf}` (reproducible via `docs/compare_photoion_cross_sections.py`).

Benchmarks. H I and He II use the exact hydrogenic form; He I uses the two-term ATES fit, which agrees with the Verner+1996 (VFKY96) single-shell fit to within 14% over 24.6–500 eV. The He 2³S broken PL reproduces its parent Norcross/p-winds table to ~10% and — checked against the **TOPbase / Opacity Project** R-matrix cross section (extracted from the CDS server; raw 922-point table saved to `data/topbase_HeI_2_3S_photo.dat`) — agrees to ~8% in the smooth near-threshold region (5–20 eV). TOPbase additionally resolves a 40–55 eV autoionizing-resonance forest (peaks to 3223 Mb); their **net effect on the 2³S rate is only** ~1–2% once the few numerically under-resolved spikes are excluded (the raw table must *not* be fed into a broadband rate integral — a quadrature artifact). Extending the cross section above 59 eV raises the 2³S rate by only 0.1–1.2%.

Change. `sigma_HeI23S` is extended with a minimal, fully-connected edit (`src/modules/functions/cross_sec.f90`): the broken PL is kept unchanged up to the 45.6 eV bump; the **last power-law segment is re-aimed** so it runs continuously from the bump peak to a power-law fit of the TOPbase smooth background at the junction $E_J = 70$ eV,

$$\sigma_{\text{tail}}(E) = 10^{5.4623} E^{-2.9964} \text{ Mb} \quad (E > 70 \text{ eV}),$$

(the last-segment slope changes only $-3.039 \rightarrow -2.658$; the intercept is derived in-code for exact continuity). The old 59 eV cutoff is dropped, so 59–1240 eV now carries the TOPbase tail instead of zero. The result is C^0 -continuous at *both* joins (45.6 eV and 70 eV); the original pre-modification formula is kept in comments. A unified full-range PCHIP node table and a tail-only power law are documented as alternatives, but this is the recommended drop-in (smallest rate change, retains the validated threshold-to-bump shape exactly).

Validation. A controlled before/after on WASP-52b (`fxuv0p25_he98_L1`, Include He23S? True, ϵ Eri SED with X-rays), same IC, only `cross_sec.f90` differing: the bulk wind is unchanged ($\log \dot{M} = 11.49$ both), and the He 10830 line changes by $< 0.001\%$ in every metric (line-center depth $15.2675\% \rightarrow 15.2675\%$; EW $0.0877 \rightarrow 0.0877 \text{ \AA}$) — below the Newton convergence noise. The extension is more physically correct yet observationally harmless for He 10830, as predicted by the +0.65% rate estimate (even smaller here because the ϵ Eri SED is soft).

Recombination cross-check. A companion memo `docs/recombination_coefficients.{tex,pdf}` (`compare_recombination.p`) documents the H/He/He 2³S recombination: H II, He II, He III use the Hui & Gnedin (1997) case-B fits ($\approx$ Draine 2011 / Verner & Ferland to ~2%); He II→He I splits into Benjamin+1999 singlet/triplet (Oklopčić & Hirata 2018) when the 2³S network is on. The He 2³S triplet rate is identical to the latest escape models (Benjamin+1999). Newer nebular-community data (Porter+2012/2013, Del Zanna & Storey 2022) differ ~5–15% but are not yet adopted by escape codes; H/He case-B needs no update.

Touched files. `src/modules/functions/cross_sec.f90` (`sigma_HeI23S`: re-aimed last segment + TOPbase tail, 59 eV cutoff removed, original kept in comments). New docs: `photoion_cross_sections.{tex,pdf}`, `compare_photoion_cross_sections.py`, `data/topbase_HeI_2_3S_photo.dat`, `recombination_coefficients.{tex,pdf}`, `compare_recombination.py`, and figures under `docs/figures/`.

22. He I (1^1S) photoionization $\rightarrow$ Verner+1996 by default, with ATES-rate switches (2026-07-01)

Motivation. The He I *ground-state* (1^1S , singlet) photoionization used the legacy ATES two-term fit, which runs $\sim 12\text{--}14\%$ below the Verner+1996 (VFKY96) single-shell fit in the $50\text{--}100\text{eV}$ band (`photoion_cross_sections.tex`, §4). We make Verner the default, keep the old behavior behind a switch, and add a parallel switch for the recombination data.

Singlet/triplet separation. He photoionization in EXHALE is already split by spin: `sigma_HeI` is the 1^1S **ground (singlet)** cross section (updated here to Verner — Verner’s He I *is* the 1^1S ground state), while `sigma_HeI23S` is the 2^3S **metastable (triplet)** cross section (Norcross/TOPbase, §21). The 2^1S singlet metastable is not tracked (only 2^3S , the 10830 lower level), so the Verner update correctly touches only the singlet ground state.

Change. A single backward-compatible switch in `input.inp`: `ATES_photoionization_rate: True` reverts He I (1^1S) photoionization to the legacy two-term fit (default = Verner+1996, via `sigma_VFKY96` with He I parameters $[E_{\text{th}}, E_0, \sigma_0, y_a, P, y_w, y_0, y_1] = 24.59, 13.61, 949.2, 1.469, 3.188, 2.039, 0.4434, 2.136$). No recombination switch is needed (see below).

Comparison. WASP-52b (`fxuv0p25_he98_L1`, He23S active, same converged IC). Switching He I photoionization legacy $\rightarrow$ Verner changes the He ionization balance and hence the metastable population (Table 1): unlike the $> 59\text{eV}$ tail (§21, $< 0.001\%$), the ground-state photoionization choice is a **real $\sim 3\%$ effect on He 10830** — it matters for the absolute depth and the inferred He abundance. `ATES_photoionization_rate: True` reproduces the old result exactly.

Table 1: He I photoionization legacy (ATES two-term) vs Verner+1996 (new default), WASP-52b He 10830.

metric	legacy (ATES)	Verner (new)	change
$\log \dot{M}$ [g/s]	11.49	11.48	-2%
He 10830 line-center depth	14.516 %	14.933 %	$+2.9\%$
He 10830 equivalent width	$0.0877 \text{ \AA}$	$0.0910 \text{ \AA}$	$+3.8\%$

Recombination — no update needed (EXHALE already current). The companion update (He recombination $\rightarrow$ Porter+2012/2013 / Del Zanna & Storey 2022) was investigated and found to be a *no-op*: those are *emissivity* papers — they update the He I line cascade / line ratios, *not* the total (effective) recombination to 2^3S , and provide no drop-in $\alpha(2^3\text{S})/\alpha(1^1\text{S})$. The escape-relevant quantity — total recombination into the triplet ladder, all funnelling to the metastable 2^3S — is set by (case-B total) $\times$ (triplet spin fraction $3/4$) and is robust across modern treatments: EXHALE (Benjamin+1999) has $\alpha_3 = 2.10 \times 10^{-13} T_4^{-0.778}$, $\alpha_1 = 1.54 \times 10^{-13} T_4^{-0.486}$, while the 2025 thermosphere paper (arXiv:2509.14499) uses $\alpha_3 = (3/4)(2.72 \times 10^{-13} T_4^{-0.789}) = 2.04 \times 10^{-13} T_4^{-0.789}$ and the *identical* singlet $\alpha_1 = 1.54 \times 10^{-13} T_4^{-0.486}$ — so α_3 agrees to $\sim 2\text{--}3\%$ and α_1 is identical. EXHALE’s He recombination already matches the latest practice; no change is warranted. (A trial direct sum of Cloudy’s `he_iso_recomb.dat` level-by-level radiative

recombination gave a total $\sim 2\times$ the case-A value owing to its unmapped 1642-level indexing, confirming the raw file is not a usable shortcut.) No recombination switch was added.

Touched files. `src/modules/init/parameters.f90` (ates_photoion_rate switch), `src/modules/files_IO/input_read.f90` (parse the optional key), `src/modules/functions/cross_sec.f90` (sigma_HeI branches Verner default / legacy two-term). Verified: Verner is the default (He I σ at 30/50/100 eV = 5.36/2.02/0.394 Mb); ATES_photoionization_rate: True recovers the legacy values (5.24/1.78/0.347 Mb).

23. He 2³S photoionization → two VFKY96 (Verner) forms + power-law bridge (2026-07-01)

Motivation. `sigma_HeI23S` was a bespoke five-segment broken power law (Norcross-1971 fit + TOPbase-tail extension, §21). Since EXHALE already uses the Verner+1996 (VFKY96) `sigma_VFKY96` routine for He I and all metals, the metastable triplet is now expressed in the same form, so every photoionization cross section shares one analytic kernel.

Change. (`src/modules/functions/cross_sec.f90`.) `sigma_HeI23S` is rebuilt from two VFKY96 single-shell forms joined by a log-linear (power-law) bridge, used piecewise, with $[E_{\text{th}}, E_0, \sigma_0, y_a, P, y_w, y_0, y_1]$:

- **wing A** ($E \leq x_3 \simeq 34.70$ eV, threshold → Cooper minimum): 4.78, 2.645, 20.8, 1.0×10^{12} , 3.42, 2.681, 1.956, 2.603;
- **bridge** ($34.70 < E < 45.59$ eV): straight log–log line joining wing A(x_3) to wing B(x_4) (the “power-law part in the middle”);
- **wing B** ($E \geq x_4 \simeq 45.59$ eV, bump → tail): 45.59, 49.68, 1052, 0.04393, 2.941, 1.717, 5.488×10^{-5} , 1.118.

Each wing is a $\sim 4\%$ fit to the resonance-averaged cross section; the construction is **continuous at both** x_3 **and** x_4 , and threshold ($E < 4.78$ eV) and the smooth high- E tail ($\sim E^{-7/2}$) are supplied by the wings (no hard cutoff). The former broken power law is retained, commented out, in the routine.

Verification. Rebuilt EXHALE.x; `sigma_HeI23S` reproduces the former broken power law to $\sim 4\%$ (12 eV: 1.466 vs 1.465; 30: 0.289; 45.6: 2.57 vs 2.68 — bump 4% low; 70: 0.855 vs 0.858; 100: 0.294; 300: 0.0110), continuous across the bridge (34.70 eV: 0.222 → 38 : 0.502 → 40 : 0.795 → 45.59 : 2.570 → 46 : 2.521 Mb).

Touched files. `src/modules/functions/cross_sec.f90` (`sigma_HeI23S`: two VFKY96 wings + power-law bridge; former broken PL kept commented out).

24. Temperature-dependent Penning ionization of He 2³S (2026-07-01)

Motivation. The metastable He(2³S) destruction by neutral hydrogen, He(2³S)+H, was carried as a *temperature-independent* constant $Q_{31} = 5 \times 10^{-10} \text{ cm}^3 \text{ s}^{-1}$ — the ATES/Oklopčić & Hirata (2018) value, itself the Roberge & Dalgarno

(1982) *sum* of Penning and associative ionization, as also used by Lampón et al. (2020). Taylor et al. (2025, ApJ 989:68) replaced this with a Maxwell–Boltzmann-averaged, **temperature-dependent** fit to the cross sections of Morgner & Niehaus (1979) and Cohen & Lane (1971). Because Penning ionization is often the *dominant* 2^3S loss near the base (where neutral H is abundant), this directly affects the He 10830 prediction EXHALE is built to compute.

New default rate (Taylor et al. 2025, Table 2). A two-branch power law ($\text{cm}^3 \text{s}^{-1}$, T in K):

$$Q_{31}(T) = \begin{cases} 1.9 \times 10^{-9} (300/T)^{0.07}, & T \leq 4000 \text{ K}, \\ 9.1 \times 10^{-9} (300/T)^{0.50}, & T > 4000 \text{ K}. \end{cases} \quad (4)$$

This is $\sim 3\times$ the legacy constant across the relevant 2000–9000 K range. Note the new rate is the Penning channel only ($\text{He}(2^3\text{S}) + \text{H} \rightarrow \text{He}(1^1\text{S}) + \text{H}^+ + \text{e}^-$); the associative-ionization channel ($\text{He}(2^3\text{S}) + \text{H} \rightarrow \text{HeH}^+ + \text{e}^-$) is not represented separately in EXHALE (no HeH^+), so the single Q_{31} continues to carry the dominant $2^3\text{S} + \text{H}$ destruction term $-n_{\text{HeITR}} n_{\text{HI}} Q_{31}$.

Unconditional default. The temperature-dependent rate is used *unconditionally*: the legacy constant 5×10^{-10} is no longer selectable. (The `ATES_photoionization_rate` flag continues to revert only the He I (1^1S) photoionization cross section to the legacy ATES 2-term fit; it does *not* affect the Penning rate.)

Implementation. Q_{31} became an array over the cells (it was a scalar). A new `penning_HeI_23S(T,coeff)` in `Cool_coeff.f90` evaluates the piecewise fit over the whole grid; `HeITR_coeffs` calls it unconditionally. The two callers (`ionization_equilibrium.f90`, `post_process_adv.f90`) pass $Q_{31}(j)$ per cell into `params`; the `System_HeH_TR*` residuals are unchanged (they already read Q_{31} from `params`).

Verification. Rebuilt EXHALE.x; ran HD 189733b (metals + He23S, step-capped) with the new rate and, for comparison, with the former 5×10^{-10} constant. The temperature dependence is active and physically sensible: relative to the old constant, $\text{He}(2^3\text{S})$ is reduced near the base (ratio ≈ 0.28 at $1.0R_p$, 0.92 at $1.1R_p$; column-integrated over $1\text{--}2R_p$ drops to $0.59\times$), and the profiles converge to unity above $\sim 1.3R_p$ where H is ionized and Penning switches off — the expected neutral-H-gated signature.

Touched files. `src/modules/radiation/Cool_coeff.f90` (new `penning_HeI_23S`); `src/modules/radiation/util_ion_eq.f90` (`HeITR_coeffs`: Q_{31} array + default/legacy branch); `src/modules/radiation/ionization_equilibrium.f90` and `src/modules/post_process/post_process_adv.f90` (Q_{31} array, `params` assignment in each cell).

25. He/H diffusive separation — Phase 1 (2026-07-01, experimental, default OFF)

Motivation. EXHALE is single-fluid and re-solves composition by *local* ionization equilibrium (conserving the element ratio), so He/H was frozen at the input HeH at all radii — it could not represent the diffusive He/H separation of Taylor et al. (2025) (8%→2.5%) and Xing et al. (2023). Top recommendation of `docs/methodology_aiolos_taylor_xing`.

Method. New module `species_diffusion.f90` transports the He element ratio $f = n_{\text{He}}/n_{\text{H}}$ with bulk advection and a molecular-diffusion drift of He relative to H, $\Phi = -Dn_{\text{H}}(\partial f/\partial r + fG)$, $G = (m_{\text{He}} - m_{\text{H}})g/(kT)$, $D = 1.52 \times 10^{18}(1/m_{\text{H}} + 1/m_{\text{He}})^{1/2}T^{1/2}/n$ (Banks & Kockarts 1973), plus eddy diffusion K_{zz} on the gradient term. Solved as a **fully conservative** implicit (backward-Euler tridiagonal) update of n_{He} ; the new split is written back into `f_sp` conserving mass ($\sum_s m_s f_{\text{sp},s} = 1$, metals frozen to H). Gated on `He_diffusion` (default off; `He_Kzz` sets K_{zz}). Derivation: `docs/design_hehe_diffusion.md`.

Root-cause bug (design doc §7). A local time scale named `t0` ($= R_0/v_0$) shadowed the *global temperature* `T0` (Fortran is case-insensitive) in `TK = Tcode*T0`, zeroing `TK` (floored to 1 K) and inflating the settling coefficient $G \propto 1/TK$ by $\sim 5800\times$, causing total over-settling. The two constants are numerically near-equal ($R_0/v_0 \approx 2.83 \times 10^4$ s, $T_0 \approx 2.83 \times 10^4$ K), which hid it. Fix: rename the local to `tscale`.

Result (HD 209458b, He23S+metals). Physical mild separation — $(\text{He}/\text{H})/\text{HeH} \approx 0.6\text{--}0.9$ in the inner thermosphere, falling to $\sim 0.17\text{--}0.24$ aloft — consistent in magnitude with Taylor ($\rightarrow 0.3\times$) and Xing. Stable (no NaN); `He_diffusion OFF` is byte-identical.

Phase-1 limitations. A near-base pile-up from the lagged- n_{H} nonlinearity is held by a physical limiter $f \leq \text{HeH}$ (aloft observable unaffected); proper fix is a self-consistent n_{H} inner iteration. Ambipolar field, thermal diffusion, metal diffusion, and a Newton-finished quantitative $\dot{M}$ are Phase-2. Flag stays default OFF.

Touched files. `new src/modules/functions/species_diffusion.f90`; `parameters.f90` (`he_diffusion`, `he_kzz`); `input_read.f90` (`He_diffusion`, `He_Kzz`); `EXHALE_main.f90`; `Makefile`.

26. He/H diffusion — Phase 2 (2026-07-02): ambipolar, thermal diffusion, metals

Extends §25 (design doc §7e). The Phase-1 kernel was factored into a shared subroutine `solve_1elem` (one conservative implicit advection–diffusion–settling solve for any element vs the H background), reused by He and each metal.

- **P2b — ambipolar-corrected settling** (default ON, `He_ambipolar`). The polarization field lifts ions, so the He-vs-H effective settling mass is $\Delta m_{\text{eff}} = 3 - 0.5(\bar{Z}_{\text{He}} - \bar{Z}_{\text{H}})$ (mean charges from the local ionization state): 3 at the neutral base, 2.5 in the fully ionized wind. Modest effect ($\sim 2\%$ less depletion aloft), correct direction.
- **P2c — thermal diffusion** (default `He_alphaT=0`). Adds $\alpha_T \partial \ln T / \partial r$ to the settling coefficient; off by default.
- **P2d — element-by-element metal diffusion** (default OFF, `He_metal_diffusion`). Each trace metal element diffuses independently vs n_{H} with its own mass (`melem_A`), binary D , and ambipolar correction; ion stages rescaled to the diffused total. Metals trace $\Rightarrow$ no n_{H} feedback. **Validated (HD 209458b C/N/O):** heavier elements deplete more — at $3 R_{\text{p}}$, `element/base` $\approx$ He 0.171, C 0.158, N 0.136, O 0.105 (monotonic in mass; matches Xing et al. 2023). Regression: OFF $\Rightarrow$ He/H unchanged, C/H frozen at 1.0.

- **P2a — self-consistent n_{H} : attempted, NOT adopted.** A Picard $n_{\text{He}} \leftrightarrow n_{\text{H}}$ iteration diverges (near-base settling genuinely concentrates He, driving $n_{\text{H}} \rightarrow 0$); the physical cap $f \leq \text{HeH}$ remains the correct limiter.

Touched files. `species_diffusion.f90` (solve_1elem kernel; ambipolar dmeff; thermal term; metal-diffusion loop); `parameters.f90` (he_ambipolar, he_alphaT, he_metal_diffusion); `input_read.f90` (He_ambipolar, He_alphaT, He_metal_diffusion).

27. Diffusive separation: detailed physics and numerics

This section consolidates §25–§26 into a single, more complete treatment of the He/H (and metal) diffusion, with emphasis on the two design decisions that are easy to get wrong: the $n_{\text{He}}-n_{\text{H}}$ mass coupling (§27.3) and the settling discretization (§27.5).

27.1 Governing equation

By default EXHALE is single-fluid: the hydro evolves $(\rho, \rho v, E)$ with one bulk velocity, and the composition is re-solved every step by *local* ionization equilibrium, which conserves the element ratios. Consequently the helium-to-hydrogen ratio $f \equiv n_{\text{He}}/n_{\text{H}}$ is frozen at the input value at all radii — EXHALE cannot represent diffusive separation without an added transport equation.

We transport the He element ratio f with advection at the bulk velocity plus a molecular-diffusion *settling* drift of He relative to H,

$$\frac{\partial f}{\partial t} + v \frac{\partial f}{\partial r} = \frac{1}{n_{\text{H}}} \frac{1}{r^2} \frac{\partial}{\partial r} \left[r^2 n_{\text{H}} (D + K_{\text{zz}}) \frac{\partial f}{\partial r} + r^2 n_{\text{H}} D f G \right], \quad (5)$$

with the diffusive number flux (relative to the bulk)

$$\Phi_{\text{He}} = -n_{\text{H}} \left[(D + K_{\text{zz}}) \frac{\partial f}{\partial r} + D f G \right], \quad G = \frac{\Delta m g}{kT} + \alpha_T \frac{\partial \ln T}{\partial r}. \quad (6)$$

Here $D = 1.52 \times 10^{18} (1/m_{\text{H}} + 1/m_{\text{He}})^{1/2} T^{1/2} / n$ is the Banks & Kockarts (1973) binary He-in-H diffusion coefficient ($\text{cm}^2 \text{s}^{-1}$), $g = GM_{\text{p}}/r^2$, and $\Delta m = m_{\text{He}} - m_{\text{H}}$ ($= 3$ amu in the neutral limit; §27.6). K_{zz} is the eddy diffusion coefficient (He_Kzz, default $10^9 \text{ cm}^2 \text{s}^{-1}$): it enters only the *gradient* term and drives f toward uniform (mixing), whereas the *settling* term $D f G$ is molecular only. Their competition sets the homopause: below it ($K_{\text{zz}} \gg D$) the atmosphere is mixed, above it ($D \gg K_{\text{zz}}$) helium settles. In diffusive equilibrium ($\Phi_{\text{He}} = 0$, no eddy) $\partial \ln f / \partial r = -G$, i.e. f decreases outward with the (small) He scale height $1/G$.

27.2 Conservative implicit discretization

An early *non-conservative* form (material-derivative advection $-v \partial f / \partial r$ combined with a conservative diffusive flux) proved flux-inconsistent and *drained* helium spuriously. The production scheme is instead **fully conservative**: it evolves the He number density n_{He} with a single divergence of the total face flux

$$J_{j+1/2} = \underbrace{n_{\text{He}} v}_{\text{advection}} + \Phi_{\text{He}}, \quad \frac{n_{\text{He}}^{n+1} - n_{\text{He}}^n}{\Delta t} = -\frac{1}{r^2 \Delta r} \left[r_{j+1/2}^2 J_{j+1/2} - r_{j-1/2}^2 J_{j-1/2} \right]. \quad (7)$$

Because the cold, dense base has a tiny He scale height, the diffusion is stiff and an explicit scheme is unstable; the update is therefore **implicit** (backward Euler), assembled as a tridiagonal system in n_{He}^{n+1} and solved by the Thomas algorithm. Implicit backward Euler is unconditionally stable, and its fixed point ($n_{\text{He}}^{n+1} = n_{\text{He}}^n$) is independent of Δt , so local cell-by-cell time steps are used for steady-state acceleration. The new element split is written back into `f_sp` conserving mass ($\sum_s m_s f_{\text{sp},s} = 1$), and the same kernel (`solve_1elem`) is reused for each trace metal (§27.6).

27.3 The $n_{\text{He}}-n_{\text{H}}$ mass coupling: why a self-consistent n_{H} is *not* needed

Diffusion redistributes species at *fixed* total mass density ρ (owned by the hydro), so in every cell

$$\rho = m_{\text{H}}n_{\text{H}} + m_{\text{He}}n_{\text{He}} + (\text{metals}) \implies n_{\text{H}} = \frac{\rho(1 - c_{\text{met}}) - m_{\text{He}}n_{\text{He}}}{m_{\text{H}}}. \quad (8)$$

Thus n_{H} is itself a function of n_{He} : where helium accumulates, hydrogen must be displaced. But the flux (6) *uses* n_{H} (in the coefficient and in $f = n_{\text{He}}/n_{\text{H}}$), so the solve is genuinely nonlinear in n_{He} . Two options:

- **Lagged n_{H} (adopted).** Hold n_{H} at its start-of-step value during the solve; update it afterward, in the mass-conserving rescale, for the next step.
- **Self-consistent n_{H} (“P2a”, tried, rejected).** Iterate within the step (Picard): solve n_{He} , reset n_{H} from (8), re-solve, until n_{He} and n_{H} are mutually consistent.

Why the Picard iteration diverges. Write the He mass fraction $Y = m_{\text{He}}n_{\text{He}}/\rho \in [0, 1]$. Near the base, settling genuinely concentrates helium ($n_{\text{He}} \uparrow$); enforcing (8) exactly then drives

$$n_{\text{H}} = \frac{\rho(1 - c_{\text{met}}) - m_{\text{He}}n_{\text{He}}}{m_{\text{H}}} \xrightarrow{Y \rightarrow 1} 0,$$

so $f = n_{\text{He}}/n_{\text{H}}$ and the $1/n_{\text{H}}$ flux coefficients diverge — a genuine singularity as helium approaches a pure-He layer. Because the concentration is *real* (settling against a fixed reservoir base), the iteration has a positive feedback and blows up (NaN). The lagged scheme avoids this: with n_{H} held at a finite value, $1/n_{\text{H}}$ stays bounded, and the physical cap $f \leq \text{HeH}$ (§27.4) prevents $Y \rightarrow 1$.

Why lagging is nonetheless harmless. Four independent arguments:

1. *The lag vanishes at the fixed point.* The diffusion runs inside a relaxation loop. At steady state $n_{\text{He}}^{n+1} = n_{\text{He}}^n$ and $n_{\text{H}}^{n+1} = n_{\text{H}}^n$, so the lagged $n_{\text{H}} = n_{\text{H}}^n$ used in the solve *equals* the final consistent n_{H}^{n+1} . The lag affects only the transient path, not the converged state: where the scheme is stable, lagged and self-consistent reach the *same* steady solution, and mass consistency is achieved automatically.
2. *The region where it matters is the (unphysical) pile-up.* Lagged and self-consistent differ appreciably only where $Y \rightarrow 1$, a thin layer just above the fixed reservoir base — below the He 10830 line-forming region, hence irrelevant to the observable, and in any case capped (§27.4).
3. *Helium is a minor mass fraction where the observable lives.* In the upper thermosphere, where He/H falls to $\sim 0.2-0.5 \times$ (the part that weakens 10830), $m_{\text{He}}n_{\text{He}} \ll m_{\text{H}}n_{\text{H}}$, so $n_{\text{H}} \approx \rho/m_{\text{H}}$ is nearly insensitive to n_{He} and the lag is negligible.

4. *Trace metals have no back-reaction.* Metals are $\sim 10^{-3}$ of the mass, so their diffusion does not change n_{H} measurably; for them the lagged n_{H} is effectively exact.

Hence the divergent Picard iteration is not merely inconvenient but *unnecessary*: the lagged- n_{H} solve plus the physical cap is stable and converges to the correct steady state.

27.4 The near-base pile-up and the physical cap

Settling transports helium down toward the reservoir base. Against a fixed-composition Dirichlet base it tends to pile up in the first cells above the base (numerically unbounded with a lagged n_{H}). In a settling-plus-escape column, however, He/H *cannot* exceed the reservoir value: there is no mechanism to enrich helium above the base abundance. The code therefore applies the physical limiter

$$f \leq \text{HeH} \quad (9)$$

(and, for each metal, $f_X \leq f_{X,\text{base}}$). This encodes the expectation that the reservoir *absorbs* settled helium rather than letting it build a pure-He layer — so the cap is a physical boundary condition, not merely a numerical band-aid. The aloft separation (the observable) sits well below HeH and is unaffected by the cap.

27.5 Settling discretization: a Péclet hybrid

The settling term DfG is an advective drift (velocity $w_s = DG$, directed inward). Central differencing of an advective term is stable and accurate only at low Péclet number, $|DG| \Delta r < 2(D + K_{zz})$. For light helium this always holds (small G), and central differencing is used (least numerical diffusion). For *heavy metals* (e.g. Fe, mass 56) G is $\sim 18\times$ larger; central differencing then loses diagonal dominance and the tridiagonal solve NaNs. The kernel therefore uses a **Péclet-based central/upwind hybrid**:

$$\text{settling drift: } \begin{cases} \text{central} & |DG| \Delta r < 2(D + K_{zz}) \quad (\text{light species}), \\ \text{first-order upwind} & \text{otherwise} \quad (\text{stiff heavy-metal settling}). \end{cases} \quad (10)$$

Upwind keeps the matrix diagonally dominant for any settling strength; the metal-stage rescale additionally skips cells where the element is negligible and clamps the scale factor to avoid $0 \times \infty$. These changes were required to run WASP-121b (Fe) stably; on HD 209458b (light C/N/O), the hybrid selects central and reproduces the He-only result.

27.6 Ambipolar field, thermal diffusion, and metals

In the ionized wind the polarization field $eE = -(1/n_e)dp_e/dr$ (Koskinen et al. 2013) lifts ions; for an H^+ background it supports $\sim$ half a proton weight per unit charge, so a species of charge Z has effective mass $m - Zm_{\text{H}}/2$. The He-vs-H relative settling mass becomes

$$\Delta m_{\text{eff}} = 3 - \frac{1}{2}(\bar{Z}_{\text{He}} - \bar{Z}_{\text{H}}), \quad (11)$$

with $\bar{Z}$ the mean charges from the local ionization state — $\Delta m_{\text{eff}} = 3$ at the neutral base, 2.5 in the fully ionized $\text{He}^{++}/\text{H}^+$ wind (He_ambipolar, default on). Thermal diffusion adds $\alpha_T \partial \ln T / \partial r$ to G (He_alphaT, default 0). Each trace metal (He_metal_diffusion, default off) is diffused independently against the post-rescale n_{H} with its own mass

(`melem_A`), binary D_X , and ambipolar-corrected Δm_{eff} ; being heavier than helium, metals deplete faster with altitude where diffusion competes with the wind.

27.7 The metal-rescale ratchet bug (found in review, fixed 2026-07-02)

The first version of the metal-stage rescale contained two bugs that together made metal depletion *irreversible*: (i) an $r_X \leq 1$ clamp — redundant for the metal/H $\leq$ reservoir cap (already enforced by $\min(n_X, f_{X,\text{base}} n_{\text{H}})$) but forbidding any *replenishment* of a previously depleted cell (a one-way ratchet); and (ii) a skip of cells where the element was negligible, which made exhausted cells permanent holes. During early relaxation (wind undeveloped) settling transiently depletes metals; the ratchet locked that in. **Fix:** r_X is no longer clamped above (the cap alone bounds it), and an exhausted cell that the solve replenishes is re-seeded through the neutral stage (the ionization equilibrium re-partitions the stages next step). **Retraction:** the previously reported WASP-121b “metal homopause” ($\text{Fe} \rightarrow 0$ by $\sim 1.25 R_p$, with a ~ 2500 K hotter thermosphere) was an artifact of this ratchet — with the fix, WASP-121b metals track He (even Fe is advection-dominated there, $w_s/v \sim 10^{-3}$) and the v1/v2 temperature structures agree. The HD 209458b mass ordering ($\text{He} > \text{C} > \text{N} > \text{O}$ aloft) is unaffected.

27.8 Boundary conditions and validation

The lower boundary is Dirichlet at the reservoir composition ($f = \text{HeH}$); the outer boundary carries a zero diffusive flux plus advective outflow (helium escapes with the wind). The two-planet before/after comparison (`docs/version_compare.{md,tex,pdf}`) shows the expected regimes: on HD 209458b helium itself separates ($\text{He} 10830 \text{ } 70.5\% \rightarrow 27.1\%$, $-2.6\times$) and heavier C/N/O deplete more than He aloft, whereas on WASP-121b ($\log \dot{M} \approx 13.3$) advection dominates for every species — helium *and* metals are dragged out essentially unfractiated (10830 unchanged; metal/H tracks He/H). Both are stable (no NaN) and leave $\dot{M}$ essentially unchanged (metals are trace). All diffusion flags are **default off**, so standard runs are byte-identical.

28. Lower-atmosphere connection: Tier 1–3 (2026-07-02)

Survey + proposal + implementation: `docs/lower_atmosphere_coupling.{md,tex,pdf}`. New folder `src/modules/lower_atmosphere` everything **opt-in / default off** (regression: the standard HD 209458 b 3000-step run reproduces v1.0’s $\log \dot{M} = 8.97$ exactly).

- **Tier 1** `lower_column.f90` + key `Lower column::` the Koskinen+2022 analytic isothermal- T_{eq} hypsometric column with Visscher chemical-equilibrium $\text{H}_2/\text{H}/\text{He}$; reports the derived 1- μbar base radius (equilibrium + fully-atomic bracket), base $q_{\text{H}_2}/q_{\text{H}}/q_{\text{He}}$ and μ next to the input “Planet radius”. Gate passed: Model A $r_0/R_{1\text{bar}} = 1.343$ (paper 1.34), $q_{\text{H}_2} = 0.838$ (0.84), $q_{\text{H}} = 0.027$ (0.026).
- **Tier 2 foundation** `h3p_cooling.f90` (Miller+2013 Table-5 LTE emission fits + Table-6 non-LTE factors $s(T, n_{\text{H}_2})$; anchors reproduced to $<0.5\%$) and `mol_rates.f90` (Koskinen+2022 Table-1 rates R1–R23, verified against the PDF). Remaining: the coupled `System_HeH_mol` solver + EOS.

- **Tier 2a** key Molecular base: True: EOS-only base correction removing the H₂-bound particles from `ntot_bc` via the equilibrium fit (chemistry stays atomic; crude, documented).
- **Tier 3** optional `base.inp` (`T_base/r_base/HeH_base/Kzz_base`; echoed, no-op when absent) consumed by `input_read`, + the driver `src/utls/run_lower.py` (isothermal or Guillot 2010 semi-grey $T(p)$); VULCAN photochemistry is the documented upgrade path.

Physics note from the Tier-1 gate: the chemical-equilibrium base stays strongly molecular up to $T \sim 2000$ K, so the atomic base of $T_{\text{eq}} \sim 1000\text{--}2000$ K hot Jupiters rests on *photochemical* dissociation (Moses 2011; Koskinen 2013a) — quantified motivation for the VULCAN tier.

29. Tier-2 molecular chemistry core: H₂/H₂⁺/H₃⁺/HeH⁺ (2026-07-03)

Opt-in key Molecular chemistry: True (requires He; v1 excludes trace metals). New coupled equilibrium system `System_HeH_mol.f90` (7–8 unknowns; atomic rows = EXHALE’s own rates, so the molecule-free limit reproduces the atomic systems; molecular channels = Koskinen+2022 Table-1 via `mol_rates`). σ_{H_2} (Yan+1998 Eqs. 17–19; verified against their Table 7) added to `cross_sec.f90` and wired into the opacity/photoionization/heating integrals; H₃⁺ IR cooling (Miller+2013 + non-LTE factor) enters `cool`; the EOS is molecule-aware (particle count, electrons, masses). `f_sp` grows 33→37 (output schema + Load-IC extended; columns zero when off). **Gates:** (0) molecules-off regression $\log \dot{M}=8.97$ unchanged; (1) HD 209458 b molecules-on: fully molecular base, sharp H₂ → H front at $r=1.019 R_p$, wind above it essentially atomic ($\log \dot{M}$ 8.91 vs 8.97, He 2³S peak +1.4%); (2) hot-Uranus-like: front at $1.148 R_p$, H₃⁺ active below. Caveats: local equilibrium only (no molecular advection), no Lyman–Werner photodissociation, P4/P5 and the 4.48 eV sink omitted, $\gamma=5/3$ retained. Details: `docs/lower_atmosphere_coupling.pdf` §4.

30. Newton–diffusion co-convergence, solver hygiene, and the Ly α lower-BC audit (2026-07-03)

Stall-based Newton hand-off. `He_diffusion` runs plateau the flux metric just above the hand-off threshold (observed: du frozen at 1.084×10^{-2} vs the 10^{-2} switch for 10^6 steps). The warm-up now tracks the du plateau and hands over to JFNK once stalled within $5 \times$ the switch. *Validated:* the hand-off engaged at $du = 2.2 \times 10^{-2}$; the JFNK then hit the known base-momentum holdout (worst cell $j=1$, $\|R\| \sim 7.7 \times 10^{-3}$ while the volume-weighted reference residual was 2.2×10^{-4}) and fell back gracefully — the base-cell viscosity work remains the enabler. **Newton + He_diffusion co-convergence.** Solver: Newton previously froze the diffused He/H field at the hand-off state (the JFNK residual carries no operator-split diffusion). An outer iteration (the excited-H pattern) now alternates the JFNK solve with 500 diffusion-relaxation steps at the converged wind until the He/H field drift $< 10^{-3}$ (max 5 passes, drift reported). **Riemann floor:** `p_min` floored at $10^{-30} p_{\text{max}}$ in `speed_estimate_HLLC/ROE` (no-op for healthy states). **PH_heat_HHe OMP CRITICAL removed** (distinct- j writes are race-free; values unchanged). **regression/ run_fcheck.sh:** periodic `-fcheck` regression per the 2026-07-02 review. **lya_rt audit** (vs. Huang et al. 2017): the escape-probability closure $\bar{J} = S(1 - \beta)$ is local and lacks Huang’s *absorbing bottom* (H₂ accidental-resonance sink at $N_{\text{H}_2} \sim 10^{14} \text{ cm}^{-2}$),

so $\bar{J}$ — and the $H(n=2)$ population — is likely overestimated in the bottom few scale heights; the $H\alpha$ -forming region lies above, so the TPM impact is expected limited. Documented; a bottom-loss β channel is the follow-up.

31. Tier-3 first light: VULCAN $\rightarrow$ base.inp (2026-07-03)

The public VULCAN photochemical kinetics code was cloned (`./VULCAN`; FastChem compiled) and run to steady state for HD 189733 b (SNCHO-2025 network, `use_photo=True`; 2356 steps, ~ 6 h wall). The new converter `src/utils/vulcan_to_base.py` turns the `.vul` output into an EXHALE `base.inp`: photochemical $q_{H_2}/q_H/q_{He}$ at $1\ \mu\text{bar}$, a hypsometric base radius using VULCAN's own $\mu(p)$ and $T(p)$, and the molecular mixing ratios as comments (*no* metal release — metal/alkali chemistry is outside VULCAN's scope and remains Lavvas-only). **Result** (HD 189733 b, $1\ \mu\text{bar}$): $q_{H_2} = 0.63$, $q_H = 0.23$ — photochemistry dissociates $\sim 11\times$ more H than the chemical-equilibrium column ($q_H = 0.020$), directly quantifying the Tier-1 caveat (Moses 2011; Koskinen 2013a); $r_{\text{base}} = 1.168$ vs the analytic $1.174 R_J$; $T_{\text{base}} = 863\ \text{K}$ (the Moses11 $T(p)$, cooler than the isothermal $T_{\text{eq}} = 1183\ \text{K}$). End-to-end chain demonstrated: VULCAN $\rightarrow$ converter $\rightarrow$ base.inp $\rightarrow$ EXHALE startup override (echoed).

32. VULCAN lower-atmosphere pre-step: obtained, not included in the tree (2026-07-06)

VULCAN (with FastChem inside it, as shipped) is used as a lower-atmosphere pre-step invoked by EXHALE itself. VULCAN is a third-party code and is *not* committed to the EXHALE repository (VULCAN/ is git-ignored, avoiding the embedded- repository problem); it is fetched into EXHALE/VULCAN/ by the new `src/utils/setup_vulcan.sh` (clone + two patches + FastChem build), which the driver also runs automatically on first use if the tree is missing. The key

```
Lower atmosphere: vulcan <R_1bar>
```

in `input.inp` runs the new driver `src/utils/vulcan_driver.py` at startup whenever no `base.inp` exists: it copies the VULCAN tree into `<run_dir>/vulcan_work/`, builds the planet's Guillot $T(p)/K_{zz}$ atmosphere and picks a stellar UV spectrum by host T_{eff} , compiles FastChem once, runs VULCAN to steady state (hours on the first run, cached `.vul` afterwards; `--force` to redo), converts the result to `base.inp`, and then proceeds with the wind solve — i.e. VULCAN acts as a subroutine of EXHALE. `Lower atmosphere: analytic <R_1bar>` invokes the fast chemical-equilibrium column (`run_lower.py`) instead; omitting the key disables the pre-step (fully optional). EXHALE_ROOT overrides the code-root path for relocated installs, and an existing `base.inp` always wins. Implementation: `parameters.f90` (`lower_atm_mode/lower_atm_r1bar`), `input_read.f90` (`run_lower_atm_prestep` via `execute_command_line` before `read_base_inp`), and the added `vulcan_driver.py`. Tested: analytic branch end-to-end, VULCAN branch with a cached `.vul` (HD 189733 b photochemical base $T = 863\ \text{K}$ applied), and the opt-out (no key) path. The download URLs, citation requirements (VULCAN: Tsai et al. 2017, 2021; FastChem: Stock et al. 2018, 2022) and the exact list of modifications `setup_vulcan.sh` applies are documented in the README “Obtaining VULCAN and FastChem” section and the user manual (§“Obtaining VULCAN and FastChem”).

33. TPM.py renamed to EXHALE_transit.py (2026-07-06)

The transmission-spectrum post-processor TPM.py (the acronym clashed with “Trusted Platform Module” and did not describe the tool) is renamed EXHALE_transit.py, matching the EXHALE_* convention. Run-time overrides are now EXHALE_TRANSIT_<NAME> (PATH, SAVE_PREFIX, HE_LMIN/LMAX/N, RSTAR_RSUN, ROTP, TSTAR); the old TPM_<NAME> names remain as backward-compatible fallbacks (a `_tenv()` helper checks the new name first, then the old). Output filenames are unchanged (<prefix>tpm_He10830.txt etc.) so downstream notebooks and figure scripts keep working. All in-repo references were updated; VULCAN/ untouched.

34. Sub-level-resolved Balmer opacity and LaRT 2s inclusion (2026-07-06)

The $H\alpha/H\beta$ transmission was computed as $f_{23}n_2$ with the multiplet oscillator strength ($f_{23} = 0.6407$ for $H\alpha$), which implicitly assumes the $2s:2p$ populations sit in the 1:3 statistical ratio. They do not—the metastable $2s$ ($A_{2s1s} = 8.26 \text{ s}^{-1}$) and the $\text{Ly}\alpha$ -pumped $2p$ depart strongly from it—and the two sub-levels have *different* Balmer cross sections (only $2s \rightarrow 3p$ is allowed from $2s$; $2p \rightarrow 3s, 3d$ from $2p$). Fixed to sum per sub-level, $\tau \sim f_{2s}n_{2s} + f_{2p}n_{2p}$, with $f_{2s} = 0.4349$, $f_{2p} = 0.70941$ ($H\alpha$) and 0.1028, 0.125886 ($H\beta$); the components share the Doppler-dominated profile.

EXHALE_transit.py (escape-prob path) now carries n_{2s}/n_{2p} separately; the effect is small ($\lesssim$ few %) since the in-code $J_{\text{Ly}\alpha}$ keeps $2s:2p$ near statistical in the forming region. examples/tpm_halpha_lart2d.py (LaRT path) previously used the $2p$ -only population, *discarding the 2s absorption entirely*; it now returns and sums both (the LaRT scattering rate pumps only $1s \rightarrow 2p$, entering the S_{2p} source term, while $2s$ is recombination/collisionally fed throughout the atmosphere). Effect: +21% (HD 189733 b) to +67% (WASP-121 b). The $\text{Ly}\alpha$ absorption line itself is unchanged and already correct— $1s \rightarrow 2p$ only ($f = 0.41641$ on the ground-state n_{HI} ; $1s \rightarrow 2s$ is a forbidden two-photon transition).

benchmarks/hd209/input.inp was missing Stellar Ly α flux, Stellar Teff, and Stellar radius; without the $\text{Ly}\alpha$ flux the LaRT converter read $F_{\text{Ly}\alpha} = 0 \Rightarrow L_{\text{Ly}\alpha} = 0 \Rightarrow n_{2p} \simeq 0$, a spurious $H\alpha = 0.000\%$. Added $T_{\text{eff}} = 6065 \text{ K}$, $R_* = 1.155 R_\odot$, and $F_{\text{Ly}\alpha} = 3.5 \times 10^3 \text{ erg cm}^{-2} \text{ s}^{-1}$ (quiet-solar-analog, R_*^2 scaled; provisional, since the absolute LaRT depth scales with it). With both paths on the same VULCAN-BC winds and the sub-level opacity, the escape-prob/LaRT $H\alpha$ ratios become HD 209458 b 0.84/0.43 (2.0), HD 189733 b 1.23/0.91 (1.4), WASP-121 b 4.41/2.33 (1.9), WASP-52 b 5.69/2.20 (2.6)—so the inline closure overestimates $H\alpha$ by only $\sim 1.4\text{--}2.6\times$, far less than the earlier $\sim 4\text{--}30\times$ (pre-VULCAN winds with the $2p$ -only LaRT). Paper abstract, methods, results table, summary, and figs/make_lart_fig.py were updated accordingly.

35. Molecular equilibrium seeding fix, molecular-base coupling, Tier-2 gate refresh (2026-07-16)

Frozen-atomic-base bug (diagnosed and fixed). The 8-unknown molecular equilibrium (ion_system_HeH_mol, hybrd1) is bistable in its initial guess: from a zero-molecular warm start it fails to converge at the dense, optically-thick base, and the `thereis_mol` branch never checked info—the failed state was silently adopted and then frozen indefinitely ($x_{\text{H}_2} = 0.0122$ at the base, flat across 130k steps and independent of T ; numerically *not* a root of the

coded network—it would require an H_2 photoionization rate $15\times$ the unattenuated stellar bound). The true dark-base equilibrium is strongly molecular (3-body formation R15 vs. thermal dissociation R12), which the earlier gate runs had reached only because the then-uninitialized molecular `f_sp` columns acted as an accidental nonzero seed; the composition-audit zero-init exposed the latent fragility. Fixes: (i) `info` is checked and a failed solve retries once from a physically-informed guess seeded by the chemical-equilibrium H_2 fit `q_h2_equilibrium(p,T)` (still-failing cells keep the previous state and are counted in a one-line summary warning); (ii) `set_IC` seeds the molecular H_2 column from the same fit (H nuclei conserved) instead of zero.

Molecular base coupled on. Solving molecular chemistry now implies the molecular base particle count (Molecular base: True behavior): with an atomic `ntot_bc` the base temperature is inflated by the particle-count ratio ($\sim 1.8\times$ for a fully molecular base), thermally dissociating the H_2 layer the chemistry just built. A note is printed when the coupling engages.

Tier-2 gates re-established (deterministic; inputs pinned at `docs/lower_atmosphere_figs/data_g*/input.inp`). Gate 0 (HD 209458 b atomic, 3000 steps): $\log \dot{M} = 8.97$ unchanged. Gate 1 (HD 209458 b molecular): fully molecular base ($x_{\text{H}_2} = 0.997$, base T unchanged at ~ 1464 K), sharp $\text{H}_2 \rightarrow \text{H}$ front at $r = 1.019 R_p$, $\log \dot{M} = 8.87$ vs. 8.97 atomic, $\text{He } 2^3\text{S}$ peak $+1.3\%$; stable over 12000 steps with zero solve failures. Gate 2 (hot-Uranus-like: $0.0457 M_J$, $R_p = 0.49 R_J$, $T_{\text{eq}} = 1140$ K, HD 209458 b orbit/spectrum—input recovered by matching the archived domain fingerprint $r_{\text{max}} = 4.814118$ to 6 decimals): fully molecular base, front at $r = 1.166 R_p$, H_3^+ active in the molecular layer. Atomic runs byte-identical throughout (make check + OMP identity).

Gate refresh (2026-07-23, current production defaults). After the 2026-07-17 onward changes (staged secondary ionization, `He_rec_coupling` on, `He-H` charge exchange on; §§36–39) the three gates were re-recorded at a uniform 12000-step relaxation-snapshot convention. The $\text{H}_2 \rightarrow \text{H}$ fronts and molecular base composition are essentially unchanged (front $1.019 \rightarrow 1.020$ for HD 209458 b, $1.166 \rightarrow 1.156$ for the hot Uranus; base $x_{\text{H}_2} = 0.996/0.9998$), so the dissociation-front physics is robust to the new defaults. The wind $\log \dot{M}$ now reads 10.58 for both HD 209458 b runs at 12000 steps — these are relaxation snapshots, not flux-flat converged: under the new defaults the HD 209458 b atomic gate does not converge (du reaches a fixed point $\sim 2.6 \times 10^{-6}$ while the $\rho v r^2$ mass-flux spread plateaus near 4%), so all three gates share the same step cap. $\text{He } 2^3\text{S}$ is now strongly suppressed in the molecular base below the front ($\text{He } 2^3\text{S} + \text{H}_2$ Penning destruction, $\sim 10^4\times$ at the front, rising back to the atomic value above $\sim 1.3 R_p$); the earlier “ $+1.3\%$ peak” was a whole-domain-peak artifact of the outer-wind transient. Pinned `data_g*/` refreshed; `lower_atmosphere_coupling` and `fig_mol_structure.pdf` updated.

36. Default H/He rates $\rightarrow$ Badnell+Mao / Voronov, van Hoof free-free Gaunt factor, free-free charge fix, and secondary ionization (2026-07-17)

Six coupled changes to the H/He microphysics. Items 1–2 are gated by the new `legacy_hhe_rates` flag (default False = new rates); items 3–4 (free-free) apply always; item 5 (secondary ionization) is gated by `use_sec_ion` (default True); item 6 rides along with the heating rework. Full cross-code context and the benchmark table: `docs/atomic_data_EXHALE_vs_MoCHII`.

36.1 What changed and why

1. H/He recombination default. The pre-update H/He case-B coefficients were unattributed fits inherited from ATES; a numerical check showed them to be exactly the Hui & Gnedin (1997) case-B fits. The new default is

$$\alpha^B(T) = \alpha_{\text{Badnell}}^A(T) - \alpha_1^{\text{Mao}}(T), \quad (12)$$

the Badnell (2023) radiative total minus the Mao & Kaastra (2016, A&A 587, A84) $n = 1$ coefficient, with the Badnell three-term dielectronic recombination added to α^A for He II. This matches the construction MoCHII already uses, so the two codes now share the H/He recombination coefficients. Setting `legacy_hhe_rates = True` restores the Hui & Gnedin fits exactly.

2. H/He collisional ionization default. The legacy H I / He I rates were Abel et al. (1997) (reproducing Janev et al. 1987) and He II a Hui & Gnedin (1997) fit, none source-annotated. The Voronov (1997, ADNDT 65, 1) fits,

$$k(T) = A \frac{(1 + P\sqrt{U})U^K e^{-U}}{X + U}, \quad U = \Delta E[\text{eV}]/k_B T, \quad (13)$$

are now the default, matching the metal rates (already Voronov in both codes) and MoCHII. Reverted together with item 1 by `legacy_hhe_rates`.

3. Free-free Gaunt factor. The frequency-averaged non-relativistic $\langle g_{\text{ff}} \rangle(\gamma^2)$ (with $\gamma^2 = Z_{\text{ion}}^2 \text{Ry}/k_B T$) of van Hoof et al. (2014, MNRAS 444, 420) is tabulated at 161 points in $\log_{10} \gamma^2$ from -6 to 10 in 0.1 dex steps (generator `cooling_data/gauntff_thermal_avg.py`), replacing the earlier unattributed two-branch logarithmic fit. Applies always, independent of `legacy_hhe_rates`.

4. Free-free charge-weighting correction (physical bug). The previous code weighted He II (net charge $+1$) by $Z^2 = 4$ and evaluated the metal Gaunt factor at the nuclear atomic number ($Z = 8$ for O, 26 for Fe). Free-free emission scales with the ion *net* charge, not the nuclear number: $Z_{\text{ion}} = 1$ for H II, He II, and singly-ionized metals; $Z_{\text{ion}} = 2$ for He III and doubly-ionized metals. The most visible consequence is that the free-free cooling of a He II-dominated layer drops by a factor of four. Applies always.

5. Secondary ionization (Shull & van Steenberg 1985) in the main loop. A photoelectron ejected with $E_0 = e_\nu - E_{\text{th}}$ above ~ 40 eV does not thermalize immediately: it deposits only $f_{\text{heat}}(x)$ of its excess as heat and drives secondary H I and He I ionizations, where x is the ionized fraction of the H+He nuclei,

$$f_{\text{heat}}(x) = 0.9971 (1 - (1 - x^{0.2663})^{1.3163}), \quad (14)$$

$$f_{\text{ion,H I}}(x) = 0.3908 (1 - x^{0.4092})^{1.7592}, \quad (15)$$

$$f_{\text{ion,He I}}(x) = 0.0554 (1 - x^{0.4614})^{1.6660}. \quad (16)$$

Below 40 eV the photoelectron thermalizes fully. Gated by `use_sec_ion` (default `True`); `use_sec_ion = False` is bit-identical to the legacy full-thermalization path. This differs from `wind_ae`, which combines the same SvS85 partition with Dere (2007) tables tied to its fixed spectral grid; the main loop runs on a runtime SED and applies the SvS85 species-resolved fits directly.

6. He I 2^3S triplet photoheating channel added. The metastable He I 2^3S (threshold 4.8 eV) previously contributed only opacity and its photoionization rate but no photoheating. Its photoelectron energy now enters the heating and absorbed-energy budgets and the secondary-ionization source consistently with every other absorber.

36.2 Files touched

Cool_coeff.f90 (rr_badnell, dr_HeII_badnell, rr_mao, alphaB*_new, voronov_ci/ci*_new, the 161-point gff_avg table); util_ion_eq.f90 (free-free net-charge weighting in eval_cool; secondary ionization and He 2^3S heating in PH_heat_H/PH_heat_HHe); T_equation.f90 (free-free net-charge weighting); ionization_equilibrium.f90 and post_process_adv.f90 (callers wired for the new signatures); input_read.f90, parameters.f90, write_setup_report.f90 (legacy_hhe_rates and use_sec_ion declared, defaulted, parsed, echoed); cooling_data/gauntff_thermal_avg.py (van Hoof table generator); docs/input_schema.md (keys K14b, K14c).

36.3 Quantitative check

New default rates (α_B , collisional ionization ci ; $\text{cm}^3 \text{s}^{-1}$), with the legacy value at 10^4 K for comparison. The legacy $\alpha_B(\text{HII})$ at 10^4 K is the Hui & Gnedin (1997) fit, 2.592×10^{-13} .

Coefficient	10^4 K	10^5 K	10^6 K	legacy (10^4 K)
$\alpha_B(\text{HII})$	2.606e-13	3.070e-14	2.163e-15	2.592e-13
$\alpha_B(\text{He II})$	2.807e-13	5.006e-13	1.015e-12	2.616e-13
$\alpha_B(\text{He III})$	1.536e-12	2.370e-13	2.244e-14	—
$ci(\text{HI})$	7.448e-16	3.963e-9	3.103e-8	7.247e-16

The Fortran reproduces an independent Python recomputation to all printed digits, and the legacy branch reproduces the pre-update values exactly. The new $\alpha_B(\text{HII})$ at 10^4 K differs from the legacy Hui & Gnedin value by $\sim 1\%$; He II and He III differ more, mainly because the new He II value carries the Badnell dielectronic contribution explicitly. The van Hoof thermal Gaunt factor at $T = 10^4 \text{ K}$ and $Z = 1$ is $\langle g_{\text{ff}} \rangle = 1.266$. A smoke run confirms the expected factor-of-four drop in free-free cooling across a He II-dominated layer, and `use_sec_ion = False` reproduces the legacy heating and photoionization rates bit-identically.

36.4 Effect on existing results

Because the default rates change, results computed with the new default differ from the pre-update golden and gate outputs. `legacy_hhe_rates = True` recovers the old H/He ionization balance, with two always-on exceptions: the free-free charge-weighting correction (item 4) and the van Hoof Gaunt table (item 3), both physical corrections intentionally not tied to the switch. Secondary ionization (item 5) is a further default change; `use_sec_ion = False` restores the full-thermalization rates bit-identically.

37. He recombination radiation → H I ionization coupling (use_he_rec_coupling, 2026-07-17)

EXHALE carries no diffuse radiation field, so pure case B implicitly assumes every He II → He I recombination photon is reabsorbed locally by helium ($\alpha_B = \alpha_A - \alpha_1$, ground capture dropped). In H-dominated gas that is wrong: part of the He recombination radiation — the ≥ 24.6 eV ground-capture continuum, the 584 Å resonance line, the 19.8 eV 2^3S line, and the >13.6 eV part of the He I two-photon continuum — ionizes H instead. The physics and the choice of recipe were worked out in docs/QUESTIONS_2026-07-17.md. This section documents an opt-in coupling using Draine's (2011) on-the-spot y/z parametrization. New input key He_rec_coupling (flag use_he_rec_coupling), default False → pure case B, bit-identical to the legacy path.

37.1 What changed and why

When on, the He recombination photons that ionize H are added as an extra H I photoionization rate with its photoelectron heating, and the He II recombination coefficient becomes $\alpha_B + y\alpha_1$ rather than pure case B. Two local parameters set the split:

- $y = 1/(1 + R n_{\text{HeI}}/n_{\text{HI}})$ (Draine Eq. 14.16), the fraction of the ≥ 24.6 eV ground-capture continuum absorbed by H, with the cross-section ratio $R = \sigma_{\text{He}}/\sigma_{\text{H}}$ at the 24.6 eV He I ground edge. Evaluated from the code cross sections, $R = 6.004$.
- $z = 0.67 + 0.29/(1 + n_e/n_{\text{crit}})$ (Draine Sec. 15.5), the density-dependent fraction of case-B (excited-state) cascade photons that ionize H, with the 2^3S critical density $n_{\text{crit}} = 1100 e^{1.2/T_4} T_4^{1/2} \text{ cm}^{-3}$; it interpolates between Draine's two limits (~ 0.96 low density → ~ 0.67 high density).

Atomic (case-B) mode. No explicit 2^3S : α_1 is the Mao & Kaastra (2016) ground ($1s^2$) capture and α_B the active He II case B. The He II recombination coefficient is $\alpha_{\text{eff}} = \alpha_B + y\alpha_1$; the extra H I photoionization rate is $n_{\text{HeII}}n_e[z\alpha_B + y\alpha_1]/n_{\text{HI}}$; the photoelectron heating uses the ground-channel energy $E_{\text{gnd}} = 24.6 - 13.6 = 11.0$ eV and a cascade-averaged $E_{\text{casc}} \approx 6.14$ eV.

He I 2^3S (triplet) mode. The mode that computes 10830: the decay channels are explicit, so z is not used — a 2^3S atom destroyed by photoionization or Penning emits no 19.8 eV photon, so the cascade must be summed at the actual channel rates (α_1 is the 1^1S channel). The 1^1S channel coefficient is $y\alpha_1 + 0.25\alpha_B$: the net ground capture plus the singlet-excited capture channel ($0.25\alpha_B$) the current network omits. The H-ionizing photon production is summed over channels: ground $y\alpha_1 n_{\text{HeII}}n_e$; singlet-excited $0.8521 \times 0.25\alpha_B n_{\text{HeII}}n_e$ ($0.8521 = \frac{2}{3}$ from the 584 Å line + $\frac{1}{3} \times 0.5564$ from the 2^1S two-photon continuum above 13.6 eV); 2^3S radiative decay $A_{31}n(2^3S)$ (19.8 eV line, always ionizes H); and 2^3S collisionally converted to the singlets $n_en(2^3S)(0.5564q_{31a} + q_{31b})$. The channel deposit energies are 11.0 (ground), 5.53 (singlet-excited), 6.2 (19.8 eV line), and 2.512/7.6 eV (collisionally converted two-photon / 584 Å). The two-photon constants are integrals of the Drake, Victor & Dalgarno (1969) spectral shape rather than round numbers. The 2^1S term emits a photon *pair* summing to 20.62 eV with a distribution peaking at half that, so above the H I

edge it delivers 0.5564 ionizing photons per decay carrying 8.9646 eV — a mean photon energy of 16.110 eV, hence 2.512 eV of photoelectron energy each. Treating the in-band photons as uniformly distributed instead puts the mean at 17.109 eV and the photoelectron at 3.511 eV, 40% high. `f_2q_HeI` and `Ee_2q_HeI` in `util_ion_eq.f90` carry these, and the 0.8521/5.53/6.14 combinations are formed from them so the three cannot drift apart.

The corrections are applied at the lagged (pre-solve) densities in the main ionization-equilibrium loop and re-evaluated at the advection-corrected densities in the post-processor (rate/coefficient before the advection solve, heating after). The He II recombination *cooling* is left at $kT \alpha_B$; the mismatch is at most $y \alpha_1 kT$, negligible.

37.2 Files touched

`util_ion_eq.f90` (new `he_rec_coupling` subroutine returning the corrected He II recombination coefficient, the extra HI photoionization rate, and the photoelectron heating; all zero when off); `Cool_coeff.f90` (new `alpha1_HeII_mao`, the Mao & Kaastra 2016 ground $1s^2$ capture, i.e. the α_1 subtracted in `alphaB_HeII_new`); `ionization_equilibrium.f90` and `post_process_adv.f90` (callers wired to `rcheiiB/P_HI/heat` in the main loop and to the advection ODE and `theat` in post-process); `input_read.f90`, `parameters.f90`, `write_setup_report.f90` (`use_he_rec_coupling` declared, defaulted `.false.`, parsed, echoed); `docs/input_schema.md` (key K14d).

37.3 Quantitative check (HD 209458 b)

Protocol: HD 209458 b `input.inp+metals.inp`, the converged profiles loaded as IC, then four runs {atomic, He 2^3S } $\times$ {off, on} each relaxed for 15000 steps from the same IC. The IC predates the 2026-07-17 rate update, so both branches share the same partial re-relaxation and **only the on/off difference is meaningful** (not the absolute values). Across the wind $y \sim 0.65\text{--}0.69$ and $R = \sigma_{\text{He}}/\sigma_{\text{H}}(24.6\text{ eV}) = 6.004$. Script `docs/plot_he_rec_coupling.py`; figures `docs/figures/he_rec_coupling_{atomic,heitr}.pdf`.

Quantity (on/off)	Atomic mode	He I 2^3S mode
He I	$\times 1.04$ ($1.2 R_p$) $\rightarrow \times 1.25$ plateau ($1.5\text{--}3.5 R_p$) $\rightarrow \times 1.32$ ($3 R_p$)	$\times 1.01\text{--}1.16$
He II (min.)	$\times 0.72\text{--}0.74$ ($1.1\text{--}1.2 R_p$) $\rightarrow \sim 1.0$ ($2 R_p$)	$\times 0.885$
$n(2^3S)$	— (not computed)	$\times 0.89$ (1.05), $\times 0.92$ (1.2), $\times 0.98\text{--}1.01$ ($1.5\text{--}2 R_p$)
Temperature	$\leq +1\%$ (outer $+3\%$)	$\leq +0.3\%$
$\log \dot{M}$	$9.58 \rightarrow 9.65$ ($+0.07$ dex, $\sim +17\%$)	$9.57 \rightarrow 9.61$ ($+0.04$ dex, $\sim +10\%$)

The atomic mode carries the larger He ionization-structure change (case B alone, $y = 0$, under-recombines He), and its extra H photoionization raises the mass-loss rate by $\sim 17\%$. In the 2^3S mode the 10830-forming region ($1.5\text{--}2 R_p$) sees $n(2^3S)$ within $\sim 1\text{--}2\%$ of the off case, so the He I 10830 line is essentially unchanged, while the base 2^3S drops $\sim 11\%$ (the TR-mode 1^1S channel coefficient change feeds directly into n_{HeII}). The temperature response is small in both modes. Default off is bit-identical to the pre-update path; existing goldens and gates are unaffected.




Figure 1: He recombination radiation $\rightarrow$ HI coupling, **atomic mode** (HD 209458 b). Clockwise from top left: He ionization fractions off vs. on; the on/off ratio for He I and He II; the photoheating rate; and the temperature. Turning the coupling on raises He I and lowers the He II minimum, with the $\sim 17\%$ mass-loss increase reflected in the extra photoheating.




Figure 2: Same comparison in the **He I 2^3S mode**. The extra panel and ratio curve track the metastable 2^3S density: it drops $\sim 11\%$ at the base but stays within $\sim 1\text{--}2\%$ of the off case through the 10830-forming region ($1.5\text{--}2 R_p$), so the He I 10830 line is essentially unchanged.

38. Staged activation of secondary ionization (2026-07-23)

38.1 The regression

With secondary ionization on by default (the `use_sec_ion` item of the 2026-07-17 rate update), the high-gravity WASP-121 b cases (backup/regression/wasp_full, wasp_he23off) crashed with a NaN at $\sim$ step 587: negative densities first, then a NaN base pressure and temperature. The coupling itself is physically correct in magnitude — at the shielded base the SvS85 secondary channel raises the H photoionization rate by $\sim 35\times$ over the direct rate, the expected size for the $\sim$ keV photoelectrons there. The failure is a transient-path problem, not a steady-state one: the extra base ionization deepens the base cooling valley and amplifies the base acoustic mode of the marginal cold initial condition until the hydro diverges. Two facts pin this down: (a) a run with the coupling off converges cleanly ($du < 10^{-3}$, $\log_{10} \dot{M} = 13.30$); and (b) restarting with the coupling on from that converged state is benign — no NaN, du stays $\lesssim 4 \times 10^{-2}$ during a ~ 2000 -step adjustment and re-converges. The steady state with the full physics exists and is stable; only the cold-start path to it is fatal.

38.2 The fix: converge first without the coupling, then switch it on

A runtime state flag `sec_ion_active` now controls whether the SvS85 partition is applied, separately from the `use_sec_ion` enable flag. The photoionization / photoheating routines apply the coupling only when `use_sec_ion` .and. `sec_ion_active`; when it is off the code path is bit-identical to the `use_sec_ion = False` (full-thermalization) path. The main loop keeps `sec_ion_active = False` while the wind first relaxes; the moment that first relaxation reaches a convergence/stall criterion, the flag flips to True, the convergence logicals and the mass-flux level-stability history are reset, and the run continues and re-converges with the full physics. Stops are additionally held for `N_stall` (default 2000) steps after the flip: du reacts only once the base adjustment wave driven by the new coupling has formed, so without the hold the run stopped one step after activation (activation at step 7144, spurious stop at 7145) with a state that had not yet incorporated the physics. The Newton (`Solver: Newton`) finish and the direct-solve diagnostic modes (which bypass the marching loop) switch the coupling on before their solve, and the flag is forced on after the loop so post-processing and final outputs always carry the physics even if the loop hit the iteration cap before the flip.

38.3 The Immediate override

`Secondary_ionization: Immediate` sets `use_sec_ion = True` and applies the coupling from step 0, restoring the pre-staging behavior. It is intended for A/B tests only; on the high-gravity cold-start cases it reproduces the $\sim$ step-587 NaN crash, confirming that the staged path is what removes it.

38.4 Files touched

`parameters.f90` (`sec_ion_active`, `sec_ion_immediate`); `util_ion_eq.f90` (`PH_heat_H` / `PH_heat_HHe` gate on the local `sec_on = use_sec_ion` .and. `sec_ion_active`); `EXHALE_main.f90` (init, stage flip on first convergence, Newton-finish and diagnostic-mode activation, post-loop force-on); `input_read.f90` (the Immediate value parsed).

38.5 Effect on existing results (wasp_full validation)

For Secondary_ionization: False the run is bit-identical to before (verified byte-identical on wasp_full and wasp_he23off). For the default (staged on) wasp_full: activation at step 7144, re-convergence at step 13489 ($du = 9.98 \times 10^{-4}$), and $\log_{10} \dot{M}$ moves $13.30 \rightarrow 13.22$. The $\sim -17\%$ mass-loss rate is the real secondary-ionization physics: the SvS85 partition diverts most of the hard-photon photoelectron energy from heat into ionization and excitation in the weakly ionized base, so less energy drives the wind. Immediate reproduces the old (crashing) behavior on the high-gravity cases.

39. He-triplet physics series from the Falorca & Vidotto (2026) review (2026-07-23)

The line-by-line comparison against Falorca & Vidotto (2026, arXiv:2607.18193) and García Muñoz (2025, A&A 698, A199) — full analysis in docs/falorca2026_he3d_review.md — exposed six physics gaps, all implemented and validated in this series (its validation work also surfaced the secondary-ionization NaN regression fixed in the previous section). Physical correctness is the acceptance criterion throughout; the impact sizes below are informational.

- Penning ionization products** (heh_tr_rows, System_HeH_mol, adv_implicit_HeH_TR): $\text{He}(2^3\text{S}) + \text{H}^0 \rightarrow \text{He}(1^1\text{S}) + \text{H}^+ + e^-$ now produces the H^+ (with its 6.2 eV of electron heating, $e_{\text{th,HeI}} - e_{\text{th,HeTR}} - e_{\text{th,HI}}$) it always removed the triplet for. In the EUV-shielded base (Γ_{HI} exponentially small) this is the *dominant* H-ionization channel: n_{HII} rises $2-7\times$ at $r < 1.03$ on wasp_full, while the wind and $\dot{M}$ are unchanged. The associative branch ($\rightarrow \text{HeH}^+$, $\sim 10\%$, GM25) is folded in as H^+ (documented at the code site).
- Collisional ionization in the TR systems** (heh_tr_rows b terms; ci_HeI23S in Cool_coeff.f90, the Black-1981-derived $6.41 \times 10^{-21} \sqrt{T} e^{-55338/T} / E_{23\text{S}}$, $k(10^4 \text{ K}) = 3.3 \times 10^{-10}$): the TR (Oklopčić-form) rows silently dropped *all* electron-impact ionization while eval_cool kept charging its cooling — both restored and consistent now, including CI of the metastable itself (a $\sim 1-2\%$ 2^3S loss at 10^4 K , its 4.8 eV cooling in the coio channel, and its He I / adv-row counterparts).
- Explicit- $n(2^3\text{S})$ cooling** (eval_cool + the semi-implicit energy solver): the 10830 Å collisional-excitation channel $1.16 \times 10^{-20} \sqrt{T} e^{-13179/T} n_e n_{23\text{S}}$ (Black form with the *explicit* metastable density — Falorca & Vidotto warn Black's implicit steady-state triplet is wrong here) plus the 0.80/1.40 eV q_{31a}/q_{31b} conversion ledger. The q_{13} (19.8 eV) term is deliberately excluded: the Cen-1992 He I excitation-cooling term already carries that channel. The terms are passed into solve_energy_semi_implicit (all three eval_cool calls) so they act on T ; without that wiring they were diagnostic-only. Measured: 1–5% of the total cooling on WASP-121 b (8–11% of local $\text{Ly}\alpha$ at the 2^3S peak), locally up to $\sim 28\%$ on HD 209458 b.
- He_rec_coupling default ON** (the previous section has the model): with the coupling off the TR singlet recombination used α_1 alone — neither case A nor case B — and the He recombination photons ionized/heated nothing. Validated stable on wasp_full: base $n_{23\text{S}} -10\%$, the 10830-forming region within $\sim 2\%$, $\dot{M}$ unchanged; on HD 209458 b $\dot{M}$ rises $+0.04-0.07$ dex. He_rec_coupling: False restores the legacy path (byte-identical).

5. **He $\leftrightarrow$ H charge exchange in the default set** (he_h_charge_exchange, default on; Koskinen 2013 rates via Huang 2023 Table 4): the group-B pair moved out of the cx_full-gated table assembly into dedicated he_h_cx_fvec/jac/fvec_adv routines called by *every* ionization system with He (previously only the metals systems could reach it), including the advection pair and the analytic Jacobians of the Newton systems. The effect is confined to the neutral-H base — He II $\times 0.59$ at $r = 1.05$, -1% by $r \geq 1.2$, $\dot{M}$ unchanged — exactly the $\text{He}^+ + \text{H}^0 \rightarrow \text{He}^0 + \text{H}^+$ sink where n_{H^0} dominates. cx_full still gates groups C/D. The new input keys of this and the preceding sections were also added to the parser’s known-keys list (they parsed correctly but tripped the unknown-line warning).
6. **He(2^3S)+H₂ Penning ionization** (penning_HeI23S_H2; System_HeH_mol rows 4/5/8 + 4.4 eV heating): GM25 Table A.5 (Cohen & Lane 1977) fitted as $5.3791 \times 10^{-12} T^{0.676} e^{-695.21/T}$ ($\leq 0.13\%$ over 500–10⁴ K). On the pinned molecular gates the metastable drops $39\times$ (HD 209 base) and $430\times$ (hot-Uranus, out to $r \sim 1.1$) exactly where $x_{\text{H}_2} \rightarrow 1$, with the upper wind and $\dot{M}$ unchanged — GM25’s dominant deep metastable sink, now present for Tier-2 + triplet runs. The _adv post-process is atomic-only (documented) and needs no counterpart.

Regression state after the series: Secondary_ionization: False + He_rec_coupling: False + He_H_charge_exchange: False reproduce the respective legacy paths byte-identically; the wasp goldens and the parse-dump corpus were re-snapshotted once at the end of the series (defaults: staged secondary ionization, He recombination coupling on, He $\leftrightarrow$ H charge exchange on). Converged wasp_full reference: activation at step ~ 7144 , final count ~ 13486 , $\log_{10} \dot{M} = 13.22$.

40. du triggers require a descending crossing (2026-08-11)

The $\text{du} < \text{du}_{\text{th}}$ convergence stop and the JFNK hand-off (with the secondary-ionization flip that shares its du test) now fire only after du has been seen at or above their own threshold. They start armed for a Load IC? True run and unarmed for a generated IC. The failure this removes: a transonic IC reported $\text{du}(1) = 1.85 \times 10^{-5}$ on WASP-121 b, which tripped both triggers on step 1 and ran away to a NaN by step 1220; with the guard the same input converges (step 8706, $\log_{10} \dot{M} = 13.21$). The PLM $\rightarrow$ WENO3 switch is deliberately *not* guarded, because the regression cases legitimately switch on step 1 at $\text{du}(1) = 0.37$. No new input keys. Details: docs/newton_scaling_and_base_wall.md §9.

41. JFNK line search decides on the true residual (2026-08-11)

The line search in solve_steady_jfnk accepted steps on the residual with the WENO3 smoothness weights frozen at the *previous* iterate, compared that against a merit memory which mixed the frozen measure with the true one and was refreshed only on acceptance, and reported both $\|R\|$ and the info = 0 verdict from the frozen residual. Measured on the HD 189733 b hand-off state: 7 of 14 accepted steps *raised* the true residual, by factors 1.09 to 9.05, and the stored reference stood $2.4\times$ below the true merit of the iterate the search started from — which is what produced the 12-consecutive-failure abort.

Trials are now evaluated with the weights recomputed at the trial state, the merit memory holds true merits and is rewritten at the start of every outer iteration, and the diagonal scaling is formed before the merit it scales. The frozen weights still define the inner Newton model, which is where they belong. WASP-121 b now reaches $\text{info} = 0$ ($\|R\| = 5.72 \times 10^{-5}$, $\log_{10} \dot{M} = 13.17$, superseding the du-stop 13.21 of the preceding section); `photo_deep_secion_cont` converges in 24 outer iterations instead of 59; no case exhausts its backtracks any more. No new input keys, no golden changed. Details: `docs/newton_scaling_and_base_wall.md` §10.

42. Metal-line cooling at a cold base: line trapping and the coronal-fit validity floor (2026-08-11)

Two physically wrong statements in the metal cooling were corrected. Both bite only where the gas is cold and dense — at the base of a planet whose base falls out of the coronal regime — and the wind is unaffected.

42.1 (a) The escape probability was hardwired to 1

`eval_cool` built a `beta_esc` from the lowest XUV-band *continuum* opacity across *one cell*, then overwrote it with `beta_esc = 1.0d0` and applied the full optically thin metal-line cooling everywhere. Neither quantity is a line optical depth and neither is grid-independent. (AIOLOS `chemistry.cpp:1006` instead scales the same gray depth by an arbitrary 10^8 , driving $\beta \rightarrow 0$ and switching metal-line cooling off; EXHALE had replaced that with the opposite limit.) Measured on the converged HD 189733 b base, the dominant coolant line [O I] 63 μm carries a line-center optical depth $\tau = 2.7$ through the cold layer, i.e. an escape probability of 0.16, not 1.

β is now the line-center escape probability of the line itself, computed from the *column* between the emitting cell and the top of the domain (`Cool_coeff.f90: kappa_OI63, kappa_CII158, line_escape_probability, fine_structure_escape`). The shape is the plane-parallel Doppler result of Hollenbach & McKee (1979) / de Jong, Boland & Dalgarno (1980), renormalized by a factor 2 so that $\beta(0) = 1$ *exactly*: the published form tends to $1/2$ because it counts escape through one face of a slab, whereas here the other direction is absorbed by the lower atmosphere and so removes energy from the modeled gas either way. The two branches are switched at $\tau_c = \sqrt{\pi} e^{a^2/4} = 6.967$, where they cross, so the switch is continuous in value. In the optically thin wind $\beta \rightarrow 1$, reproducing the previous behavior.

β enters as $A_{ul} \rightarrow \beta A_{ul}$ *inside* the two-level solution, not as a factor on its result: in the subcritical limit the cooling is set by the collisional excitation rate and must be independent of β , and multiplying outside would be wrong by a factor β there. Trapping is applied to [O I] 63 μm and [C II] 158 μm only, the two lines for which the code carries an explicit two-level solution; everything else stays optically thin, which is right in the wind and is the residual approximation at the base ($< 0.3\%$ of the base cooling there). A thick resonance line in a metal-rich wind (Mg II h&k) is still treated as thin — recorded, not fixed.

42.2 (b) The CHIANTI coronal fits were evaluated far below their validity floor

Every metal cooling coefficient in `Cool_coeff.f90` is a fit or a table built over 10^3 – 10^5 K. At the 236 K HD 189733 b base, 99.7% of the [O I] rate came from the fit's softest exponential, $e^{-930.111/T}$, whose 930 K corresponds to no [O I]

ground-term splitting at all (the splittings are 227.7 K and 326.6 K); C I is the same case ($e^{-2351.38/T}$ against splittings of 23.6 K and 62.4 K). Below the floor the only excitable metal transitions are the ground-term fine-structure lines, which the explicit two-level terms already carry, saturation included.

coronal_excitation_cutoff(T) now removes the coronal part below 10^3 K as $\exp[-((T_{\text{floor}}/T - 1)/w)^2]$ with $w = 0.5$, and is exactly 1 above it. Value and dT -slope are continuous at the floor (the Brent energy solve and the semi-implicit update both differentiate the cooling in T), and every coefficient is **bit-identical at and above 10^3 K**. The factor is applied to all CHIANTI-derived coefficients — the analytic C/N/O and Mg/Ca/Na/Fe forms and the 1-D/2-D tables, whose $\log_{10} T$ axis starts exactly at 3.0 and which otherwise hold their edge value indefinitely below it — and to the coronal remainder of cool_OI_ne_func / cool_CII_ne_func, but *not* to their two-level parts. The legacy AIOLOS branch (cno_cool_0) is deliberately left alone: its constant floors are crude fine-structure stand-ins, not extrapolated coronal fits. $w = 0.5$ is a modeling choice, not a measurement.

42.3 Measured

HD 189733 b, cell 1: total cooling $3.40 \times 10^{-5} \rightarrow 2.68 \times 10^{-8}$ erg cm $^{-3}$ s $^{-1}$ against a heating of 7.0×10^{-6} , so the base net rate changes sign from -2.65×10^{-5} (cooling) to $+6.97 \times 10^{-6}$ (heating). Almost all of that comes from (b); (a) alone is a factor ~ 5 on the [O I] two-level term. Sweeping the code's own cooling assembly in T at that cell's frozen densities, against the heating measured there, the local radiative balance temperature moves from **168 K to 486 K** — 168.2 K from (a) alone, 485.5 K from (b) alone, so (b) does essentially all of the work at the base — and the two cooling curves are identical from 10^3 K up. A side-by-side marching test does *not* resolve this: on a 5×10^4 -step scale the base is riding an inflow transient whose adiabatic heating dominates the radiative term in both binaries.

WASP-121 b (base at 2358 K, $\tau([\text{O I}] 63 \mu\text{m}) = 0.039$, $\beta = 0.956$) is unchanged: $\log_{10} \dot{M} = 13.17$ before and after, outer mass flux to 3×10^{-7} relative, profile maximum relative difference 8×10^{-5} . make check: mol_base_handoff (metals off) stays byte-identical (PASS); the two metal cases move by the intended physics with $\log_{10} \dot{M} = 13.22$ unchanged in both — wasp_full maximum relative difference 3.7×10^{-7} (density), outer mass flux to 2×10^{-8} ; wasp_he23off 1.6×10^{-3} (density, temperature) in the narrow band at $r = 1.012$ – 1.014 , outer mass flux to 1.2×10^{-5} . Goldens not re-snapshotted at that point. Separately noted while comparing: both goldens already contain **negative species densities** in that band (H I, O I, O II, S II; 17 cells in wasp_full, 3 in wasp_he23off) — a pre-existing defect, unrelated to this change.

42.4 Not fixed by this

The model still has no stellar optical/IR absorption and no thermal background, so nothing sets a radiative floor at T_{eq} for a shielded layer; and below ~ 700 K the guard leaves ions without an explicit two-level term (C I, N, Mg, Ca, Na, Fe) with no cooling at all — for C I that omits the real [C I] 609/370 μm lines, whose LTE rate at the HD 189733 b base is $\sim 10^{-10}$ erg cm $^{-3}$ s $^{-1}$, i.e. 10^{-4} of the local heating. Details: docs/hd189_base_checkerboard.md §10. Files:

```
src/modules/radiation/Cool_coeff.f90
src/modules/radiation/util_ion_eq.f90
src/modules/nonlinear_system_solver/T_equation.f90
```

src/modules/post_process/post_process_adv.f90

43. Base grid resolution as an input key (Base grid [dr,cells], 2026-08-11)

The resolution of the uniform region of the Mixed grid was two hardcoded locals in `define_grid.f90` ($N_{\text{low}} = 50$ cells of $drc = 2.0e-4 R_p$). It is now an input key, because that spacing is what decides whether the base carries the stationary $2\Delta r$ entropy mode diagnosed in `docs/hd189_base_checkerboard.md`: the controlling parameter is the number of cells per base density scale height, $H/\Delta r$ with $H = kT/(\mu g)$, and a $4\times$ refinement was measured to remove the mode.

```
Base grid [dr,cells]:  2.0e-4 50      # the default -- the historical grid
Base grid [dr,cells]:  5.0e-5 200     # 4x refinement at the same 0.01 R_p extent
```

The two numbers share one line, as `du_th [PLM,WEN03]` does, because they are not independent: their product is the radial extent of the uniform region ($0.01 R_p$ by default). Refining at fixed extent means dividing Δr and multiplying the cell count by the same factor; changing only one moves the junction with the stretched region. The second value may be omitted (the count then keeps its default). The key applies to `Grid type: Mixed` only; `Uniform` and `Stretched` build their grids from r_{max} and N alone, and `EXHALE_setup.out` echoes that the key is ignored for them.

`EXHALE_setup.out` also reports the resolution actually achieved, `Base scale-height resolution:  $H(T_{\text{eq}})/dr = 1/(b_0 \cdot dr_j(1))$` (b_0 is the surface Jeans parameter, $dr_j(1)$ the first cell after the Mixed-grid smoothing), and warns below 10 cells. On the default grid this is 26.9 cells for HD 189733 b against 102.3 for WASP-121 b — the margin the high-gravity planet does not have. `Grid type: Stretched` gives 1.7 cells for HD 189733 b, i.e. it is unusable at the base for this class of problem. The smoothing pass at the end of `define_grid` started at the literal 50, which was N_{low} written out; it now follows N_{low} so that the junction between the uniform and stretched regions is smoothed wherever it is.

Two costs, both structural. The CFL step is set by the smallest cell, so a $4\times$ finer base needs about $4\times$ more steps for the same physical time. And the total cell count $N = 500$ is a compile-time parameter, so the cells given to the base come out of the stretched region: on the $4.43 R_p$ HD 189733 b domain the stretch ratio goes $1.0119 \rightarrow 1.0252$ and the cell size at $1.5 R_p$ grows from 6.1×10^{-3} to $1.3 \times 10^{-2} R_p$. Refining the base therefore coarsens the wind, which is where $\dot{M}$ and the transmission spectrum are formed. Making N runtime-settable would need the static $(1-Ng:N+Ng)$ arrays to become allocatable throughout; it was not done here. Recorded while doing this: `count_max = 1000000` is also a compile-time parameter and the `EXHALE_MAXSTEPS` hook only lowers it, so a $4\times$ refined base can reach at most $1/4$ of the physical time the default grid reaches within the cap. On HD 189733 b $4\times$ is therefore not reachable to convergence without raising `count_max`; $2\times$ is, and $2\times$ is already where the mode goes.

Default is bit-exact. `dr_base` is declared with a *default-real* literal ($2.0e-4$, not $2.0d-4$) because that is what the old local was: the stored value is the single-precision neighbor of 2×10^{-4} , 2.5×10^{-8} relative below it. Declaring

the exact double moves every base cell by that amount; measured on `wasp_full` it changed the converged density by 5×10^{-10} relative (step count and $\log_{10} \dot{M} = 13.22$ unchanged), which is not a physics difference but is not byte-identity either. The literal is kept so that an `input.inp` without the key reproduces earlier runs exactly. The consequence to know: writing the default out explicitly (`Base_grid [dr,cells]: 2.0e-4 50`) does *not* reproduce it, because the list-directed read into `real*8` gives the exact double.

Checks. Byte-identity was tested differentially, against the binary built immediately before this change rather than against backup/regression/golden/ (the goldens are stale by the intended metal-cooling physics change of the preceding section). All output files of `wasp_full`, `wasp_he23off` and `mol_base_handoff` are byte-identical including headers, with the same step counts (13488 / 13482 / 12000) and $\log_{10} \dot{M}$ (13.22 / 13.22 / 10.58). WASP-121 b run to completion with the key absent: $\log_{10} \dot{M} = 13.17$, JFNK info = 0, $\|R\| = 5.7 \times 10^{-5}$ — unchanged from the preceding section. Malformed values are rejected in `define_grid` (a cell count outside $[2, N-10]$, or a uniform region that does not fit inside $r_{\max}$).

Details: docs/hd189_base_checkerboard.md §11. Files:

```
src/modules/init/parameters.f90
src/modules/init/define_grid.f90
src/modules/files_IO/input_read.f90
src/modules/files_IO/write_setup_report.f90
```

44. Base ghost temperature, runtime step cap, coronal-cutoff width as input keys (2026-08-11)

Three settings that were compile-time constants or an unexamined default became `input.inp` keys. All three default to the previous behavior, so a file without them runs exactly as before (make check: all three golden cases byte-identical).

44.1 Base ghost temperature: isothermal | continuous

The lower ghost cells pin the density to `rho_bc` and, in the legacy closure, the pressure to `ntot_bc + dp_bc`, i.e. the temperature to T_0 . **The T_0 pin has no physical backing:** the radiative equilibrium of the lower atmosphere is outside the model, so nothing in the code determines the ghost temperature, and T_{eq} is a choice rather than a boundary condition. It becomes actively wrong once the first interior cell settles far from T_0 . With the metal-line cooling of the preceding section, cell 1 on HD 189733 b sits at 494 K against $T_0 = 1183$ K: a factor-2.4 contact discontinuity, plus a $2.3 \times$ density inversion, held permanently on the boundary.

continuous imposes $dT/dr = 0$ at the base face instead. The ghost keeps the base composition — `ntot_bc` nuclei and `dp_bc` electrons at `rho_bc`, exactly the particle count the isothermal pin uses — and carries the cell-1 temperature,

$$p_{\text{gh}} = (\text{ntot_bc} + \text{dp_bc}) T_1, \quad T_1 = p_1 / (n_{\text{tot}} + n_e)_1.$$

$(n_{\text{tot}} + n_e)_1$ comes from `get_species_densities`, the single point where the code decides what counts as a particle. The ghost pressure therefore stays a differentiable function of the interior pressure (what the JFNK line search

needs), while the ionization state it divides by is lagged exactly like every other composition quantity across a hydro step. `BC_component_constrho` is the only place the lower ghost is built, and marching, the JFNK residual and the reconstruction boundary all go through it, so the closure cannot differ between code paths.

The key sets the same quantity as `Hydrostatic base: True` (the ghost pressure, which that key extrapolates from the interior gradient instead), so the two are mutually exclusive: `Hydrostatic base` is tested first and `input_read` warns when both appear. With `Base BC: pressure` the microbar target is imposed at T_0 when n_0 is derived, so under a continuous- T ghost only the base *density* stays anchored; `input_read` notes that too. The closure in effect is echoed in `EXHALE_setup.out`.

Measured on HD 189733 b in the production configuration, restarted from the converged production state and Newton-finished in a scratch copy. The baseline restart reproduces the stored state digit for digit, so the comparison is like for like.

	isothermal	continuous
T_{ghost} / T_1 [K]	1183.0 / 493.6	539.9 / 540.2
$\rho_1 / \rho_{\text{gh}}$	2.33	0.95
$A(\ln \rho)$ cells 1–12	0.0264	0.0092
$A(\ln \rho)$ cells 3–12 / 5–14	4.41e-3 / 1.62e-3	9.28e-3 / 8.95e-4
JFNK	info=0, $\ R\ =5.37\text{e-}4$	info=0, $\ R\ =2.76\text{e-}4$
$\log_{10} \dot{M}$	9.04	9.05

The boundary jump and the density inversion are gone; the extended alternating tail is gone as well (beyond cell 7 the amplitude drops $13\times$ to $420\times$), and what remains is a stronger disturbance confined to cells 2–4. Three further restart+Newton cycles reproduce each state digit for digit, so both are exact fixed points and no further decay is available from cycling. Details, window sweep and scope: `docs/hd189_base_checkerboard.md` §13. **Not** measured at that point: a cold start under the new closure (6 h, 10^6 steps), other planets, transit observables.

Adopted for production. HD189733b/`input.inp` now carries `Base ghost temperature: continuous`, and the planet folder was re-converged with it together with the ionization-root validation of the next section (`run_20260811_rootfix.log`): $\log_{10} \dot{M} = 9.05$, JFNK info = 0, $\|R\| = 1.72 \times 10^{-4}$, 2002 steps. Measured on that converged state, $T_{\text{ghost}}/T_1 = 529.4/529.6$ K, $\rho_1/\rho_{\text{gh}} = 0.95$, and $A(\ln \rho)$ over cells 1–12 is 1.9×10^{-3} , i.e. well below the 0.0092 of the isolated A/B above — the root validation appears to remove a further part of the base disturbance. The paper and `python/paper_data.py` carry 9.05 for this planet.

44.2 Max steps: <N>

`count_max` was an integer, `parameter = 1000000`. The preceding section recorded the consequence: a $4\times$ -refined base, whose CFL step is $4\times$ smaller, cannot reach a converged state within a cap that the default grid already used in full, and the `EXHALE_MAXSTEPS` hook only lowers the cap. `count_max` is now a runtime variable with the same default.

EXHALE_MAXSTEPS keeps its separate meaning (a deterministic mid-loop exit used by the regression harness); it does not raise count_max.

44.3 Coronal cutoff width: <w>

The coronal-excitation guard of the preceding section multiplies every CHIANTI-derived coronal coefficient by $\exp(-x^2)$ with $x = (T_{\text{floor}}/T - 1)/w$ below $T_{\text{floor}} = 10^3$ K. $w = 0.5$ was a hardwired parameter. It is a modeling choice, not a measured quantity, and the temperature the base settles at depends on it at the ~ 100 K level, so it is now settable, still defaulting to 0.5 (division by a variable holding 0.5d0 is the same arithmetic, hence byte-identical). Values ≤ 0 are rejected at parse time. The manual states the sensitivity where the guard is documented.

Checks. make check: wasp_full, wasp_he23off, mol_base_handoff all byte-identical with the keys absent. WASP-121 b run to completion with the keys absent: $\log_{10} \dot{M} = 13.17$, JFNK info = 0, $\|R\| = 5.686 \times 10^{-5}$, 8228 steps, and its output byte-identical to the same case run with a binary built from HEAD without these changes. The parse-dump corpus (backup/regression/run_parse_corpus.sh) gains three lines per case and needs a deliberate golden refresh; it was not refreshed there. Files:

```
src/modules/init/parameters.f90
src/modules/states/Apply_BC.f90
src/modules/functions/composition.f90
src/modules/radiation/Cool_coeff.f90
src/modules/files_IO/input_read.f90
src/modules/files_IO/write_setup_report.f90
```

45. Ionization-equilibrium roots validated against the physical simplex (2026-08-11)

Ion_species.txt carried negative number densities (H II, O II, O III, Fe I) in a thin band just above the base: 17 cells at $r = 1.0106\text{--}1.0140 R_p$ in the wasp_full golden, 15 cells in the WASP-121 b production run, 5 in HD 189733 b.

The equilibrium systems are polynomial in the stage fractions and possess roots outside the physical simplex. At the step where secondary ionization is switched on (step 7186 of wasp_full) MINPACK hybrd1 returned info = 1 on such a root — $x(\text{H II}) = -5.2 \times 10^{-6}$, the Fe stage fractions summing to 1.2209, i.e. an Fe I density 22% below zero. Nothing tested the root, and the next step's warm start re-seeded the solve from the same values, so the state reproduced itself for the remaining 6302 steps. The heating paid for it twice: the metal photoheating channel went negative (-6.6×10^{-7} erg cm $^{-3}$ s $^{-1}$, which is impossible), and the ionized fraction, clamped into [0, 1], sent the Spitzer & Scott (1985) heating fraction f_{heat} to zero, switching off every photoheating channel above the photoelectron threshold. Total heating in the band sat 93% below its neighbors. Clamping the offending fractions to zero is therefore not a repair — the clamp is exactly what produces $f_{\text{heat}} = 0$; the cure has to be a physical root, not a clipped one.

Every atomic cell solve now validates its root — each stage fraction ≥ 0 , and each element's ionized stages summing to at most its nucleus total — and on failure restarts from physically defined starting points: the uncoupled ionization balance at the incoming electron density (ionization_balance_at_fixed_ne, admissible by construction, since the stage populations of a three-stage element come out proportional to $(d_1 d_2, u_0 d_2, u_0 u_1)$), then the optically thick fully neutral limit. A stored state that is not physical no longer seeds the next solve, which is what breaks the self-sticking. A cell where no starting point produces an admissible root is left on the ionization balance and reported rather than accepted. A first attempt that converges to a physical root is accepted unchanged, so healthy cells follow exactly the same solver path as before. q_abs, an absorbed energy rate, must now be positive before it normalizes the heating efficiency.

Gates (single-threaded; a baseline make check on the same tree passed all three cases, so the differences are this change alone). wasp_full FAILs by intent: negative densities 17 cells $\rightarrow 0$, $\log_{10} \dot{M}$ 13.22 $\rightarrow$ 13.22, $T/\rho/v/p$ within 0.89%/0.77%/1.3%/0.13% of the golden away from the band, metal photoheating positive everywhere (domain minimum $-8.9 \times 10^{-7} \rightarrow +1.03 \times 10^{-7}$). wasp_he23off also FAILs: its golden looks clean but the run reports 35 roots outside the simplex over 13478 steps, previously accepted in silence; $\dot{M}$ unchanged, $T/\rho/p$ within 0.28%/0.086%/0.039%. mol_base_handoff is byte-identical. The two production runs, each re-run once with the pre-change binary and once with the fix so the comparison isolates this change, keep their mass-loss rates and step counts (WASP-121 b 13.17, 8228 steps; HD 189733 b 9.05, 2002 steps) and lose their negative densities entirely. HD 189733 b is the clearest demonstration of the self-sticking: its stored initial condition is one of the poisoned outputs, and at step 1 the run rejects exactly the 5 stored cell states that carried negative densities and resolves all of them. Goldenes were not re-snapshotted at that point.

Not covered. The molecular network keeps its own chemical-equilibrium retry and is left alone; the H-only system (no helium) and the post-processing advection solves in post_process_adv.f90 are outside this validation, and the advected profiles of HD 209458 b still carry small negative H II, He III and He 2³S entries in the first 14 cells above the base. Full account: docs/ionization_root_validation.md. Files:

```
src/modules/radiation/ionization_equilibrium.f90
src/modules/radiation/util_ion_eq.f90
src/modules/nonlinear_system_solver/newton_solver.f90
src/EXHALE_main.f90
```

46. Line trapping in the thick metal resonance lines: measured, no change (2026-08-11)

The metal-cooling section above introduced $\beta(\tau)$ for [O I] 63 μm and [C II] 158 μm and recorded the thick resonance lines of a metal-rich wind (Mg II h&k, Ca II H&K, Na I D, the Fe II UV multiplets) as an untreated case. They were measured on the converged WASP-121 b, HD 209458 b and HD 189733 b runs. **No code path changed.**

The lines are thick — Mg II k reaches $\tau_0 = 7.6 \times 10^4$ at the WASP-121 b base and stays above 10^2 through the region where Mg II carries a third of the local cooling, and over 99% of the integrated Mg II / Ca II / Fe II / Mg I cooling comes from gas with $\tau_0 > 1$. That is not the test, though. Trapping lengthens the random walk of a resonance photon; it does not destroy it. In two-level equilibrium the correction to an optically thin coronal fit is

$$S = \frac{\beta A_{ul}}{\beta A_{ul} + n_e q_{ul}}, \quad q_{ul} = 8.629 \times 10^{-6} \Upsilon / (g_u \sqrt{T}),$$

not a factor β , so it bites only above $n_{\text{crit,eff}} = \beta A_{ul} / q_{ul}$. For these permitted lines $A_{ul} \sim 10^8 \text{ s}^{-1}$, so even at $\beta \sim 10^{-6}$ the escape rate keeps $n_{\text{crit,eff}}$ at 10^9 – 10^{15} cm^{-3} , against a maximum n_e of $4.3 \times 10^9 \text{ cm}^{-3}$ in the WASP-121 b run. The forbidden fine-structure lines are the opposite case ($A_{ul} \sim 10^{-5}$ – 10^{-3} s^{-1} , $n_{\text{crit,eff}} \sim 10^0$ – 10^5 cm^{-3}), which is why they, and only they, need β .

Measured $\Delta\Lambda$ integrated over the wind: **0.0214%** of the total radiative losses of WASP-121 b, 0.0022% of HD 209458 b, 0.0015% of HD 189733 b. The largest local effect is 1.65% in the first WASP-121 b cell, whose T is pinned at T_{eq} by the boundary condition; the loss exceeds 0.1% only below $r = 1.007 R_p$. Per line it is 0.1% of the Mg II cooling, 0.015% of Ca II, 1.4% of Na I (its A_{ul} is the smallest of the set) and 0.02–0.8% of Fe II. The measurement is a conservative bound: it uses the Doppler-core escape probability, which understates β by 4–10 $\times$ for these Voigt profiles, and no turbulent broadening.

Checked in passing: the Fe II d^6D fine-structure lines at 25.99/35.35 μm , the one metal transition with an A_{ul} small enough to trap, reach only $\tau_0 \leq 0.06$, so the cool_FeII_ne statistical-equilibrium table needs no escape probability either. H I Ly α , outside this scope, has $\tau_0 = 1.3 \times 10^8$ and a cooling-weighted $S = 0.9994$ — two-level trapping is therefore not what suppresses Ly α cooling, and the destruction channels that could (photoionization of $\text{H}(n=2)$, collisional $2p \rightarrow 2s$) were not evaluated.

$\beta = 1$ for the resonance lines is kept and is now documented at the code site as a justified effective treatment with its validity condition, in the SCOPE comment of `src/modules/radiation/Cool_coeff.f90`, instead of an acknowledged omission. Full measurement: `docs/resonance_line_trapping.md`.

47. Metal electrons and a validity range for the post-process advection correction (2026-08-12)

Two defects in `post_process_adv`, both confined to the advection-corrected `*_adv.txt` output: the post-processor is the last call of the run and its arguments are `intent(in)`, so the converged wind and the equilibrium `*.txt` files cannot be affected by anything here.

The electron density inside the advection residuals was wrong. It counted only the H and He electrons, while the equilibrium residual it corrects counts the metal electrons as well (`metal_electron_sum`) and the *same* routine already used the metal-inclusive n_e for its heating, cooling and temperature solve. In a shielded base the metals are the dominant electron donors — measured $n_{e,\text{metal}}/n_e = 0.48$ – 1.00 in the base cells of all four paper planets — so the recombination terms were low by up to six orders of magnitude. Solved to machine precision with Brent, the residual *as coded* had its fixed point at $x_{\text{HII}} = 4.74 \times 10^{-8}$ in the first HD 209458 b cell, against an equilibrium 7.46×10^{-14} .

A field `adv_cell%xe_metal` now carries the metal electrons per H nucleus into all three advection residuals, with the definition `calc_ne` and `metal_electron_sum` use. It is not conditioned on `eos_include_metals`: that switch decides whether the metals enter the mass and particle budget, whereas recombination needs the true electron density. Metals off gives zero identically.

The correction was applied where it cannot be computed and is not needed. The residuals carry the *neutral* fraction and report the ion density as $(1 - x_{\text{HI}})n_h$, so the ion fraction inherits the solver's absolute resolution on x_{HI} , $\text{xtol} = \sqrt{\epsilon} = 1.5 \times 10^{-8}$. Where the equilibrium ion fraction is 10^{-13} the returned value is quantized at 10^{-8} with an arbitrary sign: the first HD 209458 b cell got $x_{\text{HII}} = -4.96 \times 10^{-8}$ where the exact root is $+4.74 \times 10^{-8}$ — right magnitude, wrong sign — and the upwind cascade carried the negative density outward over 14 cells. Those same cells are in local ionization equilibrium to

$$\text{Da} = \frac{\Delta r}{v} (\Gamma_{\text{HI}} + \alpha_{\text{HII}} n_e) = 350\text{--}2900,$$

so the equilibrium solution is the answer the ODE would give anyway. The correction is now skipped, and the converged equilibrium ionization kept, wherever the cell is (i) inflowing ($v \leq 0$ on either face: the upwind discretization has no upstream cell), (ii) equilibrium-dominated ($\text{Da} > 100$) or (iii) unrepresentable ($x_{\text{HII,eq}} < 10^{-6}$, less than 1% relative accuracy at that resolution). Conditions (i) and (ii) are statements about the flow and (iii) about the representation of the unknown, so the pre-existing `pp_metal_on` gate on the inflow condition is removed — it made a property of the discretization depend on the metal-cooling switch. The temperature loop keeps condition (i) alone, likewise ungated; a thermal Damköhler number is not evaluated. Clipping the extracted densities at zero was tested and rejected: it breaks the H nucleus budget $n_{\text{HI}} + n_{\text{HII}} = n_h$.

Gates. make check: PASS, all three cases byte-identical (the goldens compare `Hydro_ioniz.txt` and `Ion_species.txt`, which this does not write). Isolated A/B, the pristine HEAD binary and the modified one run from the same configuration: `Hydro_ioniz.txt`, `Ion_species.txt`, `Cooling_breakdown.txt` and `Excited_H.txt` byte-identical, only the `_adv` files differ.

planet	$\dot{M}$ [log g/s]	negative _adv	cells held at eq (newly)	outermost held	max change, $r > 1.05$
HD 209458 b	9.31 → 9.31	48 → 0	103 → 124 (21)	1.0316	1.4% (He III)
HD 189733 b	9.05 → 9.05	0 → 0	135 → 212 (77)	1.0380	0.37% (He III)
WASP-121 b	13.17 → 13.17	0 → 0	0 → 0 (0)	—	3.8% (He III)
WASP-52 b	11.63 → 11.63	0 → 0	3 → 57 (54)	1.0106	3.3% (He III)

Every cell whose equilibrium ion fraction is below 10^{-6} is now held at the equilibrium value, and no newly held cell lies above $r = 1.038 R_p$. The change in the wind is entirely the metal electrons: a build carrying the guard alone removes all 48 HD 209458 b negative entries by itself and changes the profiles above $r = 1.05 R_p$ by exactly zero on three of the four planets and by 4.2×10^{-5} on WASP-52 b. Transmission spectra move by at most 0.07% in peak depth (He 10830 of HD 209458 b) and 0.08% in equivalent width ($\text{H}\alpha$ of WASP-121 b).

The re-run reproduces the stored output/ of WASP-121 b and WASP-52 b byte-identically but not that of HD 209458 b and HD 189733 b, which restart from output/*_IC.txt files that had been refreshed to the previous run's converged state; re-running continues that convergence (HD 209458 b du : $1.199 \times 10^{-2} \rightarrow 3.13 \times 10^{-3}$) and moves the profiles

by up to 9% in density at unchanged $\dot{M}$. That is unrelated to this change — the pristine binary follows the same new trajectory — and is why the table above is an A/B between two binaries.

Files:

- `src/modules/post_process/post_process_adv.f90`
- `src/modules/nonlinear_system_solver/ion_cell_state.f90`
- `src/modules/nonlinear_system_solver/System_implicit_adv_{H,HeH,HeH_TR}.f90`

Full account: `docs/postprocess_advection_validity.md`.

Found and **not** fixed, marked at the code site: `System_implicit_adv_HeH` (helium on, `He 23S` off) writes its electron-impact ionization terms without the $n_e n_h$ factor that `adv_implicit_H`, `adv_implicit_HeH_TR` and the equilibrium rows all carry, which also makes them dimensionally inconsistent with the photoionization rates they are added to. None of the four paper planets uses that path.

48. Ground-term fine-structure statistical equilibrium for C I, C II, N II and O I (2026-08-12)

The CHIANTI C/N/O cooling fits carry the ground-term fine-structure (FS) lines in the optically thin, *low-density* limit. Their critical densities are of order 10^0 – 10^5 cm^{-3} , decades below the base density of an irradiated atmosphere ($n_e \sim 10^9$, $n_{\text{HI}} \sim 10^{14}$ cm^{-3}), so at the base the coronal form overstates what those lines can radiate by 2.9–8.3 decades. Two of them ([O I] 63 μm , [C II] 158 μm) already had a two-level saturation; C I, N II and the [O I] 145.5/44 μm channels did not, and C I was in consequence the dominant coolant of the whole base layer of HD 189733 b, at a value comparable to the local heating rate.

What replaces it. Every C/N/O coolant whose ground term is split now has that term solved in exact statistical equilibrium at the local (n_e, n_{HI}) ,

$$\sum_j f_j R_{ji} = f_i \sum_j R_{ij}, \quad \sum_i f_i = 1,$$

$$R_{ul} = C_{ul} + \beta_{ul} A_{ul}, \quad R_{lu} = C_{ul} \frac{g_u}{g_l} e^{-E_{ul}/k_B T}, \quad C_{ul} = n_e k_{e,ul}(T) + n_{\text{HI}} k_{\text{H},ul}(T),$$

$$W_{\text{FS}} = \sum_{u>l} f_u \beta_{ul} A_{ul} k_B E_{ul} \quad [\text{erg s}^{-1} \text{ per ion}],$$

and the coronal curve is refitted to keep only the channels that *leave* the ground term (Λ_{rem}). The coefficient handed to the assembly is $\Lambda_{\text{eff}} = W_{\text{FS}}/n_e + \Lambda_{\text{rem}}$, so the $n_e n_{\text{ion}}$ prefactor recovers the H-collision channel exactly. Treated: C I $2p^2\ ^3P$ (609.1/370.4 μm), C II $2p^2\ ^3P$ (157.7 μm), N II $2p^2\ ^3P$ (205.3/121.8 μm), O I $2p^4\ ^3P$ (63.2/145.5/44.1 μm). That is the complete set — N I and O II have a single-level 4S ground term, Mg I/II, Ca II and Na I a single ground level, Fe I is built from permitted lines only, and Fe II already uses a two-dimensional statistical-equilibrium table.

Limits: $n_e, n_{\text{HI}} \rightarrow 0$ collapses onto the ground level and reproduces the coronal within-term channel to the accuracy of the collision-strength fits (1.3–4.4%); high density saturates at the exact multilevel LTE emission (verified to 10^{-14} relative). β_{ul} enters inside the solution, not as a factor on the result, which is the only placement that is right in both

limits. All eight lines with a transition probability now get their own line-center escape probability from the outward column; previously only two did.

Two defects fixed on the way. (i) The two-level form multiplied a solution already normalized to the two-level partition function by the ground-term Boltzmann fraction of its lower level, so its LTE limit was low by $1 + (g_u/g_l)e^{-E/k_B T}$ — $2.7\times$ for [C II] $158\ \mu\text{m}$ at the 530 K base of HD 189733 b. (ii) The H-atom de-excitation rates (a constant 4.0×10^{-11} for C II, $4.2 \times 10^{-11}(T/100)^{0.67}$ for O I, both flagged approximate in the source) are replaced by fits to published quantum scattering calculations: Yan & Babb (2023, MNRAS 518, 6004) for C I and N II, Abrahamsson, Krems & Dalgarno (2007, ApJ 654, 1171) for O I, Barinova et al. (2005, ApJ 620, 537) for C II. They are $5\text{--}20\times$ larger.

What it does. The coronal_excitation_cutoff guard was standing in for the missing saturation, so its width w behaved as a tuning knob on the base temperature. It no longer does: the local balance temperature of the base cell is identical to every digit printed over $w = 0.02\text{--}1.2$ on all four planet runs, against a 500–650 K spread before.

run	base T	$\dot{M}$ at $2R_p$
HD 189733 b (warm restart, $w = 0.1$)	528.5 $\rightarrow$ 613.2 K	$2.351 \times 10^9 \rightarrow 2.850 \times 10^9$ g/s (+21.2%)
WASP-121 b (cold start)	2406.9 $\rightarrow$ 2438.8 K	$3.106 \times 10^{13} \rightarrow 3.370 \times 10^{13}$ g/s (+8.5%)
wasp_full golden case	2340.6 $\rightarrow$ 2357.3 K	$3.482 \times 10^{13} \rightarrow 3.549 \times 10^{13}$ g/s (+2.0%)

Transit depths on HD 189733 b move by +4.4% (Ca II) to +20.7% (H β). The profiles outside $r \sim 1.1 R_p$ move by a few percent; the change is a base-layer change that propagates through the density.

Files:

- src/modules/radiation/Cool_coeff.f90 — the SE solvers, the four ion coefficients, the generalized line opacity and escape routine
- src/modules/radiation/util_ion_eq.f90
- src/modules/nonlinear_system_solver/T_equation.f90
- src/modules/post_process/post_process_adv.f90
- cooling_data/fit_fs_saturation.py — generator; prints every coefficient
- cooling_data/make_cooling_doc_figures.py
- docs/cooling_formulas.tex, docs/coronal_cutoff_width.md section 7, TO_BE_DONE.md item (C)

Also fixed here, unrelated but found while reading the cooling diagnostics: examples/exhale_io.py named only four fixed channels in Cooling_breakdown.txt while the writer emits six (the collisional-excitation channel is written split by absorber), so every metal column was labeled two ions too early. python/paper_data.py parses the header and was already correct.

Still open: with the artifact gone it is plain that the base is not in local radiative balance at all — at the HD 189733 b base the total radiative cooling is 3.5×10^{-8} against a heating of 3.4×10^{-6} . Its temperature is set by the inner boundary and by the flow.

49. Base ghost built from the composition of the state it bounds (2026-08-12)

Apply_BC closes the Base ghost temperature: continuous ghost with

$$p_{\text{ghost}} = (\text{ntot_bc} + \text{dp_bc}) \frac{p(1)}{n_{\text{part_cell1}}},$$

where `n_part_cell1` is the cell-1 particle count written by `get_species_densities`. Two call sites filled the ghosts before any composition solve had run on the state in question, so they used whatever `n_part_cell1` was left over — the `input_read` placeholder `ntot_bc+dp_bc` in the first case:

- `init.f90`: the initial Apply_BC. On HD 189733 b the placeholder differs from the true cell-1 particle count by 4.4%, and the finite-volume steady residual of the loaded state (which stops right after `init`) reported the resulting contact discontinuity at the base face as a cells 1–2 spike of 60 (mass) and 120 (energy) per sound crossing time.
- `steady_newton.f90`, `eval_residual`: Apply_BC after `unpack_U`. There $F(Y)$ depended on the previous Y as well as on Y , so a finite-difference Jacobian column mixed two states.

Both now evaluate the composition of the interior first. `init.f90` also seeds `base_flux_const` from the initial state, for the same reason: with Base velocity: `massflux` that constant starts at -1 , which Apply_BC reads as “not available yet” before falling back to the valve, so the first step and any diagnostic stopping right after `init` used a different velocity boundary condition than the run. Files: `src/modules/init/init.f90`, `src/modules/time_step/steady_newton.f90`.

Checks. `n_part_cell1` is read only by the continuous branch and `base_flux_const` only by the mass-flux branch, so a run setting neither key is unaffected and make check is byte-identical; an isothermal HD 189733 b residual reproduces bit for bit. With the key on, the cells 1–2 mass residual falls from $5.97 \times 10^1, 5.92 \times 10^1$ to $6.41 \times 10^{-2}, 7.70 \times 10^{-2}$ and the energy residual from $1.12 \times 10^2, 1.20 \times 10^2$ to $1.31 \times 10^{-1}, 1.54 \times 10^{-1}$, with cell 3 outward and the wind window unchanged to every digit. Re-converging the planet from its stored state: JFNK scaled merit $\|F_s\|_2$ $1.43 \rightarrow 0.315$, `info` = 0, $\log_{10} \dot{M}$ $9.15 \rightarrow 9.14$. Measurements and term-by-term split: `docs/hd189_base_checkerboard.md` section 15.

50. $H(n=2)$ rate coefficients, the trapped- $2p$ balance, and de-excitation heating (2026-08-12)

Four corrections to the excited-hydrogen and $\text{Ly}\alpha$ modules, plus the structural change that made the first one possible. Measurement and the full number tables: `docs/lya_destruction_channels.md` section 11.

One definition of the $n=2$ rate coefficients. The $\text{Ly}\alpha$ and $n=2$ atomic data and every collisional rate coefficient existed in three copies — inside `n2_populations`, again as standalone functions in the same file, and a third time in `lya_rt.f90`. They are now one module, `src/modules/radiation/hydrogen_n2_rates.f90`, which

exports `c1s2s_rate`, `c1s2p_rate`, `c2s2p_rate`, `c2s1s_rate`, `c2p1s_rate`, `c2p2s_rate`, `alpha_B_hydrogen`, `alpha_2s_hydrogen`, `alpha_2p_hydrogen`, `n2p_destruction_rate` and the line constants. `excited_hydrogen.f90` and `lya_rt.f90` use it.

1. **Statistical weights shadowed by dummy arguments.** `n2_populations` was declared (`T`, `n1s`, `ne_l`, `Jlya`, `G2s`, `G2p`, `n2s`, `n2p`). Fortran does not distinguish `G2s` from the module parameter `g2s`, so the three detailed-balance rate coefficients written out in the routine body evaluated their weights as the Balmer photoionization rate the caller passed: $2s \rightarrow 1s$ and $2p \rightarrow 1s$ de-excitation came out 5 to 306 times too small (planet-dependent, scaling with the stellar Balmer continuum) and the $2p \rightarrow 2s$ ℓ -mixing rate 3 times too large. The dummies are now `gam_ion_2s / gam_ion_2p` and the body calls the shared rate functions, so the weights are resolved at module scope and out of a dummy's reach. Two further shadows of the same kind, both harmless, were renamed: the frequency-sample count `ng = 400` in `gamma_n2_balmer` and `heat_n2_balmer` shadowed the global ghost-cell count `Ng`; it is `n_nu`.
2. **Recombination cascade source.** $a2s \cdot ne^{**2} / a2p \cdot ne^{**2} \rightarrow a2s \cdot ne \cdot n_{HII} / a2p \cdot ne \cdot n_{HII}$. The electron recombines onto a proton; at the base the electrons come largely from helium and metals while hydrogen is still neutral, and n_e/n_{HII} reaches 2 on WASP-121 b and 5×10^6 on HD 209458 b. The Python transmission tool (`exhale_transit_lib.py`, `EXHALE_transit.py`) and `examples/tpm_alpha_lart2d.py` carried the same expression and were corrected the same way; their `n2_populations` signatures gained an `nHII` argument.
3. **Destruction channels in the `J_int` closure.** `lya_rt.f90` built the trapped field from $n_{2p} = P / (A_{2p \rightarrow 1s} \beta)$. The $2p$ budget also loses atoms to collisional de-excitation, $n = 2$ photoionization, and ℓ -mixing followed by two-photon decay; the denominator is now $A_{2p \rightarrow 1s} \beta + D$ with D from `n2p_destruction_rate`. Where β is smallest the omission over-estimated `J_int` by up to $1.5\times$. In the same module the top-down $\text{Ly}\alpha$ line-center optical depth, previously two identical inline loops, is the single routine `lya_line_center_optical_depth`; all three `jlya_mode` values call it, so the `tau_Lya` column of `Excited_H.txt` is filled in the LaRT import mode instead of being written as zero.
4. **Deexc heat now defaults on.** The comment justifying the old default — that collisional de-excitation heating overlaps the H I collisional-excitation cooling — was wrong: `coex_rate_HI` is the one-way Cen (1992) coronal rate with no de-excitation term in it, so `Hdx` is the correction to that limit, not a double count. Most of the $n = 2$ population it de-excites was pumped by $\text{Ly}\alpha$ rather than excited by a collision, which makes the term absorbed $\text{Ly}\alpha$ thermalized by a collision, the same size as the photoelectric heating the code already applied unconditionally. The key now also honors `Deexc heat: False`, which it previously ignored. **This changes results for any run that supplies a stellar T_{eff} /radius and does not set the key.**

Audited and left alone: the energy equation applies no escape-probability factor to the $\text{Ly}\alpha$ cooling. `cool` is built from the bare Cen coefficient (`Cool_coeff.f90:843` $\rightarrow$ `util_ion_eq.f90:640,787`); β appears nowhere in that path, in any `jlya_mode`. That is the correct form — 98.2–99.99% of trapped $\text{Ly}\alpha$ photons escape and the collisional destruction that would suppress the cooling is at most 0.06% — so the historical expectation of a $\sim 10\times$ transfer suppression is not in the code and should not be.

Checks. `mol_base_handoff` (no stellar T_{eff} , so the excited-H model is off) is byte-identical at every stage, as is its comparison against the stored golden. `wasp_full` and `wasp_he23off` move: $\log_{10} \dot{M}$ 13.23039 $\rightarrow$ 13.22966 and 13.23106 $\rightarrow$ 13.23026, with the largest local temperature change 0.57% (item 1), 0.081% (item 2) and 0.147% (item 3); both cases `pin Deexc heat: True` in their `input.inp`, so item 4 moves nothing in them. Their golden were already stale against the pre-fix tree and were left as they are. Re-converged from their stored initial conditions in a scratch copy, WASP-121 b gives 13.2075 $\rightarrow$ 13.2030, HD 189733 b 9.1286 $\rightarrow$ 9.1392 and HD 209458 b 9.3790 $\rightarrow$ 9.4581; the $H\alpha$ and $H\beta$ transit depths move by 0.4–1.7% ($n(2p)$, which carries the opacity, barely moves, while $n(2s)$ halves). HD 209458 b is the largest move (0.075 dex, from items 1–3) because its $\text{Ly}\alpha$ cooling is a negligible share of its radiative losses while the $n = 2$ heating terms are not; it is also the case `docs/lya_deexcitation_heating.md` recorded as going NaN at the base with the de-excitation heating on, which no longer reproduces with the corrected populations (re-converges from its stored initial condition, and a cold start survives a 20000-step cap). A build at `-O1 -fopenmp -g -fcheck=bounds,do,mem` running the `wasp_full` configuration is clean.

Files:

- `src/modules/radiation/hydrogen_n2_rates.f90` (new)
- `src/modules/radiation/excited_hydrogen.f90`, `src/modules/radiation/lya_rt.f90`
- `src/modules/init/parameters.f90`, `src/modules/files_IO/input_read.f90`, `Makefile`
- `exhale_transit_lib.py`, `EXHALE_transit.py`, `examples/tpm_halpalpha_lart2d.py`

51. Molecular chemistry and trace metals solved in one system (2026-08-13)

The parser refused `Molecular chemistry: True` together with a `metals.inp`. That refusal blocked two things: the deep (10^{-3} bar) lower-atmosphere handoff of `docs/base_composition_handoff_plan.md` §11.7, whose A/B pair carries metals and whose consistent configuration is Tier-2 molecular chemistry; and trace-metal work on sub-Neptunes, which have a molecular base by construction.

The two networks cannot be solved apart, and the reason is quantitative rather than architectural: **in the shielded molecular base the metals supply essentially all the free electrons**. Measured at the base cell of the two smoke tests below, with solar metal abundances:

case	n_e with metals [cm^{-3}]	metal share	n_e metals-off [cm^{-3}]
hot Uranus, solar 7 metals	1.20×10^8	1.000	6.29×10^4
HD 209458 b, 10^{-3} bar handoff, C/N/O	9.42×10^5	0.990	9.27×10^4

The flux that ionizes the low-ionization-potential metals — Mg I at 7.65 eV, Fe I at 7.90 eV, C I at 11.26 eV — lies below both the H I edge (13.60 eV) and the H_2 threshold (15.4 eV), so it is absorbed by neither and reaches the base almost unattenuated. On the hot-Uranus case it holds Mg and Fe about 1.6% ionized at the base, which is already enough to swamp the H/He and molecular electrons. Every recombination and electron-impact term of the molecular network scales with n_e , so solving the network at the metal-free n_e misses a factor 1900 (hot Uranus) to 10 (HD 209458 b at 10^{-3} bar) in that layer.

The merged system. `System_HeH_mol_metals` follows the pattern `System_HeH_TR_metals` established: the molecular unknowns keep their rows, the metals are appended above them, and both blocks are shared code rather than re-derived. The eight balance rows of the molecular network moved into `mol_heh_rows` (still in `System_HeH_mol`, with the Koskinen 2022 rate coefficients, which stay defined in exactly one module) and take n_e as an input, exactly as `heh_tr_rows` does; the metal rows, the metal electron sum and the stage fractions come from `ion_residual_core` unchanged. The layout is

```
x(1..3)      H II, He II, He III
x(4..7)      H2, H2+, H3+, HeH+
x(8)         He 23S (when tracked)
x(mbase..)   X+, X++ for each metal element, mbase = metal_row_base() = 8 or 9
```

with $N_{\text{eq}} = 7$ (+1 with the triplet) + $2n_{\text{melem}}$. `metal_row_base()` is the single definition of that offset, read by the residual, by the driver's seeding and root validation, and by `cx_metal_base` so that the Huang Table-4 charge exchange writes to the shifted metal rows. With `met_nelem = 0` the system reduces exactly to `System_HeH_mol`.

Root validation and seeding (§45 extended to the molecular branch). The molecular cell solve is bistable because the cell is: a molecular basin (the dense, optically thick base) and an atomic basin (the wind above the $\text{H}_2 \rightarrow \text{H}$ front). It now tries one starting point in each in turn — the warm start, the chemical-equilibrium H_2 fraction at the local (p, T) with the metal stages from their own ionization balance at the incoming n_e , then the molecule-free ionization balance of every element — and keeps a converged root only if it is also physical, on the same rule §45 applied to the atomic systems. The metal part of `ionization_balance_at_fixed_ne` was split into `metal_ionization_balance_at_fixed_ne` so the molecular retry seeds its metals from the same balance instead of a second copy of it. The third starting point is what the front cells needed: on the hot-Uranus case the coupled solve leaves 3 cells at $r = 1.21\text{--}1.22$ (the ionization front) with no admissible root at every step with two starting points, and 0 cells with three.

Two defects found in the surrounding code and fixed.

- The sub-13.6 eV energy grid of a power-law spectrum was an `if (thereis_HeITR) ... else if (thereis_lowIP_metal)`, with the comment that the triplet and metals together are unsupported so at most one applies. That has not been true since the triplet and the metals were merged. With the triplet on, the grid floor stopped at 4.80 eV, which is above the K I threshold (4.341 eV), so K I would be held spuriously neutral in a triplet run. The floor is now the lowest threshold over whichever sources are active, matching what the loaded-SED path (`sed_read`) already did. No golden carries K, so this is byte-identical there.
- The `pp_metals 2` metal re-solve passed the global N_{eq} to `hybrd1` while its residual (`ion_system_metals_pp`) writes only rows $1..3 + 2n_{\text{melem}}$. With the He 2³S triplet that already left one row of `fvec` unwritten; with the molecular layout it would leave four or five. It is now sized $3 + 2n_{\text{melem}}$ with a matching workspace, the same reasoning `lwa_adv` already carried a few lines above. Both goldens run `pp_metals 1`, so this is byte-identical there and changes results only for `pp_metals 2` runs.

Left in place at the time, with the validity range written at the code. `eval_cool` built its own electron density without the molecular ions. That is right in the hot atomic gas it was written for and wrong inside a deep molecular base, where H_3^+ can be the dominant ion; there it disagreed with the n_e the equilibrium solver itself uses (ionization_equilibrium does pass `nmol_eq` to `calc_ne`), so the n_e -scaling cooling channels were under-counted in that layer. Two definitions of the same quantity is one too many, but closing the gap changes the molecular results and belonged in its own change with its own measurement. *Closed in the next section.*

Gates. `make check` PASSes byte-identical on all three cases then in the matrix (`wasp_full`, `wasp_he23off`, `mol_base_handoff`); a baseline `make check` on the same tree before the change also passed, so the comparison is this change alone. Golden were not re-snapshotted. A build at `-O1 -fopenmp -g -fcheck=bounds,do,mem` is clean on all three coupled layouts for 200–300 steps: hot Uranus with the $\text{He } 2^3\text{S}$ triplet ($N_{\text{eq}} = 28$, metals at row 9), the same case without it ($N_{\text{eq}} = 27$, metals at row 8), and the HD 209458 b deep handoff.

Hot Uranus (the Tier-2 gate configuration of `docs/lower_atmosphere_figs/data_g2`) plus solar `C/N/O/Mg/Ca/-Na/Fe`, a 12000-step relaxation snapshot against the same case metals-off: the $\text{H}_2 \rightarrow \text{H}$ front sits at $r = 1.1577$ in both (the gate's 1.156), the base stays fully molecular in both, $\log_{10} \dot{M} = 10.58$ in both, no negative densities, and no cell is left without an admissible root (the same case with the triplet off, $N_{\text{eq}} = 27$ with the metals at row 8, is clean over a shorter 400-step smoke). What the metals change is the base chemistry: n_e rises $1900\times$ and H_3^+ falls by a factor 1350 ($7.08 \times 10^4 \rightarrow 52.5 \text{ cm}^{-3}$), its peak moving out from $r = 1.040$ to $r = 1.130$ — H_3^+ is destroyed by dissociative recombination (Koskinen R6/R7, both proportional to n_e), so it survives only where the metal electron fraction starts to fall. Metal ionization runs the expected way: $\text{C II/C} = 8 \times 10^{-10}$ at the base, 0.22 at the front, 0.43 at $1.9 R_p$; Mg III stays below 10^{-8} up to the front and reaches 2.4×10^{-3} only at $1.9 R_p$.

HD 209458 b at the 10^{-3} bar handoff (the §11.7 blocker), `C/N/O`, $\text{He } 2^3\text{S}$ on, a 4000-step bounded run: it integrates. §11.7 records that the atomic pair built at this level terminated with a NaN at $r = 1.067$ after 1846 steps with the time step collapsed to $dtu = 1.2 \times 10^{-6}$ (that run was not repeated here); this one reaches the cap with $dtu = 0.198$ and no NaN, a fully molecular base, and the front at $r = 1.0150$ (1.0169 metals-off). It is a relaxation snapshot at $du = 747$, nowhere near converged, so no mass-loss rate should be read off it.

Open item. The `_adv` post-process is still molecule-free by design (`post_process_adv.f90` header): `nh = nhi + nhii` counts only the free H nuclei and the molecular electrons are absent from its n_e . The metal stages are carried there as usual under `pp_metals`, but on that molecule-free H/He background, so `_adv` metal profiles inside the molecular layer inherit the approximation. With the `eval_cool` electron density closed in the next section, this is the one remaining place where a molecular run and a metal run still meet on the atomic assumption.

Files:

- `src/modules/nonlinear_system_solver/System_HeH_mol_metals.f90` (new)
- `src/modules/nonlinear_system_solver/System_HeH_mol.f90`
- `src/modules/radiation/ionization_equilibrium.f90`, `src/modules/radiation/util_ion_eq.f90`
- `src/modules/files_IO/input_read.f90`, `src/modules/files_IO/write_setup_report.f90` (the startup report now names the species blocks the coupled system carries and prints N_{eq})

- `src/modules/init/set_energy_vectors.f90`, `src/modules/post_process/post_process_adv.f90`,
Makefile
- `examples/16_molecular_metals/ (new)`

Documentation: `README.md`, `README_HOWTO.md`, `examples/README.md`, `docs/EXHALE_user_manual.tex` (the example table also became a longtable, which it needed already — as a float it ran past the bottom margin into the page footer), `docs/input_schema.md`, `docs/lower_atmosphere_coupling.md`, `docs/base_composition_handoff_plan.md`.

52. One electron density for the cooling and for the ionization solver (2026-08-13)

`eval_cool` rebuilt its own free-electron density and left the molecular ions out of it. The equilibrium solver in the same sweep does not: it balances ionization against `calc_ne(nhii,nheii,nheiii,ne,nm,nmol_eq)`, i.e. H^+ , He^+ , He^{++} , the metal stages under the `eos_metals` policy, **and** H_2^+ , H_3^+ , HeH^+ . A molecular run therefore carried two definitions of n_e , and the cooling used the smaller one. That is wrong wherever a molecular ion carries a non-negligible share of the charge, which is exactly the deep molecular base. Measured on the metals-off hot-Uranus Tier-2 gate, the two definitions differ by $168.5\times$ at the base cell (373 against $6.29 \times 10^4 \text{ cm}^{-3}$, H_3^+ supplying essentially all of it), $4.5\times$ at $r = 1.05$, 2.8% at $r = 1.10$, and nothing above $r \sim 1.15$; on the `mol_base_handoff` case the base ratio is $175.7\times$.

`eval_cool` now takes an optional `nmol` (H_2 , H_2^+ , H_3^+ , HeH^+ , cgs) and passes it straight to `calc_ne` — the electron density is not reconstructed anywhere inside `eval_cool`, it has one definition and one implementation. Callers that track the molecular network supply it: `ionization_equilibrium` (`nmol_eq`), the semi-implicit energy solver (`nmol_dim`, the cgs copy of the adimensional array it already hands `calc_ne/calc_ntot`), and the `Cooling_breakdown` dump. The `_adv` advection post-process deliberately does not: that reconstruction is molecule-free by design (its module header), and omitting the argument gives it the atomic charge sum it is built on. `ionization_equilibrium`'s own `calc_ne` call lost its `thereis_mol` branch at the same time — `nmol_eq` is zero for an atomic run, so one call covers both.

Every n_e -scaling channel picks the correction up: recombination, collisional ionization, collisional excitation and bremsstrahlung are linear in n_e , and the coronal metal line coefficients enter as `ne*nm*c_metal`. The saturated ground-term coefficients (Fe II, and C I/C II/N II/O I under `cno_chianti`) are $\Lambda_{\text{eff}} = W/n_e$, so the assembly's n_e cancels and they respond only through the level populations ($C_{21} \propto n_e$) — the correct behavior, and there is no double counting of n_e anywhere in `eval_cool`.

Atomic runs are exactly unchanged. `nmol_eq` and the molecular `f_sp` columns are zero, and `calc_ne` adds them as `ne + 1.0*0.0`, which is exact. `make check` is byte-identical on `wasp_full` (metals + $He\ 2^3S$) and `wasp_he23off`; a full WASP-121 b run (metals, $He\ 2^3S$, $Ly\alpha$ escape probability, excited H, transonic IC) reproduces all eight output files byte-identically — `Hydro_ioniz{,_adv}`, `Ion_species{,_adv}`, `Cooling_breakdown`, `Heating_breakdown`, `Excited_H`, `IC_dump` — at $\log_{10} \dot{M} = 13.20$.

`mol_base_handoff` changes, and that is the point of the change. At the pinned 12000-step snapshot the two runs stop at the same $du = 2.6689$ and the same dtu ; every hydrodynamic column agrees to $\leq 4 \times 10^{-8}$ relative and every

species column to $\leq 4.2 \times 10^{-6}$ (that worst case is H_3^+ at $r = 2.96$, where H_3^+ is $3 \times 10^{-19} \text{ cm}^{-3}$). The one column that moves is the reported cooling rate at the near-neutral base, $2.799 \times 10^{-20} \rightarrow 4.942 \times 10^{-18} \text{ erg cm}^{-3} \text{ s}^{-1}$ ($\times 176.5$, the n_e ratio, as expected for channels linear in n_e). $\log_{10} \dot{M}$ stays 10.58.

Why the dynamics barely notice. A molecular layer cools through the H_3^+ infrared lines, which do not scale with n_e , so multiplying the electron channels by 176 at the base moves the cooling from 2.8×10^{-20} to 4.9×10^{-18} against $1.3 \times 10^{-7} \text{ erg cm}^{-3} \text{ s}^{-1}$ of photoheating there. The correction is real and it is now the same quantity the solver balances ionization against, but it is energetically negligible in the layer where it is large.

Molecular smoke tests. Hot Uranus, the Tier-2 gate configuration, 12000-step relaxation snapshots — both runs stop at du of order unity, so these are snapshots and not converged solutions.

	metals off	metals on (solar C/N/O/Mg/Ca/Na/Fe)
$n_e(\text{base}) \text{ old} \rightarrow \text{new}$	$373 \rightarrow 6.29 \times 10^4 \text{ cm}^{-3} (\times 168.5)$	$1.2016 \times 10^8 \rightarrow 1.2016 \times 10^8 (\times 1.0000)$
cooling at the base	$3.18 \times 10^{-20} \rightarrow 5.36 \times 10^{-18} \text{ erg/cm}^3/\text{s}$	$+7 \times 10^{-7} \text{ relative}$
$\text{H}_2 \rightarrow \text{H}$ front	$1.15631 \rightarrow 1.15631$	$1.15677 \rightarrow 1.15677$
H_3^+ peak	$7.0808 \times 10^4 \rightarrow 7.0807 \times 10^4 \text{ cm}^{-3}$ at $r = 1.0396$	$52.465 \rightarrow 52.465$ at $r = 1.1304$
du at step 12000	$4.3354 \rightarrow 4.2914$	$5.8177 \rightarrow 5.8177$
$\log_{10} \dot{M}$	$10.58 \rightarrow 10.58$	$10.58 \rightarrow 10.58$

The metals-on column is the expected null: the low-ionization-potential metals hold n_e at 1.2×10^8 at that base (previous section), against 52 cm^{-3} of H_3^+ , so adding the molecular ions moves nothing (maximum relative change 7×10^{-10} in $\rho/T/p$, 3×10^{-9} in the metal stages). Metals-off, the snapshot differs by 1% in du and by sub-percent in the bulk state just above the layer ($\rho -0.7\%$, $T +0.8\%$ at $r = 1.19$), with the trace species that depend exponentially on T amplifying that into H_3^+ -16% , $\text{H}_2 -13\%$, $\text{He II} +10\%$ at the same radius. Those are differences between two snapshots of an unconverged relaxation, not a converged shift; the gate observables — front position, H_3^+ peak, mass-loss rate — are unchanged.

On the HD 209458 b 10^{-3} bar handoff with C/N/O (a 4000-step bounded snapshot, also unconverged) the base picture is the same: the metals dominate n_e , the base cooling rises 0.45%, the molecular-layer cooling by at most 1.9%, the front moves $1.01503 \rightarrow 1.01518$ and the H_3^+ peak by 0.05%. An `-01 -fopenmp -g -fcheck=bounds,do,mem build` runs all three molecular layouts (hot Uranus with and without metals, HD 209458 b deep handoff) for 300 steps with no runtime check firing.

Two related items found and recorded, not changed here.

- The H_3^+ infrared cooling is added to cool in `ionization_equilibrium`, *after* `eval_cool` returns. The JFNK steady residual uses that array, but the semi-implicit energy update (the default) rebuilds cool from `eval_cool` alone and overwrites it, so it never sees the H_3^+ term. For a molecular run the marching stage and the Newton stage therefore integrate slightly different energy equations, and the marching fixed point is not the zero of the residual the Newton solver drives down. The fix is to move the term inside `eval_cool` (which already receives `nmol`) with its own `cool_chan` column; that changes the `Cooling_breakdown.txt` schema and the molecular

results, so it belongs in its own change. The Cooling_breakdown header now states that the channel sum falls short of the cool column by that term in the molecular layer. Open item (F) of TO_BE_DONE.md.

- H_2^+ , H_3^+ and HeH^+ donate their electron to n_e but are left out of the bremsstrahlung charge sum, so that sum is smaller than n_e inside the molecular layer. They are $Z_{\text{ion}} = 1$ and belong there formally, but they only exist at $T \sim 10^3$ K, where the hot-plasma free-free expression is an extrapolation emitting at frequencies the atmosphere is not thin to, and where free-free is negligible against the H_3^+ and metal line cooling. The scope is written at the accumulator.

Files:

- src/modules/radiation/util_ion_eq.f90, src/modules/radiation/ionization_equilibrium.f90
- src/modules/time_step/energy_semi_implicit.f90
- src/modules/post_process/post_process_adv.f90 (comments only)

Documentation: docs/lower_atmosphere_coupling.md.

53. A fourth regression case for the molecular + metals system (2026-08-13)

The coupled molecular + metals system of the section before last was not guarded by anything in the matrix: mol_base_handoff runs the molecular network metals-off, and the two wasp_* cases run metals without molecules, so the merged residual, the metal_row_base() offset and the three-starting-point seeding had no golden.

backup/regression/mol_metals is that case: the same hot-Uranus Tier-2 gate configuration as mol_base_handoff (Molecular chemistry: True, its base.inp handoff, 12000 steps pinned in maxsteps) plus a solar-abundance metals.inp with the seven elements C/N/O/Mg/Ca/Na/Fe. It stops at $du = 3.1770$, $\log_{10} \dot{M} = 10.58$. DEFAULT_CASES in backup/regression/run_check.sh is now wasp_full wasp_he23off mol_base_handoff mol_metals, and the case list in the script header describes what each one guards.

Goldens for all four cases were re-snapshotted at the end of this change series, so they carry the two preceding sections. The two atomic cases are byte-identical to their previous goldens (both changes are exact no-ops without molecules); mol_base_handoff moves only in its cooling column, by the factor 176.5 of the previous section.

One harness trap, hit while doing this and now written into the script: golden does **not** re-run the cases, it copies whatever output/ currently sits in each case directory. Snapshotting straight after a code change baselines the *previous* binary. The order is check (which rebuilds and re-runs), then golden, then check again.

Files: backup/regression/run_check.sh, backup/regression/mol_metals/ (new), backup/regression/golden/ (re-snapshotted). Documentation: the workspace CLAUDE.md, README_HOWTO.md.

A Code-size summary: EXHALE vs. the original ATES (2026-06-07; updated 2026-07-02)

A line-by-line comparison of the Fortran source against the pristine upstream code (ATES/ATES-Code-main) quantifies the extent of the in-house extension. *Pure code lines (SLOC)* exclude blank lines and comments (the comment stripper is string-aware, so a `!` inside a quoted string is not treated as a comment); the line-by-line diff additionally normalizes whitespace and case, so reformatting and comment-only edits are *not* counted as changes. Updated 2026-07-02 after the Newton steady solver, the He 2^3S +metals merge, the composition module, He/H+metal diffusive separation, and the Wind-AE IC port; the *core* rows exclude the ported Wind-AE solver (listed separately).

	Fortran files	code lines (SLOC)
Original ATES (ATES-Code-main)	47	3 514
EXHALE core (excl. Wind-AE port)	62	9 245
net change (core)	+15	+5 731 ($\approx 2.6\times$)
Wind-AE IC-generator port (wind_ae/)	+36	+4 746
EXHALE total	98	13 991 ($\approx 4.0\times$)

Line-by-line diff, core (comments + blanks excluded; whitespace/case normalized): **5 970** code lines added in EXHALE (new modules + the new side of modified lines) and **239** removed from the original (deleted + the old side of modified lines), for **6 209** total changed lines. At the file level, of the 47 original files **27 were modified** and **20 are unchanged** (none were removed — the unchanged set is chiefly the 11 included MINPACK routines, the HLLC/ROE Riemann solvers with the flux/reconstruction pieces, and small utility modules), and **15 new core modules** were added:

New module	Purpose	SLOC
time_step/steady_newton.f90	JFNK steady-state solve (Newton finish)	540
radiation/charge_exchange.f90	charge-exchange network (Huang Table 4)	314
radiation/excited_hydrogen.f90	H($n=2$) model + Balmer feedback	305
functions/species_diffusion.f90	He/H (+ metal) diffusive separation	216
radiation/opacity_models.f90	opacity dispatcher (A/C/P/T)	180
nonlinear_system_solver/System_HeH_metals.f90	coupled H/He+metals ionization (+analytic Jacobian)	171
nonlinear_system_solver/newton_solver.f90	analytic-Jacobian Newton (+hybrd1 fallback)	124
files_IO/metals_input_read.f90	parse metals.inp	111
time_step/energy_semi_implicit.f90	semi-implicit energy solver	106
init/species_table.f90	metal-species metadata table	89
files_IO/opacity_input_read.f90	parse opacity.inp	85
nonlinear_system_solver/System_HeH_TR_metals.f90	merged He 2^3S +metals system	84
radiation/lya_rt.f90	Ly α escape-probability RT	68
functions/composition.f90	single-point species/EOS policy	46
time_step/steady_residual.f90	steady-residual assembly/norms	45
new-module total		2 484

So the 15 new core modules contribute 2 484 of the 5 970 added code lines; the remaining $\sim 3\,486$ added (and 239 removed) lines are edits spread across the 27 modified files — dominated by the cooling module (`Cool_coeff.f90`, 442 $\rightarrow$ 2 230 raw lines: CHIANTI metal-line fits, Fe tables, temperature-dependent Penning rate), followed by `EXHALE_transit.py`-side physics and the main driver, parameters/input parsing, IC construction, cross sections, ionization equilibrium, and the post-processor, which wire the metals, opacity dispatcher, Roche potential, transonic/auto IC, excited-H/Ly α physics, diffusion hooks, and the Brent/Newton solver upgrades into the existing solver. In total the EXHALE core is now $\approx 2.6\times$ the original code base (3 514 $\rightarrow$ 9 245 SLOC), touching a bit over half of the original files (27 of 47 modified, none deleted, plus 15 new); with the ported Wind-AE IC generator the full Fortran tree is 98 files / 13 991 SLOC ($\approx 4\times$). Beyond Fortran, the tree also carries the Python tooling (GUI, the `EXHALE_transit.py` transmission post-processor, plotting/analysis) and $\approx 20\,000$ lines of documentation under `docs/`.

B Fortran source inherited unchanged from ATES

Of the 47 upstream Fortran files (`ATES/ATES-Code-main`), **18 are byte-for-byte identical** in EXHALE — the numerics EXHALE reuses without touching. They fall into three groups: the included MINPACK nonlinear solver (the hybrid Powell hybrid method and its linear-algebra helpers, called by the ionization-equilibrium solver), the finite-volume hydrodynamics core (HLLC fluxes, reconstruction, source terms, Runge–Kutta right-hand side), and two radiation/ionization pieces. Table 2 lists them with their role. Two further files — the HLLC and ROE wave-speed estimators (`flux/speed_estimate_HLLC.f90`, `speed_estimate_ROE.f90`) — are *effectively* untouched: their original logic is preserved verbatim and only a one-line defensive floor `p_min=max(p_min, 1e-30 p_max)` (plus its comment) was appended in the 2026-07-02 code review, so no original line was altered.

Table 2: The 18 Fortran files byte-identical between the original ATES (ATES-Code-main) and EXHALE, grouped by role.

File	Purpose
<i>MINPACK nonlinear-system solver</i> (Argonne; dir nonlinear_system_solver/)	
hybrd.f90	Powell hybrid method: find a zero of n nonlinear equations in n unknowns
hybrd1.f90	Simplified user-facing driver for hybrd
fdjac1.f90	Forward-difference approximation of the Jacobian
qrfac.f90	QR factorization via Householder transformations
qform.f90	Accumulate the orthogonal factor Q from the QR factorization
r1updt.f90	Rank-one update of the QR factorization
r1mpyq.f90	Apply the Givens-rotation sequence of the rank-one update
dogleg.f90	Dogleg (trust-region) step combining Gauss–Newton and steepest descent
enorm.f90	Euclidean norm of a vector (overflow/underflow safe)
dpmpar.f90	Machine floating-point precision / range constants
<i>Hydrodynamics & grid core</i> (finite-volume Godunov)	
flux/Num_Fluxes.f90	Numerical interface flux assembly (HLLC Riemann fluxes)
functions/UW_conversions.f90	Conserved $\leftrightarrow$ primitive variable conversions
states/PLM_rec.f90	Piecewise-linear (2nd-order) cell reconstruction
states/Source.f90	Gravitational + geometric source terms
time_step/RK_rhs.f90	Runge–Kutta right-hand side (convection + source)
init/set_gravity_grid.f90	Pre-evaluate gravity at cell centers and interfaces
<i>Radiation & ionization</i>	
radiation/J_inc.f90	Energy-dependent incident stellar spectrum for photoionization
nonlinear_system_solver/System_implicit_adv_H.f90	H-only advection-coupled implicit ionization system (He/metal work went into the HeH/HeH_TR variants, leaving this untouched)