

EXHALE: Update Log

***Purpose.** This document is a running log of incremental updates to EXHALE as the code is extended toward reproducing Huang et al. (2023) WASP-121b (see Huang_update_plan.md). Each update is a self-contained section: what changed, why, which files were touched, and the quantitative check that was run. New updates are appended as new top-level sections.*

Contents

1. Runtime domain-extent option: Roche vs Spherical	11
1.1 Overview	11
1.2 The two domain modes	11
1.2.1 Gravitational potential	11
1.2.2 Outer boundary	12
1.2.3 Why spherical mode must drop the tidal terms	12
1.3 Input interface and backward compatibility	12
1.4 Implementation	12
1.5 Validation on WASP-121b	13
1.6 Status and relation to the Huang plan	14
2. Ionization network: add magnesium (Mg I / II / III)	14
2.1 Overview	14
2.2 Atomic data and rates	14
2.3 Touched files	15
2.4 Validation on WASP-121b	15
2.5 Status and relation to the Huang plan	16
3. Phase 1a: interface / data-structure refactor (no new physics)	16
3.1 Overview	16
3.2 What changed (Steps A–F)	17
3.3 Gate: bit-for-bit reproduction of the C/N/O + Mg run	18
3.4 What is deliberately <i>not</i> generalized yet	18
3.5 Touched files	18
3.6 Status and relation to the Huang plan	19
4. Phase 1b: correct C/N/O cross sections and add Si, Ca, Na, K, S	19
4.1 Overview	19
4.2 Variable-stage solver (2- and 3-stage elements in one packing)	19
4.3 Atomic data and rates (sources, not fabricated)	19
4.4 Rename: System_HeHCO → System_HeH_metals	20

4.5	Validation on WASP-121b	20
4.6	Caveats and gaps	21
4.7	Regression reference for Phase 1c	21
4.8	Status and relation to the Huang plan	21
5.	Phase 1c: add iron (Fe I / II / III)	22
5.1	Overview	22
5.2	Atomic data and rates (sources, not fabricated)	22
5.3	Touched files	23
5.4	Iron abundance	23
5.5	Validation on WASP-121b	23
5.6	Caveats and gaps	24
5.7	Regression reference for Phase 1d / Phase 2	24
5.8	Status and relation to the Huang plan	24
6.	Phase 1d: charge exchange with hydrogen (Huang Table 4)	25
6.1	Overview	25
6.2	Atomic data and rates (sources, not fabricated)	25
6.3	Touched files	26
6.4	Validation on WASP-121b	26
6.5	Caveats and gaps	27
6.6	Regression reference for Phase 2	27
6.7	Status and relation to the Huang plan	27
7.	Phase 2: metal line cooling — validation gate closed	28
7.1	Active line coolants	28
7.2	Channel-by-channel cooling diagnostic (this phase)	28
7.3	Validation on WASP-121b (the gate)	28
7.4	Touched files	29
7.5	Caveats and deferred	29
7.6	Status and relation to the Huang plan	29
8.	Phase 3a: in-code excited H(n=2) + Balmer proton source and heating	29
8.1	Physics and atomic data (sources, not fabricated)	30
8.2	Coupling into the solver	30
8.3	Runtime interface	30
8.4	Diagnostics	31
8.5	Touched files	31
8.6	Validation on WASP-121b (coupling ON vs OFF)	31
8.7	Caveats and gaps	32

8.8	Two analysis gotchas (recorded so they are not repeated)	33
8.9	Status and relation to the Huang plan	33
9.	Phase 2 refinement: Fe I line cooling (NIST f-values + Van Regemorter)	33
9.1	Atomic data and method (sources, not fabricated)	33
9.2	Touched files	34
9.3	Validation on WASP-121b	34
9.4	Status and relation to the Huang plan	34
10.	Phase 2 refinement: density-dependent Fe II line cooling (SE table)	35
10.1	Why the coronal table is wrong at the base	35
10.2	The density-dependent coefficient (multilevel statistical equilibrium)	35
10.3	Step 1 — the 2-D table generator (done)	36
10.4	Step 2 — the 2-D table and bilinear interpolator in Cool_coeff.f90 (done)	36
10.5	Step 3 — wiring the density-dependent Fe II into the cooling assembly (done)	37
10.6	Step 4 — WASP-121b run and the cooling budget (numbers for the reader)	37
11.	Post-process refinement: re-solve metal ionization in the advected wind (pp_metals 2)	38
11.1	Implementation — reuse the coupled residual with H/He pinned	39
11.2	The cell-by-cell sweep in post_process_adv	39
11.3	Clean A/B and what changes	40
12.	Phase 3b: in-code Lyα radiative transfer ($\bar{J}_{\text{Ly}\alpha}$ via escape probability)	41
12.1	The escape-probability field (lya_rt.f90 : jlya_escape_prob)	42
12.2	WASP-121b validation setup (spherical Case A, matched to Huang)	42
12.3	Trend vs Huang Fig. 11 (qualitative — exact match not expected)	43
13.	Phase 4: RLOF / tidal potential (Case B validated; Case D blocker = escape radius, transonic IC added, geometry-limited)	43
13.1	Base gravity for each case (Huang Table 3)	44
13.2	Case B — validated	44
13.3	Case D — real blocker found (escape radius), transonic IC added	44
14.	Phase 5: velocity-broadened transmission spectrum (in progress)	45
14.1	5a — metal resonance lines (Mg II, Ca II, Na I D), spherical — done	46
14.2	5b — Roche-equipotential reconstruction + broadening — implemented	46
14.3	5b — direct Case A comparison: hydro validated, metals over-extended	47
15.	Post-process advection guard at a “breathing” base (2026-06-07)	48
16.	He 2³S triplet + metals simultaneous solve, Lyα-flux guard, and convergence workflow (2026-06-08)	49

17. Metals in the bulk-gas mass/charge/particle budget (eos_metals, 2026-06-13)	52
18. Automatic initial-condition selection (IC mode: auto, 2026-06-13)	53
19. Base boundary conditions and the flux-based convergence decision (2026-06-15)	54
20. Tidal (Roche) potential with an extended outer boundary (2026-06-16)	56
21. Photoionization cross-section benchmarks and a TOPbase high-energy extension of He 2 ³ S (2026-07-01)	56
22. He I (1 ¹ S) photoionization → Verner+1996 by default, with ATES-rate switches (2026-07-01)	58
23. He 2 ³ S photoionization → two VFXY96 (Verner) forms + power-law bridge (2026-07-01)	59
24. Temperature-dependent Penning ionization of He 2 ³ S (2026-07-01)	59
25. He/H diffusive separation — Phase 1 (2026-07-01, experimental, default OFF)	60
26. He/H diffusion — Phase 2 (2026-07-02): ambipolar, thermal diffusion, metals	61
27. Diffusive separation: detailed physics and numerics	62
27.1 Governing equation	62
27.2 Conservative implicit discretization	62
27.3 The $n_{\text{He}}-n_{\text{H}}$ mass coupling: why a self-consistent n_{H} is <i>not</i> needed	63
27.4 The near-base pile-up and the physical cap	64
27.5 Settling discretization: a Péclet hybrid	64
27.6 Ambipolar field, thermal diffusion, and metals	64
27.7 The metal-rescale ratchet bug (found in review, fixed 2026-07-02)	65
27.8 Boundary conditions and validation	65
28. Lower-atmosphere connection: Tier 1–3 (2026-07-02)	65
29. Tier-2 molecular chemistry core: H ₂ /H ₂ ⁺ /H ₃ ⁺ /HeH ⁺ (2026-07-03)	66
30. Newton–diffusion co-convergence, solver hygiene, and the Ly α lower-BC audit (2026-07-03)	66
31. Tier-3 first light: VULCAN → base.inp (2026-07-03)	67
32. VULCAN lower-atmosphere pre-step: obtained, not included in the tree (2026-07-06)	67
33. TPM.py renamed to EXHALE_transit.py (2026-07-06)	68
34. Sub-level-resolved Balmer opacity and LaRT 2s inclusion (2026-07-06)	68
35. Molecular equilibrium seeding fix, molecular-base coupling, Tier-2 gate refresh (2026-07-16)	68

36. Default H/He rates → Badnell+Mao / Voronov, van Hoof free-free Gaunt factor, free-free charge fix, and secondary ionization (2026-07-17)	69
36.1 What changed and why	70
36.2 Files touched	71
36.3 Quantitative check	71
36.4 Effect on existing results	71
37. He recombination radiation → H I ionization coupling (use_he_rec_coupling, 2026-07-17)	72
37.1 What changed and why	72
37.2 Files touched	73
37.3 Quantitative check (HD 209458 b)	73
38. Staged activation of secondary ionization (2026-07-23)	76
38.1 The regression	76
38.2 The fix: converge first without the coupling, then switch it on	76
38.3 The Immediate override	76
38.4 Files touched	76
38.5 Effect on existing results (wasp_full validation)	77
39. He-triplet physics series from the Falorca & Vidotto (2026) review (2026-07-23)	77
40. du triggers require a descending crossing (2026-08-11)	78
41. JFNK line search decides on the true residual (2026-08-11)	78
42. Metal-line cooling at a cold base: line trapping and the coronal-fit validity floor (2026-08-11)	79
42.1 (a) The escape probability was hardwired to 1	79
42.2 (b) The CHIANTI coronal fits were evaluated far below their validity floor	79
42.3 Measured	80
42.4 Not fixed by this	80
43. Base grid resolution as an input key (Base_grid [dr,cells], 2026-08-11)	81
44. Base ghost temperature, runtime step cap, coronal-cutoff width as input keys (2026-08-11)	82
44.1 Base ghost temperature: isothermal continuous	82
44.2 Max steps: <N>	83
44.3 Coronal cutoff width: <w>	84
45. Ionization-equilibrium roots validated against the physical simplex (2026-08-11)	84
46. Line trapping in the thick metal resonance lines: measured, no change (2026-08-11)	85
47. Metal electrons and a validity range for the post-process advection correction (2026-08-12)	86

48. Ground-term fine-structure statistical equilibrium for C I, C II, N II and O I (2026-08-12)	88
49. Base ghost built from the composition of the state it bounds (2026-08-12)	90
50. $H(n=2)$ rate coefficients, the trapped- $2p$ balance, and de-excitation heating (2026-08-12)	90
51. Molecular chemistry and trace metals solved in one system (2026-08-13)	92
52. One electron density for the cooling and for the ionization solver (2026-08-13)	95
53. A fourth regression case for the molecular + metals system (2026-08-13)	97
54. H_3^+ infrared cooling inside <code>eval_cool</code> : one energy equation for a molecular run (2026-08-13)	97
55. The infrared field of the lower atmosphere (Base IR field) (2026-08-13)	101
56. H_2 photodissociation in the Lyman–Werner bands (Stellar LW flux) (2026-08-13)	102
57. A fifth regression case for Lyman–Werner photodissociation (2026-08-13)	105
58. JFNK convergence is judged on the iterate the solver returns; <code>load_IC</code> no longer accepts an all-zero metal column (2026-08-13)	105
59. The molecular equilibrium no longer inherits a cell’s previous composition (2026-08-13)	107
60. Damping the stagnant layer: a gated fourth difference inside the numerical flux (Low-Mach damping) (2026-08-13)	109
61. The residual below the escape radius, and what closes the HD 209458 b molecular + metals case (2026-08-15)	112
62. Runtime grid size, CODATA k_B , the layer residual on a physical scale, and two diagnostic closures (2026-08-15)	115
63. The rest of the physical constants, and why iron cannot move to Badnell (2026-08-19)	117
64. Resolved-configuration output, and the transit tool reading it (2026-08-22)	119
65. The metastable helium triplet is on by default (2026-08-22)	120
66. Primitive conversion of unfilled ghost cells (2026-08-24)	121
67. Element diffusion on the direct-steady route, the element drift metric, and the restart of a diffused state (2026-08-25)	121
68. Binary H/He element diffusion replaces the trace-helium kernel (2026-08-25)	122

69. Milestone M3: the binary operator on a wind, the LHS 1140 b helium-rich runs, and the removal of the trace kernel (2026-08-25)	125
69.1 T2b: the binary operator against the trace kernel on HD 209458 b	125
69.2 T8: the elemental face flux on a moving wind	126
69.3 M3-2: LHS 1140 b with diffusion	127
69.4 The trace kernel is removed	128
70. The base helium hole, the metal drain beside it, and a composition relaxation that converges (2026-08-25)	128
70.1 What was measured	128
70.2 (a) The advection was upwinded on face-averaged velocities	129
70.3 (b) The trace-metal solver drained the same cell	129
70.4 (c) The relaxation ran on the CFL step, and on a flow with no mass flux	130
70.5 The runs, re-run from the same restart files	130
70.6 Where the +0.21 dex comes from: metal transport, tested directly	132
70.7 The $K_{zz} = 0$ non-convergence is the wind, not the composition	132
70.8 Verification	134
71. The friction the ionized wind actually has: stage-resolved binary diffusion (2026-08-25)	134
71.1 What the coefficient is now	134
71.2 Two defects fixed with it	135
71.3 Tests	136
71.4 Measured effect	136
72. The advection correction is gated on the slowest species, not on hydrogen (2026-08-25)	137
72.1 What the gate is now	137
72.2 Measured	138
72.3 Effect on existing results	139
72.4 Re-post-processing the stored LHS 1140 b scans (2026-08-26)	139
73. Milestone M4: element diffusion through the molecular region, and the end of the molecular / He_diffusion exclusion (2026-08-26)	140
73.1 What the molecular region changes, and what it does not	140
73.2 T7, and the gate it closes	141
73.3 The exclusion, and the restart	141
73.4 End to end	142
73.5 Regression	143
73.6 The diffusion regression case (same day)	143

74. What a base.inp key is allowed to do: five categories, and elemental reservoirs that reach melem_ab	143
(2026-08-26)	
74.1 The contract	143
74.2 Elemental reservoirs	144
74.3 The gate: do the reservoirs hold?	144
74.4 Legacy	145
75. The He 2³S double count: an excited level is not a second species (2026-08-27)	146
75.1 Where the second count was	146
75.2 The fix	146
75.3 What moved	146
75.4 The gate	147
75.5 Effect on a converged wind	147
76. The lower atmosphere as a profile, not six scalars (2026-08-27)	147
76.1 The file and the key	148
76.2 K_{zz} stops being a constant	148
76.3 The scalar file is no longer a second source	149
76.4 What the run records	149
76.5 The overlap window is empty, and that is a finding	149
76.6 Tests	151
76.7 Two things fixed on the way	152
77. Two producers for the profile handoff: the Photochem adapter and its VULCAN cross-check (2026-08-27)	152
77.1 One way, and the file says so	152
77.2 What the fingerprint is over	153
77.3 Elemental accounting, and the trap that makes it a rule	153
77.4 The radius column	153
77.5 Trap 1 was met and is handled the way P1 handles it	153
77.6 Failure policy	154
77.7 Measurements	154
77.8 Fingerprint round trip	155
77.9 EXHALE reads what the adapter wrote	155
77.10 What is not done	156
78. The climate step, the cold trap, and LHS 1140 b through the profile (2026-08-27)	156
78.1 What the codes actually do about condensation, measured	156
78.2 The tropopause is solid, its pressure is not	157
78.3 What the trap does to the handoff	157

78.4	LHS 1140 b through the profile: it converges, and T-E8 closes	158
78.5	Profile route against the scalar base, and where the difference lives	158
78.6	What is not done	159
78.7	Beside it: a particle is not a hydrogen atom (vulcan_to_base.py)	159
79.	The elemental-flux closure: the composition becomes an output (2026-08-27)	160
79.1	The trial flux now reaches the chemistry	160
79.2	The window the closure is measured on, and why the obvious one fails	160
79.3	The driver	161
79.4	LHS 1140 b, closed	162
79.5	T-E8, and one thing it caught	163
79.6	T-E9: not verified, and why	163
80.	A restart onto a reservoir the handoff has moved (2026-08-27)	164
80.1	The rule	164
80.2	What it closes	165
80.3	What is unchanged, checked rather than argued	165
80.4	$r_{EI} = 1$ is not an identity, and the reason is worth stating	166
81.	Is the wind collisional where it is set? A Knudsen diagnostic, and LHS 1140 b's answer (2026-08-27)	166
81.1	The collision model is not a new one	166
81.2	Definitions, all stated	167
81.3	The result, measured 2026-08-27 on existing outputs	167
81.4	The kinetic estimate is a scale, not a bound	168
82.	The third body was a mass density: a molecular-network audit in the He-dominated limit (2026-08-27)	168
82.1	No rate coefficient was wrong	168
82.2	What was wrong was the density handed to the three-body reactions	169
82.3	A guard that was not guarding	169
82.4	The closures hold	169
82.5	Four caveats are recorded rather than patched	169
82.6	Does this move the LHS 1140 b results?	170
82.7	Goldens	170
83.	The outer region a converged wind inherits: path dependence beyond the exobase (2026-08-27)	171
83.1	The two solutions	171
83.2	Where they differ	171
83.3	The mechanism, read from that pattern	172
83.4	The physical judgement, stated first	172
83.5	What this says about the convergence criteria	173

83.6	What this says about running a scan	173
83.7	What was and was not tested	173
84.	Metals had no way back: element projection when the helium mass fraction reaches one (2026-08-28)	173
84.1	The path, in the code	174
84.2	The symptom that led here	174
84.3	Why <code>info = 0</code> and $\ R\ \sim 10^{-4}$ could not see it	174
84.4	A third failure mode, not either of the first two	175
84.5	The repair, and the abundance it uses	175
84.6	The check the pipeline did not have	176
84.7	The clamp is a limiter, and the comment that said otherwise was wrong	176
84.8	What was measured	177
84.9	Verification	178
85.	A structure scale that is a scale of the state, not of the pressure (2026-08-28)	179
86.	The drift flux has to vanish where the composition runs out: the discretization of $X(1-X)$ and the bounds on the helium mass fraction (2026-08-28)	180
86.1	Where the old discretization lost the bound	180
86.2	The discretization that keeps it	181
86.3	Newton, and why not another Picard sweep	182
86.4	What is measured	182
86.5	Golden movement	183
86.6	The physical result this was worried about: LHS 1140 b does not move	183
86.7	Verification	184
86.8	Two more things found in passing, both in the step diagnostic	184
87.	A rate coefficient that jumped by 1.57 at 4000 K: the He (2^3S) Penning rate, and the step it put in every figure (2026-08-28)	185
88.	A population solved as a difference, and a post-process that never read <code>info</code> (2026-08-29)	189
89.	A conservation check that could not see conservation: the deepest-level elemental test and the reservoir band it refused (2026-08-29)	193
90.	The equilibrium residual removed at its source, and the observable that did not move (2026-08-29)	195
91.	The clima bracket was not the cause: a Jacobian built from round-off, and a failure that predates the patch (2026-08-29)	198
92.	A column that came back 2.4 % apart in pressure: the grid the solver stops on, not the temperature it started from (2026-08-29)	201

93. Rows of one residual that differed by a million, and a third body corrected only halfway (2026-08-29)	203
A Code-size summary: EXHALE vs. the original ATES (2026-06-07; updated 2026-08-30)	209
B Fortran source inherited unchanged from ATES	211

1. Runtime domain-extent option: Roche vs Spherical

Added 2026-06-04.

1.1 Overview

ATES historically hardcodes the full Roche potential (planetary gravity plus the stellar-tidal and centrifugal terms) and truncates the radial grid at the Hill/L1 radius. For close-in planets this Roche lobe is small; for WASP-121b the outer boundary lands at $r_{\max} \approx 2.01 R_p$. Huang et al. (2023) Case A, by contrast, is a *spherical* model (no tidal terms) extending to $\sim 19 R_p$.

To support both geometries from a single executable, a runtime option was added with two modes:

- **Roche** (default): full Roche potential, domain truncated at the Hill/L1 radius. Identical to the previous ATES behavior.
- **Spherical**: pure planetary potential $\phi = -b_0/r$ (tidal and centrifugal terms dropped), domain extended to a user-specified outer radius.

The mode is selected by two optional lines appended to `input.inp`; when they are absent the code reverts to Roche, so pre-existing input files are unaffected.

1.2 The two domain modes

1.2.1 Gravitational potential

All quantities are in code-normalized units: radius r in units of the planetary radius $R_0 = R_p$, with $b_0 = GM_p\mu/(k_B T_0 R_0)$ the surface Jeans parameter, $M_r = M_*/M_p$ the star-to-planet mass ratio, and $\tilde{a} = a_{\text{orb}}/R_0$ the orbital distance in planetary radii.

Roche mode. The dimensionless potential and its gradient are

$$\phi(r) = -\frac{b_0}{r} - \frac{b_0 M_r}{\tilde{a} - r} - \frac{b_0(1+M_r)}{2\tilde{a}^3} \left(\frac{\tilde{a} M_r}{1+M_r} - r \right)^2, \quad (1)$$

$$\frac{d\phi}{dr} = \frac{b_0}{r^2} - \frac{b_0 M_r}{(\tilde{a} - r)^2} + \frac{b_0(1+M_r)}{\tilde{a}^3} \left(\frac{\tilde{a} M_r}{1+M_r} - r \right). \quad (2)$$

The three terms are planetary gravity, stellar tidal attraction, and the centrifugal potential of the co-rotating frame.

Spherical mode. Only the planetary term is kept:

$$\phi(r) = -\frac{b_0}{r}, \quad \frac{d\phi}{dr} = \frac{b_0}{r^2}. \quad (3)$$

1.2.2 Outer boundary

- **Roche:** $r_{\max} = (3M_r)^{-1/3} \tilde{a}$, the Hill-radius approximation to the L1 distance (unchanged from before).
- **Spherical:** $r_{\max}$ is set directly by the user-supplied “Outer radius [R_p]” value (in R_p , since the grid is normalized to $R_0 = R_p$).

1.2.3 Why spherical mode must drop the tidal terms

The outer boundary cannot simply be pushed past L1 while keeping the full Roche potential. Beyond L1 the Roche potential *turns over*: the effective gravity changes sign (points toward the star), so a monotonic, steadily accelerating outward wind is no longer defined there and the sonic-point/steady-state treatment breaks down. Removing the stellar-tidal and centrifugal terms restores a monotonic potential ($\phi = -b_0/r$ increases monotonically toward zero as $r \rightarrow \infty$), which is what permits the domain to extend arbitrarily far. This is exactly the spherical geometry of Huang Case A. A guard enforces a sensible configuration: selecting Spherical mode without an “Outer radius [R_p]” > 1 stops the run with an explicit error rather than silently falling back.

1.3 Input interface and backward compatibility

Two optional lines are appended at the end of `input.inp`, after the “Force start:” line:

```
Domain mode: Spherical
Outer radius [R_p]: 19.0
```

- “Domain mode:” accepts `Spherical`; any other value (or a missing line) means `Roche`.
- “Outer radius [R_p]:” is read only in `Spherical` mode; it is ignored in `Roche` mode, where the Hill radius is computed automatically.

The reader scans the trailing lines by *keyword* (using `index(line, . . .)`), not by fixed line position, so blank trailing lines and line ordering do not matter, and `input.inp` files written before this change (which lack both lines) keep the default Roche behavior. The active mode is echoed in the setup report `EXHALE_setup.out`.

1.4 Implementation

Touched files.

- `src/modules/init/parameters.f90`: new globals `spherical_domain` (logical, default `.false.`) and `r_out_user` (the outer radius in R_p).
- `src/modules/functions/grav_field.f90`: `phi` and `Dphi` branch on `spherical_domain`.

- `src/modules/files_IO/input_read.f90`: keyword scan of the trailing lines, and the $r_{\max}$ branch with the validity guard.
- `src/modules/files_IO/write_setup_report.f90`: prints the active mode to `EXHALE_setup.out`.

Gravity branch.

```
double precision function phi(r_in)
real*8, intent(in) :: r_in
if (spherical_domain) then
  phi = -b0/r_in
else
  phi = -b0/r_in &
    -b0*Mrapp/(atilde-r_in) &
    -b0*(1.0+Mrapp)/(2.0*atilde**3.0)* &
    (atilde*Mrapp/(1.0+Mrapp)-r_in)**2.0
endif
end function phi
```

1.5 Validation on WASP-121b

Two runs were made that differ *only* in the domain-extent option; `input.inp` was otherwise identical (power-law SED, $\text{He/H} = 0.0851$, trace metals C/N/O on, “Rate/2 + Mdot/2” 2-D flux approximation). Both reached steady state via the proper criterion (momentum constant, $du < du_{\text{th}}$), not a stall or a false convergence.

mode	$r_{\max}$ [R_p]	$\log_{10} \dot{M}$	sonic [R_p]	$r(50\% \text{HII})$	T_{peak} [K]
Roche (default)	2.01	12.68	1.825	1.178	9662
Spherical	19.00	11.81	5.270	1.152	10752
Huang Case A (ref.)	~ 19	12.57	—	—	~ 12000

$\dot{M}$ in $\log_{10} \text{g s}^{-1}$. Convergence took $\approx 5,392$ iterations (Roche) and $\approx 154,183$ (Spherical, the larger domain).

Observations (for evaluation, not pass/fail).

- The Roche run reproduces the prior Phase-0 Case A result exactly ($\log_{10} \dot{M} = 12.68$), confirming the new optional block does not perturb default behavior.
- Spherical T_{peak} (10,752 K) is higher than Roche (9,662 K) and closer to Huang’s $\sim 12,000$ K — the extended domain lets the profile develop past the $2-R_p$ L1 clip.
- Spherical $\dot{M}$ (11.81) is lower than both Roche (12.68) and Huang (12.57), consistent with the loss of tidal enhancement, but ~ 0.76 dex below Huang.
- The sonic point moves outward from 1.83 to 5.27 R_p (gentler acceleration without tidal forces).

Caveats. Two confounders remain before the spherical $\dot{M}/T$ can be compared to Huang cleanly: (a) only C/N/O metals are included here, not Huang’s full Mg/Fe/Si/Ca/Na/K/S set (Phase 1); and (b) the same $N = 500$ cells now span $1 \rightarrow 19R_p$, so the inner resolution is much coarser than in the Roche run — a denser-grid spherical re-run may be needed before drawing conclusions.

1.6 Status and relation to the Huang plan

This option closes the Phase-0 geometry gap: the Roche mode is the validated baseline, and the Spherical mode provides the Case A geometry needed for later comparison. The tidal RLOF physics of Huang Phase 4 is a separate, richer treatment (substellar streamline, L1 location, terminator mapping); the present Roche mode is the existing single-streamline truncation, not the full RLOF model. The Spherical mode is the immediate tool for Case A reproduction.

2. Ionization network: add magnesium (Mg I / II / III)

Added 2026-06-05.

2.1 Overview

The first trace metal beyond C/N/O, following the Huang et al. (2023) WASP-121b species set (Mg, Fe, Si, O, C, N, S, Ca, Na, K). Per the Phase-1 plan, elements are added one at a time (order Mg $\rightarrow$ Fe $\rightarrow$ Si $\rightarrow$ Ca $\rightarrow$ Na $\rightarrow$ K $\rightarrow$ S); this section covers **Mg only**. Magnesium is solved to second ionization (Mg I / Mg II / Mg III) inside the existing coupled MINPACK ionization system, raising the species count to 18 and the ratio-unknown count to 11. The change is confined to the ionization/radiation plumbing: there are no hydrodynamic or mass-density changes, and `calc_ne/calc_rho` still count only H/He — metal electrons enter solely inside the MINPACK system.

2.2 Atomic data and rates

Sourced, not fabricated:

- **Photoionization:** Verner et al. (1996) outer-shell fits, with Mg I threshold 7.646 eV and Mg II threshold 15.035 eV (`cross_sec.f90: sigma_MgI, sigma_MgII`).
- **Recombination:** Badnell RR+DR — the same source already used for C/N/O — via `lookup_rec` cases and `rec_MgII ( $\rightarrow$  Mg I) / rec_MgIII ( $\rightarrow$  Mg II)`. This is a *deliberate deviation* from the validation-rate prescription in `Huang_update_plan.md` §1.3 (Shull & van Steenberg 1982 for Mg); it corresponds to the plan’s Phase-6 *production* rates, applied early to match the existing C/N/O treatment. The Mg ion/neutral balance below should be read as the Badnell result, not a Huang-faithful one.
- **Collisional ionization:** Voronov (1997) fits, `ion_coeff_MgI` and `ion_coeff_MgII`.
- **Line cooling:** *not yet ported* — `cool_MgI = cool_MgII = 0`. No published Mg two-level/CHIANTI cooling fit is implemented (that is Phase 2). Magnesium therefore acts only as an electron donor and a photoelectric-heating channel here, not as a coolant.

- **Charge exchange Mg + H⁺**: omitted in this first cut.

Grid extension below 13.6 eV. Because Mg I ionizes at 7.646 eV, turning Mg on triggers the existing low-IP-metal branch (`thereis_lowIP_metal`), which extends the energy grid and the stellar SED down to 7.646 eV (`set_energy_vectors.f90`). This is required to photoionize Mg I, but has a side effect noted under Caveats.

2.3 Touched files

- `init/parameters.f90`: `n_species = 18`; `e_th_MgI/e_th_MgII`; `X_Mg`; `s_mgi/s_mgii`; `thereis_lowIP_metal`.
- `functions/cross_sec.f90`: `sigma_MgI`, `sigma_MgII`.
- `radiation/Cool_coeff.f90`: `Badnell` fits, `lookup_rec`, `rec_MgII/rec_MgIII`, `Voronov` ion_coeff_MgI/MgII, `cool_MgI/MgII` (= 0, flagged).
- `radiation/util_ion_eq.f90`: `PH_heat_HHe` and `eval_cool` carry Mg.
- `radiation/ionization_equilibrium.f90`: Mg densities, params slots 49–55, MINPACK unknowns 10/11 (Mg II/Mg III ratios), solution unpack.
- `init/set_energy_vectors.f90`: 7.646 eV sub-grid; allocate/fill `s_mgi/s_mgii`.
- `files_IO/input_read.f90`: `X_Mg` default; `N_eq = 11`; `metals/low-IP` flags. `files_IO/metals_input_read.f90`: Mg I abundance key. `files_IO/write_output.f90`: Mg I/II/III as `Ion_species.txt` columns 17–19.
- Threading through arrays, I/O and post-processing: `init/set_IC.f90`, `files_IO/load_IC.f90`, `init/init.f90`, `EXHALE_main.f90`, `time_step/energy_semi_implicit.f90`, `functions/utilities.f90`, `post_process/post_process_adv.f90`.

`post_process_adv.f90` also received a latent-bug fix: its advection systems fill only the H/He equations, so they now use a local `Neq_adv` rather than the global `N_eq = 11`, which would otherwise feed `hybrd1` uninitialized metal entries.

2.4 Validation on WASP-121b

Four runs with identical `input.inp` (power-law SED, He/H = 0.0851, Roche domain), differing only in which metals are present. All reached steady state via $du < du_{\text{th}}$. Mg/H is conserved at the input $X_{\text{Mg}} = 3.98 \times 10^{-5}$ across the domain.

run	metals	grid floor	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$
no metals	H/He only	13.6 eV	12.92	13722	1.703
Mg only	Mg	7.646 eV	12.92	13682	1.697
C/N/O	C, N, O	13.6 eV	12.68	9662	1.199
C/N/O + Mg	C, N, O, Mg	7.646 eV	12.76	9828	1.224
Huang Case A	full set	—	12.57	~12000	—

$\dot{M}$ in $\log_{10} \text{ g s}^{-1}$; r in R_p .

Observations (for evaluation, not pass/fail).

- **Clean Mg isolation** (Mg only vs no metals): $\Delta \log_{10} \dot{M} = +0.000$ dex, $\Delta T_{\text{peak}} = -40$ K. Adding magnesium alone leaves the wind essentially unchanged, as expected for a 4×10^{-5} trace species with no line cooling: it neither heats nor cools appreciably, and the soft-UV band it adds is optically thin without other absorbers.
- **Mg ionization structure**: predominantly Mg II (Mg^+) through the lower thermosphere, predominantly Mg III (Mg^{2+}) in the upper thermosphere (Mg III fraction ~ 0.47 at $1.4 R_p$ rising to ~ 0.81 at $1.9 R_p$); neutral Mg I stays subdominant. Same qualitative ordering as Huang+2023 Fig. 12.
- **MINPACK convergence**: the 11-unknown system converges without NaNs across the full temperature range.
- Mg-only $\dot{M}$ (12.92) is $\sim 2.2\times$ the Huang Case A spherical value (12.57), in the same range as the prior C/N/O Roche run; the metal-set and geometry confounds from Phase 0 still apply.

Caveats and gaps.

- **C I grid-extension confound** (C/N/O + Mg vs C/N/O). The $\dot{M}/T$ offset between these two (+0.080 dex, +165 K) is *not* a magnesium effect: extending the grid to 7.646 eV exposes **C I** (threshold 11.26 eV) to the new 11.26–13.6 eV photons the 13.6 eV-floor C/N/O run never sees, adding C I photoionization heating. This is why the clean Mg measurement uses Mg-only vs no-metals, where C is absent.
- **No Mg line cooling** ($\text{cool_MgI} = \text{cool_MgII} = 0$). Huang’s headline Mg result — Mg II h&k emission dominating the cooling at $1.15\text{--}1.4 R_p$ — cannot appear here by construction. That is Phase 2.
- **No Mg + H^+ charge exchange** yet.
- **Recombination is Badnell**, not the §1.3 validation rate.

2.5 Status and relation to the Huang plan

This completes the **Mg** step of Phase 1 (network expansion). The ionization structure satisfies the qualitative Fig. 12 ordering and the “ $\dot{M}$ within a factor ~ 2 of Case A” part of the gate; the thermal-balance part of Phase 1 is deferred to Phase 2 (Mg cooling). Next element in the planned order: **Fe**. Scratch runs and the comparison notebook live in `mg_validation/`.

3. Phase 1a: interface / data-structure refactor (no new physics)

3.1 Overview

Each metal added so far (C/N/O, then Mg in §2) widened the named-argument lists of `PH_heat_HHe`, `eval_cool`, `calc_column_dens_metals`, `write_output`, the MINPACK system `ion_system_HeH_metals`, and the driver

ionization_equilibrium.f90 by several arguments for each ion, and threaded the MINPACK coefficients for each cell through a hand-indexed flat params(60). This step converts that plumbing to a **species-metadata table + array-indexed (2D) data flow**, so that adding an element becomes “add rows to a table” instead of “thread N new arguments through M routines.” **No physics changes** — the network is still C/N/O + Mg, and the result is gated bit-for-bit against §2 (below).

3.2 What changed (Steps A–F)

- **A — Metadata module** (init/species_table.f90, new). One canonical table holds metadata for each ion (length $n_{\text{mion}} = 12$): f_sp column (mion_fsp), parent element (mion_elem), stage / charge² (mion_stage, mion_z2), photo-ionizable flag and photo-table column (mion_isphot, mion_iphot), threshold (mion_ethr), cooling flag (mion_iscool); and metadata for each element (length $n_{\text{melem}} = 4$): nuclear charge (melem_Z), neutral-ion index (melem_i0), top stage (melem_top). The canonical ion order (C I..Mg III = f_sp columns 7–18) is defined here once; element indices iel_C/O/N/Mg.
- **B — write_output 2D.** Metal ion densities pass as $\text{nm}(:, 1:n_{\text{mion}})$ and are written with an implied-do over the table $((\text{nm}(j,i)*n0, i = 1, n_{\text{mion}}))$, not a named write for each ion.
- **C — PH_heat_HHe + column density 2D.** Optical depth, photoheating, and the photoionization rates loop over photo-ionizable ions via $\text{sigma_tab}(:, \text{mion_iphot}(i))$ and Nm_col , returning a 2D rate array $\text{P_m}(:, 1:n_{\text{mion}})$ (inert top stages stay zero) instead of separate P_CI..P_MgII arguments, one for each ion.
- **D — eval_cool metal channels 2D.** Recombination and collisional-ionization coefficients return as 2D arrays (rec_m, aion_m); bremsstrahlung sums $\text{mion_z2}(i)*\text{GF_elem}(:, \text{mion_elem}(i))*\text{nm}(:, i)$ over ions. No separate coefficient arguments for each ion.
- **E — Generalized MINPACK system.** ion_system_HeH_metals keeps hybrd1’s fixed (N_eq, x, fvec, iflag, params) signature but (i) receives the metal coefficients for each cell through a **module-level block** in System_HeH_metals (set_metal_coeffs stores $\text{met_ntot}/g0/g1/b0/b1/a1/a2$), set by the driver before each hybrd1 call — safe because the ionization-equilibrium cell loop is serial — and (ii) assembles fvec by **looping over elements** (two balance equations each, force-zeroing absent elements). params now carries only H/He (1–11), He ITR (12–18), and the C/N/O charge-transfer rates (40–48); the old slots 49–55, one for each element, are gone. The driver builds metal densities, initial guesses, and the solution unpack by looping over $n_{\text{melem}}/n_{\text{mion}}$. The system size is now the formula $N_{\text{eq}} = 3 + 2*n_{\text{melem}}$ (= 11 for the current four elements), set in input_read.f90.
- **F — Generalized grid floor.** set_energy_vectors.f90 no longer hardwires the sub-13.6 eV floor to Mg’s 7.646 eV; it computes $e_{\text{sub_low}} = \min$ over active elements of the neutral threshold $\text{mion_ethr}(\text{melem_i0}(e))$ that lies below $e_{\text{th_HI}}$. The trigger flag thereis_lowIP_metal is set generically in input_read.f90 (any active element with a sub-13.6 eV neutral threshold), driven by a new canonical-order abundance array melem_ab (declared in parameters.f90, populated in input_read.f90). For the current set only C I (11.26 eV) and Mg I

(7.646 eV) lie below 13.6 eV and the minimum is Mg I, so the floor is unchanged at 7.646 eV when Mg is present. (Consequence: a future C/N/O-without-Mg run would now extend the grid to 11.26 eV, which it did not before — and once Phase 1b adds Na/K/Ca, whose neutral thresholds 5.14/4.34/6.11 eV sit below Mg's, the floor drops further.)

3.3 Gate: bit-for-bit reproduction of the C/N/O + Mg run

With C/N/O + Mg active and the same input `.inp/metals.inp` as §2, the refactored code reproduces the `mg_validation/cno_mg` reference to round-off: $\log_{10} \dot{M} = 12.76$; `Hydro_ioniz.txt`, `Hydro_ioniz_adv.txt`, and `Ion_species.txt` are byte-identical to the frozen reference. `Ion_species_adv.txt` differs at exactly one number — line 1 (innermost ghost cell), column 7 (the HeITR fraction): the reference holds 4.0048×10^{-5} left in uninitialized stack memory while the deterministic build writes 0.0; every other column of that line and every other line match. This is a pre-existing reference artifact (a HeITR ghost-cell slot never used in the H/He + metals run), not an effect of the refactor.

3.4 What is deliberately *not* generalized yet

- **Charge exchange with H** is still hardcoded for C/N/O only (the `fvec(1)` C,N,O sum and the `cxlo/cxup` mapping via `iel_C/O/N` in `System_HeH_metals`). Generalizing it is Phase 1d.
- **Atomic data for each ion and abundance for each element remain manual entries** — the `sigma_tab(:, 1..8)` assignments and the `sigma_XX` cross-section functions in `set_energy_vectors.f90`, the abundance scalars `X_C/X_O/X_N/X_Mg`, the `metals.inp` keys, and the `melem_ab` population block. These are the atomic-data / abundance rows expected to grow per element in Phase 1b/1c; they are data, not interface plumbing.

3.5 Touched files

- **New:** `init/species_table.f90` (Step A).
- `init/parameters.f90`: `melem_ab(:)` abundance array; generalized `thereis_lowIP_metal` comment.
- `files_IO/input_read.f90`: `N_eq = 3 + 2*n_melem`; `melem_ab` population; generalized low-IP-metal trigger.
- `files_IO/write_output.f90`: 2D metal-column writer (B).
- `radiation/util_ion_eq.f90`: `PH_heat_HHe` + column density 2D (C); `eval_cool` metal channels 2D (D).
- `nonlinear_system_solver/System_HeH_metals.f90`: module-level metal block for each cell + element-loop residual (E).
- `radiation/ionization_equilibrium.f90`: metadata-driven density build, params packing (H/He + charge transfer only), initial guesses, solution unpack (E).
- `init/set_energy_vectors.f90`: min-over-active-metals grid floor (F).

3.6 Status and relation to the Huang plan

Phase 1a is complete: the element-by-element plumbing of §2 is now table-driven. Adding the Phase-1b metals should require only metadata rows plus atomic/abundance data, with **no changes to `System_HeH_metals.f90` or `ionization_equilibrium.f90`**. Next: **Phase 1b** — Si, Ca, Na, K, S added together as a uniform-template batch. Gate scratch run: `/tmp/ates_gate/cno_mg/` vs frozen reference `mg_validation/cno_mg/output/`.

4. Phase 1b: correct C/N/O cross sections and add Si, Ca, Na, K, S

4.1 Overview

Two changes are made together here, both enabled by the Phase 1a (§3) table-driven plumbing:

1. **Corrected the C/N/O photoionization cross sections** to the full Verner, Ferland, Korista & Yakovlev (1996; VFKY96) Eq. (1) form. The previous code used the simpler Verner & Yakovlev (1995) inner-shell form (no y_0/y_1 offset/asymptote parameters) for C/N/O; that form is exact only when $y_0 = y_1 = 0$, which does not hold for the C/N/O outer shells. This is the dominant *physics* change.
2. **Added the five uniform-template metals Si, Ca, Na, K, S** in one batch. With Mg (§2) the in-code element set is now C, N, O, Mg, Si, Ca, Na, K, S — the full Huang+2023 §2.4 set *except Fe* (Phase 1c). Counts: $n_{\text{species}} = 30, n_{\text{mion}} = 24, n_{\text{melem}} = 9$.

The metal MINPACK solver was also **renamed** (below).

4.2 Variable-stage solver (2- and 3-stage elements in one packing)

Huang+2023 carries Si, Ca (like Mg) to second ionization but Na, K, S only to first. The solver keeps a uniform packing of **two unknowns for each element** ($N_{\text{eq}} = 3 + 2 \times n_{\text{melem}} = 21$ with all nine elements) and a `melem_top` flag for each element:

- `melem_top = 2` (C, O, N, Mg, Si, Ca): solve both $X^0 \leftrightarrow X^+$ and $X^+ \leftrightarrow X^{2+}$.
- `melem_top = 1` (Na, K, S): solve only $X^0 \leftrightarrow X^+$ and **pin the unused upper unknown to zero** (`fvec(ix+1) = x(ix+1)`), keeping the Jacobian non-singular. Every `ix+1` access is guarded by `melem_top` so a two-stage element never touches its neighbor's column.

Zero-abundance elements are still force-zeroed, so the same 21-unknown system solves C-only, C/N/O+Mg, or all nine without renumbering.

4.3 Atomic data and rates (sources, not fabricated)

- **Photoionization cross sections** (`functions/cross_sec.f90`): full VFKY96 Eq. (1) form (`sigma_VFKY96`) for C I/II, N I/II, O I/II, Si I/II, Ca I/II, S I; simpler Verner & Yakovlev 1995 form (`sigma_Verner96`, $y_0=y_1=0$) for Mg I/II, Na I, K I (their published fits have $y_0=y_1=0$, so the forms agree exactly). Neutral thresholds [eV]: Si I 8.152, Ca I 6.113, Na I 5.139, K I 4.341, S I 10.36.

- **Collisional ionization:** Voronov (1997) for all new ions.
- **Recombination:** Badnell RR+DR via the swappable dispatcher, as for C/N/O and Mg — the *same deliberate deviation* from the Huang_update_plan.md §1.3 validation rates already accepted for Mg (the Phase-6 production rate applied early). Read the balances as the Badnell result.
- **Line cooling:** *not* ported (`mion_iscool = .false.`), flagged To Be Checked; the five act only as electron donors and photoelectric-heating channels (Phase 2).
- **Charge exchange with H:** *not* added; still C/N/O only (Phase 1d).

4.4 Rename: System_HeHCO → System_HeH_metals

The metal solver was named after the AIOLOS heritage (AIOLOS tracked only C and O as metals $\Rightarrow$ “HCO” = H, C, O). It now carries H + He + nine metals, so the C/O-specific name was misleading. Renamed to parallel the existing System_HeH/System_HeH_TR family: `module System_HeHCO → System_HeH_metals`; `ion_system_HeHCO → ion_system_HeH_metals`; `set_hco_metals → set_metal_coeffs`; `hco_* → met_*`. Pure rename, no behavior change (§3 above is updated to the new names).

4.5 Validation on WASP-121b

Three converged runs, identical planet / SED / grid (power-law SED, He/H=0.0851, Roche domain, “Rate/2 + Mdot/2”); they differ only in code version and metal set.

key	code	metals	C/N/O σ	floor [eV]	$\log_{10} \dot{M}$	T_{peak} [K]	steps
ref	old	C,N,O,Mg	Verner96	7.646	12.76	9828	5003
A	new	C,N,O,Mg	VFKY96	7.646	12.85	12197	15382
B	new	+Si,Ca,Na,K,S	VFKY96	4.341	12.85	12135	15505

$\dot{M}$ in $\log_{10}$ g/s; Huang+2023 Case A reference $\log_{10} \dot{M} = 12.57$. ref and A share the 7.646 eV grid floor (Mg sets the minimum in the C/N/O+Mg set), so ref vs A isolates the C/N/O cross-section change with no grid-floor confound; A vs B changes both the element set and the floor (K I pulls it to 4.341 eV).

Observations (for evaluation, not pass/fail):

- **C/N/O correction (ref→A):** $\Delta \log_{10} \dot{M} = +0.090$ dex; T_{peak} rises $9828 \rightarrow 12197$ K and the peak moves outward from 1.22 to $\sim 1.90 R_p$. The corrected VFKY96 opacity deposits more photoheating; the hotter peak now sits near Huang’s ~ 12000 K. A genuine model change, not numerical.
- **Adding Si/Ca/Na/K/S (A→B):** $\Delta \log_{10} \dot{M} = +0.000$ dex, $\Delta T_{\text{peak}} = -62$ K, despite the floor dropping to 4.341 eV — at solar abundance and with no line cooling these five are radiatively near-passive donors.
- **Mg moves indirectly.** Mg’s own cross sections are byte-identical between ref and A, yet its ionization fractions shift by up to ~ 0.1 : an indirect response to the hotter, more extended profile (changed n_e , optical depth, photoionizing flux), not a change in Mg atomic data. Mg is coupled to the global state, hence *not* an isolated control.

- **Numerics.** Zero NaN/Inf; the 21-unknown system converges over the full T range; X/H for each element is conserved to $\sim 10^{-13}$ for all nine elements including the two-stage Na/K/S — direct evidence the variable-stage packing does not leak between element columns.
- **New-metal ionization (run B).** Si, Ca mostly *doubly* ionized in the upper thermosphere; Na, K, S mostly *singly* ionized (no second stage carried) — the qualitative Huang+2023 Fig. 12 ordering.

4.6 Caveats and gaps

- **No metal line cooling** for Si/Ca/Na/K/S (Phase 2); the thermal structure will not match Huang until Fe II/Mg II/Ca II cooling is in.
- **No charge exchange with H** for the new elements or Mg (Phase 1d).
- **Recombination is Badnell, not the §1.3 validation rate** (deliberate, as Mg).
- $\dot{M}$ is $\sim 1.9 \times$ Huang Case A (12.85 vs 12.57); the metal-set (Fe still missing) and geometry confounds from Phase 0 still apply.
- **Not done here:** Huang Fig. 12 is not digitized/overlaid (only the qualitative ordering is checked); the atomic-data values are not re-derived beyond conservation and smoothness.

4.7 Regression reference for Phase 1c

The accepted run A output is frozen at `phase1b_validation/.gate_ref_cno_mg_vfky96/output/` as the post-Phase-1b C/N/O+Mg baseline. When Phase 1c adds Fe, running with `Fe off` must reproduce it; because adding Fe bumps `n_species/n_mion` (and the `Ion_species.txt` column count), the Phase-1c check is a **round-off match on the shared physical columns**, not a byte-for-byte compare (the `mg_validation/cno_mg` byte-gate applied only to the physics-neutral Phase 1a refactor, where `n_species` was unchanged).

4.8 Status and relation to the Huang plan

This completes **Phase 1b**: the network carries C/N/O/Mg/Si/Ca/Na/K/S, the C/N/O cross sections are corrected to VFky96, and the variable-stage solver handles 1st- and 2nd-ionization elements uniformly. The gate’s “converges/conserves/Fig. 12 ordering/ $\dot{M}$ within $\sim 2 \times$ Case A” items are addressed by the numbers above; the thermal-balance part is deferred to Phase 2 (metal line cooling). Next: **Phase 1c** — iron (Fe I/II/III), which needs non-uniform cross-section (Zatsarinny+2019) and recombination data. Scratch runs and notebook: `phase1b_validation/Phase1b_validation.ipynb`.

5. Phase 1c: add iron (Fe I / II / III)

5.1 Overview

Iron is added as the third three-stage trace metal, completing the Huang+2023 §2.4 element set (C, N, O, Mg, Si, Ca, Na, K, S, **Fe**). No solver, cooling-loop, heating, output, or post-processing code changed: the variable-stage MINPACK packing and every metal loop are generic over `n_melem`, so adding iron is purely new metadata rows plus four new rate functions. Counts: `n_species` = 33, `n_mion` = 27, `n_melem` = 10, `n_mphot` = 17; the system grows to `N_eq` = 3 + 2*`n_melem` = 23. Iron carries `melem_top` = 2 (three stages, like C/O/N/Mg/Si/Ca).

5.2 Atomic data and rates (sources, not fabricated)

Iron is the one element whose data are *not* the uniform-template choice. Each quantity was checked against its primary source before encoding; the Voronov fit in particular was cross-checked because an earlier abandoned attempt had its P/X/K columns transposed.

- **Photoionization cross sections** (`functions/cross_sec.f90`): single-shell VFKY96 form (`sigma_VFKY96`) for Fe I and Fe II, parameters from the Verner+1996 `phfit2.f` PH2 block. Neutral thresholds [eV]: Fe I **7.902** (NIST ionization potential), Fe II **16.199**. *To Be Checked / known limitation*: Huang+2023 use the multi-shell Zatsarinny+2019 Fe I cross section; the single-shell fit here under-represents inner-shell absorption at high XUV energy, so the iron photoionization rate is a lower bound at the hardest photons (flagged inline). A multi-shell upgrade is left for later.
- **Collisional ionization**: Voronov (1997) `cfit.dat`. Fe I (`dE`=7.9, `A`=2.52e-7, `P`=0, `X`=0.701, `K`=0.25), Fe II (`dE`=16.2, `A`=2.21e-8, `P`=1, `X`=0.033, `K`=0.45).
- **Recombination**: the Huang+2023 Eqs. (5)–(6) total fits (a dielectronic Arrhenius term $A T^{-1.5} e^{-T_0/T} (1 + B e^{-T_1/T})$ plus a radiative power law $C (T/10^4)^{-\eta}$) used *directly*, not the Badnell dispatcher used for the lighter metals. Deliberate: the Huang iron rates are the validation target, so the model runs on exactly those rates; the Badnell-swap for iron is deferred to Phase 6. Fe II+*e* → Fe I: `A`=2.833e-8, `T0`=5.731e4, `B`=1.383e4, `T1`=120.4; `C`=1.248e-12, `eta`=0.485. Fe III+*e* → Fe II: `A`=1.094e-5, `T0`=1.490e4, `B`=36.74, `T1`=1.153e5; `C`=1.728e-12, `eta`=0.618.
- **Line cooling**: *not* ported (`mion_iscool` = `.false.` for all Fe stages) — Phase 2. This matters more for iron than for the lighter metals: Fe II is the *dominant* low-altitude line coolant in Huang+2023, so a metals-on thermal profile is *expected to run hotter than Huang* until Phase 2.
- **Charge exchange with H**: *not* added; still C/N/O only (Phase 1d). Iron's $\text{Fe}^+ + \text{H}^+ \leftrightarrow \text{Fe}^{2+} + \text{H}$ reactions (Huang Table 4) come in Phase 1d.

5.3 Touched files

Metadata rows + four rate functions + abundance/IC wiring; no solver, cooling, heating, output, or post-process changes:

- `init/species_table.f90` — `n_mion=27`, `n_melem=10`, `n_mphot=17`; Fe I/II/III rows (`f_sp` cols 31–33); `iel_Fe=10`; `melem_top` appends 2.
- `functions/cross_sec.f90` — `sigma_FeI`, `sigma_FeII`.
- `init/set_energy_vectors.f90` — `sigma_tab(:,16)=sigma_FeI`, `(:,17)=sigma_FeII`.
- `radiation/Cool_coeff.f90` — Huang recombination `alpha_rec_FeI/FeII_Huang` with `rec_FeII/rec_FeIII` wrappers; Voronov `ion_coeff_FeI/FeII`; dispatcher cases for FeII/FeIII recombination and FeI/FeII ionization. No `cool_coeff_metal` case for Fe (Phase 2).
- `init/parameters.f90`, `files_IO/input_read.f90`, `files_IO/metals_input_read.f90` — `n_species=33`; `X_Fe`; `metals.inp` keys FeI/Fe; `melem_ab(iel_Fe)`; `thereis_metals` trigger.
- `init/set_IC.f90`, `files_IO/load_IC.f90` — neutral-iron IC for `f_sp` cols 31–33.

5.4 Iron abundance

Asplund et al. (2009) solar Fe/H = 3.16×10^{-5} ($\log \epsilon_{\text{Fe}} = 7.50$). Huang+2023 adopt $4 \times$ solar metallicity for WASP-121b, so the validation run uses $X_{\text{Fe}} = 1.26 \times 10^{-4}$ by number relative to H; C/N/O/Mg are held at solar to isolate the iron change.

5.5 Validation on WASP-121b

Two converged runs (identical planet / SED / grid; power-law SED, He/H=0.0851, Roche domain, “Rate/2 + Mdot/2”) plus the frozen Phase-1b C/N/O+Mg baseline:

key	metals	iron	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$	steps
ref	C,N,O,Mg	frozen 1b run A	12.85	12197	1.899	15382
off	C,N,O,Mg	off ($X_{\text{Fe}}=0$)	12.85	12197	1.899	15382
on	+Fe	on, $4 \times$ solar	12.87	12223	1.915	16566

$\dot{M}$ in $\log_{10}$ g/s; Huang+2023 Case A reference $\log_{10} \dot{M} = 12.57$.

Observations (for evaluation, not pass/fail):

- **Regression, force-zeroed iron (off vs ref).** The Phase-1c binary carries three extra species and two extra MINPACK unknowns, yet run off reproduces the frozen baseline on every shared column to the round-off floor: worst relative difference 2.0×10^{-13} across the seven hydro columns and 3.7×10^{-13} (NI) across the 30 shared species. The three new Fe columns are *identically zero, no NaNs*. With $X_{\text{Fe}}=0$ the iron block reduces to identity rows in the residual (`fvec=x`, rooted at zero) and contributes exactly zero electrons, so the 23-unknown system returns the 21-unknown answer — adding iron is non-perturbative when off.

- **Iron ionization structure (on).** Fe II dominates the base (fraction 0.89 at $r = 1$, peaking ~ 0.96 near $r = 1.02$); Fe III dominates the escaping flow (0.97 at the outer edge), with the Fe II/Fe III crossover near $r \simeq 1.25$; Fe I is minor everywhere (max ~ 0.11 at the base). This is the qualitative Huang+2023 Fig. 12 arrangement — Fe I confined to the base, Fe II at the base, Fe III through the outflow.
- **Iron as an electron donor.** At $4\times$ solar, iron supplies up to $\sim 26\%$ of the free-electron density, so it feeds back on the whole coupled solution, not just its own columns.
- **Switching iron on (off $\rightarrow$ on).** $\Delta \log_{10} \dot{M} = +0.02$ dex, $\Delta T_{\text{peak}} = +26$ K ($12197 \rightarrow 12223$). Small and in the *heating* direction: with no Fe line cooling yet, iron acts only as an added electron donor / photoelectric-heating channel, so the thermosphere runs slightly hotter, not cooler.
- **Numerics.** Run on converged with zero NaNs; X/H for each element is conserved to 8.3×10^{-13} across all five active elements including iron.

5.6 Caveats and gaps

- **No iron line cooling** (`mion_iscool=.false.`) — Phase 2. Because Fe II is the dominant low-altitude coolant in Huang+2023, the run-on *temperature profile is not yet comparable to Huang* and is expected to sit above it. The Phase-1c physics target is the iron *ionization structure*, not $T(r)$.
- **No iron charge exchange with H** — Phase 1d.
- **Single-shell VFKY96 Fe I/II cross sections** under-represent inner-shell XUV vs Huang’s Zatsarinny+2019 Fe I (flagged in `cross_sec.f90`).
- **Recombination is the Huang fit**, not Badnell (deliberate; the Badnell-swap is Phase 6).
- $\dot{M}$ is still $\sim 2\times$ Huang Case A (12.87 vs 12.57); the missing metal cooling and the Phase-0 geometry confounds still apply.
- **Huang Fig. 12 is not digitized/overlaid** — only the qualitative Fe II-base / Fe III-outflow ordering is checked.

5.7 Regression reference for Phase 1d / Phase 2

The run-on output (C/N/O/Mg + $4\times$ -solar Fe) is frozen at `phase1c_validation/.gate_ref_cno_mg_fe/` (input, `metals.inp`, `run.log`, `output/`) as the post-Phase-1c iron-on baseline. Phase 1d (iron charge exchange) and Phase 2 (iron line cooling) both *change* the iron ionization and the thermal structure deliberately, so this is the “before” snapshot for those comparisons, not a bit-reproduction target.

5.8 Status and relation to the Huang plan

Phase 1c **implementation is complete** and the iron ionization structure is presented for validation against Huang+2023 Fig. 12 (numbers and figures in the notebook; no pass/fail asserted here — that judgement is the reader’s). The

ionization-network half of Phase 1 now carries the full Huang element set with verified iron atomic data; the thermal-balance half remains deferred (until Fe II cooling lands in Phase 2 the metals-on temperature profile will not match Huang). Next: **Phase 1d** — charge exchange with H for all metals (Huang Table 4), including $\text{Fe}^+ + \text{H}^+ \leftrightarrow \text{Fe}^{2+} + \text{H}$. Scratch runs and notebook: `phase1c_validation/Phase1c_validation.ipynb`.

6. Phase 1d: charge exchange with hydrogen (Huang Table 4)

6.1 Overview

Phase 1d adds the charge-exchange (charge-transfer, CT) reactions of Huang+2023 Table 4 as a single table-driven module. The default active set is the 23 **Group A** metal + H / H^+ reactions, which couple each metal’s ionization balance to the H^+ fraction. A runtime switch `cx_full 1` in `metals.inp` enables the rest of the table — **Group B** (He + H, 2 rows), **Group C** (metal + He, 6 rows), **Group D** (metal + metal, 32 rows). This *replaces* the previous hard-coded Kingdon & Ferland (1996) C/N/O + H charge-exchange terms in the solver with a generic implementation covering all ten metals plus He.

Re-sourcing C/N/O to Huang Table 4 has one deliberate consequence: Table 4 tabulates only the neutral $\leftrightarrow$ singly-ionized CT for C, N, O (rows A13–A18), so the switch *drops* the doubly-ionized $\text{C}^{2+}/\text{N}^{2+}/\text{O}^{2+} + \text{H}$ recombination terms the old KF96 implementation carried. This is an intentional behavior change to match the paper; it raises the C/N/O double-ionization fractions and reshapes the thermal profile (below). No solver structure, cooling loop, heating, output, or post-processing code changed beyond the CT wiring.

6.2 Atomic data and rates (sources, not fabricated)

All 63 rate rows are transcribed verbatim from `docs/charge_exchange_table4.md`, the image-verified transcription of Huang+2023 Table 4 (23 Group A + 2 Group B + 6 Group C + 32 Group D). The paper text says “65 reactions”; the 2-row discrepancy is immaterial and noted in the reference doc.

- **Functional forms** (radiation/charge_exchange.f90): KF96 bracket $aT_4^b(1 + ce^{dT_4})e^{-E_k/T_4}$ (kf); power law $a(T/300)^be^{-E_k/T_4}$ (gp); 4th-order $\ln T$ polynomial (p4) for the radiative-CT rows of N / S / Na / K (Lin/Zhao/-Dutta/Watanabe); and the Satta (2013) $f_{\text{Si}}(T)$ and $D_{\text{CSi}}(T)$ helpers for the Si + He and C + Si rows.
- **Spot checks at $T = 10^4$ K** (against the table): $\text{Fe} + \text{H}^+ = 4.000 \times 10^{-9}$ (RV72), $\text{Fe}^+ + \text{H} = 2.011 \times 10^{-12}$, $\text{Fe}^{2+} + \text{H} = 1.260 \times 10^{-9}$ (ND87), $\text{O} + \text{H}^+ = 2.091 \times 10^{-9}$ (near-resonant, S99) — all match.
- **Rate clamp** $[0, 10^{-6}] \text{ cm}^3 \text{ s}^{-1}$: the upper bound guards the $\ln T$ polynomials, which diverge as $T \rightarrow 1$ K (never reached); the lower bound zeroes a few KF96 bracket fits that dip slightly negative below their validity floor (e.g. $\text{Mg}^+ + \text{H}^+$ at $T \lesssim 1500$ K), where the true rate is negligible.
- **Reference-doc notes**: Ca has *no* direct H charge exchange (only Group D); K has only the forward $\text{K} + \text{H}^+$ row. Reverse rates for pairs lacking direct data are Huang’s detailed-balance estimates, carried verbatim.

6.3 Touched files

A new module plus minimal wiring; the solver’s old element-by-element C/N/O CT block is deleted, not extended:

- `radiation/charge_exchange.f90` (**new**) — the full Table-4 descriptor arrays, the `cx_rate` dispatcher, and the generic residual assembly. Entry points: `cx_init` (build the active list from `cx_full`), `cx_set_cell(T)` (rates in each cell), `cx_add_to_fvec` (add CT terms to the MINPACK residual via the `cx_fvidx` element/stage $\rightarrow$ row map), and the `cx_full` switch. The zero-abundance guard is automatic: each rate is $k n_{\text{donor}} n_{\text{acceptor}}$.
- `nonlinear_system_solver/System_HeH_metals.f90` — the hard-coded C/N/O + H charge-exchange terms are *removed*; a single call `cx_add_to_fvec(...)` now feeds CT contributions into the H, He, and every metal-stage residual row.
- `radiation/ionization_equilibrium.f90` — call `cx_set_cell(T_K(j))` before each cell’s `hybrd1` solve.
- `files_IO/metals_input_read.f90` — parses `cx_full <0|1>` from `metals.inp`.
- `files_IO/input_read.f90` — if `(thereis_metals)` call `cx_init` after the input read, before the ionization sweep.

6.4 Validation on WASP-121b

Two new converged runs (identical planet / SED / grid as Phase 1c) compared against the frozen Phase-1c iron-on output as the regression reference:

key	charge exchange	$\log_{10} \dot{M}$	T_{peak} [K]	$r(T_{\text{peak}})$	steps
ref	KF96 C/N/O + H (frozen 1c on)	12.87	12223	1.915	16566
cx	Huang Table 4 Group A (default)	12.89	12955	1.778	13942
full	Huang Table 4 full, <code>cx_full 1</code>	12.88	12951	1.765	11661

$\dot{M}$ in $\log_{10}$ g/s; Huang+2023 Case A reference $\log_{10} \dot{M} = 12.57$. All three runs converged with **zero NaNs**.

Observations (for evaluation, not pass/fail):

- **H⁺ sink magnitude.** In the ionized bulk ($x_{\text{HII}} > 0.5$) the net metal charge exchange is at most 3.7×10^{-3} of radiative recombination (median 1.0×10^{-3}) — negligible as an H⁺ sink, consistent with Huang §4.2 (“Fe–H CX lower than H photoionization and recombination”). At the near-neutral base ($x_{\text{HII}} \sim 7.8 \times 10^{-8}$) the *ratio* is large (~ 118) only because H⁺ is itself a trace species there.
- **Near-resonant balance.** $\text{O} + \text{H}^+ \leftrightarrow \text{O}^+ + \text{H}$ is near-resonant and near-equilibrium at the base (gross destroy 3.66×10^5 vs produce 3.69×10^5 , net $+3.0 \times 10^3$). The net base H⁺ sink is instead driven by the fast, one-sided $\text{Fe} + \text{H}^+ \rightarrow \text{Fe}^+ + \text{H}$ channel (net -2.9×10^4).
- **C/N/O double ionization rises** (the dropped 2+ CT). Column-max fractions ref $\rightarrow$ cx: O III 0.871 $\rightarrow$ 0.946, N III 0.897 $\rightarrow$ 0.944, C III 0.851 $\rightarrow$ 0.856.

- **Iron base shift.** Fe II base fraction $0.889 \rightarrow 0.896$; Fe III essentially vanishes at the base ($0.002 \rightarrow 0.000$): $\text{Fe} + \text{H}^+ \rightarrow \text{Fe}^+ + \text{H}$ pushes iron toward Fe^+ .
- **Thermal structure.** T_{peak} rises $12223 \rightarrow 12955$ K and moves inward ($1.915 \rightarrow 1.778 R_p$). This is the *combined* effect of (i) the C/N/O re-sourcing that drops the 2+ recombination terms and (ii) the new metal + H CT. Charge exchange adds *no* radiative-loss channel, so it cannot cool the gas; Fe II / metal line cooling remains deferred to Phase 2, and the metals-on profile therefore still runs *hotter* than Huang.
- **Group A vs full Table 4 (cx vs full).** Turning on Groups B/C/D changes the bulk temperature by at most 1.5% (worst relative difference), with $\max |\Delta T| = 556$ K at $r = 1.008$; at $r = 1.02$ the full run is 357 K cooler with more base H^+ and less He II. $\dot{M}$ moves -0.01 dex and the outflow is nearly unchanged.
- **Dynamics.** v_{out} and sonic radius: ref 18.83 km/s @ $1.884 R_p$; cx 20.28 km/s @ $1.826 R_p$; full 20.40 km/s @ $1.826 R_p$. The hotter, more inward thermal peak drives a slightly faster wind with the sonic point $\sim 0.06 R_p$ further in.

6.5 Caveats and gaps

- **No metal line cooling yet** (Phase 2). Charge exchange adds no loss term, so $T(r)$ is not comparable to Huang; the C/N/O re-sourcing makes the base hotter, not cooler.
- **C/N/O 2+ + H terms intentionally dropped** to match Table 4 — a real change from the Phase-1c (KF96) reference, and the dominant driver of the O III/N III rise rather than the new metal CT alone.
- $\dot{M}$ still $\sim 2 \times$ Huang Case A (12.88–12.89 vs 12.57); the missing metal cooling and the Phase-0 geometry confounds still apply.
- **Paper lists 65 reactions, 63 transcribed** — immaterial, documented in `charge_exchange_table4.md`.
- **No Huang figure digitized/overlaid** for the H^+ budget; only the relative-magnitude statement (net CT $\ll$ recombination in the bulk) is checked.

6.6 Regression reference for Phase 2

The default Group-A run is frozen at `phase1d_validation/.gate_ref_cno_mg_fe_cx/` (input, `metals.inp`, `run.log`, output/) as the post-Phase-1d baseline. Because the C/N/O re-sourcing deliberately shifted both the ionization and the thermal structure away from the Phase-1c frozen reference, *this* snapshot — not the Phase-1c one — is the “before” reference for Phase 2 (metal line cooling, including Fe II).

6.7 Status and relation to the Huang plan

Phase 1d **implementation is complete**: the 23 Group-A metal + H reactions are active by default and the full Table 4 sits behind `cx_full`. With this, the **ionization-network half of Phase 1 is physically complete** relative to Huang+2023 §2.4 + Table 4 — the full element set, verified atomic data, and the metal–H charge-exchange coupling are all in

place. The thermal-balance half remains deferred: until metal line cooling (Fe II in particular) lands in **Phase 2**, the metals-on temperature profile stays above Huang. Next: **Phase 2** — metal line cooling. Scratch runs and notebook: `phase1d_validation/Phase1d_validation.ipynb`.

7. Phase 2: metal line cooling — validation gate closed

Phase 2 brings the CHIANTI-based metal line cooling (Mg I/II, Ca II, Na I, Fe II, on top of the C/N/O coolants) into the coupled energy balance and **closes the Phase 2 validation gate** against Huang et al. (2023) Fig. 10. The atomic data for each coolant and the CHIANTI effective-cooling tables are documented in `Update_EXHALE_early_phase` (Parts II–III); this section records the integrated WASP-121b gate and the diagnostic used to close it.

7.1 Active line coolants

The line-cooling ions are flagged by `mion_iscool` in `species_table.f90` (canonical order): C I, C II, O I, O II, N I, N II, Mg I, Mg II, Ca II, Na I, Fe II. Each contributes $\beta_{\text{esc}} n_e n_{\text{ion}} \Lambda_{\text{ion}}(T)$ with the effective cooling coefficient Λ_{ion} interpolated from the hardcoded `cool_logL_*` tables in `Cool_coeff.f90` ($N_{\text{COOLT}} = 41$, $\log_{10} T = 3.0\text{--}5.0$, $\Delta \log_{10} T = 0.05$). In the optically thin limit $\beta_{\text{esc}} = 1$ (see `Update_EXHALE_early_phase`, Part II, “Adopted assumption: 100% escape”). N I/II carry the `iscool` flag but have no tabulated line cooling ($\Lambda = 0$), so they contribute exactly zero.

7.2 Channel-by-channel cooling diagnostic (this phase)

To close the gate quantitatively, `eval_cool` (`util_ion_eq.f90`) gained an optional `cool_chan` output — an **exact** decomposition of the total cooling rate into H/He recombination, collisional ionization, collisional excitation, bremsstrahlung, and one column per metal ion. It is read straight from the arrays `eval_cool` already builds, so the channel sum reproduces the cool total to machine precision (no offline re-derivation, no divergence risk). A companion writer `write_cool_breakdown_eq` (same module, called once from `EXHALE_main.f90` after the final equilibrium write) dumps output/`Cooling_breakdown.txt`: columns r/R_p , T , n_e , `cool_total`, the four H/He channels, then 27 metal-ion columns. The run prints `max |sum(channels)/cool - 1|` as an internal consistency check.

7.3 Validation on WASP-121b (the gate)

Converged equilibrium run (full metals, the curated `metals.inp`; same SED/grid as Phase 1d):

- Internal consistency: $\max |\sum \text{channels}/\text{cool} - 1| = 3.7 \times 10^{-16}$ (machine eps) — the breakdown is a bit-exact decomposition of the `Hydro_ioniz.txt` cool column.
- In the $1.15 < r/R_p < 1.4$ band the dominant radiative coolants are **Fe II** (~30%) and **Mg II** (~25%), jointly ~55%; metals are **57–85%** of all radiative cooling through the band. Fe II leads the inner edge (top coolant at $1.1\text{--}1.2 R_p$, 53% at $1.1 R_p$), Mg II takes over by $1.3\text{--}1.4 R_p$, and H recombination only overtakes above $\sim 1.5 R_p$ — the Fe II $\rightarrow$ Mg II handoff Huang report in their Fig. 10.
- $\dot{M}$ unchanged ($\log_{10} \dot{M} = 12.51$ g/s); zero NaNs.

This matches the Phase 2 gate (“a metal-driven cooling feature appears at $1.15\text{--}1.4 R_p$, Fig. 10”) and Huang_update_expected_results (Mg II + Fe II become the most important coolant in $1.15 < r/R_p < 1.4$).

7.4 Touched files

- `src/modules/radiation/util_ion_eq.f90` — optional `cool_chan` out-arg on `eval_cool` (exact channel-by-channel split); new `write_cool_breakdown_eq` writer; `mion_fsp` added to the `species_table` use-list.
- `EXHALE_main.f90` — use `utils_ion_eq`, only: `write_cool_breakdown_eq`; call it after the final equilibrium `write_output`.
- Cooling coefficients themselves live in `Cool_coeff.f90` (provenance in `Update_EXHALE_early_phase`, Parts II–III); unchanged here.

7.5 Caveats and deferred

- The ion-by-ion metal decomposition is exact in the default branch (`use_2lev_cool = .false.`, the WASP-121b setting). In the optional two-level fine-structure branch the terms for each ion are the resonance-line approximation and need not sum to the two-level `cool_M`.
- **Fe II** uses Boltzmann/coronal lower-level populations ($E_{\text{cut}} = 38459 \text{ cm}^{-1}$), validated vs Huang Fig. 6 coronal curve within $\sim 1.5\times$ — sufficient for the gate. The Hollenbach & McKee IR fine-structure + density-dependent two-level refinement (Huang “Sum” curve) is deferred.
- **Fe I** line cooling (NIST f -values + Van Regemorter, Huang Fig. 5) is not yet added; Fe I is a minor coolant in the gate band. (*Now added in §9; confirmed negligible — peak 3×10^{-5} of total cooling.*)
- The **CHIANTI Ly α swap** is intentionally skipped in Phase 2; the $0.35\times$ Ly α is a Phase-3 radiative-transfer suppression, not a cooling-rate change.
- Carrying metal cooling into the post-processed `_adv` profile is handled separately (runtime `pp_metals`, default frozen); see the post-process notes.

7.6 Status and relation to the Huang plan

Phase 2 is complete and the gate is closed (user-validated 2026-06-05). The metals-on thermal balance now reproduces the Huang Fig. 10 Fe II/Mg II cooling feature, closing the thermal-balance half of the network that Phase 1 left open. Next: **Phase 3** — in-code excited $\text{H}(n=2)$ + Ly α radiative transfer.

8. Phase 3a: in-code excited $\text{H}(n=2)$ + Balmer proton source and heating

Added 2026-06-05.

Phase 3a is the first half of Phase 3: an **in-code $\text{H}(n=2)$ population model** whose output feeds back into both the hydrogen ionization balance (an extra proton source) and the energy equation (Balmer photoelectric heating), fully

coupled. The second half (a Ly α radiative-transfer solve to supply a self-consistent mean intensity $\bar{J}_{\text{Ly}\alpha}(r)$) is deferred; this phase uses a **parameterized** $\bar{J}_{\text{Ly}\alpha}$ with a switch to read an external RT profile instead. Everything is **default-off**: with no stellar parameters in `input.inp` the build is byte-identical to the Phase-2 result (`use_excited_H = .false.`).

8.1 Physics and atomic data (sources, not fabricated)

- **H(2s)/H(2p) populations**: a 2×2 statistical-equilibrium system following Christie, Arras & Li (2013) — recombination cascade into $n=2$, collisional $2s \leftrightarrow 2p$ mixing, two-photon $2s \rightarrow 1s$ decay, Ly α $2p \rightarrow 1s$ with a $2p$ radiative-pumping term, and $2s/2p$ photoionization. Solved per cell (`n2_populations`).
- **Ly α mean intensity** (the $2p$ pump): parameterized as $\bar{J}_{\text{Ly}\alpha} \approx 0.1 F_{\text{LyC}} / \Delta\nu_D$ (Huang et al. 2017, Eq. 6) with a **top-down** $1/(1 + \tau_{\text{Ly}\alpha})$ **line-center attenuation** so the optically thick base ($\tau_{\text{Ly}\alpha} \sim 1.1 \times 10^8$ here) is not over-pumped. Oscillator strength $f_{\text{Ly}\alpha} = 0.4162$.
- **$n=2$ Balmer-continuum photoionization rate Γ_2** : from a **diluted stellar blackbody** above the $n=2$ threshold $E_2 = 3.40$ eV (`gamma_n2_balmer`), with $\sigma_{\text{LyC}} = 6.3 \times 10^{-18} \text{ cm}^2$ and threshold cross-section $s_2 = 1.4 \times 10^{-17} \text{ cm}^2$. Γ_2 carries the **same ξ dayside-dilution factor as ground-state EUV ionization** (`Rate/2` $\rightarrow 0.5$, `Rate/4` $\rightarrow 0.25$); for WASP-121b’s “`Rate/2 + Mdot/2`” ($\xi = 0.5$) this gives $\Gamma_2 \approx 613 \text{ s}^{-1}$, comparable to Huang’s $\approx 576 \text{ s}^{-1}$.
- **Heating**: Balmer photoelectric heating $H_{\text{pe}} = n_2 \cdot (\text{photoelectron energy} \times \Gamma_2)$ per cell (`heat_n2_balmer`); optional collisional de-excitation heating behind `incl_deexc_heat` (off by default — it overlaps the existing HI collisional-excitation cooling).

8.2 Coupling into the solver

The model runs **decoupled (lagged one outer step)** to keep it out of the cell-by-cell Newton solve: `excited_H_update` fills two global arrays (one entry per cell) from the previous converged state, and `ioniz_eq` injects them:

- `gph_balmer_HI [s-1]` is added to the HI photoionization rate `P_HI` — an extra **proton source** ($S_{\text{proton}} = \Gamma_2 n_2$).
- `heat_balmer [erg cm-3 s-1]` is added to the photoheating rate.

`excited_H_update` is called once per outer iteration in `EXHALE_main.f90` before `ioniz_eq`, and once more after convergence to refresh the diagnostics. The post-loop refresh deliberately does *not* re-solve `ioniz_eq` (the Balmer heating is $\sim 10^{-4}$ of the total, not worth desyncing T/f_{sp} from the converged hydro state).

8.3 Runtime interface

Append to `input.inp` (omit for the default Phase-2-identical behavior):

Stellar Teff [K]: 6459.0

Stellar radius [R_sun]: 1.458

These set `use_excited_H` (both must be > 0). Optional further tail lines: `Jlya RT file: <path>` switches $\bar{J}_{\text{Ly}\alpha}$ to **mode 1** (read an external RT profile and use it directly, instead of the mode-0 parameterized field); `Deexc heat: True`

enables collisional de-excitation heating. As with all trailing lines the scan is by keyword, so order and blank lines do not matter.

8.4 Diagnostics

write_excited_H dumps output/Excited_H.txt (14 columns) from the converged state: r/R_p , T , n_{HI} , n_e , $\bar{J}_{\text{Ly}\alpha}$, n_{2s} , n_{2p} , $S_{\text{proton}}(n=2)$, H_{pe} , H_{dx} , $\tau_{\text{Ly}\alpha}$, then the three **proton-budget** columns added in this phase — **col12 ground-state photoionization** ($P_{\text{HI}} n_{\text{HI}}$), **col13 collisional ionization** ($C n_e n_{\text{HI}}$), **col14 recombination** ($\alpha_B n_e n_{\text{HII}}$), all volumetric [$\text{cm}^{-3} \text{s}^{-1}$]. The three rate coefficients are captured inside ioniz_eq on the converged pass (globals gph_ground_HI, cion_HI, arec_HII), so the budget is self-consistent with the solver rather than re-derived offline.

8.5 Touched files

- init/parameters.f90 — Phase-3a globals (use_excited_H, T_star_eff, R_star, gamma2_bal, hpe2_bal, jlya_mode, jlya_rt_file); feedback arrays gph_balmer_HI, heat_balmer; diagnostics Jlya_arr, n2s_arr, n2p_arr, Sproton_arr, Hpe_arr, Hdx_arr; and the proton-budget capture arrays gph_ground_HI, cion_HI, arec_HII.
- radiation/excited_hydrogen.f90 (**new**) — excited_H_update, n2_populations, gamma_n2_balmer, heat_n2_balmer, load_jlya_rt, write_excited_H.
- radiation/ionization_equilibrium.f90 — inject gph_balmer_HI/heat_balmer under use_excited_H; capture the ground-state / collisional / recombination coefficients for the proton-budget dump.
- files_IO/input_read.f90 — scan the stellar / Jlya RT file / Deexc heat tail lines; set use_excited_H, jlya_mode.
- EXHALE_main.f90 — in-loop excited_H_update before ioniz_eq; post-loop refresh + write_excited_H.
- Makefile — excited_hydrogen.f90 in SRC.

8.6 Validation on WASP-121b (coupling ON vs OFF)

Two converged runs, same build, differing only in the use_excited_H toggle (stellar tail lines present/absent). Both reach steady state via $du < du_{\text{th}}$ with $\log_{10} \dot{M} = 12.51$ (unchanged from Phase 2 — the coupling does not destabilize the energy solver). Diagnostics in WASP-121b/{analyze_excited_H.py, analyze_proton_budget.py, compare_on_off.py}.

Proton budget (cf. Huang Figs. 11/27), volumetric rates [$\text{cm}^{-3} \text{s}^{-1}$]:

r/R_p	S_{n2}	S_{ground}	S_{recomb}	S_{n2}/S_{ground}	S_{n2}/S_{source}
1.10	5.1e4	6.6e6	6.6e6	0.008	0.008
1.30	1.2e4	4.3e5	4.4e5	0.029	0.028
1.50	4.9e3	7.6e4	8.1e4	0.064	0.060
1.80	1.9e3	1.5e4	1.7e4	0.125	0.111
2.00	2.0e3	5.9e3	7.9e3	0.34	0.25
2.01	2.9e3	5.1e3	7.9e3	0.57	0.36

Observations (for evaluation, not pass/fail):

- **Self-consistency:** at the base $S_{\text{ground}} \approx S_{\text{recomb}}$ (6.6×10^6 vs 6.6×10^6 at $1.1 R_p$), i.e. photoionization equilibrium holds; collisional ionization is negligible (T too low).
- **n=2 share:** the n=2 Balmer channel is a minor proton source below $\sim 1.5 R_p$ ($\sim 1\text{--}6\%$ of ground-state photoionization) but climbs to **$\sim 34\text{--}57\%$ of the ground-state rate ($\sim 25\text{--}36\%$ of the total proton source) near $2 R_p$** , where the gas is already ionized and EUV is attenuated (ground-state photoionization depleted) while the thin-Ly α -fed n=2 channel is not. This is the “significant proton source approaching $2 R_p$ ” behavior Huang describe, here concentrated near $2 R_p$ rather than flat across $1\text{--}2 R_p$.

Heating budget (cf. Huang Figs. 10/26): Balmer photoelectric heating is present in the budget, peaking at **$\sim 1.8\%$ of total heating** at the Ly α photosphere ($\tau_{\text{Ly}\alpha} \sim 1$ near $r \sim 2 R_p$) and $\sim 10^{-4}$ globally — a small but nonzero entry, dominating nowhere, as expected.

ON vs OFF temperature and ionization: the coupling makes the gas slightly **more ionized** (up to $\sim 24\%$ less neutral H near $2 R_p$, Δx_{HII} up to $\sim 1.1 \times 10^{-2}$) and slightly **cooler** (up to ~ 179 K at $r \sim 2 R_p$, mean ~ 42 K). The cooling is **physical, not numerical:** the extra n=2 photoionization converts $\text{HI} \rightarrow \text{HII}$, lowering n_{HI} ; the dominant EUV photoelectric heating scales with n_{HI} , so $\sim 24\%$ less n_{HI} at $r \sim 2 R_p$ means correspondingly less local EUV heating there, which the small Balmer H_{pe} does not offset.

8.7 Caveats and gaps

- **Parameterized $\bar{J}_{\text{Ly}\alpha}$ (mode 0).** The $0.1 F_{\text{LyC}}/\Delta\nu_D$ field with $1/(1 + \tau_{\text{Ly}\alpha})$ attenuation is a placeholder for the Phase-3b Ly α RT solve. Because our n=2 contribution is peaked near $2 R_p$ rather than flat across $1\text{--}2 R_p$, if Huang’s Figs. 11/27 are flatter/larger across the band the parameterized field is under-pumping n=2 in $1.1\text{--}1.8 R_p$; mode 1 (external RT $\bar{J}_{\text{Ly}\alpha}$) is the path to test this once a Ly α RT result exists.
- **Convergence stop is momentum-uniformity** ($du < du_{\text{th}}$), not a thermal/ionization steady-state criterion, so a small residual convergence-path sensitivity rides on top of the physical ON–OFF difference.
- **De-excitation heating off by default** (overlaps the existing HI coex cooling); left behind `incl_deexc_heat`.

8.8 Two analysis gotchas (recorded so they are not repeated)

- Hydro_ioniz.txt columns 6/7 (heat/cool) are **already physical** (write_output multiplies by q_0); do *not* multiply by q_0 again in analysis.
- Ion_species.txt column 1 is r/R_p , so $n_{\text{HI}}=\text{col2}$, $n_{\text{HII}}=\text{col3}$ — the recombination sink uses col3, not col2.

8.9 Status and relation to the Huang plan

Phase 3a is complete and the proton-source comparison is user-accepted (2026-06-05). The H(n=2) model is fully coupled into the ionization balance and the energy equation; the n=2 photoionization proton source reproduces the qualitative Huang Figs. 11/27 behavior and the Balmer photoelectric heating appears in the budget without destabilizing the solver. Next: **Phase 3b** — a Ly α radiative-transfer solve to replace the parameterized $\bar{J}_{\text{Ly}\alpha}$ with a self-consistent $\bar{J}_{\text{Ly}\alpha}(r)$ (feeding mode 1).

9. Phase 2 refinement: Fe I line cooling (NIST f -values + Van Regemorter)

Added 2026-06-05.

Fe I line cooling was one of the two coolants left deferred when the Phase 2 gate closed (§7, “Caveats and deferred”). This update adds it. CHIANTI v11 has **no Fe I model atom**, so — following Huang et al. (2023) §2.5 — Fe I cooling is built from NIST permitted-line oscillator strengths plus the Van Regemorter collision-strength approximation, with the lower levels populated in Boltzmann equilibrium over the metastable manifold. The new coolant flows through the same $\beta_{\text{esc}} n_e n_{\text{FeI}} \Lambda_{\text{FeI}}(T)$ channel as the other Phase-2 metals.

9.1 Atomic data and method (sources, not fabricated)

The cooling coefficient per $(n_e n_{\text{FeI, total}})$ in the optically thin limit [$\text{erg cm}^3 \text{ s}^{-1}$] is

$$\Lambda(T) = \sum_{\text{lines}} f_l(T) q_{lu}(T) \Delta E_{lu},$$

$$f_l(T) = \frac{g_l e^{-E_l/kT}}{U(T)}, \quad U(T) = \sum_{\substack{\text{even levels} \\ E < E_{\text{cut}}}} g e^{-E/kT}, \quad q_{lu} = 2.16 a^{-1.68} e^{-a} T^{-3/2} f_{lu},$$

with $a = \Delta E/kT$ and the absorption oscillator strength f_{lu} from the NIST A_{ki} , statistical weights, and wavelength (Huang Eq. 10).

- **Raw data:** the auditable NIST dumps fe1_nist_lines.tsv (permitted E1 lines: A_{ki} , g_i , g_k , E_i , E_k) and fe1_nist_levels.tsv (every level g , E for the partition function), fetched by fetch_fe1_nist.py.
- **Metastable manifold = EVEN-parity levels below the lowest odd level** ($z^7\text{D}^\circ$, $19351 \text{ cm}^{-1} = 2.40 \text{ eV}$). Below that energy there is no lower level for an E1 decay, so the even levels are genuinely metastable and hold the Fe I population. This is *not* a plain energy cut: the Fe II treatment can use an energy cut because its cut conveniently lands at the first odd level, but for Fe I a pure energy cut would wrongly include the radiatively-depopulated $z^7\text{D}^\circ$ odd levels. The default cut yields **17 even levels and 828 permitted cooling lines**.

- Upper levels of the cooling lines are odd (E1) and assumed to decay radiatively. The builder is `cooling_data/fe1_cooling.py`; it emits the 41-point `cool_logL_FeI(NC00LT)` Fortran block on the ATES grid ($\log_{10} T = 3.0\text{--}5.0$, $\Delta \log T = 0.05$).

Coefficient values (per $n_e n_{\text{FeI}}$, $\text{erg cm}^3 \text{ s}^{-1}$): $\Lambda(5 \times 10^3 \text{ K}) = 1.27 \times 10^{-22}$, $\Lambda(10^4 \text{ K}) = 1.23 \times 10^{-20}$, $\Lambda(2 \times 10^4 \text{ K}) = 1.31 \times 10^{-19}$ — the rising, then flattening shape of Huang Fig. 5.

9.2 Touched files

- `cooling_data/fetch_fe1_nist.py`, `cooling_data/fe1_cooling.py` (**new**) — raw NIST fetch and the Boltzmann-manifold + Van Regemorter builder.
- `radiation/Cool_coeff.f90` — `cool_logL_FeI(NC00LT)` table constant; new `cool_FeI(T, out)` (interpolates the table via `interp_cool_table`); dispatch cases FeI in `cool_coeff_metal` and `cool_coeff_metal_scalar`. The table-header comment block now documents Fe I alongside the other coolants.
- `init/species_table.f90` — `mion_iscool(FeI) = .true.` (canonical ion index 25), so the `eval_cool` metal-cooling loop and the channel-by-channel breakdown both pick up Fe I automatically.

9.3 Validation on WASP-121b

Re-ran the converged full-metals WASP-121b equilibrium (same SED/grid as the Phase 2 gate) with Fe I cooling on; the channel-by-channel `Cooling_breakdown.txt` then gives the Fe I share directly. Diagnostics: WASP-121b/`analyze_FeI_cooling.py` (left panel = the ported $\Lambda_{\text{FeI}}(T)$ coefficient read straight from `Cool_coeff.f90` vs the Fe II coronal curve, Huang Fig. 5 style; right panel = cooling in each channel vs radius, Huang Fig. 10 style), figure WASP-121b/`FeI_cooling_validation.png`. Observations (for evaluation, not pass/fail):

- **Fe I is a negligible coolant** along the whole profile: its peak share of the total cooling is 3.0×10^{-5} **at** $r = 1.06 R_p$ ($T = 8303 \text{ K}$), and it is far smaller elsewhere ($\sim 3 \times 10^{-8}$ of total at the base, where Fe II dominates). Adding Fe I does *not* perturb the gate-band Fe II $\rightarrow$ Mg II handoff.
- **Mdot unchanged** at $\log_{10} \dot{M} = 12.51$ (Fe I adds no measurable cooling), zero NaNs.
- This is the expected outcome: at WASP-121b base temperatures iron is almost entirely Fe II (Fe I fraction < 0.11 , §5), and Fe I's permitted lines have high excitation energies, so its Boltzmann-suppressed cooling cannot compete with the Fe II / Mg II forbidden + resonance lines. Fe I is carried for completeness and to mirror Huang's coolant list, not because it shapes the WASP-121b energy budget.

9.4 Status and relation to the Huang plan

Fe I line cooling is **implemented and its coefficient matches the Huang Fig. 5 shape/magnitude**; the WASP-121b run confirms it is radiatively negligible there, surfaced as a figure for the reader's judgment (no pass/fail asserted). One of the two Phase-2 deferred cooling items is now closed. The remaining one — the Fe II density-dependent (IR fine-structure) low-temperature refinement — is §10.

10. Phase 2 refinement: density-dependent Fe II line cooling (SE table)

Added 2026-06-05. (Implemented step-by-step; this section is updated as each step lands — currently: the 2-D table generator is built and verified.)

This is the second deferred Phase-2 cooling item (§7): replacing the **coronal** Fe II line-cooling coefficient with a **density-dependent** one that follows the Fe II level populations from the coronal (low- n_e) limit into the collisionally saturated (high- n_e , LTE) limit. It directly affects the cool, dense base of the WASP-121b thermosphere, where the coronal table overestimates Fe II cooling by $\sim 10^4 \times$.

10.1 Why the coronal table is wrong at the base

The Phase-2 Fe II table is the optically-thin **coronal** limit: only the ground level is populated and every collisional excitation is assumed to be followed by a radiative decay, so the cooling is strictly $\propto n_e$ (per $n_e n_{\text{FeII}}$ it is a function of T alone). That assumption fails at the cool, **dense** base of a hot Jupiter thermosphere, because the low- T Fe II cooling there is carried almost entirely by **forbidden** metastable / a⁶D fine-structure lines ($A \lesssim 5 \times 10^{-3} \text{ s}^{-1}$, far-IR 26–87 μm), whose critical densities

$$n_{\text{crit}} = A_{ul} / \sum_c n_c k_{ul}^c \sim 10^4\text{--}10^7 \text{ cm}^{-3} \quad (\text{electron collider})$$

sit far *below* the base $n_e \sim 10^8 \text{ cm}^{-3}$. Those lines are therefore collisionally **saturated** (thermalized to their LTE populations); the coronal coefficient overestimates their cooling by $\sim n_e / n_{\text{crit}}$.

10.2 The density-dependent coefficient (multilevel statistical equilibrium)

cooling_data/fe2_cooling.py solves the full CHIANTI multilevel statistical-equilibrium (SE) problem for the Fe II level populations at each (T, n_e) — 175 levels, electron-collisional excitation/de-excitation (Burgess–Tully descaled Υ from the .scups file) plus radiative decay (.wgfa A -values), normalized to $\sum_i n_i = 1$ — so the cooling smoothly interpolates between the two limits:

$$\frac{\Lambda_{\text{vol}}}{n_{\text{FeII}}} = \sum_{u>l} n_u A_{ul} \Delta E_{ul} \quad [\text{erg s}^{-1} \text{ per ion}], \quad \Lambda_{\text{eff}}(T, n_e) = \frac{\Lambda_{\text{vol}} / n_{\text{FeII}}}{n_e} \quad [\text{erg cm}^3 \text{ s}^{-1}].$$

At low n_e , $\Lambda_{\text{eff}} \rightarrow \Lambda_{\text{coronal}}(T)$; at high n_e it falls off $\sim 1/n_e$ as the numerator saturates. This Λ_{eff} is exactly the per- $(n_e n_{\text{FeII}})$ coefficient EXHALE already tabulates, so the existing assembly $\text{cool_M} = \beta_{\text{esc}} n_e \sum_i n_{\text{ion}}(i) \Lambda(i)$ carries it unchanged — the n_e in the prefactor times Λ_{eff} gives $\beta_{\text{esc}} n_{\text{FeII}} \times (\text{cooling per ion})$, i.e. the correct saturated base cooling.

Why electron-only SE is sufficient at the (mostly neutral) base. In the saturated limit the upper levels reach their **Boltzmann (LTE) populations regardless of which collider thermalizes them**, and the cooling of each ion becomes collider- and n_e -independent. At the base $n_e \sim 10^8$ already exceeds n_{crit} for the dominant forbidden lines, so the electron-only SE already reaches that LTE plateau; adding the (more abundant) neutral-H collider would only push the few still-unsaturated higher- n_{crit} lines further toward the same LTE limit — a second-order correction. An explicit H I / proton collider channel is flagged as a future refinement in the module.

10.3 Step 1 — the 2-D table generator (done)

fe2_cooling.py gained an **optimized** table builder `lambda_eff_table(Tgrid, ngrid)` and Fortran emitters (`fortran_grid_1d`, `fortran_block_2d`, `emit_fortran_table`; run `python3 fe2_cooling.py --table`). The optimization is required because the naive cell-by-cell solve rebuilds a cubic spline for every one of the ~ 4300 transitions at every grid cell: instead the descaled Υ is evaluated **once per transition over the whole T-grid** (`upsilon()` accepts an array T), and the rate matrix is split $M = n_e Q_{\text{coll}}(T) + A_{\text{rad}}$, so the n_e -loop only reassembles and solves. The full 41×29 table builds in ~ 3 s and matches the naive cell-by-cell solve to 1.5×10^{-15} (machine precision).

The emitted table is on the ATES temperature grid ($\text{NCOOLT} = 41$, $\log_{10} T = 3.0\text{--}5.0$, $\Delta \log T = 0.05$) $\times$ an electron-density axis `cool_logne(NCOOLNE = 29)`, $\log_{10} n_e = 0\text{--}14$, $\Delta \log = 0.5$ (brackets the WASP-121b run's $\log_{10} n_e \approx 7.9\text{--}9.7$ with the coronal plateau below and LTE saturation above). It is stored as `cool_logL_FeII_ne(NCOOLT, NCOOLNE) = log10  $\Lambda_{\text{eff}}$` .

Two built-in sanity checks (printed to stderr by `--table`):

- **Low- n_e edge reproduces the coronal table.** At $n_e = 1 \text{ cm}^{-3}$ the SE coefficient equals `cooling_effective(pop='coronal')` to **0.1%** across $T = 10^3\text{--}10^5 \text{ K}$ — so swapping the 1-D coronal table for the $n_e = 1$ column of the 2-D table is a no-op, and the Phase-2 gate's coronal Fe II validation (vs Huang Fig. 6, within $\sim 1.5\times$) carries over unchanged.
- **Base suppression.** At the WASP-121b base ($T \approx 2239 \text{ K}$, $n_e = 10^8$) the coronal coefficient 2.80×10^{-20} drops to the SE value 1.79×10^{-24} — a **15,663 $\times$ suppression**. The suppression is strongest at low T (forbidden lines) and weak at high T ($T \gtrsim 3 \times 10^4 \text{ K}$, where the resonance lines with large A dominate and stay near-coronal): e.g. only $\sim 1.4\times$ at $T = 3 \times 10^4$, $n_e = 10^8$.

10.4 Step 2 — the 2-D table and bilinear interpolator in `Cool_coeff.f90` (done)

The generated table and its interpolators are now in `src/modules/radiation/Cool_coeff.f90`:

- **Table.** `integer, parameter :: NCOOLNE = 29`, the `cool_logne(NCOOLNE)` axis ($\log_{10} n_e = 0\text{--}14$, $d\log = 0.5$), and the data block `cool_logL_FeII_ne(NCOOLT, NCOOLNE)` (a `reshape(..., [NCOOLT, NCOOLNE])` with T on the inner index and n_e on the outer), pasted straight from `fe2_cooling.py --table`. The $n_e = 1$ column matches the existing 1-D `cool_logL_FeII` to the fifth decimal (first node -19.78002 vs -19.77975).
- **Bilinear interpolator.** `interp_cool_table_2d(logL2d, T, ne, out)` — log-log bilinear in $(\log_{10} T, \log_{10} n_e)$. The T axis reuses the uniform `cool_dlogT` scheme; the n_e axis uses the uniform `dlogne = 0.5`. Both axes clamp to the nearest edge outside the table (so n_e below the grid gives the coronal plateau, above it the LTE branch). A scalar twin `interp_cool_table_2d_scalar(logL2d, Ts, nes)` mirrors it bit-for-bit for the cell-by-cell post-process solve.
- **Coolant wrapper.** `cool_FeII_ne(T, ne, out)` calls the 2-D interpolator on `cool_logL_FeII_ne`. The 1-D `cool_FeII` is kept as the low- n_e edge / scalar fallback; its header `!To Be Checked` note was rewritten to record that the density-dependent table now supersedes the coronal-only rate in the cooling assembly.

Builds clean (gfortran, full make); the new routines are additive and not yet wired, so behavior is unchanged until Step 3.

10.5 Step 3 — wiring the density-dependent Fe II into the cooling assembly (done)

The 2-D coefficient is now wired into both cooling paths so the local n_e selects the coronal $\rightarrow$ LTE-saturated branch automatically:

- **Main solver (util_ion_eq.f90, eval_cool).** n_e is already computed at the top of eval_cool (call calc_ne). Immediately after the canonical c_metal(:,i) fill loop, a short loop locates Fe II by name and overrides its column with call cool_FeII_ne(T_K, ne, ...). Nothing else changes: the existing assembly cool_M = $\beta_{\text{esc}} n_e \sum_i n_{m,i} c_{\text{metal},i}$ then multiplies by n_e , which *cancels* the $1/n_e$ in $\Lambda_{\text{eff}} = (\sum_u n_u A_{ul} \Delta E_{ul})/n_e$, leaving the correct cooling for each Fe II ion. The saturated (LTE) populations are collider-independent, so the electron-only SE solve is exact in that limit; at low n_e , Λ_{eff} reduces to the coronal rate, so optically-thin upper-atmosphere cells are unaffected. The override flows automatically into the use_2lev_cool branch (which already reads c_metal(:,26)) and into the channel-by-channel cool_chan diagnostic (which reads c_metal(:,i)), so the breakdown stays an exact decomposition of cool.
- **Post-process cell-by-cell solve (T_equation.f90).** This path computes its own local n_e (from H, He, and the frozen metal stages) and sums the metal coolants with the scalar dispatcher. The Fe II term is special-cased to call the new cool_FeII_ne_scalar(TT, ne) (a thin wrapper over interp_cool_table_2d_scalar), so the post-process temperature balances the same density-dependent Fe II cooling.
- **species_table.f90.** The !To Be Checked Fe II comment was rewritten to state that eval_cool now applies the density-dependent SE coefficient (no longer “deferred”).

Builds clean (gfortran, full make, all four touched modules relink).

10.6 Step 4 — WASP-121b run and the cooling budget (numbers for the reader)

A fresh WASP-121b run (solar metals, pp_metals 1, otherwise the same input.inp as the FeI-cooling run) was compared against the snapshot of the previous (coronal Fe II) run kept in output_coronalFeII/. Both share the same grid, so Cooling_breakdown.txt columns are differenced directly (analyze_FeII_density_correction.py $\rightarrow$ FeII_density_correction_compare.png).

At the dense base ($r = 1.0$, $T = 2358\text{K}$ fixed by the lower BC). The Fe II line cooling drops from the coronal 2.23×10^{-4} to the SE $1.71 \times 10^{-8} \text{ erg cm}^{-3} \text{ s}^{-1}$ — a $\approx 1.3 \times 10^4 \times$ suppression, consistent with the offline fixed- (T, n_e) check in Step 1 ($15,663 \times$ at $T = 2239\text{K}$, $n_e = 10^8$). Because coronal Fe II carried **99.1%** of the base cooling, the *total* radiative cooling at the base falls $\approx 122 \times$ ($2.25 \times 10^{-4} \rightarrow 1.84 \times 10^{-6}$); Fe II’s share there is now **0.9%**. The previous run’s post-process printed “35 of 502 cells fell back to eq T (stiff base band)”; with the spurious base overcooling removed that message is **gone** — the cell-by-cell T solve no longer stiffens at the base.

Transition region 1.15–1.4 R_p (Huang Fig. 10). The band-integrated metal-line cooling share in the new run is

coolant	share
Mg II	44.0%
O II	18.9%
Fe II	13.4%
O I	7.3%
C II	6.3%
Ca II	5.1%

so **Mg II is now unambiguously the dominant metal coolant** across the transition region, with Fe II demoted to a $\sim 13\%$ contributor — the qualitative ordering Huang’s Fig. 10 shows. (In the old coronal run Fe II still topped the budget out to $\sim 1.2 R_p$.)

Global response. Removing the base overcooling raises the heating efficiency, so the steady-state mass-loss rate rises from $\log_{10} \dot{M} = 12.51$ to **12.80** g/s ($\approx 2\times$). The self-consistent wind is correspondingly denser and, above the base, slightly *cooler* (e.g. $T(1.05)$: $7810 \rightarrow 5744$ K) — the freed energy goes into mechanical escape and adiabatic expansion rather than being radiated, the expected energy-limited behavior. Whether this $\dot{M}$ and the transition-region budget match Huang’s WASP-121b quantitatively is left to the reader’s judgment against the paper; the figure and table above are surfaced for that comparison, not asserted as a pass.

Caveat on the figure’s panel (b). The “suppression factor” there is the *ratio of Fe II cooling between the two runs*, not a fixed-condition coefficient ratio. At the base (BC-fixed T) it equals the coefficient suppression ($\sim 1.3 \times 10^4$); above $\sim 1.2 R_p$ the coronal and SE coefficients converge (the hot resonance lines dominate and stay near-coronal), so the residual ratio there reflects the changed — denser — wind profile, not the coefficient.

11. Post-process refinement: re-solve metal ionization in the advected wind (pp_metals 2)

Added 2026-06-05.

The `_adv` profiles that EXHALE_transit.py reads are built by `post_process_adv.f90`, which re-solves the *temperature in each cell* at the advection-corrected structure (the wind is advected, then its $T(r)$ is re-converged against heating/cooling). Until now the metal *ion* densities carried into those profiles were handled in one of two ways, selected at runtime by `pp_metals` in `metals.inp` ($\rightarrow$ `pp_metal_mode`):

- `pp_metals 0` — **metal-free legacy** `_adv` (the pre-metals behavior).
- `pp_metals 1` — **frozen equilibrium** metals (the default): the converged equilibrium metal stages are carried into `_adv` unchanged, so they track the *equilibrium* structure, not the (slightly hotter, more ionized) advected wind.

This step implements the deferred third option:

- `pp_metals 2` — **re-solve**: the metal ionization balance is recomputed in each cell at the advection-corrected H/He and the post-process temperature, so the metal stage split tracks the advected wind self-consistently.

11.1 Implementation — reuse the coupled residual with H/He pinned

The metals are pure equilibrium species (no advection equations of their own): in the main loop they are solved *coupled* with H/He through `ion_system_HeH_metals + hybrd1` (photo + collisional ionization + recombination + charge exchange, no time derivative). “Re-solve” therefore means re-evaluating that same equilibrium balance at the post-process local T and the advected H/He — *not* a new physics path.

Rather than duplicate the residual, mode 2 reuses it verbatim through a thin wrapper, `ion_system_metals_pp` (new, in `post_process_adv.f90`):

```
subroutine ion_system_metals_pp(N_in,x,fvec,iflag,params)
...
call ion_system_HeH_metals(N_in,x,fvec,iflag,params) ! full coupled residual
fvec(1) = x(1) - pp_xHII_fix      ! pin HII fraction
fvec(2) = x(2) - pp_xHeII_fix     ! pin HeII fraction
fvec(3) = x(3) - pp_xHeIII_fix    ! pin HeIII fraction
end subroutine
```

The wrapper overwrites only the three H/He rows with identity equations that pin H/He to their advection-corrected fractions (stored in module-level `pp_xHII_fix` / `pp_xHeII_fix` / `pp_xHeIII_fix`), while leaving the $2n_{\text{elem}} = 20$ metal-stage rows untouched. So `hybrd1` moves only the metal unknowns; H/He stay fixed at the advected wind. Two properties fall out for free:

- **Self-consistent n_e** . The reused residual computes the electron density *internally* from the pinned H/He plus the metals it is currently solving — exactly as the equilibrium solver does — so the re-solve sees the same n_e (metals) feedback, not a frozen n_e .
- **Charge exchange with the advected H**. The Kingdon & Ferland $\text{H} \leftrightarrow \text{metal}$ couplings inside `ion_system_HeH_metals` see the *pinned* H^+/He^+ densities, so the metal balance reacts to the advected hydrogen ionization, not the equilibrium one.

11.2 The cell-by-cell sweep in `post_process_adv`

Mode 2 adds one sweep over cells, placed after the post-advection `calc_ne` and before the photoheating refresh, inside the same k -loop that re-converges T , so metals, H/He and T are driven to a *joint* fixed point. Per cell j (skipping $n_h(j) \leq 0$):

1. `cx_set_cell(T_K(j))` — charge-exchange rate coefficients at the local T .

2. Build the coefficients for each element (canonical order) from the existing post-process rate arrays: $\text{meg_g0} = P_m$ (photoionization), $\text{meg_b0} = \text{aion_m_pp}$ (collisional ionization), $\text{meg_a1} = \text{rec_m_pp}$ (recombination), and the second-stage $g1/b1/a2$ where $\text{melem_top} \geq 2$; hand them to `set_metal_coeffs`.
3. Pin $\text{pp_xHII_fix} = n_{\text{HII}}/n_h$, $\text{pp_xHeII_fix} = n_{\text{HeII}}/n_{he}$, $\text{pp_xHeIII_fix} = n_{\text{HeIII}}/n_{he}$; set `params(7) =  $n_h$` , `params(8) =  $n_{he}$` .
4. Seed `sys_x` with the pinned H/He fractions plus the current metal split and call `hybrd1(ion_system_metals_pp, . . .)`.
5. **Conserve element totals:** the pre-solve totals `nm_tot_pp(:,im)` (summed over each element's stages once, before the k -loop) are multiplied by the re-solved fractions to write back `nm_w`. Only the *stage split* changes; the element abundance is untouched.

Supporting edits: the use block pulls in `ion_system_HeH_metals`, `set_metal_coeffs`, and `cx_set_cell`; `params` was widened from dimension(25) to dimension(60) to match the residual's dummy length (only entries 1–11 are consumed by the metal solve); module-level `pp_x*_fix` and locals `nm_tot_pp/meg_*/i0`, `top`, `im` were added. Builds clean (gfortran, full make).

11.3 Clean A/B and what changes

Because the **equilibrium** solve in the main RK loop is mode-independent and deterministic, a full `pp_metals 2` run reproduces `Hydro_ioniz.txt`, `Ion_species.txt`, and `Cooling_breakdown.txt` **byte-identically** to the `pp_metals 1` run (verified with `diff -q`); *only* the `*_adv.txt` profiles differ. So the A/B isolates exactly the post-process metal treatment. The mode-1 default is preserved in `output_ppm1_dd/`, the mode-2 run in `output_ppm2_resolve/`, and the canonical `output/` is left at the mode-1 default (metals.inp reset to `pp_metals 1`); `analyze_metal_resolve.py` differences the two and writes `metal_resolve_compare.png`.

Singly-ionized fraction X^+/X_{tot} in the `_adv` wind (eq frozen $\rightarrow$ re-solve). Only the elements that move appreciably are tabulated (Mg and Na are unchanged to $\leq 2\%$); the full output for each element is in the analysis script.

r/R_p	T_{adv} [K]	O	N	Ca	Fe
1.05	5798	.00 $\rightarrow$.00	.01 $\rightarrow$.02	.64 $\rightarrow$.59	1.00 $\rightarrow$ 1.00
1.10	7984	.02 $\rightarrow$.01	.12 $\rightarrow$.13	.65 $\rightarrow$.60	.99 $\rightarrow$.99
1.15	9044	.10 $\rightarrow$.06	.38 $\rightarrow$.40	.66 $\rightarrow$.59	.99 $\rightarrow$.99
1.20	10035	.37 $\rightarrow$.26	.51 $\rightarrow$.47	.61 $\rightarrow$.54	.93 $\rightarrow$.96
1.30	10274	.34 $\rightarrow$.34	.35 $\rightarrow$.34	.38 $\rightarrow$.38	.39 $\rightarrow$.40
1.40	10680	.23 $\rightarrow$.23	.23 $\rightarrow$.23	.25 $\rightarrow$.25	.18 $\rightarrow$.19

Largest density change. The element-by-element scan of $\max |\Delta n(X^+)/n(X^+)|$ over the wind 1.0–1.48 R_p :

element	max rel. change	at r/R_p	$n(X^+)$ frozen $\rightarrow$ resolve
Fe	+9.5%	1.232	$1.19 \times 10^5 \rightarrow 1.30 \times 10^5$
Ca	-11.9%	1.191	$1.24 \times 10^4 \rightarrow 1.09 \times 10^4$
Na	-6.2%	1.000	$3.50 \times 10^6 \rightarrow 3.28 \times 10^6$
C	-2.1%	1.216	$1.22 \times 10^6 \rightarrow 1.19 \times 10^6$
Mg	-2.0%	1.000	$9.82 \times 10^7 \rightarrow 9.62 \times 10^7$

The corrections are **modest and concentrated in the 1.1–1.25 R_p ionization transition**, where T_{adv} runs a few hundred K hotter than T_{eq} : **Fe II** is *higher* across the 1.18–1.27 R_p shoulder (blue above red in the figure), **Ca II** is *lower* by ~ 5 –7 points through 1.05–1.25 R_p , and **every element converges back to the equilibrium split** in the outer wind ($\gtrsim 1.27 R_p$) where T_{adv} rejoins T_{eq} . The post-process printed **no** stiff-base T -fallback message in mode 2 (the density-dependent Fe II cooling of the previous section already removed the base overcooling).

Tiny-denominator artifacts (not physical). The same scan reports **O** +6538% at $r = 1.000$ and **N** +17.9% at $r = 1.013$. These are at the dense base where O and N are essentially *fully neutral* (O II fraction ≈ 0.00 in the table), so $n(\text{O II})$ is a vanishingly small number; a tiny absolute shift ($n^+ \approx 9.7 \times 10^2 \rightarrow 6.4 \times 10^4 \text{ cm}^{-3}$, against an O total $\sim 10^8$) reads as a huge *relative* change but is negligible in absolute and in any observable. They are denominator artifacts, not a physical re-solve effect.

Reader’s judgment, not a pass. The re-solve makes the metal stage split in the `_adv` profiles consistent with the advected wind rather than the equilibrium structure, with the physically expected sign (more ionization where the advected wind is hotter) and magnitude (single-digit-percent, transition-localized). The default remains `pp_metals 1` (frozen): the deliberate, documented choice (the equilibrium split is the converged one, and EXHALE_transit.py’s lines are dominated by H and He). Mode 2 is provided for studies that need the metal ionization to track the advected wind. The table and `metal_resolve_compare.png` are surfaced for that comparison, not asserted as a validation pass.

12. Phase 3b: in-code $\text{Ly}\alpha$ radiative transfer ($\bar{J}_{\text{Ly}\alpha}$ via escape probability)

Phase 3a drove the $\text{H}(n=2)$ population with a *parameterized* $\text{Ly}\alpha$ mean intensity $\bar{J}_{\text{Ly}\alpha} = 0.1 F_{\text{LyC}} / \Delta\nu_D / (1 + \tau_{\text{Ly}\alpha})$ (mode 0). That estimate peaks the $n=2$ population near $2R_p$ and *under-pumps* it across 1.1–1.8 R_p , leaving a `jlya_mode=1` hook (`load_jlya_rt`) to ingest a real RT field. Phase 3b supplies the real $\bar{J}_{\text{Ly}\alpha}(r)$. Three findings reshaped the original plan (a Monte Carlo with an offline two-pass and a β -on-cooling term):

1. **Method $\rightarrow$ escape probability, not Monte Carlo.** The WASP-121b column is extremely thick: $\tau_0(\text{Ly}\alpha) \approx 5.5 \times 10^7$ at the base (Voigt $a \approx 5 \times 10^{-4}$, so $a\tau_0 \approx 10^4 \gg 1$). A faithful acceleration-free MC needs $\sim \tau_0 \approx 10^7$ – 10^8 scatterings/photon (intractable), and core-skipping is both forbidden and incompatible with the core-dominated pumping field. For $a\tau_0 \gg 1$ the Neufeld (1990)/Harrington analytic wing-escape solution is accurate, so we use a local escape-probability source function.
2. **In-code, not offline.** Being a cheap formula, $\bar{J}_{\text{Ly}\alpha}$ is evaluated every timestep inside the run (`jlya_mode=2`, the same lagged-explicit decoupling as the Phase-3a Balmer feedback); a single run is self-consistent. The

jlya_mode=1 option to read an external (e.g. real MC) profile is kept.

3. **No β -on-cooling term.** From Huang §2.5/Fig. 7 directly: the “Ly α cooling 10 $\times$ lower” is the Black (1981) vs CHIANTI *collisional-excitation rate-coefficient* difference, not photon trapping. EXHALE already uses the CHIANTI-consistent Cen coex_rate_HI (§7), so there is no 10 $\times$ to apply. The trapping feedback is instead the Phase-3a H($n=2$) heating (photoelectric Hpe + collisional de-excitation Hdx), which strengthens once $\bar{J}_{\text{Ly}\alpha}$ is correct. The scaffolding β -on-coex hook was reverted.

12.1 The escape-probability field (lya_rt.f90 : jlya_escape_prob)

Per cell, from the converged T , $n_{\text{HI}} = n_{1s}$, n_{HII} , n_e :

- $\Delta\nu_D = \nu_{\text{Ly}\alpha} \sqrt{2kT/m_{\text{H}}/c}$; Voigt $a = (A_{2p1s}/4\pi)/\Delta\nu_D$; optical depth $\tau(r) = \int_r^\infty n_{1s}\sigma_0 dr'$, $\sigma_0 = C_{\text{Ly}\alpha}/\Delta\nu_D$.
- Wing escape probability $\beta = \pi^{-1/4} \sqrt{a/\tau} (\leq 1)$: escape once the photon diffuses to $x_* = (a\tau/\sqrt{\pi})^{1/2}$ where the wing optical depth is unity.
- Internal sources $P = \alpha_B n_e n_{\text{HII}} + q_{1s2p} n_e n_{\text{HI}}$ (recombination cascade + electron-impact $1s \rightarrow 2p$). The escape-probability closure $\bar{J} = S(1 - \beta)$ with the trapped pile-up $n_{2p} = P/(A_{2p1s}\beta)$ gives $J_{\text{int}} = (2h\nu^3/c^2)(g_{1s}/g_{2p})P(1 - \beta)/(A_{2p1s}\beta n_{1s})$. The $(1 - \beta)$ factor makes $J_{\text{int}} \rightarrow S$ when thick and $\rightarrow 0$ when thin ($\beta \rightarrow 1$), instead of $S \propto 1/n_{1s}$ diverging as the gas ionizes in the outer wind (which spuriously raised $\bar{J}$ outward).
- Stellar beam (trapped): Ly α is *scattered*, not destroyed, so it is not killed by $e^{-\tau}$; the broad stellar line (half-width $\Delta v_s \approx \pm 70 \text{ km s}^{-1}$) penetrates through its wings. A photon at offset x reaches r if $|x| > x_1 = \sqrt{a\tau/\sqrt{\pi}}$; for a Gaussian stellar profile of width $X_s = \Delta v_s/v_{\text{th}}$, $T_*(r) = \text{erfc}(x_1/(\sqrt{2}X_s))$. The profile-averaged mean intensity uses the *stellar* line width $\Delta\nu_* = \nu_0 \Delta v_s/c$ (constant), NOT the planetary $\Delta\nu_D$ ($\bar{J} = \int J_\nu \phi d\nu \approx J_\nu(\nu_0)$ for a stellar field flat across the narrow planetary line; dividing by $\Delta\nu_D$ makes J_{star} rise as the wind cools). A bounded buildup $E = 1 + (\text{boost} - 1)(1 - \beta)T_*$ enhances the penetrating beam in the thick wind and $\rightarrow 1$ when thin or at the deep edge, avoiding both the $1/\beta$ over-count (which blew the $n=2$ source to 6300 $\times$ ground-state) and a flat cap. Hence $J_{\text{star}} = \xi F_* T_* E / (4\pi \Delta\nu_*)$, ξ the dayside dilution. $\bar{J}_{\text{Ly}\alpha} = J_{\text{int}} + J_{\text{star}}$.

New input.inp lines (default-off $\Rightarrow$ byte-identical to 3a; a 3b-0 scaffolding step verified this, frozen reference output_pre3b_baseline/):

```
Jlya escape-prob: True           # jlya_mode = 2 (in-line escape-prob RT)
Stellar Lya flux [erg/cm2/s]: 1.0e5 # incident stellar Lya at the planet (x0.35)
Lya stellar halfwidth [km/s]: 70   # broad stellar line Dv_s (penetration + Dnu_star)
Lya stellar boost [-]: 5           # bounded trapping-buildup cap (left un-tuned)
Deexc heat: True                  # Phase-3a collisional de-excitation heating
```

12.2 WASP-121b validation setup (spherical Case A, matched to Huang)

Huang’s Fig. 11 is **Case A** (spherical, no RLOF), so the comparison must match his geometry *and* radial frame — which exposed a setup error: the run had been Roche-truncated at $\sim 2R_p$, and EXHALE’s base ($r=1$) sits at **1 μbar**

($n_0 = 3.1 \times 10^{12}$, $T = 2358$ K), but Huang defines R_p at the **4 mbar transit radius** and places his $1 \mu\text{bar}$ hydro bottom at $\sim 1.25 R_p$ (Fig. 9 top-pressure axis / Fig. 12 green band; the earlier “ $1.46 R_p$ ” note was wrong). The setup was corrected to spherical Case A with R_0 set to Huang’s $1 \mu\text{bar}$ radius ($R_0 = 1.25 \times 1.766 = 2.2075 R_J$, planet mass unchanged $\Rightarrow$ correct base gravity) and outer radius $15.2 (= 19 R_p^{\text{Huang}})$; EXHALE r maps as $R_{\text{Huang}} = 1.25 r$. The Case-A $\dot{M} = 0.077 M_p/\text{Gyr}$ is within $\sim 1.5\times$ of Huang’s 0.052 (Table 3).

12.3 Trend vs Huang Fig. 11 (qualitative — exact match not expected)

Huang’s $\bar{J}_{\text{Ly}\alpha}$ is a **plane-parallel Monte Carlo** RT; ours is an analytic escape-probability formula, so the two cannot agree quantitatively — a quantitative comparison is only sensible once *all* Phase-3b features (incl. the velocity term) are in. Qualitatively (lya_caseA_huang11.py, not in the tree; the same Fig. 11 case-A comparison is now WASP-121b/Lya_vs_huang11.py), both show a significant inner peak (EXHALE ~ 0.6 at 1.35 , Huang ~ 0.5 at $1.6 R_p$) and comparable magnitude (0.05–0.6). EXHALE has a dip at $\sim 1.8 R_p$ and rises gently outward where Huang gently declines; the outward rise traces to the $n2s$ recombination cascade in the Phase-3a $\text{H}(2s/2p)$ model ($n=2$ stays recombination-populated while ground-state H vanishes in the ionized wind), *not* $\bar{J}$, which after the two fixes below is physically sound (J_{int} peaks then declines; $J_{\text{star}} \approx \text{const}$ outward). $\text{H}(n=2)$ heating stays $\sim 1\%$ of total with $\text{Hdx} \gtrsim \text{Hpe}$ near the n_2 peak (cf. Fig. 10).

Two $\bar{J}$ bugs fixed during the trend check. (1) J_{star} used the planetary $\Delta\nu_D$ instead of the stellar width $\Delta\nu_*$, making $\bar{J}$ rise as the wind cooled; (2) J_{int} lacked the $(1-\beta)$ closure, so S diverged ($\propto 1/n_{1s}$) in the ionized wind.

Wind-velocity (Sobolev) escape — implemented, negligible here. $\beta_{\text{tot}} = 1 - (1 - \beta_{\text{static}})(1 - \beta_{\text{Sob}})$ with $\tau_S = C_{\text{sob}} n_{1s} / |dv/dr|$ is now in jlya_escape_prob (the outflow $v(r)$ threaded through excited_H_update). For WASP-121b it changes $\bar{J}$ by $< 1\%$: where it could help (thick inner wind) $\tau_S \sim \tau_{\text{static}} \sim 10^7$ so $\beta_{\text{Sob}} \sim 10^{-7}$, and the thin outer wind already has $\beta_{\text{static}} \approx 1$ — the velocity gradient and the $\text{Ly}\alpha$ opacity scale together ($\propto n_{1s}$), opening no new escape route. **Phase 3b is thus feature-complete.** *Deferred per the trend-first guidance:* the post-all-features quantitative comparison and the $n2s$ outward-rise revisit (a Phase-3a $2s$ effect, not the RT); boost, Δv_s , and the $\pi^{-1/4}$ prefactor are left un-tuned.

13. Phase 4: RLOF / tidal potential (Case B validated; Case D blocker = escape radius, transonic IC added, geometry-limited)

Phase 4 adds the Roche-lobe-overflow (RLOF) tidal physics that drives Huang’s enhanced mass loss — **Case B** $\dot{M} = 0.32$, **Case D** $= 1.03 M_p/\text{Gyr}$ — the regime that matches the observed transit depths (spherical Case A underestimates them). **EXHALE already has the 1D Roche potential** (grav_field.f90, the spherical_domain=.false. branch: planet + stellar-tidal + centrifugal), the L1-truncated domain, and the tidal momentum source (Source.f90) — Caldiroli’s ATES-v2 Roche mode. So the substellar tidal hydrodynamics is already in place.

13.1 Base gravity for each case (Huang Table 3)

The key to each case is the $1\ \mu\text{bar}$ base radius R_0 . Huang’s Table 3 lists $\log g @ 1\ \mu\text{bar} = GM_p/r_{1\mu\text{bar}}^2$ per case — exactly the R_0 -setting quantity ($R_0 = \sqrt{GM_p/g}$):

Case	tidal	$M_p (M_J)$	dT	XUV/Ly α /Fe	$\log g$	$R_0 (R_J)$
A	N	1.1824	0	1/1/1	2.84	2.058
B	Y	1.1824	0	1/1/1	2.70	2.418
D	Y	1.1204	350	0.74/0.35/4	2.62	2.581

R_0 grows $A \rightarrow B \rightarrow D$ as the tidal force + the $dT+350$ lower-atmosphere heating puff the $1\ \mu\text{bar}$ level outward (weaker base gravity), driving the RLOF.

13.2 Case B — validated

WASP-121b in Roche mode (tidal on) + the corrected base + metals + Ly α mode-2 gives $\dot{M} = 0.337 M_p/\text{Gyr} \approx$ Huang’s Case B (0.32), within $\sim 5\%$. With the base-radius fix both cases agree (Case A $0.077 \approx 1.5 \times$ Huang’s 0.052; Case B $0.337 \approx 1.05 \times$ Huang’s 0.32). EXHALE’s existing Roche potential thus captures the RLOF enhancement — no geometric-reconstruction code was needed for the Case B $\dot{M}$ gate.

13.3 Case D — real blocker found (escape radius), transonic IC added

Case D’s base ($R_0 = 2.581 R_J$, $\log g = 2.62$) sits at $\sim 76\%$ of the L1 distance: the L1-truncated Roche domain reaches only $r_{\text{max}} = 1.353 R_p$. Re-investigating the earlier “won’t launch” report uncovered the *actual* blocker, and it is *not* the initial condition.

Root cause — the escape radius lies outside the L1-truncated domain. `define_grid` sets the constant-momentum convergence window $[j_{\text{min}} : N]$ from the **escape radius** (Escape radius `[R_p]:`, default 1.50). It scans for the first cell with $r(j) \geq r_{\text{esc}}$; if *no* cell qualifies — which happens precisely when $r_{\text{esc}} > r_{\text{max}}$ — the DO loop falls through leaving $j = N + N_g + 1$, so $j_{\text{min}} > N$ and the window is **empty**. Fortran `maxval/minval` over a zero-size array return $\mp \text{HUGE}$, so $\text{mom}_{\text{max}} = -\text{HUGE}$, $\text{mom}_{\text{min}} = +\text{HUGE} \Rightarrow du = \text{Inf}$, and the time-derivative norm $dtu = -\text{HUGE}$. Since $dtu = -\text{HUGE} < dtu_{\text{th}}$, the loop reports **“steady state reached” at step 0** — a *false* convergence unrelated to the wind. The earlier bracket (“works $R_0 \leq 2.30$, fails $R_0 \geq 2.40$ ”) is exactly r_{max} crossing the 1.50 escape radius at $R_0 \approx 2.37$, confirming the escape radius, not the IC, as the discriminant. (The earlier “isothermal IC cannot launch the deep-RLOF wind” diagnosis was therefore wrong.)

Fixes. (1) A guard in `define_grid.f90`: if $j_{\text{min}} > N$, re-anchor the window to mid-domain ($r \geq 1 + \frac{1}{2}(r_{\text{max}} - 1)$) and print a loud WARNING — so this can never again masquerade as convergence. (2) For Case D, set Escape radius `[R_p]:` 1.20 (inside the 1.353 domain). With a valid window the diagnostics become finite and meaningful, and **Case D launches and runs** — from *both* the hydrostatic and the transonic IC. (The earlier $du = |(\text{mom}_{\text{max}} -$

$\text{mom}_{\min}) / \max(\text{mom}_{\min}, 10^{-30})$ floor at EXHALE_main.f90:308 is retained as defense-in-depth against a genuine transient zero-momentum cell, but it was *not* what unblocked Case D — with an empty window $\text{mom}_{\min} = +\text{HUGE}$, not zero, so the floor never engaged.)

Transonic-wind IC (implemented, Transonic IC: True). set_IC.f90 gained an optional transonic isothermal-wind profile, solved from the EXHALE potential via the algebraic (Bernoulli) integral of the steady isothermal-wind equation,

$$\frac{1}{2}v^2 - c^2 \ln v - 2c^2 \ln r + \phi(r) = B, \quad c^2 = (1 + \delta p_{bc}) / \rho_{bc},$$

with B fixed by the sonic point ($v = c_s$ where $D\phi(r_c) = 2c^2/r_c$, located by bisection). At each radius the relation has a subsonic ($v < c_s$) and a supersonic ($v > c_s$) root; the wind takes the subsonic root for $r < r_c$ and the supersonic root for $r > r_c$ (1-D bisection on each monotone branch), and ρ then follows from mass conservation $\rho v r^2 = \dot{M}$ pinned to the base BC. Because that IC already satisfies $\rho v r^2 = \text{const}$, $du \approx 0$ at step 0 would itself trip the momentum-constant exit before the cold wind heats; the flag therefore **auto-enables force_start**. If no interior sonic point exists it **falls back to the hydrostatic IC**. Default **off** $\Rightarrow$ existing cases byte-identical (Case B re-verified at $\dot{M} = 0.337$).

Case D outcome — launches, but the 1-D L1-truncated domain undershoots. With the escape-radius fix the wind launches and relaxes, but it **never reaches a sonic point inside the domain**: the heated metal wind is **deeply subsonic everywhere** ($\text{Mach} \lesssim 0.05$ — rising to ~ 0.050 at $r \approx 1.25$, then *falling* to ~ 0.040 at L1). With a subsonic outer boundary the zero-gradient outflow BC is only marginally well-posed and $\rho v r^2$ stays non-constant (peaks $\sim 1.7\times$ the L1 value), so the run settles into a **breathing limit cycle**: after the $\sim 10^2$ heating transient damps, du keeps oscillating in $\sim 0.04\text{--}0.8$ (touching but never crossing the strict $du_{\text{th}} = 0.02$) for as long as it runs. The mass flux is correspondingly time-variable — the quasi-steady $\dot{M}$ oscillates over $\sim 0.3\text{--}0.6 M_p/\text{Gyr}$, well *below* Huang’s 1.03, the deficit set by the L1 truncation cutting off the still-accelerating subsonic wind before it reaches the sonic point. The transonic IC behaves identically to the hydrostatic IC here: both launch the same subsonic flow. (Warm-starting via load_IC reintroduces a base NaN — the documented grid re-read inconsistency — so it is not a workaround.)

Conclusion. Case D’s deep RLOF ($\dot{M} = 1.03$) is **not reachable in EXHALE’s 1-D L1-truncated Roche mode**: the wind is subsonic at L1, so the truncated domain structurally undershoots. Capturing it needs the 3-D geometric treatment (Phase 5) or a domain extended past L1 — *not* a different IC. The transonic IC and the escape-radius guard are correct, necessary infrastructure (they remove the false convergence and let the deep-RLOF wind launch and run with meaningful diagnostics) but are not sufficient on their own for Case D’s full mass-loss rate. The geometric reconstruction (substellar $\rightarrow$ terminator mapping, triaxial radii Table 2, the 4/9 mass-loss conversion) feeds Phase 5 (§14), and also supplies the geometry this subsonic L1-truncated wind needs.

14. Phase 5: velocity-broadened transmission spectrum (in progress)

Phase 5 is the headline deliverable — the NUV/optical **transit-depth spectrum** the whole upgrade exists to match (Huang Table 3 / Figs. 17–24). It extends the existing spherical post-processor EXHALE_transit.py in two stages: **5a**

adds the metal resonance lines (spherical, gated on Case A); **5b** replaces the spherical geometry with the 3-D Roche-equipotential reconstruction + velocity broadening (gated on Case D). Pure Python post-processing — no Fortran change.

Line	Case A	Case D	Observed
Mg II $\lambda 2796$ (4 Å)	0.182	0.302	0.309 ± 0.036
Ca II K $\lambda 3934$	0.199	0.278	0.281 ± 0.009
H α $\lambda 6563$	0.201	0.185	0.186 ± 0.003
H β $\lambda 4861$	0.174	0.135	0.143 ± 0.005
Na D2 $\lambda 5890$	0.152	0.147	0.147 ± 0.002

These are the **effective transit radius** R_{eff}/R_* , not an absorption percent (the expected-results doc states the unit is R_p/R_*). Our code returns the line-center absorption fraction h , mapped to the same quantity via $R_{\text{eff}}/R_* = \sqrt{(R_p/R_*)^2 + h}$.

14.1 5a — metal resonance lines (Mg II, Ca II, Na I D), spherical — done

EXHALE_transit.py gained Mg II $\lambda 2796$, Ca II K $\lambda 3934$, Na I D2 $\lambda 5890$ as an **isolated appended block** (the validated He 10830 / Ly α / H α / H β pipeline above it is untouched). These are resonance lines whose lower level is the ion ground state; at $\sim 10^4$ K the excited fine-structure levels are Boltzmann-negligible, so the lower-level density is the ion density itself ($n_{\text{lower}} \approx n_{\text{ion}}$). The ion densities live past the H/He block of Ion_species_adv.txt, in mion_fsp order (species_table.f90): numpy columns **17 = Mg II**, **23 = Ca II**, **25 = Na I**. The block reuses the existing spherical chord / Voigt (wofz) / trapezoid- τ / disk-average / instrument convolution path (resonance_depth()).

Case A result (line-center % / 4 Å-band %) and diagnosis. Mg II: 26.1 / 2.3 ($R_{\text{eff}} = 4.2 R_p$); Ca II K: 5.5 / 0.33 ($1.9 R_p$); H α : 3.9 ($1.6 R_p$); Na D2: 0.8 / 0.06 ($0.7 R_p$). Two real gaps: (1) **Mg II already reaches Huang's extent** ($R_{\text{eff}} = 4.2 R_p$ vs his clustered $\sim 3.6 R_p$) — the metal treatment is sound — but **Ca II / Na I / H α are too optically thin** ($R_{\text{eff}} = 0.7\text{--}1.9 R_p$), a line-specific ionization / abundance / $n = 2$ question, not geometry. (2) The lines are **too narrow**: the Case A wind reaches only ~ 29 km/s, so a 26%-deep core washes out to $\sim 2\%$ over a 4 Å (± 215 km/s) bin, whereas Huang's lines stay deep *across* the bin — they are velocity-broadened to hundreds of km/s by the fast Roche outflow + rotation, i.e. **Phase-5b physics**. (The existing H α tool shows the same shortfall, so it is systematic.) *Units — resolved*: Huang's table is R_{eff}/R_* (radius ratio; the expected-results doc header states the unit is R_p/R_* , not absorption percent). The Case A spherical run over-extends ($15 R_p$ domain $\Rightarrow$ Mg II reaches $4.2 R_p$, $R_{\text{eff}}/R_* = 0.51$ vs Huang 0.182); the proper test is the L1-confined Roche run in R_{eff}/R_* units (§5b).

14.2 5b — Roche-equipotential reconstruction + broadening — implemented

New module roche_recon.py. Assuming a uniform state on Roche-equipotential surfaces, the field at any (x, y, z) equals the 1-D substellar state at the equivalent radius r_{eff} where $\phi_{\text{sub}}(r_{\text{eff}}) = \phi_{3D}(x, y, z)$. The dimensionless Roche potential (R_p units) mirrors grav_field.f90:

$$\phi(x, y, z) = -\frac{1}{r} - \frac{q}{r_s} - \frac{1+q}{2a^3} [(x-x_{\text{cm}})^2 + y^2], \quad q = M_*/M_p, \quad a = \text{orbit}[R_p].$$

Tidal terms $\rightarrow 0$ recovers $\phi = -1/r$ (spherical), so the triaxial reconstruction reduces to the 5a spherical case for Case A. The module gives `roche_phi`, the L1 root of $d\phi_{\text{sub}}/dr = 0$, a monotone $\phi_{\text{sub}} \rightarrow r_{\text{eff}}$ inverse (ReconMap), the triaxial radii, and the Eggleton lobe radius. EXHALE_transit.py gained a `geometry='triaxial'` switch (`triaxial_depth()`): the transit LOS runs along $+x$; the state at each 3-D point is the substellar state at r_{eff} ; and the **LOS velocity in each sector** $v_{\text{sub}}(r_{\text{eff}})x/r - \Omega y$ (wind + tidally-locked rotation, Huang Eq. 16) is integrated over **20 angular sectors**.

Demonstrated on the converged Case B (Roche) run. Geometry self-consistent: $L1 = 1.574 R_p$; $R_{\text{RL}}/L1 = 0.714 \approx 2/3$; tidally triaxial $R_{px} = 1.05 > R_{py} = 0.919 > R_{pz} = 0.893$ (substellar bulge, polar compression).

Result — the model reproduces Huang’s transit radii. Converting the Case B triaxial depths to R_{eff}/R_* (the table unit):

Line	Huang A	Huang D	Case B (this work)
Mg II	0.182	0.302	0.234
Ca II	0.199	0.278	0.222
Na D2	0.152	0.147	0.163
H α	0.201	0.185	0.239
H β	0.174	0.135	0.232

Case B is the *intermediate* model (between Case A and Case D), and the **metals (Mg II, Ca II, Na) land right between Huang’s Case A and Case D** — the lobe-confined reconstruction correctly reproduces his transit radii in magnitude and trend. The Balmer lines run $\sim 25\%$ high (H α 0.239 vs the expected ~ 0.19), consistent with the 5a finding that $n = 2/T$ is slightly over-predicted.

Correction. An earlier draft claimed the lobe-confined model “undershoots Huang because it omits the beyond-L1 escaping tail.” That was wrong twice over: (1) Huang’s reconstruction is *also* lobe-confined (equipotential mapping to the Roche-lobe boundary, §1.8) — neither model has a beyond-L1 tail, so it cannot be the difference; and (2) the apparent shortfall was a **units error** (our absorption *percent* vs Huang’s R_{eff}/R_*). In the correct units the two agree (table above). The small triaxial-vs-spherical reduction (terminator compression $R_{py} < R_{px}$) is minor; Huang’s area-preserving $R_{py}R_{pz} = R_p^2$ normalization would offset it.

14.3 5b — direct Case A comparison: hydro validated, metals over-extended

A clean spherical Case A at the *correct* 1 μbar base $R_0 = 2.058 R_J$ (Table 3 $\log g = 2.84$; the old output_caseA used $R_0 = 1.766$, the 4 mbar transit radius) was run to convergence.

- **Hydro / base validated:** $\dot{M} = 0.056 M_p/\text{Gyr} \approx$ Huang’s Case A 0.052 (the wrong base gave a different $\dot{M}$) — base and wind are right.
- **Metal transit radii — consistent once compared like-for-like** (Huang quotes Mg II in a 4 Å NUV bin; Ca II / Na are optical line-center): Mg II (4 Å) **0.220** vs Huang 0.182 ($\times 1.2$); Na D2 (LC) **0.175** vs 0.152 ($\times 1.15$); Ca II K (LC) **0.306** vs 0.199 ($\times 1.5$, high).

The “over-extension” was a comparison-definition artifact, not physics. An earlier draft reported $\text{Mg II } R_{\text{eff}}/R_* = 0.60$ ($3.3 \times$ Huang) — that compared the *line-center* R_{eff} against Huang’s $4 \text{ \AA}$ NUV-bin value. In the same $4 \text{ \AA}$ bin $\text{Mg II} = 0.22 \approx$ Huang’s 0.182; Na matches too; only Ca II is genuinely $\sim 1.5 \times$ high.

The metal ionization is correct and is NOT the lever. All channels are wired and physical (Verner+1996 σ , Badnell RR+DR, Voronov, Huang+2023 charge exchange; $a_\tau = 0$). The Mg II/III balance is a close competition between a weak photoionization ($\Gamma \approx 2 \times 10^{-3} \text{ s}^{-1}$, weak because $\sigma_{\text{MgII}} \approx 0.24 \text{ Mb}$ at threshold) and $\text{RR} \cdot n_e$; DR ($\sim 5 \times 10^5 \text{ K}$ resonances) and collisional ionization ($e^{-15 \text{ eV}/kT}$ at 10^4 K) are negligible. A test showed $\times 30$ **more Mg II photoionization moves R_{eff} only $0.57 \rightarrow 0.32$** — the resonance line is so optically thick that ionizing it harder cannot confine it, so the ionization is a dead end as a lever (and is not the cause once the bin is matched).

Status: the full Phase-5 pipeline is implemented and runs end-to-end; the hydro/ $\dot{M}$ reproduces Huang (Case A 0.056, Case B 0.337); and the metal transit radii match Huang once compared in his own bins (Mg II $4 \text{ \AA}$ 0.22 vs 0.182, Na 0.175 vs 0.152). Remaining items are minor: **Ca II $\sim 1.5 \times$ high** (line-center; plausibly Ca II $\rightarrow$ Ca III or the Ca abundance), the slightly-high Balmer $n = 2$, the $R_{\text{py}} R_{\text{pz}} = R_p^2$ normalization (methodological), a terminating Case D run, and phase5_transmission.ipynb. Comparisons must use Huang’s line-by-line definition ($4 \text{ \AA}$ bin for the NUV metals, line-center for the optical lines) — the two differ by $\sim 3 \times$ for the optically-thick Mg II.

15. Post-process advection guard at a “breathing” base (2026-06-07)

Symptom. The spherical-vs-Roche example figure (user manual Fig. 4, WASP-121b) showed a sharp temperature spike at $r \approx 1.1 R_p$ in the *advection-corrected* (`_adv`) Roche profile, with a sawtooth band just below it. The converged equilibrium (eq) profile is smooth there, so the spike is a post-processing artifact, not a feature of the hydro solution.

Diagnosis (two compounding causes).

1. *A breathing, inflowing base (1-D limitation).* WASP-121b nearly fills its Roche lobe, so the 1-D wind is subsonic at L_1 and the dense lower atmosphere settles into a slowly recirculating state with small *negative* (inflow) velocities and a stagnation point ($v = 0$) near the wind base ($r \approx 1.1$). This is intrinsic to spherical symmetry, not a bug.
2. *A spurious hot root in the metal-cooled post-processor.* The advection post-processor re-solves the energy/ionization balance in each cell assuming an outflow (each cell upwinds from the next-inner cell). In the dense base the non-monotonic metal line-cooling curve gives the energy equation a second, spurious *hot* root that the Newton (MINPACK) solver can land on; the upwind coupling then cascades one bad cell outward, producing the sawtooth and the overshoot spike. The pre-existing $2 \times$ band guard around eq T did not catch the $\sim 1.4 \times$ spike.

Fix (option c: eq fallback where the flow is not a clean outflow). In `post_process_adv.f90`, wherever $v \leq 0$ (the breathing inflow base) the post-processor now *skips* the advection correction and keeps the converged eq ionization and temperature. Three guards were added, all gated on `pp_metal_on` so the metals-off mode stays byte-identical: the no-He and He H/He advection loops (set the H/He fractions to eq and `cycle` when $v(j) \leq 0$ or $v(j-1) \leq 0$) and the cell-by-cell temperature solve (set $T_{\text{out}} = T_{\text{in}}$ and `cycle` when $v \leq 0$). Keeping the inner cells at eq removes the spurious

root *and* breaks the upwind cascade at its source, while the genuine `_adv` correction is retained in the outflow above the stagnation point.

Validation (PP-only re-run on the converged Case B state). Rebuilt and re-ran the post-processor (Do only PP on the converged IC):

- eq reproduced: $\max |\Delta T|/T = 2.6 \times 10^{-3}$, $\max |\Delta v| = 0.015 \text{ km s}^{-1}$; $\log_{10} \dot{M} = 13.38$ unchanged ($\approx 0.34 M_p \text{ Gyr}^{-1}$); Mg II reproduced to 1.6% in the line-forming region.
- `_adv` base oscillation eliminated: temperature sign changes in $[1.02, 1.10] R_p$ dropped from **36 to 0**; the spike is gone; the eq $\rightarrow$ `_adv` handoff at $v = 0$ is continuous (0.7% step) and the few-percent wind correction is preserved.
- Spherical Case A ($v > 0$ everywhere) is untouched — the guard never triggers, so its `_adv` is byte-equivalent.

Cause 1 (the breathing base) remains a 1-D limitation but is now represented there by the smooth eq solution; resolving the inflow itself would require a 2-D/3-D model.

Documentation updates (this session). The user manual (`EXHALE_user_manual.tex`) gained: §1.3 *The gravitational potential* (spherical, Roche substellar $\phi(r)$ and $d\phi/dr$, and the 3-D off-axis form, matching `grav_field.f90`); §3.5 *Spectral coverage for metal photoionization* (neutral-metal thresholds down to K I at 4.34 eV/2856 Å, and how the power-law sub-grid extends below 13.6 eV vs. a loaded SED that must itself reach those wavelengths); §8 *Atomic and molecular data* (a by-process source table — Asplund09 abundances, Verner96 photoionization, Voronov97 collisional ionization, Hui & Gnedin97 H/He + Badnell06 RR/adf48 DR metals + Shull & Van Steenberg82 Ca I + Huang23 Fe, Huang23 charge exchange, AIOLOS+CHIANTI cooling, Christie13/Draine11/Huang17 excited-H, NIST line data; and the explicit note that the code is purely atomic/ionic with *no* molecular data); and the §4 *breathing-base caveat* documenting the issue and fix above. The Fig. 4 example was regenerated from the corrected `_adv` output.

16. He 2³S triplet + metals simultaneous solve, Ly α -flux guard, and convergence workflow (2026-06-08)

Motivation. Until now the He metastable triplet (He 2³S, “HeITR”) and the trace-metal ionization network could not be solved in the same run: both placed an extra unknown at `x(4)` of the coupled ionization residual, and `input_read.f90` allocated the metal unknowns only if `(thereis_metals .and. .not. thereis_HeITR)`. Modeling HD189733b with the triplet *and* solar metals requires both at once.

Merged solver (factual). New module `nonlinear_system_solver/System_HeH_TR_metals.f90` (`ion_system_HeH_TR_metals`, 84 SLOC). The H/He/triplet rows ($x_1 = \text{HII}$, $x_2 = \text{HeII}$, $x_3 = \text{HeIII}$, $x_4 = \text{He } 2^3\text{S}$) are taken verbatim from `System_HeH_TR`; the metal rows are taken verbatim from `System_HeH_metals` and shifted to $x(5+2(e-1))$. The two blocks share one electron density $n_e = n_{\text{HII}} + n_{\text{HeII}} + 2n_{\text{HeIII}} + \sum_e (n_{X^+} + 2n_{X^{++}})$ (the neutral triplet is excluded), and charge exchange is added through `cx_add_to_fvec`. Solved with `hybrd1`. Supporting edits: `charge_exchange.f90` gained a settable base index `cx_metal_base` (default 4; the merged dispatch sets 5) so the metal charge-exchange rows land at the

right offset; `input_read.f90` sets $N_{\text{eq}} = 4 + 2n_{\text{melem}}$ when both are active; `ionization_equilibrium.f90` uses a generalized metal base index (`mbase = 4`, or 5 with HeITR) in the initial guess / solution scatter and a both-on dispatch branch; the new source was added to the Makefile. The radiation grid, cross sections, photo-heating, IC setup and post-processing already handled both species; only the equilibrium *solver* needed merging.

Regression & physical check. With metals on / triplet off the run is byte-identical to the pre-merge binary over 8 048 steps (the new branch is not taken). With both on, the coupled solve is clean (no NaN) and was exercised end-to-end on WASP-121b with the full stack (triplet + metals + excited-H + $\text{Ly}\alpha$), reaching a flat converged solution (peak $\text{He } 2^3\text{S}$ fraction $\sim 5 \times 10^{-6}$).

$\text{Ly}\alpha$ escape-probability input guard (factual). `input_read.f90`: if `Jlya escape-prob: True (jlya_mode=2)` but `Stellar Lya flux [erg/cm2/s]:` is missing or ≤ 0 , the run now stops with an explanatory message rather than silently dropping the stellar beam (in mode 2 $\bar{J}_{\text{Ly}\alpha} = \bar{J}_{\text{int}} + \bar{J}_*$ with $\bar{J}_* \propto F_{\text{Ly}\alpha,*}$, so $F_{\text{Ly}\alpha,*} = 0$ would make the halfwidth/boost settings silent no-ops).

Convergence threshold and Reconstruction scheme:-controlled two-stage reconstruction (factual).

- `du_th` is now a runtime variable (was a compile-time parameter) with default $1.0\text{e-}3$ — the original ATES-Code-main value. EXHALE had carried $2.0\text{e-}2$, at which a run flagged “converged” can still show $\sim 2\%$ spatial spread in the supersonic mass flux $\rho v r^2$.
- The single-stage vs. two-stage choice is now set by the `Reconstruction scheme: line`. `Reconstruction scheme: PLM+WENO3` selects a two-stage run: with the input line `du_th [PLM,WENO3]: <du_plm> <du_final>` supplying both thresholds (`du_th_plm` and `du_th`), `EXHALE_main.f90` starts in PLM and switches `rec_method` to WENO3 once $du < du_{\text{plm}}$ (or PLM stalls), then converges at $du < du_{\text{final}}$. This automates the manual `PLM → Load IC+WENO3` workflow recommended in the ATES README. `Reconstruction scheme: PLM` or `WENO3` is single-stage and uses only the first `du_th` value (the second is ignored).
- CFL is now runtime and settable via the input line `CFL: <value>`.

Other source-term changes (kept, with caveats). `energy_semi_implicit.f90` now damps the implicit energy update with $\partial F / \partial T = 1 + c |dC/dT|$ (was $\max(0, dC/dT)$), adding damping on the falling cooling branch; it removed one oscillation in testing but did *not* by itself fix the convergence behavior below. `Cool_coeff.f90` now interpolates the cooling tables with monotone PCHIP (C^1) rather than linear (C^0); this was implemented to test an interpolation-kink hypothesis (not supported, below) and is retained as a smoothness improvement.

Convergence — provisional observations (not conclusions). *The following is a research log, deliberately tentative; several intermediate “non-convergence” readings in earlier sessions were later traced to methodology, not physics.* Two artifacts are recorded so they are not repeated: an over-short step cap (weak-wind cases judged “flooded” near 25k steps while du was still descending — HD209458b later reached $du < 2 \times 10^{-2}$ near 50k steps), and a run-script bug that killed the backgrounding subshell rather than the `EXHALE.x` child, so runs continued orphaned and the reported du

was a premature snapshot (use `kill -x EXHALE.x`). Because of these, the earlier framing that “metal line cooling destabilizes the wind” should be treated as *not established*. With the two-stage strict-du_th pipeline we currently see, across planets:

- **WASP-121b** (inflated, strongly irradiated): converges quickly and flat ($\sim 0.16\%$ supersonic mass-flux spread, vs $\sim 2\%$ for single-stage PLM). This case looks solid.
- **HD209458b**: reaches a small $du \sim 0.02\text{--}0.025$ that then oscillates rather than dropping to 10^{-3} . *Tentatively* usable as a quasi-steady state: a He $2^3\text{S}+7$ -metal snapshot from the two-stage pipeline gives a smooth transonic wind with $T_{\text{max}} \approx 6.8 \times 10^3 \text{ K}$ and $\log_{10} \dot{M} \approx 10.1$ (median over the supersonic region), consistent with the literature $\sim 10^{10} \text{ g s}^{-1}$; see `heitr_metals_dev/HD209458b_quasisteady.ipynb`.
- **HD189733b** (strongly bound, $\beta_0 \approx 192$): du oscillates around ~ 2 in a sustained limit cycle and does not reach the threshold — the problematic case.

A plausible (but unproven) reading is that the difficulty tracks how weakly driven / strongly bound the wind is, rather than the metal content per se.

Base-breathing hypothesis and the fixes that did *not* work. The HD189733b oscillation resembles the known base-breathing behavior, and the lower boundary (`Apply_BC.f90`, `BC_component_constrho`) hard-pins the ghost density and pressure to cold reservoir values with a one-way velocity valve $\max(v_1, 0)$ — a configuration that can reflect acoustic waves. Three variants were tried and reverted:

Test	Outcome
Lower CFL ($0.6 \rightarrow 0.2$)	du floor barely moved ($\sim 2.8 \rightarrow \sim 1.9$)
Zero-gradient base velocity (drop the valve)	du still ~ 2.3
Riemann-invariant base velocity $v = v_1 + \frac{2(c_b - c_1)}{\gamma - 1}$	did not help HD189733b <i>and</i> broke WASP-121b

So the base-breathing reading is *consistent with* the observations but is **not demonstrated**, and these simple BC variants are not solutions. Candidate routes if pursued (both substantial, uncertain payoff): a characteristic / non-reflecting (NSCBC-style) base BC; a steady-state Newton / BVP solver (which would not orbit a time-marching limit cycle if a steady solution exists); or accepting HD189733b as a marginal, time-averaged quasi-steady case. A separate, cheaper idea for *speed* (not the oscillation) is local cell-by-cell time-stepping, since the global dt is currently set by the smallest base cell (`eval_dt.f90`). Full working notes: `docs/heitr_metals_and_convergence_notes.md`; solver-option discussion: `docs/numerical_methods.md`.

Code-size note. The merge adds one Fortran module (`System_HeH_TR_metals.f90`, 84 SLOC), so the 2026-06-07 file tally of 57 becomes 58; the remaining changes are small edits to `charge_exchange.f90`, `input_read.f90`, `ionization_equilibrium.f90`, `parameters.f90`, `EXHALE_main.f90`, `energy_semi_implicit.f90` and `Cool_coeff.f90`.

17. Metals in the bulk-gas mass/charge/particle budget (eos_metals, 2026-06-13)

Motivation. Until this update the trace metals were strictly *passive* species: they contributed to heating/cooling and were solved in the coupled ionization equilibrium, but they were excluded from the mass density, the base density ρ_{bc} , and the electron/particle densities used everywhere *outside* the ionization solver (EOS temperature, opacities, cooling calls, energy update). The only self-consistent place was the charge balance *inside* the MINPACK residuals (System_HeH_metals.f90). At solar abundances the missing metal mass is $\sim 1.3\%$ of the H/He mass — small but systematic, and inconsistent with solving the metal ionization to high precision. Recorded in TO_BE_DONE.md on 2026-06-12; implemented 2026-06-13.

Runtime interface. A new key eos_metals 0|1 in metals.inp (documented in inputdata/metals.inp.example). 1 (the new default) includes the metals in the bulk-gas budget; 0 restores the legacy trace approximation exactly. Metals-off runs (no metals.inp) are identical either way.

What changed (factual).

- **Composition constants** (input_read.f90): the mass per hydrogen nucleus is now $\mu_H \equiv \text{mass_per_H} = 1 + 4Y + \sum_X a_X A_X$ (with $Y = \text{HeH}$ and a_X, A_X the metal abundance and atomic weight), and the base density becomes $\rho_{bc} = \mu_H / (1 + Y)$ in units of $m_H n_0$. Atomic weights live in species_table.f90:melem_A. The normalization n_0 keeps its H/He-nuclei meaning, so $n_H = n_0 / (1 + Y)$ is unchanged. (Correction to the original TO_BE_DONE entry: v_0, p_0, b_0 are m_H -based normalizations and never contained a mean molecular weight — they were never affected.)
- **Mass density:** calc_rho takes an optional nm argument and adds $\sum A_X n_{X,i}$ over all stages; passed from ionization_equilibrium.
- **Electron density:** calc_ne takes an optional nm and adds $\sum q_i n_{X,i}$; passed at every site where nm is available — composition (main-loop EOS), ionization_equilibrium (opacity/cooling), util_ion_eq (eval_cool + cooling dump), energy_semi_implicit, excited_hydrogen, post_process_adv. (The scalar T_equation already carried metal electrons via pp_nm_cell.)
- **Particle density:** calc_ntot likewise adds the metal nuclei, so the EOS $T = p / (n_{\text{tot}} + n_e)$ sees them.
- **Ghost pressure / dp_bc:** the lower-BC ghost pressure is now $n_{\text{tot_bc}} + \text{dp_bc}$ with $n_{\text{tot_bc}} = (1 + Y + \sum_X a_X) / (1 + Y)$ — so $T = T_0$ holds *exactly* at the base — and dp_bc includes the metal electrons $\sum q_i n_{X,i} / n_0$ at the first cell. set_IC/load_IC use mass_per_H/ntot_bc consistently. Full formulas: EXHALE_BC_and_IC.tex.

Validation (WASP-121b regression matrix).

- eos_metals 0 is **byte-identical** to the pre-change goldens for both regression cases (wasp_full 7296 steps, wasp_he23off 7300) — the legacy path is untouched.

- `eos_metals 1` (new default) converges cleanly (7145/7142 steps, $\sim 2\%$ fewer); profile shifts are at the expected $\sim 1\%$ mass scale: median $|\Delta n| \sim 2.4\%$, $|\Delta p| \sim 2.3\%$, $|\Delta T| \sim 0.9\%$ (max $|\Delta T| = 132$ K in the breathing-base region $r \sim 1.09$). The *number* flux nvr^2 drops 1.1% ($\log_{10} \dot{M}$: 13.31 $\rightarrow$ 13.30), which offsets the +1.2% mass per particle — the *mass*-loss rate is essentially unchanged.
- Goldens re-snapshotted under the new default; a repeat run reproduces them byte-identically (single-thread determinism intact).

Still open (small). The `use_2lev_cool` legacy cooling branch remains AIOLOS-hard-coded (default off), and the cell-by-cell scalar `T_equation` uses `pp_nm_cell` only in the post-process path (the time-marching cell-by-cell solve receives n_e via the now-metal-aware `energy_semi_implicit`).

18. Automatic initial-condition selection (IC mode: auto, 2026-06-13)

Motivation. The code now carries three IC families — cold hydrostatic (default), transonic isothermal wind (added for the Phase-4 deep-RLOF Case D, where the cold IC false-converges), and the hot-Parker warm seed (benchmarked: no speedup) — and the choice was manual. Kubyshkina+2018 build an IC for each planet automatically; the design study is `docs/auto_ic_design.md` and the literature comparison `docs/code_comparison.tex`.

The selector (one exact probe, no tunable threshold). `select_IC_auto` (`set_IC.f90`) evaluates the *cold* base sound speed in code units, $c_{\text{cold}}^2 = (n_{\text{tot,bc}} + \Delta p_{\text{bc}}) / \rho_{\text{bc}}$, and calls the existing `find_sonic` bisection for a critical point $\phi'(r_c) = 2c^2/r_c$ in the *actual* potential (Roche or spherical). If an interior sonic point exists, the cold isothermal atmosphere already wants to blow a transonic wind, so the transonic IC is selected; otherwise the cold hydrostatic IC is kept. One criterion catches both regimes that need the transonic start: deep RLOF (through the Roche/L1 topology) and low-gravity boil-off ($r_c \simeq b_0 / (2c_0^2)$) in a point-mass potential). The Jeans parameter b_0 (Kubyshkina's Λ) is logged as a diagnostic but is *not* a decision variable; the hot-Parker seed is never auto-selected.

Runtime interface. New input line `IC mode: <cold|transonic|hot_parker|auto>`. The named modes simply set the legacy flags; `auto` defers the decision to `set_IC`. The legacy keys (`Transonic IC:`, `Hot Parker IC:`) take precedence over `auto`, and the setup report (`EXHALE_setup.out`) records the family actually in effect, with a “(chosen automatically: IC mode = auto)” tag. Touched files: `parameters.f90` (`ic_mode`), `input_read.f90` (key parsing), `set_IC.f90` (`select_IC_auto`), `write_setup_report.f90`.

Validation gates (all run 2026-06-13).

Gate	Result
A-1 no IC mode key	byte-identical regression (defaults untouched)
A-2 IC mode: cold	byte-identical to A-1
A-3 spherical tutorial, auto	$b_0 = 83.2$, no interior $r_c \rightarrow$ cold; IC_dump byte-identical
B-1 Case D (deep RLOF), auto	$r_c = 1.297$ (just inside $L1 = 1.353$) $\rightarrow$ transonic, same as the manual key
B-2 boil-off planet ($0.0189M_J$, $0.446R_J$, 1100 K)	$b_0 = 8.5$, $r_c = 5.2$ (as predicted) $\rightarrow$ transonic; wind launches (initial residual 0.19 vs 14 cold)
B-3 wasp_full (Roche), auto	$\rightarrow$ transonic; du -stop in 5876 vs 7145 steps (−18%)

Honest speed verdict. The −18% step count is *Roche/deep-RLOF-specific* (the steady wind there is close to the cold isothermal transonic solution). For the mild spherical boil-off case B-2 the cold control *also* launched and converged first ($\log_{10} \dot{M} = 11.20$), while the auto run relaxed monotonically toward the same flux from $\sim 2\times$ above — the cold- c^2 Parker profile *overestimates* the flux of a weakly heated wind. So auto is not harmful (same attractor) but not faster there.

Key byproduct (a trap for all $\dot{M}$ work). While adjudicating gate B-3 it was found that the $du < 10^{-3}$ stopping criterion is **path-dependent at the $\sim 5\%$ level**: the cold-start and transonic-start wasp_full runs both “converge” by du but sit 4.8% apart in mass flux, each only 1–2% flux-flat. Newton-finishing *both* states (Load IC + Solver: Newton) collapses them onto the *same* fixed point (flux ratio 0.9996, flatness 0.005%). Practical rule, now standing: **for quantitative $\dot{M}$ always finish with Solver: Newton**; never compare bare du -stops at the percent level.

Deferred. A Phase-C retry ladder (escalate cold $\rightarrow$ warm seed on NaN/stall) is designed in auto_ic_design.md §6 but *not* implemented — the B-2 result suggests failure-triggered escalation is the right shape if a real case ever needs it. IC formulas and the selection logic are also documented in EXHALE_BC_and_IC.tex §5 and the user manual.

19. Base boundary conditions and the flux-based convergence decision (2026-06-15)

Context. A systematic study of how the lower boundary and the stabilizers interact with convergence on HD 209458b is collected in docs/EXHALE_BC_and_IC.tex (test matrix, two stabilizer tables, and a convergence-criterion history). This entry records the resulting default changes to the code; the companion document holds the full evidence.

Shapiro filter: off by default. The optional Shapiro (1970) spatial filter — ported from CETIMB (Koskinen et al. 2013a), whose use of it was found through Huang et al. (2023), to damp the same gravity-unbalanced base sound-wave instability — damps the transient “breathing” of the base cell, but on a clean cold-start it also drives the wind into a slow infall and prevents convergence. Isolated in a 2×2 (filter $\times$ base-valve) sweep, the filter — not the valve — was the cause. It is therefore *off* by default (shapiro_eps < 0) and opt-in only via Shapiro filter: <eps> [every].

Base boundary condition: density vs. pressure. The base can now be anchored either on density (the legacy fixed n_0) or on a fixed base pressure, selected by Base BC: `density|pressure [p_ubar]` (default density; pressure fixes p_{base} in μbar and derives n_0). In practice the converged wind is nearly insensitive to which anchor is used — the real lever on $\dot{M}$ is the base *density* itself, not the BC form.

The flux-based convergence decision (history). EXHALE measures two distinct quantities at the base: the flux metric du (the radial spread of $\rho v r^2$, which is exactly the constant-mass-flux criterion of ATES, $\Delta\dot{M}/\dot{M} < 10^{-3}$, and of CETIMB, $F_c = \rho v r^2$ altitude-flat), and the steady residual $\|R\| = \|\partial_t u\|$ over mass/momentum/energy, a *stricter* EXHALE-specific check. The convergence rule evolved: before the Newton finish it was purely flux-based; the Newton work then added the stronger $\|R\|$ criterion; but the reference codes converge on flux alone. The standing decision, adopted here, is therefore **flux-based** — $du < du_{\text{th}}$ is the acceptance test, and $\|R\|$ is computed and reported *for reference only* (it gates the stop only if `Resid tol:` is set).

Volume-weighted residual norm. When $\|R\|$ is reported, the default is now a cell-volume-weighted relative norm, $\sum_j |R_j| r_j^2 \Delta r_j / \sum_j |u_j| r_j^2 \Delta r_j$ (`resid_vol`, settable `Resid norm: vol|Linf`), rather than the bare L_∞ maximum. On the non-uniform radial grid the L_∞ max is dominated by the few tiny near-base cells and overstates the residual; the volume weighting gives a physically representative number. The infall diagnostic used in the test matrix was likewise switched to a *radius-weighted* negative-velocity fraction (weighted by Δr over the linear radius range), so that sub- 10 m s^{-1} near-base noise is not counted as bulk infall.

du -keyed Newton/JFNK hand-off. The production Newton finish (`Solver: Newton`) now hands off from the marching warm-up to the JFNK steady solve when the *flux* metric drops below `newton_du_switch` (default 10^{-2}), consistent with the flux-based decision and adjustable via an optional third token, `Solver: Newton [du_switch]`. This replaces the earlier residual-keyed hand-off ($\|R\| < 5 \times 10^{-2}$, variable `newton_R_switch`, now removed); the two are equivalent in practice, but keying on du avoids tying the hand-off to the stricter residual. The JFNK solve still *targets* $\|R\| < \text{resid_th}$ (10^{-3}); on failure it keeps its best iterate and falls back to marching with the du -based stops re-armed.

The full-physics residual floor (base $j = 1$). With $\text{He } 2^3\text{S} + \text{metals}$, the du -keyed hand-off plus the volume-weighted norm drove the JFNK residual from ~ 0.17 down to $\sim 2 \times 10^{-3}$ — about two orders of magnitude better — but it then *stalls* (`info` $\neq 0$) on a localized momentum imbalance at the base cell $j = 1$ (the “breathing base”), short of 10^{-3} . The wind is nonetheless flux-converged (e.g. HD 209458b full physics: $\log_{10} \dot{M}$ steady, $\rho v r^2$ spread $< 1\%$), which is a converged run by the reference standard. Closing the last factor of ~ 2 in $\|R\|$ requires physically damping the base momentum — the explicit-viscosity task at the top of `TO_BE_DONE.md`.

Touched files. `parameters.f90` (`shapiro_eps`, `base_bc_mode/base_p_ubar`, `base_v_massflux`, `resid_vol`, `newton_du_switch`; `newton_R_switch` removed); `input_read.f90` (`Base BC:`, `Resid norm:`, `Solver: Newton [du_switch]` parsing); `steady_residual.f90` (`residual_norms_vol`); `steady_newton.f90` (`resid_relnorm vol/L $\infty$` branch); `EXHALE_main.f90` (du -keyed JFNK trigger, always-on flux/residual diagnostic). Full detail and the test matrix: `docs/EXHALE_BC_and_IC.{tex,pdf}`.

20. Tidal (Roche) potential with an extended outer boundary (2026-06-16)

Motivation. The two domain modes were coupled to two different gravitational potentials. Domain mode: Spherical extends the grid to a user Outer radius but *drops* the stellar-tidal and centrifugal terms (pure $-b_0/r$); the default Roche mode keeps the full tidal potential but truncates the domain at the inner Lagrange point L_1 . Neither matched the escape models of Yan et al. (2022) / Huang et al. (2023), which retain the tidal force *and* integrate out to $\sim 10R_p$ (well past L_1). For WASP-52b, for example, $L_1 \simeq 2.53R_p$, so Roche mode truncated the domain far short of the $10R_p$ used by Yan, cutting off the outer He 10830 / H α absorbing layers.

Change. The outer-boundary extent is now decoupled from the choice of potential. In Roche mode (no Domain mode line, so the full tidal potential of grav_field.f90 is retained), supplying an explicit Outer radius [R_p]: > 1 sets r_max to that value instead of the L_1 radius:

```
else ! Roche mode: full tidal potential retained
  r_max = (3.0*Mrapp)**(-1.0/3.0)*atilde ! L1 (default)
  if (r_out_user .gt. 1.0d0) r_max = r_out_user ! NEW: extend past L1
endif
```

This reproduces the Yan/Huang “tidal force + extended domain” setup with a two-line change and no new keyword (the existing Outer radius is simply honored in Roche mode too).

Caveats and usage. Beyond L_1 the Roche potential is past its maximum (net outward force), so the 1-D radial flow there is a spherical approximation to what is physically an L_1 funnel — the same approximation Yan/Huang make. The cold hydrostatic IC is invalid past L_1 , so pair the extended domain with IC mode: auto (validated: the auto selector finds an interior cold sonic point and seeds a transonic wind). The tidal singularity of the $-b_0 q/(\tilde{a}-r)$ term sits at the star ($r = \tilde{a} \gg 10$), so a $10R_p$ domain is numerically safe. Domain mode: Spherical (no tidal) remains available for Huang Case-A-like comparisons.

Touched files. input_read.f90 (honor r_out_user in the Roche branch of the r_max assignment). Documented in docs/EXHALE_user_manual.tex (§ potential and the input-options table).

21. Photoionization cross-section benchmarks and a TOPbase high-energy extension of He 2³S (2026-07-01)

Motivation. The He 2³S (metastable triplet) photoionization cross section sigma_HeI23S — the drain term that sets the He 10830 lower-level population — was a four-segment broken power-law fit to Norcross (1971), hard-cut to zero above 59.2 eV. Two questions: (i) are EXHALE’s H/He/He 2³S photoionization cross sections consistent with modern atomic data, and (ii) does the unphysical 59 eV cutoff matter? Both are settled in the new memo docs/photoion_cross_sections.{tex,pdf} (reproducible via docs/compare_photoion_cross_sections.py).

Benchmarks. H I and He II use the exact hydrogenic form; He I uses the two-term ATES fit, which agrees with the Verner+1996 (VFKY96) single-shell fit to within 14% over 24.6–500 eV. The He 2³S broken PL reproduces its parent Norcross/p-winds table to ~10% and — checked against the **TOPbase / Opacity Project** R-matrix cross section (extracted from the CDS server; raw 922-point table saved to data/topbase_HeI_2_3S_photo.dat) — agrees to ~8% in the smooth near-threshold region (5–20 eV). TOPbase additionally resolves a 40–55 eV autoionizing-resonance forest (peaks to 3223 Mb); their **net effect on the 2³S rate is only** ~1–2% once the few numerically under-resolved spikes are excluded (the raw table must *not* be fed into a broadband rate integral — a quadrature artifact). Extending the cross section above 59 eV raises the 2³S rate by only 0.1–1.2%.

Change. sigma_HeI23S is extended with a minimal, fully-connected edit (src/modules/functions/cross_sec.f90): the broken PL is kept unchanged up to the 45.6 eV bump; the **last power-law segment is re-aimed** so it runs continuously from the bump peak to a power-law fit of the TOPbase smooth background at the junction $E_J = 70$ eV,

$$\sigma_{\text{tail}}(E) = 10^{5.4623} E^{-2.9964} \text{ Mb} \quad (E > 70 \text{ eV}),$$

(the last-segment slope changes only $-3.039 \rightarrow -2.658$; the intercept is derived in-code for exact continuity). The old 59 eV cutoff is dropped, so 59–1240 eV now carries the TOPbase tail instead of zero. The result is C^0 -continuous at *both* joins (45.6 eV and 70 eV); the original pre-modification formula is kept in comments. A unified full-range PCHIP node table and a tail-only power law are documented as alternatives, but this is the recommended drop-in (smallest rate change, retains the validated threshold-to-bump shape exactly).

Validation. A controlled before/after on WASP-52b (fxuv0p25_he98_L1, Include He23S? True, ϵ Eri SED with X-rays), same IC, only cross_sec.f90 differing: the bulk wind is unchanged ($\log \dot{M} = 11.49$ both), and the He 10830 line changes by $< 0.001\%$ in every metric (line-center depth $15.2675\% \rightarrow 15.2675\%$; EW $0.0877 \rightarrow 0.0877 \text{ \AA}$) — below the Newton convergence noise. The extension is more physically correct yet observationally harmless for He 10830, as predicted by the +0.65% rate estimate (even smaller here because the ϵ Eri SED is soft).

Recombination cross-check. A companion memo docs/recombination_coefficients.{tex,pdf} (compare_recombination.p documents the H/He/He 2³S recombination: H II, He II, He III use the Hui & Gnedin (1997) case-B fits ($\approx$ Draine 2011 / Verner & Ferland to ~2%); He II→He I splits into Benjamin+1999 singlet/triplet (Oklopčić & Hirata 2018) when the 2³S network is on. The He 2³S triplet rate is identical to the latest escape models (Benjamin+1999). Newer nebular-community data (Porter+2012/2013, Del Zanna & Storey 2022) differ ~5–15% but are not yet adopted by escape codes; H/He case-B needs no update.

Touched files. src/modules/functions/cross_sec.f90 (sigma_HeI23S: re-aimed last segment + TOPbase tail, 59 eV cutoff removed, original kept in comments). New docs: photoion_cross_sections.{tex,pdf}, compare_photoion_cross_sections.py, data/topbase_HeI_2_3S_photo.dat, recombination_coefficients.{tex,pdf}, compare_recombination.py, and figures under docs/figures/.

22. He I (1^1S) photoionization $\rightarrow$ Verner+1996 by default, with ATES-rate switches (2026-07-01)

Motivation. The He I *ground-state* (1^1S , singlet) photoionization used the legacy ATES two-term fit, which runs $\sim 12\text{--}14\%$ below the Verner+1996 (VFKY96) single-shell fit in the $50\text{--}100\text{eV}$ band (`photoion_cross_sections.tex`, §4). We make Verner the default, keep the old behavior behind a switch, and add a parallel switch for the recombination data.

Singlet/triplet separation. He photoionization in EXHALE is already split by spin: `sigma_HeI` is the 1^1S **ground (singlet)** cross section (updated here to Verner — Verner’s He I *is* the 1^1S ground state), while `sigma_HeI23S` is the 2^3S **metastable (triplet)** cross section (Norcross/TOPbase, §21). The 2^1S singlet metastable is not tracked (only 2^3S , the 10830 lower level), so the Verner update correctly touches only the singlet ground state.

Change. A single backward-compatible switch in `input.inp`: `ATES_photoionization_rate: True` reverts He I (1^1S) photoionization to the legacy two-term fit (default = Verner+1996, via `sigma_VFKY96` with He I parameters $[E_{\text{th}}, E_0, \sigma_0, y_a, P, y_w, y_0, y_1] = 24.59, 13.61, 949.2, 1.469, 3.188, 2.039, 0.4434, 2.136$). No recombination switch is needed (see below).

Comparison. WASP-52b (`fxuv0p25_he98_L1`, He23S active, same converged IC). Switching He I photoionization legacy $\rightarrow$ Verner changes the He ionization balance and hence the metastable population (Table 1): unlike the $> 59\text{eV}$ tail (§21, $< 0.001\%$), the ground-state photoionization choice is a **real $\sim 3\%$ effect on He 10830** — it matters for the absolute depth and the inferred He abundance. `ATES_photoionization_rate: True` reproduces the old result exactly.

Table 1: He I photoionization legacy (ATES two-term) vs Verner+1996 (new default), WASP-52b He 10830.

metric	legacy (ATES)	Verner (new)	change
$\log \dot{M}$ [g/s]	11.49	11.48	-2%
He 10830 line-center depth	14.516 %	14.933 %	$+2.9\%$
He 10830 equivalent width	$0.0877 \text{ \AA}$	$0.0910 \text{ \AA}$	$+3.8\%$

Recombination — no update needed (EXHALE already current). The companion update (He recombination $\rightarrow$ Porter+2012/2013 / Del Zanna & Storey 2022) was investigated and found to be a *no-op*: those are *emissivity* papers — they update the He I line cascade / line ratios, *not* the total (effective) recombination to 2^3S , and provide no drop-in $\alpha(2^3\text{S})/\alpha(1^1\text{S})$. The escape-relevant quantity — total recombination into the triplet ladder, all funnelling to the metastable 2^3S — is set by (case-B total) $\times$ (triplet spin fraction $3/4$) and is robust across modern treatments: EXHALE (Benjamin+1999) has $\alpha_3 = 2.10 \times 10^{-13} T_4^{-0.778}$, $\alpha_1 = 1.54 \times 10^{-13} T_4^{-0.486}$, while the 2025 thermosphere paper (arXiv:2509.14499) uses $\alpha_3 = (3/4)(2.72 \times 10^{-13} T_4^{-0.789}) = 2.04 \times 10^{-13} T_4^{-0.789}$ and the *identical* singlet $\alpha_1 = 1.54 \times 10^{-13} T_4^{-0.486}$ — so α_3 agrees to $\sim 2\text{--}3\%$ and α_1 is identical. EXHALE’s He recombination already matches the latest practice; no change is warranted. (A trial direct sum of Cloudy’s `he_iso_recomb.dat` level-by-level radiative

recombination gave a total $\sim 2\times$ the case-A value owing to its unmapped 1642-level indexing, confirming the raw file is not a usable shortcut.) No recombination switch was added.

Touched files. `src/modules/init/parameters.f90` (ates_photoion_rate switch), `src/modules/files_IO/input_read.f90` (parse the optional key), `src/modules/functions/cross_sec.f90` (sigma_HeI branches Verner default / legacy two-term). Verified: Verner is the default (He I σ at 30/50/100 eV = 5.36/2.02/0.394 Mb); ATES_photoionization_rate: True recovers the legacy values (5.24/1.78/0.347 Mb).

23. He 2³S photoionization → two VFKEY96 (Verner) forms + power-law bridge (2026-07-01)

Motivation. `sigma_HeI23S` was a bespoke five-segment broken power law (Norcross-1971 fit + TOPbase-tail extension, §21). Since EXHALE already uses the Verner+1996 (VFKEY96) `sigma_VFKEY96` routine for He I and all metals, the metastable triplet is now expressed in the same form, so every photoionization cross section shares one analytic kernel.

Change. (`src/modules/functions/cross_sec.f90`.) `sigma_HeI23S` is rebuilt from two VFKEY96 single-shell forms joined by a log-linear (power-law) bridge, used piecewise, with $[E_{\text{th}}, E_0, \sigma_0, y_a, P, y_w, y_0, y_1]$:

- **wing A** ($E \leq x_3 \simeq 34.70$ eV, threshold → Cooper minimum): 4.78, 2.645, 20.8, 1.0×10^{12} , 3.42, 2.681, 1.956, 2.603;
- **bridge** ($34.70 < E < 45.59$ eV): straight log–log line joining wing A(x_3) to wing B(x_4) (the “power-law part in the middle”);
- **wing B** ($E \geq x_4 \simeq 45.59$ eV, bump → tail): 45.59, 49.68, 1052, 0.04393, 2.941, 1.717, 5.488×10^{-5} , 1.118.

Each wing is a $\sim 4\%$ fit to the resonance-averaged cross section; the construction is **continuous at both** x_3 **and** x_4 , and threshold ($E < 4.78$ eV) and the smooth high- E tail ($\sim E^{-7/2}$) are supplied by the wings (no hard cutoff). The former broken power law is retained, commented out, in the routine.

Verification. Rebuilt EXHALE.x; `sigma_HeI23S` reproduces the former broken power law to $\sim 4\%$ (12 eV: 1.466 vs 1.465; 30: 0.289; 45.6: 2.57 vs 2.68 — bump 4% low; 70: 0.855 vs 0.858; 100: 0.294; 300: 0.0110), continuous across the bridge (34.70 eV: 0.222 → 38 : 0.502 → 40 : 0.795 → 45.59 : 2.570 → 46 : 2.521 Mb).

Touched files. `src/modules/functions/cross_sec.f90` (`sigma_HeI23S`: two VFKEY96 wings + power-law bridge; former broken PL kept commented out).

24. Temperature-dependent Penning ionization of He 2³S (2026-07-01)

Motivation. The metastable He(2³S) destruction by neutral hydrogen, He(2³S)+H, was carried as a *temperature-independent* constant $Q_{31} = 5 \times 10^{-10} \text{ cm}^3 \text{ s}^{-1}$ — the ATES/Oklopčić & Hirata (2018) value, itself the Roberge & Dalgarno

(1982) *sum* of Penning and associative ionization, as also used by Lampón et al. (2020). Taylor et al. (2025, ApJ 989:68) replaced this with a Maxwell–Boltzmann-averaged, **temperature-dependent** fit to the cross sections of Morgner & Niehaus (1979) and Cohen & Lane (1971). Because Penning ionization is often the *dominant* 2^3S loss near the base (where neutral H is abundant), this directly affects the He 10830 prediction EXHALE is built to compute.

New default rate (Taylor et al. 2025, Table 2). A two-branch power law ($\text{cm}^3 \text{s}^{-1}$, T in K):

$$Q_{31}(T) = \begin{cases} 1.9 \times 10^{-9} (300/T)^{0.07}, & T \leq 4000 \text{ K}, \\ 9.1 \times 10^{-9} (300/T)^{0.50}, & T > 4000 \text{ K}. \end{cases} \quad (4)$$

This is $\sim 3\times$ the legacy constant across the relevant 2000–9000 K range. Note the new rate is the Penning channel only ($\text{He}(2^3\text{S})+\text{H} \rightarrow \text{He}(1^1\text{S})+\text{H}^++\text{e}^-$); the associative-ionization channel ($\text{He}(2^3\text{S})+\text{H} \rightarrow \text{HeH}^++\text{e}^-$) is not represented separately in EXHALE (no HeH^+), so the single Q_{31} continues to carry the dominant $2^3\text{S}+\text{H}$ destruction term $-n_{\text{HeITR}} n_{\text{HI}} Q_{31}$.

Unconditional default. The temperature-dependent rate is used *unconditionally*: the legacy constant 5×10^{-10} is no longer selectable. (The `ATES_photoionization_rate` flag continues to revert only the He I (1^1S) photoionization cross section to the legacy ATES 2-term fit; it does *not* affect the Penning rate.)

Implementation. Q_{31} became an array over the cells (it was a scalar). A new `penning_HeI_23S(T,coeff)` in `Cool_coeff.f90` evaluates the piecewise fit over the whole grid; `HeITR_coeffs` calls it unconditionally. The two callers (`ionization_equilibrium.f90`, `post_process_adv.f90`) pass $Q_{31}(j)$ per cell into `params`; the `System_HeH_TR*` residuals are unchanged (they already read Q_{31} from `params`).

Verification. Rebuilt EXHALE.x; ran HD 189733b (metals + He23S, step-capped) with the new rate and, for comparison, with the former 5×10^{-10} constant. The temperature dependence is active and physically sensible: relative to the old constant, $\text{He}(2^3\text{S})$ is reduced near the base (ratio ≈ 0.28 at $1.0R_p$, 0.92 at $1.1R_p$; column-integrated over $1\text{--}2R_p$ drops to $0.59\times$), and the profiles converge to unity above $\sim 1.3R_p$ where H is ionized and Penning switches off — the expected neutral-H-gated signature.

Touched files. `src/modules/radiation/Cool_coeff.f90` (new `penning_HeI_23S`); `src/modules/radiation/util_ion_eq.f90` (`HeITR_coeffs`: Q_{31} array + default/legacy branch); `src/modules/radiation/ionization_equilibrium.f90` and `src/modules/post_process/post_process_adv.f90` (Q_{31} array, `params` assignment in each cell).

25. He/H diffusive separation — Phase 1 (2026-07-01, experimental, default OFF)

Motivation. EXHALE is single-fluid and re-solves composition by *local* ionization equilibrium (conserving the element ratio), so He/H was frozen at the input HeH at all radii — it could not represent the diffusive He/H separation of Taylor et al. (2025) (8%→2.5%) and Xing et al. (2023). Top recommendation of `docs/methodology_aiolos_taylor_xing`.

Method. New module `species_diffusion.f90` (deleted 2026-08-25, section 69; replaced by `src/modules/functions/binary_el`) transports the He element ratio $f = n_{\text{He}}/n_{\text{H}}$ with bulk advection *and* a molecular-diffusion drift of He relative to H, $\Phi = -D n_{\text{H}}(\partial f / \partial r + f G)$, $G = (m_{\text{He}} - m_{\text{H}})g / (kT)$, $D = 1.52 \times 10^{18} (1/m_{\text{H}} + 1/m_{\text{He}})^{1/2} T^{1/2} / n$ (Banks & Kockarts 1973), plus eddy diffusion K_{zz} on the gradient term. Solved as a **fully conservative** implicit (backward-Euler tridiagonal) update of n_{He} ; the new split is written back into `f_sp` conserving mass ($\sum_s m_s f_{\text{sp},s} = 1$, metals frozen to H). Gated on `He_diffusion` (default off; `He_Kzz` sets K_{zz}). Derivation: `docs/design_hehe_diffusion.md`.

Root-cause bug (design doc §7). A local time scale named `t0` ($= R_0/v_0$) shadowed the *global temperature* `T0` (Fortran is case-insensitive) in `TK = Tcode*T0`, zeroing `TK` (floored to 1 K) and inflating the settling coefficient $G \propto 1/TK$ by $\sim 5800\times$, causing total over-settling. The two constants are numerically near-equal ($R_0/v_0 \approx 2.83 \times 10^4$ s, $T_0 \approx 2.83 \times 10^4$ K), which hid it. Fix: rename the local to `tscale`.

Result (HD 209458b, He23S+metals). Physical mild separation — $(\text{He}/\text{H})/\text{HeH} \approx 0.6\text{--}0.9$ in the inner thermosphere, falling to $\sim 0.17\text{--}0.24$ aloft — consistent in magnitude with Taylor ($\rightarrow 0.3\times$) and Xing. Stable (no NaN); `He_diffusion OFF` is byte-identical.

Phase-1 limitations. A near-base pile-up from the lagged- n_{H} nonlinearity is held by a physical limiter $f \leq \text{HeH}$ (aloft observable unaffected); proper fix is a self-consistent n_{H} inner iteration. Ambipolar field, thermal diffusion, metal diffusion, and a Newton-finished quantitative $\dot{M}$ are Phase-2. Flag stays default OFF.

Touched files. new `src/modules/functions/species_diffusion.f90`; `parameters.f90` (`he_diffusion`, `he_kzz`); `input_read.f90` (`He_diffusion`, `He_Kzz`); `EXHALE_main.f90`; `Makefile`.

26. He/H diffusion — Phase 2 (2026-07-02): ambipolar, thermal diffusion, metals

Extends §25 (design doc §7e). The Phase-1 kernel was factored into a shared subroutine `solve_1elem` (one conservative implicit advection–diffusion–settling solve for any element vs the H background), reused by He and each metal.

- **P2b — ambipolar-corrected settling** (default ON, `He_ambipolar`). The polarization field lifts ions, so the He-vs-H effective settling mass is $\Delta m_{\text{eff}} = 3 - 0.5(\bar{Z}_{\text{He}} - \bar{Z}_{\text{H}})$ (mean charges from the local ionization state): 3 at the neutral base, 2.5 in the fully ionized wind. Modest effect ($\sim 2\%$ less depletion aloft), correct direction.
- **P2c — thermal diffusion** (default `He_alphaT=0`). Adds $\alpha_T \partial \ln T / \partial r$ to the settling coefficient; off by default.
- **P2d — element-by-element metal diffusion** (default OFF, `He_metal_diffusion`). Each trace metal element diffuses independently vs n_{H} with its own mass (`melem_A`), binary D , and ambipolar correction; ion stages rescaled to the diffused total. Metals trace $\Rightarrow$ no n_{H} feedback. **Validated (HD 209458b C/N/O):** heavier elements deplete more — at $3 R_{\text{p}}$, `element/base` $\approx$ He 0.171, C 0.158, N 0.136, O 0.105 (monotonic in mass; matches Xing et al. 2023). Regression: OFF $\Rightarrow$ He/H unchanged, C/H frozen at 1.0.

- **P2a — self-consistent n_{H} : attempted, NOT adopted.** A Picard $n_{\text{He}} \leftrightarrow n_{\text{H}}$ iteration diverges (near-base settling genuinely concentrates He, driving $n_{\text{H}} \rightarrow 0$); the physical cap $f \leq \text{HeH}$ remains the correct limiter.

Touched files. `species_diffusion.f90` (solve_1elem kernel; ambipolar dmeff; thermal term; metal-diffusion loop); `parameters.f90` (he_ambipolar, he_alphaT, he_metal_diffusion); `input_read.f90` (He_ambipolar, He_alphaT, He_metal_diffusion).

27. Diffusive separation: detailed physics and numerics

This section consolidates §25–§26 into a single, more complete treatment of the He/H (and metal) diffusion, with emphasis on the two design decisions that are easy to get wrong: the $n_{\text{He}}-n_{\text{H}}$ mass coupling (§27.3) and the settling discretization (§27.5).

27.1 Governing equation

By default EXHALE is single-fluid: the hydro evolves $(\rho, \rho v, E)$ with one bulk velocity, and the composition is re-solved every step by *local* ionization equilibrium, which conserves the element ratios. Consequently the helium-to-hydrogen ratio $f \equiv n_{\text{He}}/n_{\text{H}}$ is frozen at the input value at all radii — EXHALE cannot represent diffusive separation without an added transport equation.

We transport the He element ratio f with advection at the bulk velocity plus a molecular-diffusion *settling* drift of He relative to H,

$$\frac{\partial f}{\partial t} + v \frac{\partial f}{\partial r} = \frac{1}{n_{\text{H}}} \frac{1}{r^2} \frac{\partial}{\partial r} \left[r^2 n_{\text{H}} (D + K_{\text{zz}}) \frac{\partial f}{\partial r} + r^2 n_{\text{H}} D f G \right], \quad (5)$$

with the diffusive number flux (relative to the bulk)

$$\Phi_{\text{He}} = -n_{\text{H}} \left[(D + K_{\text{zz}}) \frac{\partial f}{\partial r} + D f G \right], \quad G = \frac{\Delta m g}{kT} + \alpha_T \frac{\partial \ln T}{\partial r}. \quad (6)$$

Here $D = 1.52 \times 10^{18} (1/m_{\text{H}} + 1/m_{\text{He}})^{1/2} T^{1/2} / n$ is the Banks & Kockarts (1973) binary He-in-H diffusion coefficient ($\text{cm}^2 \text{s}^{-1}$), $g = GM_{\text{p}}/r^2$, and $\Delta m = m_{\text{He}} - m_{\text{H}}$ ($= 3$ amu in the neutral limit; §27.6). K_{zz} is the eddy diffusion coefficient (He_Kzz, default $10^9 \text{ cm}^2 \text{s}^{-1}$): it enters only the *gradient* term and drives f toward uniform (mixing), whereas the *settling* term $D f G$ is molecular only. Their competition sets the homopause: below it ($K_{\text{zz}} \gg D$) the atmosphere is mixed, above it ($D \gg K_{\text{zz}}$) helium settles. In diffusive equilibrium ($\Phi_{\text{He}} = 0$, no eddy) $\partial \ln f / \partial r = -G$, i.e. f decreases outward with the (small) He scale height $1/G$.

27.2 Conservative implicit discretization

An early *non-conservative* form (material-derivative advection $-v \partial f / \partial r$ combined with a conservative diffusive flux) proved flux-inconsistent and *drained* helium spuriously. The production scheme is instead **fully conservative**: it evolves the He number density n_{He} with a single divergence of the total face flux

$$J_{j+1/2} = \underbrace{n_{\text{He}} v}_{\text{advection}} + \Phi_{\text{He}}, \quad \frac{n_{\text{He}}^{n+1} - n_{\text{He}}^n}{\Delta t} = -\frac{1}{r^2 \Delta r} \left[r_{j+1/2}^2 J_{j+1/2} - r_{j-1/2}^2 J_{j-1/2} \right]. \quad (7)$$

Because the cold, dense base has a tiny He scale height, the diffusion is stiff and an explicit scheme is unstable; the update is therefore **implicit** (backward Euler), assembled as a tridiagonal system in n_{He}^{n+1} and solved by the Thomas algorithm. Implicit backward Euler is unconditionally stable, and its fixed point ($n_{\text{He}}^{n+1} = n_{\text{He}}^n$) is independent of Δt , so local cell-by-cell time steps are used for steady-state acceleration. The new element split is written back into `f_sp` conserving mass ($\sum_s m_s f_{\text{sp},s} = 1$), and the same kernel (`solve_1elem`) is reused for each trace metal (§27.6).

27.3 The $n_{\text{He}}-n_{\text{H}}$ mass coupling: why a self-consistent n_{H} is *not* needed

Diffusion redistributes species at *fixed* total mass density ρ (owned by the hydro), so in every cell

$$\rho = m_{\text{H}}n_{\text{H}} + m_{\text{He}}n_{\text{He}} + (\text{metals}) \implies n_{\text{H}} = \frac{\rho(1 - c_{\text{met}}) - m_{\text{He}}n_{\text{He}}}{m_{\text{H}}}. \quad (8)$$

Thus n_{H} is itself a function of n_{He} : where helium accumulates, hydrogen must be displaced. But the flux (6) *uses* n_{H} (in the coefficient and in $f = n_{\text{He}}/n_{\text{H}}$), so the solve is genuinely nonlinear in n_{He} . Two options:

- **Lagged n_{H} (adopted).** Hold n_{H} at its start-of-step value during the solve; update it afterward, in the mass-conserving rescale, for the next step.
- **Self-consistent n_{H} (“P2a”, tried, rejected).** Iterate within the step (Picard): solve n_{He} , reset n_{H} from (8), re-solve, until n_{He} and n_{H} are mutually consistent.

Why the Picard iteration diverges. Write the He mass fraction $Y = m_{\text{He}}n_{\text{He}}/\rho \in [0, 1]$. Near the base, settling genuinely concentrates helium ($n_{\text{He}} \uparrow$); enforcing (8) exactly then drives

$$n_{\text{H}} = \frac{\rho(1 - c_{\text{met}}) - m_{\text{He}}n_{\text{He}}}{m_{\text{H}}} \xrightarrow{Y \rightarrow 1} 0,$$

so $f = n_{\text{He}}/n_{\text{H}}$ and the $1/n_{\text{H}}$ flux coefficients diverge — a genuine singularity as helium approaches a pure-He layer. Because the concentration is *real* (settling against a fixed reservoir base), the iteration has a positive feedback and blows up (NaN). The lagged scheme avoids this: with n_{H} held at a finite value, $1/n_{\text{H}}$ stays bounded, and the physical cap $f \leq \text{HeH}$ (§27.4) prevents $Y \rightarrow 1$.

Why lagging is nonetheless harmless. Four independent arguments:

1. *The lag vanishes at the fixed point.* The diffusion runs inside a relaxation loop. At steady state $n_{\text{He}}^{n+1} = n_{\text{He}}^n$ and $n_{\text{H}}^{n+1} = n_{\text{H}}^n$, so the lagged $n_{\text{H}} = n_{\text{H}}^n$ used in the solve *equals* the final consistent n_{H}^{n+1} . The lag affects only the transient path, not the converged state: where the scheme is stable, lagged and self-consistent reach the *same* steady solution, and mass consistency is achieved automatically.
2. *The region where it matters is the (unphysical) pile-up.* Lagged and self-consistent differ appreciably only where $Y \rightarrow 1$, a thin layer just above the fixed reservoir base — below the He 10830 line-forming region, hence irrelevant to the observable, and in any case capped (§27.4).
3. *Helium is a minor mass fraction where the observable lives.* In the upper thermosphere, where He/H falls to $\sim 0.2-0.5 \times$ (the part that weakens 10830), $m_{\text{He}}n_{\text{He}} \ll m_{\text{H}}n_{\text{H}}$, so $n_{\text{H}} \approx \rho/m_{\text{H}}$ is nearly insensitive to n_{He} and the lag is negligible.

4. *Trace metals have no back-reaction.* Metals are $\sim 10^{-3}$ of the mass, so their diffusion does not change n_{H} measurably; for them the lagged n_{H} is effectively exact.

Hence the divergent Picard iteration is not merely inconvenient but *unnecessary*: the lagged- n_{H} solve plus the physical cap is stable and converges to the correct steady state.

27.4 The near-base pile-up and the physical cap

Settling transports helium down toward the reservoir base. Against a fixed-composition Dirichlet base it tends to pile up in the first cells above the base (numerically unbounded with a lagged n_{H}). In a settling-plus-escape column, however, He/H *cannot* exceed the reservoir value: there is no mechanism to enrich helium above the base abundance. The code therefore applies the physical limiter

$$f \leq \text{HeH} \quad (9)$$

(and, for each metal, $f_X \leq f_{X,\text{base}}$). This encodes the expectation that the reservoir *absorbs* settled helium rather than letting it build a pure-He layer — so the cap is a physical boundary condition, not merely a numerical band-aid. The aloft separation (the observable) sits well below HeH and is unaffected by the cap.

27.5 Settling discretization: a Péclet hybrid

The settling term DfG is an advective drift (velocity $w_s = DG$, directed inward). Central differencing of an advective term is stable and accurate only at low Péclet number, $|DG| \Delta r < 2(D + K_{zz})$. For light helium this always holds (small G), and central differencing is used (least numerical diffusion). For *heavy metals* (e.g. Fe, mass 56) G is $\sim 18\times$ larger; central differencing then loses diagonal dominance and the tridiagonal solve NaNs. The kernel therefore uses a **Péclet-based central/upwind hybrid**:

$$\text{settling drift: } \begin{cases} \text{central} & |DG| \Delta r < 2(D + K_{zz}) \quad (\text{light species}), \\ \text{first-order upwind} & \text{otherwise} \quad (\text{stiff heavy-metal settling}). \end{cases} \quad (10)$$

Upwind keeps the matrix diagonally dominant for any settling strength; the metal-stage rescale additionally skips cells where the element is negligible and clamps the scale factor to avoid $0 \times \infty$. These changes were required to run WASP-121b (Fe) stably; on HD 209458b (light C/N/O), the hybrid selects central and reproduces the He-only result.

27.6 Ambipolar field, thermal diffusion, and metals

In the ionized wind the polarization field $eE = -(1/n_e)dp_e/dr$ (Koskinen et al. 2013) lifts ions; for an H^+ background it supports $\sim$ half a proton weight per unit charge, so a species of charge Z has effective mass $m - Zm_{\text{H}}/2$. The He-vs-H relative settling mass becomes

$$\Delta m_{\text{eff}} = 3 - \frac{1}{2}(\bar{Z}_{\text{He}} - \bar{Z}_{\text{H}}), \quad (11)$$

with $\bar{Z}$ the mean charges from the local ionization state — $\Delta m_{\text{eff}} = 3$ at the neutral base, 2.5 in the fully ionized $\text{He}^{++}/\text{H}^+$ wind (He_ambipolar, default on). Thermal diffusion adds $\alpha_T \partial \ln T / \partial r$ to G (He_alphaT, default 0). Each trace metal (He_metal_diffusion, default off) is diffused independently against the post-rescale n_{H} with its own mass

(`melem_A`), binary D_X , and ambipolar-corrected Δm_{eff} ; being heavier than helium, metals deplete faster with altitude where diffusion competes with the wind.

27.7 The metal-rescale ratchet bug (found in review, fixed 2026-07-02)

The first version of the metal-stage rescale contained two bugs that together made metal depletion *irreversible*: (i) an $r_X \leq 1$ clamp — redundant for the metal/H $\leq$ reservoir cap (already enforced by $\min(n_X, f_{X,\text{base}} n_{\text{H}})$) but forbidding any *replenishment* of a previously depleted cell (a one-way ratchet); and (ii) a skip of cells where the element was negligible, which made exhausted cells permanent holes. During early relaxation (wind undeveloped) settling transiently depletes metals; the ratchet locked that in. **Fix:** r_X is no longer clamped above (the cap alone bounds it), and an exhausted cell that the solve replenishes is re-seeded through the neutral stage (the ionization equilibrium re-partitions the stages next step). **Retraction:** the previously reported WASP-121b “metal homopause” ($\text{Fe} \rightarrow 0$ by $\sim 1.25 R_p$, with a ~ 2500 K hotter thermosphere) was an artifact of this ratchet — with the fix, WASP-121b metals track He (even Fe is advection-dominated there, $w_s/v \sim 10^{-3}$) and the v1/v2 temperature structures agree. The HD 209458b mass ordering ($\text{He} > \text{C} > \text{N} > \text{O}$ aloft) is unaffected.

27.8 Boundary conditions and validation

The lower boundary is Dirichlet at the reservoir composition ($f = \text{HeH}$); the outer boundary carries a zero diffusive flux plus advective outflow (helium escapes with the wind). The two-planet before/after comparison (`docs/version_compare.{md,tex,pdf}`) shows the expected regimes: on HD 209458b helium itself separates ($\text{He } 10830 \text{ } 70.5\% \rightarrow 27.1\%$, $-2.6\times$) and heavier C/N/O deplete more than He aloft, whereas on WASP-121b ($\log \dot{M} \approx 13.3$) advection dominates for every species — helium *and* metals are dragged out essentially unfractiated (10830 unchanged; metal/H tracks He/H). Both are stable (no NaN) and leave $\dot{M}$ essentially unchanged (metals are trace). All diffusion flags are **default off**, so standard runs are byte-identical.

28. Lower-atmosphere connection: Tier 1–3 (2026-07-02)

Survey + proposal + implementation: `docs/lower_atmosphere_coupling.{md,tex,pdf}`. New folder `src/modules/lower_atmosphere` everything **opt-in / default off** (regression: the standard HD 209458 b 3000-step run reproduces v1.0’s $\log \dot{M} = 8.97$ exactly).

- **Tier 1** `lower_column.f90` + key `Lower column::` the Koskinen+2022 analytic isothermal- T_{eq} hypsometric column with Visscher chemical-equilibrium $\text{H}_2/\text{H}/\text{He}$; reports the derived 1- μbar base radius (equilibrium + fully-atomic bracket), base $q_{\text{H}_2}/q_{\text{H}}/q_{\text{He}}$ and μ next to the input “Planet radius”. Gate passed: Model A $r_0/R_{1\text{bar}} = 1.343$ (paper 1.34), $q_{\text{H}_2} = 0.838$ (0.84), $q_{\text{H}} = 0.027$ (0.026).
- **Tier 2 foundation** `h3p_cooling.f90` (Miller+2013 Table-5 LTE emission fits + Table-6 non-LTE factors $s(T, n_{\text{H}_2})$; anchors reproduced to $<0.5\%$) and `mol_rates.f90` (Koskinen+2022 Table-1 rates R1–R23, verified against the PDF). Remaining: the coupled `System_HeH_mol` solver + EOS.

- **Tier 2a** key Molecular base: True: EOS-only base correction removing the H₂-bound particles from `ntot_bc` via the equilibrium fit (chemistry stays atomic; crude, documented).
- **Tier 3** optional `base.inp` (`T_base/r_base/HeH_base/Kzz_base`; echoed, no-op when absent) consumed by `input_read`, + the driver `src/utls/run_lower.py` (isothermal or Guillot 2010 semi-grey $T(p)$); VULCAN photochemistry is the documented upgrade path.

Physics note from the Tier-1 gate: the chemical-equilibrium base stays strongly molecular up to $T \sim 2000$ K, so the atomic base of $T_{\text{eq}} \sim 1000\text{--}2000$ K hot Jupiters rests on *photochemical* dissociation (Moses 2011; Koskinen 2013a) — quantified motivation for the VULCAN tier.

29. Tier-2 molecular chemistry core: H₂/H₂⁺/H₃⁺/HeH⁺ (2026-07-03)

Opt-in key Molecular chemistry: True (requires He; v1 excludes trace metals). New coupled equilibrium system `System_HeH_mol.f90` (7–8 unknowns; atomic rows = EXHALE’s own rates, so the molecule-free limit reproduces the atomic systems; molecular channels = Koskinen+2022 Table-1 via `mol_rates`). σ_{H_2} (Yan+1998 Eqs. 17–19; verified against their Table 7) added to `cross_sec.f90` and wired into the opacity/photoionization/heating integrals; H₃⁺ IR cooling (Miller+2013 + non-LTE factor) enters `cool`; the EOS is molecule-aware (particle count, electrons, masses). `f_sp` grows 33→37 (output schema + Load-IC extended; columns zero when off). **Gates:** (0) molecules-off regression $\log \dot{M}=8.97$ unchanged; (1) HD 209458 b molecules-on: fully molecular base, sharp H₂ → H front at $r=1.019 R_p$, wind above it essentially atomic ($\log \dot{M}$ 8.91 vs 8.97, He 2³S peak +1.4%); (2) hot-Uranus-like: front at $1.148 R_p$, H₃⁺ active below. Caveats: local equilibrium only (no molecular advection), no Lyman–Werner photodissociation, P4/P5 and the 4.48 eV sink omitted, $\gamma=5/3$ retained. Details: `docs/lower_atmosphere_coupling.pdf` §4.

30. Newton–diffusion co-convergence, solver hygiene, and the Ly α lower-BC audit (2026-07-03)

Stall-based Newton hand-off. `He_diffusion` runs plateau the flux metric just above the hand-off threshold (observed: du frozen at 1.084×10^{-2} vs the 10^{-2} switch for 10^6 steps). The warm-up now tracks the du plateau and hands over to JFNK once stalled within $5 \times$ the switch. *Validated:* the hand-off engaged at $du = 2.2 \times 10^{-2}$; the JFNK then hit the known base-momentum holdout (worst cell $j=1$, $\|R\| \sim 7.7 \times 10^{-3}$ while the volume-weighted reference residual was 2.2×10^{-4}) and fell back gracefully — the base-cell viscosity work remains the enabler. **Newton + He_diffusion co-convergence.** Solver: Newton previously froze the diffused He/H field at the hand-off state (the JFNK residual carries no operator-split diffusion). An outer iteration (the excited-H pattern) now alternates the JFNK solve with 500 diffusion-relaxation steps at the converged wind until the He/H field drift $< 10^{-3}$ (max 5 passes, drift reported). **Riemann floor:** `p_min` floored at $10^{-30} p_{\text{max}}$ in `speed_estimate_HLLC/ROE` (no-op for healthy states). **PH_heat_HHe OMP CRITICAL removed** (distinct- j writes are race-free; values unchanged). **regression/ run_fcheck.sh:** periodic `-fcheck` regression per the 2026-07-02 review. **lya_rt audit** (vs. Huang et al. 2017): the escape-probability closure $\bar{J} = S(1 - \beta)$ is local and lacks Huang’s *absorbing bottom* (H₂ accidental-resonance sink at $N_{\text{H}_2} \sim 10^{14} \text{ cm}^{-2}$),

so $\bar{J}$ — and the $H(n=2)$ population — is likely overestimated in the bottom few scale heights; the $H\alpha$ -forming region lies above, so the TPM impact is expected limited. Documented; a bottom-loss β channel is the follow-up.

31. Tier-3 first light: VULCAN $\rightarrow$ base.inp (2026-07-03)

The public VULCAN photochemical kinetics code was cloned (`./VULCAN`; FastChem compiled) and run to steady state for HD 189733 b (SNCHO-2025 network, `use_photo=True`; 2356 steps, ~ 6 h wall). The new converter `src/utils/vulcan_to_base.py` turns the `.vul` output into an EXHALE `base.inp`: photochemical $q_{H_2}/q_H/q_{He}$ at $1\ \mu\text{bar}$, a hypsometric base radius using VULCAN's own $\mu(p)$ and $T(p)$, and the molecular mixing ratios as comments (*no* metal release — metal/alkali chemistry is outside VULCAN's scope and remains Lavvas-only). **Result** (HD 189733 b, $1\ \mu\text{bar}$): $q_{H_2} = 0.63$, $q_H = 0.23$ — photochemistry dissociates $\sim 11\times$ more H than the chemical-equilibrium column ($q_H = 0.020$), directly quantifying the Tier-1 caveat (Moses 2011; Koskinen 2013a); $r_{\text{base}} = 1.168$ vs the analytic $1.174 R_J$; $T_{\text{base}} = 863\ \text{K}$ (the Moses11 $T(p)$, cooler than the isothermal $T_{\text{eq}} = 1183\ \text{K}$). End-to-end chain demonstrated: VULCAN $\rightarrow$ converter $\rightarrow$ base.inp $\rightarrow$ EXHALE startup override (echoed).

32. VULCAN lower-atmosphere pre-step: obtained, not included in the tree (2026-07-06)

VULCAN (with FastChem inside it, as shipped) is used as a lower-atmosphere pre-step invoked by EXHALE itself. VULCAN is a third-party code and is *not* committed to the EXHALE repository (VULCAN/ is git-ignored, avoiding the embedded- repository problem); it is fetched into EXHALE/VULCAN/ by the new `src/utils/setup_vulcan.sh` (clone + two patches + FastChem build), which the driver also runs automatically on first use if the tree is missing. The key

```
Lower atmosphere: vulcan <R_1bar>
```

in `input.inp` runs the new driver `src/utils/vulcan_driver.py` at startup whenever no `base.inp` exists: it copies the VULCAN tree into `<run_dir>/vulcan_work/`, builds the planet's Guillot $T(p)/K_{zz}$ atmosphere and picks a stellar UV spectrum by host T_{eff} , compiles FastChem once, runs VULCAN to steady state (hours on the first run, cached `.vul` afterwards; `--force` to redo), converts the result to `base.inp`, and then proceeds with the wind solve — i.e. VULCAN acts as a subroutine of EXHALE. `Lower atmosphere: analytic <R_1bar>` invokes the fast chemical-equilibrium column (`run_lower.py`) instead; omitting the key disables the pre-step (fully optional). EXHALE_ROOT overrides the code-root path for relocated installs, and an existing `base.inp` always wins. Implementation: `parameters.f90` (`lower_atm_mode/lower_atm_r1bar`), `input_read.f90` (`run_lower_atm_prestep` via `execute_command_line` before `read_base_inp`), and the added `vulcan_driver.py`. Tested: analytic branch end-to-end, VULCAN branch with a cached `.vul` (HD 189733 b photochemical base $T = 863\ \text{K}$ applied), and the opt-out (no key) path. The download URLs, citation requirements (VULCAN: Tsai et al. 2017, 2021; FastChem: Stock et al. 2018, 2022) and the exact list of modifications `setup_vulcan.sh` applies are documented in the README “Obtaining VULCAN and FastChem” section and the user manual (§“Obtaining VULCAN and FastChem”).

33. TPM.py renamed to EXHALE_transit.py (2026-07-06)

The transmission-spectrum post-processor TPM.py (the acronym clashed with “Trusted Platform Module” and did not describe the tool) is renamed EXHALE_transit.py, matching the EXHALE_* convention. Run-time overrides are now EXHALE_TRANSIT_<NAME> (PATH, SAVE_PREFIX, HE_LMIN/LMAX/N, RSTAR_RSUN, ROTP, TSTAR); the old TPM_<NAME> names remain as backward-compatible fallbacks (a `_tenv()` helper checks the new name first, then the old). Output filenames are unchanged (<prefix>tpm_He10830.txt etc.) so downstream notebooks and figure scripts keep working. All in-repo references were updated; VULCAN/ untouched.

34. Sub-level-resolved Balmer opacity and LaRT 2s inclusion (2026-07-06)

The $H\alpha/H\beta$ transmission was computed as $f_{23}n_2$ with the multiplet oscillator strength ($f_{23} = 0.6407$ for $H\alpha$), which implicitly assumes the $2s:2p$ populations sit in the 1:3 statistical ratio. They do not—the metastable $2s$ ($A_{2s1s} = 8.26\text{ s}^{-1}$) and the $\text{Ly}\alpha$ -pumped $2p$ depart strongly from it—and the two sub-levels have *different* Balmer cross sections (only $2s \rightarrow 3p$ is allowed from $2s$; $2p \rightarrow 3s, 3d$ from $2p$). Fixed to sum per sub-level, $\tau \sim f_{2s}n_{2s} + f_{2p}n_{2p}$, with $f_{2s} = 0.4349$, $f_{2p} = 0.70941$ ($H\alpha$) and 0.1028, 0.125886 ($H\beta$); the components share the Doppler-dominated profile.

EXHALE_transit.py (escape-prob path) now carries n_{2s}/n_{2p} separately; the effect is small ($\lesssim$ few %) since the in-code $J_{\text{Ly}\alpha}$ keeps $2s:2p$ near statistical in the forming region. examples/tpm_halpha_lart2d.py (LaRT path) previously used the $2p$ -only population, *discarding the 2s absorption entirely*; it now returns and sums both (the LaRT scattering rate pumps only $1s \rightarrow 2p$, entering the S_{2p} source term, while $2s$ is recombination/collisionally fed throughout the atmosphere). Effect: +21% (HD 189733 b) to +67% (WASP-121 b). The $\text{Ly}\alpha$ absorption line itself is unchanged and already correct— $1s \rightarrow 2p$ only ($f = 0.41641$ on the ground-state n_{HI} ; $1s \rightarrow 2s$ is a forbidden two-photon transition).

benchmarks/hd209/input.inp was missing Stellar Lya flux, Stellar Teff, and Stellar radius; without the $\text{Ly}\alpha$ flux the LaRT converter read $F_{\text{Ly}\alpha} = 0 \Rightarrow L_{\text{Ly}\alpha} = 0 \Rightarrow n_{2p} \simeq 0$, a spurious $H\alpha = 0.000\%$. Added $T_{\text{eff}} = 6065\text{ K}$, $R_* = 1.155 R_{\odot}$, and $F_{\text{Ly}\alpha} = 3.5 \times 10^3\text{ erg cm}^{-2}\text{ s}^{-1}$ (quiet-solar-analog, R_*^2 scaled; provisional, since the absolute LaRT depth scales with it). With both paths on the same VULCAN-BC winds and the sub-level opacity, the escape-prob/LaRT $H\alpha$ ratios become HD 209458 b 0.84/0.43 (2.0), HD 189733 b 1.23/0.91 (1.4), WASP-121 b 4.41/2.33 (1.9), WASP-52 b 5.69/2.20 (2.6)—so the inline closure overestimates $H\alpha$ by only $\sim 1.4\text{--}2.6\times$, far less than the earlier $\sim 4\text{--}30\times$ (pre-VULCAN winds with the $2p$ -only LaRT). Paper abstract, methods, results table, summary, and figs/make_lart_fig.py were updated accordingly.

35. Molecular equilibrium seeding fix, molecular-base coupling, Tier-2 gate refresh (2026-07-16)

Frozen-atomic-base bug (diagnosed and fixed). The 8-unknown molecular equilibrium (ion_system_HeH_mol, hybrd1) is bistable in its initial guess: from a zero-molecular warm start it fails to converge at the dense, optically-thick base, and the `thereis_mol` branch never checked info—the failed state was silently adopted and then frozen indefinitely ($x_{\text{H}_2} = 0.0122$ at the base, flat across 130k steps and independent of T ; numerically *not* a root of the

coded network—it would require an H_2 photoionization rate $15\times$ the unattenuated stellar bound). The true dark-base equilibrium is strongly molecular (3-body formation R15 vs. thermal dissociation R12), which the earlier gate runs had reached only because the then-uninitialized molecular `f_sp` columns acted as an accidental nonzero seed; the composition-audit zero-init exposed the latent fragility. Fixes: (i) `info` is checked and a failed solve retries once from a physically-informed guess seeded by the chemical-equilibrium H_2 fit `q_h2_equilibrium(p,T)` (still-failing cells keep the previous state and are counted in a one-line summary warning); (ii) `set_IC` seeds the molecular H_2 column from the same fit (H nuclei conserved) instead of zero.

Molecular base coupled on. Solving molecular chemistry now implies the molecular base particle count (Molecular base: True behavior): with an atomic `ntot_bc` the base temperature is inflated by the particle-count ratio ($\sim 1.8\times$ for a fully molecular base), thermally dissociating the H_2 layer the chemistry just built. A note is printed when the coupling engages.

Tier-2 gates re-established (deterministic; inputs pinned at `docs/lower_atmosphere_figs/data_g*/input.inp`). Gate 0 (HD 209458 b atomic, 3000 steps): $\log \dot{M} = 8.97$ unchanged. Gate 1 (HD 209458 b molecular): fully molecular base ($x_{\text{H}_2} = 0.997$, base T unchanged at ~ 1464 K), sharp $\text{H}_2 \rightarrow \text{H}$ front at $r = 1.019 R_p$, $\log \dot{M} = 8.87$ vs. 8.97 atomic, $\text{He } 2^3\text{S}$ peak $+1.3\%$; stable over 12000 steps with zero solve failures. Gate 2 (hot-Uranus-like: $0.0457 M_J$, $R_p = 0.49 R_J$, $T_{\text{eq}} = 1140$ K, HD 209458 b orbit/spectrum—input recovered by matching the archived domain fingerprint $r_{\text{max}} = 4.814118$ to 6 decimals): fully molecular base, front at $r = 1.166 R_p$, H_3^+ active in the molecular layer. Atomic runs byte-identical throughout (make check + OMP identity).

Gate refresh (2026-07-23, current production defaults). After the 2026-07-17 onward changes (staged secondary ionization, `He_rec_coupling` on, `He-H` charge exchange on; §§36–39) the three gates were re-recorded at a uniform 12000-step relaxation-snapshot convention. The $\text{H}_2 \rightarrow \text{H}$ fronts and molecular base composition are essentially unchanged (front $1.019 \rightarrow 1.020$ for HD 209458 b, $1.166 \rightarrow 1.156$ for the hot Uranus; base $x_{\text{H}_2} = 0.996/0.9998$), so the dissociation-front physics is robust to the new defaults. The wind $\log \dot{M}$ now reads 10.58 for both HD 209458 b runs at 12000 steps — these are relaxation snapshots, not flux-flat converged: under the new defaults the HD 209458 b atomic gate does not converge (du reaches a fixed point $\sim 2.6 \times 10^{-6}$ while the $\rho v r^2$ mass-flux spread plateaus near 4%), so all three gates share the same step cap. $\text{He } 2^3\text{S}$ is now strongly suppressed in the molecular base below the front ($\text{He } 2^3\text{S} + \text{H}_2$ Penning destruction, $\sim 10^4\times$ at the front, rising back to the atomic value above $\sim 1.3 R_p$); the earlier “ $+1.3\%$ peak” was a whole-domain-peak artifact of the outer-wind transient. Pinned `data_g*/` refreshed; `lower_atmosphere_coupling` and `fig_mol_structure.pdf` updated.

36. Default H/He rates $\rightarrow$ Badnell+Mao / Voronov, van Hoof free-free Gaunt factor, free-free charge fix, and secondary ionization (2026-07-17)

Six coupled changes to the H/He microphysics. Items 1–2 are gated by the new `legacy_hhe_rates` flag (default False = new rates); items 3–4 (free-free) apply always; item 5 (secondary ionization) is gated by `use_sec_ion` (default True); item 6 rides along with the heating rework. Full cross-code context and the benchmark table: `docs/atomic_data_EXHALE_vs_MoCHII`.

36.1 What changed and why

1. H/He recombination default. The pre-update H/He case-B coefficients were unattributed fits inherited from ATES; a numerical check showed them to be exactly the Hui & Gnedin (1997) case-B fits. The new default is

$$\alpha^B(T) = \alpha_{\text{Badnell}}^A(T) - \alpha_1^{\text{Mao}}(T), \quad (12)$$

the Badnell (2023) radiative total minus the Mao & Kaastra (2016, A&A 587, A84) $n = 1$ coefficient, with the Badnell three-term dielectronic recombination added to α^A for He II. This matches the construction MoCHII already uses, so the two codes now share the H/He recombination coefficients. Setting `legacy_hhe_rates = True` restores the Hui & Gnedin fits exactly.

2. H/He collisional ionization default. The legacy H I / He I rates were Abel et al. (1997) (reproducing Janev et al. 1987) and He II a Hui & Gnedin (1997) fit, none source-annotated. The Voronov (1997, ADNDT 65, 1) fits,

$$k(T) = A \frac{(1 + P\sqrt{U})U^K e^{-U}}{X + U}, \quad U = \Delta E[\text{eV}]/k_B T, \quad (13)$$

are now the default, matching the metal rates (already Voronov in both codes) and MoCHII. Reverted together with item 1 by `legacy_hhe_rates`.

3. Free-free Gaunt factor. The frequency-averaged non-relativistic $\langle g_{\text{ff}} \rangle(\gamma^2)$ (with $\gamma^2 = Z_{\text{ion}}^2 \text{Ry}/k_B T$) of van Hoof et al. (2014, MNRAS 444, 420) is tabulated at 161 points in $\log_{10} \gamma^2$ from -6 to 10 in 0.1 dex steps (generator `cooling_data/gauntff_thermal_avg.py`), replacing the earlier unattributed two-branch logarithmic fit. Applies always, independent of `legacy_hhe_rates`.

4. Free-free charge-weighting correction (physical bug). The previous code weighted He II (net charge $+1$) by $Z^2 = 4$ and evaluated the metal Gaunt factor at the nuclear atomic number ($Z = 8$ for O, 26 for Fe). Free-free emission scales with the ion *net* charge, not the nuclear number: $Z_{\text{ion}} = 1$ for H II, He II, and singly-ionized metals; $Z_{\text{ion}} = 2$ for He III and doubly-ionized metals. The most visible consequence is that the free-free cooling of a He II-dominated layer drops by a factor of four. Applies always.

5. Secondary ionization (Shull & van Steenberg 1985) in the main loop. A photoelectron ejected with $E_0 = e_\nu - E_{\text{th}}$ above ~ 40 eV does not thermalize immediately: it deposits only $f_{\text{heat}}(x)$ of its excess as heat and drives secondary H I and He I ionizations, where x is the ionized fraction of the H+He nuclei,

$$f_{\text{heat}}(x) = 0.9971 (1 - (1 - x^{0.2663})^{1.3163}), \quad (14)$$

$$f_{\text{ion,HI}}(x) = 0.3908 (1 - x^{0.4092})^{1.7592}, \quad (15)$$

$$f_{\text{ion,HeI}}(x) = 0.0554 (1 - x^{0.4614})^{1.6660}. \quad (16)$$

Below 40 eV the photoelectron thermalizes fully. Gated by `use_sec_ion` (default `True`); `use_sec_ion = False` is bit-identical to the legacy full-thermalization path. This differs from `wind_ae`, which combines the same SvS85 partition with Dere (2007) tables tied to its fixed spectral grid; the main loop runs on a runtime SED and applies the SvS85 species-resolved fits directly.

6. He I 2^3S triplet photoheating channel added. The metastable He I 2^3S (threshold 4.8 eV) previously contributed only opacity and its photoionization rate but no photoheating. Its photoelectron energy now enters the heating and absorbed-energy budgets and the secondary-ionization source consistently with every other absorber.

36.2 Files touched

Cool_coeff.f90 (rr_badnell, dr_HeII_badnell, rr_mao, alphaB*_new, voronov_ci/ci*_new, the 161-point gff_avg table); util_ion_eq.f90 (free-free net-charge weighting in eval_cool; secondary ionization and He 2^3S heating in PH_heat_H/PH_heat_HHe); T_equation.f90 (free-free net-charge weighting); ionization_equilibrium.f90 and post_process_adv.f90 (callers wired for the new signatures); input_read.f90, parameters.f90, write_setup_report.f90 (legacy_hhe_rates and use_sec_ion declared, defaulted, parsed, echoed); cooling_data/gauntff_thermal_avg.py (van Hoof table generator); docs/input_schema.md (keys K14b, K14c).

36.3 Quantitative check

New default rates (α_B , collisional ionization ci ; $\text{cm}^3 \text{s}^{-1}$), with the legacy value at 10^4 K for comparison. The legacy $\alpha_B(\text{HII})$ at 10^4 K is the Hui & Gnedin (1997) fit, 2.592×10^{-13} .

Coefficient	10^4 K	10^5 K	10^6 K	legacy (10^4 K)
$\alpha_B(\text{HII})$	2.606e-13	3.070e-14	2.163e-15	2.592e-13
$\alpha_B(\text{He II})$	2.807e-13	5.006e-13	1.015e-12	2.616e-13
$\alpha_B(\text{He III})$	1.536e-12	2.370e-13	2.244e-14	—
$ci(\text{HI})$	7.448e-16	3.963e-9	3.103e-8	7.247e-16

The Fortran reproduces an independent Python recomputation to all printed digits, and the legacy branch reproduces the pre-update values exactly. The new $\alpha_B(\text{HII})$ at 10^4 K differs from the legacy Hui & Gnedin value by $\sim 1\%$; He II and He III differ more, mainly because the new He II value carries the Badnell dielectronic contribution explicitly. The van Hoof thermal Gaunt factor at $T = 10^4 \text{ K}$ and $Z = 1$ is $\langle g_{\text{ff}} \rangle = 1.266$. A smoke run confirms the expected factor-of-four drop in free-free cooling across a He II-dominated layer, and `use_sec_ion = False` reproduces the legacy heating and photoionization rates bit-identically.

36.4 Effect on existing results

Because the default rates change, results computed with the new default differ from the pre-update golden and gate outputs. `legacy_hhe_rates = True` recovers the old H/He ionization balance, with two always-on exceptions: the free-free charge-weighting correction (item 4) and the van Hoof Gaunt table (item 3), both physical corrections intentionally not tied to the switch. Secondary ionization (item 5) is a further default change; `use_sec_ion = False` restores the full-thermalization rates bit-identically.

37. He recombination radiation → H I ionization coupling (use_he_rec_coupling, 2026-07-17)

EXHALE carries no diffuse radiation field, so pure case B implicitly assumes every He II → He I recombination photon is reabsorbed locally by helium ($\alpha_B = \alpha_A - \alpha_1$, ground capture dropped). In H-dominated gas that is wrong: part of the He recombination radiation — the ≥ 24.6 eV ground-capture continuum, the 584 Å resonance line, the 19.8 eV 2^3S line, and the >13.6 eV part of the He I two-photon continuum — ionizes H instead. The physics and the choice of recipe were worked out in docs/QUESTIONS_2026-07-17.md. This section documents an opt-in coupling using Draine's (2011) on-the-spot y/z parametrization. New input key `He_rec_coupling` (flag `use_he_rec_coupling`), default `False` → pure case B, bit-identical to the legacy path.

37.1 What changed and why

When on, the He recombination photons that ionize H are added as an extra H I photoionization rate with its photoelectron heating, and the He II recombination coefficient becomes $\alpha_B + y\alpha_1$ rather than pure case B. Two local parameters set the split:

- $y = 1/(1 + R n_{\text{HeI}}/n_{\text{HI}})$ (Draine Eq. 14.16), the fraction of the ≥ 24.6 eV ground-capture continuum absorbed by H, with the cross-section ratio $R = \sigma_{\text{He}}/\sigma_{\text{H}}$ at the 24.6 eV He I ground edge. Evaluated from the code cross sections, $R = 6.004$.
- $z = 0.67 + 0.29/(1 + n_e/n_{\text{crit}})$ (Draine Sec. 15.5), the density-dependent fraction of case-B (excited-state) cascade photons that ionize H, with the 2^3S critical density $n_{\text{crit}} = 1100 e^{1.2/T_4} T_4^{1/2} \text{ cm}^{-3}$; it interpolates between Draine's two limits (~ 0.96 low density → ~ 0.67 high density).

Atomic (case-B) mode. No explicit 2^3S : α_1 is the Mao & Kaastra (2016) ground ($1s^2$) capture and α_B the active He II case B. The He II recombination coefficient is $\alpha_{\text{eff}} = \alpha_B + y\alpha_1$; the extra H I photoionization rate is $n_{\text{HeII}} n_e [z\alpha_B + y\alpha_1]/n_{\text{HI}}$; the photoelectron heating uses the ground-channel energy $E_{\text{gnd}} = 24.6 - 13.6 = 11.0$ eV and a cascade-averaged $E_{\text{casc}} \approx 6.14$ eV.

He I 2^3S (triplet) mode. The mode that computes 10830: the decay channels are explicit, so z is not used — a 2^3S atom destroyed by photoionization or Penning emits no 19.8 eV photon, so the cascade must be summed at the actual channel rates (α_1 is the 1^1S channel). The 1^1S channel coefficient is $y\alpha_1 + 0.25\alpha_B$: the net ground capture plus the singlet-excited capture channel ($0.25\alpha_B$) the current network omits. The H-ionizing photon production is summed over channels: ground $y\alpha_1 n_{\text{HeII}} n_e$; singlet-excited $0.8521 \times 0.25\alpha_B n_{\text{HeII}} n_e$ ($0.8521 = \frac{2}{3}$ from the 584 Å line + $\frac{1}{3} \times 0.5564$ from the 2^1S two-photon continuum above 13.6 eV); 2^3S radiative decay $A_{31} n(2^3S)$ (19.8 eV line, always ionizes H); and 2^3S collisionally converted to the singlets $n_e n(2^3S)(0.5564 q_{31a} + q_{31b})$. The channel deposit energies are 11.0 (ground), 5.53 (singlet-excited), 6.2 (19.8 eV line), and 2.512/7.6 eV (collisionally converted two-photon / 584 Å). The two-photon constants are integrals of the Drake, Victor & Dalgarno (1969) spectral shape rather than round numbers. The 2^1S term emits a photon *pair* summing to 20.62 eV with a distribution peaking at half that, so above the H I

edge it delivers 0.5564 ionizing photons per decay carrying 8.9646 eV — a mean photon energy of 16.110 eV, hence 2.512 eV of photoelectron energy each. Treating the in-band photons as uniformly distributed instead puts the mean at 17.109 eV and the photoelectron at 3.511 eV, 40% high. `f_2q_HeI` and `Ee_2q_HeI` in `util_ion_eq.f90` carry these, and the 0.8521/5.53/6.14 combinations are formed from them so the three cannot drift apart.

The corrections are applied at the lagged (pre-solve) densities in the main ionization-equilibrium loop and re-evaluated at the advection-corrected densities in the post-processor (rate/coefficient before the advection solve, heating after). The He II recombination *cooling* is left at $kT \alpha_B$; the mismatch is at most $y \alpha_1 kT$, negligible.

37.2 Files touched

`util_ion_eq.f90` (new `he_rec_coupling` subroutine returning the corrected He II recombination coefficient, the extra HI photoionization rate, and the photoelectron heating; all zero when off); `Cool_coeff.f90` (new `alpha1_HeII_mao`, the Mao & Kaastra 2016 ground $1s^2$ capture, i.e. the α_1 subtracted in `alphaB_HeII_new`); `ionization_equilibrium.f90` and `post_process_adv.f90` (callers wired to `rcheiiB/P_HI/heat` in the main loop and to the advection ODE and `theat` in post-process); `input_read.f90`, `parameters.f90`, `write_setup_report.f90` (`use_he_rec_coupling` declared, defaulted `.false.`, parsed, echoed); `docs/input_schema.md` (key K14d).

37.3 Quantitative check (HD 209458 b)

Protocol: HD 209458 b `input.inp+metals.inp`, the converged profiles loaded as IC, then four runs {atomic, He 2^3S } $\times$ {off, on} each relaxed for 15000 steps from the same IC. The IC predates the 2026-07-17 rate update, so both branches share the same partial re-relaxation and **only the on/off difference is meaningful** (not the absolute values). Across the wind $y \sim 0.65\text{--}0.69$ and $R = \sigma_{\text{He}}/\sigma_{\text{H}}(24.6\text{eV}) = 6.004$. Script `docs/plot_he_rec_coupling.py`; figures `docs/figures/he_rec_coupling_{atomic,heitr}.pdf`.

Quantity (on/off)	Atomic mode	He I 2^3S mode
He I	$\times 1.04$ ($1.2 R_p$) $\rightarrow \times 1.25$ plateau ($1.5\text{--}3.5 R_p$) $\rightarrow \times 1.32$ ($3 R_p$)	$\times 1.01\text{--}1.16$
He II (min.)	$\times 0.72\text{--}0.74$ ($1.1\text{--}1.2 R_p$) $\rightarrow \sim 1.0$ ($2 R_p$)	$\times 0.885$
$n(2^3S)$	— (not computed)	$\times 0.89$ (1.05), $\times 0.92$ (1.2), $\times 0.98\text{--}1.01$ ($1.5\text{--}2 R_p$)
Temperature	$\leq +1\%$ (outer +3%)	$\leq +0.3\%$
$\log \dot{M}$	$9.58 \rightarrow 9.65$ (+0.07 dex, $\sim +17\%$)	$9.57 \rightarrow 9.61$ (+0.04 dex, $\sim +10\%$)

The atomic mode carries the larger He ionization-structure change (case B alone, $y = 0$, under-recombines He), and its extra H photoionization raises the mass-loss rate by $\sim 17\%$. In the 2^3S mode the 10830-forming region ($1.5\text{--}2 R_p$) sees $n(2^3S)$ within $\sim 1\text{--}2\%$ of the off case, so the He I 10830 line is essentially unchanged, while the base 2^3S drops $\sim 11\%$ (the TR-mode 1^1S channel coefficient change feeds directly into n_{HeII}). The temperature response is small in both modes. Default off is bit-identical to the pre-update path; existing goldens and gates are unaffected.




Figure 1: He recombination radiation $\rightarrow$ HI coupling, **atomic mode** (HD 209458 b). Clockwise from top left: He ionization fractions off vs. on; the on/off ratio for He I and He II; the photoheating rate; and the temperature. Turning the coupling on raises He I and lowers the He II minimum, with the $\sim 17\%$ mass-loss increase reflected in the extra photoheating.




Figure 2: Same comparison in the **He I 2^3S mode**. The extra panel and ratio curve track the metastable 2^3S density: it drops $\sim 11\%$ at the base but stays within $\sim 1\text{--}2\%$ of the off case through the 10830-forming region ($1.5\text{--}2 R_p$), so the He I 10830 line is essentially unchanged.

38. Staged activation of secondary ionization (2026-07-23)

38.1 The regression

With secondary ionization on by default (the `use_sec_ion` item of the 2026-07-17 rate update), the high-gravity WASP-121 b cases (backup/regression/wasp_full, wasp_he23off) crashed with a NaN at $\sim$ step 587: negative densities first, then a NaN base pressure and temperature. The coupling itself is physically correct in magnitude — at the shielded base the SvS85 secondary channel raises the H photoionization rate by $\sim 35\times$ over the direct rate, the expected size for the $\sim$ keV photoelectrons there. The failure is a transient-path problem, not a steady-state one: the extra base ionization deepens the base cooling valley and amplifies the base acoustic mode of the marginal cold initial condition until the hydro diverges. Two facts pin this down: (a) a run with the coupling off converges cleanly ($du < 10^{-3}$, $\log_{10} \dot{M} = 13.30$); and (b) restarting with the coupling on from that converged state is benign — no NaN, du stays $\lesssim 4 \times 10^{-2}$ during a ~ 2000 -step adjustment and re-converges. The steady state with the full physics exists and is stable; only the cold-start path to it is fatal.

38.2 The fix: converge first without the coupling, then switch it on

A runtime state flag `sec_ion_active` now controls whether the SvS85 partition is applied, separately from the `use_sec_ion` enable flag. The photoionization / photoheating routines apply the coupling only when `use_sec_ion` .and. `sec_ion_active`; when it is off the code path is bit-identical to the `use_sec_ion = False` (full-thermalization) path. The main loop keeps `sec_ion_active = False` while the wind first relaxes; the moment that first relaxation reaches a convergence/stall criterion, the flag flips to True, the convergence logicals and the mass-flux level-stability history are reset, and the run continues and re-converges with the full physics. Stops are additionally held for `N_stall` (default 2000) steps after the flip: du reacts only once the base adjustment wave driven by the new coupling has formed, so without the hold the run stopped one step after activation (activation at step 7144, spurious stop at 7145) with a state that had not yet incorporated the physics. The Newton (Solver: Newton) finish and the direct-solve diagnostic modes (which bypass the marching loop) switch the coupling on before their solve, and the flag is forced on after the loop so post-processing and final outputs always carry the physics even if the loop hit the iteration cap before the flip.

38.3 The Immediate override

Secondary_ionization: Immediate sets `use_sec_ion = True` and applies the coupling from step 0, restoring the pre-staging behavior. It is intended for A/B tests only; on the high-gravity cold-start cases it reproduces the $\sim$ step-587 NaN crash, confirming that the staged path is what removes it.

38.4 Files touched

parameters.f90 (`sec_ion_active`, `sec_ion_immediate`); util_ion_eq.f90 (`PH_heat_H` / `PH_heat_HHe` gate on the local `sec_on = use_sec_ion` .and. `sec_ion_active`); EXHALE_main.f90 (init, stage flip on first convergence, Newton-finish and diagnostic-mode activation, post-loop force-on); input_read.f90 (the Immediate value parsed).

38.5 Effect on existing results (wasp_full validation)

For Secondary_ionization: False the run is bit-identical to before (verified byte-identical on wasp_full and wasp_he23off). For the default (staged on) wasp_full: activation at step 7144, re-convergence at step 13489 ($du = 9.98 \times 10^{-4}$), and $\log_{10} \dot{M}$ moves $13.30 \rightarrow 13.22$. The $\sim -17\%$ mass-loss rate is the real secondary-ionization physics: the SvS85 partition diverts most of the hard-photon photoelectron energy from heat into ionization and excitation in the weakly ionized base, so less energy drives the wind. Immediate reproduces the old (crashing) behavior on the high-gravity cases.

39. He-triplet physics series from the Falorca & Vidotto (2026) review (2026-07-23)

The line-by-line comparison against Falorca & Vidotto (2026, arXiv:2607.18193) and García Muñoz (2025, A&A 698, A199) — full analysis in docs/falorca2026_he3d_review.md — exposed six physics gaps, all implemented and validated in this series (its validation work also surfaced the secondary-ionization NaN regression fixed in the previous section). Physical correctness is the acceptance criterion throughout; the impact sizes below are informational.

- Penning ionization products** (heh_tr_rows, System_HeH_mol, adv_implicit_HeH_TR): $\text{He}(2^3\text{S}) + \text{H}^0 \rightarrow \text{He}(1^1\text{S}) + \text{H}^+ + e^-$ now produces the H^+ (with its 6.2 eV of electron heating, $e_{\text{th,HeI}} - e_{\text{th,HeTR}} - e_{\text{th,HI}}$) it always removed the triplet for. In the EUV-shielded base (Γ_{HI} exponentially small) this is the *dominant* H-ionization channel: n_{HII} rises $2-7\times$ at $r < 1.03$ on wasp_full, while the wind and $\dot{M}$ are unchanged. The associative branch ($\rightarrow \text{HeH}^+$, $\sim 10\%$, GM25) is folded in as H^+ (documented at the code site).
- Collisional ionization in the TR systems** (heh_tr_rows b terms; ci_HeI23S in Cool_coeff.f90, the Black-1981-derived $6.41 \times 10^{-21} \sqrt{T} e^{-55338/T} / E_{23\text{S}}$, $k(10^4 \text{ K}) = 3.3 \times 10^{-10}$): the TR (Oklopcic-form) rows silently dropped *all* electron-impact ionization while eval_cool kept charging its cooling — both restored and consistent now, including CI of the metastable itself (a $\sim 1-2\%$ 2^3S loss at 10^4 K , its 4.8 eV cooling in the coio channel, and its He I / adv-row counterparts).
- Explicit- $n(2^3\text{S})$ cooling** (eval_cool + the semi-implicit energy solver): the $10830 \text{ \AA}$ collisional-excitation channel $1.16 \times 10^{-20} \sqrt{T} e^{-13179/T} n_e n_{23\text{S}}$ (Black form with the *explicit* metastable density — Falorca & Vidotto warn Black's implicit steady-state triplet is wrong here) plus the 0.80/1.40 eV q_{31a}/q_{31b} conversion ledger. The q_{13} (19.8 eV) term is deliberately excluded: the Cen-1992 He I excitation-cooling term already carries that channel. The terms are passed into solve_energy_semi_implicit (all three eval_cool calls) so they act on T ; without that wiring they were diagnostic-only. Measured: 1–5% of the total cooling on WASP-121 b (8–11% of local $\text{Ly}\alpha$ at the 2^3S peak), locally up to $\sim 28\%$ on HD 209458 b.
- He_rec_coupling default ON** (the previous section has the model): with the coupling off the TR singlet recombination used α_1 alone — neither case A nor case B — and the He recombination photons ionized/heated nothing. Validated stable on wasp_full: base $n_{23\text{S}} -10\%$, the 10830-forming region within $\sim 2\%$, $\dot{M}$ unchanged; on HD 209458 b $\dot{M}$ rises $+0.04-0.07$ dex. He_rec_coupling: False restores the legacy path (byte-identical).

5. **He $\leftrightarrow$ H charge exchange in the default set** (he_h_charge_exchange, default on; Koskinen 2013 rates via Huang 2023 Table 4): the group-B pair moved out of the cx_full-gated table assembly into dedicated he_h_cx_fvec/jac/fvec_adv routines called by *every* ionization system with He (previously only the metals systems could reach it), including the advection pair and the analytic Jacobians of the Newton systems. The effect is confined to the neutral-H base — He II $\times 0.59$ at $r = 1.05$, -1% by $r \geq 1.2$, $\dot{M}$ unchanged — exactly the $\text{He}^+ + \text{H}^0 \rightarrow \text{He}^0 + \text{H}^+$ sink where n_{H^0} dominates. cx_full still gates groups C/D. The new input keys of this and the preceding sections were also added to the parser’s known-keys list (they parsed correctly but tripped the unknown-line warning).
6. **He(2^3S)+H₂ Penning ionization** (penning_HeI23S_H2; System_HeH_mol rows 4/5/8 + 4.4 eV heating): GM25 Table A.5 (Cohen & Lane 1977) fitted as $5.3791 \times 10^{-12} T^{0.676} e^{-695.21/T}$ ($\leq 0.13\%$ over 500–10⁴ K). On the pinned molecular gates the metastable drops $39\times$ (HD 209 base) and $430\times$ (hot-Uranus, out to $r \sim 1.1$) exactly where $x_{\text{H}_2} \rightarrow 1$, with the upper wind and $\dot{M}$ unchanged — GM25’s dominant deep metastable sink, now present for Tier-2 + triplet runs. The _adv post-process is atomic-only (documented) and needs no counterpart.

Regression state after the series: Secondary_ionization: False + He_rec_coupling: False + He_H_charge_exchange: False reproduce the respective legacy paths byte-identically; the wasp goldens and the parse-dump corpus were re-snapshotted once at the end of the series (defaults: staged secondary ionization, He recombination coupling on, He $\leftrightarrow$ H charge exchange on). Converged wasp_full reference: activation at step ~ 7144 , final count ~ 13486 , $\log_{10} \dot{M} = 13.22$.

40. du triggers require a descending crossing (2026-08-11)

The $\text{du} < \text{du}_{\text{th}}$ convergence stop and the JFNK hand-off (with the secondary-ionization flip that shares its du test) now fire only after du has been seen at or above their own threshold. They start armed for a Load IC? True run and unarmed for a generated IC. The failure this removes: a transonic IC reported $\text{du}(1) = 1.85 \times 10^{-5}$ on WASP-121 b, which tripped both triggers on step 1 and ran away to a NaN by step 1220; with the guard the same input converges (step 8706, $\log_{10} \dot{M} = 13.21$). The PLM $\rightarrow$ WENO3 switch is deliberately *not* guarded, because the regression cases legitimately switch on step 1 at $\text{du}(1) = 0.37$. No new input keys. Details: docs/newton_scaling_and_base_wall.md §9.

41. JFNK line search decides on the true residual (2026-08-11)

The line search in solve_steady_jfnk accepted steps on the residual with the WENO3 smoothness weights frozen at the *previous* iterate, compared that against a merit memory which mixed the frozen measure with the true one and was refreshed only on acceptance, and reported both $\|R\|$ and the info = 0 verdict from the frozen residual. Measured on the HD 189733 b hand-off state: 7 of 14 accepted steps *raised* the true residual, by factors 1.09 to 9.05, and the stored reference stood $2.4\times$ below the true merit of the iterate the search started from — which is what produced the 12-consecutive-failure abort.

Trials are now evaluated with the weights recomputed at the trial state, the merit memory holds true merits and is rewritten at the start of every outer iteration, and the diagonal scaling is formed before the merit it scales. The frozen weights still define the inner Newton model, which is where they belong. WASP-121 b now reaches $\text{info} = 0$ ($\|R\| = 5.72 \times 10^{-5}$, $\log_{10} \dot{M} = 13.17$, superseding the du-stop 13.21 of the preceding section); `photo_deep_secion_cont` converges in 24 outer iterations instead of 59; no case exhausts its backtracks any more. No new input keys, no golden changed. Details: `docs/newton_scaling_and_base_wall.md` §10.

42. Metal-line cooling at a cold base: line trapping and the coronal-fit validity floor (2026-08-11)

Two physically wrong statements in the metal cooling were corrected. Both bite only where the gas is cold and dense — at the base of a planet whose base falls out of the coronal regime — and the wind is unaffected.

42.1 (a) The escape probability was hardwired to 1

`eval_cool` built a `beta_esc` from the lowest XUV-band *continuum* opacity across *one cell*, then overwrote it with `beta_esc = 1.0d0` and applied the full optically thin metal-line cooling everywhere. Neither quantity is a line optical depth and neither is grid-independent. (AIOLOS `chemistry.cpp:1006` instead scales the same gray depth by an arbitrary 10^8 , driving $\beta \rightarrow 0$ and switching metal-line cooling off; EXHALE had replaced that with the opposite limit.) Measured on the converged HD 189733 b base, the dominant coolant line [O I] 63 μm carries a line-center optical depth $\tau = 2.7$ through the cold layer, i.e. an escape probability of 0.16, not 1.

β is now the line-center escape probability of the line itself, computed from the *column* between the emitting cell and the top of the domain (`Cool_coeff.f90`: `kappa_OI63`, `kappa_CII158`, `line_escape_probability`, `fine_structure_escape`). The shape is the plane-parallel Doppler result of Hollenbach & McKee (1979) / de Jong, Boland & Dalgarno (1980), renormalized by a factor 2 so that $\beta(0) = 1$ *exactly*: the published form tends to $1/2$ because it counts escape through one face of a slab, whereas here the other direction is absorbed by the lower atmosphere and so removes energy from the modeled gas either way. The two branches are switched at $\tau_c = \sqrt{\pi} e^{a^2/4} = 6.967$, where they cross, so the switch is continuous in value. In the optically thin wind $\beta \rightarrow 1$, reproducing the previous behavior.

β enters as $A_{ul} \rightarrow \beta A_{ul}$ *inside* the two-level solution, not as a factor on its result: in the subcritical limit the cooling is set by the collisional excitation rate and must be independent of β , and multiplying outside would be wrong by a factor β there. Trapping is applied to [O I] 63 μm and [C II] 158 μm only, the two lines for which the code carries an explicit two-level solution; everything else stays optically thin, which is right in the wind and is the residual approximation at the base ($< 0.3\%$ of the base cooling there). A thick resonance line in a metal-rich wind (Mg II h&k) is still treated as thin — recorded, not fixed.

42.2 (b) The CHIANTI coronal fits were evaluated far below their validity floor

Every metal cooling coefficient in `Cool_coeff.f90` is a fit or a table built over 10^3 – 10^5 K. At the 236 K HD 189733 b base, 99.7% of the [O I] rate came from the fit's softest exponential, $e^{-930.111/T}$, whose 930 K corresponds to no [O I]

ground-term splitting at all (the splittings are 227.7 K and 326.6 K); C I is the same case ($e^{-2351.38/T}$ against splittings of 23.6 K and 62.4 K). Below the floor the only excitable metal transitions are the ground-term fine-structure lines, which the explicit two-level terms already carry, saturation included.

coronal_excitation_cutoff(T) now removes the coronal part below 10^3 K as $\exp[-((T_{\text{floor}}/T - 1)/w)^2]$ with $w = 0.5$, and is exactly 1 above it. Value and dT -slope are continuous at the floor (the Brent energy solve and the semi-implicit update both differentiate the cooling in T), and every coefficient is **bit-identical at and above 10^3 K**. The factor is applied to all CHIANTI-derived coefficients — the analytic C/N/O and Mg/Ca/Na/Fe forms and the 1-D/2-D tables, whose $\log_{10} T$ axis starts exactly at 3.0 and which otherwise hold their edge value indefinitely below it — and to the coronal remainder of cool_OI_ne_func / cool_CII_ne_func, but *not* to their two-level parts. The legacy AIOLOS branch (cno_cool_0) is deliberately left alone: its constant floors are crude fine-structure stand-ins, not extrapolated coronal fits. $w = 0.5$ is a modeling choice, not a measurement.

42.3 Measured

HD 189733 b, cell 1: total cooling $3.40 \times 10^{-5} \rightarrow 2.68 \times 10^{-8}$ erg cm $^{-3}$ s $^{-1}$ against a heating of 7.0×10^{-6} , so the base net rate changes sign from -2.65×10^{-5} (cooling) to $+6.97 \times 10^{-6}$ (heating). Almost all of that comes from (b); (a) alone is a factor ~ 5 on the [O I] two-level term. Sweeping the code's own cooling assembly in T at that cell's frozen densities, against the heating measured there, the local radiative balance temperature moves from **168 K to 486 K** — 168.2 K from (a) alone, 485.5 K from (b) alone, so (b) does essentially all of the work at the base — and the two cooling curves are identical from 10^3 K up. A side-by-side marching test does *not* resolve this: on a 5×10^4 -step scale the base is riding an inflow transient whose adiabatic heating dominates the radiative term in both binaries.

WASP-121 b (base at 2358 K, $\tau([\text{O I}] 63 \mu\text{m}) = 0.039$, $\beta = 0.956$) is unchanged: $\log_{10} \dot{M} = 13.17$ before and after, outer mass flux to 3×10^{-7} relative, profile maximum relative difference 8×10^{-5} . make check: mol_base_handoff (metals off) stays byte-identical (PASS); the two metal cases move by the intended physics with $\log_{10} \dot{M} = 13.22$ unchanged in both — wasp_full maximum relative difference 3.7×10^{-7} (density), outer mass flux to 2×10^{-8} ; wasp_he23off 1.6×10^{-3} (density, temperature) in the narrow band at $r = 1.012$ – 1.014 , outer mass flux to 1.2×10^{-5} . Goldens not re-snapshotted at that point. Separately noted while comparing: both goldens already contain **negative species densities** in that band (H I, O I, O II, S II; 17 cells in wasp_full, 3 in wasp_he23off) — a pre-existing defect, unrelated to this change.

42.4 Not fixed by this

The model still has no stellar optical/IR absorption and no thermal background, so nothing sets a radiative floor at T_{eq} for a shielded layer; and below ~ 700 K the guard leaves ions without an explicit two-level term (C I, N, Mg, Ca, Na, Fe) with no cooling at all — for C I that omits the real [C I] 609/370 μm lines, whose LTE rate at the HD 189733 b base is $\sim 10^{-10}$ erg cm $^{-3}$ s $^{-1}$, i.e. 10^{-4} of the local heating. Details: docs/hd189_base_checkerboard.md §10. Files:

```
src/modules/radiation/Cool_coeff.f90
src/modules/radiation/util_ion_eq.f90
src/modules/nonlinear_system_solver/T_equation.f90
```

src/modules/post_process/post_process_adv.f90

43. Base grid resolution as an input key (Base grid [dr,cells], 2026-08-11)

The resolution of the uniform region of the Mixed grid was two hardcoded locals in `define_grid.f90` ($N_{\text{low}} = 50$ cells of $drc = 2.0e-4 R_p$). It is now an input key, because that spacing is what decides whether the base carries the stationary $2\Delta r$ entropy mode diagnosed in `docs/hd189_base_checkerboard.md`: the controlling parameter is the number of cells per base density scale height, $H/\Delta r$ with $H = kT/(\mu g)$, and a $4\times$ refinement was measured to remove the mode.

```
Base grid [dr,cells]:  2.0e-4 50      # the default -- the historical grid
Base grid [dr,cells]:  5.0e-5 200     # 4x refinement at the same 0.01 R_p extent
```

The two numbers share one line, as `du_th [PLM,WEN03]` does, because they are not independent: their product is the radial extent of the uniform region ($0.01 R_p$ by default). Refining at fixed extent means dividing Δr and multiplying the cell count by the same factor; changing only one moves the junction with the stretched region. The second value may be omitted (the count then keeps its default). The key applies to `Grid type: Mixed` only; `Uniform` and `Stretched` build their grids from r_{max} and N alone, and `EXHALE_setup.out` echoes that the key is ignored for them.

`EXHALE_setup.out` also reports the resolution actually achieved, `Base scale-height resolution:  $H(T_{\text{eq}})/dr = 1/(b_0 \cdot dr_j(1))$` (b_0 is the surface Jeans parameter, $dr_j(1)$ the first cell after the Mixed-grid smoothing), and warns below 10 cells. On the default grid this is 26.9 cells for HD 189733 b against 102.3 for WASP-121 b — the margin the high-gravity planet does not have. `Grid type: Stretched` gives 1.7 cells for HD 189733 b, i.e. it is unusable at the base for this class of problem. The smoothing pass at the end of `define_grid` started at the literal 50, which was N_{low} written out; it now follows N_{low} so that the junction between the uniform and stretched regions is smoothed wherever it is.

Two costs, both structural. The CFL step is set by the smallest cell, so a $4\times$ finer base needs about $4\times$ more steps for the same physical time. And the total cell count $N = 500$ is a compile-time parameter, so the cells given to the base come out of the stretched region: on the $4.43 R_p$ HD 189733 b domain the stretch ratio goes $1.0119 \rightarrow 1.0252$ and the cell size at $1.5 R_p$ grows from 6.1×10^{-3} to $1.3 \times 10^{-2} R_p$. Refining the base therefore coarsens the wind, which is where $\dot{M}$ and the transmission spectrum are formed. Making N runtime-settable would need the static $(1-Ng:N+Ng)$ arrays to become allocatable throughout; it was not done here. Recorded while doing this: `count_max = 1000000` is also a compile-time parameter and the `EXHALE_MAXSTEPS` hook only lowers it, so a $4\times$ refined base can reach at most $1/4$ of the physical time the default grid reaches within the cap. On HD 189733 b $4\times$ is therefore not reachable to convergence without raising `count_max`; $2\times$ is, and $2\times$ is already where the mode goes.

Default is bit-exact. `dr_base` is declared with a *default-real* literal ($2.0e-4$, not $2.0d-4$) because that is what the old local was: the stored value is the single-precision neighbor of 2×10^{-4} , 2.5×10^{-8} relative below it. Declaring

the exact double moves every base cell by that amount; measured on `wasp_full` it changed the converged density by 5×10^{-10} relative (step count and $\log_{10} \dot{M} = 13.22$ unchanged), which is not a physics difference but is not byte-identity either. The literal is kept so that an `input.inp` without the key reproduces earlier runs exactly. The consequence to know: writing the default out explicitly (`Base_grid [dr,cells]: 2.0e-4 50`) does *not* reproduce it, because the list-directed read into `real*8` gives the exact double.

Checks. Byte-identity was tested differentially, against the binary built immediately before this change rather than against `backup/regression/golden/` (the goldens are stale by the intended metal-cooling physics change of the preceding section). All output files of `wasp_full`, `wasp_he23off` and `mol_base_handoff` are byte-identical including headers, with the same step counts (13488 / 13482 / 12000) and $\log_{10} \dot{M}$ (13.22 / 13.22 / 10.58). WASP-121 b run to completion with the key absent: $\log_{10} \dot{M} = 13.17$, JFNK info = 0, $\|R\| = 5.7 \times 10^{-5}$ — unchanged from the preceding section. Malformed values are rejected in `define_grid` (a cell count outside $[2, N-10]$, or a uniform region that does not fit inside $r_{\max}$).

Details: `docs/hd189_base_checkerboard.md` §11. Files:

```
src/modules/init/parameters.f90
src/modules/init/define_grid.f90
src/modules/files_IO/input_read.f90
src/modules/files_IO/write_setup_report.f90
```

44. Base ghost temperature, runtime step cap, coronal-cutoff width as input keys (2026-08-11)

Three settings that were compile-time constants or an unexamined default became `input.inp` keys. All three default to the previous behavior, so a file without them runs exactly as before (`make check`: all three golden cases byte-identical).

44.1 Base ghost temperature: isothermal | continuous

The lower ghost cells pin the density to `rho_bc` and, in the legacy closure, the pressure to `ntot_bc + dp_bc`, i.e. the temperature to T_0 . **The T_0 pin has no physical backing:** the radiative equilibrium of the lower atmosphere is outside the model, so nothing in the code determines the ghost temperature, and T_{eq} is a choice rather than a boundary condition. It becomes actively wrong once the first interior cell settles far from T_0 . With the metal-line cooling of the preceding section, cell 1 on HD 189733 b sits at 494 K against $T_0 = 1183$ K: a factor-2.4 contact discontinuity, plus a $2.3 \times$ density inversion, held permanently on the boundary.

continuous imposes $dT/dr = 0$ at the base face instead. The ghost keeps the base composition — `ntot_bc` nuclei and `dp_bc` electrons at `rho_bc`, exactly the particle count the isothermal pin uses — and carries the cell-1 temperature,

$$p_{\text{gh}} = (\text{ntot_bc} + \text{dp_bc}) T_1, \quad T_1 = p_1 / (n_{\text{tot}} + n_e)_1.$$

$(n_{\text{tot}} + n_e)_1$ comes from `get_species_densities`, the single point where the code decides what counts as a particle. The ghost pressure therefore stays a differentiable function of the interior pressure (what the JFNK line search

needs), while the ionization state it divides by is lagged exactly like every other composition quantity across a hydro step. `BC_component_constrho` is the only place the lower ghost is built, and marching, the JFNK residual and the reconstruction boundary all go through it, so the closure cannot differ between code paths.

The key sets the same quantity as `Hydrostatic base: True` (the ghost pressure, which that key extrapolates from the interior gradient instead), so the two are mutually exclusive: `Hydrostatic base` is tested first and `input_read` warns when both appear. With `Base BC: pressure` the microbar target is imposed at T_0 when n_0 is derived, so under a continuous- T ghost only the base *density* stays anchored; `input_read` notes that too. The closure in effect is echoed in `EXHALE_setup.out`.

Measured on HD 189733 b in the production configuration, restarted from the converged production state and Newton-finished in a scratch copy. The baseline restart reproduces the stored state digit for digit, so the comparison is like for like.

	isothermal	continuous
T_{ghost} / T_1 [K]	1183.0 / 493.6	539.9 / 540.2
$\rho_1 / \rho_{\text{gh}}$	2.33	0.95
$A(\ln \rho)$ cells 1–12	0.0264	0.0092
$A(\ln \rho)$ cells 3–12 / 5–14	4.41e-3 / 1.62e-3	9.28e-3 / 8.95e-4
JFNK	info=0, $\ R\ =5.37\text{e-}4$	info=0, $\ R\ =2.76\text{e-}4$
$\log_{10} \dot{M}$	9.04	9.05

The boundary jump and the density inversion are gone; the extended alternating tail is gone as well (beyond cell 7 the amplitude drops $13\times$ to $420\times$), and what remains is a stronger disturbance confined to cells 2–4. Three further restart+Newton cycles reproduce each state digit for digit, so both are exact fixed points and no further decay is available from cycling. Details, window sweep and scope: `docs/hd189_base_checkerboard.md` §13. **Not** measured at that point: a cold start under the new closure (6 h, 10^6 steps), other planets, transit observables.

Adopted for production. HD189733b/`input.inp` now carries `Base ghost temperature: continuous`, and the planet folder was re-converged with it together with the ionization-root validation of the next section (`run_20260811_rootfix.log`): $\log_{10} \dot{M} = 9.05$, JFNK info = 0, $\|R\| = 1.72 \times 10^{-4}$, 2002 steps. Measured on that converged state, $T_{\text{ghost}}/T_1 = 529.4/529.6$ K, $\rho_1/\rho_{\text{gh}} = 0.95$, and $A(\ln \rho)$ over cells 1–12 is 1.9×10^{-3} , i.e. well below the 0.0092 of the isolated A/B above — the root validation appears to remove a further part of the base disturbance. The paper and `python/paper_data.py` carry 9.05 for this planet.

44.2 Max steps: <N>

`count_max` was an integer, `parameter = 1000000`. The preceding section recorded the consequence: a $4\times$ -refined base, whose CFL step is $4\times$ smaller, cannot reach a converged state within a cap that the default grid already used in full, and the `EXHALE_MAXSTEPS` hook only lowers the cap. `count_max` is now a runtime variable with the same default.

EXHALE_MAXSTEPS keeps its separate meaning (a deterministic mid-loop exit used by the regression harness); it does not raise count_max.

44.3 Coronal cutoff width: <w>

The coronal-excitation guard of the preceding section multiplies every CHIANTI-derived coronal coefficient by $\exp(-x^2)$ with $x = (T_{\text{floor}}/T - 1)/w$ below $T_{\text{floor}} = 10^3$ K. $w = 0.5$ was a hardwired parameter. It is a modeling choice, not a measured quantity, and the temperature the base settles at depends on it at the ~ 100 K level, so it is now settable, still defaulting to 0.5 (division by a variable holding 0.5d0 is the same arithmetic, hence byte-identical). Values ≤ 0 are rejected at parse time. The manual states the sensitivity where the guard is documented.

Checks. make check: wasp_full, wasp_he23off, mol_base_handoff all byte-identical with the keys absent. WASP-121 b run to completion with the keys absent: $\log_{10} \dot{M} = 13.17$, JFNK info = 0, $\|R\| = 5.686 \times 10^{-5}$, 8228 steps, and its output byte-identical to the same case run with a binary built from HEAD without these changes. The parse-dump corpus (backup/regression/run_parse_corpus.sh) gains three lines per case and needs a deliberate golden refresh; it was not refreshed there. Files:

```
src/modules/init/parameters.f90
src/modules/states/Apply_BC.f90
src/modules/functions/composition.f90
src/modules/radiation/Cool_coeff.f90
src/modules/files_IO/input_read.f90
src/modules/files_IO/write_setup_report.f90
```

45. Ionization-equilibrium roots validated against the physical simplex (2026-08-11)

Ion_species.txt carried negative number densities (H II, O II, O III, Fe I) in a thin band just above the base: 17 cells at $r = 1.0106\text{--}1.0140 R_p$ in the wasp_full golden, 15 cells in the WASP-121 b production run, 5 in HD 189733 b.

The equilibrium systems are polynomial in the stage fractions and possess roots outside the physical simplex. At the step where secondary ionization is switched on (step 7186 of wasp_full) MINPACK hybrd1 returned info = 1 on such a root — $x(\text{H II}) = -5.2 \times 10^{-6}$, the Fe stage fractions summing to 1.2209, i.e. an Fe I density 22% below zero. Nothing tested the root, and the next step's warm start re-seeded the solve from the same values, so the state reproduced itself for the remaining 6302 steps. The heating paid for it twice: the metal photoheating channel went negative (-6.6×10^{-7} erg cm $^{-3}$ s $^{-1}$, which is impossible), and the ionized fraction, clamped into [0, 1], sent the Spitzer & Scott (1985) heating fraction f_{heat} to zero, switching off every photoheating channel above the photoelectron threshold. Total heating in the band sat 93% below its neighbors. Clamping the offending fractions to zero is therefore not a repair — the clamp is exactly what produces $f_{\text{heat}} = 0$; the cure has to be a physical root, not a clipped one.

Every atomic cell solve now validates its root — each stage fraction ≥ 0 , and each element's ionized stages summing to at most its nucleus total — and on failure restarts from physically defined starting points: the uncoupled ionization balance at the incoming electron density (ionization_balance_at_fixed_ne, admissible by construction, since the stage populations of a three-stage element come out proportional to $(d_1 d_2, u_0 d_2, u_0 u_1)$), then the optically thick fully neutral limit. A stored state that is not physical no longer seeds the next solve, which is what breaks the self-sticking. A cell where no starting point produces an admissible root is left on the ionization balance and reported rather than accepted. A first attempt that converges to a physical root is accepted unchanged, so healthy cells follow exactly the same solver path as before. q_abs, an absorbed energy rate, must now be positive before it normalizes the heating efficiency.

Gates (single-threaded; a baseline make check on the same tree passed all three cases, so the differences are this change alone). wasp_full FAILs by intent: negative densities 17 cells $\rightarrow 0$, $\log_{10} \dot{M}$ 13.22 $\rightarrow$ 13.22, $T/\rho/v/p$ within 0.89%/0.77%/1.3%/0.13% of the golden away from the band, metal photoheating positive everywhere (domain minimum $-8.9 \times 10^{-7} \rightarrow +1.03 \times 10^{-7}$). wasp_he23off also FAILs: its golden looks clean but the run reports 35 roots outside the simplex over 13478 steps, previously accepted in silence; $\dot{M}$ unchanged, $T/\rho/p$ within 0.28%/0.086%/0.039%. mol_base_handoff is byte-identical. The two production runs, each re-run once with the pre-change binary and once with the fix so the comparison isolates this change, keep their mass-loss rates and step counts (WASP-121 b 13.17, 8228 steps; HD 189733 b 9.05, 2002 steps) and lose their negative densities entirely. HD 189733 b is the clearest demonstration of the self-sticking: its stored initial condition is one of the poisoned outputs, and at step 1 the run rejects exactly the 5 stored cell states that carried negative densities and resolves all of them. Goldenes were not re-snapshotted at that point.

Not covered. The molecular network keeps its own chemical-equilibrium retry and is left alone; the H-only system (no helium) and the post-processing advection solves in post_process_adv.f90 are outside this validation, and the advected profiles of HD 209458 b still carry small negative H II, He III and He 2³S entries in the first 14 cells above the base. Full account: docs/ionization_root_validation.md. Files:

```
src/modules/radiation/ionization_equilibrium.f90
src/modules/radiation/util_ion_eq.f90
src/modules/nonlinear_system_solver/newton_solver.f90
src/EXHALE_main.f90
```

46. Line trapping in the thick metal resonance lines: measured, no change (2026-08-11)

The metal-cooling section above introduced $\beta(\tau)$ for [O I] 63 μm and [C II] 158 μm and recorded the thick resonance lines of a metal-rich wind (Mg II h&k, Ca II H&K, Na I D, the Fe II UV multiplets) as an untreated case. They were measured on the converged WASP-121 b, HD 209458 b and HD 189733 b runs. **No code path changed.**

The lines are thick — Mg II k reaches $\tau_0 = 7.6 \times 10^4$ at the WASP-121 b base and stays above 10^2 through the region where Mg II carries a third of the local cooling, and over 99% of the integrated Mg II / Ca II / Fe II / Mg I cooling comes from gas with $\tau_0 > 1$. That is not the test, though. Trapping lengthens the random walk of a resonance photon; it does not destroy it. In two-level equilibrium the correction to an optically thin coronal fit is

$$S = \frac{\beta A_{ul}}{\beta A_{ul} + n_e q_{ul}}, \quad q_{ul} = 8.629 \times 10^{-6} \Upsilon / (g_u \sqrt{T}),$$

not a factor β , so it bites only above $n_{\text{crit,eff}} = \beta A_{ul} / q_{ul}$. For these permitted lines $A_{ul} \sim 10^8 \text{ s}^{-1}$, so even at $\beta \sim 10^{-6}$ the escape rate keeps $n_{\text{crit,eff}}$ at 10^9 – 10^{15} cm^{-3} , against a maximum n_e of $4.3 \times 10^9 \text{ cm}^{-3}$ in the WASP-121 b run. The forbidden fine-structure lines are the opposite case ($A_{ul} \sim 10^{-5}$ – 10^{-3} s^{-1} , $n_{\text{crit,eff}} \sim 10^0$ – 10^5 cm^{-3}), which is why they, and only they, need β .

Measured $\Delta\Lambda$ integrated over the wind: **0.0214%** of the total radiative losses of WASP-121 b, 0.0022% of HD 209458 b, 0.0015% of HD 189733 b. The largest local effect is 1.65% in the first WASP-121 b cell, whose T is pinned at T_{eq} by the boundary condition; the loss exceeds 0.1% only below $r = 1.007 R_p$. Per line it is 0.1% of the Mg II cooling, 0.015% of Ca II, 1.4% of Na I (its A_{ul} is the smallest of the set) and 0.02–0.8% of Fe II. The measurement is a conservative bound: it uses the Doppler-core escape probability, which understates β by 4–10 $\times$ for these Voigt profiles, and no turbulent broadening.

Checked in passing: the Fe II d^6D fine-structure lines at 25.99/35.35 μm , the one metal transition with an A_{ul} small enough to trap, reach only $\tau_0 \leq 0.06$, so the cool_FeII_ne statistical-equilibrium table needs no escape probability either. H I Ly α , outside this scope, has $\tau_0 = 1.3 \times 10^8$ and a cooling-weighted $S = 0.9994$ — two-level trapping is therefore not what suppresses Ly α cooling, and the destruction channels that could (photoionization of $\text{H}(n=2)$, collisional $2p \rightarrow 2s$) were not evaluated.

$\beta = 1$ for the resonance lines is kept and is now documented at the code site as a justified effective treatment with its validity condition, in the SCOPE comment of `src/modules/radiation/Cool_coeff.f90`, instead of an acknowledged omission. Full measurement: `docs/resonance_line_trapping.md`.

47. Metal electrons and a validity range for the post-process advection correction (2026-08-12)

Two defects in `post_process_adv`, both confined to the advection-corrected `*_adv.txt` output: the post-processor is the last call of the run and its arguments are `intent(in)`, so the converged wind and the equilibrium `*.txt` files cannot be affected by anything here.

The electron density inside the advection residuals was wrong. It counted only the H and He electrons, while the equilibrium residual it corrects counts the metal electrons as well (`metal_electron_sum`) and the *same* routine already used the metal-inclusive n_e for its heating, cooling and temperature solve. In a shielded base the metals are the dominant electron donors — measured $n_{e,\text{metal}}/n_e = 0.48$ – 1.00 in the base cells of all four paper planets — so the recombination terms were low by up to six orders of magnitude. Solved to machine precision with Brent, the residual *as coded* had its fixed point at $x_{\text{HII}} = 4.74 \times 10^{-8}$ in the first HD 209458 b cell, against an equilibrium 7.46×10^{-14} .

A field `adv_cell%xe_metal` now carries the metal electrons per H nucleus into all three advection residuals, with the definition `calc_ne` and `metal_electron_sum` use. It is not conditioned on `eos_include_metals`: that switch decides whether the metals enter the mass and particle budget, whereas recombination needs the true electron density. Metals off gives zero identically.

The correction was applied where it cannot be computed and is not needed. The residuals carry the *neutral* fraction and report the ion density as $(1 - x_{\text{HI}})n_h$, so the ion fraction inherits the solver's absolute resolution on x_{HI} , $\text{xtol} = \sqrt{\epsilon} = 1.5 \times 10^{-8}$. Where the equilibrium ion fraction is 10^{-13} the returned value is quantized at 10^{-8} with an arbitrary sign: the first HD 209458 b cell got $x_{\text{HII}} = -4.96 \times 10^{-8}$ where the exact root is $+4.74 \times 10^{-8}$ — right magnitude, wrong sign — and the upwind cascade carried the negative density outward over 14 cells. Those same cells are in local ionization equilibrium to

$$\text{Da} = \frac{\Delta r}{v} (\Gamma_{\text{HI}} + \alpha_{\text{HII}} n_e) = 350\text{--}2900,$$

so the equilibrium solution is the answer the ODE would give anyway. The correction is now skipped, and the converged equilibrium ionization kept, wherever the cell is (i) inflowing ($v \leq 0$ on either face: the upwind discretization has no upstream cell), (ii) equilibrium-dominated ($\text{Da} > 100$) or (iii) unrepresentable ($x_{\text{HII,eq}} < 10^{-6}$, less than 1% relative accuracy at that resolution). Conditions (i) and (ii) are statements about the flow and (iii) about the representation of the unknown, so the pre-existing `pp_metal_on` gate on the inflow condition is removed — it made a property of the discretization depend on the metal-cooling switch. The temperature loop keeps condition (i) alone, likewise ungated; a thermal Damköhler number is not evaluated. Clipping the extracted densities at zero was tested and rejected: it breaks the H nucleus budget $n_{\text{HI}} + n_{\text{HII}} = n_h$.

Gates. make check: PASS, all three cases byte-identical (the goldens compare `Hydro_ioniz.txt` and `Ion_species.txt`, which this does not write). Isolated A/B, the pristine HEAD binary and the modified one run from the same configuration: `Hydro_ioniz.txt`, `Ion_species.txt`, `Cooling_breakdown.txt` and `Excited_H.txt` byte-identical, only the `_adv` files differ.

planet	$\dot{M}$ [log g/s]	negative _adv	cells held at eq (newly)	outermost held	max change, $r > 1.05$
HD 209458 b	9.31 → 9.31	48 → 0	103 → 124 (21)	1.0316	1.4% (He III)
HD 189733 b	9.05 → 9.05	0 → 0	135 → 212 (77)	1.0380	0.37% (He III)
WASP-121 b	13.17 → 13.17	0 → 0	0 → 0 (0)	—	3.8% (He III)
WASP-52 b	11.63 → 11.63	0 → 0	3 → 57 (54)	1.0106	3.3% (He III)

Every cell whose equilibrium ion fraction is below 10^{-6} is now held at the equilibrium value, and no newly held cell lies above $r = 1.038 R_p$. The change in the wind is entirely the metal electrons: a build carrying the guard alone removes all 48 HD 209458 b negative entries by itself and changes the profiles above $r = 1.05 R_p$ by exactly zero on three of the four planets and by 4.2×10^{-5} on WASP-52 b. Transmission spectra move by at most 0.07% in peak depth (He 10830 of HD 209458 b) and 0.08% in equivalent width ($\text{H}\alpha$ of WASP-121 b).

The re-run reproduces the stored output/ of WASP-121 b and WASP-52 b byte-identically but not that of HD 209458 b and HD 189733 b, which restart from output/*_IC.txt files that had been refreshed to the previous run's converged state; re-running continues that convergence (HD 209458 b du : $1.199 \times 10^{-2} \rightarrow 3.13 \times 10^{-3}$) and moves the profiles

by up to 9% in density at unchanged $\dot{M}$. That is unrelated to this change — the pristine binary follows the same new trajectory — and is why the table above is an A/B between two binaries.

Files:

- src/modules/post_process/post_process_adv.f90
- src/modules/nonlinear_system_solver/ion_cell_state.f90
- src/modules/nonlinear_system_solver/System_implicit_adv_{H,HeH,HeH_TR}.f90

Full account: docs/postprocess_advection_validity.md.

Found and **not** fixed, marked at the code site: System_implicit_adv_HeH (helium on, He 2³S off) writes its electron-impact ionization terms without the $n_e n_h$ factor that adv_implicit_H, adv_implicit_HeH_TR and the equilibrium rows all carry, which also makes them dimensionally inconsistent with the photoionization rates they are added to. None of the four paper planets uses that path. (Since fixed, 2026-08-12: the terms now carry ionhi*xe*n_h; commit b0d44bc, TO_BE_DONE item (B).)

48. Ground-term fine-structure statistical equilibrium for C I, C II, N II and O I (2026-08-12)

The CHIANTI C/N/O cooling fits carry the ground-term fine-structure (FS) lines in the optically thin, *low-density* limit. Their critical densities are of order 10^0 – 10^5 cm⁻³, decades below the base density of an irradiated atmosphere ($n_e \sim 10^9$, $n_{\text{HI}} \sim 10^{14}$ cm⁻³), so at the base the coronal form overstates what those lines can radiate by 2.9–8.3 decades. Two of them ([O I] 63 μm , [C II] 158 μm) already had a two-level saturation; C I, N II and the [O I] 145.5/44 μm channels did not, and C I was in consequence the dominant coolant of the whole base layer of HD 189733 b, at a value comparable to the local heating rate.

What replaces it. Every C/N/O coolant whose ground term is split now has that term solved in exact statistical equilibrium at the local (n_e, n_{HI}),

$$\sum_j f_j R_{ji} = f_i \sum_j R_{ij}, \quad \sum_i f_i = 1,$$

$$R_{ul} = C_{ul} + \beta_{ul} A_{ul}, \quad R_{lu} = C_{ul} \frac{g_u}{g_l} e^{-E_{ul}/k_B T}, \quad C_{ul} = n_e k_{e,ul}(T) + n_{\text{HI}} k_{\text{H},ul}(T),$$

$$W_{\text{FS}} = \sum_{u>l} f_u \beta_{ul} A_{ul} k_B E_{ul} \quad [\text{erg s}^{-1} \text{ per ion}],$$

and the coronal curve is refitted to keep only the channels that *leave* the ground term (Λ_{rem}). The coefficient handed to the assembly is $\Lambda_{\text{eff}} = W_{\text{FS}}/n_e + \Lambda_{\text{rem}}$, so the $n_e n_{\text{ion}}$ prefactor recovers the H-collision channel exactly. Treated: C I $2p^2\ ^3P$ (609.1/370.4 μm), C II $2p^2\ ^3P$ (157.7 μm), N II $2p^2\ ^3P$ (205.3/121.8 μm), O I $2p^4\ ^3P$ (63.2/145.5/44.1 μm). That is the complete set — N I and O II have a single-level 4S ground term, Mg I/II, Ca II and Na I a single ground level, Fe I is built from permitted lines only, and Fe II already uses a two-dimensional statistical-equilibrium table.

Limits: $n_e, n_{\text{HI}} \rightarrow 0$ collapses onto the ground level and reproduces the coronal within-term channel to the accuracy of the collision-strength fits (1.3–4.4%); high density saturates at the exact multilevel LTE emission (verified to 10^{-14}

relative). β_{ul} enters inside the solution, not as a factor on the result, which is the only placement that is right in both limits. All eight lines with a transition probability now get their own line-center escape probability from the outward column; previously only two did.

Two defects fixed on the way. (i) The two-level form multiplied a solution already normalized to the two-level partition function by the ground-term Boltzmann fraction of its lower level, so its LTE limit was low by $1 + (g_u/g_l)e^{-E/k_B T}$ — $2.7\times$ for [C II] 158 μm at the 530 K base of HD 189733 b. (ii) The H-atom de-excitation rates (a constant 4.0×10^{-11} for C II, $4.2 \times 10^{-11}(T/100)^{0.67}$ for O I, both flagged approximate in the source) are replaced by fits to published quantum scattering calculations: Yan & Babb (2023, MNRAS 518, 6004) for C I and N II, Abrahamsson, Krems & Dalgarno (2007, ApJ 654, 1171) for O I, Barinovs et al. (2005, ApJ 620, 537) for C II. They are 5–20 $\times$ larger.

What it does. The coronal_excitation_cutoff guard was standing in for the missing saturation, so its width w behaved as a tuning knob on the base temperature. It no longer does: the local balance temperature of the base cell is identical to every digit printed over $w = 0.02$ –1.2 on all four planet runs, against a 500–650 K spread before.

run	base T	$\dot{M}$ at $2R_p$
HD 189733 b (warm restart, $w = 0.1$)	528.5 $\rightarrow$ 613.2 K	$2.351 \times 10^9 \rightarrow 2.850 \times 10^9$ g/s (+21.2%)
WASP-121 b (cold start)	2406.9 $\rightarrow$ 2438.8 K	$3.106 \times 10^{13} \rightarrow 3.370 \times 10^{13}$ g/s (+8.5%)
wasp_full golden case	2340.6 $\rightarrow$ 2357.3 K	$3.482 \times 10^{13} \rightarrow 3.549 \times 10^{13}$ g/s (+2.0%)

Transit depths on HD 189733 b move by +4.4% (Ca II) to +20.7% (H β). The profiles outside $r \sim 1.1 R_p$ move by a few percent; the change is a base-layer change that propagates through the density.

Files:

- src/modules/radiation/Cool_coeff.f90 — the SE solvers, the four ion coefficients, the generalized line opacity and escape routine
- src/modules/radiation/util_ion_eq.f90
- src/modules/nonlinear_system_solver/T_equation.f90
- src/modules/post_process/post_process_adv.f90
- cooling_data/fit_fs_saturation.py — generator; prints every coefficient
- cooling_data/make_cooling_doc_figures.py
- docs/cooling_formulas.tex, docs/coronal_cutoff_width.md section 7, TO_BE_DONE.md item (C)

Also fixed here, unrelated but found while reading the cooling diagnostics: examples/exhale_io.py named only four fixed channels in Cooling_breakdown.txt while the writer emits six (the collisional-excitation channel is written split by absorber), so every metal column was labeled two ions too early. python/paper_data.py parses the header and was already correct.

Still open: with the artifact gone it is plain that the base is not in local radiative balance at all — at the HD 189733 b base the total radiative cooling is 3.5×10^{-8} against a heating of 3.4×10^{-6} . Its temperature is set by the inner boundary and by the flow.

49. Base ghost built from the composition of the state it bounds (2026-08-12)

Apply_BC closes the Base ghost temperature: continuous ghost with

$$p_{\text{ghost}} = (\text{ntot_bc} + \text{dp_bc}) \frac{p(1)}{n_{\text{part_cell1}}},$$

where `n_part_cell1` is the cell-1 particle count written by `get_species_densities`. Two call sites filled the ghosts before any composition solve had run on the state in question, so they used whatever `n_part_cell1` was left over — the `input_read` placeholder `ntot_bc+dp_bc` in the first case:

- `init.f90`: the initial Apply_BC. On HD 189733 b the placeholder differs from the true cell-1 particle count by 4.4%, and the finite-volume steady residual of the loaded state (which stops right after `init`) reported the resulting contact discontinuity at the base face as a cells 1–2 spike of 60 (mass) and 120 (energy) per sound crossing time.
- `steady_newton.f90`, `eval_residual`: Apply_BC after `unpack_U`. There $F(Y)$ depended on the previous Y as well as on Y , so a finite-difference Jacobian column mixed two states.

Both now evaluate the composition of the interior first. `init.f90` also seeds `base_flux_const` from the initial state, for the same reason: with Base velocity: `massflux` that constant starts at -1 , which Apply_BC reads as “not available yet” before falling back to the valve, so the first step and any diagnostic stopping right after `init` used a different velocity boundary condition than the run. Files: `src/modules/init/init.f90`, `src/modules/time_step/steady_newton.f90`.

Checks. `n_part_cell1` is read only by the continuous branch and `base_flux_const` only by the mass-flux branch, so a run setting neither key is unaffected and make check is byte-identical; an isothermal HD 189733 b residual reproduces bit for bit. With the key on, the cells 1–2 mass residual falls from $5.97 \times 10^1, 5.92 \times 10^1$ to $6.41 \times 10^{-2}, 7.70 \times 10^{-2}$ and the energy residual from $1.12 \times 10^2, 1.20 \times 10^2$ to $1.31 \times 10^{-1}, 1.54 \times 10^{-1}$, with cell 3 outward and the wind window unchanged to every digit. Re-converging the planet from its stored state: JFNK scaled merit $\|F_s\|_2$ $1.43 \rightarrow 0.315$, `info` = 0, $\log_{10} \dot{M}$ $9.15 \rightarrow 9.14$. Measurements and term-by-term split: `docs/hd189_base_checkerboard.md` section 15.

50. $H(n=2)$ rate coefficients, the trapped- $2p$ balance, and de-excitation heating (2026-08-12)

Four corrections to the excited-hydrogen and $\text{Ly}\alpha$ modules, plus the structural change that made the first one possible. Measurement and the full number tables: `docs/lya_destruction_channels.md` section 11.

One definition of the $n=2$ rate coefficients. The $\text{Ly}\alpha$ and $n=2$ atomic data and every collisional rate coefficient existed in three copies — inside `n2_populations`, again as standalone functions in the same file, and a third time in `lya_rt.f90`. They are now one module, `src/modules/radiation/hydrogen_n2_rates.f90`, which

exports `c1s2s_rate`, `c1s2p_rate`, `c2s2p_rate`, `c2s1s_rate`, `c2p1s_rate`, `c2p2s_rate`, `alpha_B_hydrogen`, `alpha_2s_hydrogen`, `alpha_2p_hydrogen`, `n2p_destruction_rate` and the line constants. `excited_hydrogen.f90` and `lya_rt.f90` use it.

1. **Statistical weights shadowed by dummy arguments.** `n2_populations` was declared (`T`, `n1s`, `ne_l`, `Jlya`, `G2s`, `G2p`, `n2s`, `n2p`). Fortran does not distinguish `G2s` from the module parameter `g2s`, so the three detailed-balance rate coefficients written out in the routine body evaluated their weights as the Balmer photoionization rate the caller passed: $2s \rightarrow 1s$ and $2p \rightarrow 1s$ de-excitation came out 5 to 306 times too small (planet-dependent, scaling with the stellar Balmer continuum) and the $2p \rightarrow 2s$ ℓ -mixing rate 3 times too large. The dummies are now `gam_ion_2s / gam_ion_2p` and the body calls the shared rate functions, so the weights are resolved at module scope and out of a dummy's reach. Two further shadows of the same kind, both harmless, were renamed: the frequency-sample count `ng = 400` in `gamma_n2_balmer` and `heat_n2_balmer` shadowed the global ghost-cell count `Ng`; it is `n_nu`.
2. **Recombination cascade source.** $a2s \cdot ne^{**2} / a2p \cdot ne^{**2} \rightarrow a2s \cdot ne \cdot n_{HII} / a2p \cdot ne \cdot n_{HII}$. The electron recombines onto a proton; at the base the electrons come largely from helium and metals while hydrogen is still neutral, and n_e/n_{HII} reaches 2 on WASP-121 b and 5×10^6 on HD 209458 b. The Python transmission tool (`exhale_transit_lib.py`, `EXHALE_transit.py`) and `examples/tpm_halpha_lart2d.py` carried the same expression and were corrected the same way; their `n2_populations` signatures gained an `nHII` argument.
3. **Destruction channels in the `J_int` closure.** `lya_rt.f90` built the trapped field from $n_{2p} = P / (A_{2p \rightarrow 1s} \beta)$. The $2p$ budget also loses atoms to collisional de-excitation, $n = 2$ photoionization, and ℓ -mixing followed by two-photon decay; the denominator is now $A_{2p \rightarrow 1s} \beta + D$ with D from `n2p_destruction_rate`. Where β is smallest the omission over-estimated `J_int` by up to $1.5\times$. In the same module the top-down $\text{Ly}\alpha$ line-center optical depth, previously two identical inline loops, is the single routine `lya_line_center_optical_depth`; all three `jlya_mode` values call it, so the `tau_Lya` column of `Excited_H.txt` is filled in the LaRT import mode instead of being written as zero.
4. **Deexc heat now defaults on.** The comment justifying the old default — that collisional de-excitation heating overlaps the H I collisional-excitation cooling — was wrong: `coex_rate_HI` is the one-way Cen (1992) coronal rate with no de-excitation term in it, so `Hdx` is the correction to that limit, not a double count. Most of the $n = 2$ population it de-excites was pumped by $\text{Ly}\alpha$ rather than excited by a collision, which makes the term absorbed $\text{Ly}\alpha$ thermalized by a collision, the same size as the photoelectric heating the code already applied unconditionally. The key now also honors `Deexc heat: False`, which it previously ignored. **This changes results for any run that supplies a stellar T_{eff} /radius and does not set the key.**

Audited and left alone: the energy equation applies no escape-probability factor to the $\text{Ly}\alpha$ cooling. `cool` is built from the bare Cen coefficient (`Cool_coeff.f90:843` $\rightarrow$ `util_ion_eq.f90:640,787`); β appears nowhere in that path, in any `jlya_mode`. That is the correct form — 98.2–99.99% of trapped $\text{Ly}\alpha$ photons escape and the collisional destruction that would suppress the cooling is at most 0.06% — so the historical expectation of a $\sim 10\times$ transfer suppression is not in the code and should not be.

Checks. `mol_base_handoff` (no stellar T_{eff} , so the excited-H model is off) is byte-identical at every stage, as is its comparison against the stored golden. `wasp_full` and `wasp_he23off` move: $\log_{10} \dot{M}$ 13.23039 $\rightarrow$ 13.22966 and 13.23106 $\rightarrow$ 13.23026, with the largest local temperature change 0.57% (item 1), 0.081% (item 2) and 0.147% (item 3); both cases `pin Deexc heat: True` in their `input.inp`, so item 4 moves nothing in them. Their golden were already stale against the pre-fix tree and were left as they are. Re-converged from their stored initial conditions in a scratch copy, WASP-121 b gives 13.2075 $\rightarrow$ 13.2030, HD 189733 b 9.1286 $\rightarrow$ 9.1392 and HD 209458 b 9.3790 $\rightarrow$ 9.4581; the $H\alpha$ and $H\beta$ transit depths move by 0.4–1.7% ($n(2p)$, which carries the opacity, barely moves, while $n(2s)$ halves). HD 209458 b is the largest move (0.075 dex, from items 1–3) because its $\text{Ly}\alpha$ cooling is a negligible share of its radiative losses while the $n = 2$ heating terms are not; it is also the case `docs/lya_deexcitation_heating.md` recorded as going NaN at the base with the de-excitation heating on, which no longer reproduces with the corrected populations (re-converges from its stored initial condition, and a cold start survives a 20000-step cap). A build at `-O1 -fopenmp -g -fcheck=bounds,do,mem` running the `wasp_full` configuration is clean.

Files:

- `src/modules/radiation/hydrogen_n2_rates.f90` (new)
- `src/modules/radiation/excited_hydrogen.f90`, `src/modules/radiation/lya_rt.f90`
- `src/modules/init/parameters.f90`, `src/modules/files_IO/input_read.f90`, `Makefile`
- `exhale_transit_lib.py`, `EXHALE_transit.py`, `examples/tpm_halpalpha_lart2d.py`

51. Molecular chemistry and trace metals solved in one system (2026-08-13)

The parser refused `Molecular chemistry: True` together with a `metals.inp`. That refusal blocked two things: the deep (10^{-3} bar) lower-atmosphere handoff of `docs/base_composition_handoff_plan.md` §11.7, whose A/B pair carries metals and whose consistent configuration is Tier-2 molecular chemistry; and trace-metal work on sub-Neptunes, which have a molecular base by construction.

The two networks cannot be solved apart, and the reason is quantitative rather than architectural: **in the shielded molecular base the metals supply essentially all the free electrons**. Measured at the base cell of the two smoke tests below, with solar metal abundances:

case	n_e with metals [cm^{-3}]	metal share	n_e metals-off [cm^{-3}]
hot Uranus, solar 7 metals	1.20×10^8	1.000	6.29×10^4
HD 209458 b, 10^{-3} bar handoff, C/N/O	9.42×10^5	0.990	9.27×10^4

The flux that ionizes the low-ionization-potential metals — Mg I at 7.65 eV, Fe I at 7.90 eV, C I at 11.26 eV — lies below both the H I edge (13.60 eV) and the H_2 threshold (15.4 eV), so it is absorbed by neither and reaches the base almost unattenuated. On the hot-Uranus case it holds Mg and Fe about 1.6% ionized at the base, which is already enough to swamp the H/He and molecular electrons. Every recombination and electron-impact term of the molecular network scales with n_e , so solving the network at the metal-free n_e misses a factor 1900 (hot Uranus) to 10 (HD 209458 b at 10^{-3} bar) in that layer.

The merged system. `System_HeH_mol_metals` follows the pattern `System_HeH_TR_metals` established: the molecular unknowns keep their rows, the metals are appended above them, and both blocks are shared code rather than re-derived. The eight balance rows of the molecular network moved into `mol_heh_rows` (still in `System_HeH_mol`, with the Koskinen 2022 rate coefficients, which stay defined in exactly one module) and take n_e as an input, exactly as `heh_tr_rows` does; the metal rows, the metal electron sum and the stage fractions come from `ion_residual_core` unchanged. The layout is

```
x(1..3)      H II, He II, He III
x(4..7)      H2, H2+, H3+, HeH+
x(8)         He 23S (when tracked)
x(mbase..)   X+, X++ for each metal element, mbase = metal_row_base() = 8 or 9
```

with $N_{\text{eq}} = 7$ (+1 with the triplet) + $2n_{\text{melem}}$. `metal_row_base()` is the single definition of that offset, read by the residual, by the driver's seeding and root validation, and by `cx_metal_base` so that the Huang Table-4 charge exchange writes to the shifted metal rows. With `met_nelem = 0` the system reduces exactly to `System_HeH_mol`.

Root validation and seeding (§45 extended to the molecular branch). The molecular cell solve is bistable because the cell is: a molecular basin (the dense, optically thick base) and an atomic basin (the wind above the $\text{H}_2 \rightarrow \text{H}$ front). It now tries one starting point in each in turn — the warm start, the chemical-equilibrium H_2 fraction at the local (p, T) with the metal stages from their own ionization balance at the incoming n_e , then the molecule-free ionization balance of every element — and keeps a converged root only if it is also physical, on the same rule §45 applied to the atomic systems. The metal part of `ionization_balance_at_fixed_ne` was split into `metal_ionization_balance_at_fixed_ne` so the molecular retry seeds its metals from the same balance instead of a second copy of it. The third starting point is what the front cells needed: on the hot-Uranus case the coupled solve leaves 3 cells at $r = 1.21\text{--}1.22$ (the ionization front) with no admissible root at every step with two starting points, and 0 cells with three.

Two defects found in the surrounding code and fixed.

- The sub-13.6 eV energy grid of a power-law spectrum was an `if (thereis_HeITR) ... else if (thereis_lowIP_metal)`, with the comment that the triplet and metals together are unsupported so at most one applies. That has not been true since the triplet and the metals were merged. With the triplet on, the grid floor stopped at 4.80 eV, which is above the K I threshold (4.341 eV), so K I would be held spuriously neutral in a triplet run. The floor is now the lowest threshold over whichever sources are active, matching what the loaded-SED path (`sed_read`) already did. No golden carries K, so this is byte-identical there.
- The `pp_metals 2` metal re-solve passed the global N_{eq} to `hybrd1` while its residual (`ion_system_metals_pp`) writes only rows $1..3 + 2n_{\text{melem}}$. With the He 2³S triplet that already left one row of `fvec` unwritten; with the molecular layout it would leave four or five. It is now sized $3 + 2n_{\text{melem}}$ with a matching workspace, the same reasoning `lwa_adv` already carried a few lines above. Both goldens run `pp_metals 1`, so this is byte-identical there and changes results only for `pp_metals 2` runs.

Left in place at the time, with the validity range written at the code. `eval_cool` built its own electron density without the molecular ions. That is right in the hot atomic gas it was written for and wrong inside a deep molecular base, where H_3^+ can be the dominant ion; there it disagreed with the n_e the equilibrium solver itself uses (ionization_equilibrium does pass `nmol_eq` to `calc_ne`), so the n_e -scaling cooling channels were under-counted in that layer. Two definitions of the same quantity is one too many, but closing the gap changes the molecular results and belonged in its own change with its own measurement. *Closed in the next section.*

Gates. `make check` PASSes byte-identical on all three cases then in the matrix (`wasp_full`, `wasp_he23off`, `mol_base_handoff`); a baseline `make check` on the same tree before the change also passed, so the comparison is this change alone. Golden were not re-snapshotted. A build at `-O1 -fopenmp -g -fcheck=bounds,do,mem` is clean on all three coupled layouts for 200–300 steps: hot Uranus with the He 2^3S triplet ($N_{\text{eq}} = 28$, metals at row 9), the same case without it ($N_{\text{eq}} = 27$, metals at row 8), and the HD 209458 b deep handoff.

Hot Uranus (the Tier-2 gate configuration of `docs/lower_atmosphere_figs/data_g2`) plus solar C/N/O/Mg/Ca/Na/Fe, a 12000-step relaxation snapshot against the same case metals-off: the $\text{H}_2 \rightarrow \text{H}$ front sits at $r = 1.1577$ in both (the gate's 1.156), the base stays fully molecular in both, $\log_{10} \dot{M} = 10.58$ in both, no negative densities, and no cell is left without an admissible root (the same case with the triplet off, $N_{\text{eq}} = 27$ with the metals at row 8, is clean over a shorter 400-step smoke). What the metals change is the base chemistry: n_e rises $1900\times$ and H_3^+ falls by a factor 1350 ($7.08 \times 10^4 \rightarrow 52.5 \text{ cm}^{-3}$), its peak moving out from $r = 1.040$ to $r = 1.130$ — H_3^+ is destroyed by dissociative recombination (Koskinen R6/R7, both proportional to n_e), so it survives only where the metal electron fraction starts to fall. Metal ionization runs the expected way: C II/C = 8×10^{-10} at the base, 0.22 at the front, 0.43 at $1.9 R_p$; Mg III stays below 10^{-8} up to the front and reaches 2.4×10^{-3} only at $1.9 R_p$.

HD 209458 b at the 10^{-3} bar handoff (the §11.7 blocker), C/N/O, He 2^3S on, a 4000-step bounded run: it integrates. §11.7 records that the atomic pair built at this level terminated with a NaN at $r = 1.067$ after 1846 steps with the time step collapsed to $dtu = 1.2 \times 10^{-6}$ (that run was not repeated here); this one reaches the cap with $dtu = 0.198$ and no NaN, a fully molecular base, and the front at $r = 1.0150$ (1.0169 metals-off). It is a relaxation snapshot at $du = 747$, nowhere near converged, so no mass-loss rate should be read off it.

Open item. The `_adv` post-process is still molecule-free by design (`post_process_adv.f90` header): `nh = nhi + nhii` counts only the free H nuclei and the molecular electrons are absent from its n_e . The metal stages are carried there as usual under `pp_metals`, but on that molecule-free H/He background, so `_adv` metal profiles inside the molecular layer inherit the approximation. With the `eval_cool` electron density closed in the next section, this is the one remaining place where a molecular run and a metal run still meet on the atomic assumption.

Files:

- `src/modules/nonlinear_system_solver/System_HeH_mol_metals.f90` (new)
- `src/modules/nonlinear_system_solver/System_HeH_mol.f90`
- `src/modules/radiation/ionization_equilibrium.f90`, `src/modules/radiation/util_ion_eq.f90`
- `src/modules/files_IO/input_read.f90`, `src/modules/files_IO/write_setup_report.f90` (the startup report now names the species blocks the coupled system carries and prints N_{eq})

- `src/modules/init/set_energy_vectors.f90`, `src/modules/post_process/post_process_adv.f90`,
Makefile
- `examples/16_molecular_metals/ (new)`

Documentation: `README.md`, `README_HOWTO.md`, `examples/README.md`, `docs/EXHALE_user_manual.tex` (the example table also became a longtable, which it needed already — as a float it ran past the bottom margin into the page footer), `docs/input_schema.md`, `docs/lower_atmosphere_coupling.md`, `docs/base_composition_handoff_plan.md`.

52. One electron density for the cooling and for the ionization solver (2026-08-13)

`eval_cool` rebuilt its own free-electron density and left the molecular ions out of it. The equilibrium solver in the same sweep does not: it balances ionization against `calc_ne(nhii,nheii,nheiii,ne,nm,nmol_eq)`, i.e. H^+ , He^+ , He^{++} , the metal stages under the `eos_metals` policy, **and** H_2^+ , H_3^+ , HeH^+ . A molecular run therefore carried two definitions of n_e , and the cooling used the smaller one. That is wrong wherever a molecular ion carries a non-negligible share of the charge, which is exactly the deep molecular base. Measured on the metals-off hot-Uranus Tier-2 gate, the two definitions differ by $168.5\times$ at the base cell (373 against $6.29 \times 10^4 \text{ cm}^{-3}$, H_3^+ supplying essentially all of it), $4.5\times$ at $r = 1.05$, 2.8% at $r = 1.10$, and nothing above $r \sim 1.15$; on the `mol_base_handoff` case the base ratio is $175.7\times$.

`eval_cool` now takes an optional `nmol` (H_2 , H_2^+ , H_3^+ , HeH^+ , cgs) and passes it straight to `calc_ne` — the electron density is not reconstructed anywhere inside `eval_cool`, it has one definition and one implementation. Callers that track the molecular network supply it: `ionization_equilibrium` (`nmol_eq`), the semi-implicit energy solver (`nmol_dim`, the cgs copy of the adimensional array it already hands `calc_ne/calc_ntot`), and the `Cooling_breakdown` dump. The `_adv` advection post-process deliberately does not: that reconstruction is molecule-free by design (its module header), and omitting the argument gives it the atomic charge sum it is built on. `ionization_equilibrium`'s own `calc_ne` call lost its `thereis_mol` branch at the same time — `nmol_eq` is zero for an atomic run, so one call covers both.

Every n_e -scaling channel picks the correction up: recombination, collisional ionization, collisional excitation and bremsstrahlung are linear in n_e , and the coronal metal line coefficients enter as `ne*nm*c_metal`. The saturated ground-term coefficients (Fe II, and C I/C II/N II/O I under `cno_chianti`) are $\Lambda_{\text{eff}} = W/n_e$, so the assembly's n_e cancels and they respond only through the level populations ($C_{21} \propto n_e$) — the correct behavior, and there is no double counting of n_e anywhere in `eval_cool`.

Atomic runs are exactly unchanged. `nmol_eq` and the molecular `f_sp` columns are zero, and `calc_ne` adds them as `ne + 1.0*0.0`, which is exact. `make check` is byte-identical on `wasp_full` (metals + $He\ 2^3S$) and `wasp_he23off`; a full WASP-121 b run (metals, $He\ 2^3S$, $Ly\alpha$ escape probability, excited H, transonic IC) reproduces all eight output files byte-identically — `Hydro_ioniz{,_adv}`, `Ion_species{,_adv}`, `Cooling_breakdown`, `Heating_breakdown`, `Excited_H`, `IC_dump` — at $\log_{10} \dot{M} = 13.20$.

`mol_base_handoff` changes, and that is the point of the change. At the pinned 12000-step snapshot the two runs stop at the same $du = 2.6689$ and the same dtu ; every hydrodynamic column agrees to $\leq 4 \times 10^{-8}$ relative and every

species column to $\leq 4.2 \times 10^{-6}$ (that worst case is H_3^+ at $r = 2.96$, where H_3^+ is $3 \times 10^{-19} \text{ cm}^{-3}$). The one column that moves is the reported cooling rate at the near-neutral base, $2.799 \times 10^{-20} \rightarrow 4.942 \times 10^{-18} \text{ erg cm}^{-3} \text{ s}^{-1}$ ($\times 176.5$, the n_e ratio, as expected for channels linear in n_e). $\log_{10} \dot{M}$ stays 10.58.

Why the dynamics barely notice. A molecular layer cools through the H_3^+ infrared lines, which do not scale with n_e , so multiplying the electron channels by 176 at the base moves the cooling from 2.8×10^{-20} to 4.9×10^{-18} against $1.3 \times 10^{-7} \text{ erg cm}^{-3} \text{ s}^{-1}$ of photoheating there. The correction is real and it is now the same quantity the solver balances ionization against, but it is energetically negligible in the layer where it is large.

Molecular smoke tests. Hot Uranus, the Tier-2 gate configuration, 12000-step relaxation snapshots — both runs stop at du of order unity, so these are snapshots and not converged solutions.

	metals off	metals on (solar C/N/O/Mg/Ca/Na/Fe)
$n_e(\text{base}) \text{ old} \rightarrow \text{new}$	$373 \rightarrow 6.29 \times 10^4 \text{ cm}^{-3} (\times 168.5)$	$1.2016 \times 10^8 \rightarrow 1.2016 \times 10^8 (\times 1.0000)$
cooling at the base	$3.18 \times 10^{-20} \rightarrow 5.36 \times 10^{-18} \text{ erg/cm}^3/\text{s}$	$+7 \times 10^{-7} \text{ relative}$
$\text{H}_2 \rightarrow \text{H}$ front	$1.15631 \rightarrow 1.15631$	$1.15677 \rightarrow 1.15677$
H_3^+ peak	$7.0808 \times 10^4 \rightarrow 7.0807 \times 10^4 \text{ cm}^{-3}$ at $r = 1.0396$	$52.465 \rightarrow 52.465$ at $r = 1.1304$
du at step 12000	$4.3354 \rightarrow 4.2914$	$5.8177 \rightarrow 5.8177$
$\log_{10} \dot{M}$	$10.58 \rightarrow 10.58$	$10.58 \rightarrow 10.58$

The metals-on column is the expected null: the low-ionization-potential metals hold n_e at 1.2×10^8 at that base (previous section), against 52 cm^{-3} of H_3^+ , so adding the molecular ions moves nothing (maximum relative change 7×10^{-10} in $\rho/T/p$, 3×10^{-9} in the metal stages). Metals-off, the snapshot differs by 1% in du and by sub-percent in the bulk state just above the layer ($\rho -0.7\%$, $T +0.8\%$ at $r = 1.19$), with the trace species that depend exponentially on T amplifying that into H_3^+ -16% , $\text{H}_2 -13\%$, $\text{He II} +10\%$ at the same radius. Those are differences between two snapshots of an unconverged relaxation, not a converged shift; the gate observables — front position, H_3^+ peak, mass-loss rate — are unchanged.

On the HD 209458 b 10^{-3} bar handoff with C/N/O (a 4000-step bounded snapshot, also unconverged) the base picture is the same: the metals dominate n_e , the base cooling rises 0.45%, the molecular-layer cooling by at most 1.9%, the front moves $1.01503 \rightarrow 1.01518$ and the H_3^+ peak by 0.05%. An `-01 -fopenmp -g -fcheck=bounds,do,mem build` runs all three molecular layouts (hot Uranus with and without metals, HD 209458 b deep handoff) for 300 steps with no runtime check firing.

Two related items found and recorded, not changed here.

- The H_3^+ infrared cooling is added to cool in `ionization_equilibrium`, *after* `eval_cool` returns. The JFNK steady residual uses that array, but the semi-implicit energy update (the default) rebuilds cool from `eval_cool` alone and overwrites it, so it never sees the H_3^+ term. For a molecular run the marching stage and the Newton stage therefore integrate slightly different energy equations, and the marching fixed point is not the zero of the residual the Newton solver drives down. The fix is to move the term inside `eval_cool` (which already receives `nmol`) with its own `cool_chan` column; that changes the `Cooling_breakdown.txt` schema and the molecular

results, so it belongs in its own change. The Cooling_breakdown header now states that the channel sum falls short of the cool column by that term in the molecular layer. Open item (F) of TO_BE_DONE.md.

- H_2^+ , H_3^+ and HeH^+ donate their electron to n_e but are left out of the bremsstrahlung charge sum, so that sum is smaller than n_e inside the molecular layer. They are $Z_{\text{ion}} = 1$ and belong there formally, but they only exist at $T \sim 10^3$ K, where the hot-plasma free-free expression is an extrapolation emitting at frequencies the atmosphere is not thin to, and where free-free is negligible against the H_3^+ and metal line cooling. The scope is written at the accumulator.

Files:

- src/modules/radiation/util_ion_eq.f90, src/modules/radiation/ionization_equilibrium.f90
- src/modules/time_step/energy_semi_implicit.f90
- src/modules/post_process/post_process_adv.f90 (comments only)

Documentation: docs/lower_atmosphere_coupling.md.

53. A fourth regression case for the molecular + metals system (2026-08-13)

The coupled molecular + metals system of the section before last was not guarded by anything in the matrix: mol_base_handoff runs the molecular network metals-off, and the two wasp_* cases run metals without molecules, so the merged residual, the metal_row_base() offset and the three-starting-point seeding had no golden.

backup/regression/mol_metals is that case: the same hot-Uranus Tier-2 gate configuration as mol_base_handoff (Molecular chemistry: True, its base.inp handoff, 12000 steps pinned in maxsteps) plus a solar-abundance metals.inp with the seven elements C/N/O/Mg/Ca/Na/Fe. It stops at $du = 3.1770$, $\log_{10} \dot{M} = 10.58$. DEFAULT_CASES in backup/regression/run_check.sh is now wasp_full wasp_he23off mol_base_handoff mol_metals, and the case list in the script header describes what each one guards.

Goldens for all four cases were re-snapshotted at the end of this change series, so they carry the two preceding sections. The two atomic cases are byte-identical to their previous goldens (both changes are exact no-ops without molecules); mol_base_handoff moves only in its cooling column, by the factor 176.5 of the previous section.

One harness trap, hit while doing this and now written into the script: golden does **not** re-run the cases, it copies whatever output/ currently sits in each case directory. Snapshotting straight after a code change baselines the *previous* binary. The order is check (which rebuilds and re-runs), then golden, then check again.

Files: backup/regression/run_check.sh, backup/regression/mol_metals/ (new), backup/regression/golden/ (re-snapshotted). Documentation: the workspace CLAUDE.md, README_HOWTO.md.

54. H_3^+ infrared cooling inside eval_cool: one energy equation for a molecular run (2026-08-13)

Closes item (F) of TO_BE_DONE.md, opened by the audit in the section before last.

The H_3^+ infrared cooling was added to cool in ionization_equilibrium, *after* eval_cool returned. Only some of the consumers of cool saw it:

term in cool	rebuilt by	returned by ioniz_eq (JFNK steady residual, the explicit-energy option, the in-loop $\ R\ $ monitor, the Hydro_ioniz.txt cool column)
H/He recombination	yes	yes
collisional ionization (H I, He I, He II, He 2^3S)	yes	yes
collisional excitation (H I Ly α , He I, He II, He 2^3S 10830 + $2^1S/2^1P$ conversion)	yes	yes
bremsstrahlung (H^+ , He^+ , He^{++} , metal ions)	yes	yes
metal line cooling (27 ions, saturated ground-term fine structure included)	yes	yes
H_3^+ infrared (Miller et al. 2013)	NO	yes

heat has no such split: it is assembled once in ioniz_eq (photoheating by each absorber including H_2 and the metals, Shull & van Steenberg secondary-ionization partition, Balmer photoelectric and Ly α de-excitation, He-recombination coupling, He $2^3S + H$ and He $2^3S + H_2$ Penning) and passed into energy_semi_implicit as intent(in), never rebuilt. H_3^+ was the only asymmetric term, and it was the whole asymmetry.

Three consequences followed from that one row.

- **The marching fixed point was not the zero of the residual the Newton solver drives down.** The marching relaxed $\Gamma = \Lambda_{\text{atomic+metal}}$; the residual demanded $\Gamma = \Lambda_{\text{atomic+metal}} + \Lambda_{H_3^+}$.
- **The two energy-solver options integrated different equations.** Energy solver: Explicit used the ioniz_eq array directly and did carry the term; the semi-implicit default (the production setting) did not.
- **The cool column of Hydro_ioniz.txt meant different things depending on how the run stopped.** A du-stop left the semi-implicit array (no H_3^+) in place; the JFNK finish calls ioniz_eq after solve_steady_jfnk and left the array that has it.

Size of the missing term. On the metals-off hot-Uranus Tier-2 gate (backup/regression/mol_base_handoff) the H_3^+ lines carry **more than 99% of the total radiative cooling from the base out to the $H_2 \rightarrow H$ front** (Cooling_breakdown H3p_IR / cool_total > 0.5 over $r = 0.9998\text{--}1.1856$, > 0.99 over $0.9998\text{--}1.1557$), peaking at $9.35 \times 10^{-7} \text{ erg cm}^{-3} \text{ s}^{-1}$ at $r = 1.0413$, which is **$7.3\times$ the local photoheating**. Below the front every atomic channel is exponentially suppressed at $T \sim 1.3 \times 10^3 \text{ K}$, so the marching stage was integrating that layer with essentially no radiative loss at all: the cooling it used was smaller than the true total by a factor 1.1×10^2 to 1.2×10^{11} over $r = 1.00\text{--}1.156$. With solar trace metals (mol_metals) the metal lines dominate instead and H_3^+ reaches only 22% of the local cooling, in a thin band $r = 1.066\text{--}1.193$.

The fix. eval_cool computes the H_3^+ cooling itself, from the nmol array it has taken since the section before last, using the same h3p_cooling module (one definition of the Miller et al. 2013 fits and the Table-6 non-LTE factor;

eval_cool uses it, nothing is re-derived). The post-return addition in ionization_equilibrium is gone, and a comment at the eval_cool call states why nothing may be added there. cool_chan gains column 7 for the channel, so the decomposition still reproduces cool exactly (measured $\max|\sum_{\text{channels}}/\text{cool}-1| = 3.3 \times 10^{-16}$ on the molecular gate, 5.5×10^{-16} on wasp_full). Callers that model a molecule-free gas (post_process_adv, and T_equation which solves for the _adv temperature) omit nmol and therefore still get the atomic charge sum *and* no H_3^+ term — the same approximation, stated at both call sites.

A side effect worth naming: energy_semi_implicit builds its damping derivative dC/dT by finite-differencing eval_cool in temperature, so the steep temperature dependence of the H_3^+ emission now enters the implicit temperature update instead of being invisible to it.

Output schema. Cooling_breakdown.txt gains col11 H3p_IR; the 27 metal-ion columns move from 11–37 to 12–38. It is written for every run and is zero without the molecular network. python/paper_data.py reads the column names from the header and follows automatically; examples/exhale_io.py has a fixed list, which now carries H3p (and is named COOL_GAS_CHANNELS, since it is no longer only the H/He channels).

Atomic runs are exactly unchanged. nmol(:, 3) is zero, the term is skipped cell by cell, and the total is $\dots + 0.0$. make_check is byte-identical on wasp_full (14060 steps, metals + He 2³S) and wasp_he23off (14037 steps), with the same step counts as their goldens.

The two molecular cases change, which is the point. Both were re-run at their pinned 12000-step convention; the goldens were **not** re-snapshotted.

	metals off (mol_base_handoff)	metals on (mol_metals)
du at step 12000	2.6689 $\rightarrow$ 2.6274	3.1770 $\rightarrow$ 3.1769
$\log_{10} \dot{M}$	10.9043 $\rightarrow$ 10.9040	10.9105 $\rightarrow$ 10.9105
$\text{H}_2 \rightarrow \text{H}$ front ($2n_{\text{H}_2}/n_{\text{H}} = 0.5$)	1.16087 $\rightarrow$ 1.16039	1.16246 $\rightarrow$ 1.16246
H_3^+ peak	$7.0237 \times 10^4 \rightarrow 7.0285 \times 10^4 \text{ cm}^{-3}$ at $r = 1.0419$	52.139 $\rightarrow$ 52.140 at $r = 1.1355$
base T	1213.42 $\rightarrow$ 1213.42 K	1213.46 $\rightarrow$ 1213.46 K
max T change, whole grid	+1.4% at $r = 1.93$ (wind transient)	$+5.2 \times 10^{-5}$ at $r = 1.20$
T change where H_3^+ dominates	−0.53% at $r = 1.078$, mean −0.20% over $r \leq 1.156$	-3.0×10^{-5} at $r = 1.183$
reported cooling at the base	$4.94 \times 10^{-18} \rightarrow 5.97 \times 10^{-7} \text{ erg cm}^{-3} \text{ s}^{-1}$ ($\times 1.2 \times 10^{11}$)	unchanged
max cooling change	$\times 1.23 \times 10^{11}$ at $r = 1.0021$	+28% at $r = 1.1597$

Why the gate observables barely move, and why that is not evidence the term is small. The molecular layer is thermally massive: its H_3^+ cooling time $\frac{3}{2}nkT/\Lambda_{\text{H}_3^+}$ is $3.7 \times 10^6 \text{ s}$ at $r = 1.04$ and $2.2 \times 10^7 \text{ s}$ at the base, i.e. **200–2000 sound-crossing times** $t_s = R_0/v_0$. A 12000-step snapshot advances a small fraction of that, so the layer has only begun to respond: T has fallen by 0.1–0.5% through $r = 1.02$ –1.12 and is still falling. The front, the H_3^+ peak and $\dot{M}$ are unchanged to $< 0.1\%$ *at this step budget*, so the recorded Tier-2 gate numbers stand — but the converged molecular

solution now has a different thermal structure than the pre-change code would have produced, and reaching it needs a step budget the gate convention does not provide. When this was written no molecular configuration in the tree had converged (examples/15_molecular had run 135694 steps to $du = 2.9 \times 10^{-2}$), so it was stated as a limit, not a result; the paragraph after next closes it.

What was verified and what was not. The residual/Newton path is unchanged by the move: EXHALE_RESIDUAL=1 on one frozen molecular state gives a **bit-identical** output/residual_profile.txt before and after, so only the marching side moved. What could *not* be shown at the time was the H_3^+ term disappearing from the steady energy residual: on the only molecular states then available — 12000-step relaxation snapshots with a base-breathing transient, volume-weighted $\|R\| \sim 3$ — the cell-by-cell energy residual in the molecular layer is $1.5\text{--}7\times$ larger than $\Lambda_{H_3^+}/q_0$, so the systematic offset is buried in the transient. That check needs a converged molecular run.

Closed 2026-08-13: the check on a converged molecular run. Restarting the metals-off hot-Uranus snapshot with Load IC? True, Solver: Newton $5.0e-2$ and Max steps: 150000 reaches info = 0 after a JFNK finish at step 22556 — the first Newton-grade molecular solution the code has produced. In the 180 cells where H_3^+ carries more than half the cooling (the cells at $r \leq 1.001$, whose residual sets $\|R\|$, excluded) the ratio of the cell-by-cell energy residual to $\Lambda_{H_3^+}/q_0$ falls from **4.42 at the 12000-step snapshot (sign 89+/166–) to 0.70 at the default $\|R\| < 10^{-3}$ solution (173+/79–) to 0.0017 once the same state is re-solved at Resid tol: $1.0e-5$ (25+/79–)**: at the converged solution the residual carries no offset of the size of the H_3^+ term, and no preferred sign.

Two things follow. First, the default $\|R\| < 10^{-3}$ is not a converged molecular run — the norm is set by two or three cells above the base that carry $\sim 10^3$ times the residual of anything else, so the molecular layer is still cooling when the solver stops; Resid tol: $1.0e-5$ is required, and a third solve from that state reproduces it (T to 0.14%, front to five digits, $\|R\| = 6.7 \times 10^{-7}$). Second, an A/B against the pre-change binary isolates the term directly: marching both binaries 12000 steps from the same converged state, the temperature difference $T(\text{before}) - T(\text{after})$ is proportional to the local $\Lambda_{H_3^+}/(\frac{3}{2}nk)$ with a single constant — over the 155 cells where $\Lambda_{H_3^+}$ is within two decades of its peak the ratio has median 1.27 and 10–90% range 1.13–1.32, and over all 187 cells where H_3^+ carries more than half the cooling the sign is right in 181. The whole difference between the two marching energy equations is the H_3^+ term and nothing else. The physics of the converged solution, including the 190–300 K molecular layer, is in docs/lower_atmosphere_coupling.md and in the next section.

An -O1 -fopenmp -g -fcheck=bounds,do,mem build runs both molecular configurations for 300 steps with no runtime check firing.

Files: src/modules/radiation/util_ion_eq.f90, src/modules/radiation/ionization_equilibrium.f90, src/modules/post_process/post_process_adv.f90 (comments only), examples/exhale_io.py, python/paper_data.py, python/make_struct_figures.py. Documentation: README.md, docs/EXHALE_user_manual.tex (section 4.3 rewritten — it still described four H/He channels and columns 9–35 for the metals, which had been wrong since the collisional-excitation split; a radiative-cooling row for H_3^+ added to the atomic-data table), docs/lower_atmosphere_coupling.md, TO_BE_DONE.md.

55. The infrared field of the lower atmosphere (Base IR field) (2026-08-13)

Narrows item (G) of TO_BE_DONE.md.

The infrared coolants of a molecular base were radiating into a vacuum they are not in. The ground-term fine-structure lines carried an escape probability built from the *outward* column only, with the downward direction dismissed as absorbed by the lower atmosphere; the H_3^+ bands carried no transfer at all. Absorption of what that reservoir radiates back was in neither.

Measured first. Line-center depths on the converged metals-on hot-Uranus solution, with the level populations and Doppler widths the cooling module itself uses:

direction	[C I] 609 μm	[C I] 370 μm	[O I] 63 μm	[O I] 145 μm
outward, base cell	0.088	0.172	2.01	0.54
outward, coldest cell ($r = 1.083$)	1.4e-3	2.6e-3	0.028	4.6e-3
one scale height below the base	0.046	0.092	1.08	0.32
pressure where the downward $\tau = 1$	205 μbar	106 μbar	17 μbar	37 μbar

The base sits at 9.0 μbar , so every line that carries the cooling is black downward within one to three scale heights of unresolved atmosphere, while outward the collapsed layer is thin ($\beta \geq 0.90$ for C I, ≥ 0.63 for O I above $r = 1.05$). Trapping was never the missing piece; the incident field was. The ratio of the incident mean intensity to the line source function, $\frac{1}{2}B_\nu(T_{\text{base}})/B_\nu(T)$, passes 1 below $T \sim 610$ K and reaches 3.4–5.7 at the coldest cell: those lines were cooling a gas they should have been heating.

And it refutes the diagnosis in item (G) for the metals-on case. The specific entropy p/ρ^γ rises monotonically outward through the whole collapsed layer, 1.00 at the base to 1.67 at $r = 1.05$ and 5.69 at $r = 1.083$; the gas is net heated everywhere and is cold only because ρ has fallen by 10^2 . Its cooling is 36% of the local photoheating at $r = 1.01$ and 2–16% above $r = 1.02$, and $\Gamma - \Lambda$ is balanced by $\rho v [d\varepsilon/dr + p d(1/\rho)/dr]$ to a few percent. And the metals-off solution, with no fine-structure cooling at all, collapses further (115 K at $r = 1.033$, total cooling 3.8×10^{-12} erg cm $^{-3}$ s $^{-1}$). The metals-on layer is an adiabatic expansion, not a radiative collapse. Radiation does dominate in one place: the metals-off base, where H_3^+ carries 100% of the cooling and the entropy *falls* by a factor 3 between the base and $r = 1.02$.

What was implemented. Base IR field: True (default False) has the lower atmosphere radiate $B_\nu(T_0)$ over the sky fraction $f = 1 - \sqrt{1 - (R_p/r)^2}$ and both coolant families absorb it.

- Fine structure: the field enters *inside* the ground-term statistical equilibrium as the photon occupation number $\bar{n} = \beta_1(\tau_{\text{dn}})f/(e^{E_{\text{int}}/T_0} - 1)$, with radiative rates $\beta A(1 + \bar{n})$ down and $\beta A(g_u/g_l)\bar{n}$ up, so the returned power is the net one. The escape probability becomes two-sided, $\beta = \beta_1(\tau_{\text{up}}) + \beta_1(\tau_{\text{dn}})$, β_1 the single-face Hollenbach & McKee (1979) / de Jong et al. (1980) form; with the field off it is $2\beta_1(\tau_{\text{up}})$, the previous value, bit for bit. Each line stops cooling at its own radiative equilibrium temperature — 575.8 K ([C I] 609 μm) to 641.7 K ([O I] 44 μm) for $T_0 = 1140$ K and half-sky coverage.
- H_3^+ : Miller et al. (2013) fit the *total* emission, so the exchange is closed on one effective band, the ν_2 fundamental at 2521.3 cm $^{-1}$ ($E/k = 3627.5$ K): $\Lambda_{\text{net}} = \Lambda_{\text{emit}}(T)[1 - \bar{n}e^{E/T}]$, equilibrium at 936.1 K at the base. The approximations and their range are written at h3p_net_cooling_rate.

- Not applied, with the reason at the code: Fe II (a precomputed statistical-equilibrium table, which cannot take a field as an argument), the coronal remainders and every other metal ion (fits to CHIANTI sums with no line list; their lowest terms are at $E/k \gtrsim 1.6 \times 10^4$ K, and the heating they would add was estimated at $\sim 0.4\%$ of the losses of this layer), and the legacy use_2lev_cool branch.

Effect. Restarting the two converged Resid tol: $1.0\text{e-}5$ molecular solutions with the switch on, both returning info = 0:

	metals off, off $\rightarrow$ on	metals on, off $\rightarrow$ on
T at $r = 1.005$	626.8 $\rightarrow$ 910.0 K	1099.0 $\rightarrow$ 1098.9 K
T at $r = 1.02$	299.6 $\rightarrow$ 859.4 K	928.9 $\rightarrow$ 928.9 K
T at $r = 1.03$	140.2 $\rightarrow$ 840.2 K	820.0 $\rightarrow$ 820.1 K
coldest cell	115.1 K @ 1.0334 $\rightarrow$ 229.7 K @ 1.0947	187.9 $\rightarrow$ 189.1 K @ 1.0831
H ₂ $\rightarrow$ H front	1.03582 $\rightarrow$ 1.08833	1.07976 $\rightarrow$ 1.07979
$\log_{10} \dot{M}$ [g/s], spherical	10.539 $\rightarrow$ 10.617	10.596 $\rightarrow$ 10.596

The metals-off layer settles *on* the H₃⁺ radiative equilibrium curve: predicted 911/886/874 K at $r = 1.005/1.02/1.03$ with the local sky fraction, solution 910/859/840 K, just below by the margin the expansion takes out. Its H₃⁺ channel is now a heating term (-1.1×10^{-7} erg cm⁻³ s⁻¹ at $r = 1.02$). The solution is a fixed point: re-solving returns info = 0 at $\|R\| = 8.8 \times 10^{-6}$ and reproduces T below $r = 1.3$ to 10^{-4} . The metals-on layer does not move, which is the measurement above, not a failure of the closure.

Atomic runs are unaffected: the switch is off by default, and even on it touches only channels that carry no cooling in an atomic wind. make check is byte-identical on all four cases.

Files: src/modules/radiation/Cool_coeff.f90, src/modules/lower_atmosphere/h3p_cooling.f90, src/modules/radiation/util_ion_eq.f90, src/modules/nonlinear_system_solver/T_equation.f90, src/modules/post_process/post_process_adv.f90, src/modules/init/parameters.f90, src/modules/files_I0/input_read.f90, src/modules/files_I0/write_setup_report.f90. Documentation: docs/lower_atmosphere_coupling.md, docs/EXHALE_user_manual.tex, TO_BE_DONE.md.

56. H₂ photodissociation in the Lyman–Werner bands (Stellar LW flux) (2026-08-13)

The molecular network had no photodissociation of neutral H₂. The Koskinen et al. (2022) Table-1 photo-rates start at the 15.4 eV photoionization edge, so below that the only H₂ losses were thermal (R12), electron impact (R14) and ion chemistry. docs/base_composition_handoff_plan.md §5 named the omission as the reason EXHALE’s base comes out more molecular than any photochemical code gives, and recorded that pinning q_{H_2} at the base (Route 1) patches over it. This adds the missing term.

The band is not in the code’s own radiation field, and had to be. Checked in the source: set_energy_vectors builds the photon grid from 13.6 eV up and reaches below the H I edge only when the He 2³S metastable (4.8 eV) or a low-IP metal is active, and then only to that species’ threshold; read_sed selects SED rows by the same e_low. For

a numerical SED the 11.2–13.6 eV band is therefore not read at all, and for a power-law SED (every Tier-2 gate case) it is the XUV power law extrapolated downward, which has nothing to do with a star’s FUV. Even where the grid does reach into the band, its opacity is continuum photoionization, whereas Lyman–Werner absorption is a forest of saturated lines. The band flux is a separate input,

Stellar LW flux [erg/cm²/s]: 343.0 # 912–1110 Å, integrated, at the planet

default 0 = off, and inert (with a warning) when Molecular chemistry is off.

Rate. Draine & Bertoldi (1996), ApJ 468, 269 (published version). For their flat- F_λ spectrum ($u_\nu \propto \nu^{-2}$, their eq. 24) at $\chi = 1$ the band photon flux is $F = 1.208 \times 10^7 \text{ cm}^{-2} \text{ s}^{-1}$ (their Table 1) and the unshielded dissociation rate is $\zeta_{\text{diss}}(0) = 4.17 \times 10^{-11} \text{ s}^{-1}$ (their Table 2 and Fig. 7 caption), so the rate per band photon is an effective cross section $\sigma_{\text{LW}} = 3.452 \times 10^{-18} \text{ cm}^2$. With the mean photon energy of a flat- F_λ band, $2hc/(912+1110 \text{ Å}) = 12.2635 \text{ eV}$,

$$k_{\text{LW,thin}} = 1.757 \times 10^{-7} F_{\text{LW}} \quad [\text{s}^{-1}, F_{\text{LW}} \text{ in erg cm}^{-2} \text{ s}^{-1}]. \quad (17)$$

Their Table 1 was reproduced numerically from their eqs. (22)–(24) to confirm the flux convention (Habing 1.2220×10^7 against their 1.222×10^7 ; ν^{-2} 1.2084×10^7 against 1.208×10^7 ; Draine 1978 1.2313×10^7 against 1.232×10^7). The band-shape sensitivity is $\pm 25\%$: the same arithmetic for the soft Draine 1978 field gives $2.685 \times 10^{-18} \text{ cm}^2$.

Self-shielding is their eq. (37), the fit to the full multiline calculation *including line overlap*, valid over $10^{14} < N_{\text{H}_2} < 3 \times 10^{21} \text{ cm}^{-2}$, with b the thermal H_2 Doppler parameter $(2kT/m_{\text{H}_2})^{1/2}$. N_{H_2} is the star-ward column, built by the same `calc_column_dens_one` radial integration and `opa_pf` weighting as every other absorber column, from the incoming (pre-solve) H_2 density. There is no dust (EXHALE’s metals are atomic and trace), the trace-metal continuum in the band is $\tau \sim 4 \times 10^{-23} N_{\text{H}}$ and negligible where the shielding matters, and the H Lyman-series lines inside the band are inside DB96’s own fit and not treated separately.

Chemistry and heating. $\text{H}_2 + h\nu \rightarrow \text{H} + \text{H}$ enters the H_2 balance row of the coupled system next to the H_2 photoionization (`mol_heh_rows`, shared by `System_HeH_mol` and `System_HeH_mol_metals`), so it is solved with everything else. Both products are neutral H, which the H-nucleus closure supplies; no other row changes. The fragments carry 0.4 eV of kinetic energy (Black & Dalgarno 1977, ApJS 34, 405, p. 418, “the yield is about 0.4 eV per atom pair”), added to heat. The 4.48 eV bond energy is paid by the photon, not by the gas, and is not a sink of this channel. `output/Lyman_Werner.txt` (r , T , x_{H_2} , n_{H_2} , N_{H_2} , f_{shield} , k_{LW} , heating) is written whenever a molecular run carries a band flux.

Band flux for a planet. Integrated externally from the VULCAN stellar spectra in the tree (EXHALE/VULCAN/atm/stellar_flux/, stellar surface) over 91.2–111.0 nm and diluted by $(R_*/a)^2$. The `VULCAN_run_*` spectra directories are empty since the 2026-08-09 loss, so no HD 209458 spectrum exists here and the solar spectrum stands in: $343 \text{ erg cm}^{-2} \text{ s}^{-1}$ at 0.048 AU behind a $1.155 R_\odot$ star, $k_{\text{LW,thin}} = 6.03 \times 10^{-5} \text{ s}^{-1}$. The same file returns a solar Ly α irradiance of $6.12 \text{ erg cm}^{-2} \text{ s}^{-1}$ at 1 AU against an observed 6–8, which is the calibration check. HD 189733 b comes out at 600 (Moses11 spectrum) or 1830 (Bourrier 2020) $\text{erg cm}^{-2} \text{ s}^{-1}$ — a factor 3 that dwarfs the $\pm 25\%$ of the band shape.

Effect. A/B on the key alone, hot-Uranus gate (metals off, the `mol_base_handoff` configuration restarted with Solver: Newton 5.0e-2, Resid tol: 1.0e-5, Base IR field: True), both returning `info = 0`:

	LW off	LW on (343)
residual norm at exit	8.15e-6	4.86e-6
H ₂ → H front ($x_{\text{H}_2} = 0.5$)	1.08806	1.08622
x_{H_2} at the base	0.99921	0.99883
x_{H_2} at $r = 1.10$	2.54e-2	1.51e-4
H ₃ ⁺ peak [cm ⁻³]	3.19e5 @ 1.0079	3.20e5 @ 1.0025
$\log_{10} \dot{M}$ [g/s], spherical	10.6196	10.6248

and on HD 209458 b (examples/15_molecular + Base IR field: True, both members at Resid tol: $2.0\text{e-}5$ because the band-on run stalls the JFNK at 1.5×10^{-5} on the base-adjacent momentum row, item (A) of TO_BE_DONE.md; both info = 0, both mass fluxes flat to 4.1×10^{-3}): front 1.01578 → 1.01290, x_{H_2} at $r = 1.02$ $2.37 \times 10^{-4} \rightarrow 4.22 \times 10^{-6}$, $\log_{10} \dot{M}$ 10.5711 → 10.5600.

The self-shielding does exactly what it should: at the hot-Uranus base $N_{\text{H}_2} = 4.3 \times 10^{21} \text{ cm}^{-2}$ gives $f_{\text{shield}} = 9.8 \times 10^{-7}$ and $k_{\text{LW}} = 5.9 \times 10^{-11} \text{ s}^{-1}$, while above $r = 1.15$ the column has fallen below 10^{10} cm^{-2} , $f_{\text{shield}} = 1.000$ and k_{LW} recovers the unshielded $6.026 \times 10^{-5} \text{ s}^{-1}$.

It does not close the base-composition gap, and the measurement says why. The base q_{H_2} moves from 0.8618 to 0.8612 against the 0.75 the case's base.inp carries — 0.5% of the gap — even though Lyman–Werner becomes the largest single H₂ loss at the base ($5.9 \times 10^{-11} \text{ s}^{-1}$ against 2.5×10^{-11} for H⁺ + H₂). The two facts are consistent because the base partition is a formation–destruction balance in which the three-body R15 sets $n_{\text{H}_2} \propto n_{\text{H}}^2$, so $n_{\text{H}}/n_{\text{H}_2}$ scales as the square root of the destruction rate: the measured $2.3\times$ rise in destruction raises the atomic fraction only $1.5\times$. Reaching $q_{\text{H}_2} = 0.75$ needs a destruction rate $4100\times$ larger. The unshielded band would supply it many times over; the column stops it. The extra atomic H photochemical codes find at $1 \mu\text{bar}$ comes from catalytic cycles on O, OH and H₂O, and EXHALE's H₂/H₂⁺/H₃⁺/HeH⁺ network has nowhere to put them. Route 1 (q_H2_base) therefore stays the only way EXHALE can carry a photochemical base partition, and is now a documented modeling choice rather than a patch over a missing term.

make check is byte-identical on all four cases with the key absent.

Files, by directory under src/modules/: lower_atmosphere/ — new lyman_werner.f90, mol_rates.f90; init/ — parameters.f90; files_IO/ — input_read.f90, write_setup_report.f90, write_output.f90; nonlinear_system_solver/ — ion_cell_state.f90, System_HeH_mol.f90, System_HeH_mol_metals.f90; radiation/ — ionization_equilibrium.f90, util_ion_eq.f90; plus Makefile. Documentation: docs/lower_atmosphere_coupling.md, docs/base_composition_handoff_plan.md, docs/EXHALE_user_manual.tex, docs/input_schema.md, README.md, README_HOWTO.md, examples/exhale_io.py.

57. A fifth regression case for Lyman–Werner photodissociation (2026-08-13)

Nothing in the four-case matrix exercised the code the section 56. key turns on: with Stellar LW flux absent every case leaves $F_{\text{LW_star}} = 0$, so the star-ward H_2 column, the Draine & Bertoldi self-shielding factor, the LW loss row in the H_2 budget, the 0.4 eV fragment heating and the Lyman_Werner.txt writer all ran in no case.

backup/regression/mol_lyman_werner is that case: a copy of mol_base_handoff (the hot-Uranus Tier-2 gate, Molecular chemistry: True, the same base.inp handoff, 12000 steps pinned in maxsteps) plus the single line Stellar LW flux [erg/cm2/s]: 343.0 — the HD 209458 b solar-proxy band flux of the preceding section. It stops at $du = 2.1841$, $dtu = 4.8072 \times 10^{-3}$, clearly apart from the LW-off mol_base_handoff ($du = 2.6274$, $dtu = 5.7385 \times 10^{-3}$): the case does pass through the new code path. DEFAULT_CASES in backup/regression/run_check.sh is now wasp_full wasp_he23off mol_base_handoff mol_metals mol_lyman_werner, and the script header describes what the case guards.

The golden was taken in the order the fourth-case section's trap requires — check (which rebuilds and re-runs the case), then golden, then the full check again — and the final run is byte-identical on all five cases.

Files: backup/regression/run_check.sh, backup/regression/mol_lyman_werner/ (new), backup/regression/golden/mol_lyman_werner/ (new). Documentation: the workspace CLAUDE.md, README_HOWTO.md.

58. JFNK convergence is judged on the iterate the solver returns; load_IC no longer accepts an all-zero metal column (2026-08-13)

Two things came out of trying to converge HD 209458 b with molecular chemistry AND solar C/N/O (examples/16_molecular_metals), which fails with the prescription that works for its metals-off twin examples/15_molecular.

The diagnosis is in docs/hd209_metal_stagnation.md; what follows is what changed in the code, and one thing that was measured and deliberately *not* changed.

The measurement. solve_steady_jfnk accepts a step on the whole-domain scaled merit $\|F/D\|_2$ over cells $1 \dots N$ and decides info on $\|R\|$, the volume-weighted relative residual over the escape region $[j_{\min} : N]$ alone. On this configuration the two move in opposite directions. Warm-restarted from its own stalled state, the solve reached $\|R\| = 3.958 \times 10^{-5}$ at outer iteration 2 and then, over 115 further iterations, cut the merit from 5.67×10^{-1} to 1.58×10^{-2} — a factor 36 — while $\|R\|$ rose to 9.5×10^{-5} . The reason is a shell just above the base that stops flowing once the metal cooling is on: at $r = 1.02 R_p$ cooling exceeds heating by a factor 9.7 with C/N/O against heating exceeding cooling by a factor 33 without them, the mean speed over $r = 1.015\text{--}1.030$ falls from 207 to 6–24 cm s^{-1} , and an undamped $2dr$ contact mode grows there whose amplitude reaches several times the mean velocity. That shell dominates the whole-domain merit and is invisible to $\|R\|$.

Aligning the merit's region was tried and rejected, for the second time. Replacing $\|F/D\|_2$ over $1 \dots N$ by the volume-weighted RMS of F/D over $[j_{\min}:N]$ breaks solves that converge today: examples/15 goes from `info = 0` at $\|R\| = 9.542 \times 10^{-6}$ in 272 outer iterations to `info = 2` with a best of 1.171×10^{-4} in 30, and the HD 209458 b base.inp continuation photo_deep_secion_cont from `info = 0` at 7.284×10^{-4} in 4 iterations to `info = 2` at 3.249×10^{-3} in 15. This reproduces on two more cases the measurement that rejected the same change on 2026-08-11 (docs/newton_scaling_and_base_wall.md §4), and for the reason given there: the Newton step is computed from the full residual over all rows, so a merit that ignores most of those rows rejects the steps that step takes. The merit stays whole-domain, and the code now says so at the site with the numbers.

What did change: info is decided on the iterate that is returned. The solver already tracked the best iterate by $\|R\|$, already returned it when the solve failed, and already tested $\|R\| < \text{resid_tol}$ at the top of every outer iteration — so an iterate that satisfies the tolerance is normally caught the moment it is produced. Two gaps remained. The restore of the best iterate was gated on `info  $\neq$  0`, and the tolerance was never tested after the loop, so a solve whose last allowed iteration met the tolerance fell out with `info = 1`. Both close with one line, `if (rnorm < resid_tol) info = 0` after the loop, with the best-iterate restore made unconditional. It matters because EXHALE_main treats `info  $\neq$  0` as a failure, disables the Newton finish for the rest of the run, and falls back to a du-stopped march — discarding a state that satisfied the convergence criterion.

Scope, measured, not inferred: this changes no trajectory, and it is not what unlocks examples/16. The stalled examples/16 warm restart was re-run at `Resid tol: 5.0e-5` with the binary before and the binary after: both hand off at marching step 9363, both take 2 outer iterations, both print the same two iterate lines to every digit, and both end `info = 0` at $\|R\| = 3.958 \times 10^{-5}$. The examples/16 failures at `Resid tol: 1.0e-5` are therefore not a bookkeeping artifact either — that configuration genuinely does not reach 10^{-5} , for the reasons in the memo. What the fix removes is the remaining way for a returned state that satisfies the tolerance to be labelled a failure.

load_IC no longer accepts a metal column that is present but zero. It decided whether the restart file carried an element from the presence of its columns alone. The schema-2 writer emits the metal columns unconditionally, so a metals-OFF output carries them with the value zero; a metals-ON restart from such a file therefore ran with **zero metal density everywhere** while the base boundary condition still counted metals in the mass and particle budget under eos_metals. It was silent, and it produced confident nonsense: restarting the converged examples/15 (metals off) solution as a metals-on run reported $\|R\| = 9.54 \times 10^{-6}$ — it was reproducing the metals-off solution under a metals-on label. An element is now taken from the file only if every stage has a column AND the loaded stages are not identically zero everywhere; otherwise it is built from the abundance exactly as `set_IC` does, $n_X = \text{melem_ab} \cdot n_H$ all in the neutral stage, and the substitution is reported on stdout and in EXHALE_setup.out. The rebuilt density is inserted before `calc_rho`, so the restart keeps the mass closure $\sum_i f_i A_i = 1$ that a cold start has by construction; for the same reason the legacy headerless-file path now carries its metal mass in ρ as well. *Measured* on that restart: the three elements are reported, the base density at $j = 1$ goes from 2.0618×10^{14} to $2.0804 \times 10^{14} \text{ m}_H \text{ cm}^{-3}$ (0.9 %, the metal mass), and the reported $\|R\|$ goes from the false 9.54×10^{-6} to 7.36×10^{-2} — the honest distance of a metals-off wind from the metals-on steady state. A metals-off restart, and a restart whose file does carry metal densities, both give a

byte-identical residual profile to before.

Gates. make check is 5/5 byte-identical (no regression case reaches the JFNK hand-off and none uses Load IC, so neither change can touch them). examples/15 re-run from cold with the changed binary reproduces its reference output **byte-identically** — info = 0, $\|R\| = 9.542 \times 10^{-6}$, 272 outer iterations, 25837 marching steps, $\log_{10} \dot{M} = 10.23$, $H_2 = H\ I$ front at $1.0084 R_p$. The hot-Uranus molecular + solar-metals configuration (backup/regression/mol_metals run to the Newton finish instead of its 12000-step regression snapshot) converges: info = 0, $\|R\| = 3.342 \times 10^{-6}$ in 23 outer iterations, 34537 marching steps, $\log_{10} \dot{M} = 10.30$. photo_deep_secion_cont converges: info = 0, $\|R\| = 7.284 \times 10^{-4}$ in 4 iterations, $\log_{10} \dot{M} = 9.69$.

And examples/16 itself now has a Newton-grade solution on the warm path. Restarted from the converged examples/15 state — exactly the restart the load_IC fix repairs — with Resid tol: $5.0e-5$ it converges: info = 0, $\|R\| = 3.337 \times 10^{-5}$ in 30 outer iterations, $\log_{10} \dot{M} = 9.65$. At 10^{-5} it does not, and the cold path converges at neither tolerance — 112920 marching steps to the hand-off against 25837 for the metals-off twin, then 281 outer iterations to a best $\|R\|$ of 3.623×10^{-4} with the worst scaled residual at $r = 1.021$, inside the stagnant shell. Both remain open, in docs/hd209_metal_stagnation.md.

Files: src/modules/time_step/steady_newton.f90, src/modules/files_I0/load_IC.f90, src/modules/files_I0/write_setup_report.f90. Documentation: docs/hd209_metal_stagnation.md (new), TO_BE_DONE.md item (A).

59. The molecular equilibrium no longer inherits a cell's previous composition (2026-08-13)

ioniz_eq solves the ionization equilibrium — and, under Molecular chemistry: True, the dissociation equilibrium with it — cell by cell, and the composition it returns is what the cooling, the photoionization columns and the steady residual are all built on. When the coupled molecular network produced no root inside the physical simplex in a cell, it used to hand that cell its own *incoming* fractions back and warn (molecular equilibrium failed at N cells (kept previous state)).

That is wrong on its own terms. The composition returned is then not the equilibrium of the cell at the state being evaluated; it is whatever the previous evaluation happened to leave there, and the error is not small. In the JFNK finish it also means that the residual changes definition between outer iterations: f_sp is overwritten by each accepted trial, so the map the next iteration differentiates and line-searches differs from the previous one, by a finite amount, at exactly those cells (docs/hd209_metal_stagnation.md section 3, where the same section corrects an earlier claim — within one outer iteration the Jacobian probes and the line-search trials do all start from the same composition).

What the cells were failing on. Measured with a counting build on the HD 209458 b molecular + solar C/N/O warm restart (40 marching steps, single-threaded): 224 cell solves ended below the solver tolerance, 223 of them keeping a physical but unconverged root and 1 inheriting its previous composition. The rejected roots were not on another branch. At step 1, cell $j = 117$ ($r = 1.0294$, $T = 1652$ K), the first attempt converged (info = 1) to a root violating the simplex

only by $x(\text{He II}) = -3.8 \times 10^{-6}$ and $x(\text{H}_2) = -4.5 \times 10^{-4}$, and the molecular-basin retry returned a strongly molecular root, $x(\text{H}_2) = 0.998$, violating it only by $x(\text{He II}) = -1.2 \times 10^{-8}$ — round-off around a fully neutral helium. What fails is a root resolved onto a *face* of the simplex, not an unphysical branch.

Two changes, both in `ionization_equilibrium.f90`.

1. *The molecular-basin retry seed is now a state.* It used to combine the chemical-equilibrium H_2 fraction at the local (p, T) with the **previous state's** H II, He II, He III and 2^3S fractions — a pair that can put more H nuclei into H_2 and H II together than the cell has, i.e. a starting point that is not a composition at all. The new seed, `dissociation_ionization_balance_at_fixed_ne`, is the limit of the network in which the couplings to the molecular ions are dropped: the H nuclei split between H_2 and atomic H by the same Koskinen et al. (2022) Eq. 11 mixing ratio the molecular base boundary condition uses, the atomic remainder and every other element in their own ionization balance at the incoming n_e , $\text{H}_2^+/\text{H}_3^+/\text{HeH}^+$ at zero. It is inside the simplex by construction and depends only on the local state.
2. *A cell whose roots all left the simplex keeps the closest one, clamped.* `element_budget_violation` measures how far a root lies outside the allowed states (the largest negative stage fraction, and the largest amount by which one element's tracked stages exceed its nuclei); the least-offending root is kept and `clamp_fractions_to_element_budget` moves it onto the boundary — negative fractions to zero, and any element still over its nuclei scaled to sum exactly to them. Given the measurement above this is the treatment of a root the solver has resolved onto a face and delivered as a small number of either sign: the clamped state is that root, to the accuracy the solve reached. The cell's previous composition is never used.

The warning is now every molecular equilibrium root left the physical simplex at N cell(s), clamped onto the element budget, and the run-wide count is reported at the end of the run next to the `ioniz-eq roots` line. A cell that reaches the solver tolerance on a physical root takes exactly the path it took before, so the change is inert wherever the solve is healthy.

One thing was tried and reverted. `ionization_fractions_physical` rejects a root for a component below -10^{-10} , and the rejected roots above sit at -10^{-8} , inside the solver's own $\text{xtol} = \sqrt{\epsilon_{\text{mach}}}$. Widening the test to that band is defensible on paper and worse in practice: it lets a cell stop on a root just outside a face instead of retrying from another starting point, and the retried root is the better one. With the band widened, `examples/15` goes from `info = 0` at $\|R\| = 6.105 \times 10^{-6}$ in 89 outer iterations to `info = 2` at 1.112×10^{-4} in 133. The band stays at 10^{-10} and the measurement is recorded at the site.

Gates. `make check` is 4/5 byte-identical. The fifth, `mol_metals`, moves: one cell in the 12000-step run takes a retry it did not take before (`ioniz-eq roots` restarts $9 \rightarrow 10$), and the difference that survives to the output is at round-off — at most 1.3×10^{-11} relative in every `Hydro_ioniz.txt` column and in every H, He and metal density, 1.1×10^{-7} in H_2 and H_2^+ , and 2.3×10^{-4} in H_3^+ in the far wind where H_3^+ is 10^{-17} cm^{-3} against a peak of 52. $\log_{10} \dot{M}$ is 10.58 before and after. The golden was re-snapshotted for that case only, in the order `check → golden → check`. No molecular regression case reaches the clamp ($\emptyset$ cell(s) in all three).

examples/15 (HD 209458 b, molecular, cold) gets *better*, not merely different: $\text{info} = 0$ at $\|R\| = 6.105 \times 10^{-6}$ in 89 outer iterations against $\text{info} = 0$ at 9.542×10^{-6} in 272, hand-off at marching step 25839 against 25837, $\log_{10} \dot{M} = 10.24$ against 10.23, $H_2 = HI$ front at $1.0098 R_p$ against 1.0084. The hot-Uranus molecular + solar-metals configuration holds: $\text{info} = 0$ at $\|R\| = 5.826 \times 10^{-6}$ in 23 outer iterations against 3.342×10^{-6} in 23, same 34537 marching steps, $\log_{10} \dot{M} = 10.30$ unchanged. Runtime checks (`-O1 -fcheck=bounds,do,mem`) are clean on `mol_metals`, `mol_base_handoff` and on the warm restart that does exercise the clamp.

examples/16 (HD 209458 b, molecular AND solar C/N/O), the case the memo is about, does **not** improve — it gets worse. The warm restart that reached $\text{info} = 0$ at $\|R\| = 3.337 \times 10^{-5}$ in 30 outer iterations now bottoms at 7.931×10^{-5} in 67 and reports $\text{info} = 2$, identically at `Resid tol` 5×10^{-5} , 2×10^{-5} and 10^{-5} ; the cold path bottoms at 1.081×10^{-3} against 5.038×10^{-4} before. Each half of the change on its own converges that restart — retry seed alone at 3.059×10^{-5} , clamp alone at 4.853×10^{-5} — and each half on its own fails examples/15, which the two together converge in a third of the iterations it used to take. Both HD 209458 b cases are marginal solves and an $O(1)$ change to a handful of cells' composition moves them either way; what does not move is that their worst scaled residual sits in the stagnant shell at $r = 1.021$ throughout (docs/hd209_metal_stagnation.md section 6). Removing the inheritance is justified on its own terms, not by what it does to these two solves.

Files: `src/modules/radiation/ionization_equilibrium.f90`, `src/EXHALE_main.f90`. Documentation: docs/hd209_metal_stagnation.md sections 3 and 5, TO_BE_DONE.md item (A).

60. Damping the stagnant layer: a gated fourth difference inside the numerical flux (Low-Mach damping) (2026-08-13)

The HLLC flux resolves the middle wave of the Riemann fan exactly. The jump that wave carries — the entropy jump, at constant pressure and velocity — moves at the contact speed $S^* \simeq v$, and the dissipation the solver applies to it is proportional to $|S^*|$. It therefore vanishes as the flow stagnates. The acoustic families keep their damping at $|v| \pm c_s$, so a stationary cell-to-cell entropy pattern is a discretely undamped mode of the scheme, and in a near-hydrostatic layer the velocity field is tied to it by continuity and by the force balance, so it carries the same pattern.

HD 209458 b with molecular chemistry and solar C/N/O is where that stops being academic. The metal cooling exceeds the photoionization heating by a factor ~ 10 at $r \simeq 1.02 R_p$ and stops the flow: the mean $|v|$ over $r = 1.015$ – 1.030 falls to 6 – 24 cm s^{-1} against $\sim 207 \text{ cm s}^{-1}$ in the metals-off twin ($M \sim 3 \times 10^{-5}$ against a sound speed of $\sim 4.5 \text{ km s}^{-1}$), and the alternating component of v grows to 4.7–10 times the local mean $|v|$ where the metals-off solution carries 0.001. Every failed steady solve of that configuration puts its worst scaled residual in that shell (docs/hd209_metal_stagnation.md sections 1 and 6).

The term. Low-Mach damping: `<eps> [M_th]` adds a fourth-difference (Jameson, Schmidt & Turkel 1981) stress to the numerical **momentum** flux at each interior face,

$$D_p^{j+1/2} = \varepsilon_4 g(M) \rho_f \lambda_f (v_{j+2} - 3v_{j+1} + 3v_j - v_{j-1}), \quad \lambda_f = \frac{1}{2} [(|v| + c_s)_j + (|v| + c_s)_{j+1}], \quad (18)$$

and, because a stress does work, $D_E^{j+1/2} = v_f D_p^{j+1/2}$ to the **total energy** flux — the same pairing viscous_conduction uses for the physical Navier–Stokes stress. Without it the kinetic energy the stress removes would come out of the internal energy instead of being converted into it. With it the pair conserves total energy exactly, and the internal-energy change is $-D_p dv/dr$, which has either sign cell by cell but integrates to $+\int \mu_4 (d^2 v/dr^2)^2 dr \geq 0$: kinetic energy into heat, never the reverse. The undivided third difference is $\Delta r^3 d^3 v/dr^3 + O(\Delta r^4)$, so the stress is $\tau = -\mu_4 d^3 v/dr^3$ with $\mu_4 = \varepsilon_4 g \rho c_s \Delta r^3$ — the acoustic momentum diffusivity $\rho c_s \Delta r$ an upwind scheme already carries, reapplied at fourth order.

The mass flux is not touched, and neither is the energy flux other than through the work term. The plain JST form, a fourth difference on the conserved variables themselves, is inadmissible here and was measured to be so: the physical mass flux ρv and energy flux $v(E+p)$ both vanish with the velocity while $\Delta r^3 d^3 \rho/dr^3$ and $\Delta r^3 d^3 E/dr^3$ do not, because ρ and E are stratified over a barely resolved scale height. Their dissipative divergences come out 10^4 – 10^5 times the physical mass-flux divergence and $\sim 5 \times 10^3$ times the radiative source — they would rewrite the base velocity and thermal structure rather than damp an oscillation. The momentum flux $\rho v^2 + p$ does not vanish at stagnation (it tends to p), which is why the momentum form stays small there.

The gate.

$$g = [\max(0, 1 - M_f^2/M_{th}^2)]^2, \quad M_f^2 = \frac{1}{2} (M_j^2 + M_{j+1}^2) \quad (19)$$

is 1 at $M = 0$, exactly 0 for $M_f \geq M_{th}$, and C^1 at the threshold so the steady residual stays differentiable for the Newton solve. It is written in M^2 rather than $|v|/c_s$ because $|v|$ has a kink at $v = 0$ and the stagnant layer is exactly where v changes sign.

Discrete consistency is the point of the design. The stress is added inside RK_rhs, the one routine the marching loop and assemble_residual — hence the JFNK steady solver — both call. The equation the Newton residual measures is the equation the marching loop relaxes, and the fixed point of one is the zero of the other. This is exactly what the Shapiro filter lacks: it is applied to the marching state only, so the Newton residual never sees it and the mode it suppresses during marching is still an undamped mode of the system Newton solves.

Both fluxes are set to zero at the base face ($j = 0$) and the outer face ($j = N$), so nothing is injected or removed through the boundaries; those are also the faces whose four-cell stencil would reach outside the ghost layer. Cell j then depends on cells $j-2 \dots j+2$, which is exactly the stencil the WENO3 residual already has (WL(:, j) reads cells $j-1 \dots j+1$, WR(:, j) cells $j \dots j+2$), so the banded Jacobian bandwidth kl_jac = ku_jac = 8 is unchanged.

Stability. On the $2\Delta r$ mode $v_j = (-1)^j a$ the fourth difference is $16a$, so the mode decays at $16\varepsilon_4 \lambda/\Delta r$; the explicit step is $\Delta t = \text{CFL} \Delta r/\lambda$, so it is damped by $16\varepsilon_4 \text{CFL}$ per step and explicit stability requires $\varepsilon_4 < 1/(16\text{CFL})$ — 0.10 at the default CFL = 0.6. input_read warns when the bound is violated. The classical JST range $1/64$ – $1/32$ sits inside it.

How big the term actually gets. A numerical dissipation is admissible only where it is negligible against the physical fluxes, and neither the fourth difference nor the gate is a bound on that by itself, so any run that uses the key now reports

it: `contact_mode_dissipation_magnitude` prints the peak $|D_p|/|\rho v^2 + p|$ over the interior faces, where it occurs, and the outermost face at which the gate is still open. Measured on the converged profiles (recomputing the same quantities in cgs from `Hydro_ioniz.txt`, independently of the Fortran, agrees to the printed digits):

run	ε_4	peak $ D_p / \rho v^2 + p $	at	gate open to	first $M > 0.1$
examples/15 (metals off)	2×10^{-2}	1.83×10^{-4}	1.0005	$1.043 R_p$	$1.76 R_p$
examples/16 warm	1×10^{-2}	8.36×10^{-5}	1.0005	$1.108 R_p$	$1.99 R_p$
examples/16 warm	5×10^{-3}	3.94×10^{-5}	1.0005	$1.105 R_p$	$1.99 R_p$
hot Uranus (mol. + metals)	2×10^{-2}	3.70×10^{-5}	1.0005	$1.083 R_p$	—

The peak is at the first interior face, where the base boundary condition forces the steepest velocity gradient in the domain, and it is 0.02% or less of the physical momentum flux there; through the stagnant shell itself it is 10^{-7} or below. The first cell with $M > 0.1$ is at 1.76 – $2.00 R_p$, so the wind, the $r_{\text{esc}} = 2 R_p$ window over which $\|R\|$ is measured and the $N - 20$ cell at which $\dot{M}$ is evaluated all sit where the term is identically zero.

Gates. `make check` is **5/5 byte-identical** (`wasp_full`, `wasp_he23off`, `mol_base_handoff`, `mol_metals`, `mol_lyman_werner`); no golden was touched, and none should have been, since none of those inputs carries the key. A `-O1 -fcheck=bounds,do,mem` build is clean on two key-ON cases (the HD 209458 b molecular + C/N/O warm restart, which exercises the open gate, and the hot-Uranus cold start, which exercises the closed one).

What it does to the case it was written for. HD 209458 b with molecular chemistry *and* solar C/N/O, warm-started from the converged examples/15 state, `Resid tol: 1.0e-5`, `Solver: Newton 5.0e-2`, `Max steps: 150000`, same binary and same restart files in every row:

ε_4	hand-off step	info	best $\ R\ $	outer it	$\log_{10} \dot{M}$	$\text{H}_2 = \text{H I front}$
off	48060	1	5.216×10^{-5}	500 (cap)	—	—
5.0×10^{-3}	48062	0	9.934×10^{-6}	162	9.65	1.0112
1.0×10^{-2}	48007	0	9.987×10^{-6}	202	9.67	1.0117
2.0×10^{-2}	65836	2	2.529×10^{-3}	237	—	—
4.0×10^{-2}	66015	2	2.870×10^{-3}	66	—	—

This is the first `info = 0` at `Resid tol: 1.0e-5` this configuration has produced. The two converged rows stop at 9.9×10^{-6} because the solver tests the tolerance at the top of each outer iteration and exits on the first iterate below it, so those are first crossings and not a knife edge; the key-off control, from the same restart, spent its whole 500-iteration budget and plateaued at 5.2×10^{-5} . Nothing diverges in the rows that do not converge: each falls back to time-marching and runs out the 150000-step cap at a *du* plateau of 1.0×10^{-3} to 2.6×10^{-3} , $\log_{10} \dot{M}$ 9.58–9.66. Those are marching states, not Newton-grade ones, which is why the table leaves their $\dot{M}$ blank.

The coefficient has a working range, and it is the bottom of the classical one. At $\varepsilon_4 \geq 2 \times 10^{-2}$ this solve is *worse*, and the reason shows up before the Newton phase: the stress changes the base solution enough to slow the *du* descent (*du* at step 40000 is 0.139 with the key off and 0.360 at $\varepsilon_4 = 2 \times 10^{-2}$), so the hand-off comes 18000 steps later from a

different state. The key is opt-in and the pair has to be run on any new configuration; 2×10^{-2} is not a recommended default, and there is no default.

Sensitivity, honestly. Halving ε_4 from 10^{-2} to 5×10^{-3} moves $\log_{10} \dot{M}$ by 0.02 and the $H_2 = HI$ front by $0.0005 R_p$, and in the wind ($r > 1.2 R_p$) ρ by $\leq 7.4\%$, v by $\leq 3.7\%$ and T by $\leq 0.9\%$. Inside the stagnant layer the two converged states differ a great deal (ρ by 137% at $r = 1.013$, T by 38% at $r = 1.015$) and differ in what the layer looks like: the alternating component of v over $r = 1.015$ – 1.030 is 0.024 at $\varepsilon_4 = 10^{-2}$ and 4.29 at 5×10^{-3} . Both satisfy $\|R\| < 10^{-5}$ on the wind window, which is the point — the convergence measure does not look at that layer, so it cannot pin it down.

The configurations that already converge do not move. Cold starts, Resid tol: $1.0e-5$, $\varepsilon_4 = 2 \times 10^{-2}$ where on:

case	key	hand-off	info	$\ R\ $	outer it	$\log_{10} \dot{M}$	front
examples/15	off	25839	0	6.105×10^{-6}	89	10.24	1.0098
examples/15	on	25839	0	3.642×10^{-6}	100	10.24	1.0095
hot Uranus	off	34537	0	5.826×10^{-6}	23	10.30	1.0774
hot Uranus	on	34537	0	6.252×10^{-6}	23	10.30	1.0774

Both hand off at the same marching step and reach the same $\dot{M}$; on the hot Uranus the entire profile moves by at most 7.3×10^{-4} relative in ρ and 7.1×10^{-4} in T . On examples/15 the wind above $1.2 R_p$ moves by at most 1.2% in ρ , 0.5% in v and 0.16% in T ; the first cells above the base move more (ρ up to 16%, T 6.9% below $1.05 R_p$), and those are cells that carry an odd-even velocity oscillation 86 times the local mean with the key *off*. The cold examples/16 start (Resid tol: $5.0e-5$) improves and still fails: info = 1 at 9.307×10^{-4} with $\varepsilon_4 = 2 \times 10^{-2}$ against 1.081×10^{-3} with the key off, from the identical hand-off step 112920 — the gate is closed through the cold relaxation, when the flow is still fast everywhere, and the term only acts at the end.

Files: src/modules/flux/low_mach_dissipation.f90 (new), src/modules/time_step/RK_rhs.f90, src/modules/init/parameters.f90, src/modules/files_IO/input_read.f90, src/modules/files_IO/write_setup_report.f90, src/EXHALE_main.f90, Makefile. Documentation: docs/hd209_metal_stagnation.md sections 7 and 8, docs/EXHALE_user_manual.tex, docs/input_schema.md K31b, README.md, TO_BE_DONE.md item (A).

61. The residual below the escape radius, and what closes the HD 209458 b molecular + metals case (2026-08-15)

Section 60 left the examples/16 configuration converging from a warm restart at two damping coefficients and failing from a cold one, and left open whether the acceptance window should be widened to include the layer the failures live in. Both questions, and the seed the warm prescription depends on, were measured on 2026-08-15. The full account is docs/hd209_metal_stagnation.md section 9; this section records the code change and the conclusions.

The change: the layer's residual is printed, not tested. resid_relnorm_below_escape and write_resid_below_escape in src/modules/time_step/steady_newton.f90 apply the **same** volume-weighted

relative norm that `resid_relnorm` uses for the convergence measure to the complementary region `[1:j_min-1]`, i.e. below the escape radius, and locate the cell carrying the largest cell-wise relative residual $|F_{k,j}|/|u_{k,j}|$ there, reporting j , $r(j)$ and which of the three rows it is. One line is printed at four points: PTC start, PTC end, JFNK start, JFNK end. The JFNK end line is emitted *before* the best-iterate restore, because that is the last point at which F and u are a consistent pair; it therefore refers to the last iterate visited, which is also the returned state whenever no restore happens.

The value is never tested and nothing in the solve depends on it. The whole change is **+95 lines in that one file**, output only, and `make check` is **5/5 byte-identical** (`wasp_full`, `wasp_he23off`, `mol_base_handoff`, `mol_metals`, `mol_lyman_werner`) — no output file moves, only stdout grows.

Widening the acceptance window is rejected by what it prints. With the diagnostic in place the question is a measurement rather than a design argument. The norm over `[1:401]` ($j_{\min} = 402$):

state	norm below r_{esc}	worst cell
examples/16 converged, $\varepsilon_4 = 5 \times 10^{-3}$	0.195	$j = 75, r = 1.0153$
examples/16 converged, $\varepsilon_4 = 1 \times 10^{-2}$	2.352	$j = 5, r = 1.0010$
examples/16 converged, $5 \times 10^{-3} + \text{LW}$	0.939	$j = 3$
examples/15 converged (metals off)	9.99	$j = 2$
warm restart, at JFNK start	0.60–1.21	$j = 6\text{--}15$, momentum
cold $\varepsilon_4 = 5 \times 10^{-3}$, at failure	594	$j = 97, r = 1.0218$
cold $\varepsilon_4 = 1 \times 10^{-2}$, at failure	1599	$j = 112, r = 1.0273$

A converged state carries a relative residual of 0.2 to 10 below the escape radius, four to six orders above `Resid tol`, and the largest of them belongs to the metals-off `examples/15` solution that nobody doubts. That is not about metals: the layer is near-hydrostatic ($|v| \sim 10 \text{ cm s}^{-1}$), so the denominator of the momentum row’s relative residual, $|\rho v|$, collapses wherever the flow is slow, right state or not. Widening the window on this norm would report `info = 2` for every configuration that converges today without determining the layer any better. The window stays `[j_min:N]`; the diagnostic makes the uncontrolled part visible instead. The failure rows sit three orders above the converged ones, so the number does separate good from bad — it is the absolute scale that is not usable as a tolerance.

The cold path is closed. Scanning Low-Mach damping on the cold start gives `info = 1` at 1.081×10^{-3} (off), 9.953×10^{-4} (5×10^{-3}), 1.025×10^{-3} (1×10^{-2}) and 9.307×10^{-4} (2×10^{-2}) — one $\sim 10^{-3}$ floor, a factor 100 above the tolerance, and the two coefficients that work on the warm restart do nothing here that off does not. The cold hand-off state appears to lie outside the Newton basin the warm restart reaches, and the coefficient is not the lever. The cold path for this configuration is abandoned.

The warm prescription is pinned to its seed. It needs one thing section 60 did not say: the restart must come from an `examples/15` state converged with the **current** binary. Seed A is the copy `examples/15_molecular/output` held until 2026-08-15, written before the section 59 change (kept as `output_pre_item3_20260813`); seed B is `examples/15 re-`

converged with the current binary, which the example directory holds since 2026-08-15. Same binary and same keys otherwise:

seed	ε_4	hand-off	info	best $\ R\ $	$\log_{10} \dot{M}$
A (stored)	5.0×10^{-3}	50645	2	1.303×10^{-4}	—
A (stored)	1.0×10^{-2}	50629	2	5.770×10^{-5}	—
A + LW 343	5.0×10^{-3}	50189	0	9.799×10^{-6}	9.68
B (current)	5.0×10^{-3}	48062	0	9.934×10^{-6}	9.65
B (current)	1.0×10^{-2}	48007	0	9.987×10^{-6}	9.67
B + LW 343	5.0×10^{-3}	47965	0	9.900×10^{-6}	9.68

Seed B reproduces the section 60 table exactly, hand-off step included (the re-convergence that produces it also reproduces the section 60 control: step 25839, info = 0, $\|R\| = 6.105 \times 10^{-6}$, $\log_{10} \dot{M} = 10.24$, front at $1.0097 R_p$). Seed A does not: the hand-off comes ~ 2600 steps later, from a different state, and the line search then collapses ($\lambda \sim 10^{-6}$) and aborts. The section 59 change moved the examples/15 solution by more than this marginal solve tolerates in its starting point, so a restart file written by an older binary is no longer the same seed. The Lyman–Werner member converges from both seeds — the LW coupling appears to widen the basin here rather than narrow it.

What the weak determination of the layer costs. The two converged solutions ($\varepsilon_4 = 5 \times 10^{-3}$ and 1×10^{-2} , both info = 0 at Resid tol: 1.0×10^{-5}) agree in the wind — $\Delta\rho/\rho$ +5.1% at $1.3 R_p$, +3.4% at 1.5, +2.6% at 2.0, $\Delta v/v \leq 1.6\%$, $\Delta T/T \leq 0.4\%$, $\log_{10} \dot{M}$ 9.65 against 9.67 — and disagree inside the layer, where the $10^{-2}/5 \times 10^{-3}$ density ratio runs 0.78 at $r = 1.005$, 1.31 at 1.013, 1.41 at 1.020 and 1.32 at 1.030, T differs by +20% at 1.005, +38% at 1.015 and -7.7% at 1.020, and the alternating component of v over $r = 1.015$ – 1.030 is 0.392 against 0.028 on a mean $|v|$ of 12–16 cm s $^{-1}$. That does not stay internal. Run through EXHALE_transit.py, the theoretical peak excess absorption is 40.19% against 40.26% in He I 10830 (0.2% relative), 2.098% against 2.199% in H α (+4.8%), 1.046% against 1.134% in H β (+8.4%), with Ly α saturated in both. He I 10830 forms in the wind and does not care which layer sits underneath it; part of the H($n = 2$) absorption is formed in the layer, so the weak determination appears to propagate to the Balmer depths at the 5–8% relative level. Quote those depths with the pair, not from a single coefficient.

Lyman–Werner with metals on. The Stellar LW flux [erg/cm 2 /s]: 343.0 A/B on seed B at $\varepsilon_4 = 5 \times 10^{-3}$, both members info = 0:

	LW off	LW on (343)
$\log_{10} \dot{M}$	9.65	9.68
$H_2 = \text{HI front}$	1.01108	1.01026
T at $r = 1.02$ [K]	1801	1470
$x(H_2)$ at $r = 1.02$	4.0×10^{-6}	2.1×10^{-5}
H_3^+ peak [cm^{-3}]	1.93×10^5 at 1.0006	1.99×10^5 at 1.0006
HeH^+ peak [cm^{-3}]	0.285	0.053
C II fraction at 1.02 / 1.20	0.197 / 0.449	0.180 / 0.430
$H\alpha$ peak	2.098%	2.297%
$H\beta$ peak	1.046%	1.221%
He I 10830 peak	40.19%	40.33%

The self-shielding is the section 56 physics unchanged: at the base $N(H_2) = 2.68 \times 10^{21} \text{ cm}^{-2}$, $f_{\text{shield}} = 2.11 \times 10^{-6}$, $k_{\text{LW}} = 1.27 \times 10^{-10} \text{ s}^{-1}$; above $r = 1.10 R_p$ the shielding is gone ($f_{\text{shield}} = 1.000$, $k_{\text{LW}} = 6.026 \times 10^{-5} \text{ s}^{-1}$). $x(H_2)$ at $r = 1.02$ is *higher* with the band on, which appears to be the layer temperature rather than the chemistry: that layer is 330 K colder, and the temperature dependence of the H_2 balance outweighs the added destruction. The front moves down by $0.0008 R_p$, consistent with more destruction where the shielding has lifted. Caveat: the band changes the $H\alpha$ peak by +9.5% relative and $H\beta$ by +17%, only about twice the numerical spread above, so those changes are not quotable as observational predictions without the coefficient pair run alongside.

Files: `src/modules/time_step/steady_newton.f90` (diagnostic only). Documentation: `docs/hd209_metal_stagnation.md` section 9, `examples/README.md`, `README_HOWTO.md`, `TO_BE_DONE.md` item (A).

62. Runtime grid size, CODATA k_B , the layer residual on a physical scale, and two diagnostic closures (2026-08-15)

One working series. The verification was staged so that every piece except the k_B value is proven a no-op on the regression matrix, and the golden refresh at the end carries the k_B change alone.

Grid cells: — the cell count is a runtime quantity. N was a compile-time parameter (500); it is now set by the optional key `Grid cells: <N>` with default 500, so an `input.inp` without the key is unchanged. The grid-sized module arrays of `global_parameters`, `lya_rt`, `excited_hydrogen` and `ionization_equilibrium` became allocatable (lower bound 1- N_g kept) and are allocated once parsing has resolved the key; the main program's state vectors likewise, zeroed on allocation to reproduce their former static-storage start. Procedure-local grid arrays were already automatic. `write_setup_report` echoes the resolved N and `load_IC` counts the IC records first, so a restart against an IC written at a different N reports the mismatch instead of an end-of-file error. This closes `TO_BE_DONE` item (E): the refinement pairs (Base grid [dr , cells], Grid cells) = (1.0e-4 100, 500), (5.0e-5 200, 658), (2.5e-5 400, 916), (1.25e-5 800, 1375) hold the upper stretch from `input.inp` alone. Gates: `make` and `make wind_ae_ic` build; `make check 5/5 byte-identical` against the pre-change goldens; a 750-cell run writes $750 + 2N_g$ rows; `-fcheck clean`

at 500 and 750 cells; an H-only A/B (the one path the matrix does not cover, where a parameter zero array had to become a local) is byte-identical.

k_B is the CODATA value; the lyman_werner copy is synced. $kb_erg = 1.38 \times 10^{-16}$ (the truncated ATES literal) is now the SI-exact 1.380649×10^{-16} . The judgment first: the old value was wrong at 4.7×10^{-4} relative and every thermal quantity passes through it; the smallness of the effect is not an argument for keeping it. The dependency-free copy kb_lw in `lyman_werner.f90` is synced. Measured against the old goldens: the WASP-121 b cases re-converge at count 14065/14040 (from 14060/14037), fields move by a few 10^{-4} – 3×10^{-3} relative, $\log_{10} \dot{M}$ by the last printed digit ($13.22 \rightarrow 13.23$). **The goldens and the parse corpus were re-snapshotted at the end of the series and carry exactly this change.** Deliberately not touched (recorded so the list is explicit): kb_eV , hp_erg , hp_eV , μ , G_c and the astronomical constants — a full constants pass would be its own measured rebaseline.

The below-escape momentum residual has a physical scale. Section 61's diagnostic divided every row by $|u|$; for the momentum row that is $|\rho v|$, which collapses in the quasi-hydrostatic layer and made converged states read 0.2–10. The momentum row is now divided by the gravitational force density $|\rho| |G\phi_i(j) - G\phi_i(j-1)| / \Delta r_j$ (the discrete potential difference of the source term), so the number is the fractional violation of hydrostatic balance; the three rows print separately. Print-only, make check 5/5 byte-identical. Re-measured on the section 61 states (`docs/hd209_metal_stagnation.md` section 10): the converged examples/15 (old scale 9.99) and examples/16 $\varepsilon_4 = 5 \times 10^{-3}$ (0.195) and 1×10^{-2} (2.352) solutions all hold the layer at 10^{-5} – 1.4×10^{-4} of ρg . The layer is not momentum-unbalanced; what stays uncontrolled between the pair members is *which* hydrostatic stratification it settles on.

Heating_breakdown.txt is consistent for a molecular run. The dump reconstructed atomic heating only (and said so in a NOTE); its electron density omitted the molecular ions. It now rebuilds the H_2 photoheating (the H_2 density enters `PH_heat_HHe`, so column and opacity are the solver's), the $He(2^3S)+H_2$ Penning and the Lyman–Werner heating exactly as `ionization_equilibrium` adds them, uses the same `calc_ne`, and appends `col15 heat_He23S_H2_Penning`, `col16 heat_H2_LW`. Measured on `mol_lyman_werner`: $\max |\sum \text{cols } 5\text{--}16 / \text{heat_total} - 1| = 2 \times 10^{-15}$, all three molecular channels nonzero. Atomic runs unchanged; the file is not part of the regression comparison.

Lya absorbing bottom: — a downward loss channel (default off). The section 30 audit follow-up: Huang et al. (2017) close their Monte Carlo with a pure-absorber bottom (H_2 accidental-resonance true absorption), while the local closure had no lower boundary. With the key on, a third channel joins the wing and Sobolev escapes: the same Neufeld form on the depth from the cell *down* to the bottom of the domain, $\beta_{bot} = \min(1, \pi^{-1/4} \sqrt{a/\tau_{bot}})$, combined in parallel and feeding the existing β_{tot} , so the loss acts on the internal field and the trapped stellar beam consistently. First look (HD 189733 b benchmark configuration, 300-step snapshot, qualitative): $\bar{J}$ and n_{2p} drop together — 0 at the innermost ghost, $\times 0.24$ at $r = 1.0002$, back to within 1.5% by $r = 1.023$, exactly 1 above 1.05 — confined to the bottom few scale heights, as the audit expected. Off path executes the original expression unchanged (byte-identical pass below).

windae du-trigger arming, measured. The 2026-08-11 descending-crossing guard had never been measured on IC mode: windae, whose ~ 3 -step false du stop originally forced the “always Newton-finish” rule. On examples/12_windae_ic_hd209 (2000-step cap): du at step 2 is 1.43×10^{-4} — below every threshold on the untouched generated IC — and nothing fires; the du stop arms at step 23 (1.027×10^{-3}) and the Newton hand-off at step 222 (1.002×10^{-2}), both on the first ascending crossing. A du-stop $\dot{M}$ still carries the usual path spread; the Newton finish stays the prescription.

Verification ledger.

stage	tree	make check
runtime N	that change only	5/5 byte-identical
+ layer scale + load_IC guard		5/5 byte-identical
+ breakdown fix + Lya bottom (off), old k_B	all but k_B	5/5 byte-identical
+ CODATA kb_erg	full series	FAIL vs old goldens as measured above; goldens re-snapshotted; final 5/5

63. The rest of the physical constants, and why iron cannot move to Badnell (2026-08-19)

Two items from the “recorded, blocks nothing” list, taken together because the first needs a golden refresh and the second turned out to need none.

The constants pass. Section 62.2 updated k_B and listed what was deliberately left alone; those are now done, against CODATA 2018 and IAU 2015:

constant	was	now	rel. change
kb_eV	8.6167×10^{-5}	$8.617333262 \times 10^{-5}$	7.4×10^{-5}
mu	1.673×10^{-24}	$1.67353284 \times 10^{-24}$	3.2×10^{-4}
Gc	6.67259×10^{-8}	6.67430×10^{-8}	2.6×10^{-4}
hp_erg	6.62620×10^{-27}	$6.62607015 \times 10^{-27}$	2.0×10^{-5}
hp_eV	4.1357×10^{-15}	$4.135667696 \times 10^{-15}$	7.8×10^{-6}

π , erg2eV, c , parsec, AU and R_J were already exact to the printed digits. mu is the mass of the hydrogen *atom* ($m_p + m_e - 13.6\text{eV}/c^2$), which is what the density normalization means; the comment now says so. The same values appear in four other places, all synced: m_h2 in lyman_werner.f90, Gcgs in lower_column.f90, and the Wind-AE constants in wae_params.f90, wae_soe.f90 and wae_exhale_input.f90. The Wind-AE files port an external code, so the original defs.h values are written into the comment there; against them the shifts are 2.6×10^{-4} , 6.5×10^{-6} and 1.4×10^{-4} , so the port now agrees with the C/Python reference at that level instead of bit-exactly, and restoring the three lines reproduces the gates verbatim (backup/regression/windae_oracle, working copy only). The helium atomic mass was left at Wind-AE’s $6.6464790722 \times 10^{-24}$, which is the CODATA He atom mass to 10^{-6} .

Measured effect: fields move 10^{-3} – 5×10^{-3} relative and $\log_{10} \dot{M}$ does not move in the printed digits (wasp_full 13.23; counts 14065 $\rightarrow$ 14059, 14040 $\rightarrow$ 14036). The one outlier is the photoheating of the last interior cell of wasp_he23off ($r = 1.6056$), which changes by 97% because the old value was an isolated dip below its neighbour; the rest of that file agrees to 0.2%. Goldens and parse corpus re-snapshotted, make check 5/5. The four paper models were re-converged (all info = 0, $\log_{10} \dot{M}$ unchanged at 9.47 / 9.14 / 11.70 / 13.20); the LaRT fields were not re-run, their inputs moving an order of magnitude below the Monte Carlo noise of a 2×10^6 -photon run.

Iron recombination: Phase 6 is closed, not deferred. The code carried a note that switching iron from the Huang et al. (2023) Eqs. (5)–(6) fits to Badnell RR+DR was “deferred to a later pass”. It cannot be done, for a reason outside this code: **the data do not exist**. Fe I is produced by recombining Fe II (Mn-like, 25 electrons) and Fe II by Fe III (Cr-like, 24), while the Badnell dielectronic-recombination project reaches the phosphorus sequence (15 electrons) in its latest instalment (paper XVI, 2022). Neither iron stage is covered — nor is the K-like sequence the Ca I comment in the same file already records as missing. The other nine metals are on Badnell, which is what the phase asked for. The Huang plan’s own Phase 1c prescribes exactly this fallback “where Badnell coverage is thin”, so the code was following the plan rather than lagging it. What remains unbuilt is only that gate’s bookkeeping: a runtime key selecting one rate source for all metals, and a Huang-vs-Badnell delta table; neither applies to iron.

Documentation pass on the md/tex corpus. The sweep that accompanied section 62 predates 63.1/63.2, so a second pass propagated them and re-checked the reference documents against the code. Findings beyond the expected propagation: the manual and the README presented the iron swap as pending (both now carry the isoelectronic reason, and the manual names the Ca I parallel); the README described line trapping as covering “the two fine-structure coolants” in a two-level scheme, when it has been the eight ground-term lines of C I/C II/N II/O I ($n_{\text{fsl}} = 8$) in exact statistical equilibrium (cooling_formulas.tex and the manual were already right); EXHALE_BC_and_IC.tex quoted $\log_{10} \dot{M} = 9.31$ for HD 209458 b as the value “under the present code”, which was the 2026-08-11 number against 9.47 today; four line-number citations in atomic_data_EXHALE_vs_MoCHII.tex pointed at the wrong lines and were replaced by symbol names; the transit resolving powers were undocumented, and Ca II/Na I default to *whatever* Instr_res_Ha is, so an H α override silently moves both doublets (README_HOWTO.md now tabulates the seven defaults and states the trap); and wind_ae_solver.tex claimed a bit-exact base-BC gate that 63.1 turned into a 10^{-6} – 3×10^{-4} agreement. Two mechanical cross-checks passed in both directions: all 56 keys parsed by input_read.f90 appear in the manual, and every key-like entry in the manual is still parsed. PDFs rebuilt and inspected: manual 54 pp., this log 118, BC/IC 23, Wind-AE 7, atomic data 7; no errors, no undefined references, and the surviving overfull boxes are pre-existing table rows away from the edited pages.

What the constants pass had still missed, and the golden refresh it needed. Section 63.1 swept the .f90 files for the constants *named* in parameters.f90. Rescanning for their *values* found four more sites. One is in the solution path: Cool_coeff.f90 computed He II recombination cooling as $1.38\text{e-}16 * T * \alpha_B$ in two places. That literal is the Boltzmann constant, not a fit coefficient — the $k_B T \alpha_B$ form *is* the Hui & Gnedin (1997) He II expression, while H II and He III use their fitted case-B rates — so it was the truncated ATES value section 62.2 had replaced everywhere else, leaving this coefficient 4.7×10^{-4} low. It is kb_erg now. The other three are report-only or write generated files:

`mh_g` in `lower_column.f90` and `ICW_mH` in `wae_ic_writer.f90` were the proton mass where the hydrogen atom is the code's mass unit, and `m_amu_g` in `species_diffusion.f90` was labeled “amu” while holding the proton mass — with masses on the $H = 1$ scale there, the right multiplier is the hydrogen atom, neither the proton nor u . Fifteen constants in ten Python tools were brought to the same values, including `src/utls/glob.py` (the GUI's mirror of the global block) and the three-digit G in `exhale_transit_lib.py`; the `cooling_data` scripts deliberately keep the rounded $k_B = 1.38 \times 10^{-16}$ the CHIANTI fits were made with.

The verification was staged so the two kinds of change stay separable. The report-only fixes alone give `make check 5/5` byte-identical, which is what proves `mh_g` never reaches a solution. Adding the `kb_erg` fix fails all five, as it must: against the old goldens the fields move by $3.7\text{--}3.8 \times 10^{-6}$ (max, in the `cool` column — the edited term) on the two atomic cases with median $\sim 4 \times 10^{-7}$, and by $\leq 2 \times 10^{-4}$ with median $\sim 10^{-8}$ on the molecular ones, which carry almost no He II. Step counts and $\log_{10} \dot{M}$ (13.23) do not move. Goldens re-snapshotted, `make check` back to 5/5; the parse corpus is unchanged (125 cases OK) because none of these constants enters a derived setup value. The four paper models were not re-converged: 3.7×10^{-6} on the hottest, most ionized case is three decades below the printed precision of every number in `paper/`.

64. Resolved-configuration output, and the transit tool reading it (2026-08-22)

`EXHALE_setup.out` records the resolved configuration in prose, but `EXHALE_transit.py` re-parsed `input.inp` only — so a `base.inp` handoff that overrides the planet radius, base temperature, or He/H reached the wind and not the transit geometry. Measured on benchmarks/wasp52: the wind used $1.40492 R_J / 1181.2 \text{ K} / \text{He/H } 0.0959$ (from `base.inp`) while the transit read $1.27 R_J / 1304 \text{ K} / 0.0204$ from `input.inp`. The radius enters chord lengths, absorbing areas, and the rotation velocity, so this was a real transit-side error for every handoff run.

Fix, per docs/lhs1140b_lower_atmosphere_plan_new.md Phase B:

- **EXHALE_resolved.out** (new, machine-readable): written next to `EXHALE_setup.out` at startup by `write_resolved_config` (`write_setup_report.f90`) after `input.inp` + `base.inp` resolution. Key-value lines: `planet_radius_RJ`, `planet_mass_MJ`, `equilibrium_temperature_K`, `HeH_number_ratio`, `orbital_distance_AU`, `star_mass_Msun`, `base_inp_present`.
- `exhale_transit_lib.read_input_params` prefers that file when it sits next to `input.inp` (returns `resolved=True`, plus the resolved HeH); absent file falls back to `input.inp` exactly as before. A present-but-unreadable file is an error, not a silent fallback. `EXHALE_transit.py` prints which source it used.
- **he_line_metrics.py** (new, repo root): the He 10830 line metrics of Cherubim et al. (2026) — three Gaussians with shared width and shift; blended-red depth, blue depth, red/blue amplitude ratio, FWHM of the blended feature — with a synthetic-profile selftest (`python3 he_line_metrics.py --selftest`). `EXHALE_transit.py` now fits the instrument-convolved He curve (air frame) and writes `tpm_He10830_metrics.txt`; a fit failure warns and does not stop the run.

Verified: a scratch run with a `base.inp` overriding radius/temperature/He shows the transit consuming $1.50 R_J / 900 \text{ K}$ (the override), and the fallback path still returns the `input.inp` values; `make check` byte-identity, since the new file is additive and the compared outputs are untouched.

65. The metastable helium triplet is on by default (2026-08-22)

`Include He23S` was a mandatory key initialized to `.false.`, so every run carried whatever the file it was copied from happened to say. Across the tree that produced 28 runs with the triplet off, almost none of them off on purpose: the setting was inherited when the case was created for something else entirely (a solver, an IC, a domain). A helium-bearing wind has the metastable level whether or not the input file remembers to ask for it, and the level is not a diagnostic add-on — it carries its own photoionization, Penning ionization with H and H₂, recombination, and collisional-excitation cooling, so leaving it out changes the electron budget, the energy budget, and the 10830 Å observable. Off is the special case, not on.

- `parameters.f90: thereis_HeITR = .true.`
- `input_read.f90`: the key became **optional** (it was read through `req`, so the initializer never actually applied to a valid run). A file that omits the line gets the triplet; `Include He23S? False` is the deliberate opt-out and is what the HeITR-off regression branch uses. `thereis_He false` still forces it off, as before.
- The Tk interface (`src/utis/EXHALE_interface_functions.py`) follows: the `checkboxbutton` variable starts at 1, the reset button restores it to 1 instead of clearing it, and `load_input` treats the key as optional with the same semantics as the Fortran reader — on unless the file says `False`. Since a file may now omit the line, the sequential reader checks the label before consuming the line, so an older file without it no longer shifts `Load IC`, `Do only PP`, and `Force start` by one.

Verified: a one-step run with the key absent prints “Including helium triplet chemistry” in the setup report, and `Include He23S? False` does not; the Tk reader parses the tutorial input with and without the line to the same `Load IC / Do only PP` values. `make check` is unaffected — every `backup/regression/*/input.inp` sets the key explicitly, including the two `wasp_he23off*` cases that exist to exercise the off branch.

The `examples/01–12` ladder keeps `Include He23S? False` on purpose: each folder is its base plus exactly one line, and atomic helium is the baseline that leaves `06_he23s` a one-line difference against `03_newton` instead of a folder that differs in nothing. Outside the ladder the triplet is on — including `examples/13_lower_atmosphere/`, which had been running two of its four planets with the triplet and two without. That split mattered: the $\text{He } 2^3\text{S} + \text{H}_2$ Penning channel is active only when a molecular run also tracks the triplet, so one example was doing different physics on different planets. All four now carry `True`; none of them has stored output.

The runs whose stored output predates the flip are listed in `docs/he23s_default_recompute_list.md`; nothing there has been recomputed.

66. Primitive conversion of unfilled ghost cells (2026-08-24)

A runtime-checked JFNK run of the LHS 1140b He/H = 1000 case (Phase C audit, LHS1140b/exhale/audit_summary.md) ended with gfortran reporting a signalling IEEE_INVALID_FLAG. Rebuilt with `-ffpe-trap=invalid,zero,overflow`, the solver died in its first residual evaluation at `UW_conversions.f90:58`, `W(2,:) = U(2,:)/U(1,:)`: `U_to_W` converts the whole array, ghost cells included, and `Apply_BC` called it on a state whose ghosts were still zero — the comment two lines above the call in `steady_newton.f90` had already noted that the ghosts are zero at that point. The `0/0` produced NaNs in the ghost primitives that `Apply_BC_W` overwrote on the next line, so no result was ever affected; but every steady solve tripped an invalid operation, which is what had kept the trapping variant of the `-fcheck` regression off the steady-solver path. With that site fixed the trap moved to the two line-search trial conversions in `steady_newton.f90` (call `U_to_W(utry, Wtry)` in the hybrid and JFNK searches), where `unpack_U` fills only the interior of `utry` and only `Wtry(1,1:N)` and `Wtry(3,1:N)` are then read.

- `UW_conversions.f90`: new `U_to_W_interior`, converting cells $1 \dots N$ through `U_to_W_comp` and zeroing the ghost columns — one definition for every site that converts a state before its ghosts are set.
- `Apply_BC.f90`: `U_to_W` $\rightarrow$ `U_to_W_interior`; the ghosts are outputs of `Apply_BC_W`, never inputs.
- `steady_newton.f90`: both line-search trial conversions $\rightarrow$ `U_to_W_interior`.

Verified: under `-O1 -g -fcheck=bounds,do,mem -ffpe-trap=invalid,zero,overflow` the He/H = 1000 JFNK now completes, `info=0` $\|R\| = 7.929 \times 10^{-4}$ (identical to the untrapped run) with no signalling exception, and so does the H-rich solar_gj699 case (`info=0` $\|R\| = 6.433 \times 10^{-4}$). `make check`: 5/5 byte-identical (all five cases), as the discarded ghost values require.

67. Element diffusion on the direct-steady route, the element drift metric, and the restart of a diffused state (2026-08-25)

Three defects on the `He_diffusion` path, all of them independent of the transport formulation and all of them blocking the Phase-D milestones of `docs/binary_diffusion_design.md` (sections 7.3, 7.4).

The steady solve and the diffusion relaxation are now one routine. The steady residual (`steady_newton.f90`) carries no diffusion: the He/H field is moved by the operator-split `he_diffusion_step`, so a single steady solve freezes the composition at the state it was handed. The marching hand-off answered that with an outer Picard loop (JFNK $\rightarrow$ 500 relaxation steps $\rightarrow$ `ioniz_eq` $\rightarrow$ drift test, at most five passes), but the direct-steady route — `EXHALE_PTC=1`, the route the LHS 1140b cases take — ran a bare `solve_steady_jfnk / solve_steady_ptc`, so a diffused wind could not be brought to a steady state there at all. That loop is now `steady_wind_with_element_diffusion` (contained in `EXHALE_main`), called from both routes; the solver, its iteration budget and its initial pseudo-time are its arguments, so the direct-steady route reaches it with JFNK or with pseudo-transient continuation. `dt_loc`, which the relaxation step needs, is built by the `eval_dt` that route already called. With `He_diffusion` off the loop runs exactly once and the body is the single solve plus the state refresh both sites performed before.

The drift metric counts elements. It read $(f_sp(:,3)+f_sp(:,4)+f_sp(:,5))/(f_sp(:,1)+f_sp(:,2))$, which is He I + He II + He III over H I + H II: it omitted the He 2³S triplet and every molecular species, and so did not measure the same quantity that `species_diffusion.f90` transports (its own `sumHe` includes the triplet). Replaced by `element_ratio_HeH` in `composition.f90`, which counts nuclei with the `bsp_nH` / `bsp_nHe` weights of `species_table.f90` over all ten base species — H₂ and H₂⁺ two H nuclei, H₃⁺ three, HeH⁺ one of each. It is the same element count `load_IC` applies to a restart file, and it is correct in the molecular region without a second definition.

A diffused state can be restarted. `load_IC.f90` rescaled every cell to the global HeH whenever the file's composition differed from it by more than 10⁻⁶, on the ground that the input file is the authority on composition. With `He_diffusion` on that erases the answer: the cell-by-cell split is what the diffusion operator produced. The rescale is now confined to the base cells (1-Ng to 1), where the operator itself holds a Dirichlet reservoir composition; every cell above keeps the helium fraction it was written with. With the flag off the whole column is rescaled as before, bit for bit. The HeH⁺ refusal is unchanged — molecular chemistry and `He_diffusion` remain an excluded pair until the molecular closure is validated (memo section 5, milestone M4).

No diffusion operator changed; `species_diffusion.f90` is untouched.

Verified. `make check==> REGRESSION PASS` (all cases byte-identical), 5/5, which is the evidence that the `He_diffusion-off` path is unmoved.

The restart was measured against the stored `examples/14_diffusion` output under `EXHALE_DUMP_IC=1`, which writes the state as loaded, before the first ionization-equilibrium solve. The loaded He/H profile reproduces the file's to 3.3×10^{-16} relative, worst cell, over all 504 cells; the same file through the pre-change binary comes back flattened to the input ratio ($r = 1.05 R_p$: file 1.616×10^{-2} , loaded 8.333×10^{-2} , a factor 5.2). The direct-steady route with `EXHALE_PTC=1` `EXHALE_PTC_JFNK=1` now runs the outer loop — five passes, `info=0` on every one, drift falling 9.40×10^{-3} , 7.74×10^{-3} , 5.79×10^{-3} , 4.64×10^{-3} , 4.03×10^{-3} . It reaches the pass limit rather than the 10⁻³ drift test, which is what a five-pass Picard budget gives from a state that was not converged to begin with (the stored output is the unconverged JFNK `info=2` one); the same test from a converged state is part of the baseline capture below.

The pre-change binary is kept as `backup/phase_d_baseline/EXHALE_prephaseD.x`, with the reference runs it produces, as the memo's T2 comparison.

68. Binary H/He element diffusion replaces the trace-helium kernel (2026-08-25)

The composition operator of `He_diffusion` is rewritten from the trace-helium-in-hydrogen kernel to the two-component (binary) formulation of `docs/binary_diffusion_design.md`, sections 2–4. New module `src/modules/functions/binary_element_diffusion.f90`, routine `element_diffusion_step`; the old module is kept for the duration of milestone M2 with its routine renamed `he_diffusion_step_trace` and its body untouched, reachable only through `EXHALE_DIFFUSION_KERNEL=trace`, so the two can be compared (test T2a). Both call sites in `EXHALE_main.f90` go through one dispatcher, `element_composition_step`.

What changed in the physics. *The transported variable.* The trace kernel evolved the ratio $f = n_{\text{He}}/n_{\text{H}}$, which diverges as hydrogen becomes the minor element — exactly the limit a helium-rich atmosphere needs — and it wrote the diffusive flux as $n_{\text{H}}Dd(n_{\text{He}}/n_{\text{H}})/dr$, which divides by a vanishing background. The new operator transports the helium **mass fraction** $X = \rho_{\text{He}}/\rho$ of two components, hydrogen (with the trace metals slaved to it) and helium, and is bounded at both ends of the composition axis. The gradient coefficient of the memo, $A = (m_1 m_{\text{He}} n / \bar{m}) D_{12} (dx/dX)$, collapses to $A = \rho D_{12}$, so nothing in the operator divides by an element density. The settling flux carries the prefactor $\rho X(1-X)$, which switches settling off as either element is exhausted, while the gradient flux does not vanish there — a helium-free cell next to a helium-bearing one still receives helium.

The cap is gone. The trace kernel clamped $f \leq \text{HeH}$ over the whole domain, and in the stored examples/14_diffusion output the profile sat exactly on that clamp between 1.2 and $2R_p$. The clamp existed because the trace flux has no reason to vanish when hydrogen runs out; the binary flux does. It is removed, together with the base pile-up limiter. A pile-up, if one appears, is physics for the base boundary condition to answer, not something to clamp. What remains is a round-off assertion ($0 \leq X \leq 1$), and the bounds it asserts are a property of the discretization: the backward-Euler matrix is a strictly diagonally dominant M-matrix, with the Péclet hybrid on the settling term chosen so that it never flips the sign of an off-diagonal.

The ambipolar field is computed, not assumed. The previous effective settling mass $3 - 0.5(\bar{Z}_{\text{He}} - \bar{Z}_{\text{H}})$ is the hydrogen-plasma value $eE = m_{\text{H}}g/2$ written as a constant; in a fully ionized helium plasma the massless-electron balance gives $eE = \frac{4}{3}m_{\text{H}}g$ instead, and the relative settling mass $5/3$, not 2.5 . The field is now evaluated from the solved electron pressure of the cell, $eE = -(1/n_e)d(n_e kT)/dr$, in its logarithmic form $-kT d \ln(n_e T)/dr$ (exact for a barometric stratification, the same second-order central difference elsewhere). The three analytic limits are reproduced to round-off (test T9): neutral 3, H^+ plasma 2.5 — so the validated helium-trace behaviour is unchanged — and He^{++} plasma $5/3$.

Advection. The operator receives only the new ρ , v , T and dt , not the hydro's Runge–Kutta face mass fluxes, so a second conservative advection of ρX could not be made consistent with the hydro's own continuity step. X is therefore transported in the advective form, $\rho \partial X / \partial t + \rho v \partial X / \partial r = -\nabla \cdot (r^2 J) / r^2$, with the velocity divergence cancelled analytically in the row assembly rather than formed and subtracted (that cancellation is catastrophic whenever the step is long compared with the advective crossing time). A uniform X is then preserved exactly whatever the hydro does, and at a steady state, where $\rho v r^2$ is constant, the advective and conservative forms are the same equation. In a transient the helium mass is conserved only to the order of the hydro's own truncation; this is the stated limitation of the phase.

Write-back. The species vector is projected onto the new element totals by a nonnegative projection: helium-only species scaled by r_{He} , hydrogen-only species (H_2 , H_2^+ , H_3^+ included) and the slaved metals by r_{H} , HeH^+ by $\min(r_{\text{H}}, r_{\text{He}})$, and the shortfall of the element with the larger factor deposited into its neutral ground species. Every operation is a multiplication by a nonnegative factor or an addition, so no negative intermediate can arise, and both element totals are met exactly. The old write-back used one factor per element and inferred the metal mass as $1 - m_{\text{H}}f_{\text{H}} - m_{\text{He}}f_{\text{He}}$, which counts molecular mass as metal.

Trace metals keep the previous element-by-element treatment behind `He_metal_diffusion` (a trace element cannot feed back on the background it moves through); only the background definition is now the nucleus count over all species rather than $\text{H I} + \text{H II}$.

User-visible change. He_Kzz now defaults to 0.0 instead of $10^9 \text{ cm}^2 \text{ s}^{-1}$. With the old default every diffusion run silently carried a constant eddy term; an eddy coefficient is a property of the atmosphere being modelled, so it is stated rather than inherited. An input that relied on the default must now say He_Kzz: 1.0e9; examples/14_diffusion already does. The four stored run directories that carried He_diffusion: True without the key — benchmarks/hd209, docs/version_compare/v2_diff, docs/version_compare/wasp_v2_diff, scratch_fs/hd189 — now state He_Kzz: 1.0e9 explicitly, so each input records the eddy coefficient its stored output was produced with; no output was regenerated. docs/input_schema.md K19 and README_HOWTO.md are updated. examples/exhale_io.py gains a derived Run.heh_profile, the He/H element ratio with the same nucleus weights the operator uses, so the analysis reads one definition.

Verification. A standalone driver, src/tests/diffusion_tests.f90, builds with make diffusion_tests and runs the acceptance tests of the memo's section 6 on synthetic columns and on a stored wind; all 17 checks pass.

id	measured
T1a	closed isothermal column relaxes to the barometric separation $dx/dr = -x(1-x)(m_{\text{He}} - m_1)g/kT$: max error of logit x 4.56×10^{-6} at $N = 400$; helium mass drift 2.0×10^{-16} at a step of about the cell diffusion time
T1b	Dirichlet base, integrated boundary flux accounts for the mass change to 2.2×10^{-11} relative
T3	He/H = 1000 with an outflow: finite, $0 \leq X \leq 1$, hydrogen rises (X falls below the reservoir value), no cap
T4	closed column, arbitrary initial $X(r)$, 10^4 steps: helium mass drift 2.4×10^{-14} (eos_metals on) and 3.7×10^{-14} (off), metals present in both
T5	two-component mass closure $m_1 n_{\text{H}} + m_{\text{He}} n_{\text{He}} = \rho$ after the write-back: 4.0×10^{-15} relative, worst cell over 200 steps
T6	uniform X , $g = 0$, compressing and expanding $v(r)$: 4.4×10^{-14} relative departure after 500 steps
T9	relative settling mass 3, 2.5, 5/3 to 8.3×10^{-13} in the three analytic limits; finite and continuous across a resolved ionization front
T10	second order in dr (observed 1.998); the steady state is step-size independent to 1.2×10^{-7} over two decades of dt
T2a	the operator increment against the trace kernel on a stored LHS 1140 b wind falls off first order in the helium fraction: 2.28×10^{-3} , 3.00×10^{-4} , 3.04×10^{-5} , 3.05×10^{-6} at He/H = 0.0833, 10^{-2} , 10^{-3} , 10^{-4} — a factor 10 per decade, and 3.0×10^{-6} at 10^{-4} against the memo's 10^{-3} gate. Over 500 relaxation steps the two profiles agree to 4.2×10^{-5} , 7.4×10^{-6} , 7.8×10^{-7} , 7.8×10^{-8} on the same sequence

T2a is measured with the wind's ρ and T and $v = 0$. The stored wind is not steady below $r \sim 1.2$ — its mass flux $\rho v r^2$ varies by orders of magnitude there and the base cell has $v < 0$ — so with that velocity the conservative advection of the trace kernel and the advective form differ by $X \nabla \cdot (r^2 \rho v)/r^2$, a difference of the **state** rather than of the trace limit, which does not fall off with He/H (the driver reports that column too). The comparison of the transport closure, which is what changed, is therefore made at $v = 0$.

The measured helium-mass drift of the closed column grows with dt (2.0×10^{-16} , 6.2×10^{-12} , 3.1×10^{-10} , 1.1×10^{-7} at $dt_{\text{code}} = 10^1$, 10^3 , 10^5 , 10^7 over 1000 steps). This is the round-off of the tridiagonal solve amplified by the conditioning of the implicit step, not a leak: the discretization telescopes exactly, and the drift falls to round-off as soon as the step is comparable with the cell diffusion time. The driver prints the table.

T0: make check is 5/5 byte-identical, which is the evidence that the He_diffusion-off path is unmoved by any of this. T7 (the molecular closure) is milestone M4 and T8 (the moving-wind face flux on a converged LHS 1140 b run) waits on the baseline capture; neither is claimed here.

69. Milestone M3: the binary operator on a wind, the LHS 1140 b helium-rich runs, and the removal of the trace kernel (2026-08-25)

Milestone M3 of docs/binary_diffusion_design.md: the binary element diffusion operator of section 68 is run on converged winds — the HD 209458 b examples/14_diffusion configuration against the captured baseline of the kernel it replaces (T2b) and the three LHS 1140 b helium-rich cases — the elemental face flux is measured on a moving wind (T8), and the trace kernel is then deleted.

69.1 T2b: the binary operator against the trace kernel on HD 209458 b

Both sides are the same configuration (HD 209458 b, He/H = 0.0833, Include He23S? True, He_metal_diffusion: True, PLM+WENO3), Newton finished, with the same five-pass diffusion outer loop. The baseline (backup/phase_d_baseline/kzz0, kzz1e9) was marched from a cold start with the pre-Phase-D binary; the binary-operator runs (new_kzz0, new_kzz1e9) restart from those converged states on the direct-steady route (EXHALE_PTC=1 EXHALE_PTC_JFNK=1), which is the route section 67 gave a diffusion loop. This is a report, not a gate (the memo's section 6): the baseline profile sits on the $f_{\text{He}} \leq \text{HeH}$ cap over part of the domain, so a difference there is the cap, not a defect of either side.

quantity	trace, $K_{zz} = 0$	binary, $K_{zz} = 0$	trace, $K_{zz} = 10^9$	binary, $K_{zz} = 10^9$
JFNK info / final $\ R\ $	0 / 7.737×10^{-4}	0 / 8.043×10^{-4}	0 / 8.017×10^{-4}	0 / 8.308×10^{-4}
outer passes / final drift	5 / 1.22×10^{-2}	5 / 1.29×10^{-1}	5 / 3.88×10^{-3}	5 / 3.19×10^{-2}
$\log_{10} \dot{M}$ [g s^{-1}]	9.7882	9.7883	9.7909	9.7910
He 10830 red depth [%]	9.7970	9.8134	9.7958	9.8118
(He/H)/HeH at $1.05 R_p$	0.62354	0.62354	0.62806	0.62808
at $1.10 R_p$	0.69771	0.69744	0.70193	0.70170
at $1.20 R_p$	0.95514	0.95515	0.95536	0.95536
at $1.50 R_p$	0.97542	0.97437	0.97555	0.97447
at $2.00 R_p$	1.00000	1.00111	1.00000	1.00112
at $3.00 R_p$	0.98628	0.98712	0.98643	0.98725
at $4.00 R_p$	0.85734	0.85834	0.85806	0.85922
cells within 10^{-6} of HeH	45 ($r = 1.00 \dots 2.44$)	3 (the base)	45 ($r = 1.00 \dots 2.44$)	3 (the base)

Three things are worth reading off it.

The flat stretch was the cap. Forty-five cells of the trace profile sat exactly on HeH between 1.0 and $2.44 R_p$. With the binary operator only the three Dirichlet base cells are on that value, and the profile crosses *above* it — 1.0011 at $2 R_p$ — which the clamp forbade. Helium is mildly enriched there relative to the reservoir, and the enrichment is small.

The base depletion is not a cap artifact. At $1.05 R_p$ both kernels give 0.6235; the two profiles differ by at most 4×10^{-3} anywhere above $1.05 R_p$ after 2500 composition steps of the binary operator. The depletion of the first cells (down to 10^{-3} of the reservoir ratio in the cell at $1.0004 R_p$, in *both* kernels) is not the limiter and not the lag of the trace form.

Corrected in section 70 (2026-08-25). That last depletion was read here as settling. It is not: the barometric He–H separation scale of that base is 24 cell widths, so a 550-fold drop across one cell cannot be barometric, and the settling flux measured at those faces is 2–20% of the gradient flux. It was the face-averaged upwinding of the advective term, which decouples a cell whose two face velocities straddle zero. The two kernels agreed because they shared it. Section 70 also supersedes the drift and profile numbers of the tables in this section.

Neither side is composition-converged. Both leave the outer loop at its five-pass ceiling, and the binary operator’s drift is an order of magnitude larger than the trace kernel’s. That drift is the maximum *relative* change of He/H over a pass and it is dominated by those near-base cells, where He/H is $\sim 10^{-3}$ of the reservoir value: over the same five passes the profile above $1.05 R_p$ moved by less than 0.4%. The observable moved as little: $\dot{M}$ by 10^{-4} dex and the He 10830 red depth by 0.017 percentage points ($9.797\% \rightarrow 9.813\%$). The five-pass ceiling is unchanged by this milestone; making the outer loop reach drift $< 10^{-3}$ needs either many more passes or a drift metric that is not dominated by a helium-free base cell, and is left open. (Answered in section 70: neither — the relaxation was being stepped at the hydro CFL step, four to six orders below the composition time scale, so it moved the base by $\sim 10^{-6}$ of the way per step. With the step taken from the composition time scale the $K_{zz} = 10^9$ case converges to 1.2×10^{-4} absolute in ten passes of ~ 30 steps.) load_IC kept the diffused profile across the restart, which is what section 67 added: the run reports top cell He/H = $7.0843\text{E-}02$ against the input $8.3333\text{E-}02$ and pins only the base cells to the reservoir value.

69.2 T8: the elemental face flux on a moving wind

With EXHALE_DIFFUSION_CHECK=1 the operator writes diffusion_faceflux.txt in the run directory at every step (replaced each time, so the file holds the state the run ended on): the diffusive face flux J , the advective face flux, their sum as $F_{\text{He}} = 4\pi r^2(\rho X v + J)$, the total mass flux $4\pi r^2 \rho v$ and the relative settling mass, all taken from the same face coefficients the implicit solve used — the file reports the discrete fluxes, not a re-derivation of them.

On new_kzz0 (HD 209458 b, $K_{zz} = 0$), with the spread measured exactly as the code’s own du is, $(\max - \min)/|\text{mean}|$ over a radial window:

window	spread of F_{He}	spread of $\rho v r^2$
$r \geq 2.0 R_p$ (the code’s escape window [j_min:N])	5.13×10^{-3}	7.19×10^{-3}
$r \geq 1.5 R_p$	1.68×10^{-2}	7.71×10^{-3}
$r \geq 1.2 R_p$	5.14×10^{-2}	1.25×10^{-2}
$r \geq 1.1 R_p$	3.84×10^{-1}	7.77×10^{-2}

T8 passes: in the window the code itself uses to declare the wind steady, the elemental helium flux is as flat as the mass flux (flatter, in fact); over the wider windows the two spreads stay within a factor 5, i.e. the same order of magnitude, and both are dominated by the same unsteady base region below $\sim 1.1 R_p$. The diffusive part is 12% of the advective part at most above $1.2 R_p$, so this is a real test of the sum and not of advection alone.

69.3 M3-2: LHS 1140 b with diffusion

Three cases in LHS1140b/exhale/, each the seed case's `input.inp` plus `He_diffusion: True` (`He_Kzz` at its new default 0, `He_ambipolar` at its default on), restarted from the seed's converged output and finished on the direct-steady route by `finish_case.sh` (JFNK $\rightarrow$ post-processing pass $\rightarrow$ EXHALE_transit.py). All three converged, `info = 0` on every pass.

	heh0p55_diff	heh1000_diff	heh0p06_gj699_diff
seed	heh0p55 (GJ 1132 SED)	heh1000	heh0p06_gj699
He/H reservoir	0.55	1000	0.06
JFNK info / $ R $	$0 / 4.439 \times 10^{-4}$	$0 / 7.929 \times 10^{-4}$	$0 / 8.482 \times 10^{-4}$
outer passes / final drift	$5 / 4.73 \times 10^{-2}$	$5 / 5.17 \times 10^{-2}$	$5 / 2.20 \times 10^{-2}$
$\log_{10} \dot{M}$, seed $\rightarrow$ diffused	7.7642 $\rightarrow$ 7.7645	7.7442 $\rightarrow$ 7.7437	8.5885 $\rightarrow$ 8.5883
(He/H)/HeH at $1.05 R_p$	0.9997	0.9973	0.9996
at $1.2 R_p$	0.9996	1.0020	0.9996
at $2 R_p$	1.0008	1.0030	1.0010
at $5 R_p$	1.0048	1.0074	1.0041
at $10 R_p$	1.0085	1.0101	1.0050
at $20 R_p$	1.0084	1.0102	1.0034
profile extremes	0.683 at $29.0 R_p$, 1.009 at $14.5 R_p$	0.626 at $29.0 R_p$, 1.195 at $1.03 R_p$	0.827 at $29.0 R_p$, 1.004 at $6.7 R_p$
H ionization front ($x_{\text{HII}} = 0.5$)	$9.43 R_p$	$1.03 R_p$	$5.66 R_p$
He 10830 red depth [%], seed $\rightarrow$ diffused	4.2847 $\rightarrow$ 4.2875	73.389 $\rightarrow$ 73.321	3.8335 $\rightarrow$ 3.8453

The separation is small on these winds. Between the base and $20 R_p$ the element ratio departs from the reservoir value by $\sim 1\%$ at most; only in the last few cells below the $30 R_p$ outer boundary does it fall to 0.63–0.83. The mass-loss rate and the He 10830 depth move by less than 0.1%. This is the expected ordering rather than a null result: these are fast, low-gravity winds, and the advection time across the domain is short against the diffusion time, so the wind carries the mixture out before it can separate. Whether the residual 5×10^{-2} drift would grow the separation if the outer loop were allowed to run to drift $< 10^{-3}$ is not answered here, and the five-pass ceiling was not changed for this milestone.

The ambipolar settling mass on the solved states. From the `dmeff` column of `diffusion_faceflux.txt`, above $1.05 R_p$: 2.56–3.00 (heh0p55_diff), 2.76–3.05 (heh0p06_gj699_diff) and 1.65–5.47 (heh1000_diff). The two hydrogen-rich cases stay near the neutral value 3 because helium is largely neutral there and the H^+ term enters through $\bar{Z}_{\text{He}} - \bar{Z}_1$. The helium-rich case spans the widest range and touches 1.67 at $1.11 R_p$ — the He^{++} value 5/3 to three digits — but **it does not get there as the He^{++} limit**: $x(\text{He III})$ reaches only 0.28 at the top of the domain in that run, so the 5/3 limit of the memo's (3a) is not sampled by these winds and remains guarded by the T9 unit test alone. Below $1.05 R_p$ `dmeff` is large and of either sign (down to -43 in the base cell at $1.03 R_p$), where the electron pressure gradient is steep and the diffusive flux itself is negligible.

LHS1140b/exhale/finish_case.sh gains one optional variable, EXHALE_BIN, which pins the executable a case is run with (these three were run with a snapshot binary while the tree was being rebuilt). With it unset the script behaves

exactly as before.

69.4 The trace kernel is removed

With T2a (section 68) and T2b recorded, `species_diffusion.f90` is deleted, with it the `EXHALE_DIFFUSION_KERNEL=trace` selector and the `element_composition_step` dispatcher in `EXHALE_main.f90`: both call sites now call `element_diffusion_step` directly. In the test driver T2a can no longer be measured — its comparison partner is gone — so it prints the M2 measurement as a recorded result and counts no verdict; the driver is 15 checks, all passing. `make check` stays 5/5 byte-identical (no regression case turns diffusion on, so this is the statement that the `He_diffusion-off` path is untouched).

Not done in this milestone: T7 and the molecular closure (milestone M4, the `error stop` on molecular chemistry + `He_diffusion` stands), and the regression case with `He_diffusion: True` that section 7.5 of the memo asks for at the end of the series.

70. The base helium hole, the metal drain beside it, and a composition relaxation that converges (2026-08-25)

Three defects on the element-diffusion path, all found on the milestone-M3 runs of section 69, all in the same place: the first cells above the base.

70.1 What was measured

On backup/phase_d_baseline/new_kzz0 (HD 209458 b, $K_{zz} = 0$, JFNK info = 0) the helium ratio of the first free cell above the base sat at $(\text{He}/\text{H})/\text{HeH} = 1.8 \times 10^{-3}$ between two cells at 1.0 — a 550-fold drop across one cell width, where the barometric He–H separation scale of that base is $H = kT/(\Delta mg) = 0.0045 R_p$, **24 cell widths**. A 550-fold drop across one cell is not barometric physics. Every metal was emptied from the same cell by $\sim 10^3$ (C I $2.3 \times 10^7 \text{ cm}^{-3}$ against 2.6×10^{10} in its neighbours), and with them the metal-line cooling of that cell. The dip sat exactly where the breathing base alternates the sign of the velocity: $v = (-796, +63, +28, -8, -9, +5) \text{ cm s}^{-1}$ over cells 1 to 6.

Rebuilding the operator's face coefficients from the converged state gives the row terms of the first free cell (cell 2), per unit X , in $\text{g cm}^{-3} \text{ s}^{-1}$:

cell	(He/H)/HeH	diffusive aa	diffusive cc	advective aa	advective cc
2	2.4×10^{-3}	-2.97×10^{-17}	-2.92×10^{-17}	0	0
3	4.5×10^{-1}	-2.78×10^{-17}	-3.22×10^{-17}	-4.53×10^{-15}	0
4	7.2×10^{-1}	-3.09×10^{-17}	-3.36×10^{-17}	-9.94×10^{-16}	-8.46×10^{-16}

and the face fluxes it balances, in $\text{g cm}^{-2} \text{ s}^{-1}$:

face	gradient J	settling J	total J	$\rho_f v_f$
1/2	1.451×10^{-11}	-3.193×10^{-13}	1.420×10^{-11}	-7.65×10^{-8}
3/2	-6.098×10^{-12}	-1.459×10^{-13}	-6.244×10^{-12}	8.78×10^{-9}
5/2	-4.114×10^{-12}	-3.575×10^{-13}	-4.472×10^{-12}	1.93×10^{-9}
7/2	-2.011×10^{-12}	-4.723×10^{-13}	-2.484×10^{-12}	-1.64×10^{-9}

The settling flux is 2–20% of the gradient flux and the Péclet number of the settling hybrid is 0.02 everywhere at the base, so the hybrid is on its central branch and the settling is fully resolved: **the depletion is not settling**. Both advective off-diagonals of that row are *identically zero*, so the cell was held by its diffusive coupling alone — and that coupling gives a relaxation time $\rho/(|aa|+|cc|) = 3.3 \times 10^6$ s against the ~ 2 s hydro step the relaxation was being driven at. Whatever composition the cell happened to hold was frozen there.

70.2 (a) The advection was upwinded on face-averaged velocities

face_coefficients built the advective coefficients from $v_f = 0.5(v_j + v_{j+1})$ and chose the donor per face. A cell whose two face velocities straddle zero ($v_f(j-1) < 0 < v_f(j)$), which is what the numbers above give: -366 and $+45$ cm s $^{-1}$) then donates at both faces and receives at neither, and its advective term vanishes identically — even though its own $v_j = +63$ cm s $^{-1}$ is large and its donor, the Dirichlet base itself, sits one cell away.

The advective term is now the one-sided upwind difference on the **cell** velocity,

$$\rho_j v_j \frac{X_j - X_{j-1}}{r_j - r_{j-1}} \quad \text{for } v_j \geq 0, \quad \rho_j v_j \frac{X_{j+1} - X_j}{r_{j+1} - r_j} \quad \text{for } v_j < 0, \quad (20)$$

which contributes +cadv to the diagonal and –cadv to the donor and to nothing else: a uniform X stays exact (T6), the row sum of the advective part is still zero, the M-matrix argument is unchanged, and the flux-minus-divergence construction the old form needed is gone. On the cell above, the coupling it restores is 6.43×10^{-15} , **216 times** the diffusive coupling that was holding the cell.

New unit test **T11**: a column seeded with the measured 500-fold hole in the first free cell and a velocity that alternates in sign cell by cell — so every face average vanishes while the cell velocities do not — must refill the hole. It gives 6.7×10^{-5} relative departure from the reservoir with the new form and 2.8×10^{-2} (28 times the 10^{-3} tolerance) with the old one. The driver is now 16 checks.

70.3 (b) The trace-metal solver drained the same cell

solve_trace_element_in_hydrogen transported the metal **density** conservatively with the same face-averaged velocities. In a conservative form the configuration above is not a decoupling but a drain: the cell has an outflow term at both faces and an inflow term at neither. That divergence has no counterpart in the hydro's own ρ , whose face fluxes come from the Riemann solver.

The solver now transports the **mixing ratio** $f_X = n_X/n_H$ in the same advective form as helium, so a uniform f_X is preserved under any velocity field and no spurious divergence can create or destroy the element. The $f_X \leq f_X(\text{base})$ cap in the write-back went with it: settling piles an element up as readily as it depletes one, and the cap is the same

limiter the memo (section 1) rejects for helium. In the runs below no metal exceeds its reservoir ratio anywhere, so the cap removal changes nothing in them by itself; what changes the metal profile is the transport.

70.4 (c) The relaxation ran on the CFL step, and on a flow with no mass flux

steady_wind_with_element_diffusion called the operator 500 times per pass with the hydro's local CFL step dt_{loc} . The implicit operator is unconditionally stable, so nothing tied its step to a sound-crossing time, while the base composition relaxes on $\Delta r^2/D_{12} \sim 10^6$ s against $dt_{loc} \sim 2$ s. That is why every case in section 69 left the outer loop at its five-pass ceiling with the drift falling by only ~ 0.75 per pass. The drift metric compounded it: $\max |\Delta \text{He}/\text{H}|/(\text{He}/\text{H})$ is a *relative* measure, and it was dominated by exactly the emptied cells where it means nothing. The relaxation is now `relax_element_composition` in the diffusion module:

- the step is the composition time scale, $\min(\Delta r^2/(D_{12} + K_{zz}), \Delta r/\max(|v|, w_s))$ per cell with $w_s = D_{12}|G|$, grown geometrically ($\times 1.5$, capped at 10^{12} times the start) so the late steps are direct steady solves, with the coefficients Picard-updated at every step;
- the measure is absolute, $\max_j |X_j^{\text{new}} - X_j^{\text{old}}|/X_{\text{base}}$, for the inner stop (10^{-12} per step) and for the outer pass stop (10^{-3} , at most 20 passes);
- the advecting flow is the wind's **steady mass flux**, $\rho v = \dot{m}_{\text{steady}}/r^2$ with $\dot{m}_{\text{steady}}$ the median of $r^2 \rho v$ over the escape window $[j_{\text{min}}:N]$. The advective form is the conservative equation only where ρ and v satisfy continuity, and below $1.02 R_p$ the converged states do not: the spread of $r^2 \rho v$ there is 4.8×10^2 (new_kzz0), 7.4×10^3 (new_kzz1e9) and 7.3×10^3 (heh0p55_diff) times its own median, against 1.7×10^{-2} to 4.6×10^{-1} over 1.1 – $2 R_p$. Relaxed on that field the operator converges to the composition of a flow that neither conserves mass nor exists. Solving the operator's steady state offline on the M3 new_kzz0 wind: with the cell velocities the base takes a five-fold cliff at the fourth cell and a plateau at 0.10 of the reservoir ratio; with the steady flux the first free cells are flat to 1% and the profile leaves the base on the barometric scale.

The marching call keeps dt_{loc} and the cell velocities: there the wind is genuinely transient and both are the consistent choice.

`element_ratio_HeH` is no longer used by the outer loop itself; it is kept as the element count of the diagnostic below, which writes one profile per pass.

70.5 The runs, re-run from the same restart files

backup/phase_d_baseline/new_kzz0_fix, new_kzz1e9_fix and LHS1140b/exhale/heh0p55_diff_fix: the section-69 configurations, the same output/*_IC.txt restart files, the same direct-steady route (EXHALE_PTC=1 EXHALE_PTC_JFNK=1 EXHALE_PTC_DTAU0=1.0). The M3 directories are untouched.

	$K_{zz} = 0$ M3	$K_{zz} = 0$ now	$K_{zz} = 10^0$ M3	$K_{zz} = 10^9$ now	LHS heh0p55 M3	LHS now
outer passes	5 (ceiling)	20 (ceiling)	5 (ceiling)	10	5 (ceiling)	6
final drift	1.29×10^{-1} rel	1.71×10^{-1} abs	3.19×10^{-2} rel	1.18×10^{-4} abs	4.73×10^{-2} rel	1.67×10^{-4} abs
relaxation steps per pass	500	24–33	500	28–31	500	28–400
(He/H)/HeH, first free cell	1.51×10^{-4}	0.9971	0.8293	1.0000	1.0001	0.9400
next four cells	0.381, 0.662, 0.810, 0.871	0.9939, 0.9960, 0.9994, 1.0025	0.7146, 0.6549, 0.6371, 0.6385	1.0000×4	0.9990, 0.9998, 1.0001, 1.0000	0.8818, 0.8510, 0.8143, 0.7695
at $1.05 R_p$	0.6235	0.9085	0.6282	0.9823	0.9997	4.9×10^{-4}
at $1.10 / 1.20 R_p$	0.6977 / 0.9551	0.8342 / 0.7986	0.6967 / 0.9553	0.9268 / 0.8962	0.9995 / 0.9996	$9 \times 10^{-5} / 2 \times 10^{-5}$
at $1.50 / 2.00 R_p$	0.9744 / 1.0011	0.7912 / 0.7974	0.9745 / 1.0011	0.8903 / 0.8959	1.0001 / 1.0008	~ 0
at $3.00 / 4.00 R_p$	0.9871 / 0.8583	0.7935 / 0.7145	0.9871 / 0.8603	0.8923 / 0.8279	1.0021 / 1.0035	~ 0
$\log_{10} \dot{M}$ [g s^{-1}]	9.7883	10.396	9.7910	9.998	7.7645	7.6065
He 10830 red depth [%]	9.8134	6.4032	9.8118	8.2115	4.2875	1.0×10^{-4}

The base is what the fix was for and the base is fixed. On both HD 209458 b cases the first free cells now sit within 0.6% of the reservoir ratio and the profile leaves them monotonically, instead of the isolated 550-fold hole followed by a five-cell recovery. On LHS 1140 b the first free cells fall 0.940, 0.882, 0.851, 0.814, 0.770 — ratios of 0.94–0.96 per cell against the $\exp(-\Delta r/H) = 0.93$ of that column’s barometric scale, i.e. the base now leaves the reservoir on the barometric scale, which is the criterion this change was measured against. The metals follow helium: C/H in the first free cell was 0.57 of the reservoir with $K_{zz} = 10^9$ and is 1.0000 now, and above $1.2 R_p$ it settles to ~ 0.56 where the removed cap pinned it at 1.0000.

Three things this exposes rather than fixes, and none is claimed here:

- **The separation is much larger than M3 reported** — 0.80 instead of 0.98 of the reservoir ratio at $1.2\text{--}3 R_p$ for $K_{zz} = 0$, and on LHS 1140 b helium is gone above $1.05 R_p$ instead of flat at the reservoir. That is what a relaxation that reaches its steady state gives; the M3 profiles were the restart files barely moved. For the LHS 1140 b column it is also the expected physics of $K_{zz} = 0$: at 226 K and $g = 1.8 \times 10^3 \text{ cm s}^{-2}$ the homopause of a hydrogen-helium mixture with no eddy mixing sits *at* the base, and the He 10830 line goes with it ($4.29\% \rightarrow 10^{-4}\%$). Those cases need a He_Kzz argued from the planet, not the default 0.
- **The $K_{zz} = 0$ HD 209458 b case does not co-converge.** JFNK returns info = 0 on every one of the twenty passes, but the pass drift settles into a limit cycle instead of falling, so its $\log_{10} \dot{M}$ of 10.396 is not a converged number — it is one branch of a period-2 cycle whose other branch is 9.867. The subsection after next shows the cycle is in the wind, not in the composition. The $K_{zz} = 10^9$ case converges (1.18×10^{-4} at pass 10) and moved by +0.21 dex against M3; the next subsection attributes that.

- The wind itself is better converged than M3's: the spread of $r^2 \rho v$ over $1.1\text{--}2 R_p$ is 1.15×10^{-2} ($K_{zz} = 0$) against M3's 1.81×10^{-1} .

70.6 Where the +0.21 dex comes from: metal transport, tested directly

backup/phase_d_baseline/new_kzz1e9_nometdiff is new_kzz1e9_fix with one key changed, He_metal_diffusion: False, from the same restart files and on the same route. It is the control that separates the two halves of the fix: helium transport and relaxation on one side, metal transport on the other. Note what the flag does and does not do — with metal diffusion off the metals are not reset to the reservoir ratio, they are **frozen** at whatever the restart file carried, which here is the M3 profile including its base hole. So the control is “the new helium side, the old metal profile”.

	M3 (new_kzz1e9)	metal diffusion off	full fix
outer passes / final drift	5 (ceiling) / 3.19×10^{-2} rel	3 / 1.23×10^{-4} abs	10 / 1.18×10^{-4} abs
$\log_{10} \dot{M}$ [g s ⁻¹]	9.7910	9.8160	9.9976
He 10830 red depth [%]	9.8118	9.1081	8.2115

C/H relative to the base cell (O/H is the same to four digits in the first eight rows), first eight rows and then interpolated:

	row 1–3	4	5	6	7	8	1.05	1.10	1.20	1.50	2.00	3.00	4.00
M3	1.0000	0.5694	0.6266	0.6213	0.6255	0.6379	0.628	0.696	0.953	0.975	1.000	0.944	0.573
metal diffusion off	1.0000	0.5694	0.6266	0.6213	0.6255	0.6379	0.628	0.696	0.953	0.975	1.000	0.944	0.573
full fix	1.0000	1.0000	1.0000	1.0000	0.9999	0.9999	0.909	0.668	0.569	0.551	0.569	0.565	0.426

The metal-diffusion-off column is identical to M3's, which is the check that the flag froze the restart profile as intended, and the (He/H)/HeH base cells of that run are 1.0000 across, which is the check that the helium side of the fix is fully active in it.

The attribution holds. Switching the metal transport off returns $\log_{10} \dot{M}$ to 9.8160 and the He 10830 depth to 9.11%, i.e. essentially back to M3's 9.7910 and 9.81% and away from the full fix's 9.9976 and 8.21%. Of the +0.207 dex, +0.182 is the metal transport and +0.025 is the helium side. (The decision rule as posed had this the other way round: if the metals were *not* responsible, switching them off would have left $\dot{M}$ near 9.998. It did not.)

One correction to the wording above: the dominant part of that +0.182 dex is not the base hole. This control cannot split the metal change into its two pieces — the base drain and the settling above the base that the removed cap previously forbade — but the base hole spans five cells at $1.0004\text{--}1.0010 R_p$ while the settling changes C/H from $\sim 0.95\text{--}1.00$ to ~ 0.55 over $1.2\text{--}4 R_p$, where nearly all of the metal-line cooling column sits. On the size of the regions the settling has to be the larger term. It is a reasoned inference, not a measurement.

70.7 The $K_{zz} = 0$ non-convergence is the wind, not the composition

The outer Picard loop — wind at fixed composition, composition at fixed wind — has no damping of its own, so the composition update is now under-relaxed,

$$X \leftarrow X_{\text{old}} + \omega (X_{\text{relaxed}} - X_{\text{old}}) \quad (21)$$

with ω starting at 0.5 and halved (floor 0.125) on any pass whose drift failed to fall. The blend is projected back onto the species vector by the same nonnegative projection element_diffusion_step ends on, so both element totals are still met exactly. The drift that the convergence test and the halving rule read is the **undamped** distance $\max |X_{\text{relaxed}} - X_{\text{old}}| / X_{\text{base}}$, so a small ω cannot buy a false convergence by making the applied step small. EXHALE_DIFF_OMEGA=<val> pins the starting value for experiment, and EXHALE_DIFFUSION_CHECK=1 now also writes output/diffusion_pass_profiles.txt, one composition and mass-flux profile per outer pass — the question a non-converging loop raises is about the profiles, not about the scalar drift.

It does not converge, and the reason is not the composition. Two runs from the M3 restart files, both $K_{zz} = 0$:

	start $\omega = 0.5$ (new_kzz0_fix)	start $\omega = 1$ (new_kzz0_undamped)
drift by pass	9.97×10^{-1} , 4.99×10^{-1} , then the pass-3 JFNK stagnates (info = 2, line search collapses to $\lambda = 9.5 \times 10^{-7}$ at $r = 1.009$, energy row) and the loop exits on its solver-failure guard	falls to $\omega = 0.125$ by pass 11 and then holds a period-2 cycle: 7.62×10^{-2} , 6.84×10^{-2} , 8.10×10^{-2} , 7.31×10^{-2} , 8.43×10^{-2} , 7.60×10^{-2} , 8.95×10^{-2} , 7.97×10^{-2} (passes 13–20)
final state	not converged, $\ R\ = 2.4 \times 10^{-3}$, $\log_{10} \dot{M} = 9.838$	20 passes, every one info = 0

The profiles written each pass say what is cycling. Over passes 3–20 the median $r^2 \rho v$ in the escape window alternates between **2.1–2.2** $\times 10^{-7}$ on odd passes and **5.9–6.8** $\times 10^{-7}$ on even ones — a factor 2.8 in mass flux — while each branch is individually flat to 10^{-3} and every JFNK solve returns info = 0 with $\|R\|$ between 9.0×10^{-5} and 9.6×10^{-4} , all inside the 10^{-3} target. The composition, meanwhile, barely moves between the two:

$r [R_p]$	odd (pass 19)	even (pass 20)
1.010	0.9968	0.9974
1.05	0.9645	0.9674
1.10	0.9065	0.9123
1.20	0.8816	0.8893
1.50	0.8768	0.8849
2.00	0.8815	0.8894
3.00	0.8785	0.8865
4.00	0.8219	0.8332
$\log_{10} \dot{M}$	9.867	10.376

The largest odd-even difference in (He/H)/HeH above $1.05 R_p$ is 0.0117, and each parity is nearly stationary from one occurrence to the next (0.0028 odd to odd, 0.0024 even to even). So this is not a composition chasing a wind: it is **two wind solutions, 2.8 $\times$ apart in mass flux, that the residual test cannot tell apart, sitting on compositions that differ by $\sim 1\%$** . Damping the composition cannot remove it, and the measurement above is the evidence — ω reached its 0.125 floor and the cycle held at full amplitude.

Two things follow, neither of them settled here. The $\log_{10} \dot{M} = 10.396$ quoted for this case in the table above is one branch of that cycle (the other is 9.867, close to the 9.998 of the converged $K_{zz} = 10^9$ case and to the 9.816 of the

metal-diffusion-off control). And a residual norm that accepts two states a factor 2.8 apart in $\dot{M}$ is a weak convergence test for this configuration, which is the standing caveat behind the flux-based convergence decision of section 19 ($\|R\|$ is reported for reference only) rather than anything the diffusion introduced.

70.8 Verification

diffusion_tests.x: 16/16 (T11 new; the other 15 unchanged). make check: 5/5 byte-identical — no regression case turns He_diffusion on, and the he_diffusion guard at the top of the operator means the flag-off path never enters any of this.

71. The friction the ionized wind actually has: stage-resolved binary diffusion (2026-08-25)

Element diffusion used one coefficient over the whole domain, the neutral Banks & Kockarts hard sphere. That is the wrong friction above the ionization front, and it is wrong in the direction that matters: a helium or metal *ion* moving through a proton gas is held by Coulomb collisions, whose momentum-transfer cross section at $T \sim 10^4$ K is $\sim 10^{-12}$ cm² against the $\sim 10^{-15}$ cm² of a hard sphere. With a neutral coefficient everywhere the solver let helium and every metal keep settling through the ionized wind, and once the $f_X \leq f_X(\text{base})$ limiter went away (section 70) nothing stopped them: on the HD 209458 b $K_{zz} = 10^9$ wind iron and calcium reached *zero* by $1.1 R_p$ and magnesium 2.5% of its base ratio by $4 R_p$.

Koskinen et al. (2013, section 2.1) solve the same Chapman & Cowling equation with collision terms that “account for neutral-neutral, resonant and non-resonant ion-neutral, and Coulomb collisions”, and report (their section 3.2.2) that Coulomb collisions “are much more efficient in preventing diffusive separation than collisions with neutral H.”

This is a correction to the physics, not an option: there is no input key for it.

71.1 What the coefficient is now

Each element is resolved into its ionization stages and the friction is built pair by pair, every pair being the Chapman-Enskog first approximation $D = 3kT / (16n\mu\Omega^{(1,1)})$ evaluated with the cross section that pair actually has (binary_element_diffusion.f90):

pair	routine	coefficient	scaling
neutral-neutral	hard_sphere_pair_diffusion	$1.52 \times 10^{18} (1/A_s + 1/A_t)^{1/2} T^{1/2} / n$	$T^{1/2} / n$
ion-neutral, non-resonant	ion_neutral_pair_diffusion	$1/D = 1/D_{\text{pol}} + 1/D_{\text{hs}},$ $D_{\text{pol}} = kT / (2.21 \pi e n (\alpha_n \mu)^{1/2})$	T/n then $T^{1/2}/n$
ion-ion	coulomb_pair_diffusion	$3(kT)^{5/2} / (4(2\pi\mu)^{1/2} n (Z_s Z_t e^2)^2 \ln \Lambda)$	$T^{5/2} / n$

and the element-element coefficient is the stage-fraction harmonic sum (frictions add; Blanc’s law across stages)

$$\frac{1}{D_{\text{eff}}} = \sum_{s \in \text{element A}} \sum_{t \in \text{element B}} \frac{y_s y_t}{D_{st}}. \quad (22)$$

Resonant charge exchange $H^+ + H$ is deliberately absent: both partners carry the same element, and a binary diffusion coefficient is driven by the friction *between* the two elements, so every ion-neutral pair that appears is non-resonant. HeH^+ carries both elements and stays out of both carrier lists, as it already was out of the mean charges.

The ion-neutral pair carries two channels. The polarization (Langevin) coefficient is the long-range limit: it keeps the induced-dipole attraction and no repulsive core, so its friction falls as T^{-1} where a rigid core would hold it at $T^{-1/2}$. Taken alone it would make an ion-neutral pair *more* mobile than a neutral-neutral one of the same masses above $\sim 1.5 \times 10^3$ K, which is unphysical — opening a second channel cannot make a pair more mobile. The two are momentum-transfer cross sections of the same encounter, so to first order their $Q^{(1)}$ add, hence their collision integrals add, hence their frictions add; and since $D = 3kT/(16n\mu\Omega^{(1,1)})$ is linear in $1/\Omega$, adding frictions is adding inverse coefficients, $1/D_{in} = 1/D_{pol} + 1/D_{hs}$. That is the same rule the stage mixture uses. Both limits come out exactly: $D_{in} \rightarrow D_{pol}$ where the polarization channel dominates and $\rightarrow D_{hs}$ where the core does, crossing at **1504 K** for an H/He pair. At the 10^4 K of the transition layer D_{pol} is 2.6–4.7 times D_{hs} and D_{in} sits at 0.82 of D_{hs} , i.e. 0.18–0.28 of the polarization value alone.

Sources. The hard-sphere form is the one already in the code (Banks & Kockarts 1973). The polarization constant 2.21π is the one behind the non-resonant ion-neutral collision frequency of Schunk & Nagy (*Ionospheres*, eq. 4.88) and the Coulomb coefficient is the standard Chapman-Enskog result (Paquette et al. 1986; Schunk & Nagy eq. 4.142 give the equivalent collision frequency). Neither textbook is in references/, so **both numerical constants were re-derived from the Chapman-Enskog collision integral rather than copied**, and the derivation is written into the comment beside each routine. The check that the chain is consistent is that the rigid-sphere collision integral put through the same Chapman-Enskog expression returns the Banks & Kockarts prefactor exactly, with a collision diameter of $2.7 \text{ \AA}$.

Polarizabilities are the recommended static dipole values of Schwerdtfeger & Nagle (2019), in atomic units, with $\alpha(H) = 4.5 a_0^3$ the exact nonrelativistic ground-state value; H and He come out at 0.667 and $0.205 \text{ \AA}^3$, the values Schunk & Nagy tabulate for aeronomy.

Validity caveat, recorded at the coefficient. The polarization limit keeps the induced-dipole attraction and no repulsive core, so its friction falls as T^{-1} where a rigid core would hold it at $T^{-1/2}$. It is the smaller of the two coefficients (stronger friction) below $\sim 1.5 \times 10^3$ K and the larger above — by 2.6 to 4.7 at 10^4 K, where the true coefficient is the smaller because the core contribution adds to the collision integral. The consequence is confined to the partially ionized transition layer; both asymptotic limits are exact.

71.2 Two defects fixed with it

- **The trace-metal loop assumed the hydrogen-plasma ambipolar field.** It built its settling coefficient from $\Delta m_{\text{eff}} = (A_X - 1) - 0.5(\bar{Z}_X - \bar{Z}_1)$, the constant $eE = m_H g/2$ of a pure proton plasma — the same assumption decision D3 removed for helium at M2, still standing here. The loop now reads the field settling_coefficient

computes from the electron pressure gradient, so there is one ambipolar field in the module and no second, cruder copy. With He_ambipolar off the field is zero and the loop reduces to the pure-gravity form as before.

- The metal loop also used a single hard-sphere X -in-H coefficient; it now runs the same stage mixture, metal stages against hydrogen carriers.

71.3 Tests

src/tests/diffusion_tests.f90 gains **T12**, the three limits of D_{eff} on a prescribed isothermal column at 10^4 K, each against a closed form written out independently in the test rather than against the module's own pair routines: all neutral reproduces Banks & Kockarts to 1.2×10^{-16} relative, fully ionized reproduces the $\text{He}^{++}/\text{H}^+$ Coulomb coefficient to round-off, half ionized reproduces the four-term Blanc average (whose two ion-neutral terms are themselves the combined form) to round-off. T12d adds the two temperature limits of that combined form: within 2.5×10^{-2} of the polarization channel at 1 K, within 1.2×10^{-3} of the hard sphere at 10^9 K, and never above either. ./diffusion_tests.x: **21 passed, 0 failed**. make check: **5/5 byte-identical** — the operator is entered only with He_diffusion on and no regression case sets it.

EXHALE_DIFFUSION_CHECK=1 adds D_{eff} , the neutral coefficient of the same cell, their ratio and the dominant stage pair to diffusion_faceflux.txt.

71.4 Measured effect

The full tables are docs/binary_diffusion_design.md section 9.3. In short, on two winds restarted from the same converged state and Newton-finished:

- $D_{\text{eff}}/D_{\text{neutral}}$ is 1.00 in the neutral base, 0.16 at the front ($1.2 R_p$), 9.6×10^{-3} at $1.5 R_p$ and 7.6×10^{-5} at $4 R_p$ on HD 209458 b — the suppression *grows* outward because the Coulomb coefficient carries $T^{5/2}$ against the hard sphere's $T^{1/2}$.
- On HD 209458 b every metal mixing ratio becomes **flat above the front**, which is what a steady wind must do once diffusion loses to advection. Iron goes from 0.000 to 0.107 of its base ratio and calcium from 0.000 to 0.019 — the difference between having and not having Ca II H&K and Fe II in the transmission spectrum. $\log_{10} \dot{M}$ moves $10.0106 \rightarrow 10.0561$ (+0.046 dex) and the He 10830 red depth $8.20\% \rightarrow 7.85\%$. The control is a run of the same configuration from the same converged state with a binary built from this commit's parent, so the friction is the only difference.
- The ion-neutral channel decides how much of that survives. Between 1.0 and $1.2 R_p$ hydrogen and helium are still neutral (measured $x(\text{H}^+) = 0.000, 0.004, 0.072$ at $1.05, 1.10, 1.20 R_p$), so the combined form barely touches the *helium* profile — but the low-ionization-potential metals are already ionized there (Fe 0.999, Mg 0.943, Ca 0.269, C 0.329, Na 0.053, O 0.003 ionized at $1.10 R_p$), and going from the polarization channel alone to the combined form raises their frozen plateaus in exactly that order: Fe $\times 3.8$, Ca $\times 2.3$, Mg $\times 1.23$, C $\times 1.07$, Na $\times 1.00$, O $\times 0.97$.

- On LHS 1140 b at $\text{He}/\text{H} = 0.55$ nothing moves: $\log_{10} \dot{M}$ 7.6065 both, the temperature identical to 9×10^{-16} relative, He 10830 within 0.2%, and the composition relaxation exits on its first pass. That wind separates its helium inside the first cell above the base, where the gas is neutral and $D_{\text{eff}} = D_{\text{neutral}}$ by construction, and above that there is no helium left for the Coulomb suppression to hold.

72. The advection correction is gated on the slowest species, not on hydrogen (2026-08-25)

`post_process_adv.f90` replaces the local ionization balance of a cell by the steady advection-ionization ODE, and skips cells where that replacement can carry no information. One of the three skip conditions is a Damkohler number,

$$\text{Da} = \frac{\Delta r}{v} (P_{\text{HI}} + \alpha_{\text{HII}} n_e) > 100 \quad \longrightarrow \quad \text{keep the equilibrium state} \quad (23)$$

built from the **hydrogen** ionization/recombination rate. The correction it gates, however, carries the whole solved species vector — $\text{H I}/\text{H II}$, $\text{He I}/\text{He II}/\text{He III}$ and, when `Include He23S?` is on, the $\text{He } 2^3\text{S}$ metastable — so one species' clock decided the fate of all of them. The metastable relaxes orders of magnitude more slowly than hydrogen ($A_{31} = 1.272 \times 10^{-4} \text{ s}^{-1}$ alone sets a 2.2-hour floor, against sub-second photoionization for H in an irradiated wind), so wherever hydrogen was equilibrated and the metastable was not, the gate froze $n(2^3\text{S})$ at its equilibrium value in cells where it is in fact advected.

That is physically wrong independently of how often it bites, and it bit visibly on the LHS 1140 b run `heh10_diff_kzz1e8` ($\text{He}/\text{H} = 10$), where the equilibrium solution has a sharp hydrogen ionization front: 115 of 503 cells above the front kept the equilibrium ionization, and the `_adv` metastable profile stepped from 5.15 to 223 cm^{-3} across the single cell boundary $r = 6.083 \rightarrow 6.172 R_p$ (1.64 dex). The transit synthesis reads the `_adv` profiles, so the run reported a red-pair depth of 83.7% and $\text{EW} = 23.4 \% \text{ \AA}$ — a number set by the discontinuity. The limitation was recorded in `LHS1140b/kzz_decision.md` section 6.2 and the run was excluded from that memo's crossing fit.

72.1 What the gate is now

The advection systems (`adv_implicit_H`, `adv_implicit_HeH`, `adv_implicit_HeH_TR`) solve the whole H/He vector at once, so there is no gate to be had one row at a time: the honest statement is that a cell may be pinned to equilibrium only when **every** population it solves is equilibrated. The Damkohler number is therefore formed with the **slowest** relaxation rate of the solved set, each rate being the total rate at which that population is destroyed and re-formed:

population	relaxation rate ν
$\text{H I} / \text{H II}$	$P_{\text{HI}} + (a_{\text{ion_HI}} + \alpha_{\text{HII}}) n_e$
$\text{He I} / \text{He II}$	$P_{\text{HeI}} + (a_{\text{ion_HeI}} + \alpha_{\text{HeII}} + \alpha_{\text{HeI23S}}) n_e$
$\text{He II} / \text{He III}$	$P_{\text{HeII}} + (a_{\text{ion_HeII}} + \alpha_{\text{HeIII}}) n_e$
$\text{He}(2^3\text{S})$	$A_{31} + P_{\text{HeITR}} + (q_{31a} + q_{31b} + a_{\text{ion_HeITR}}) n_e + Q_{31} n_{\text{HI}}$

$\text{Da} = (\Delta r/v) \min(\nu)$, compared with the same threshold of 100. The $\text{He}(2^3\text{S})$ row is term by term the loss side of `fvec(4)` of `adv_implicit_HeH_TR` — radiative decay, photoionization, collisional de-excitation to 2^1S and 2^1P ,

electron-impact ionization, and Penning ionization on neutral H — and every coefficient is read from the arrays `HeITR_coeffs` and `eval_cool` already fill in the same routine. No rate is redefined. The He I row now also carries `rcheiTR`, because with the triplet on `HeITR_coeffs` overwrites `rcheiIB` with the singlet channel alone and the He I row of the residuals adds the two. The two other skip conditions — inflow $v \leq 0$ on either face, and $x_{\text{HII,eq}} < 10^{-6}$ — are untouched, so the breathing-base guard behaves exactly as before.

The hydrogen rate now includes collisional ionization `a_ion_HI n_e`, which the old expression omitted, for consistency with the three rows beside it. It can only raise Da_H , and Da_H is no longer the quantity that decides anything on its own.

72.2 Measured

Diagnostic build dumping Da per cell, LHS 1140 b `heh10_diff_kzz1e8`, post-processing pass 10:

$r [R_p]$	Da_H , H-only gate	Da_H , new binary	Da of the slowest species
6.083	6.1×10^{-4}	6.1×10^{-4}	6.1×10^{-4}
6.172	3.2×10^2	6.0×10^{-4}	2.8×10^{-3} (H-only run), 6.0×10^{-4} (new)
6.946	7.8×10^5	4.8×10^{-4}	$2.3 \times 10^{-3} / 4.8 \times 10^{-4}$
7.602	4.6×10^7	1.1×10^2	$2.1 \times 10^{-3} / 1.7 \times 10^{-3}$
7.956	1.7×10^8	7.2×10^2	$1.9 \times 10^{-3} / 1.6 \times 10^{-3}$

The metastable Damkohler number is $\sim 2 \times 10^{-3}$ throughout that region: the $\text{He } 2^3\text{S}$ population there is advection-dominated by five orders of magnitude, in exactly the cells the hydrogen gate pinned to equilibrium. The two Da_H columns differ because the old gate is self-reinforcing — a pinned cell holds the fully ionized equilibrium state, whose n_e and P_HI then keep Da_H above the threshold for the cell above it.

Effect on that run (both binaries run from the same converged state with `Do_only_PP: True`, in scratch copies; the case directory itself was not re-run):

quantity	H-only gate	slowest-species gate
cells kept at equilibrium	115 of 503	15 of 503
outermost such cell	$30.0 R_p$	$1.0034 R_p$
$\max \text{dlog}_{10} n(2^3\text{S}) $ between adjacent cells, $r > 1.05$	1.636 ($5.15 \rightarrow 223$ at 6.13)	0.386 ($4.99 \rightarrow 12.1$ at 7.43)
He 10830 red depth [%]	83.69	21.52
red EW [% Å]	23.39	5.69
$\log_{10} \dot{M}$ [g/s]	7.76	7.76 (the wind solution is not re-run)

The remaining 15 pinned cells all sit at $r \leq 1.0034 R_p$, the shielded base where conditions (i) and (iii) are meant to act. The largest remaining variation in the metastable profile, 0.386 dex over $r = 7.38 \rightarrow 7.49 R_p$, is not a switch: it is where the advected hydrogen finally ionizes (x_HI falls from 0.80 to 5×10^{-4} over three cells as Da_H crosses 1) and the Penning loss `Q31 n_HI` on the metastable switches off with it. The `_adv` profiles are continuous across it, and `hybrd1` returns `info = 1` in every cell of the region, including those with $\text{Da}_\text{H} \sim 10^5$.

The corrected run's EW of 5.69 % Å is still above what the He/H = 2.7–4.0 sequence of `kzz_decision.md` section 6.1 extrapolates (~ 2.8 % Å at He/H = 10 on that sequence's local log-log slope of 0.73), so the He-rich end appears to be

leaving the regime that sequence describes for its own reasons; what the fix removes is the discontinuity, not the run's distance from the trend.

72.3 Effect on existing results

Three reference cases were re-post-processed with both binaries in scratch copies — LHS1140b/exhale/heh0p55 (diffusion off), heh2p13_diff_kzz1e9 (one of the crossing brackets), and the H-rich standard backup/phase_d_baseline/new_kzz1e9_d3b (HD 209458 b, $K_{zz} = 10^9$, metals on). In all three, Hydro_ioniz_adv.txt and Ion_species_adv.txt come out **byte-identical** to the H-only-gate binary, and the number of pinned cells is unchanged (42, 30 and 153 respectively): in those runs the pinned cells are held by the inflow or $x_{\text{HII,eq}}$ condition, which the change does not touch, not by the Damkohler condition. He 10830 red depth and EW therefore move by 0.00%, and no stored result needs regenerating.

make check passes 5/5 — as it must, since the goldens compare Hydro_ioniz.txt and Ion_species.txt, which the post-process never writes; running it is the check that the change stays inside the post-process.

72.4 Re-post-processing the stored LHS 1140 b scans (2026-08-26)

The three reference cases above were the wrong sample for the He-rich end. All eighteen diffusion-off cases of LHS1140b/exhale/ — the eleven of the GJ 1132 composition scan and the seven _gj699 companions — were therefore re-post-processed in place, each from its own untouched JFNK solution of record (output/*_IC.txt) through one Do only PP: True pass and both transit syntheses; the pre-fix products are kept beside each case in pre_gate72_adv/. Five cases changed, all helium-dominated: heh10, heh100, heh1000, heh10000 and heh1000_gj699 lose 89–93% of their red-pair depth and equivalent width (heh10 75.97% and 21.53 % Å to 7.42% and 2.270 % Å; heh1000_gj699 105.1% and 35.87 % Å to 9.53% and 3.479 % Å, the fitted depth above 100% being the pre-fix saturated trough). The count of _adv cells still holding exactly the equilibrium metastable density falls from 173–437 of 504 to 13–22 in those five. The remaining thirteen move by at most 0.06% and keep their cell counts exactly, so the earlier sample's conclusion holds wherever the hydrogen front is not sharp; the 0.06% is a drift in the equilibrium solve between the binary that wrote those outputs (2026-08-24) and this one, at most 4×10^{-4} relative, not the gate.

Nothing the memo docs/lhs1140b_exhale_vs_pwinds.tex concludes moves: the equivalent-width crossing stays at He/H = 0.550 (0.541 with turbulence) on the GJ 1132 proxy and 0.060 on the GJ 699 one, because both are interpolated inside $\text{He}/\text{H} \leq 1$, and so does the matched broadening kernel, 22.30 km s⁻¹. What moves is the He-rich end of the two composition tables, the bump table's He/H = 1000 rows, and the figures built from them. heh10_diff_kzz1e6, the He-rich bracket of the K_{zz} crossing table, was checked in a scratch copy and is unchanged, so tab:kzzcross stands as published; heh10_diff_kzz1e8 was already excluded from it.

One artifact is not the gate's and survives it: heh100 still steps 3.8 dex in $n(2^3\text{S})$ across a single cell at $r = 1.059 R_p$, on the hydrogen ionization front, where the base velocity is still sign-changing. The shell is thin enough (0.008 R_p against a 6.8 R_p stellar radius) that it can account for at most 0.04% of that case's 5.13% red depth.

73. Milestone M4: element diffusion through the molecular region, and the end of the molecular / He_diffusion exclusion (2026-08-26)

input_read.f90 refused Molecular chemistry: True together with He_diffusion: True — an error stop that docs/binary_diffusion_design.md section 5 put there deliberately, to be lifted by a validated closure and not by deleting the check. The closure is now built and the check is gone.

73.1 What the molecular region changes, and what it does not

Equation (1) of the memo transports **elements**, and the chemistry moves hydrogen between H, H⁺, H₂, H₂⁺ and H₃⁺ without changing the hydrogen element's mass, so the transport equation is the same above and below the molecular front. What the species set changes is the closure: helium's collision partners are then a mixture of carriers, and gravity, the ambipolar field and the friction all act on **particles**, not on nuclei. Four things follow, three of them already anticipated by section 5 and one not.

1. **Friction: Blanc's law over the carriers.** The stage sum $1/D_{\text{eff}} = \sum_s \sum_t y_s y_t / D_{st}$ of section 2.6 already ran over the carrier lists {HI, HII, H2, H2⁺, H3⁺} and {HeI, HeII, HeIII, HeTR}, each carrier with its own mass, charge and polarizability (H₂ enters at mass 2 and with $\alpha(\text{H}_2) = 5.315 a_0^3 = 0.787 \text{ \AA}^3$, the Schwerdtfeger & Nagle 2019 recommended value, alongside the H and He entries that were already there). It reduces to $D(\text{He}, \text{H})$ in the atomic region and to $D(\text{He}, \text{H}_2)$ in a fully molecular one.
2. **Forces per carrier.** G now uses the mean carrier mass $m_{c1} = m_1 * m_{\text{car}}$ and the mean carrier charge per carrier, where m_{car} is the hydrogen nuclei bound into one collision partner. In the atomic region $m_{\text{car}} = 1$ and this is the old nucleus form exactly; in an H₂ base the relative settling mass is $4-2 = 2$ instead of $4-1 = 3$. The trace-metal loop settles a metal against the same mean carrier mass.
3. **The carrier density in the pair coefficients.** n in the Chapman-Enskog $D_{st} = 3kT / (16n\mu\Omega^{(1,1)})$ is a density of *colliding particles*. It was the nucleus density; it is now the carrier density, which is the same number wherever the hydrogen is atomic and smaller below a molecular front.
4. **The mole-fraction driver the chemistry carries** — the term section 5 did not state, and the one that makes the other three a force rather than bookkeeping. The transported variable is the mass fraction X ; the Chapman-Cowling driving force is the gradient of the **mole** fraction x among the carriers. With $\psi = n_{1c}/n_{\text{H}}$ (1 atomic, 1/2 fully H₂),

$$\text{logit}(x) = \text{logit}(X) + \ln\left(\frac{m_1}{m_{\text{He}}\psi}\right), \quad (24)$$

$$\frac{d\text{logit}(x)}{dr} = \frac{d\text{logit}(X)}{dr} - \frac{d\ln\psi}{dr}. \quad (25)$$

The gradient coefficient is untouched — $A = (m_1 m_{\text{He}} n / \bar{m}) D (dx/dX)$ collapses to ρD with the carriers exactly as it does with the nuclei, for any ψ , the two carrier factors cancelling — so the whole content of the term is one addition to the settling coefficient, $G \rightarrow G - d\ln\psi/dr$. The discretization, the Péclet hybrid and the M-matrix argument are unchanged, and the term is identically zero wherever the hydrogen is atomic.

It is a real driver: where hydrogen turns molecular going down, each nucleus is spread over fewer collision partners, helium's mole fraction rises, and helium diffuses down that gradient — outward across the molecular front — even at a uniform mass fraction. Its size is set by the sharpness of the front, and across a front where ψ runs $1/2 \rightarrow 1$ it is a factor 2 in the He/H nucleus ratio (test T7d below).

The eddy term carries none of this: eddy mixing transports the mixture as a whole and has no preferred species, so it acts on dX/dr alone.

73.2 T7, and the gate it closes

Four checks in `src/tests/diffusion_tests.f90`, each against a closed form typed out in the test rather than read back from the module:

id	check	result
T7a	forced atomic column: D_{eff} vs $D(\text{He}, \text{H})$	1.13×10^{-16} relative
T7b	forced fully molecular column: D_{eff} vs $D(\text{He}, \text{H}_2)$	1.57×10^{-16} relative
T7c	prescribed molecular column, three constant K_{zz} : the homopause read off the relaxed profile vs the radius where $D_{\text{eff}} = K_{zz}$	$3.4 \times 10^{-3} R_p$, about half a cell, at all three
T7d	same column with $g = 0$, where $-d \ln \psi / dr$ is the only driver left	the mole fraction levels out to 6.2×10^{-5} while the He/H nucleus ratio moves by 0.999

T7c reads the homopause off the profile without assuming where it is: at the zero-flux steady state of a motionless column, $\rho(D+K)dX/dr = -\rho X(1-X)DG$, so $-[d \log(X)/dr]/G$ is exactly $D/(D+K)$ and the radius where it passes $1/2$ is the radius where $D = K$. The three K_{zz} were chosen to put that radius at 1.5 , 2.0 and $2.5 R_p$, and the measured profile put it at 1.4993 , 2.0013 and 2.5010 .

T7d is the sharpest statement of the closure: with no gravity, no field and no eddy, diffusion levels the **mole** fraction of helium among the collision partners, not its mass fraction, so a column that is molecular below and atomic above ends with an He/H nucleus ratio that differs by the factor ψ across the front and a mole fraction that is flat.

make diffusion_tests && ./diffusion_tests.x: **28 passed, 0 failed** (21 before this change; the four new ones and T7c three times). Rebuilt with `-O1 -g -fcheck=bounds,do,mem -ffpe-trap=invalid,zero,overflow` — the new carrier arrays are new array bounds — the same 28 pass with no runtime error.

73.3 The exclusion, and the restart

input_read.f90 no longer refuses the pair; the paragraph there now says why it does not. docs/input_schema.md, README_HOWTO.md and docs/oxygen_chemistry_new_plan.md are corrected where they stated the exclusion.

load_IC.f90 refused to restart *any* molecular state whose composition differed from the input He/H, because HeH+ carries a nucleus of each element and a single rescaling factor cannot set both counts. With He_diffusion on the only cells rescaled at all are the base and its inner ghosts (the column above keeps its diffused split, section 7.4 of the memo), so those cells are now projected onto their two element totals the same way binary_element_diffusion

writes back — HeH+ by the smaller of the two factors, the shortfall deposited into the neutral ground species. Without He_diffusion the whole column would have to be rescaled and the refusal stands.

Found while reading that block and fixed with it: the helium nucleus count load_IC builds to decide whether a restart file's composition matches the input left **HeTR out** (HeI + HeII + HeIII + HeH+), while the rescaling a line below scales HeTR like every other helium species. The metastable triplet carries a helium nucleus — $\text{bsp_nHe(HeTR)} = 1$, and every other element count in the code includes it — so a restart of a He23S state read its own He/H slightly low and rescaled by that error. It is now in the count. No golden moves: no regression case restarts.

73.4 End to end

backup/regression/mol_base_handoff (the hot Uranus, Molecular chemistry + a base.inp handoff) re-run in a scratch copy with He_diffusion: True added, EXHALE_MAXSTEPS=12000 as in the regression:

- runs to completion with no NaN and no negative density;
- the two-component mass closure $m_1 n_H + m_{\text{He}} n_{\text{He}} = \rho$ holds to 4.4×10^{-16} relative over the whole grid, read back out of Ion_species.txt and Hydro_ioniz.txt;
- $\log_{10} \dot{M} = 10.58$, the same as the diffusion-off case at the same step count;
- (He/H)/HeH stays within 1.3×10^{-3} of the reservoir everywhere.

The flat composition is the physics of this planet, not an inert operator. He_Kzz = 1e9 comes from the case's own base.inp and is three decades above the $D_{\text{eff}} \sim 8 \times 10^5 \text{ cm}^2/\text{s}$ of the molecular base, so the column is eddy-mixed; but the same run repeated with He_Kzz = 0 and Kzz_base 0.0 comes out flat as well, within 4×10^{-3} of the reservoir, and the reason is in the diagnostic: at the base face the diffusive helium flux is $-3.1 \times 10^{-13} \text{ g cm}^{-2} \text{ s}^{-1}$ against an advective -3.5×10^{-7} , six orders of magnitude smaller. This is a hot Uranus losing $10^{10.6} \text{ g/s}$; the wind sweeps the composition along faster than either diffusion or settling can separate it. What the run demonstrates is that the operator carries the molecular composition without NaN, without a negative density and with the element budget closed to round-off — the separation physics itself is exercised by T7 and by the atomic-region cases of sections 69–71.

That the molecular carriers are the ones being used is visible in diffusion_faceflux.txt (EXHALE_DIFFUSION_CHECK=1) of the same run, which names the stage pair carrying most of the friction:

$r [R_p]$	dominant pair	relative settling mass	$D_{\text{eff}}/D_{\text{neutral}}$
1.000	HeI-H2	2.0005	0.7747
1.107	HeI-H2	2.062	0.7855
1.169	HeI-H2	2.520	0.8775
1.196	HeI-HI	2.923	0.9783
1.212	HeI-HI	3.000	0.9999
1.691	HeII-HII	3.067	2.52×10^{-3}

The relative settling mass runs from 2 to 3 across the molecular front, exactly the carrier weight of point 2 above, and $D_{\text{eff}}/D_{\text{neutral}}$ from $\sqrt{0.75/1.25} = 0.7746$ — the He-H₂ hard sphere against the He-H one — to 1, before the Coulomb suppression of section 71 takes over above the ionization front.

73.5 Regression

make check: 5/5 byte-identical at the time of the change. The operator is entered only with He_diffusion on, which no case set then, and every molecular quantity added here (carrier lists, mean carrier mass, carrier density, $d\ln\psi/dr$) collapses onto its nucleus form identically when the molecular columns are zero.

73.6 The diffusion regression case (same day)

Byte-identity of the old cases only shows that the operator stays out of the way when it is off; nothing in make check ran it at all (docs/binary_diffusion_design.md section 7.5 had been carrying that as an open item since the design was written). The sixth default case, backup/regression/mol_diffusion, closes it. It is mol_base_handoff — the Tier-2 hot-Uranus gate with molecular chemistry on and its pinned base.inp copied unchanged — plus He_diffusion: True and He_Kzz: 1.0e9, run as the same 12000-step relaxation snapshot. One case covers the widest path: molecular carriers below the front, the stage-resolved pairs through it, the projection back onto f_sp, and the Coulomb friction of section 71 above the ionization front, which is why this gate was chosen over the atomic HD 209458 b run the design had sketched.

The measured effect is the one the face-flux diagnostic above predicts for this planet — present but small (the same hot Uranus, whose wind sweeps the composition along faster than diffusion separates it). Against mol_base_handoff at the same step count, Hydro_ioniz.txt and Ion_species.txt both differ, while

mol_base_handoff	final: count=12000	du= 2.6292E+00	dtu= 5.7467E-03	log10 Mdot = 10.58
mol_diffusion	final: count=12000	du= 2.6290E+00	dtu= 5.7474E-03	log10 Mdot = 10.58

so the golden pins the operator's arithmetic, not a large dynamical signature. make check is 6/6 with the five existing goldens untouched.

74. What a base.inp key is allowed to do: five categories, and elemental reservoirs that reach melem_ab (2026-08-26)

base.inp carried six keys and no statement of what any of them meant for the wind. That is P2 of docs/oxygen_chemistry_new_plan.md, and the reason it matters is not tidiness: q_H2_base had been described in places as pinning the base composition, when all it does is set the base particle count, and the handoff had no way at all to carry an element — the quantity EXHALE actually transports.

74.1 The contract

Every key now belongs to one of five categories, and the category states what the value may do. The table lives in three places that must agree: read_base.inp's header comment, docs/input_schema.md section 2c, and here.

Key	Category	Reaches
comments only	provenance	nothing (the machine-readable keys are A1a/P0, not yet built)
T_base, r_base, p_base, q_H2_base	EOS boundary	T0, R0, p_base_bar, q_h2_base
HeH_base, <El>_H_base	elemental reservoir	HeH, X_<El> → melem_ab
(none today)	initial guess	-
Kzz_base	boundary constraint	he_kzz
q_H2O, q_CO, ...	diagnostic	nothing; metadata, not physics

q_H2_base is an **EOS anchor**: comp_ntot_bc removes the H nuclei bound into H₂ at the base, and that is the whole of its effect. Nothing holds H₂ at that value, no H₂ profile is seeded from it, and the molecular network moves away from it in the first cell. The A1c species keys stay comments because no part of the code reads them.

74.2 Elemental reservoirs

<El>_H_base is new: C_H_base, N_H_base, O_H_base, S_H_base and the same form for Mg Si Ca Na K Fe, each the El/H **nuclei** ratio at the handoff level. It overrides metals.inp for that element, and it turns the metal system on by itself when it is the only nonzero abundance. Hydrogen and helium already had this: HeH_base.

The change that makes it work is one of order. thereis_metals, melem_ab and thereis_lowIP_metal were derived immediately after the core input.inp block, roughly seven hundred lines before read_base_inp ran, so nothing the handoff said about an element could have been heard. They are now derived in the composition block after read_base_inp, beside the thereis_He / HeITR / molecular reconciliation — the one place where every composition flag has its final value. Nothing between the two points touched any of them, which is why the move is a pure reordering.

read_base_inp also matches its keys the way the input.inp reader does: lbl_match (anchored, separator-terminated) instead of a bare index(line, key) > 0 substring test anywhere in the line. lbl_match/is_sep moved out of input_read's internal scope to module scope for that. One file, one matching rule.

src/utls/vulcan_to_base.py now writes the elemental keys it can support: C, N, O and S, summed over **every** carrier at the handoff level, since photochemistry moves nuclei between molecules without creating or destroying them. Mg/Si/Ca/Na/K/Fe are still metals.inp's job — metal and alkali chemistry and condensation are outside VULCAN's networks. The same carrier sum corrects HeH_base, which used to be q_He/(q_H + 2 q_H2) and so lost the hydrogen bound in H₂O, CH₄ and NH₃; on the HD 209458 b run that is 0.09698 → 0.096915, a 0.07% correction and, at 1 microbar, the whole of the difference.

74.3 The gate: do the reservoirs hold?

src/utls/element_budget.py <run_dir> answers it from the output profiles. For every element it compares, cell by cell,

$$n_{\text{El}}/n_{\text{H}} \quad \text{against} \quad (\text{El}/\text{H})_{\text{resolved}} \quad (26)$$

with n_H counting hydrogen nuclei in every carrier (HI, HII, 2 H₂, 2 H₂⁺, 3 H₃⁺, HeH⁺), n_{He} counting HeI (which already includes the triplet), HeII, HeIII, HeH⁺, and hydrogen itself closing against the mass density, $\rho = \text{mass_per_H} * n_H$. The abundances it compares against are the run's own resolved values: EXHALE_resolved.out now carries mass_per_H, ntot_bc and one abundance_<El> line per element, all of them written to full double precision so the check has no round-off floor of its own.

Measured on a copy of the mol_metals regression case (hot Uranus, molecular chemistry + solar trace metals) whose base.inp was given O_H_base 9.80e-4 — twice the metals.inp value — and S_H_base 1.32e-5, an element metals.inp does not list at all:

```
# element budget: <scratch>/mol_metals_elem
# 504 cells, profiles Ion_species.txt, tolerance 1.0e-08 (relative)
elem      reservoir max |n_El/n_H      at r[Rp] verdict
           (El/H)    - res|/res
H      rho/mass_per_H      1.982e-04      4.81412      FAIL
H*     triplet-corr.      3.056e-13      1.11481      ok
NOTE: rho - mass_per_H*n_H equals 4*n(He 2^3S) to 0.1%; the triplet mass is counted twice in calc_rho (
      HeI already includes it). The H* row removes it.
He      7.930000e-02      1.269e-12      1.11481      ok
C      2.690000e-04      8.132e-13      1.04248      ok
Ca     2.190000e-06      7.915e-13      1.04248      ok
Fe     3.160000e-05      8.177e-13      1.04308      ok
Mg     3.980000e-05      8.310e-13      1.04248      ok
N      6.760000e-05      8.037e-13      1.04308      ok
Na     1.740000e-06      8.091e-13      1.04074      ok
O      9.800000e-04      1.803e-12      1.04553      ok
S      1.320000e-05      8.220e-13      1.04074      ok

1 element(s) miss their reservoir by more than 1.0e-08
```

The nine elemental ratios close to round-off (at most 1.8×10^{-12} relative over 504 cells). Hydrogen does not, and the reason is not the handoff: the mismatch is exactly $4n(\text{He}2^3\text{S})$, the triplet mass counted a second time in calc_rho, because HeI already contains the triplet (util_ion_eq forms the singlet as nhei - nheiTR). Removing it, hydrogen closes to 3.1×10^{-13} as well — the H* row. The double count is a pre-existing defect of the mass and particle book-keeping, not of P2; it is left standing here because fixing it moves every He23S: True golden, which is a change of its own. Its size on this run is 2.0×10^{-4} of ρ at $4.8 R_p$ and zero in the molecular base, where the triplet population is.

The EXHALE_resolved.out of that run shows abundance_O = 9.80e-4 and abundance_S = 1.32e-5, i.e. the handoff overrode metals.inp for oxygen and activated sulfur.

74.4 Legacy

A base.inp without an element key behaves exactly as before: make check is 6/6 byte-identical, goldens untouched.

75. The He 2^3S double count: an excited level is not a second species (2026-08-27)

Section 74 measured it and left it standing: ρ carried the mass of every metastable helium atom twice. The He I density the code transports, `nhei` (`f_sp` column `isp_HeI`), is the **total** He I population — the solver sets it as $n_{\text{HeI}}(1 - x_{\text{HeII}} - x_{\text{HeIII}})$ and the photoionization rates recover the singlet where they need it as $n_{\text{HeIS}} = n_{\text{HeI}} - n_{\text{HeITR}}$ (`util_ion_eq`). The 2^3S column is a **level population inside that number**, not a species beside it. Every budget sum nevertheless added it a second time with mass $4m_{\text{H}}$, one gas particle and one helium nucleus.

75.1 Where the second count was

Seven places, all reading the triplet column as if it were an independent species:

Place	What was counted twice
<code>calc_rho</code> (<code>utilities.f90</code>)	$4m_{\text{H}}$ of mass $\rightarrow \rho$
<code>calc_ntot</code> (<code>utilities.f90</code>)	one gas particle $\rightarrow n_{\text{tot}}$, hence $p = (n_{\text{tot}} + n_e)T$
<code>element_ratio_HeH</code> (<code>composition.f90</code>)	one He nucleus $\rightarrow$ the He/H element ratio
<code>element_nucleus_counts</code> and the <code>gotHe</code> shortfall of <code>project_elements</code> (<code>binary_element_diffusion.f90</code>)	one He nucleus $\rightarrow$ the two element totals the diffusion operator conserves
<code>hecar_isp</code> (<code>binary_element_diffusion.f90</code>)	one collision partner $\rightarrow$ <code>carHe</code> , the helium friction
<code>nHe_l</code> / <code>gotHe_l</code> of <code>load_IC</code>	one He nucleus $\rightarrow$ the He/H a restart file is judged to carry
<code>Run._NUC_HE</code> (<code>examples/exhale_io.py</code>)	one He nucleus $\rightarrow$ <code>heh_profile</code>

`calc_ne` was never wrong: the triplet is neutral, so its `bsp_charge` is zero and the term it contributed was zero.

75.2 The fix

`species_table` gains `bsp_is_excited_level(n_bsp)`, `.true.` only for HeTR, with the rule written beside it: a flagged column is a sub-population of another column, so every budget sum skips it while the transport and the scaling of the level itself do not. The nucleus- and charge-counting loops (`composition`, `binary_element_diffusion`, `diffusion_tests`) test the flag; `hecar_isp` loses its HeTR row, whose mass, charge and polarizability were a copy of HeI's anyway; `load_IC` and `exhale_io` drop the triplet term from their helium sums.

`calc_rho` and `calc_ntot` no longer take `nheiTR` **as an argument at all**. Removing the dummy is what makes the defect unrepeatable: there is no longer a triplet number in scope to add. The `thereis_HeITR` branch of `energy_semi_implicit`'s `calc_ntot` call disappears with it.

75.3 What moved

`make check: wasp_he23off` (He 2^3S off) stays byte-identical, as it must — nothing else changed. The five He23S: True cases move, and the size of the move is the triplet mass fraction of each run:

case	max abs. rel. $d\rho$	max abs. rel. dT	$d\dot{M}/\dot{M}$	max $4n(2^3\text{S})/\rho$
wasp_full	1.24×10^{-6}	1.32×10^{-7}	-6.5×10^{-8}	1.29×10^{-6}
wasp_he23off	0	0	0	0
mol_base_handoff	4.01×10^{-4}	1.59×10^{-4}	$+3.7 \times 10^{-4}$	1.95×10^{-4}
mol_metals	2.57×10^{-4}	6.77×10^{-5}	$+2.1 \times 10^{-4}$	1.98×10^{-4}
mol_lyman_werner	1.39×10^{-4}	1.27×10^{-4}	$+4.4 \times 10^{-5}$	2.92×10^{-4}
mol_diffusion	5.35×10^{-4}	2.73×10^{-4}	$+5.4 \times 10^{-4}$	1.94×10^{-4}

WASP-121 b is hot enough that its 2^3S population is a part in 10^6 of the mass; the hot-Uranus cases carry a few parts in 10^4 . Step counts are unchanged in all six. The golden of the five moved cases were re-snapshotted (check $\rightarrow$ golden $\rightarrow$ check, 6/6 PASS); wasp_he23off was left alone.

75.4 The gate

src/utils/element_budget.py on the mol_metals outputs, before and after, same EXHALE_resolved.out:

before:	H	rho/mass_per_H	1.979e-04	at 4.81412	FAIL
after:	H	rho/mass_per_H	2.975e-13	at 1.11481	ok

Hydrogen now closes against the mass density at round-off, like the nine elemental ratios beside it, with no correction of any kind. The H* triplet-corrected row and its NOTE are gone from the script: there is nothing left to correct.

75.5 Effect on a converged wind

LHS 1140 b, heh0p55 (He/H = 0.55, He 2^3S on, metals off), a scratch copy re-finished with JFNK from the same seed under both binaries, then the same post-processing and transit pass:

	before	after
JFNK	info=0, $\ R\ = 4.326 \times 10^{-4}$	info=0, $\ R\ = 4.277 \times 10^{-4}$
$\dot{M}$ [g/s]	6.47929×10^7	$6.48008 \times 10^7 (+1.2 \times 10^{-4})$
He 10830 red depth [%]	4.279701	4.279674 (-6.3×10^{-6})

The wind is the same wind: ρ differs by at most 2.5×10^{-4} (at the base, $r = 1.0006$), the triplet mass fraction of this run peaks at 2.6×10^{-4} near $11 R_p$, and the observable is unmoved at the fifth digit. The stored run directories were not regenerated.

76. The lower atmosphere as a profile, not six scalars (2026-08-27)

Milestone E1 of docs/phase_e_flux_closure_design.md. The lower-atmosphere handoff has been a set of single-level scalars in base.inp since Tier 3, and that shape is what makes the composition an input rather than an output: HeH is one number, the diffusion operator imposes it as the Dirichlet reservoir at the base, and nothing in the handoff carries a gradient, so no flux condition between the two models can even be stated. E1 replaces the shape. It does not yet close the flux; that is E4.

76.1 The file and the key

One new opt-in key,

```
Lower atmosphere profile: lower_atmosphere_profile.dat
```

names a plain-text file: a block of # key value header lines, a # columns: schema line, then a fixed-column table running deep to shallow. Same conventions as every other EXHALE product, so `examples/exhale_io.py` reads it (`load_lower_atmosphere_profile`) with the loader everything else uses. The full schema is `docs/input_schema.md` section 2d; the header of the shipped example is

```
# EXHALE lower-atmosphere profile (docs/phase_e_flux_closure_design.md section 2)
# solution_id 43902c714fe3462149dfe4e5e259553ff3327e405f2a0039176bb5668ee56770
# source_code analytic
# source_version examples/17_lower_profile/make_example_profile.py
# mechanism none (synthetic H2/H/He column, no reaction network)
# stellar_flux none (no photochemistry was solved)
# p_match_bar 9.99999999999999955E-07
# p_top_bar 1.00000000000000002E-08
# p_deep_bar 1.00000000000000002E-03
# trial_flux_H 5.0000000000000000E+10
# trial_flux_He 0.0000000000000000E+00
# iteration 0
# reached_steady_state T
# notes synthetic example column; T(p) prescribed, no climate solution, no cold trap
# units p[bar] r[R_J] T[K] n_tot[cm^-3] rho[g/cm^3] Kzz[cm^2/s] q_*[-] X_*[El/H nuclei] F_*[g/s]
# columns: p r T n_tot rho Kzz q_H2 q_H X_He X_C X_N X_O F_H F_He q_H20
```

`p_match_bar` is the producer's statement of where the two models meet, and EXHALE puts its base there: T_0 , R_0 , `p_base_bar`, `q_h2_base`, `HeH` and every $X_{<El>}$ are the profile's columns interpolated at that pressure, through the same doors the scalar keys use (`set_element_abundance` for the elements, direct assignment for the rest). `n_tot` and `rho` are carried, not imposed: the base density is still n_0 of `input.inp`, and the base particle count and mass still follow from T_0 , `HeH` and `q_H2` through `comp_ntot_bc / comp_mass_per_H`. Nothing in `composition.f90` changed.

Interpolation is linear in $\log p$ — the variable both models solve on; the two codes' radius scales are equal only if their hydrostatic integrations agree — and is never extrapolated. A target that falls on a level returns that level's value bit for bit.

Columns the reader has no consumer for are kept, not dropped, and every column is found by name: a producer's element list is ordered differently from run to run.

76.2 K_{zz} stops being a constant

`he_kzz` was one scalar read at three places inside the diffusion operator: the diffusive time scale that sets the relaxation step, the binary operator's gradient-plus-eddy face coefficient, and the trace-metal kernel's $D + K$. `docs/eddy_diffusion_kzz.tex` and `LHS1140b/kzz_decision.md` had already recorded why a constant is the wrong shape — the molecular coefficient rises by nearly five decades between the base and $1.2 R_p$, and the homopause, the one consequential thing K_{zz} does, is where $K_{zz} = D$, a property of two profiles.

The operator now reads `kzz_cell(1-Ng:N+Ng)`, filled once by `eddy_diffusion_on_grid` after `define_grid`: from the profile's Kzz column interpolated onto the grid when one is given, and from the scalar `he_kzz` in every cell otherwise. The two face sites take $0.5 \times (kzz_cell(j) + kzz_cell(j+1))$, matching how `Df` and `Gf` are already face-averaged beside them.

`he_kzz` and its key `He_Kzz` keep their names: the constant is what a run states when it has no profile. With a profile in use the key is inert and a warning says so.

76.3 The scalar file is no longer a second source

With a profile in use, every `base.inp` key of the EOS-boundary, elemental-reservoir and boundary-constraint categories is refused with an error `stop` naming the key and its category; provenance comments and diagnostic keys are unchanged. This removes the same-solution problem by construction for those keys instead of checking it.

What is left to check is the pair as a whole, and it is now a refusal rather than the message it used to be: a `base.inp` beside a profile must carry a `# solution_id <hash>` comment, and it must equal the profile's. Either condition failing stops the run with both ids printed.

76.4 What the run records

`EXHALE_resolved.out` gains `lower_profile_present` and, with a profile in use, the file name, `solution_id`, source and version, `p_match_bar`, `p_top_bar`, the trial fluxes, the iteration index and whether the profile is iterable at all (a hand-written profile with no iteration index runs; it just cannot be driven by the closure loop). It is written on the direct steady route (`EXHALE_PTC=1`) as well now, which it was not before — `element_budget.py` and the closure both read it, and every LHS 1140b case uses that route.

`write_element_flux_profile` gained the hydrogen element, $F_H = 4\pi r^2 (\rho(1-X)v - J)$, on the same faces and from the same X and J the step used, so $F_H + F_{He}$ is the face mass flux identically. The file moved from `./diffusion_faceflux.txt` to `./output/element_flux_profile.txt`, with the other products, and is written unconditionally when a profile is in use rather than only under `EXHALE_DIFFUSION_CHECK=1`.

76.5 The overlap window is empty, and that is a finding

Section 3.4 of the design asks for one number per element: the elemental flux over the overlap window, as a median and a relative radial spread, measured on the escape window `[j_min:N]` intersected with the profile's coverage. That is implemented and it works — a run with `Escape radius [R_p]: 1.02` reports

```
lower_profile_flux_state  measured
lower_profile_flux_nface  59
```

— but `j_min` is the first cell with $r \geq r_{\text{esc}}$, and r_{esc} is $2R_p$ in every configuration that carries diffusion, while a profile reaching from the microbar match to 10^{-8} bar covers about $0.05R_p$ above the base. Two decades of pressure is a few scale heights, not a radius doubling. On every realistic configuration the two intervals do not intersect and the run reports

lower_profile_flux_state window_empty

which is what section 3.4 requires it to do rather than quote a base-cell number. The consequence is for E4, not for E1: the window the closure needs is between the base sound-wave region ($\sim 1.02 R_p$) and the profile top, and its lower edge has to be stated as such instead of borrowed from $j_{\min}$. The design document now carries this as an open E4 item.

76.6 Tests

id	what was run	result
T-E1	examples/17_lower_profile/ lower_atmosphere_profile.dat read by the Fortran reader (EXHALE_PARSE_DUMP=1) and by exhale_io. load_lower_atmosphere_profile, then re-emitted	15 columns carried (including the four the code has no consumer for), re-emission bitwise identical to the file, max abs round-trip error 0 over every column. The reader's match values equal the loader's: T0 = 1.4500000000000000E+003, R0 = 9.794530965504000E+009 (= 1.401 R _j exactly), HeH = 8.333333300000000E-002, q_h2_base = 1.196136002775885E-001, p_base_bar = 1.000000000000000E-006
T-E2	make check, no Lower atmosphere profile: key anywhere	==> REGRESSION PASS (all cases byte-identical) — 6/6, 12 file comparisons, goldens untouched
T-E3	mol_diffusion re-run from a scratch copy with base.inp deleted, He_Kzz removed, and a profile whose Kzz column is uniformly 1.0e9, EXHALE_MAXSTEPS=12000 OMP_NUM_THREADS=1	Hydro_ioniz.txt and Ion_species.txt byte-identical to golden/mol_diffusion.Kzz(r) on the grid: 1.000E+09 to 1.000E+09 cm ² /s; the parse dump differs from the base.inp run in exactly one variable, he_kzz (0 with the profile, 1e9 from Kzz_base), because the operator now reads kzz_cell
T-E4	four refusals and one acceptance, each from a scratch run	T_base / O_H_base / Kzz_base beside a profile stop with the key and its category named; a mismatched solution_id stops with both ids printed; a base.inp with no solution_id stops; a base.inp carrying only a matching solution_id and diagnostic comments runs
(schema)	six malformed profiles	each refused with its own message: no overlap (p_top_bar >= p_match_bar), missing solution_id, missing required column, table not strictly decreasing in p, p_match_bar outside the coverage, a short data row
(unit)	make diffusion_tests && ./diffusion_tests.x	28 passed, 0 failed — the T7c homopause tests set he_kzz and now refill kzz_cell through the same routine the code uses

make check was run on the binary built from the state of the tree at the time; the edits made after that build are confined to the profile module behind its `if (len_trim(lap_file) .eq. 0) return`, to `apply_lower_atmosphere_profile` behind `if (.not. lap_in_use) return`, and to comments, so none of them is on the path the six cases execute.

76.7 Two things fixed on the way

- The face-flux table's # columns: line carried labels with spaces in them ($F_{\text{He}}=4\pi r^2 (F_{\text{adv}}+J)[\text{g/s}]$), so it could not be split into one token per column the way every other EXHALE product's schema line can. It is now one whitespace-free token per column and the formulas moved to the comment lines above it.
- LHS1140b/make_memo_figures.py and LHS1140b/exhale/kzz_scan_table.py read D_{eff} out of that table by fixed position (usecols=(1, 8)), which the new F_H column would have shifted silently. Both now find the file under either name and the column by name, falling back to position 8 for a table written before the schema line was fixed. Stored run outputs were not regenerated.

77. Two producers for the profile handoff: the Photochem adapter and its VULCAN cross-check (2026-08-27)

Milestone E2 of docs/phase_e_flux_closure_design.md. Section 76 gave EXHALE a reader for the lower-atmosphere profile and an example file built by hand; what was missing was a producer. E2 supplies two, behind one command line:

- src/utls/photochem_to_lower_profile.py — solves a gas-giant photochemistry model to steady state on a prescribed $T(p)$, $K_{zz}(p)$ column and writes the handoff. This is the production path.
- src/utls/vulcan_to_lower_profile.py — the same schema and the same fingerprint from a finished VULCAN .vul output. This is the cross-check arm, not a second production path: VULCAN has no climate model, so it can never be the route away from a prescribed $T(p)$ (docs/vulcan_photochem_comparison.md P1.7).
- src/utls/lower_profile_schema.py — everything that is a property of the FORMAT rather than of either chemistry code: the shared command line, the elemental accounting, the hydrostatic radius, the level insertion that puts an exact node at the matching pressure, the solution_id and the two files that carry it. It exists so the two adapters cannot drift apart; a schema written twice is a schema that will be written differently twice.

The scalar generator src/utls/vulcan_to_base.py stays. It writes the single-level base.inp special case that four regression cases pin, and the two must not be read into the same run: with a profile in use EXHALE refuses every scalar physics key beside it (section 76).

77.1 One way, and the file says so

E2 is the one-way milestone. The trial elemental fluxes are recorded in the header and in the fingerprint, and they are **not** imposed as an upper boundary condition on the chemistry; imposing them, measuring the wind's own elemental flux over the overlap and iterating the pair is E4. The F_H and F_He columns therefore carry the stated trial values, and the header notes line says so in the file itself, so that no reader can mistake a boundary value for a measurement:

```
# notes one-way E2 handoff: F_H and F_He are the STATED trial fluxes, not
    measured; no climate step, T(p) prescribed.
```

77.2 What the fingerprint is over

`solution_id` is a sha256 over the lower model's configuration: the source code and version, the mechanism and thermodynamic files by their own sha256, the stellar flux file by its sha256 and the dilution applied, the prescribed $T(p)$ file, the planet mass and reference radius, the zenith angle, the top-of-atmosphere pressure, the elemental abundance vector, the matching pressure and the two trial fluxes. The closure *iteration index* is deliberately outside it: two iterations handed the same configuration and the same trial fluxes are the same solution, and the fingerprint is what says so. The same id is written into the paired `base.inp`, which carries **no physics key at all** — the provenance comment is its whole content, and any scalar key in it would stop the run by design.

77.3 Elemental accounting, and the trap that makes it a rule

Element columns are E/H **nuclei** ratios summed over every carrier, the lesson `vulcan_to_base.py` already records: photochemistry moves nuclei between molecules without creating or destroying them, so the nuclei ratio is the conserved quantity the wind inherits and the only one `melem_ab` can hold.

How the carriers are counted differs by code, and it has to. VULCAN labels are formulas once the ionization sign and the excited-state suffix are stripped (H3^+ , CH2_1), so they are parsed. Photochem labels are **not**: O1D and N2D are excited states, and a formula parser reads them as an oxygen with a deuterium attached. The Photochem arm therefore takes the composition from the model's own matrix (`dat.species_composition`), indexed by name against `dat.atoms_names` — which is also the trap P1.8 item 4 records, since that atom order changes from run to run. Condensed carriers are excluded by default: the gas EXHALE's base inherits is the gas phase, and the difference between the two sums is the cold trap, which is E3's subject.

The adapter's own conservation check (design section 4.2) is stated at the deepest level, because that is the only level where the elemental ratios are still the input ones — transport and the upper boundary move nuclei with height, which is the whole point of the closure. The design proposed a tolerance of 10^{-10} . Measured on HD 209458 b, the departure is 1.6×10^{-7} with the Zahnle H/He/N/O/C set and 2.7×10^{-5} with the VULCAN NCHO network converted through `vulcan2yaml`, i.e. chemistry-solver noise and not a miscount, so the default is 10^{-4} : loose enough not to refuse a converged solution, tight enough that a miscounted carrier (an $O(1)$ or percent-level error) is still caught. The measured worst departure is printed on every run.

77.4 The radius column

`r` is the adapter's own hydrostatic integration from a stated reference level (`--r-ref` at `--p-ref`, 1 bar by default, the transit-radius convention), RK4 in $\ln p$ with radius-dependent gravity, refined 64x per level so the answer does not depend on how the chemistry code spaced its grid. The mean molecular weight enters as `mu * m_amu`, the mass of a particle of `mu` atomic mass units.

77.5 Trap 1 was met and is handled the way P1 handles it

`gasgiants` requires $3 * \text{TOA_pressure_avg} < \text{P_climate_top}$. On the HD 209458 b column the supplied $T(p)$ stops at $2.152 \times 10^{-2} \text{ dyn/cm}^2$ while the requested TOA was 10^{-2} , so the requirement fails on the file as given. The

adapter keeps the stated TOA and cuts the climate column back to $3.05 \times \text{TOA}$, which is what `read_climate` does in `docs/p1_matched_comparison.py`; the alternative — silently lowering the TOA to fit the file — was tried first and rejected, because it changes a stated input without saying so. Trap 2 (the thermodynamic data capping the temperature) is handled by **truncating** the column from the top down, never by clipping T , because a clipped $T(p)$ is a profile nobody solved; the truncation is recorded in notes, and if it falls below `p_match_bar` the run is refused. HD 209458 b needed no truncation: its 5891 K top was accepted.

77.6 Failure policy

No profile is written when the handoff would not be one. `reached_steady_state = False` or `gave_up = True`, a temperature truncation that reaches below the match, a climate cut that leaves fewer than ten levels, `p_top_bar` $\geq$ `p_match_bar`, a matching pressure outside the table, fewer than two levels after the requested window, and the elemental conservation check are all refusals that print the reason and write nothing.

77.7 Measurements

HD 209458 b, the P1 configuration ($M_p = 0.720 M_J$, $r(1 \text{ bar}) = 1.36 R_J$, $R_* = 1.155 R_\odot$, $a = 0.0480 \text{ AU}$, Gueymard solar spectrum diluted, zenith 48 deg, TOA 10^{-2} dyn/cm^2 , VULCAN's own `atm_HD209_Kzz.txt` as the prescribed $T(p)$, $K_{zz}(p)$), matching level 1 microbar. Both Photochem arms reached steady state; the VULCAN arm is its stored HD209 output.

at <code>p_match = 1e-6 bar</code>	VULCAN / NCHO	Photochem / NCHO	Photochem / Zahnle H,He,N,O,C	PC/VUL, same network
T [K]	2332.25	2332.28	2332.43	1.0000
r [R_J]	1.495029	1.494127	1.494022	0.9994
n_{tot} [cm^{-3}]	3.1098×10^{12}	3.1095×10^{12}	3.1124×10^{12}	0.9999
ρ [g cm^{-3}]	9.4403×10^{-12}	9.6497×10^{-12}	9.7202×10^{-12}	1.0222
K_{zz} [$\text{cm}^2 \text{ s}^{-1}$]	5.000×10^{11}	5.000×10^{11}	5.000×10^{11}	1.0000
q_{H2}	0.42366	0.45363	0.46266	1.0707
q_{H}	0.44987	0.41710	0.40726	0.9271
X_{He}	9.6915×10^{-2}	9.6912×10^{-2}	9.6912×10^{-2}	1.0000
X_{C}	2.7745×10^{-4}	2.7754×10^{-4}	2.7754×10^{-4}	1.0003
X_{N}	8.1809×10^{-5}	8.1831×10^{-5}	8.1828×10^{-5}	1.0003
X_{O}	5.0329×10^{-4}	6.0608×10^{-4}	6.0608×10^{-4}	1.2042
q_{H2O}	2.2073×10^{-4}	3.4312×10^{-4}	3.5306×10^{-4}	1.5545

The matched-network hydrogen ratio is $q_{\text{H}}(\text{VULCAN})/q_{\text{H}}(\text{Photochem}) = 1.079$, which reproduces the **1.08** P1 measured for the code factor on this planet's 2331 K base (`docs/vulcan_photochem_comparison.md` P1.7) — the adapters carry the same disagreement the direct comparison found, and add none of their own. T , n_{tot} and the two element ratios the codes were given identically agree to 10^{-4} or better.

X_{O} and q_{H2O} differ by 1.20x and 1.55x for a reason that is not a code difference: the stored VULCAN HD209 run was configured with $\text{O}/\text{H} = 5.03\text{e-}4$, against the Lodders 2009 $6.06\text{e-}4$ the Photochem arms were given. That is

a difference of input, and it is what the X_{El} columns are for — it would have been invisible in a comparison of molecular mixing ratios alone.

77.8 Fingerprint round trip

test	result
the same input twice, VULCAN arm	same solution_id (4f80f048...), and the profile files are byte-identical
the same input twice, Photochem arm (two full chemistry solves)	same solution_id (f6232692...), and the profile files are byte-identical : the chemistry solve reproduces, so the fingerprint and the table agree
--trial-flux-H 5e10 instead of 0	0a4d2051..., different
--p-match 2e-6 instead of 1e-6	ceed0b58..., different
--iteration 3 instead of 0	unchanged, by design: the iteration index is bookkeeping, not configuration
the paired base.inp	carries the same id as the profile, and no physics key

77.9 EXHALE reads what the adapter wrote

A scratch case built the way examples/17_lower_profile/ is (the same input.inp, the Photochem HD 209458 b profile and its paired base.inp in place of the synthetic ones), run single-threaded with EXHALE_PARSE_DUMP=1 and then with EXHALE_MAXSTEPS=20:

```
(lower_atmosphere_profile) Reading lower_atmosphere_profile.dat ..
  base.inp: solution_id matches the profile.
  profile: T0 ->    2332.4 K          profile: q_H2(base) ->  0.46266
  profile: R0 ->    1.4940 R_J        profile: He/H ->  0.09691
  profile: p_base ->  1.00E-06 bar    profile: C/H ->  2.775E-04
  profile: source photochem, closure iteration 0
(lower_atmosphere_profile) K_zz(r) on the grid:  5.000E+11 to  5.000E+11 cm^2/s
```

Every echoed value is the profile's own number at the matching level ($T = 2332.43$ K, $r = 1.494022$ R_J, $q_{\text{H}_2} = 4.626610\text{e-}1$, $X_{\text{He}} = 9.691225\text{e-}2$, $X_{\text{C}} = 2.775366\text{e-}4$), and EXHALE_resolved.out carries the provenance: lower_profile_solution_id f6232692..., source_code photochem, p_top_bar 7.44e-9, iterable T, steady T. The paired base.inp — provenance comments only — was accepted, and the id check passed rather than being skipped.

The run itself is a 20-step relaxation snapshot and its $\dot{M}$ is not a result. One thing it does show: lower_profile_flux_state window_empty, the same finding E1 recorded — with Escape radius [R_p]: 2.00 the escape window [j_min:N] and the profile's coverage do not intersect, so the elemental flux has no window to be measured on. That is the open item E4 inherits, restated here on a real photochemical profile rather than a synthetic one.

77.10 What is not done

The LHS 1140 b application and T-E8 (element_budget.py closing on a profile run) are E2's stated gate and neither is done here; the adapters were exercised on HD 209458 b because that is where P1's measurement exists to check them against. The climate step, the tropopause and the cold-trap water are E3, and the flux closure is E4.

78. The climate step, the cold trap, and LHS 1140 b through the profile (2026-08-27)

Milestone E3 of docs/phase_e_flux_closure_design.md, and with it the two gates E2 left open: the LHS 1140 b application and T-E8.

Until now the handoff's $T(p)$ was an input. --climate makes it a solution:

```
<photochem 0.8.4 env>/bin/python \
  src/utis/photochem_to_lower_profile.py LHS1140b/lower_profile \
  --mp 0.0176220 --r-ref 0.157692 --p-match 1e-6 --kzz-const 1.0e9 \
  --stellar-flux LHS1140b/sed/lhs1140_sed_gj1132_at_b.txt \
  --flux-at-planet --wavelength-unit A \
  --atoms H,He,N,O,C --abundances He=2.09 \
  --climate --climate-p-deep 20.0 --toa 1.0e-2 --boa-pressure-factor 1.0
```

src/utis/radiative_convective_column.py is the new module: Photochem's clima (AdiabatClimate) with solve_for_T_trop on, so the stratospheric temperature is the skin temperature of the solution and not a number chosen for it, and with the background gas picked by abundance rather than by name — which is what lets one code path cover both a hydrogen atmosphere and the helium-rich one the LHS 1140 b line implies. The elemental vector is partitioned into the carriers a cool H/He atmosphere holds in equilibrium (O in H₂O, C in CH₄, N in N₂, He atomic, the rest H₂) with the validity range written at the code and a refusal above 1000 K, because that partition is the equilibrium one only below the CO/CH₄ and N₂/NH₃ transitions. Nothing downstream inherits it: the photochemistry computes its own composition from the same elemental vector.

78.1 What the codes actually do about condensation, measured

The design proposed enabling the mechanism's H₂O particle and setting rainout to the tropopause. Both halves turned out wrong for a gas giant, in opposite directions, and both were measured rather than assumed:

- the H₂O particle needs no enabling. The Zahnle set restricted to H/He/N/O/C already carries H2Oaer, and asking for water-condensation: true is **refused** — “Either “fix-water-in-troposphere” or “water-condensation” is turned on in the settings file, but the reaction mechanism already implements H2O condensation via a particle.” The cold trap is applied by the chemistry solver already; the settings switch must stay false.
- rainout is a surface process: gas-rainout: true is refused for want of a rainfall-rate, and a gas-giant model has no surface for rain to fall to. It is not how a hydrogen or helium atmosphere traps its water.

Section 5 of the design now records this correction beside the procedure.

78.2 The tropopause is solid, its pressure is not

Scanned over the one free parameter of the climate solve, its deep boundary:

deep boundary [bar]	T_{deep} [K]	P_{trop} [bar]	T_{trop} [K]	$f_{\text{H}_2\text{O}}$ above the tropopause
1	242.4	0.488	186.0	3.33×10^{-7}
3	308.5	0.765	185.8	2.06×10^{-7}
10	439.9	0.974	185.4	1.49×10^{-7}
20	555.9	1.031	185.0	1.32×10^{-7}
30, 50	refused: the pseudoadiabats leaves the thermodynamic data			

T_{trop} moves by 1 K over a factor of 20 in the deep boundary, against the skin temperature $T_{\text{eq}}/2^{1/4} = 190$ K. P_{trop} moves by a factor of two, and the trapped water moves with it, since at fixed T_{trop} the trapped mixing ratio is $p_{\text{sat}}(T_{\text{trop}})/P_{\text{trop}}$.

Against Cherubim et al. (2026). Their closed-form estimate is $f_{\text{H}_2\text{O}} \sim 7$ ppm at a 0.1 bar tropopause with $T_{\text{skin}} = 194$ K. The climate solution gives **0.13–0.33 ppm at a tropopause of 0.5–1.0 bar** — a factor 20–50 less water at a tropopause 5–10x deeper. The saturation physics agrees; what differs is where the tropopause is put, and ours is a solved level in a helium-dominated column ($\mu = 3.6$ at $\text{He}/\text{H} = 2.09$) rather than an assumed one. Reported as a result about the climate model, not tuned toward the published number.

78.3 What the trap does to the handoff

Gas-phase oxygen falls from $\text{O}/\text{H} = 6.062 \times 10^{-4}$ at the deep boundary to 4.957×10^{-7} at the matching level, a factor of 1223, while C/H , N/H and He/H cross the tropopause unchanged: CH_4 and N_2 do not condense at 185 K. The profile now always computes the elemental sums twice, over the gas phase and over every carrier including the condensates, and carries both in the header notes — at the match the two coincide, because the condensate is left behind at the tropopause. In this handoff the cold trap shows as a depletion with height, not as a condensed reservoir sitting at the base.

The 20 bar deep boundary is not a physical choice: it is the shallowest one at which the photochemical model bottom fits inside the climate column. On an atmosphere this cold the chemistry never equilibrates within any column clima can reach, so `determine_quench_levels` returns the deepest level of the grid and Photochem's default `BOA_pressure_factor = 5` then asks for a bottom five times deeper than the column it was handed. `--boa-pressure-factor` exposes it and the adapter's refusal text names it.

The prescribed-column path of E2 is unchanged where it matters: re-run on the HD 209458 b configuration of section 77 it reproduces that section's numbers exactly ($T = 2332.43$ K, $r = 1.494022$ R_J, $q_{\text{H}_2} = 4.626610\text{e-}1$, $X_{\text{He}} = 9.691225\text{e-}2$). Its `solution_id` does change, because `boa_pressure_factor` is now part of the configuration the fingerprint is over, which is the fingerprint doing its job: a profile made with a different model bottom is a different solution. Files already on disk are untouched.

78.4 LHS 1140 b through the profile: it converges, and T-E8 closes

LHS1140b/lower_profile/lower_atmosphere_profile.dat (102 levels, 16.73 bar to 1.10×10^{-8} bar, solution_id 51302c76...) drives a direct-steady JFNK run, seeded from exhale/heh2_diff_kzz1e9, to info = 0 with the element-diffusion outer loop converged (composition drift 9.1×10^{-4} at pass 8). What the reader put in agrees digit for digit with what the file says at the match, and EXHALE_resolved.out carries it:

planet_radius_RJ	1.623994484710662E-001
equilibrium_temperature_K	1.850492232658783E+002
HeH_number_ratio	2.092134697681555E+000
abundance_C	2.777529772010253E-004
abundance_O	4.956907073555432E-007
abundance_N	8.187442081270519E-005
lower_profile_solution_id	51302c7647a319bdfbd142d3b4bf0a367b25d2dc07e2675c8da3cd430dc63177

lower_profile_notes is now echoed there too, so the climate record — deep boundary, tropopause, cold-trap water, both elemental sums — travels with the run and not only with the file (lap_notes widened to 1000 characters to hold it).

T-E8. src/utils/element_budget.py closes on every element the profile carries: C 1.44×10^{-14} , N 1.47×10^{-14} , O 1.45×10^{-14} , against a 10^{-8} tolerance, and on H (0.0) and He (2.1×10^{-16}) at the base. *All element budgets close within 1.0×10^{-8} .*

The H and He rows needed a fix in the checker, not in the run. With He_diffusion on, He/H is a *solved profile* — 2.0921 at the base, 1.206 at $1.15 R_p$, 0.167 at $30 R_p$ — and the resolved HeH_number_ratio is the reservoir the base is held at, not a column invariant. The tool was testing it cell by cell and calling the separation the operator exists to produce a failure. EXHALE_resolved.out now states he_diffusion, and with it on the H and He rows are checked at the base cell, with the solved separation printed beside them. The trace elements are unaffected and stay column-wide: nothing diffuses them.

The Fortran side of E3 is two writes and one string length — lap_notes widened to 1000 characters, echoed into EXHALE_resolved.out, and he_diffusion stated there — none of which any solver reads. make check is **6/6 byte-identical**, goldens untouched.

78.5 Profile route against the scalar base, and where the difference lives

Beside exhale/heh2p13_diff_kzz1e9, the converged scalar-base case at the same eddy coefficient. Every row below is a JFNK info = 0 run finished the same way, red being the He 10830 red-pair depth from he_line_metrics.py:

case	what it adds	$\log_{10} \dot{M}$	red
heh2p13_diff_kzz1e9	scalar base, He/H = 2.13, T ₀ = 226 K	7.80	4.396
control A	HeH_base = 2.0921	7.81	4.393
control B	+ T_base = 185.05 K	7.81	4.389
control C	+ r_base = 0.162399, q_H2_base = 0.19265	7.86	4.985
control D	+ C/N/O_H_base from the profile	7.46	1.517
profile route	the profile itself	7.46	1.517

Two things fall out. Control D and the profile route agree, which is the consistency statement the single-source rule of section 2.4 is supposed to buy: handed the same base state, the two routes are the same run. (*They agreed to every digit printed here. Re-solved on the section-87 coefficient the agreement is 0.01 dex in $\log_{10} \dot{M}$ and 1 per cent in the line — inside the workflow’s own 0.3–1.5 per cent path spread, and the two rungs were solved at different residual tolerances. Agreement to every printed digit was the tolerance, not the physics.*) And the difference from the scalar LHS 1140 b case is **not** the base temperature — dropping T₀ from 226 K to the solved 185.05 K moves $\log_{10} \dot{M}$ by less than 0.01 dex — but the **elemental carbon, nitrogen and oxygen the profile brings with it**, which the scalar runs never had: −0.40 dex in $\dot{M}$ and a 3.3x shallower helium line, metal-line cooling of a 185 K base at C/H = 2.8e−4. (*Re-measured on the García Muñoz rate of section 87 (LHS1140b/exhale/basemetals_gm25/), that step is −0.37 dex and a factor 2.89 in the red-pair depth at R = 68,000 (2.88 at the R = 80,000 these were read at). The two numbers here also belong to different pairs of runs: −0.40 dex is the C → D rung, while the factor quoted as 2.9 elsewhere is ms → D, the whole ladder, which is 2.59 on the new coefficient.*) p_base is 1 microbar on both sides and is not a candidate.

78.6 What is not done

lower_profile_flux_state window_empty again: with Escape radius [R_p]: 2.00 the escape window and the profile’s coverage do not intersect, so T-E5 still has no window to be measured on. That is E4’s, together with the flux iteration itself. The composition here remains an input (He/H = 2.09, the diffusion-limit crossing of LHS1140b/kzz_decision.md section 6), not yet an output.

78.7 Beside it: a particle is not a hydrogen atom (vulcan_to_base.py)

Unrelated to E3 and found while reading the hydrostatic integrations. src/utils/vulcan_to_base.py integrated $dr/d\ln p = -kT/(\mu m_H g)$ with the hydrogen ATOM mass, while its mu_profile is a mean molecular weight in **amu**, built from VULCAN’s own species masses (H₂ = 2.016, He = 4.003, H₂O = 18.02). The mass of one particle is $\mu * m_{\text{amu}}$; using m_H makes every particle 0.8% heavier and the scale height 0.8% shorter. Measured on the HD 209458 b VULCAN output, r_base moves from 1.49388 to **1.49504 R_J** — 0.078% of the planet radius, 0.86% of the column above 1 bar, which is the 0.8% arriving where it should. Fixed; lower_profile_schema.py already used m_{amu} and recorded the discrepancy rather than copying it.

The pinned base.inp of the regression cases were deliberately **not** regenerated: they are frozen inputs, and the generator’s fix does not retroactively change a run that was made with the old file.

`run_lower.py` and `lower_column.f90` multiply the same m_H by a mean molecular weight built from integer mass numbers ($\text{mass} = 1 + 4 \text{ f_He}$), where $m_H = 1.008 \text{ amu}$ is the better of the two constants for the H_2 that dominates that mixture. They are left alone; the mixed convention is noted here so the next reader does not “fix” them into being worse.

79. The elemental-flux closure: the composition becomes an output (2026-08-27)

Milestone E4 of `docs/phase_e_flux_closure_design.md`, and with it the last three of Phase E’s tests, T-E5 through T-E7.

Until now the lower model was told what to do and never heard back. E2 and E3 carried `trial_flux_H` and `trial_flux_He` in the profile header but only as statements: nothing imposed them on the chemistry, so the handoff was one-way and He/H was still an input the wind reproduced. E4 closes the loop in both directions.

79.1 The trial flux now reaches the chemistry

`photochem_to_lower_profile.py` imposes each element’s trial flux at the photochemical model top with `set_upper_bc(species, 'flux', flux=...)`, the elemental flux converted at the model’s own top radius,

$$\phi_{\text{El}} [\text{nuclei}/\text{cm}^2/\text{s}] = \Phi_{\text{El}} [\text{g}/\text{s}] / m_{\text{El}} / (4\pi r_{\text{top}}^2) \quad (27)$$

and placed on the dominant carrier: helium on He, hydrogen on H_2 at $\phi_H/2$, two nuclei to the molecule. The dominant-carrier step is measured rather than assumed — the atomic-H share of the hydrogen nuclei at the model top is 1.0×10^{-4} in the converged LHS 1140b solution — and the adapter refuses a solution in which that share has risen past 1 per cent, because the flux would then have been put on the wrong carrier. The sign convention was read off Photochem’s own right-hand side (`photochem_evoatmosphere_rhs.f90`, `rhs(k) = ... - var%upper_flux/dz`): positive is a loss at the model top, which is this code’s outward. A zero trial flux does not call `set_upper_bc` at all, so E2’s and E3’s one-way profiles stay reproducible; that was checked by reproducing `solution_id 51302c76...` bitwise.

That the condition takes is checked every iteration instead of trusted. The adapter measures the model’s own top-of-atmosphere elemental fluxes with `gas_fluxes()` and writes them into the header as `measured_flux_H` and `measured_flux_He`. On the reference arm the imposed $1.800000 \times 10^7 \text{ g/s}$ comes back as $1.800162 \times 10^7 \text{ g/s}$, a departure of 9×10^{-5} ; the closed-top solution’s own numerical floor is 3.0 g/s in hydrogen and 18.2 g/s in helium, seven decades below.

--p-top-bar states the shallowest level to carry directly in bar, and refuses a top at or above `p_match_bar` (no overlap) or a simultaneous --toa, which would be the same level said twice in two units.

79.2 The window the closure is measured on, and why the obvious one fails

E1 left an open item: the escape window `[j_min:N]` and the profile’s coverage never intersect, so `lower_profile_flux_state` was `window_empty` on every configuration. E4 replaced the single window with two,

and the choice between them is a measurement.

The overlap window is the interval both models describe. Its upper edge is $r(p_{\text{top}})$; its lower edge is now measured — the first face at which the 5-face moving spread of $r^2\rho v$ falls below 10 per cent, which is where the standing base sound wave stops dominating — with a $1.02 R_p$ fallback that `lower_profile_flux_r_lo_source` reports when no face qualifies. Both edges, the face count and the reduction go into `EXHALE_resolved.out`.

On LHS 1140 b that window cannot carry the closure, and extending the profile is not the fix. Where it is empty it misses by $1.6 \times 10^{-4} R_p$ ($r_{\text{lo}} = 1.0099998$, $r_{\text{hi}} = 1.009840$); where it is not empty it holds 12–18 faces between 1.007 and $1.0098 R_p$ and the flux across them has a radial spread of **13 in F_{H} and 35 in F_{He}** , with a median of $1.4\text{--}1.8 \times 10^8$ g/s against the 3.16×10^7 g/s the wind actually removes. The reason is geometric: from the match to the profile top the column covers 1.96 decades of pressure in $0.0098 R_p$, about $0.005 R_p$ per decade, while the wind is still 6x above its far-field flux at $1.05 R_p$ and does not settle within 10 per cent until $\sim 1.5 R_p$. Reaching $1.5 R_p$ on that scale would ask the photochemical column for of order 100 further decades of pressure. Mapping p_{top} through EXHALE’s own $r(p)$ rather than the profile’s — the two hydrostatic scales differ, the wind being hot where the cold column is not — moves the upper edge to $1.0734 R_p$ and makes the window non-empty, but its spread is then 5.6 to 7.3. Both mappings say the same thing: the interval the two models share lies wholly inside the region where the wind’s mass flux is not yet its own.

The steady-flux window, $r \geq r_{\text{esc}}$, is what the closure uses. At a steady state the elemental flux $4\pi r^2(\rho X v + J)$ does not depend on radius, so a flux measured there *is* the flux through the matching level; that identity, and not convenience, is what makes it a statement about the handoff. On the reference arm it holds 191 faces and gives $F_{\text{H}} = 1.8143 \times 10^7$ g/s, $F_{\text{He}} = 1.3204 \times 10^7$ g/s with radial spreads of 0.44 and 0.46 per cent, flat to the same tolerance as the mass flux itself (0.45 per cent). That is T-E5, met in the window where it can be met. The driver reads the overlap window first every iteration and substitutes the steady one only with the reason written into `closure.log` and the window named in the history table; it never substitutes silently.

The spread depends on how far the wind was converged, and that was measured rather than assumed: at `Resid tol: 1.0e-3`, the value `finish_case.sh` uses, the same solution’s steady-window spread is **7.4 per cent** (F_{H} 7.1, F_{He} 7.9) and would fail the closure’s own validity test; continuing the same solution to `1.0e-4` brings it to **0.50 per cent** and moves $\log_{10} \dot{M}$ from 7.47 to 7.50. The 7 per cent was under-convergence, not a floor, and the closure runs the wind at `1.0e-4`.

79.3 The driver

`src/utils/element_flux_closure.py` owns the loop and nothing else. Damped Picard on both elements with $\omega = 0.5$, halved to a floor of 0.125 on any iteration whose residual fails to fall, the convergence test reading the **undamped** residual $|F_{\text{measured}} - \Phi_{\text{trial}}|/|F_{\text{measured}}|$, `tol = 0.05` and $k_{\text{max}} = 8$. Both hydrogen and helium are iterated: holding helium at zero, as the design first proposed, would leave the H:He partition — the whole question on this planet — as a diagnostic rather than a condition.

Section 6.2's validity rule is a statement about the *window*, not the *iterate*: a window whose spread exceeds `tol` cannot decide a residual at `tol` whatever that residual is, so `spread > tol` stops the run as UNRESOLVED with both spreads printed, every iteration and independently of how far the *iterate* is.

Everything that belongs to the planet rather than the iteration lives in one JSON configuration (`--print-config-template`), so the driver carries no planet constants; `LHS1140b/exhale/flux_closure/lhs1140b_closure.json` is the LHS 1140 b one. `closure_history.txt` carries one row per iteration with both trial and measured fluxes, both residuals, ω , the window used and its spreads, He/H at the match, $\log_{10} \dot{M}$, the EXHALE info, and the profile's `solution_id` — the fingerprint chain, so the sequence is auditable after the fact. `--resume` continues from the last completed iteration and `--dry-run` exercises the reader against an existing run directory.

Two departures from `finish_case.sh`, both measured. `Resid tol: 1.0e-4`, for the reason above. And the post-processing pass runs **without** the EXHALE_PTC* variables: `finish_case.sh` sets them inline on the JFNK command alone, and carried into the second pass they make the binary re-enter the steady solver and stop before writing the advection-corrected profiles, so `*_adv.txt` and the mass-loss rate never appear. The first LHS 1140 b closure run reported `log10 Mdot = nan` and that is what it was.

79.4 LHS 1140 b, closed

Three arms, from $\Phi_{\text{ref}} = (1.8 \times 10^7, 1.3 \times 10^7)$ g/s in (H, He) and from $0.3 \times$ and $3 \times$ that, all seeded from the same converged wind, all with He/H = 2.09 as the *starting* reservoir.

start	k	F_{H} [g/s]	F_{He} [g/s]	He/H at the match	$\log_{10} \dot{M}$	He 10830 red depth [%]	EW [mÅ]
1.0 x	0	1.81433×10^7	1.32040×10^7	2.0923516	7.500	1.699	4.743
0.3 x	5	1.82031×10^7	1.34266×10^7	2.0923486	7.500	1.721	4.804
3.0 x	6	1.82119×10^7	1.34710×10^7	2.0923602	7.500	1.725	4.818

The three converged He/H agree to 5.5×10^{-6} , $\log_{10} \dot{M}$ to better than the 0.005 the log prints, F_{H} to 3.8×10^{-3} and F_{He} to 2.0×10^{-2} — all far inside the 5 per cent criterion, so **T-E7 passes and the closure is single-valued on this planet**. ω never halved in any arm: the residual fell on every iteration, so the map is a plain contraction here and the damping rule was not exercised. The iteration counts are exactly what a contraction at $\omega = 0.5$ predicts, the residual halving each step, which also says the 3x arm used six of its eight allowed iterations with no margin: a start further than an order of magnitude out would need a larger k_{max} or a secant acceleration.

The result to read, though, is what the closure returns. He/H at the match comes back at 2.09235 against the 2.09 the reservoir was started from — a change of 1.1×10^{-3} , and only 2×10^{-4} of that from the escape flux itself. The composition is now an output, and on this planet the output is the *well-mixed* value: at $K_{\text{zz}} = 10^9$ cm²/s eddy mixing homogenizes the column all the way to 1 microbar, so a 1.8×10^7 g/s hydrogen escape imposed at the model top moves the matching-level He/H by 2×10^{-4} while moving it by 1.8×10^{-2} at the model top itself. It is the eddy coefficient, not the escape flux, that sets the composition at the match. That is a result about this configuration and not a general one, and it makes LHS1140b/kzz_decision.md section 0's adopted K_{zz} the quantity the answer now rests on.

Above the match the wind fractionates strongly, as it did before: the elemental fluxes carry $F_{\text{He}}/F_{\text{H}} = 0.73$ by mass against the base reservoir's 8.31, and He/H falls from 2.09 at the base to 0.183 at $30 R_p$. The closure does not weaken that; it says the separation happens in the wind and not below the match.

The flux budget. $F_{\text{H}} + F_{\text{He}}$ reproduces the mass flux identically — exactly, in the two arms whose medians fall on the same face — because the binary operator's X is the helium mass fraction of the H+He mixture and the trace metals ride inside ρ . So the measured F_{H} is hydrogen *plus* metals, overstating elemental hydrogen by the metal mass share, which is 4.79×10^{-4} here (mass_per_H = 9.3739 amu, C/N/O at 2.778×10^{-4} , 8.188×10^{-5} , 4.955×10^{-7}). That is two decades below tol and is stated rather than corrected.

79.5 T-E8, and one thing it caught

The element budget closes for helium at the base cell, where the reservoir is the boundary condition, to 1.5×10^{-15} . Hydrogen, carbon, nitrogen and oxygen miss by 5.6×10^{-8} , 1.2×10^{-4} , 1.2×10^{-4} and 3.3×10^{-4} against the checker's 10^{-8} , and the miss is not noise: **it equals the change in the elemental reservoir between the seeding solution and the run.** load_IC restores every metal density the restart file carries and rebuilds an element from melem_ab only when the file does not carry it at all (load_IC.f90:283-334), so an iteration that changes a reservoir keeps the seed's value everywhere above the base and EXHALE_resolved.out then reports a reservoir the run did not use above the base cell. It was seen at full size first: an iteration seeded across a 22 per cent change in O/H reported exactly 17.8 per cent, which is that change expressed on the new reservoir.

This was **not** fixed at the time, because the fix changes restart semantics, which is a user-visible behavior and expensive to reverse; it was recorded for a decision. Its size on the result reported here is small and measured: the closure moves reservoirs by at most 5×10^{-4} between iterations, and the coolant that matters here is C I, which carries 94 per cent of the radiated energy and is stale by 1.2×10^{-4} ; oxygen, the worst row at 3.3×10^{-4} , carries under 10^{-4} of the cooling on this planet. The decision was taken and the fix implemented on 2026-08-27: section 80.. The numbers quoted in this section are what the code gave **before** it, and are left as they were measured.

79.6 T-E9: not verified, and why

Not verified: VULCAN did not converge on this cold He-rich column (longdy rising $2.65 \rightarrow 2.75$ against yconv_cri = 0.01, dt pinned at 1.85×10^3 s, the model clock at 1.42×10^7 s, limiting cell nz = 72 and H2O — the cold trap — exiting on “Maximal allowed steps exceeded (10000)”; stopped 2026-08-27.

The cross-check arm was built on the production profile's own $T(p)$ and K_{zz} and never produced a profile, so there is no VULCAN composition at the match, no wind run on one, and no cross-code difference to state beside the P1 spread. LHS1140b/exhale/flux_closure/vulcan_arm/ holds only an input.inp and the seeded initial condition; the binary was not launched. The full record, including the second run that reached step ~500 of 30000 before being stopped and therefore says nothing either way, is LHS1140b/lower_profile_vulcan/vulcan_work/STOPPED.txt.

Two things found on the way that outlive the attempt. First, a real defect in src/utls/vulcan_to_lower_profile.py, now fixed: sch.formula_elements strips the phase suffix, so H2O_l_s parsed as a second H2O and condensed carriers

were counted as gas, which made the cold trap invisible in `X_0` whatever `--count-condensates` was set to. The gas phase is now selected by the `.vul`'s own `atm['gas_indx']`, condensate is excluded from μ always as VULCAN itself does, and the flag now selects. The Photochem adapter and `lower_profile_schema.py` are untouched, so the production arm and its `solution_id 5130...` are unaffected. Second, a configuration trap worth stating because it silently produces a wrong atmosphere rather than an error: `vulcan_cfg.He_H` is NOT read on the FastChem `ini_mix = 'EQ'` path (`build_atm.ini_fc` builds the element file from `atom_list` minus H), so helium must be in `atom_list` or the run keeps the solar He 10.9864 and integrates at $\text{He}/\text{H} = 0.097$. The first run did exactly that.

Whether VULCAN can be made to converge here is open. It is not a blocker for Phase E: Photochem is the production path precisely because it carries the climate model this planet needs, and T-E5 through T-E7 are closed without it.

80. A restart onto a reservoir the handoff has moved (2026-08-27)

Section 79. left T-E8 closing for helium and missing for hydrogen, carbon, nitrogen and oxygen by 5.6×10^{-8} , 1.2×10^{-4} , 1.2×10^{-4} and 3.3×10^{-4} , and named the cause: `load_IC` restores every metal density the restart file carries, so a flux-closure iteration that moves an elemental reservoir keeps the seed's value everywhere above the base cell while the boundary condition already uses the new one. The fix was deferred there because it changes restart semantics, which is user-visible. The decision was taken on 2026-08-27; this section is what was implemented and what it measures.

80.1 The rule

An element is renormalized as it is loaded if, and only if, the lower-atmosphere handoff states its reservoir — a `<El>_H_base` key of `base.inp`, or an elemental ratio of the file named by `Lower atmosphere profile:`. For such an element the whole loaded column is multiplied by one factor,

$$r_{\text{El}} = \frac{(\text{El}/\text{H})_{\text{stated}}}{(\text{El}/\text{H})_{\text{loaded state, base cell}}}, \quad (28)$$

the same factor for every ionization stage. The normalization moves and nothing else does: the ionization split of each cell and the radial shape of the loaded profile are preserved exactly, because a factor common to all stages and all cells cancels out of every ratio the state is made of. El/H at the base cell is counted in nuclei — the element summed over its stages against the hydrogen nuclei carried by H I , H II , H_2 , H_2^+ , H_3^+ and HeH^+ — which is the convention of `element_ratio_HeH` and of `src/utls/element_budget.py`, so the factor is measured in the same units the budget check reads it back in.

Helium is not part of this rule and keeps its own, older convention: the base cells are set to the input He/H , and with `He_diffusion` on the column above keeps the diffused split, because that split is the state being restarted rather than a defect of the file.

What decides “the handoff states it” is not a test on the numbers. Both readers set an abundance through one routine, `set_element_abundance` in `input_read.f90`, and passing through that door is what marks the element in the new `melem_from_handoff(n_melem)` of `global_parameters`. An abundance that reached the run from `metals.inp` alone never passes through it, so it is not marked and its column is loaded untouched. There is no tolerance and no `abs(r - 1) > tiny` guard: the branch is entered or it is not, and when it is not entered the arithmetic is never performed.

Each rescaled element prints its own line, with the loaded ratio, the stated ratio and the factor, so the factor can be compared against the reservoir change it is supposed to equal:

```
(load_IC) C/H at the base cell: restart file 2.777530E-04, handoff 2.777859E-04; column rescaled by 1.000118E+00
(load_IC) O/H at the base cell: restart file 4.956907E-07, handoff 4.955292E-07; column rescaled by 9.996741E-01
(load_IC) N/H at the base cell: restart file 8.187442E-05, handoff 8.188411E-05; column rescaled by 1.000118E+00
```

80.2 What it closes

The test is the last iteration of the LHS 1140 b closure of section 79., LHS1140b/exhale/flux_closure/hi/k06, re-run from the same seed (k05/output/) on the same profile, once with the previous binary and once with this one. The three factors above are the run’s own log. They are the measurement, not an illustration: 1.000118 and 0.9996741 are the same numbers as the 1.183×10^{-4} and 3.259×10^{-4} the budget was missing by, which is what the section 79. diagnosis said they would be.

element	budget before	budget after
H (against rho/mass_per_H)	5.633×10^{-8}	1.543×10^{-15}
He	1.910×10^{-15}	1.698×10^{-15}
C	1.183×10^{-4}	3.688×10^{-14}
N	1.183×10^{-4}	3.724×10^{-14}
O	3.260×10^{-4}	4.017×10^{-14}

element_budget.py goes from “4 element(s) miss their reservoir by more than 1.0e-08” to “all element budgets close within 1.0e-08”. Hydrogen closes with the rest and for the same reason: it is checked against rho/mass_per_H, and the metal mass a stale reservoir leaves in ρ is exactly what put it 5.6×10^{-8} out.

The effect on the physical result is small, as section 79. predicted from the size of the reservoir change. The steady mass flux at the top of the domain moves by -2.4×10^{-5} in relative terms ($\log_{10} \dot{M} = 7.50$ either way), and the He 10830 line, run through EXHALE_transit.py on both solutions, moves by 4.7×10^{-5} in relative terms in the red-component depth ($1.725492 \rightarrow 1.725411$) and by at most 1.2×10^{-6} in the normalized line profile.

80.3 What is unchanged, checked rather than argued

Two restarts that must not move, both re-run with the two binaries and compared bitwise on output/Hydro_ioniz.txt and output/Ion_species.txt:

- LHS1140b/exhale/heh0p55 reloaded and advanced one step — a restart with no handoff of any kind. Identical.
- backup/regression/wasp_full seeded from its own converged output with Load_IC? True and advanced one step — a restart that carries a full metal state, from metals.inp and not from a handoff, so it exercises the metals block of load_IC with every element unmarked. Identical.

The binary the comparison was made against is kept as backup/EXHALE_pre_handoff_restart_rescale.x.

make check is unaffected by construction and by measurement: all seven default cases are cold starts (Load_IC? False), so none of them reaches this code at all, and the full matrix — wasp_full, wasp_he23off, mol_base_handoff,

mol_metals, mol_lyman_werner, mol_diffusion, lower_profile — comes back byte-identical (==> REGRESSION PASS (all cases byte-identical), 2026-08-27).

80.4 $r_{\text{El}} = 1$ is not an identity, and the reason is worth stating

The gate is the flag, not the value of the factor, and that is deliberate: $r_{\text{El}} = 1$ does **not** leave the state bitwise where it was. Restarting backup/regression/wasp_full from its own converged output twice with the same binary — once with the abundances coming from metals.inp alone, once with a base.inp restating those same seven numbers as <El>_H_base — prints 1.000000E+00 for every element and still gives a different state after one step. The factors are 1 only to about 10^{-14} : the seed's base-cell EL/H is what the ionization solver converged to, and it reproduces the reservoir to its own tolerance rather than to the last bit. After one step the departure is 1.8×10^{-16} in the median and 1.5×10^{-6} at its worst, and the worst cell is O III at $n = 11 \text{ cm}^{-3}$ near the base, six orders below that ion's own peak in the same column — the solver's tolerance seen from the other side, not a change in the solution.

So the byte-identity claim is exactly the one the flag supports and no wider: **an element the handoff does not state, and a restart with no handoff at all, never enter the arithmetic**, and those are what were compared bitwise against the pre-change binary above. An element the handoff does state is renormalized whenever it is loaded, which is the user-visible change of restart semantics this section is about.

81. Is the wind collisional where it is set? A Knudsen diagnostic, and LHS 1140 b's answer (2026-08-27)

Until now the repository had no way to ask whether a converged solution is a solution of the physical problem. A hydrodynamic wind is a continuum solution, valid only where the gas is collisional on the scale over which the flow varies — above all through the critical point, which is where the topology, and with it the mass flux, is fixed. Nothing measured that. src/utils/collisional_validity.py (new, Python post-processing; it reads an existing run directory and changes nothing, and no Fortran source was touched) computes it. Definitions, provenance and the full result tables: docs/collisional_validity.md.

81.1 The collision model is not a new one

The four momentum-transfer limits are the ones the binary element-diffusion operator already uses, in the same Chapman–Enskog first approximation: rigid-sphere neutral–neutral (Banks & Kockarts 1973), non-resonant ion–neutral with the induced-dipole and rigid-core frictions added, and screened Coulomb with the Spitzer $\ln \Lambda$ for ion–ion and ion–electron. Python cannot call the Fortran, so the four routines are transcribed from src/modules/functions/binary_element_diffusion.f90 with the source line cited at each definition and every constant taken from that file rather than re-chosen; the two are meant to be read side by side. The step from a diffusion coefficient to a collision frequency is the definition of the binary coefficient itself,

$$\nu_{st} = \frac{n_t kT}{\mu_{st} \hat{D}_{st}}, \quad \hat{D} = nD, \quad (29)$$

in which the total carrier density cancels.

Two channels are deliberately missing, and both make the verdict *more* conservative by lengthening the mean free path: resonant charge exchange ($H^+ + H$, $He^+ + He$), which the element-diffusion operator excludes on the grounds that both partners carry the same element, and electron–neutral momentum transfer, which no element-transport coefficient needs. The electron Knudsen number is therefore quoted only with that caveat, and the bulk is built from the heavy particles alone.

81.2 Definitions, all stated

$\lambda_s = \bar{v}_s / \nu_s$ with the harmonic sum over partners; the structure scale $L = \min(H_p, L_v, r)$ with $H_p = |d \ln p / dr|^{-1}$ measured from the solution and $L_v = |v / (dv/dr)|$ admitted only above Mach 0.01 (below that it is the distance to a stagnation point, and the base sound-wave layer would otherwise drive Kn to infinity); $Kn_s = \lambda_s / L$. The **exobase** is $Kn_{\text{bulk}} = 1$; the **critical point** is the sonic point with the code’s own sound speed ($c_s = \sqrt{\gamma p / \rho}$, $\gamma = 1.666666666667$, `eval_dt.f90:29, parameters.f90:385`), not a re-definition; the **critical region** runs from peak volumetric heating to the sonic point, and “collisional” means $\max Kn < 0.1$ across it. Coupling times $1/\nu_s$ and the electron–ion energy-coupling time are reported against the flow and heating times, with the note that they are a *different* test: a strongly subsonic flow gives hundreds of collisions per flow time even at $Kn \sim 1$, and it is the Knudsen number that decides continuum validity, because the terms the closure drops scale with the gradient over a mean free path.

81.3 The result, measured 2026-08-27 on existing outputs

Three representative LHS 1140 b solutions — LHS1140b/exhale/heh0p55 (diffusion off), heh2p13_diff_kzz1e9 (diffusion at the adopted K_{zz}), and the converged flux closure `flux_closure/hi/k06` — against the HD 209458 b control backup/phase_d_baseline/new_kzz1e9_d3b:

case	critical point	exobase	$Kn = 0.1$ at	max Kn , crit. region	verdict
heh0p55	none in domain (max Mach 0.540)	25.22	6.42	2.2	unvalidated
heh2p13_diff_kzz1e9	none in domain (max Mach 0.565)	26.66	6.79	1.9	unvalidated
<code>flux_closure/hi/k06</code>	none in domain (max Mach 0.550)	19.76	5.53	3.0	unvalidated
HD 209458 b new_kzz1e9_d3b	4.085	above 4.15	—	0.021	validated

Radii in R_p ; the LHS 1140 b domain ends at $30 R_p$.

LHS 1140 b’s wind never reaches its critical point inside the domain. It is subsonic to $30 R_p$ in all three solutions, so the mass flux is set at the outer boundary — where Kn_{bulk} is already 1.3–2.3. Against the p-winds retrieval, which places its isothermal Parker sonic point at 8–9.5 R_p ; that radius lies *inside* the exobase computed here (19.8–26.7 R_p), but Kn_{bulk} there is already 0.14–0.30. So the answer is neither “the sonic point is outside the exobase” nor “the critical region is collisional” — it is **transitional**. Neutral hydrogen sets the number; the ions and electrons are Coulomb-held two to five orders of magnitude more tightly. HD 209458 b is the control at a factor of ~ 100 lower Kn , and it validates cleanly.

81.4 The kinetic estimate is a scale, not a bound

Jeans escape through the computed exobase (Chamberlain & Hunten 1987 eq. 7.2.5) gives $1.0\text{--}2.0 \times 10^7$ g/s, a factor about three below the continuum $\dot{M}$ of $10^{7.47}\text{--}10^{7.81}$. But $\lambda_J(\text{H}) = 0.78\text{--}1.45$ at that exobase: it is barely gravitationally bound, the Jeans integral is most of the Maxwellian rather than its tail, and the atmosphere is in blow-off, so the number does not bound anything. What it does say is that the continuum result is within a factor of a few of what a collisionless outer atmosphere at the same density and temperature would lose — the continuum failure is not hiding an order of magnitude. On HD 209458 b, $\lambda_J = 10.8$ and the hydrodynamic rate exceeds Jeans by 10^4 , which is the signature of a driven wind rather than an evaporating one.

The numbers above are read from the advection-corrected profiles, the ones the transit tool consumes. Read from the equilibrium pair instead (--eq) the verdicts are unchanged — still no critical point in the LHS 1140 b domain, still $Kn \sim 1$ where the outer boundary sets the flux, still collisional through HD 209458 b's critical point — but the exobase moves above $30 R_p$, because the equilibrium far field is more ionized and Coulomb collisions there are far stronger than the neutral ones. An exobase radius quoted from this tool has to say which pair it came from.

This closes the plan's baseline row 11 the way that row demanded: an uncollisional critical region makes the hydrodynamic $\dot{M}$ **unvalidated, not overestimated**, and the tool's `validity_statement` says exactly that, with the numbers it rests on. What remains unsettled is the LHS 1140 b rate itself, which needs a kinetic or transitional-flow calculation this repository does not contain. The He 10830 line-forming layer is not implicated: it sits at $1.1\text{--}3 R_p$, where $Kn \leq 0.021$ in every case.

82. The third body was a mass density: a molecular-network audit in the He-dominated limit (2026-08-27)

Phase F of the LHS 1140 b plan asks for an audit of the molecular network where it is least tested — a helium-dominated envelope. R16–R20 and R23, the shared H–He charge-exchange path, the two He 2^3S Penning channels, and the electron and H-nucleus closures were checked reaction by reaction against the publisher PDFs, and the closures were measured on runs spanning $\text{He}/\text{H} = 0.079$ to 1000. The record, with the full reaction table and the numbers quoted below, is `docs/molecular_chemistry_audit_he_rich.md`.

82.1 No rate coefficient was wrong

All 23 Koskinen et al. (2022) Table-1 entries are transcribed verbatim and correctly, R16–R20 and R23 included; so are the Huang et al. (2023) Table-4 B1/B2 charge-exchange pair, the Taylor et al. (2025) $\text{He}(2^3\text{S})+\text{H}$ fit and the García Muñoz (2025) $\text{He}(2^3\text{S})+\text{H}_2$ fit, which reproduces its four tabulated points to 0.13 % (re-measured). Every one of the eight balance rows of `mol_heh_rows` carries each reaction in the rows it belongs to, with the right weight and sign.

82.2 What was wrong was the density handed to the three-body reactions

R12 ($\text{H}_2 + \text{M}$), R13 ($\text{H}^+ + \text{H}_2 + \text{M}$) and R15 ($\text{H} + \text{H} + \text{M}$) are written as two-body coefficients times the third-body density, and the interface that carries it, `ion_cell_state%ntot`, says “total particle density”. What `ionization_equilibrium` actually passed was `n_in_dim =  $\rho/m_{\text{H}}$` . That is a *mass* density in units of m_{H} — `calc_rho` weights every species by `bsp_mass`, so an H_2 molecule contributes 2 and a helium atom 4 — and it over-counts the third bodies by the mean particle mass. At the hot-Uranus molecular base the factor is 2.3; in a helium-dominated base it approaches 4. The error therefore grew along exactly the axis this audit was asked to examine. The same ρ/m_{H} was also used as the gas pressure of the chemical-equilibrium seed that restarts a failed molecular cell, where the EOS’s own law is $p = (n_{\text{tot}} + n_e)k_{\text{B}}T$. Both now use `calc_ntot` (one particle per species, electrons excluded) and, for the pressure, that sum plus n_e .

The consequence is visible where it should be. On the `mol_diffusion` configuration re-run at $\text{He}/\text{H} = 1, 10$ and 1000 (only the He/H number ratio and the `HeH_base` of the handoff changed, 12000 marching steps, `OMP_NUM_THREADS=1`), the $\text{He}/\text{H} = 10$ arm no longer aborts on the marching loop’s NaN detector and completes all 12000 steps, and the number of cells whose molecular root leaves the physical simplex and has to be clamped onto the element budget falls from 37 to 4 at $\text{He}/\text{H} = 1000$ and from 181 to 169 at $\text{He}/\text{H} = 10$. The $\text{He}/\text{H} = 1$ arm still aborts, later (5538 steps against 3402); it is not diagnosed here.

82.3 A guard that was not guarding

The molecular systems close neutral helium as $n_{\text{He}}(1-x_2-x_3)-n_{\text{HeH}^+}$, but the admissibility test `ionization_fractions_physical`, the distance measure `element_budget_violation` and the clamp `clamp_fractions_to_element_budget` all summed the helium budget as $x_2 + x_3$ (plus the metastable) alone. A root with $x_2 + x_3 = 1$ and any HeH^+ was accepted although its neutral helium is negative. All three now carry the HeH^+ nucleus, converted with this cell’s $n_{\text{H}}/n_{\text{He}}$. Nothing in the measured runs reaches that state — HeH^+/He stays below 10^{-10} at every He/H tested — so the change moves no result; it closes a hole rather than fixing a symptom.

82.4 The closures hold

With molecules on, charge neutrality $n_e = n_{\text{H}^+} + n_{\text{H}_2^+} + n_{\text{H}_3^+} + n_{\text{HeH}^+} + n_{\text{He}^+} + 2n_{\text{He}^{++}}$ closes against the n_e the code writes to 4×10^{-16} at every He/H from 0.079 to 1000, the H-nucleus/mass closure to 4×10^{-16} , and `src/utls/element_budget.py` closes both elements to 6×10^{-14} at $\text{He}/\text{H} = 1000$. Element conservation is not what the helium-dominated limit breaks. What it does break is the solve: the molecular `hybrd1` system balances helium in rows of order n_{He}^2 and the molecules in rows of order n_{H}^2 , so at $\text{He}/\text{H} = 1000$ the two blocks of one residual vector differ by 10^6 with no row scaling applied, and isolated cells fall into the atomic basin next to neighbours at $f_{\text{H}_2} = 0.88$. The correlation with He/H is measured; the cause is not established.

82.5 Four caveats are recorded rather than patched

Each would replace a published rate with one from a different compilation. (i) The identity of M: Koskinen et al. never say what their “ n ” counts, and a monatomic third body is generally less efficient than H_2 , so R12/R13/R15 with

helium as M are upper bounds. (ii) HeH^+ formation: Table 1 has only $\text{He}^+ + \text{H}_2$ (R20), and not $\text{H}_2^+ + \text{He} \rightarrow \text{HeH}^+ + \text{H}$, which needs no He^+ and is therefore the route that grows with the helium fraction (García Muñoz 2025 Table A.6 gives 1.0×10^{-11} at 2000 K rising to 1.5×10^{-10} at 10^4 K); the $\sim 10\%$ associative branches of both Penning channels end in HeH^+ too and are not resolved. Against the same tables R16 is $3.4\text{--}8.6\times$ low, R18 $1.2\times$ high, R19 $1.4\text{--}2.6\times$ low, R20 $14\times$ high and R17 up to 1.9×10^4 high. (iii) The $\text{H} \leftrightarrow \text{He}$ charge-exchange pair does not satisfy detailed balance: $k(\text{H}^+ + \text{He})/k(\text{He}^+ + \text{H})$ should be $4\exp(-127500/T)$ and the published fits give 10^2 times that at 10^4 K, a discrepancy weighted by $n_{\text{He}}/n_{\text{H}}$ in the He^+ balance. (iv) The Taylor two-branch Penning fit steps by $1.57\times$ at its own 4000 K break point, which is the published fit and not a transcription error. Each is noted at the code site as well.

82.6 Does this move the LHS 1140 b results?

No: **no LHS 1140 b run in the repository has molecular chemistry on**. None of the 170 `input.inp` files under LHS1140b/ sets `Molecular chemistry`, and none of the 193 `Ion_species.txt` files carries an $\text{H}_2/\text{H}_2^+/\text{H}_3^+/\text{HeH}^+$ column. Every statement corrected above lives inside a `thereis_mol` branch, so the $\text{He}/\text{H} = 2.09$ solution, the closure ladder and the `heh2p13_diff_kzz1e9` arm are untouched and no re-convergence is needed. What can be measured on those stored solutions is the size of the factor that *would* have been wrong had they been molecular — the ratio of the mis-supplied ρ/m_{H} to the EOS particle count $p/(k_{\text{B}}T) = n_{\text{tot}} + n_e$ at the base cell:

stored solution	He/H	$(\rho/m_{\text{H}})/(n_{\text{tot}} + n_e)$ at the base
LHS1140b/exhale/flux_closure/hi/k06	2.09	3.03
LHS1140b/exhale/heh2p13_diff_kzz1e9	2.13	3.04
LHS1140b/exhale/flux_closure/heh9p7/k03	9.7	3.72
backup/regression/golden/mol_diffusion	0.0793	2.27

82.7 Goldens

The matrix splits exactly along the `thereis_mol` branch: `wasp_full`, `wasp_he23off` and `lower_profile` PASS byte-identical, and the four molecular cases move. Their shift on the 12000-step snapshots, largest relative difference over the column against the previous golden:

case	ρ	p	T	heat	cool	H_2 front ($f = 0.5$)	$\log_{10} \dot{M}$
mol_base_handoff	1.4 %	2.6 %	1.2 %	1.6 %	3.0 %	1.1597 $\rightarrow$ 1.1597	10.58 $\rightarrow$ 10.58
mol_metals	3.0 %	5.0 %	2.2 %	3.5 %	6.7 %	1.1617 $\rightarrow$ 1.1617	10.58 $\rightarrow$ 10.58
mol_lyman_werner	0.34 %	0.53 %	0.21 %	0.44 %	0.71 %	1.1304 $\rightarrow$ 1.1304	10.58 $\rightarrow$ 10.58
mol_diffusion	1.4 %	2.6 %	1.2 %	1.6 %	3.0 %	1.1597 $\rightarrow$ 1.1597	10.58 $\rightarrow$ 10.58

The velocity column moves by more ($2.5\text{--}8\times$) but only in the base sound-wave layer of these unconverged relaxation snapshots, where v is small and oscillating. The H_2 front does not move at all, and the base H_2 fraction shifts in the fourth decimal: these are H-rich hot-Uranus bases already at $f_{\text{H}_2} \simeq 1$, and a fraction pinned at unity cannot respond to a $2.3\times$ change in the third-body density — it responds in ρ , p and T instead. That is the same statement the He-rich scan makes from the other side: the term becomes decisive only once helium dilutes the H_2 . Goldens re-snapshotted for the four moved cases at the end of the series, and make `check` re-run after.

83. The outer region a converged wind inherits: path dependence beyond the exobase (2026-08-27)

The LHS 1140 b closure ladder (LHS1140b/kzz_decision.md Sect. 8) climbs a reservoir axis by seeding each arm from the converged wind of the arm below. One rung was reached by a large jump instead of a step, and it converged to a different solution — different only where the continuum equations no longer hold. This section records what separates the two, why neither convergence test noticed, and what it means for reading any of these winds.

83.1 The two solutions

Both are flux-closure arms at the same reservoir $\text{He}/\text{H} = 12.0$ over the same photochemical lower atmosphere. One, LHS1140b/exhale/flux_closure/heh12/, was seeded from the 8.0 arm — a $1.50\times$ jump in reservoir — and closed at $k = 4$. The other, .../heh12_ctl/, was seeded from the 11.1 arm, a $1.08\times$ step, and closed at $k = 2$.

	seeded across a $1.50\times$ jump	seeded one $1.08\times$ step
He/H at the $1\ \mu\text{bar}$ match	12.013522	12.013520
cold-trap O/H of the column	1.0656×10^{-6}	1.0660×10^{-6}
exhale_info	0	0
achieved wind $\ R\ $	2.529×10^{-4}	3.864×10^{-4}
steady-window flux spread	1.74 %	3.30 %
$\log_{10} \dot{M}$ [g/s]	7.420	7.410
He 10830 red depth [%]	4.630	3.856
He 10830 red-pair EW [% Å]	1.3126	1.1181

Identical base, identical matching-level composition to seven digits, both converged and both inside every tolerance the driver applies — and a 17 per cent difference in the observable the whole study is fitted to. Neither residual separates them: on the coefficient in force when this was written the outlier reported the *smaller* of the two, and re-solved on the García Muñoz rate of section 87 the order reverses (jump-seeded 2.533×10^{-4} against the control's 1.249×10^{-4} , both info = 0, both inside the 4.0×10^{-4} target). Which one reports less is not a property of the pair; that $\|R\|$ does not rank them is (LHS1140b/exhale/xuvkzz_gm25/results.txt).

83.2 Where they differ

Nowhere the line window or the closure window looks. Two measures that do not depend on any sampling choice locate it. r_{drop} is the first radius outside $1.5R_p$ at which T falls below half of $T(12R_p)$; the other is the share of the He 2^3S radial column that lies outside $10R_p$. Both are monotonic along the ladder, and both single out one point.

reservoir He/H	$r_{\text{drop}} [R_p]$	$T(12R_p)$ [K]	He 2^3S column outside $10R_p$
2.09	30.00	893	0.0119
3.0	27.63	1035	0.0116
5.0	23.76	1310	0.0105
8.0	20.80	1632	0.0104
9.7	19.79	1789	0.0108
10.3	19.79	1834	0.0108
10.7	19.47	1863	0.0109
11.1	19.47	1891	0.0109
12.0, one step	19.47	1951	0.0112
12.0, one jump	25.83	1957	0.0219

The wind gets hotter and more compact as the reservoir rises: r_{drop} moves inward from 30.0 to 19.5 R_p without a reversal, and the outer share of the metastable column stays in 0.0104–0.0119 over a factor 5.7 in composition. The jump-seeded solution sits at 25.8 R_p where its own reservoir gives 19.5, and carries twice the outer column share of any other rung, while its $T(12R_p)$ matches its control to 0.3 per cent. It is the same wind with a hotter, more extended outer atmosphere on top.

83.3 The mechanism, read from that pattern

The outer state is inherited from the seed and the JFNK finish does not re-solve it. The density out there contributes almost nothing to the residual norm, so $\|R\|$ cannot separate the two states; the elemental-flux spread the closure tests is measured on a steady window that is flat in both. He 10830 is optically thick at these columns, so a hot, extended outer region fills the line wings and lifts the equivalent width while holding about one per cent of the metastable column. Small mass, large opacity, invisible to both convergence tests.

83.4 The physical judgement, stated first

The 20–30 R_p region in which the two solutions differ is where the collisional-validity diagnostic of Sect. 81. places the exobase, 19.8–26.7 R_p , and where $Kn > 1$. *The fluid solution is not valid in exactly the region that makes the difference.* Two consequences follow, and neither is comfortable.

- The ground for discarding the jump-seeded solution is **continuity of the trend, not a physical criterion**. Nothing in the equations being solved prefers one outer state over the other, because the equations do not apply there.
- By the same argument, the absolute equivalent widths of **every** rung of that ladder carry a systematic error the closure tolerance does not catch. The size is small — the region outside $10R_p$ is about one per cent of the metastable column — but it is not zero, and it is not bounded by any test currently applied.

83.5 What this says about the convergence criteria

`exhale_info = 0`, a wind residual under target, and a flat elemental flux across the steady window are together sufficient to certify the solution *inside* the collisional region, and silent outside it. That is not a defect of the criteria: a residual norm weighted by density cannot see a region that carries no density, and a flux-spread test measured over a window is blind to what the window excludes. It is a limit on what “converged” means for a subsonic wind whose domain extends past its own exobase. Where an observable is optically thick and integrates the whole column — He 10830 here, Ly- α more so — the certified region and the region that sets the observable are not the same region.

83.6 What this says about running a scan

A ladder scan has to be climbed one step at a time. The step that worked here was $1.08\times$ in reservoir; the $1.50\times$ jump did not. There is no way to detect the failure from the solver’s own output, so the protection is procedural: keep the step small, and check a sampling-free outer diagnostic (r_{drop} , or the outer share of whatever column the observable integrates) for monotonicity along the scan.

83.7 What was and was not tested

Path independence was checked directly at $\text{He}/\text{H} = 3.0$, from a seed $3.7\times$ away in composition and in the downward direction the ladder never takes: it returns the same r_{drop} to the digit, with everything else inside the closure tolerance. The 5.0, 8.0 and 2.09 arms were not checked this way. Determinism was checked separately at 10.31 (`.../heh10p3_ctl/`, same seed, re-run): `Hydro_ioniz_adv.txt` and `Ion_species_adv.txt` reproduce bitwise, which says nothing about seed dependence. **No code was changed by this section**; the ladder table, the crossing and the figure were re-derived on the control arm, and the record is `LHS1140b/kzz_decision.md` Sect. 8.3 and `docs/lhs1140b_exhale_vs_pwinds.tex`, “The outer region a converged wind inherits”.

84. Metals had no way back: element projection when the helium mass fraction reaches one (2026-08-28)

The judgement first. The element projection of the binary H/He diffusion operator destroyed trace-metal nuclei with no sink and no way back, and that is physically wrong whatever it costs. In a cell whose helium mass fraction reaches $X = 1$ the hydrogen count $n_{\text{H}}^{\text{new}} = (1-X)\rho/m_1$ is exactly zero and the projection multiplied the metals by $r_{\text{H}} = n_{\text{H}}^{\text{new}}/n_{\text{H}}^{\text{old}} = 0$. Zeroing them *at that instant* is right: in this two-component split the metals are part of component 1, slaved to hydrogen at a fixed metal/H, so a cell with no hydrogen carries no metals and component 1 carries no mass. What is wrong is the step after. Hydrogen returns to that cell through the shortfall deposit at the end of the same projection; **the metals do not**, and nothing later can lift them, because every writer to the metal columns of `f_sp` on this path multiplies what is already there. The result is a contiguous band of cells holding exactly zero carbon, nitrogen and oxygen between cells at the reservoir ratio — a state these equations cannot produce, since they contain no sink for a metal nucleus. The fix is that the metals return **with** the hydrogen.

84.1 The path, in the code

element_diffusion_step solves for the helium mass fraction X and hands the new element totals to project_elements (src/modules/functions/binary_element_diffusion.f90). There, per cell,

```
nucH_new = (1.0d0 - Xhe(j))*msum(j)/m_1
if (nucH_old(j) .gt. 1.0d-30) then
    rH = nucH_new/nucH_old(j)
else
    rH = 0.0d0
endif
...
do im = 1, n_mion ! metals slaved to hydrogen
    f_sp(j,mion_fsp(im)) = f_sp(j,mion_fsp(im))*rH
enddo
...
if (nucH_new .gt. gotH) &
    f_sp(j,isp_HI) = f_sp(j,isp_HI) + (nucH_new - gotH)
```

Step k : X reaches 1, so $\text{nucH_new} = 0$, $rH = 0$, and every hydrogen species and every metal ion of the cell is multiplied by zero. Step $k+1$: nucH_old is now zero, so $rH = 0$ again for want of a ratio to preserve, while $\text{nucH_new} > 0$ — and the deposit on the last line puts the hydrogen back as neutral H. There is no such line for the metals. Their zero is final.

84.2 The symptom that led here

LHS1140b/exhale/branch_diagnosis/ records the measurement (no code was changed by that work). At $K_{zz} = 10^7$ the reservoir scan reported two converged states at the same reservoir He/H, $\log_{10} \dot{M}$ 7.54 against 7.87 and He 10830 red-pair equivalent width 0.4287 against 1.3248 %Å, both with $\text{info} = 0$. The hot one carries 172 cells, 1.0333–1.6835 R_p , at exactly zero C, N and O in every stage, with the same elements at C/H = 2.639×10^{-4} immediately above and below. On the cool solution the C I lines carry 95.6 % of the volume-integrated cooling over 1.0–1.2 R_p ; deleting the carbon drops the integrated cooling of that band by a factor 23 at $0.91 \times$ the heating and lifts $T_{\max}$ from 3899 K to 5879 K. Refilling C, N and O at the handoff's own reservoir ratio and re-solving the wind with the same handoff collapsed the hot state onto the cool branch ($\log_{10} \dot{M} = 7.530$, $T_{\max} = 3890$ K, EW = 0.4065 %Å). Of the 550 output*/Ion_species.txt files in the tree, 130 carry such a band; 128 are in LHS1140b/exhale/kzz_profile_scan, the other two are runs already on record as failed. **No regression golden and no backup/regression/ case output carries one**, which is why the repair below moves nothing.

84.3 Why $\text{info} = 0$ and $\|R\| \sim 10^{-4}$ could not see it

Both states report $\text{info} = 0$ and a residual of a few times 10^{-4} , and both are right to. **The steady residual is a residual of the hydro and energy equations**, and their solution with the metals removed is a perfectly good solution of the equations as posed: the cooling function simply has fewer terms. Nothing in that norm asks whether the composi-

tion it was evaluated on is a composition the transport equations can reach. The elemental-flux closure does test the composition, but on the radial spread of the elemental *fluxes* over a window at $2R_p$ and above — outside the band. The handoff's own elemental conservation check (`src/utls/lower_profile_schema.py`) is stated on the *chemistry* profile at its deepest level, not on the wind. **Elemental conservation was nowhere checked for the wind**, which is the gap this section closes.

Sect. 83. is the same shape of failure at a different place: there two converged solutions differed only outside the exobase, where the residual norm carries no density and the flux window does not reach, and He 10830 saw the difference because it is optically thick and integrates the whole column. The statement the two sections share is that `info = 0` plus a residual under target plus a flat elemental flux certify the solution **on the quantities and in the region those tests are evaluated on**, and are silent elsewhere. What is new here is that the silence is not about a region: it is about a *conserved quantity that was never formed into a test at all*.

84.4 A third failure mode, not either of the first two

- **Sect. 70.** (HD 209458 b, $K_{zz} = 0$): two wind solutions a factor 2.8 apart in $\dot{M}$, both `info = 0`, alternating in a period-2 limit cycle of the outer composition/wind Picard loop, on compositions $\sim 1\%$ apart. Every backup/phase_d_baseline/ run was checked for the band and **none has one**, so that bistability is not this defect and stays open. (Caveat: only the final state of each run is on disk; the alternating passes of `new_kzz0_undamped` were not saved with species densities.)
- **Sect. 83.** (LHS 1140 b, jump-seeded rung): two solutions that agree below the exobase to seven digits and differ only above it, where the fluid equations do not apply. Here the two states differ from $r = 1.0004R_p$ upward — by 30 % in T by $1.0013R_p$ — and the difference is made in the collisional part of the atmosphere ($Kn_{\text{bulk}} \sim 10^{-3}$ at T_{max}).
- **This section:** a state that is not a solution of the equations at all, because an element has been deleted from a band of cells by the discretization and nothing can put it back.

The three should not be merged.

84.5 The repair, and the abundance it uses

The metals are a *ratio*, not a density: what the projection has to preserve for them is n_X/n_H . Multiplying by r_H does that in a cell that had hydrogen, and that branch is untouched. Where n_H^{old} is zero the ratio is 0/0 — the cell carries no metal/H of its own — and the operator supplies one:

```
if (nucH_old(j) .gt. 1.0d-30) then
  do im = 1, n_mion
    f_sp(j,mion_fsp(im)) = f_sp(j,mion_fsp(im))*rH
  enddo
else if (thereis_metals) then
  do ie = 1, n_melem
```

```

    i0e = melem_i0(ie)
    f_sp(j,mion_fsp(i0e)) = melem_ab(ie)*nuch_new
    do k = 1, melem_top(ie)
        f_sp(j,mion_fsp(i0e+k)) = 0.0d0
    enddo
enddo
endif

```

melem_ab is not one candidate among several. It is the metal/H that **this operator's own mass budget already assumes**: the mass per hydrogen nucleus of component 1 is $m_1 = \text{mass_per_H_nucleus_without_He}() = m_H + \sum_X \text{melem_ab}(X)A_X$, so the two-component closure $m_1 n_H + m_{\text{He}} n_{\text{He}} = \rho$ that the write-back is required to meet is stated at that ratio and no other. It is also the reservoir ratio set_IC builds a cold start from, the one load_IC rebuilds an element the restart file does not carry from, and the one Sect. 80. rescales a restart column onto. The alternatives were considered and rejected: the cell's own previous ratio is exactly what does not exist there, and a neighbour's ratio would make the result depend on the grid.

All of it goes into the **neutral ground stage**, for the same reason the hydrogen and helium shortfalls go into H I and He I: a cell that held no hydrogen held no ionization split either, ioniz_eq owns the split within an element and resolves it from the element total on the next call, and this is already the convention of set_IC, of load_IC, and of the trace-metal re-seed branch of element_diffusion_step itself.

The repair prevents the band from being *created*. It does not reconstruct one that a restart file already carries: there $n_H^{\text{old}} > 0$ in the emptied cells, so the preserving branch is the one that runs, and it faithfully preserves a ratio of zero. That is the right behaviour for a restart — silently rewriting loaded data would be worse — and it is why the census below is unconditional.

84.6 The check the pipeline did not have

metal_hydrogen_ratio_departure now runs at the end of every diffusion step. It measures $\max |(n_X/n_H)/\text{melem_ab}-1|$ over the cells 1.. N and the elements the reservoir states, and it counts the (cell, element) pairs holding no metal nuclei at all while hydrogen is present. The departure is a magnitude and not an error — he_metal_diffusion moves the metals relative to hydrogen on purpose, and a restart may carry a settled column — so it is only reported, under the existing EXHALE_DIFFUSION_CHECK=1 switch, together with the rest of the step diagnostic. The count is different in kind: it is zero on any state these equations can reach, so it raises a warning on stdout the first time it is nonzero **whether or not the diagnostic is on**. No new input key was added.

84.7 The clamp is a limiter, and the comment that said otherwise was wrong

The projection was preceded by

```

! Round-off assertion (NOT a limiter): the M-matrix argument in the
! header gives 0 <= X <= 1 per step; only round-off can leave the range.
where (Xhe .lt. 0.0d0) Xhe = 0.0d0

```

where (Xhe .gt. 1.0d0) Xhe = 1.0d0

That claim had never been measured, because the diagnostic printed the *clamped* X , which cannot tell an assertion from a limiter — both print 1.0. Measured now, on the reproduction run below: X exceeds 1 before the clamp in **140 of 7962 steps, by up to 0.520** (i.e. $X = 1.52$), and 105 of those excursions exceed 0.1. The lower end holds: $\min X$ stayed above 0 in every step of both runs. **The clamp is a limiter on the upper side.**

The reason is in the discretization, not in round-off. The settling flux $-\rho D_{12} G X (1-X)$ is linearized with the $(1-X)$ factor lagged, $W_f = \rho_f D_f G_f (1-X_f^{\text{lag}})$, and entered as $-W_f X$. In all three branches of the Péclet hybrid the face coefficients satisfy $\text{PdL}(f) + \text{PdR}(f) = -W_f$, so the row sum of the backward-Euler matrix is

$$\frac{\rho_j}{\Delta t} - K_j \left(r_{j+1/2}^2 W_j - r_{j-1/2}^2 W_{j-1} \right)$$

and **not** $\rho_j / \Delta t$. A uniform $X = 1$ is therefore not a fixed point of the discrete row unless the lagged settling coefficient is divergence-free, and the comparison principle that the header's M-matrix argument invokes does not apply — at either end. What separates the two ends here is measurement, not proof: the lower one held in every step of both runs, the upper one did not. The header comment and the corresponding sentence of docs/binary_diffusion_design.md Sect. 3 (“a clip at round-off (1e-15) is kept only as an assertion, not as a limiter”) are corrected to state the measurement. **What is not done here.** Making the upper bound hold by construction means changing the linearization of the settling term — for instance solving the $X(1-X)$ product with both factors at the new level, or symmetrizing the Picard sweep — and that changes every diffusion-on result, including two regression goldens. It is a formulation decision, it is recorded here, and it is left open. What this section does is remove the *consequence* that made the clamp destructive: with the metals returning with the hydrogen, a cell that touches $X = 1$ recovers instead of being permanently emptied. The $X = 1$ excursions themselves still happen — 140 of them in the repaired run, against 142 in the unrepaired one.

84.8 What was measured

The step kzz1e7/heh3p8 (clean) $\rightarrow$ kzz1e7/heh4p2, re-taken with the same binary in both columns, the projection branch being the only difference. This is the step LHS1140b/exhale/branch_diagnosis/metal_dropout_reproduce.py takes, run with EXHALE_DIFFUSION_CHECK=1 and with the He I 10830 synthesis added; both runs were made outside the repository, so no recorded output in it was overwritten.

	before	after
zero C/N/O cells, Ion_species.txt	157 (1.0429–1.6716 R_p)	0
zero C/N/O cells, Ion_species_adv.txt	157	0
exhale_info	0	0
$\log_{10} \dot{M}$	7.52	7.52
$T_{\max}$	4182.6 K at 1.358 R_p	4185.6 K at 1.358 R_p
He 10830 red-pair EW	0.5306 %Å	0.5249 %Å
diffusion steps	7916	7962
steps whose X reached 1	142	140
max X before the clamp	1.521	1.520
steps with a vanished (cell, element)	7541, up to 471 pairs	0
max metal/H departure from <code>melem_ab</code>	1.0 (an element gone)	8.1×10^{-13}

The unrepaired column reproduces the recorded diagnosis exactly — 157 cells over the same 1.0429–1.6716 R_p , 7916 diffusion steps, 142 of them touching $X = 1$ — so the added diagnostics are numerically inert and the projection branch is the sole cause of the difference. The repaired solution matches the cool branch continued to this reservoir by refilling and re-solving (`branch_diagnosis/cool_branch_continuation.txt`: $\log_{10} \dot{M} = 7.520$, $T_{\max} = 4184.4$ K, EW = 0.5287 %Å).

Note what the $T_{\max}$ and EW rows do *not* say: in this particular step the band happened to cost only 3 K and 1 % of the equivalent width. The factor-23 cooling loss and the 3899 → 5879 K jump quoted above belong to the recorded scan arms, where the band sits where the C I cooling is decisive. The band is a defect wherever it appears; how much it moves the answer depends on where it lands.

84.9 Verification

- `make check`: 7/7 byte-identical. Nothing moves, and the reason is structural rather than lucky: the preserving branch $n_{\text{H}}^{\text{old}} > 0$ is bitwise the previous code, and no golden state ever empties a cell of hydrogen.
- `make diffusion_tests`: 31/31 (28 as before, plus 3 assertions of the new T13). **T13 fails on the pre-fix projection** — 21 of 21 band cells hold hydrogen and no metals, metal/H departure 1.0 — which is the point of adding it. It seeds a closed metal-bearing column with a band of pure helium, so $n_{\text{H}} = 0$ and $X = 1$ exactly, lets diffusion refill the band, and requires that no cell holding hydrogen holds zero metals and that n_X/n_{H} equals `melem_ab` to 10^{-12} (measured 5.3×10^{-15}).
- `run_fcheck.sh` (`-fcheck=bounds,do,mem`) clean. It had to be repaired first, for two reasons that predate this change and are unrelated to it: its repo root was still one level up from `backup/regression/` (so every `make` in it failed with “No rule to make target ‘distclean’”), and it copies `HD209458b/input.inp` into an empty scratch directory without turning off the `Load IC?` True that file now carries, so `load_IC` stopped on a missing restart and the script reported that error `stop backtrace` as a runtime check firing. With both fixed the HD 209458 b bounded run is clean ($\log_{10} \dot{M} = 9.60$).

- That case runs with element diffusion **off**, so it does not reach the code this section changes. Two runs were added for that: `diffusion_tests.x` built with the same flags (31/31, and T13 is the test that enters the new branch), and the `lower_profile` regression case — element diffusion together with metals — bounded to 400 steps. Both clean, no trap.

Not done, and left to a decision. The 128 contaminated arms of LHS1140b/exhale/kzz_profile_scan are unchanged on disk. The repair prevents the band from forming; it does not repair a stored state, and re-running that scan is a data decision, not a code one.

The open item above is closed by section 86. (2026-08-28): the two factors of the drift product $X(1-X)$ are now taken from opposite sides of the face and both at the new time level, so the flux shuts off at either end of the composition axis as it does in the continuum, and the $X > 1$ excursions are gone. The clamp is an assertion at both ends.

85. A structure scale that is a scale of the state, not of the pressure (2026-08-28)

`src/utls/collisional_validity.py` divides a mean free path by a local structure scale L to get the Knudsen number that decides whether a wind solution is a continuum solution. Until now $L = \min(H_p, L_v, r)$, and below Mach 0.01 the velocity term switched itself off, so in a subsonic region — LHS 1140 b runs at Mach 1.6×10^{-4} near the base — L was the pressure scale height alone.

The pressure is the one field that can go flat while the gas does not. $p = nkT$; across a heating front ρ falls and T rises with nearly the same logarithmic slope, so $d \ln p / dr$ passes through a broad near-zero while the state changes fast. Measured in LHS1140b/exhale/heh0p55: H_p runs 9.13×10^6 cm at $1.0244R_p$, 2.55×10^8 cm at $1.0453R_p$ (its maximum, 3.05×10^8 , is at 1.0486) and 1.41×10^8 cm at $1.0648R_p$, while at $1.0453R_p$ H_T is 1.80×10^7 and H_ρ 1.68×10^7 cm — fifteen times shorter — and $\text{Kn}(\text{HI})$ showed a hole of $5.4 \times (3.67 \times 10^{-5}, 6.82 \times 10^{-6}, 3.56 \times 10^{-5})$ at those three radii) in a quantity that is monotonic on either side.

The fix is the state vector. The continuum closure expands about a local Maxwellian, which is fixed by (n, T, u) ; those logarithmic gradients plus the spherical divergence are the complete set, and the change of the state over a distance is a vector whose length is the sum in quadrature,

$$\frac{1}{L^2} = \left(\frac{d \ln \rho}{dr} \right)^2 + \left(\frac{d \ln T}{dr} \right)^2 + \left(\frac{|dv/dr|}{c_s} \right)^2 + \frac{1}{r^2}, \quad L = \max \left(\frac{1}{\sqrt{\dots}}, \Delta r_{\text{cell}} \right). \quad (30)$$

Pressure is not in it: it is a derived field, and adding it would count the density and temperature gradients twice wherever they do not cancel while contributing nothing exactly where they do. The quadrature sum, not a minimum over the same lengths, because a minimum jumps discontinuously to the next length whenever the field it is taken from flattens — the failure above in general form — whereas in quadrature a flat field contributes zero and L still cannot exceed the shortest individual scale. The velocity enters as $|dv/dr|/c_s$ and not $|v^{-1} dv/dr|$ because the term the closure drops is the viscous stress, whose size relative to the pressure is $\sim \lambda |dv/dr|/\bar{v}$: a bulk-velocity change distorts the distribution function in proportion to the thermal speed, not to the local bulk speed. That removes the Mach-number floor — `MACH_FLOOR` is gone — which existed only because normalizing by v diverges wherever the base sound-wave layer crosses $v = 0$. The

floor at the local cell width says a gradient claiming structure below one cell is measuring the mesh; it binds in 3 of 500 cells of one of the four runs, all at $r < 1.002R_p$, and in no critical region.

No conclusion moves. Same runs, same outputs, nothing rebuilt or re-run. HD 209458 b stays validated through its critical point (sonic point $4.085R_p$, which L does not enter; max Kn in the critical region $0.0214 \rightarrow 0.0186$), all three LHS 1140 b solutions stay unvalidated with no critical point in the domain (max Kn $1.9\text{--}3.0 \rightarrow 2.0\text{--}3.1$). What moves is at the 10–30% level and in both directions: the exobase $19.8\text{--}26.7 \rightarrow 18.8\text{--}25.4R_p$, Kn_{bulk} at the outer boundary $1.3\text{--}2.3 \rightarrow 1.4\text{--}2.4$, Kn_{bulk} at the $8\text{--}9.5R_p$ where the p-winds retrieval puts its sonic point $0.14\text{--}0.30 \rightarrow 0.13\text{--}0.29$, and the He 10830 line-forming layer at $1.1\text{--}3R_p$ $\text{Kn} \leq 0.021 \rightarrow \leq 0.019$. The one large change is inside $1.1R_p$, where the old H_p sat on its isobaric plateau: Kn_{bulk} at $1.1R_p$ goes from $4 \times 10^{-5}\text{--}1.2 \times 10^{-4}$ to $2.5\text{--}4.2 \times 10^{-4}$. Mach maxima, $\log_{10}\dot{M}$ and the coupling times are independent of L and unchanged. docs/collisional_validity.md carries the recomputed tables and a row-by-row account of what moved; the numbers quoted in section 81. are the ones from the old scale and are superseded by it. Figure docs/figures/lhs1140b_knudsen.pdf regenerated (that block of LHS1140b/make_memo_figures.py only). Post-processing diagnostic, no Fortran touched, goldens untouched.

86. The drift flux has to vanish where the composition runs out: the discretization of $X(1-X)$ and the bounds on the helium mass fraction (2026-08-28)

The judgement first. The drift (settling) term of the binary H/He transport equation carries the factor $X(1-X)$ and vanishes at both ends of the composition axis, because a cell with no helium has none to send and a cell with no hydrogen has nothing to send it in exchange. The discretization lagged the $(1-X)$ factor at the previous iterate, which throws that shutoff away: the discrete flux into a cell already at $X = 1$ was not zero, and the step overfilled it. Measured, X reached 1.52. That is not a small inaccuracy, it is a discrete operator that does not have the property the continuum flux has, and the clamp that hid it was doing the work of a limiter. The drift product is now taken with its two factors from **opposite sides of the face** — helium from the cell it leaves, hydrogen from the cell it enters — and both at the new time level, and the resulting nonlinear step is solved by Newton. The excursion outside $[0, 1]$ is gone: over every case measured the solve now stays strictly inside the range.

Section 84. fixed the *consequence* (a cell driven to $X = 1$ lost its metals permanently) and left the cause open in as many words: “Making the upper bound hold by construction means changing the linearization ... it is left open.” This section closes it.

86.1 Where the old discretization lost the bound

The flux is (module header of src/modules/functions/binary_element_diffusion.f90)

$$J = -\rho(D_{12} + K_{zz}) \frac{dX}{dr} - \rho D_{12} G X(1-X), \quad (31)$$

and the old face form entered the drift as one lagged coefficient times one implicit factor,

$$W_f = \rho_f D_f G_f (1 - X_f^{\text{lag}}), \quad J_{\text{drift}}(f) = -W_f X(f), \quad (32)$$

with $X(f)$ central or upwind by the Péclet switch. Two things follow. The row sum of the backward-Euler matrix is $\rho/\Delta t - K(r_R^2 W_j - r_L^2 W_{j-1})$ and not $\rho/\Delta t$, so no comparison principle applies (section 84. established this much). More to the point, and this is the physical statement: **a face between two cells both at $X = 1$ still carries a drift flux**, $-W_f$, unless the lag already knows that $X_f = 1$. The endpoint states of the composition axis are not stationary states of the discrete operator, while they are exact stationary states of the continuum flux. A cell driven toward pure helium therefore kept receiving helium after it had run out of hydrogen to exchange, and the clamp turned the overshoot into an exact 1 — which is where the metal band of section 84. came from.

86.2 The discretization that keeps it

The drift is a **counter-flow**: the helium mass flux $-BX(1-X)$ (with $B = \rho D_{12}G$) is matched by an equal and opposite hydrogen flux, since the two components close, $J_1 = -J_{\text{He}}$. A donor-cell rule for a counter-flow has to take each element's mass fraction from the cell **that element** leaves, and the two elements leave from opposite sides. With $B \geq 0$ helium drifts inward, so helium is donated by cell $j+1$ and hydrogen by cell j :

$$J_{\text{drift}}(f) = -B_f X(j+1)[1-X(j)] \quad B_f \geq 0 \quad (\text{helium inward}), \quad (33)$$

$$J_{\text{drift}}(f) = -B_f X(j)[1-X(j+1)] \quad B_f < 0 \quad (\text{helium outward}). \quad (34)$$

Both factors are at the new time level. A donor at $X = 0$ sends nothing; an **acceptor** at $X = 1$ receives nothing. Those are the two statements that bound X , and they are the discrete image of the two limits the continuum factor $X(1-X)$ expresses.

The Péclet hybrid stays, with its switch now on the drift coefficient itself, $|B_f|\Delta r \leq 2(A_f + E_f)$, and the central branch evaluates the product at the face average, $X_f(1-X_f)$. That branch is inside the bound as well, and for a reason worth writing down: with $X(j) = 1$,

$$X_f(1-X_f) = \frac{[1+X(j+1)][1-X(j+1)]}{4} \leq \frac{1-X(j+1)}{2}, \quad (35)$$

so the Péclet condition $|B_f| \leq 2A_f/\Delta r$ bounds the drift flux by the gradient flux $A_f[1-X(j+1)]/\Delta r$ that is carrying helium *out* of that cell at the same face. The central branch cannot reverse the sign of the total face flux at a cell sitting on the bound. (The switch is stricter than the old one by the factor $1-X$, so slightly more faces are upwind.)

The bounds. Let X solve the implicit step and suppose it first touches 1 in cell m , every other cell still inside $[0, 1]$. Then the time term $\rho(X_m - X_m^{\text{old}})/\Delta t \geq 0$; the gradient flux leaves m at both faces because X_m is the maximum; the upwind advection contributes $\rho|v|(X_m - X_{\text{donor}})/\Delta r \geq 0$; and the drift flux is an outflow or exactly zero at both faces by the rule above. Every term of the row has the same sign, and their sum is the row residual, which is zero — so no term can be strictly positive and $X_m > 1$ is impossible. $X \geq 0$ is the same statement, because the scheme is **exactly symmetric** under $X \rightarrow 1-X$, $B \rightarrow -B$: the donor and the acceptor exchange roles and the central product is even about $X_f = 1/2$. That symmetry is the binary symmetry of the transport equation itself — there is one independent diffusive flux and the two components carry it in opposite directions — and the old lagged form broke it, which is why one end held in practice and the other did not.

The row sum is still not $\rho/\Delta t$ and does not need to be: the argument is made cell by cell on the residual of the nonlinear step, not from a comparison principle for a linear system.

86.3 Newton, and why not another Picard sweep

The step is now nonlinear in X^{new} , so it is solved by Newton: `composition_residual` builds the residual and the two face-flux slopes, and `solve_mass_fraction` assembles the tridiagonal Jacobian and takes the correction. The Jacobian is an M-matrix for any iterate in $[0, 1]$ — the slopes are $\partial J / \partial X_{\text{left}} \geq 0$ and $\partial J / \partial X_{\text{right}} \leq 0$ in all three branches — so the correction is well posed, and the nonlinearity is quadratic, so the iteration converges in one to three passes. A pass that does not reduce the residual is halved.

A Picard sweep on a lagged factor was the alternative and is what was there. It does not restore the bound, whatever the sweep count: at every sweep the flux at a face between two cells at $X = 1$ is $-W_f \neq 0$, and only the *converged* fixed point has the shutoff. The two-sweep version in use was not converged — measured below, the LHS 1140 b step needs one to three Newton passes to reach its residual floor — and it is the unconverged iterate that was clamped. Reducing Δt was not considered a fix: the bound is a property the operator should have at any step, the implicit operator is otherwise unconditionally stable, and the relaxation deliberately runs at 10^5 cell diffusion times to reach the composition's steady state at all (section 27.).

Convergence measure. The Newton residual is measured *relative to the size of the terms of its own row*. At a step long against the cell diffusion time the time term and the flux divergence each exceed their difference by $\sim \Delta t / t_{\text{diff}} \sim 10^5$, so an absolute residual bottoms out at the round-off of the larger term (measured: $\sim 10^{-11}$ in units of X) and cannot be a convergence test. The relative residual reaches $\sim 7 \times 10^{-12}$ on LHS 1140 b in one to three passes; below that the round-off of the tridiagonal solve, amplified by the same conditioning, is all that is left, so an already-small residual that no longer halves is taken as converged. The smallness test is not one number: the reachable floor is set by the conditioning of the tridiagonal solve, which is 10^{-12} in a wind and $\sim 10^{-5}$ in the $K_{zz} = 2 \times 10^{12}$ homopause column of test T7, where the time term is negligible against the eddy term. It is therefore either an absolute 10^{-8} or a drop of 10^{-6} from the residual the step started at, whichever is reached first. That floor test is **not** applied while the residual is still large: the bounds belong to the converged step, and stopping early loses them — with the floor test applied unconditionally, T14 fails at an excursion of 7×10^{-3} .

86.4 What is measured

The excursion outside $[0, 1]$ *before* the clip, which section 84. added to the step diagnostic and which is now also exposed as the module variables `he_fraction_over_one` and `he_fraction_under_zero` so a test can read the solve instead of its clip:

	before	after
T14, 400 steps, closed column at $b_0 = 20$, $\Delta t = 10^6$	max $X = \mathbf{1.514}$ ($+5.136 \times 10^{-1}$)	max $X = 0.99908$ (-9.2×10^{-4})
T13, 200 steps from a band seeded at $X = 1$ exactly	-8.584×10^{-3}	-8.587×10^{-3}
LHS 1140 b <code>heh2p13_diff_kzz1e9</code> , JFNK finish	-1.050×10^{-1}	-1.050×10^{-1}

The excursion is negative in every case now measured, i.e. the solve stays strictly inside $[0, 1]$ and never reaches the clip at all. It is *not* the round-off-sized positive number one would quote if some cell sat exactly on the bound: in these runs the solution approaches the bound and stops short of it, which is the behaviour the shutoff produces. The

clip lines are therefore assertions at both ends, and they are kept, with the excursion measured every step rather than asserted.

The old code fails T14 by 0.514 — essentially the 0.520 recorded in section 84. from a wholly different run, which is the reproduction that identifies the mechanism.

86.5 Golden movement

The change moves every diffusion-on result, so it moves two of the seven regression cases. The other five — `wasp_full`, `wasp_he23off`, `mol_base_handoff`, `mol_metals`, `mol_lyman_werner` — run with `He_diffusion` off, never enter the module and stay **byte-identical**.

Largest relative difference over the column, against the previous golden, on the two 12000-step snapshots that do run element diffusion:

case	ρ	p	T	heat	cool	v	$\log_{10} \dot{M}$
<code>mol_diffusion</code>	$10^{-11} \%$	$2 \times 10^{-11} \%$	$9 \times 10^{-12} \%$	$10^{-11} \%$	$2 \times 10^{-11} \%$	$1.4 \times 10^{-8} \%$	$10.58 \rightarrow 10.58$
<code>lower_profile</code>	$4 \times 10^{-9} \%$	$5 \times 10^{-9} \%$	$3 \times 10^{-9} \%$	$7 \times 10^{-9} \%$	$9 \times 10^{-9} \%$	$5.4 \times 10^{-7} \%$	$9.58 \rightarrow 9.58$

The worst-moving species is H3+ at $1.9 \times 10^{-4} \%$ (`mol_diffusion`) and HeITR at $1.3 \times 10^{-8} \%$ (`lower_profile`); the H₂ front of `mol_diffusion` is unchanged at $1.16036 R_p$ and `lower_profile` carries no molecular column.

These two cases do not exercise the bound at all, and that is why they barely move: X stays in $[0.2406, 0.2408]$ through all 12000 steps of `mol_diffusion` and in $[0.2468, 0.2496]$ of `lower_profile`, both at $K_{zz} = 10^9$ where the eddy term dominates the drift, and the largest excursion the diagnostic reports is -0.759 and -0.750 — X is nowhere near either end. What moved them is the second-order difference between the old and the new evaluation of the drift product plus the Picard sweep's non-convergence, both at round-off amplified by the step conditioning here. Golden re-snapshotted for the two, and make check re-run after: 7/7.

Newton cost, measured on the same runs: `mol_diffusion` converges in one pass in all 12000 steps (largest relative residual 5.9×10^{-13}); `lower_profile` in one pass in 11807 steps and two in 193 (9.8×10^{-13}). The wall time is unchanged within the noise of a shared machine.

86.6 The physical result this was worried about: LHS 1140 b does not move

The composition of LHS1140b/exhale/heh2p13_diff_kzz1e9 is the one the He I 10830 line is read from, so the fix was re-taken on it: the case re-solved from its own converged state, JFNK finish, post-processing pass and the He I 10830 synthesis at WINERED HIRES-Y ($R = 68,000$), once with a binary carrying the old linearization and once with this one, both outside the repository so that no recorded output was overwritten.

	old linearization	this section
exhale_info / $\ R\ $	0 / 4.285e-3	0 / 4.285e-3
$\log_{10} \dot{M}$	7.80	7.80
He 10830 red-pair EW (10832.60–10834.20 Å vacuum)	1.1404754 %Å	1.1404754 %Å
red depth	4.186959 %	4.186959 %
FWHM	0.2557731 Å	0.2557731 Å
X range over the run	0.6503–0.8950	0.6504–0.8950
largest excursion before the clip	-1.050×10^{-1}	-1.050×10^{-1}

The two solutions agree to 1.2×10^{-9} relative in ρ , v , p and T over the whole column (the one exception is He III at $1.0103R_p$, 3.5×10^{-3} relative, on a population that is 6×10^{-13} of the helium there). **This wind never approaches** $X = 1$, so the clamp never fired on it and the fix has nothing to remove; what is left is the second-order difference in the drift term, and at this K_{zz} it is invisible.

On the reservoir crossing: the equivalent width moves by -4.1×10^{-13} in relative terms, and with the local slope $d \ln \text{EW} / d \ln (\text{He}/\text{H}) = 0.3819$ measured over the upper rungs of the closure ladder that is $d \ln (\text{He}/\text{H}) = -1.1 \times 10^{-12}$, i.e. the crossing stays at $\text{He}/\text{H} = 11.730000$. The adopted result is untouched.

What this does not say. The recorded tpm_He10830.txt of that case gives EW 1.1289 %Å against the 1.1405 measured here; the difference is *not* this change, it is everything the working tree has accumulated since that file was written (section 84. among it). Both columns of the table above were produced with binaries differing only in the drift linearization, which is what makes them a controlled comparison.

86.7 Verification

- make check: five cases byte-identical, mol_diffusion and lower_profile moved by the amounts tabled above; goldens re-snapshotted for those two and make check re-run, 7/7.
- make diffusion_tests: 35/35 (31 as before, plus T14's three assertions and one added to T13). **T14 fails on the old linearization** at an excursion of +0.514, which is the point of adding it: it drives a closed column hard enough that the base reaches pure helium, requires that the run actually reach the boundary ($\max X \geq 0.99$, else the assertion is vacuous), and reads he_fraction_over_one / he_fraction_under_zero, the excursions the operator measures *before* its own clip. Reading the clipped X instead would have shown 1.0 in both columns and proved nothing.
- run_fcheck.sh (-fcheck=bounds, do, mem) clean. That case runs with element diffusion off, so as in section 84. two further bounded runs cover the code this section changes: diffusion_tests.x built with the same flags, and the lower_profile regression case bounded to 400 steps. Both clean, no trap.

86.8 Two more things found in passing, both in the step diagnostic

- The diagnostic printed $\max |J_{\text{He}} + J_1|$ as $\max(\text{sumJ}, \text{abs}(Jf(j) - Jf(j)))$ — an identical zero in the clothes of a measurement. A binary mixture has one independent diffusive flux and the hydrogen component carries $-J_{\text{He}}$

by construction, so there is nothing there to measure; the field is removed and the reason is written where it used to be printed.

- The trace-metal solver ends on where $(f_X \text{ .lt. } 0.0d0) \ f_X = 0.0d0$ under the comment “Round-off assertion, not a limiter: the M-matrix keeps $f_X \geq 0$ ” — the same unmeasured claim this section removes from the helium equation. That one is a different case and is left as it is: the trace equation is *linear* in f_X (the trace limit drops the $1 - f_X$ factor), and an M-matrix sign pattern with a nonnegative right-hand side is a real argument for $f_X \geq 0$. But it is an argument, so the excursion is now measured too (trace_ratio_under_zero, reported with the step diagnostic) instead of being asserted.

87. A rate coefficient that jumped by 1.57 at 4000 K: the He (2^3S) Penning rate, and the step it put in every figure (2026-08-28)

The judgement first. A rate coefficient is a continuous function of temperature. The He (2^3S)+H Penning rate in Cool_coeff.f90 was not:

$$\begin{aligned} T \leq 4000 \text{ K} : \quad & 1.9 \times 10^{-9} (300/T)^{0.07} \rightarrow 1.5849 \times 10^{-9} \text{ at } 4000 \text{ K} \\ T > 4000 \text{ K} : \quad & 9.1 \times 10^{-9} (300/T)^{0.50} \rightarrow 2.4921 \times 10^{-9} \text{ at } 4000 \text{ K} \end{aligned}$$

a factor **1.5724** across the seam. That is Taylor et al. (2025), ApJ 989:68, Table 2, transcribed correctly — the code was not wrong about what the paper says. The published fit is what cannot be used, and “it is the published value” is not a defence of an expression that is discontinuous in a state variable. It is replaced by the García Muñoz (2025) form, which is continuous everywhere.

Checking the fit against its own paper made the case decisive. Taylor et al.’s Figure 19 (bottom right) plots the quantity Table 2 is supposed to represent: their Maxwell–Boltzmann average of the Cohen & Lane (1971) and Morgner & Niehaus (1979) cross sections. Digitized here, that curve rises to 1.50×10^{-9} near 1400 K and falls to 1.29×10^{-9} at 4000 K. The tabulated low branch decreases monotonically from 1.83 to 1.58×10^{-9} and, being a power law, cannot produce a maximum at all:

T [K]	Fig. 19 (their rate)	Table 2 low branch	ratio	new rate
500	1.135	1.833	1.62	1.024
1000	1.467	1.746	1.19	1.207
1400	1.502	1.706	1.14	1.271
2000	1.470	1.664	1.13	1.321
4000	1.291	1.585	1.23	1.356

(all in $10^{-9} \text{ cm}^3 \text{ s}^{-1}$).

Above 4000 K the high branch is worse: 2.23×10^{-9} at 5000 K, above *every* cross-section determination collected in García Muñoz (2025) Fig. 5, whose curves lie between about 1.3 and 1.7×10^{-9} there. The appendix of Taylor et al. says how this happened — they fitted “an equation already implemented into our KPP routine for simplicity” with `scipy.optimize.curve_fit`, in two pieces, with nothing tying the pieces together at the join.

The replacement. García Muñoz (2025), A&A 698, A199, Fig. 5 publishes a closed form for the Maxwell–Boltzmann average of the Movre & Meyer (1997) total ionization cross sections,

$$k = 10^{-9} \exp(c/T + d_1 \ln T + d_2 (\ln T)^2 + d_3 (\ln T)^3) \quad [\text{cm}^3 \text{s}^{-1}],$$

$c = -8.64804\text{E}+01$, $d_1 = -2.86766\text{E}-01$, $d_2 = +8.68445\text{E}-02$, $d_3 = -5.73001\text{E}-03$, quoted there as accurate to better than 2% from 200 to 10^4 K. It reproduces that paper’s own Table A.5 (1.02, 1.32, 1.35, 1.27×10^{-9} at 500, 2000, 5000, 10^4 K) to 0.4%, measured here. Note the paper’s network file writes the third term as $\exp(c)$; only the c/T of the Fig. 5 caption reproduces Table A.5, and the code comment records that.

This is not a case of trading one paper’s fit for another’s. Taylor’s Figure 19 and the García Muñoz curve agree to 10–20%; it was the *tabulated* fit that was the outlier. Re-digitizing Movre & Meyer’s Fig. 10 cross sections independently and Maxwell-averaging them reproduces Table A.5 to 5–10%, the accuracy of reading a figure. Movre & Meyer give cross sections only, over 0.01–1.0 eV, and no rate coefficient — the 200 K end of the fit leans on García Muñoz’s low-energy $E^{-0.2}$ extrapolation (15% of the integral weight at 200 K) and the 10^4 K end on the high-energy E^{-1} one (40% of the weight at 10^4 K), which is why the validity range is stated where the expression is written.

The older prescriptions are outliers for identifiable reasons and are not candidates: Roberge & Dalgarno (1982)’s temperature-independent 5×10^{-10} (what EXHALE used before Taylor) sits far below every cross-section calculation, and Fassia et al. (1998)’s $7.5 \times 10^{-10}(T/300)^{1/2}$ is Bell (1970)’s single 300 K point stretched with hard-sphere scaling, which the energy dependence of the measured cross section contradicts.

The H₂ channel and the 0.9:0.1 branching. Both partners ionize the metastable through two channels that share one cross section, Penning ($\text{He}(2^3\text{S}) + \text{X} \rightarrow \text{He}(1^1\text{S}) + \text{X}^+ + \text{e}^-$) and associative ($\rightarrow \text{HeH}^+ + \text{fragment} + \text{e}^-$), and García Muñoz splits the total “an average 0.9:0.1”. The published network file applies exactly that split through the amplitudes of two rows per channel (0.9e–9 / 0.1e–9 for H; 4.86740e–12 / 5.40822e–13 for H₂, summing to 5.408222e–12), which settles what the paper text leaves implicit. EXHALE had charged 100% of both totals to Penning. It now follows the split, and the distinction matters in a way a single scale factor does not:

- the **total** removes the metastable — both channels quench it — so every sink term keeps the sum. Using 0.9 there would have made the metastable density 10% too high, and the metastable density is the observable;
- $f_{\text{penning_HeI23S}} = 0.9$ scales only the terms that leave a lasting proton or H₂⁺ behind and that deposit the Penning exothermicity;
- the molecular system carries HeH⁺ with its own balance row, so the remaining 0.1 is now a real HeH⁺ source there. The atomic systems have no HeH⁺ row and drop it, on the stated ground that in an H₂-poor gas HeH⁺ dissociatively recombines back to He + H faster than it reacts onward, so the branch is a metastable sink but not a lasting proton or electron source. Where H₂ is abundant that cycle is broken by $\text{HeH}^+ + \text{H}_2 \rightarrow \text{H}_3^+ + \text{He}$, which is why the molecular system does not take the same shortcut.

The H₂ rate itself is now the published $5.408222 \times 10^{-12} T^{0.675388} \exp(-696.275/T)$ rather than a fit made here to the four Table A.5 points. The old fit was accurate (0.13%) but reproduced the *total*, so charging all of it to Penning overstated H₂⁺ production from this reaction by 1/0.9.

Naming. penning_HeI_23S and penning_HeI23S_H2 returned totals, not Penning rates, and had inconsistent names for the same physics. They are now ioniz_HeI23S_H and ioniz_HeI23S_H2, both elemental functions, named for what they compute.

This step was the discontinuity in the LHS 1140 b memo figures. The detector in LHS1140b/exhale/metastable_step_diagnosis flags a cell whose neighbour ratio departs by more than 10% from the geometric mean of the ratios on either side. Before and after, on the same three runs:

run	flags before (eq / adv)	flags after
heh0p55	6 / 6	0 / 0
heh2p13_diff_kzz1e9	6 / 6	0 / 0
flux_closure/heh11p1/k01	4 / 8	0 / 4

The four that remain in k01 are at $r = 19.15, 19.47, 28.57$ and $29.05 R_p$, were present before with the same values, and are outer-boundary features, not Penning $T=4000$ ones. All three runs do cross 4000 K, so the test is live: the inner crossing is at 1.1775, 1.1084 and $1.1123 R_p$. At the crossing cell of heh0p55 the neighbour ratio $n(2^3S)_{j+1}/n(2^3S)_j$ went

before: 1.0199 1.0194 **0.6651** 1.0232 1.0227 (a 33% drop in one cell)

after: 1.0190 1.0185 1.0180 1.0175 1.0171

and the other two runs behave the same way ($0.6681 \rightarrow 1.0242$, $0.6968 \rightarrow 1.0352$).

What moves. The new coefficient is *lower* than the old one everywhere in the wind (ratio new/old 0.69 at 1000 K, 0.86 just below 4000 K, 0.54 just above, 0.81 at 10^4 K), so the metastable lives longer and the 10830 line gets stronger. Removing the step is the small part of this; the change in level is the large part.

Regression matrix, current binary against the goldens as they stand ($\max |\text{new} - \text{golden}| / \max |\text{golden}|$ over the column, pointwise relative maximum in parentheses):

case	ρ	v	p	T	heat	$\log_{10} \dot{M}$
wasp_full	8.6e-05 (3.3e-03)	7.8e-04 (1.8e-02)	4.2e-05 (1.5e-03)	6.4e-04 (1.3e-03)	2.6e-03 (5.8e-03)	13.23 $\rightarrow$ 13.23
wasp_he23off	0	0	0	0	0	13.23 $\rightarrow$ 13.23
mol_base_handoff	2.2e-06 (2.2e-02)	1.1e-02 (2.2e+00)	4.3e-06 (3.9e-02)	1.5e-02 (1.7e-02)	6.1e-03 (2.7e-02)	10.58 $\rightarrow$ 10.57
mol_metals	1.6e-05 (2.3e-02)	1.1e-02 (1.4e-01)	1.3e-06 (4.0e-02)	1.5e-02 (1.7e-02)	5.4e-03 (2.7e-02)	10.58 $\rightarrow$ 10.58
mol_lyman_werner	6.3e-05 (2.9e-02)	1.5e-02 (8.4e-01)	7.8e-06 (4.8e-02)	1.8e-02 (2.0e-02)	7.8e-03 (3.3e-02)	10.58 $\rightarrow$ 10.58
mol_diffusion	2.2e-06 (2.2e-02)	1.1e-02 (1.9e+00)	4.3e-06 (3.9e-02)	1.5e-02 (1.7e-02)	6.1e-03 (2.7e-02)	10.58 $\rightarrow$ 10.57
lower_profile	3.2e-06 (2.8e-03)	2.5e-02 (2.6e+00)	4.6e-06 (3.5e-03)	3.0e-02 (3.9e-02)	3.3e-03 (8.8e-02)	9.58 $\rightarrow$ 9.58

wasp_he23off is byte-identical, which is the control: with the metastable off the coefficient is never evaluated. The large pointwise v ratios are cells where the golden velocity passes through zero, not large absolute changes — the scale-relative column is the one to read. The H_2 front does not move ($1.0181 R_p$ before and after in all four molecular cases), and $\log_{10} \dot{M}$ moves by at most 0.01 dex.

The species that moves is the one the change is about: peak $n(\text{He}2^3S)$ rises by 3.1% (wasp_full), 6.9% (lower_profile) and 22–30% (the molecular cases). Peak HeH^+ rises by a factor 116–710 in the molecular

cases ($2.6 \times 10^{-1} \rightarrow 3.4 \times 10^1 \text{ cm}^{-3}$ in `mol_base_handoff`), which is the newly resolved associative branch — previously the only HeH^+ source in those runs was $\text{He}^+ + \text{H}_2$, and this one is larger.

Goldens are NOT refreshed here. The matrix is measured, not re-snapshotted; the refresh belongs at the end of the change series. make `diffusion_tests 35/35` and `run_fcheck.sh CLEAN` with the change in.

The LHS 1140 b numbers move, and the helium-abundance crossings move with them. Re-solving the three reference winds (the flux-closure case with the wind re-converged on a fixed profile, no closure iteration) gives:

run	EW before	EW after	FWHM before	after	$\log_{10} \dot{M}$
heh0p55	1.10884	1.30991 (+18.1%)	0.2573	0.2583	7.76, unchanged
heh2p13_diff_kzz1e9	1.12887	1.33805 (+18.5%)	0.2553	0.2568	7.80, unchanged
flux_closure/heh11p1/k01	1.08532	1.16816 (+7.6%)	0.2825	0.2840	7.40, unchanged

(red-pair equivalent width in % $\text{\AA}$).

The mass-loss rate is untouched: this is a trace-level population, not an energy-budget term. The crossings with the measured equivalent width, solved on re-computed composition ladders rather than from the local slope (the +18% shift is well outside the slope’s range of validity), move down:

ladder	before	after	1σ
well mixed, diffusion off	0.55	0.424	0.409–0.440
$K_{zz} = 10^9$ (adopted)	2.09	1.626	1.573–1.685
flux closure	11.73	9.317	8.792–10.050

The local-slope estimate (slopes 0.709, 0.921, 0.382) gives 0.435, 1.738 and 9.676 for comparison; the direct ladder solve is the number to use. Two caveats stand on the flux-closure entry: 9.317 is an interpolation across existing arms, not a re-run closure, and confirming it means solving a closure arm at a reservoir near 9.3, where the cold-trap composition differs arm by arm. The stored He/H grid runs elsewhere in LHS1140b/exhale are still on the old coefficient and were not re-solved.

`docs/lhs1140b_exhale_vs_pwinds.tex` is *not* edited here. Every quantitative statement in it that depends on $n(2^3S)$ is now stale — the three crossings above, the Penning share of the metastable sink (89%/60% at $\text{He}/\text{H} = 2.09/10.31$) and the lifetimes in its Table 4, its EW table, and the two places that name the rate as “the Taylor et al. (2025) fit”. Its figures that show the He 10830 profile or $n(2^3S)$ need regenerating, and `figures/lhs1140b_ew_vs_heh.pdf` in particular carries the crossing.

Files. `src/modules/radiation/Cool_coeff.f90` (both rate functions renamed and replaced, `f_penning_HeI23S` added); `ion_residual_core.f90` (Penning branch on the H^+ row, total kept in the triplet row); `System_HeH_mol.f90` (branch on the H^+ and H_2^+ rows, associative branch into the HeH^+ row, total on the H_2 and triplet sinks); `System_implicit_adv_HeH_TR.f90`; `util_ion_eq.f90`; `ionization_equilibrium.f90`; `post_process_adv.f90` (heating terms). `docs/molecular_chemistry_audit_he_rich.md` section 4.4 updated, including the withdrawal of its claim that the step entered the Jacobian — the coefficient is constant within a cell, so it appeared only as a jump in the loss rate between neighbouring cells. `references/garcia_munoz_2025_network/` holds the published network file the branched amplitudes come from.

88. A population solved as a difference, and a post-process that never read info (2026-08-29)

The judgement first. Two defects in the advection-corrected post-process, both numerical, both fixed here.

1. **The ground-singlet helium population was solved as a difference.** `System_implicit_adv_HeH_TR` carried the *summed* He I fraction as x_2 and the metastable as x_4 , and formed the singlet as $x_2 - x_4$. Where helium is heavily ionized and the little neutral helium left sits mostly in the metastable, that difference loses every significant digit and finally evaluates to exactly zero. The summed He I row then degenerates — it no longer contains the unknown it is supposed to determine — and `hybrd1` returns `info = 4` with a residual it cannot reduce. Fixed by solving for the two populations, $n(1^1S)/n_{\text{He}}$ and $n(2^3S)/n_{\text{He}}$, and forming the *sum* where the total is wanted. A sum has no such failure mode.
2. **No return code of the post-process was ever read.** `post_process_adv` calls `hybrd1` three times per cell — the advection ionization system, the metal stage re-solve (`pp_metals = 2`), the energy equation — and wrote the returned iterate out in every case, converged or not. Fixed: a cell whose solve did not converge keeps the equilibrium state it came in with, and the number of such cell solves is reported.

The run that exposed both, `LHS1140b/exhale/heh0p55_diff_ctrl`, is **not made physical by either fix**, and the reason is a third thing, left untouched here and written up in § 88.4: its post-process drives P_{HeI} to $7 \times 10^{15} \text{ s}^{-1}$.

88.1 A population, not a difference

The advection systems integrate the ionization ODE upwind across one cell. The triplet system solved for $x_1 = n_{\text{HI}}/n_{\text{H}}$, $x_2 = n_{\text{HeI}}/n_{\text{He}}$ (summed), $x_3 = n_{\text{HeII}}/n_{\text{He}}$, $x_4 = n(2^3S)/n_{\text{He}}$, and built the singlet, which every rate in rows 2 and 4 needs, as $x_2 - x_4$. Nothing keeps the two apart. In the helium-settled control run the difference collapses pass by pass until x_2 and x_4 are the same double, at which point the singlet is exactly zero, row 2 reduces to a balance between the metastable and He II alone, and the solver stalls. Measured on that run, before the change: **310 of 4660 advection cell solves returned info = 4** (ten post-process passes, `OMP_NUM_THREADS=1`); the other three runs checked returned `info = 1` everywhere.

The unknowns are now the two populations themselves, and row 2 is the ground-singlet balance — the summed He I row minus the metastable row of row 4, which is the same system written in different variables:

$$f_2 = x_{s,\text{old}} - x_s + c_1 \left[x_{\text{HeII}} \alpha_{\text{HeII}} n_e - x_s (P_{\text{HeI}} + (b_{\text{HeI}} + q_{13}) n_e) + x_t ((q_{31a} + q_{31b}) n_e + A_{31} + x_{\text{HI}} Q_{31} n_{\text{H}}) \right]. \quad (36)$$

Loss of the singlet: photoionization, electron-impact ionization, and collisional excitation q_{13} into the metastable. Gain: recombination of He II into the singlet ladder (α_{HeII} ; the α_{2^3S} channel feeds the metastable and is charged to row 4), and every route by which the metastable comes back to the ground state — collisional de-excitation $q_{31a} + q_{31b}$, the $2^3S \rightarrow 1^1S$ decay A_{31} , and the $\text{He}(2^3S) + \text{H}^0$ ionizing collisions Q_{31} , whose Penning branch leaves $\text{He}(1^1S) + \text{H}^+ + e^-$ and whose associative branch makes HeH^+ that dissociatively recombines back to ground-state helium. The He–H charge-exchange pair stays where it was, on rows 1 and 2, with the *summed* He I as the reactant density: that is the accounting

the summed-He I form carried and the one the equilibrium systems make, neither of which gives the metastable a charge-exchange channel of its own.

post_process_adv now carries nheiS and nheiTR as the two profiles it solves for and forms $nhei = nheiS + nheiTR$ wherever a routine wants the total. The difference survives in exactly one place, the entry point, where the equilibrium solution hands over a summed He I and a metastable; it is well conditioned there because the equilibrium metastable is bounded by its own balance, whose loss side is dominated by A_{31} .

What the change does to the answer. The reparameterization is a change of variables, so a converged solution must not move. Measured, PP-only re-runs at OMP_NUM_THREADS=1 against a baseline built from the same tree without the change:

run	info $\neq 1$ solves, before $\rightarrow$ after	max relative change of the _adv species
heh2p13_diff_kzz1e9 ($K_{zz} = 10^9$)	0 $\rightarrow$ 0	2.6×10^{-9} (He II), median 3×10^{-14}
flux_closure/heh11p1/k01	0 $\rightarrow$ 0	2.0×10^{-9} (He II), median 3×10^{-15}
heh0p55 (no diffusion)	0 $\rightarrow$ 0	9.7×10^{-4} (He III), median 1.3×10^{-4}
heh0p55_diff_ctrl ($K_{zz} = 0$)	310 $\rightarrow$ 0	0.25 (He I and 2^3S , at $2.7 R_p$)

The first two are roundoff: the solve is the same solve. The heh0p55 entry is *not* the reparameterization — it is the new info guard of § 88.2 firing on one cell of the energy equation in passes 5 and 7. The last is the run whose solution was not a solution.

Where the cancellation went. In the control run at $r = 2.53 R_p$, tenth pass, the same cell before and after:

	before	after
unknowns	$x_2 = x_4 = 1.089669 \times 10^{-5}$ (singlet exactly 0)	$x_s = 2.53 \times 10^{-22}$, $x_t = 1.098 \times 10^{-5}$ ($x_s/x_t = 2.3 \times 10^{-17}$)
info	4	1
residual	($-9.5e-4$, $1.22e-2$, $7.7e-6$, $-3.5e-4$)	($4.3e-17$, $2.4e-8$, $-3.3e-18$, $9.6e-18$)

The largest residual falls from 1.2×10^{-2} to 2.4×10^{-8} , and the singlet keeps its relative accuracy seventeen decades below the metastable, where the difference form had nothing left. One-cell steps in the _adv metastable profile of that run, counted with the local-trend detector of § 87, go from **9 flagged cells to 1** — and that one, at $r = 29.05$, is the outer-boundary feature the equilibrium profile carries too.

Its neighbours are the more instructive record: cells 334 and 335 returned info = 1 with $f_2 = 1.17 \times 10^{-2}$ and 1.31×10^{-2} . hybrd1's info = 1 is a test on the *increment* between iterates, not on the residual, so a solver frozen against a degenerate row reports success. **The info check of § 88.2 would not have caught those cells.** Only the reparameterization does.

88.2 The return code

post_process_adv made three hybrd1 calls and read the return code of none of them:

call	what it solves	on failure, before
<code>adv_implicit_H / _HeH / _HeH_TR</code>	the advection ionization system of the cell	the <code>iterate</code> was written to <code>nhi..nheiTR</code> and fed the upwind cascade of the next cell
<code>ion_system_metals_pp</code>	the metal stage split at fixed H/He (<code>pp_metals = 2</code>)	the <code>iterate</code> replaced the equilibrium split
<code>T_equation</code>	the advected energy equation (legacy MINPACK path)	the <code>iterate</code> became $T_{\text{out}}(j)$, subject only to the metal-mode band reject

A non-converged `iterate` satisfies neither the advection balance it was asked to solve nor the equilibrium balance it started from, so it is not a state of the gas. The equilibrium solution of that same cell is, and it is already what the three validity conditions of § 72 fall back to. Each call now keeps the equilibrium state on `info ≠ 1` — the ionization of the cell, the equilibrium metal split, the converged equilibrium temperature — and counts it. The counts are summed over the ten passes, not read off the last one: a cell reverted in an early pass feeds that pass’s upwind cascade whether or not the last pass converges everywhere.

Measured on four LHS 1140b runs, the guard is quiet: the advection system fires nowhere once § 88.1 is in, and the energy equation on one cell in two of the ten passes of `heh0p55` and of the He 2³S-off variant of `heh2p13_diff_kzz1e9` — worth 10^{-4} to 10^{-3} in the `_adv` profiles of those runs. The metal guard is untested by them: `pp_metals` defaults to the frozen mode (1), so the re-solve block is one those runs never enter.

The guard accepts `info = 1` and rejects everything else. That is a statement about the increment, not about the residual, and § 88.1 shows the two can disagree. A residual test would be the stronger guard; it needs a threshold in the units of the rows (a fraction per cell crossing), which is a choice, not a measurement, and it is not made here.

A third, smaller thing in the same block. With helium off (`thereis_He = .false.`) the post-process left `nhei`, `nheii`, `nheiii`, `nheiTR` — and, with this change, `nheiS` — undefined until the helium-free branch of the ionization loop zeroes them at its *end*. The first pass reads them before that, in `nhe`, in the electron sum and in `eval_cool`. They are now zeroed where the helium arrays are initialized. Helium-on runs are bit-identical across the change (checked on `heh2p13_diff_kzz1e9`, all three `_adv` files).

88.3 Why nothing caught this

A run’s convergence tests never look at the post-process.

- *du* and the steady residual $\|R\|$ are momentum and energy metrics of the *wind*, evaluated in the time loop. The loop is over before `post_process_adv` is called; there is no metric downstream of it.
- `ionization_equilibrium` validates the root it accepts — physical simplex, up to three starting points, `ok_rank`, and a one-line sweep summary of reseeds, retries, unphysical roots and failures (§ 45). The advection-corrected post-process, solving a system of the same size with the same solver, validated nothing.
- `make check` bit-compares `Hydro_ioniz.txt` and `Ion_species.txt`. Those are the *equilibrium* files, written before `post_process_adv` runs. **No golden case compares an `_adv` file at all**, so nothing in the regression can

move when the post-process does — which is also why the change documented here leaves the goldens byte-identical (7/7, 14/14 files), and why that result carries no information about the `_adv` profiles. The A/B tables above are what stands in for it.

That leaves the `_adv` files checked only by whoever plots them. The two defects were found that way: an order-of-magnitude one-cell step in a metastable profile.

88.4 What the control run actually shows, and what is not fixed here

With § 88.1 in place the solver converges everywhere in `heh0p55_diff_ctrl` and still returns a neutral-helium column five orders of magnitude below the equilibrium one. The singlet is not being lost to cancellation any more; it is being destroyed by the rate the post-process hands the solver. Traced pass by pass at $r = 2.53 R_p$:

pass	1	2	3	5	10
$P_{\text{HeI}} [\text{s}^{-1}]$	2.5e-4	3.7e-2	5.3e0	1.1e5	6.9×10^{15}
$x(1^1S)$	6.8e-3	4.7e-5	3.3e-7	1.6e-11	2.5e-22

A photoionization rate of $7 \times 10^{15} \text{ s}^{-1}$ is not a rate. Two things produce it, and neither is touched here.

(a) A volumetric secondary-ionization rate divided by the density of the species it ionizes. `util_ion_eq.f90` adds the Shull & van Steenberg (1985) secondary channel as $P_{\text{HeI}}(j) = P_{\text{HeI}}(j) + R_{\text{secHeI}}/\max(n_{\text{HeI}}(j), 1.0\text{d-99})$. R_{secHeI} is the energy deposited per unit volume times the SvS85 branching into He I ionization: a *volumetric* rate, and one carrying SvS85's own fixed helium abundance. Dividing it by the actual $n(1^1S)$ gives the right per-atom rate only at that abundance, and diverges as the singlet disappears. In this run helium has settled to $n_{\text{He}}/n_{\text{H}} = 6 \times 10^{-6}$ at $1.36 R_p$, four decades below the SvS85 composition, and the ten post-process passes close the loop: a smaller singlet raises its own destruction rate, which shrinks it further. The $\max(\dots, 1\text{e-99})$ is what stops the divergence, at a rate of order 10^{60} . The physical scaling is the opposite one — the secondary He I ionization rate per unit volume should fall with the helium abundance, roughly as $n_{\text{HeI}}\sigma_{\text{HeI}}/\sum_a n_a\sigma_a$, leaving the per-atom rate near its solar-composition value.

(b) In a Do only PP: True run the equilibrium solve and the post-process use different physics. The SvS85 coupling is staged: `sec_ion_active` starts false and is switched on once the wind has first converged (§ 38). `EXHALE_main.f90` also sets it unconditionally *after* the time loop, so the post-process always runs with the full coupling. With `Do only PP` the loop exits after one iteration, so the single `ionization_equilibrium` call runs with `sec_on = F` and all twenty-one `PH_heat_HHe` calls of the post-process run with `sec_on = T` (printed and counted). The equilibrium P_{HeI} at $1.36 R_p$ of the control run is $2.4 \times 10^{-7} \text{ s}^{-1}$; the post-process's first pass, at the same densities, $4.3 \times 10^{-5} \text{ s}^{-1}$. The `_adv` files of such a run are not a correction to the equilibrium files next to them: they are a different model.

(c) The same difference, in the radiation. `util_ion_eq.f90` builds its own singlet as $n_{\text{HeI}} = n_{\text{HeI}} - n_{\text{HeITR}}$ and uses it for the He I column density, the He I photoionization integrand and the denominator in (a). The state vector of the whole code carries (summed He I, metastable) and reconstructs the singlet wherever it is needed — `species_table.f90` says so in as many words. § 88.1 fixes the one system where the difference was catastrophic;

carrying $n(1^1S)$ as the state variable throughout would touch the global species vector, the output schema and every reader of it, and is not attempted here.

Six of the 354 stored LHS1140b runs carry the signature of (a) — $n_{\text{HeI}} \leq n(2^3S)$ in the `_adv` file above $1.5R_p$: `heh0p55_diff_ctrl`, `heh0p55_diff_d3`, `heh0p55_diff_fix` (171 cells of 233 each), `heh0p55_diff_kzz1e6` (36), `xuv0p01_heh2p13` (25) and `xuv0p10_heh2p13` (28). All six are helium-settled: $K_{zz} \leq 10^6$, or a reduced XUV flux. The $K_{zz} = 0$ arm of Fig. `lhs1140b_kzz_profiles.pdf` and its quoted peak $n(2^3S) = 0.046 \text{ cm}^{-3}$ are read from the first of them.

Files. `src/modules/nonlinear_system_solver/System_implicit_adv_HeH_TR.f90` (the two populations as unknowns, the ground-singlet row); `System_implicit_adv_HeH.f90` (the same naming, no arithmetic change: without the metastable all He I is the singlet); `ion_cell_state.f90` (`xhei_old` $\rightarrow$ `xheiS_old`); `src/modules/post_process/post_process.f90` (`nheiS` and `nheiTR` as the profiles solved for, `nhei` formed as their sum, the three info guards, their counters and reports, and `pin_cell_to_equilibrium`, which the validity conditions and the guards share). `docs/postprocess_advection_validity.f90` carries the same items next to the 2026-08-12 pair. `docs/lhs1140b_exhale_vs_pwinds.tex` is *not* edited.

Found in passing and fixed, outside the post-process: `LHS1140b/make_memo_figures.py` and `LHS1140b/exhale/kzz_scan_table.py` formed the elemental helium density as `HeI + HeII + HeIII + HeITR`, counting the metastable twice — it is a level of He I and is already inside the He I column, which `species_table.f90` states in capitals. Both now sum the three ion stages. Measured on the four K_{zz} arms of the memo figure, the quoted elemental ratios move by at most 2×10^{-6} relative ($0.318246 \rightarrow 0.318244$ at $5R_p$, $K_{zz} = 10^{10}$), so nothing was regenerated.

89. A conservation check that could not see conservation: the deepest-level elemental test and the reservoir band it refused (2026-08-29)

Read with § 90. The diagnosis below stands, but its two decisions were superseded the same day. Photochem was corrected at the source, so the `--abundance-tol` default went $10^{-3} \rightarrow 10^{-10}$ and the three `clima` failures noted at the end were fixed (two others appeared).

The judgement first. The profile adapter's deepest-level elemental check was refusing converged LHS 1140 b columns on **the residual of Photochem's chemical-equilibrium solver**, not on any error of the photochemistry, the transport or this code's carrier summation. At the level the check is stated, Photochem pins every gas species of the handed-over mixture with `bc_type='press'` (`gasgiants._initialize_atmosphere`), so that level carries its equilibrium initialization unchanged and **no photochemical or transport leak can reach the check at all**. What it reads instead is how far `composition_at_metallicity` got, and that routine states `# Do not enforce convergence`. The `--abundance-tol` default is raised from 10^{-4} to 10^{-3} — an interim value for the solver as it then was; § 90 tightens it to 10^{-10} once the solver is fixed. Full diagnosis, with every measurement: `docs/deep_level_elemental_check.md`.

89.1 The symptom

Reservoir He/H = 8.5, 8.55, 8.6, 8.65 and 8.7 could not be run in the closure ladder of `LHS1140b/exhale/crossings_gm25/`: the N/H ratio at the deepest level departed from the input by 1.02, 1.22, 1.43, 1.67 and 1.94×10^{-4} against the 10^{-4}

tolerance. Two closure lineages with different trial escape fluxes give the same numbers to the three digits the refusal prints, and a standalone chemistry-only scan reproduces them a third time — as expected, since the trial flux enters only at the upper boundary and the refused level is Dirichlet.

89.2 The level is a Dirichlet boundary, and the solver behind it is used unconverged

Measured at $\text{He/H} = 8.6$: chemical equilibrium plus the quench overwrite, the photochemical model's initial state, and the converged steady state agree at the deepest level to **all 13 digits printed** ($X_{\text{N}} = 8.185823483630 \times 10^{-5}$ in all three; deep $q_{\text{NH}_3} = 8.995632484 \times 10^{-6}$ converged against 8.995632×10^{-6} initial). Called on its own at the same P and T , `equilibrate.ChemEquiAnalysis.solve` returns `converged = True` at every reservoir swept while its own `molfracs_atoms_gas` misses the requested N/H by up to 1.72×10^{-4} .

89.3 Why nitrogen and nothing else

In all 38 runs of the scan $\text{dev}_{\text{N}}/\text{dev}_{\text{He}} = 4072$ (4069–4086 over the whole set, 4072 in 32 of them), against $1/(3\text{N/H}) = 1/(3 \times 8.18465 \times 10^{-5}) = 4073$. Deep nitrogen is essentially all NH_3 — on the $\text{He/H} = 8.6$ profile the N_2 nuclei share of X_{N} is 1.0×10^{-5} and HCN 's is 2.2×10^{-25} — so **one excess N rides on three excess H**: He, C and O carry the common hydrogen error alone, N carries it plus the relative error of the equilibrium NH_3 mole fraction. The check was testing that one mole fraction 4072 times more tightly than anything else, which nobody chose.

89.4 Not a branch change, and not a special band or planet

Across $8.45 \rightarrow 8.75$ the deep H_2 , H_2O , CH_4 , NH_3 , N_2 and He all move monotonically by less than 1 % with no jump; the residual jumps by a factor 38 between 8.71 (1.99×10^{-4}) and 8.72 (5.2×10^{-6}), 1500 between 9.15 and 9.20, and 5800 between 10.25 and 10.30. At the old tolerance the same scan also refuses $\text{He/H} = 6.0$ (1.31×10^{-4}) and 10.20–10.25 ($1.5\text{--}2.0 \times 10^{-4}$); the tree already held refusals at $\text{He/H} = 14.643$, 15.090 (2.12×10^{-4} , the largest recorded) and 25.0 (LHS1140b/exhale/kzz_profile_scan/); and at **solar composition** the same solver misses N/H by 1.50×10^{-4} at 800 K and 1.14×10^{-4} at 1600 K. The two HD 209458 b departures that fixed the original default (1.6×10^{-7} , 2.7×10^{-5} ; § 77) were a fortunate sample, not a floor.

89.5 No accounting error

Recomputed from the written profile, the diagnostic columns account for **100.00 %** of X_{N} at the deepest level. They account for 93.9 % at the 1 μbar match, and any earlier reading of a shortfall in the carrier sum refers to *that* level and to the deliberately short `q_*` diagnostic column set — the check reads neither.

89.6 The change

`src/utls/lower_profile_schema.py`: the `--abundance-tol` default goes $10^{-4} \rightarrow 10^{-3}$, and both the option help and the comment at the check now state what the level can and cannot test. 10^{-3} leaves a factor 4.7 over the worst departure on record and gives up next to no detection power: each element sits in one dominant carrier there, so a real

miscount is an $O(1)$ error, and the smallest conceivable one — dropping N_2 — is 1.0×10^{-5} of X_N in the band at issue and was already invisible at 10^{-4} . Over the 38 profiles the N_2 nuclei share runs 4.2×10^{-6} ($\text{He}/\text{H} = 12$, deep $T = 379$ K) to 4.0×10^{-3} ($\text{He}/\text{H} = 1$, deep $T = 588$ K), rising with the deep temperature, so where dropping it would be a real loss it stays above 10^{-3} . Nothing else was changed: not the check, not its location, not the refusal policy. **Superseded by § 90:** the headroom argument was a statement about how far the 0.8.4 equilibrium solver missed, and once that was fixed the default went to 10^{-10} .

89.7 What it means for the LHS 1140 b crossings

Against `crossings_gm25/results.txt`: the flux-closure crossing at $\text{He}/\text{H} = 9.048$ is untouched — the 9.0 and 9.1 arms both ran and bracket it — and so is the -1σ end 10.77, bracketed by 10.6 and 11.0. The $+1\sigma$ **end 8.49 was interpolated across the un-runnable band**, a 3.4 % gap in reservoir between the 8.45 and 8.75 arms; with the tolerance raised those five rungs can be run and that end replaced by measurement. That was done, on the corrected build: § 90. Because the refusals at 6.0, 10.20–10.25, 14.643, 15.090 and 25.0 are the same trap at other reservoir values, and nothing in the pattern predicts where the next one falls, any future ladder can meet it anywhere.

89.8 Seen in passing, not investigated

`clima`'s `surface_temperature_bg_gas` root solve fails outright at three of the 41 reservoir values attempted — $\text{He}/\text{H} = 0.0969$ (with Failed to compute heat capacity), 8.1 and 15 — so those compositions produce no column at all. That is a **different** failure from the refusal above (this one happens before a column exists, the refusal after), and it is recorded as item (M) of `TO_BE_DONE.md`. All three are fixed by the bounded, scaled climate solves of § 90; the corrected build refuses two other compositions instead, $\text{He}/\text{H} = 9.4$ and 9.5, and § 91 shows that is the same underlying solver defect moving address rather than a new one — 0.8.4 refuses two compositions of its own on the same grid.

Files. `src/utis/lower_profile_schema.py`; `docs/deep_level_elemental_check.md` (the diagnosis); `LHS1140b/exhale/nh_re` (the measurement record: README, four scripts, five tables, 38 adapter runs under scan/). No EXHALE run was made for any of it and no golden was touched; the adapter's chemistry is unchanged and `make check` does not exercise this path.

90. The equilibrium residual removed at its source, and the observable that did not move (2026-08-29)

The judgement first. The residual § 89 diagnosed was fixed where it came from — in Photochem's chemical-equilibrium path, not by relaxing this code's check — and **it does not move the observable**. Re-measured end to end on the corrected build with 11 arms, the LHS 1140 b flux-closure crossing of the He I 10830 equivalent width is at $\text{He}/\text{H} = 9.0484$ (chord) / **9.0461** (quadratic), the same to every digit the photochem 0.8.4 ladder printed; on seven like-for-like arms the equivalent width moves by at most 1.2×10^{-6} in relative terms and $\log_{10} \dot{M}$ is unchanged in its fourth decimal. The one EXHALE change is that `--abundance-tol` goes back to its design value, $10^{-3} \rightarrow 10^{-10}$. Records: `LHS1140b/exhale/crossings_pc090/results.txt`, `docs/photochem_solver_modification_implementation.md`.

90.1 What was changed, and where

Nothing in the EXHALE Fortran. The corrected chemistry is a modified Photochem v0.9.0 source tree, built as photochem 0.9.0 into a dedicated conda environment:

- **Equilibrate** now tests each positive requested elemental abundance against its own relative tolerance ($\text{abs}(b_0(i_{\text{atom}})) \cdot \text{mass_tol}$) instead of a common absolute threshold scaled by $\text{MAXVAL}(b_0)$, and no longer excludes elements at or below 10^{-6} . This is the actual cause of § 89.
- **gasgiants.py** exposes `equilibrium_mass_tol` (default 10^{-12}), assigns it to `ChemEquiAnalysis.mass_tol`, verifies the returned elemental vector against the requested one, raises rather than silently keeping the last result when all five temperature retries fail, and restores `use_prev_guess` in a `finally` block.
- **clima** v0.7.5 gets a bracketed background-pressure solve, a $\log_{10}(T_{\text{surface}} - T_{\text{trop}})$, $\log_{10} T_{\text{trop}}$ temperature parameterization, and normalized residual acceptance, replacing unconstrained MINPACK solves.

Both dependency changes are reproducible patches applied through Photochem's CMake configuration. Diagnosis: `docs/photochem_solver_modification_investigation.md`.

90.2 A blocking defect, found only by running a column

The first version of the strict closure test **rejected every composition**, including ones that ran on 0.8.4. It compared the requested elemental vector with `gas.molfracs_atoms`, the total over gas and condensate, at every level. With `rainout_condensed_atoms=True` the vector handed to the next level is the gas-only one, and `gas.molfracs_atoms` started from the previous level through `use_prev_guess` still carries that level's condensate: measured on the LHS 1140b column at $\text{He}/\text{H} = 9$, the total reads 1.1×10^2 against the request where the same state solved from scratch closes to 5.7×10^{-14} . Identical on 0.8.4 and on the patched build, so it is Equilibrate's existing reporting behavior, not something the Fortran patch created.

The test now forms the residual **only at a level whose `molfracs_atoms_condensate` is entirely zero** — which is every level deep enough to set an elemental handoff, and where the residual it exists to catch appears. Why the defect survived the first round of tests is worth recording: those tests ran **one level for each of the 38 saved deep states** and never carried a column, so `use_prev_guess` chaining and condensing levels were never reached.

90.3 Measured

Re-solved standalone, the 38 saved deep states of LHS1140b/exhale/nh_refusal_diagnosis/scan/ give a largest relative elemental residual of $3.539391003 \times 10^{-13}$ and a median of $1.915134717 \times 10^{-14}$, against $2.018197773 \times 10^{-4}$ for the same measure on 0.8.4 — all 38 converged and independently accepted. Over the 68 handoff writes the corrected build made under LHS1140b/exhale/ (43 in `scan_pc090/`, 25 in `crossings_pc090/`), the largest deepest-level departure is 3.28×10^{-13} and the median 1.73×10^{-14} .

90.4 The tolerance

src/utls/lower_profile_schema.py: --abundance-tol default $10^{-3} \rightarrow 10^{-10}$, a factor 305 above the worst hand-off write measured. Detection power is untouched — each element sits in one dominant carrier at that level, so a genuine miscount is an $O(1)$ error, and the smallest conceivable one, dropping N_2 , is 10^{-5} to 4×10^{-3} of X_N . **A 0.8.4 handoff is refused at this default, and that is the intended signal:** over the 43 stored 0.8.4 handoffs of crossings_gm25/ the departures run 3.48×10^{-8} (best) to 8.50×10^{-5} (worst), so even the best of them misses 10^{-10} by a factor 348. Reproducing a stored 0.8.4 result needs an explicit --abundance-tol on the command line, which is how a deliberate reproduction states that it accepts that residual; the scans under nh_refusal_diagnosis/ already pass --abundance-tol 1.0 and are unaffected.

90.5 What it does to the LHS 1140 b crossings

quantity	0.8.4, 16 arms	corrected build, 11 arms
crossing, chord	9.0484	9.0484
crossing, quadratic	9.0461	9.0461
1σ band	8.4923–10.7716	8.4649–10.7357

The crossing does not move. The handoff profile does: at He/H = 9 the deepest-level N/H departure goes from 2.1×10^{-5} to 1.2×10^{-14} and trace species such as q_H2O, q_CO and q_HCN move by up to 10^{-2} relative — none of which reaches He I 10830, a line set by H, He and the temperature. **The 1σ band moves for a different reason than the solver:** its low end used to be interpolated across the un-runnable band He/H = 8.5–8.7, which now runs (departures $1.4\text{--}3.2 \times 10^{-14}$), so real arms at 8.5 and 8.6 carry that end and it goes $8.4923 \rightarrow 8.4649$; the high end goes $10.7716 \rightarrow 10.7357$ because the 10.6 arm was not seeded the same way in the two ladders, a 1.1×10^{-3} seed difference in equivalent width, the size of the seed spread the stored ladder already measured.

Also opened: solar composition. He/H = 0.0969 is runnable end to end for the first time. It was blocked not by the tolerance but by the climate step, which now solves (deep boundary 720.5 K, against 432.4 K at He/H = 8.1 and 393.1 K at 15 — the other two recovered compositions).

90.6 Two compositions the corrected build refuses, and what they turned out to be

He/H = 9.4 and 9.5 raise Could not bracket background pressure in make_profile_bg_gas on this build and produce columns on 0.8.4. This section first recorded that as a regression of the new bracketing scan, with the cause undiagnosed. **Section 91 measured it and that reading was wrong.** The scan is not the cause — the root lies inside the interval it searches, with one monotone sign change in the last scan step — and the same defect is present in 0.8.4, which fails at 9.37 and 9.45 on the same grid where this build fails at 9.40 and 9.50. Both builds are tripped by a forward-difference Jacobian formed at a step below the flux residual's round-off jitter. **The net accounting of “three failures removed, one introduced” written here does not hold:** it compared against 0.8.4 columns that had not been run at the matching reservoir values, and 9.45 in particular fails on 0.8.4. Neither value is a bracket arm of the crossing. Diagnosis, repair and status: § 91; item (M) of TO_BE_DONE.md.

90.7 Scope

No EXHALE Fortran source or binary changed, no golden was refreshed, and `make check` was not run: this path is not in the regression matrix. Not tested: condensate-rich and sulfur-network equilibrium sweeps — note that after the fix of § 90.2 a condensing level is accepted on the solver's own converged flag, so the closure test does not speak for such a level at all. The stored 0.8.4 ladder was not re-run, and $\text{He/H} = 9.5$ cannot be repeated on this build. The 0.8.4 environment is untouched and remains the reproduction path for every profile whose header reads `# source_version photochem 0.8.4`.

Files. `src/utis/lower_profile_schema.py` (the default); `docs/photochem_solver_modification_investigation.md` (diagnosis), `docs/photochem_solver_modification_implementation.md` (what was built and measured), `docs/deep_level_elemental_check.md` §§ 8–10 (the tolerance, the crossings, the climate solve); `LHS1140b/exhale/crossings_pc0` (the 11-arm re-measurement) and `LHS1140b/exhale/nh_refusal_diagnosis/scan_pc090/` (the 45-point reservoir scan on the corrected build). The Photochem changes live outside this repository, in the local `photochem/` tree.

91. The clima bracket was not the cause: a Jacobian built from round-off, and a failure that predates the patch (2026-08-29)

The judgement first. The two compositions § 90 reported as a regression are not one. `clima`'s new bracketing scan does not fail to find a root at $\text{He/H} = 9.4$ and 9.5 — the root is inside the scanned interval, and the residual crosses zero once, monotonically, inside the last scan step. The scan fails because the *outer* `surface_temperature_bg_gas` solve hands it a surface temperature of 6.5×10^8 K, where every scan point is thermodynamically invalid. That step comes from a forward-difference Jacobian whose difference step is smaller than the residual's own round-off jitter, and **the 0.8.4 build has the same defect at the same rate** — it fails at 9.37 and 9.45 where the corrected build fails at 9.40 and 9.50. The patch changed which compositions are unlucky, not how many. Record: `LHS1140b/exhale/clima_bracket_diagnosis/` (README, 12 scripts, six sweep and variant tables), measured 2026-08-29 against both installed builds with nothing in `photochem/` modified.

91.1 The bracket is not the cause

`make_profile_bg_gas` scans 49 points over $\log_{10} P$ spanning 12 decades below the requested surface pressure, so for the LHS 1140 b ladder's 20 bar deep boundary the interval is $\log_{10} P = [-4.6990, 7.3010]$. At $\text{He/H} = 9.4$ the 0.8.4 solution sits at $\log_{10} P_{\text{bg}} = 7.278533$, inside it. So does every neighbor: over 9.30–9.60 the 0.8.4 background pressure moves only from 7.278297 to 7.278990, always inside (`sweep_old_9p30_9p60.txt`). Inside one failing call the residual is monotone with exactly one sign change, in the last scan step (`scan_residual.py`) — a grid of 49 points cannot step over it.

What the scan actually reports is that it had **no valid point at all**. The trial surface temperature it is called with is 6.5×10^8 K (`outer_trace.py`), and at that temperature `make_profile` refuses all 49 points with `Failed to compute heat capacity: the thermodynamic polynomials stop at 6000 K (T_ceiling.py)`. With no valid point there is no sign

change to find, and the message `Could not bracket background pressure in make_profile_bg_gas` is the only thing the routine can say.

91.2 The Jacobian is round-off

`surface_temperature_bg_gas` solves the two-component system (flux closure, skin temperature) with MINPACK `hybrd1`, which fixes `epsfcn = 0`; `fdjac1` then forms $\varepsilon = \sqrt{\max(\text{epsfcn}, \epsilon_{\text{mach}})}$ and steps $h = \varepsilon|x|$. At the ladder's start, $T_{\text{surf}} = 400$ K and $T_{\text{trop}} = 120$ K, that is 3.6×10^{-8} in the log variable, a temperature step of 2.4×10^{-5} K. Over that step the normalized flux residual moves by 3.8×10^{-8} , and its own jitter is $\sim 3 \times 10^{-8}$ (`jacobian_noise.py`, `noise_source.py`) — the difference is the noise, not the derivative.

The resulting entries have the wrong magnitude and, at these compositions, the wrong **sign**: at $\text{He}/\text{H} = 9.4$ the cross term $\partial f_1 / \partial x_2$ comes out +52.1 where a resolved step gives −0.553 (`first_step.py`). Because that cross term is tens of times the diagonal $\partial f_1 / \partial x_1 = -1.21$, an error in the small, noisy diagonal is amplified into a first step of several decades in $\log_{10} T_{\text{surf}}$, which is how a 400 K guess becomes 6.5×10^8 K in one iteration.

91.3 The failure predates the patch; only its address changed

The same measurement on 0.8.4 (`first_step_old.py`) shows the identical mechanism at $\text{He}/\text{H} = 9.37$: its noisy Jacobian sends the first trial to $T_{\text{surf}} = 0.10$ K, which is exactly the T_{surf} is less than T_{trop} refusal that build reports there. The patch reparameterized the solve as $\log_{10}(T_{\text{surf}} - T_{\text{trop}})$, which makes that particular refusal unreachable, so the same bad step now surfaces at the other end as a bracket message. The failure rates are the same:

grid	0.9.0 + patch	0.8.4
He/H 9.30–9.60, step 0.01 (31 points)	9.40, 9.50	9.37, 9.45
He/H 8.00–11.00, step 0.05 (61 points)	9.40, 9.50, 10.45	8.10, 9.45, 9.85, 10.15

(`sweep_new_9p30_9p60.txt`, `sweep_old_9p30_9p60.txt`, `sweep_new_wide.txt`, `sweep_old_wide.txt`.) **The net accounting of “three failures removed, one introduced” that §§ 89, 90 and `deep_level_elemental_check.md` carried does not hold** — it counted the corrected build's failures against a 0.8.4 column that had not been run at the same reservoir values. In particular the 9.45 entry those tables record as not attempted on 0.8.4 is a **failure** on 0.8.4. Nothing about the composition marks a failing point. The solution is smooth across them: the deep temperature falls monotonically from 422.72470 K at $\text{He}/\text{H} = 9.30$ to 420.61128 K at 9.60, with the failures at 9.40 and 9.50 sitting on that trend. The failure is a discrete event of the solve, not a boundary in the physics.

91.4 The workaround, before the repair was installed

`--climate-t-deep-guess` (default 400.0, in `src/utils/photochem_to_lower_profile.py`) is the only user-facing control over where the outer solve starts. At both 9.4 and 9.5, every guess tried other than exactly 400 K — 300, 350, 380, 420, 450, 500, 600 — solves, and each returns the 0.8.4 answer to five decimals (`tguess_scan.py`). A composition that refused could therefore be recovered without rebuilding. The installed environment no longer refuses any composition on either grid (91.5), so the key is a diagnostic rather than a workaround.

91.5 The repair, tested in replication and then installed

`fix_variants.py` drives the patched residual through the same MINPACK `hybrd` from Python, reproducing the installed build's converged deep temperature to 5×10^{-5} K wherever it solves, and then varies one thing at a time over the 61-point grid:

variant	solved of 61	failed or stalled
as patched	54	9.40, 9.50, 9.80, 10.05, 10.25, 10.70, 10.80
initial trust-region factor 1	54	the same seven
trial T_{surf} confined to $(T_{\text{trop}}, 6000 \text{ K})$	55	8.50, 9.55, 9.80, 9.85, 10.20, 10.35
$\text{epsfcn} = 10^{-4}$	61	none

epsfcn is not the step itself: with $\varepsilon = \sqrt{\max(\text{epsfcn}, \epsilon_{\text{mach}})}$ and $h = \varepsilon|x|$, $\text{epsfcn} = 10^{-4}$ is a relative step of 10^{-2} in the log variable, 16.2 K at the starting point against 2.4×10^{-5} K at the default. The Jacobian it produces is smooth and nearly composition-independent, $[[-1.21, -0.55], [5.0 \times 10^{-4}, -3.557]]$ across the whole 9.30–9.55 range, against entries running from +52 to –42 at the default step. Applied to 0.8.4's own parameterization it also solves all 61 points (`old_epsfcn.py`), so the repair is independent of what the patch changed. The recovered temperatures are 422.00813 K (9.40), 421.30374 K (9.50) and 415.15428 K (10.45) — on the smooth trend, and equal to the 0.8.4 values at 9.40 and 9.50 to all printed digits.

Confining the trial temperature to the range of the thermodynamic data helps but does not fix the solve on its own — it converts one failure mode into another, because the tropopause variable stays unbounded.

Installed, and measured there. In `clima` the repair is `hybrd` in place of `hybrd1` in `AdiabaticClimate_simple_solver`, the routine both `surface_temperature_bg_gas` and `surface_temperature` drive, so that epsfcn can be set at all; $\text{epsfcn} = 10^{-8}$, `factor = 100`, every other setting `hybrd1`'s. 10^{-8} and not the 10^{-4} of the table above: a relative step of 10^{-4} in the log variable, 0.16 K at the start, is about 10^4 above the jitter and still reproduces a converged central difference to 0.12 percent on the flux row, where 10^{-4} is a 16 K step whose flux row is 9 percent off the derivative. Both solve the grid; 10^{-8} is the middle of the window in which neither error source is active. The reasoning is in the patch at the constant's declaration.

The installed library carries it. Read out of it — `AdiabaticClimate_simple_solver` disassembled, the stack arguments of its MINPACK call followed to the constants they point at — it calls `__minpack_module_MOD_hybrd` with $\text{epsfcn} = 10^{-8}$ and `factor = 100.0` (LHS1140b/exhale/clima-epsfcn-installed/installed-epsfcn.py). Measured on that build, not in replication:

grid	installed	as patched, before
He/H 8.00–11.00, step 0.05 (61 points)	61 / 61	54 / 61
He/H 9.30–9.60, step 0.01 (31 points)	31 / 31	29 / 31

Every composition either build refused now solves — 8.10, 9.37, 9.40, 9.45, 9.50, 9.80, 9.85, 10.05, 10.15, 10.25, 10.45, 10.70, 10.80 — and the temperatures the replication predicted are what the installed build returns: 422.00813 K at 9.40, 421.30375 at 9.50 and 415.15429 at 10.45, against 422.00813, 421.30374 and 415.15428 above, agreeing to the last printed digit. The deep temperature is monotone in composition across the whole grid, 433.26144 K at 8.00 falling to 411.97203 K at 11.00.

What it moves. Nothing in the observable. The two arms of the LHS 1140 b elemental-flux-closure ladder that bracket the He I 10830 crossing were re-run end to end on the corrected environment against the stored arms of LHS1140b/exhale/crossings_pc090, which are the same photochem 0.9.0 source without the repair — same seed, same parent, same binary, same transit configuration. Both closure iterations converge at $k = 1$ on residuals equal to the stored ones to all five decimals, and the equivalent width moves by 5×10^{-8} in relative terms:

arm	EW without the repair	with it	relative	$\log_{10} \dot{M}$
9.0	1.1068538680	1.1068539227	$+4.94 \times 10^{-8}$	$7.39 \rightarrow 7.39$
9.1	1.1092155021	1.1092155563	$+4.88 \times 10^{-8}$	$7.39 \rightarrow 7.39$

The crossing moves from 9.0484199298 to 9.0484176256 (chord) and from 9.0460519944 to 9.0460494223 (quadratic), 2.3×10^{-6} and 2.6×10^{-6} , so the 9.0484 and 9.0461 of record stand. Every quality indicator of that ladder is unchanged digit for digit. The climate solve itself moves a composition that already solved by 2.1×10^{-6} K at He/H = 9.0 (424.9534068835 against 424.9534089414 K). Record: LHS1140b/exhale/clima_epsfcn_installed/results.txt.

The source tree carries the repair; a clone outside this repository does not. photochem/src/dependencies/patches/clima-bound holds `real(dp)`, `parameter :: epsfcn = 1.0e-8_dp (md5 3c3c40cc)`, which is also what `src/utls/photochem_exhale.patch` writes, and `src/utls/setup_photochem.sh` checks that line by name and passes. The tree reproduces the installed library. What is a revision behind is the separate clone at the workspace level, `ExoAtmosphere/photochem/ (md5 6734053e)`, where the modification was first developed: it still writes `hybrd1` and has no `epsfcn`. That clone is outside this repository and nothing here reads it.

91.6 What is not established

The $\sim 4 \times 10^{-7}$ relative jitter in ISR – OLR was not traced to a particular part of the radiative transfer. It is not the inner pressure solve: the background pressure returned is bit-identical across the steps that produce the difference (`noise_source.py`). Whether the same trap is reachable at other planets, other K_{zz} or other deep pressures was not tested — only LHS 1140 b's configuration was scanned. No EXHALE Fortran source, binary or golden was touched, and `make check` does not exercise this path.

Files. LHS1140b/exhale/clima_bracket_diagnosis/ (the diagnosis); LHS1140b/exhale/clima_epsfcn_installed/ (the repair measured on the installed environment); `src/utls/photochem_exhale.patch` (the current Clima patch, `epsfcn` included); `src/utls/photochem_to_lower_profile.py` (`--climate-t-deep-guess`); item (M) of `TO_BE_DONE.md`.

92. A column that came back 2.4 % apart in pressure: the grid the solver stops on, not the temperature it started from (2026-08-29)

The judgement first. The handoff-profile sensitivity left open at the end of § 91 — a deep boundary temperature 2.1×10^{-6} K apart giving a column 2.4 % apart in pressure, 15 % in one trace species and 0.5 % in the elemental O/H — is **the extent of Photochem's altitude grid at the instant the run is declared steady, and nothing**

else. It is numerical, it is not an amplification of 2.1×10^{-6} K, and it is not the chemistry's convergence tolerance. **It does not limit the closure chain:** the base state EXHALE reads at the matching level reproduces to 6.5×10^{-8} in the He/H the closure iterates on, against `element_flux_closure.py`'s `tol = 0.05`. One defect was found on the way and corrected: `insert_level` broke $n = p/(k_B T)$ at exactly one level of the table, the matching level. Full diagnosis, with every measurement: `docs/lower_profile_deep_boundary_sensitivity.md`; record `LHS1140b/exhale/deep_temperature_sensitivity/`.

Reproduced inside one build. The initial guess handed to the climate root solve is not physics — it is where MINPACK starts — so ten guesses over 300–450 K are ten solves of the same problem. They return the deep boundary temperature over a range of 2.121×10^{-6} K (424.95340688 to 424.95340876), the tropopause to 2.5×10^{-8} relative and the water above it to 6.6×10^{-8} ; re-run in a fresh process the ladder is bit-identical. Seven full adapter runs at seven of those guesses reproduce every number the record of § 91 left open: the pressure at fixed level index spreads by 2.470×10^{-2} , `q_OH` by 0.58, and the oxygen handed over at the matching level by 5.207×10^{-3} . `g400a` and `g400b` are the same command run twice and write byte-identical files, so none of this is run-to-run scatter.

Everything is two-valued, and the label is the model top. Over all eleven runs the profile top is either 1.0934×10^{-8} bar or 1.12075×10^{-8} bar and never between; the oxygen at the match is either 8.8461×10^{-7} or 8.8892×10^{-7} . Within a group the columns agree to 10^{-6} or better. The two groups are two grids: with the inserted node removed so the indices line up, the pressure difference is *zero at the pinned bottom* (16.73098 bar in both) and grows linearly in level index to 2.496×10^{-2} at the top — the signature of a grid laid out uniformly in *altitude*, $dz = (\text{top_atmos} - \text{bottom})/nz$ (`vertical_grid`, `photochem/src/photochem_eqns.f90`), with two different `top_atmos`. Nothing pins `top_atmos` at the exit: `robust_step` (`photochem/photochem/extensions/gasgiants.py`) re-pins it only every `freq_update_TOA = 1000` internal steps and then accepts any top pressure within a **factor 3** of the requested one. The 2.1×10^{-6} K decides which phase of that cycle the integration exits on, and no more.

Excluded by measurement. (i) *An amplification of the temperature:* sorted by the deep boundary temperature the seven runs interleave — `g401` at 424.95340782 K carries the lower model top, `g400_1em4` at 424.95340736 K the higher — and a pair 3.0×10^{-7} K apart within one group agrees to 4.2×10^{-6} in O/H while a pair 4.6×10^{-7} K apart across the groups differs by 5.2×10^{-3} . (ii) *The photochemical stopping tolerance:* four runs repeated with `conv_longdy` $10^{-2} \rightarrow 10^{-4}$, `conv_longdydt` $10^{-6} \rightarrow 10^{-8}$ and `equilibrium_time` $10^{17} \rightarrow 10^{22}$ s give the same two groups and the same spreads (2.466×10^{-2} in pressure, 5.199×10^{-3} in O/H) and reproduce their loose-tolerance twins to 4×10^{-7} . (iii) *A second chemical solution:* within a group the composition reproduces to 10^{-6} across guesses spanning 300–450 K and a factor 10^5 in stopping time. (iv) *A cold-trap cell switching in the climate solve:* the climate grid is fixed by its two stated pressures and 60 layers and does not move at all, and the tropopause reproduces to 2.6×10^{-8} . (v) *Round-off amplification:* the code is bit-deterministic and the outcome is two discrete states.

The matching pressure never moved. It is an exact node inserted by `lower_profile_schema.insert_level` and reads $1.000000000000 \times 10^{-6}$ bar in every profile written. The 2.4 % is a comparison at fixed level index; EXHALE reads the file as a table over pressure.

What the chain actually inherits. At the matching level, over the seven runs: T 2.7×10^{-9} , r 4.8×10^{-9} , q_{H2} 5.7×10^{-8} , **He/H** 6.5×10^{-8} , C/H 3.1×10^{-7} , N/H 3.2×10^{-7} , O/H 5.2×10^{-3} . `n_tot` and `rho` are carried and explicitly not imposed (`input_read.f90`). The only place the 2.4 % reaches EXHALE at all is `lap_r_top_RJ`, the upper edge of

the elemental flux window (`binary_element_diffusion.f90`), which moves by $5.2 \times 10^{-5} R_p$. Against `tol = 0.05` on $|F - \Phi|/|F|$, and against the 0.084–0.085 flux-window spread the two LHS 1140b arms carry, the handoff's own reproducibility is six orders below the tolerance in He/H and does not enter the H residual at all. Worth stating: the closure re-runs the adapter in every iteration directory, so the oxygen reservoir carries a 0.5 % step between iterations of one ladder — nothing in the He I 10830 chain reads it, but a result that read the oxygen would inherit it.

The change. `src/utls/lower_profile_schema.py`: `insert_level` interpolated every column linearly in $\log p$, including `n_tot` and `rho`, which are not free — $n = p/(k_B T)$ holds at every node the solution was computed on. Measured on these columns the identity held at 101 of 102 levels and was violated at **exactly one, the matching level**, by 7.7×10^{-4} and 1.24×10^{-3} . It now recomputes `n_tot` from the inserted pressure and temperature and `rho` from the interpolated mean molecular weight. Verified by rerunning one case: exactly two cells of the file change, the header and `solution_id` are unchanged, $n_{\text{tot}}/(p/k_B T) - 1$ over the whole table goes from 7.706×10^{-4} to 2.2×10^{-16} , and the run-to-run spread of `n_tot` at the matching level falls from 4.646×10^{-4} to 4.682×10^{-10} . Also corrected: the module docstring of `src/utls/element_flux_closure.py` stated the convergence test as $|F - \Phi|/\max(|F|, |\Phi|)$ while the function implements — and its own docstring documents — $|F - \Phi|/|F|$.

Not repaired, and proposed instead. The grid state at the exit is Photochem's. Re-pinning the grid and re-testing convergence before `reached_steady_state` is set is the right repair and needs a rebuild of `photochem/`; doing the same from the adapter after `photochemical_steady_state` returns (`pc.update_vertical_grid`, then step to convergence again) needs no rebuild but changes the converged answer of every future handoff column, against which the stored `crossings_gm25` and `crossings_pc090` ladders were measured. Recorded as a proposal, not made.

Scope. No EXHALE Fortran source or binary changed, no golden was refreshed and `make check` was not run: the `lower_profile` regression case reads a stored profile and never runs the adapter, so the schema change cannot reach it. Not measured: another reservoir value, another planet, another K_{zz} , and whether the two grid groups are the only two.

Files. `docs/lower_profile_deep_boundary_sensitivity.md` (the diagnosis); `LHS1140b/exhale/deep_temperature_sensitivity` (the record: README, six scripts, three tables, thirteen adapter runs); `src/utls/lower_profile_schema.py` (`insert_level`); `src/utls/element_flux_closure.py` (the docstring). The observation this answers: `LHS1140b/exhale/crossings` § 7 and § 91 above.

93. Rows of one residual that differed by a million, and a third body corrected only halfway (2026-08-29)

The judgement first. The molecular-network audit of 2026-08-27 left two items open in the molecular network: a NaN abort at He/H = 1 that it did not diagnose, and the suspicion that the molecular `hybrd1` system fails at high helium because its rows are not scaled. Both are settled here, and they are **not the same problem**.

- **The rows were the problem they were suspected of being, and scaling them fixes it.** Every row of the molecular system is now divided by its own turnover rate — the rate at which the species that row balances is produced or destroyed in that cell — so what `hybrd1` minimizes is the *relative* imbalance of each species balance instead of an absolute reaction rate. The fraction of solver attempts ending on MINPACK's `info = 4` ("iteration is not

making good progress”) falls from **54.6 % to 0.01 %** at He/H = 1000, from 41.3 % to 0.01 % at 100 and from 36.3 % to 0.26 % at 10.

- **The third-body fix of the earlier section had reached only one of its three reactions.** `ionization_equilibrium` set `ieq_cell%ntot` from `calc_ntot`, which is what the R12 term of the residual reads, but still called `set_mol_coeffs(T, n_in_dim)` — and R13 and R15 fold the third body into *their coefficient*, so both were still evaluated with ρ/m_{H} . Corrected; the goldens move a second time, for the same reason they moved the first.
- **The He/H = 1 NaN is in the hydrodynamics, not in the network**, and it is now located exactly: face $j = 232$, $r = 1.1398 R_p$, where the reconstructed right state carries $p = -3.65 \times 10^{-2}$ in code units and the first invalid operation of the whole run is `aR = sqrt(g*pR/rhoR)` (`src/modules/flux/Num_Fluxes.f90:46`). Row scaling does not move it by one step. It is recorded as an open item, not patched, because the patch belongs outside the molecular branch.

The measurement

backup/regression/mol_diffusion (hot Uranus, molecules + He 2^3S + binary H/He diffusion + a base.inp hand-off), copied to scratch with **only** He/H number ratio in `input.inp` and `HeH_base` in `base.inp` changed together, `OMP_NUM_THREADS=1`, `EXHALE_MAXSTEPS=12000`. Three binaries: the tree as found, the tree with the third body corrected, and the tree with the third body corrected **and** the rows scaled. The info histogram is a new run-wide diagnostic (`ieq_n_mol_info`), reported next to the clamped-cell count; “fail” is `info = 4` as a fraction of all molecular hybrid attempts.

build	He/H	steps	NaN	clamped	info=1	info=4	fail	$\log_{10} \dot{M}$
as found	1	12000	no	614	6031444	56627	0.93 %	11.06
as found	10	12000	no	166	5405217	2169365	28.6 %	11.04
as found	100	12000	no	212	4972511	3499546	41.3 %	11.27
as found	1000	12000	no	16	4360657	5242809	54.6 %	11.11
+ third body	1	3406	yes	507	1677801	148735	8.1 %	8.91
+ third body	10	8827	yes	274	3754138	2139333	36.3 %	8.85
+ third body	100	12000	no	216	5051123	3285000	39.4 %	11.26
+ third body	1000	12000	no	7	4368719	5222557	54.5 %	11.11
+ row scaling	1	3406	yes	337	1702668	48098	2.8 %	8.91
+ row scaling	10	5474	yes	259	2758854	7098	0.26 %	8.87
+ row scaling	100	12000	no	408	6048842	670	0.01 %	11.26
+ row scaling	1000	12000	no	38	6048079	450	0.01 %	11.09

These are relaxation snapshots followed by the steady solver, not a composition study; the $\dot{M}$ column is a diagnostic of where each run went. What it does show is that the row scaling is a change of *path*: at He/H = 100 and 1000, where every arm completes, $\dot{M}$ moves by 0.01 and 0.02 dex while the failed-attempt fraction falls by three and a half orders of magnitude. A positive diagonal scaling of a residual has exactly the same zeros, so this is what should happen.

The clamped-cell count does not follow the info count, and that is worth saying plainly. At $\text{He}/\text{H} = 100$ it rises, 216 to 408, and at 1000, 7 to 38. The clamp counts cells where *no* starting point produced an admissible root; with the rows scaled, `hybrd1` now reaches its tolerance on nearly every attempt, and a converged root lying on a face of the simplex (a fully dissociated or fully ionized species) is rejected by the admissibility test where a non-converged root that happened to sit inside it was accepted before. The counts are rare either way — 408 clamps against 6.05 million cell solves, 7×10^{-5} — and the solution does not move with them. The suspicion recorded earlier was that row scaling drives the simplex failures; the measurement says it drives the **solver** failures, by three and a half orders of magnitude, and that the residual clamps are a separate, much smaller population.

What the scale is

A row of this system is a production–loss balance in $\text{cm}^{-3} \text{s}^{-1}$, and its magnitude is set by the element that carries it: the helium rows run as n_{He} times a rate, the hydrogen and molecular rows as n_{H} times a rate, with a second density inside every bilinear term. At $\text{He}/\text{H} = 1000$ the two blocks of one residual vector differ by about 10^6 . MINPACK's `hybrd1` scales the *variables* (its internal `diag`) and never the rows: the dogleg step minimizes $\|J \delta x + f\|_2$, in which a block 10^6 below the other carries no weight, so the unknowns that block determines — the molecular fractions — are left wherever the helium block puts them.

The scale each row is divided by is **the sum of the magnitudes its terms reach when every species is set to the whole of its element**: each rate coefficient of the row times the cell's n_{H} , n_{He} , n_e and n_{tot} . It is an upper bound on the row, strictly positive wherever the row carries any reaction, and — this is the part that matters for the solver — **constant across the cell's solve**, so the scaled residual is the same smooth function of x that the finite-difference Jacobian of `hybrd1` assumes. The electron density used is the cell's incoming n_e and not the neutrality bound $n_{\text{H}} + 2n_{\text{He}}$: in the shielded molecular base the bound overstates the electron terms by many orders and would scale those rows into insignificance, which is the failure this is meant to remove.

`set_mol_turnover_rates(System_HeH_mol)` fills rows 1–8 and `set_mol_metal_turnover_rates(System_HeH_mol_metals)` the metal rows above them, so the whole system reaches the solver equilibrated; a metal row runs as n_X times a rate and a trace element carries $n_X \sim 10^{-4} n_{\text{H}}$, so it sits below the helium block for the same reason the molecular one does. The identity rows `metal_rows` writes for an absent element, and for the X^{++} of a two-stage element, are already dimensionless and keep a scale of 1. The routine mirrors `mol_heh_rows` term by term and says so at the code site: a reaction added to a row has to be added to its scale as well, and a missing term only makes the scale a weaker bound — it cannot make the root wrong.

The third body, the other two reactions

The earlier account is right about what the third body is and wrong about how far the fix reached. R12 takes n_M in the residual, through `ieq_cell%ntot`, and that path was corrected. R13 and R15 are written in `mol_rates` as *two-body-equivalent* coefficients, $k = 3.2 \times 10^{-29} n$ and $k = 8 \times 10^{-33} (300/T)^{0.6} n$, so the third body enters when the coefficient is built — in `set_mol_coeffs`, which `ionization_equilibrium` still called with $n_{\text{in_dim}} = \rho/m_{\text{H}}$. Both now get `n_tot(j)`, the same particle density the R12 term reads. R13 is the H_3^+ source at the base and R15 the H_2 source, so

this is not a small correction to a small term: it is the pair that builds the molecular base, and it was running with a third body $2.3\times$ too dense at the hot-Uranus base and up to $4\times$ in a helium-dominated one.

The corrected third body is what brings the NaN abort back. The tree as found does not abort at any He/H tested — the earlier abort at step 5538 is absent from it — and with R13/R15 corrected the He/H = 1 arm aborts at step 3406 and the He/H = 10 arm at 8827. The physically correct third body is kept; the abort it exposes is diagnosed below and recorded as an open item.

Where the He/H = 1 NaN is

Reproduced in a build with `-fcheck=bounds,do,mem` and `-ffpe-trap=invalid,zero,overflow` (the detector `backup/regression/run_fcheck.sh` uses, built to `EXHALE_fcheck.x` so it does not overwrite the production binary). The first invalid operation of the run is at step 3406:

Program received signal SIGFPE

```
#3 __numerical_fluxes_MOD_num_flux at src/modules/flux/Num_Fluxes.f90:46
#4 __rk_integration_MOD_rk_rhs      at src/modules/time_step/RK_rhs.f90:62
```

Line 46 is `aR = sqrt(g*pR/rhoR)`. A second build, instrumented to print the face states, names the face and the state that reaches it:

```
face j = 232   r_edg = 1.13983 R_p
WL (rho,v,p) = 2.9894   26.707   +0.23329
WR (rho,v,p) = 2.8916   26.816   -0.036499
cell 231 (rho, rho v, E) = 2.7948   75.246   1013.68
cell 232                = 2.8916   77.540   1040.21
cell 233                = 5.8744  152.863   1988.84
```

The reconstructed right state has a **negative pressure**, and HLLC takes the square root of it. The layer it sits in explains how: below $r = 1.139$ the run has developed a cold hypersonic shell — T about 650 K, v about 10^7 cm s⁻¹, **Mach 50 to 60** — riding on a dense hot wall at $r = 1.144$ (ρ up by $20\times$, T 9100 K). At Mach 60 the thermal pressure is 2×10^{-4} of the total energy the scheme carries, so the pressure recovered as $(\gamma - 1)(E - \rho v^2/2)$ is the difference of two numbers that agree to four digits, and a reconstruction across the jump at the next cell tips it below zero. Nothing between the reconstruction and the Riemann solver tests the reconstructed pressure for positivity.

This is a hydrodynamic robustness gap, not a molecular one: it is identical, to the step and to the face, with and without the row scaling, and the positivity floor that would close it belongs in the reconstruction, outside the `thereis_mol` branch this series is confined to. Recorded as `TO_BE_DONE.md` item (O).

Containment, and the goldens

Every change is inside `thereis_mol` — the two `set_mol*_turnover_rates` calls and the corrected `set_mol_coeffs` argument sit in the `if (thereis_mol)` cell-state block of `ioniz_eq`, the scaling is applied only in the two molecular

residual routines, and the info histogram is counted only in the molecular attempt loop. Two measurements say so rather than the code reading alone:

- `make check` passes byte-identical on `wasp_full`, `wasp_he23off` and `lower_profile` — the three non-molecular cases — against goldens that predate the series.
- The LHS 1140b arm `LHS1140b/exhale/crossings_gm25/rs_heh0p425`, re-solved in a scratch copy from its stored restart with the pre-series binary and with the new one, writes **byte-identical** `Hydro_ioniz.txt`, `Ion_species.txt`, both `_adv` files, both `_IC` files, `Cooling_breakdown.txt` and `Heating_breakdown.txt`. The stored run directory was not touched.

The four molecular cases move, and the third body is what moves them. Each was also re-run with the third body corrected but the rows *unscaled*, so the two changes can be told apart. Largest relative difference over the column against the previous golden, with the median in brackets (these are 12000-step relaxation snapshots, so the maximum sits in the base sound-wave layer or in the thin outer cells and the median is the better measure of the profile):

case	ρ	p	T	cool	H ₂ front	$\log_{10} \dot{M}$
mol_base_handoff	9.1 % (0.20)	14.9 % (0.32)	6.1 % (0.087)	53.5 % (1.8)	1.1617 → 1.1577	10.57 → 10.58
mol_metals	6.5 % (0.16)	9.8 % (0.34)	6.2 % (0.093)	68.3 % (0.67)	1.1637 → 1.1597	10.58 → 10.58
mol_lyman_werner	27.6 % (0.57)	45.3 % (0.71)	17.4 % (0.32)	59.8 % (2.4)	1.1321 → 1.1271	10.58 → 10.49
mol_diffusion	9.1 % (0.20)	14.9 % (0.32)	6.1 % (0.087)	53.5 % (1.8)	1.1617 → 1.1577	10.57 → 10.58

The row scaling contributes almost none of that. Measured on its own — the third-body-corrected build against the third-body-corrected **and** scaled build, same case, same steps — the largest difference over the column is ρ 0.019 %, p 0.0066 %, T 0.018 % on `mol_base_handoff`; 0.28 %, 0.46 %, 0.23 % on `mol_metals`; 0.013 %, 0.020 %, 0.0079 % on `mol_lyman_werner`; and 0.017 %, 0.0068 %, 0.016 % on `mol_diffusion`. The H₂ front does not move at all in any of the four.

That is the judgement the goldens were refreshed on. The move is the physics correction, and it is the correction the earlier audit intended and described: R13 builds H₃⁺ at the base and R15 builds H₂, both linearly in the third-body density, which was $2.3\times$ too large at this hot-Uranus base. The H₂ front shifts by $0.004 R_p$ and the mass-loss rate by 0.01 dex on three of the four cases and by 0.09 dex on the Lyman-Werner one, where the photodissociation front and the three-body formation front sit on top of each other and the base is therefore the most sensitive. The row scaling, in contrast, leaves the solution where it was to a few parts in 10^4 while removing the solver failures — which is what a diagonal rescaling of a residual should do, and is the evidence that it is a change of path and not of physics.

Goldens re-snapshotted for the four molecular cases at the end of the series and `make check` re-run afterwards.

Files changed

- `src/modules/nonlinear_system_solver/System_HeH_mol.f90` — `mol_inv_turnover` and `set_mol_turnover_rates`; the scaling applied at the end of `ion_system_HeH_mol`.

- `src/modules/nonlinear_system_solver/System_HeH_mol_metals.f90` — `set_mol_metal_turnover_rates` for the metal block; the scaling applied at the end of `ion_system_HeH_mol_metals`.
- `src/modules/radiation/ionization_equilibrium.f90` — `set_mol_coeffs` given `n_tot(j)` instead of `n_in_dim(j)`; the two turnover-scale calls; `ieq_n_mol_info`, the run-wide `hybrd1` exit-code histogram of the molecular solves.
- `src/EXHALE_main.f90` — the info histogram reported at the end of a run, next to the clamped-cell count.
- `TO_BE_DONE.md` — new item (O), the negative reconstructed pressure.
- `docs/molecular_chemistry_audit_he_rich.md` — section 7, closing both open items of its section 6.

A Code-size summary: EXHALE vs. the original ATES (2026-06-07; updated 2026-08-30)

A line-by-line comparison of the Fortran source against the pristine upstream code (ATES/ATES-Code-main) quantifies the extent of the in-house extension. *Pure code lines (SLOC)* exclude blank lines and comments (the comment stripper is string-aware, so a `!` inside a quoted string is not treated as a comment); the line-by-line diff additionally normalizes whitespace and case, so reformatting and comment-only edits are *not* counted as changes. The measurement script is `src/utils/codesize.py`, which reprints every number of this appendix and the byte-identical list of Table 2. Updated 2026-08-30, i.e. after the binary element-diffusion layer and the lower-atmosphere profile handoff (sections 64–80) and the defect-fix and solver series of sections 81–93; the *core* rows exclude the ported Wind-AE solver (listed separately). Relative to the 2026-08-15 edition, `species_diffusion.f90` was deleted with the Phase-D rewrite and is superseded by `binary_element_diffusion.f90`.

	Fortran files	code lines (SLOC)
Original ATES (ATES-Code-main)	47	3 514
EXHALE core (excl. Wind-AE port)	75	16 272
net change (core)	+28	+12 758 ($\approx 4.6\times$)
Wind-AE IC-generator port (<code>wind_ae/</code>)	+36	+4 746
EXHALE total	111	21 018 ($\approx 6.0\times$)

Line-by-line diff, core (comments + blanks excluded; whitespace/case normalized): **13 457** code lines added in EXHALE (new modules + the new side of modified lines) and **699** removed from the original (deleted + the old side of modified lines), for **14 156** total changed lines. At the file level, of the 47 original files **34 were modified** and **13 are unchanged** (none were removed — the unchanged set is the 10 included MINPACK routines plus three small physics/setup pieces; the full list is Table 2), and **28 new core modules** were added:

New module	Purpose	SLOC
functions/binary_element_diffusion.f90	conservative binary H/He element diffusion (stage-resolved pair coefficients, donor-cell drift, metals slaved to H)	1164
tests/diffusion_tests.f90	test driver of the diffusion operator (analytic and conservation checks)	1025
time_step/steady_newton.f90	JFNK steady-state solve (PTC warm-up, line search, below-escape diagnostic)	636
files_IO/lower_atmosphere_profile.f90	lower-atmosphere profile reader: log- p interpolation, elemental reservoirs, $K_{zz}(p)$	367
radiation/charge_exchange.f90	charge-exchange network (Huang Table 4)	366
radiation/excited_hydrogen.f90	H($n=2$) model + Balmer feedback	288
time_step/viscous_conduction.f90	molecular viscosity + heat conduction (Crank–Nicolson)	264
nonlinear_system_solver/System_HeH_mol.f90	coupled H/He + molecular ($H_2/H_2^+/H_3^+/HeH^+$) system, turnover-scaled rows	210
radiation/opacity_models.f90	opacity dispatcher (A/C/P/T)	180
nonlinear_system_solver/System_HeH_metals.f90	coupled H/He+metals ionization (+analytic Jacobian)	150
nonlinear_system_solver/newton_solver.f90	analytic-Jacobian Newton (+hybrd1 fallback)	128
nonlinear_system_solver/ion_residual_core.f90	shared ionization-residual assembly of the System_* family	127
functions/composition.f90	single-point species/EOS policy	119
init/species_table.f90	metal + molecular species metadata table	112
files_IO/metals_input_read.f90	parse metals.inp	111
time_step/energy_semi_implicit.f90	semi-implicit energy solver	108
nonlinear_system_solver/System_HeH_mol_metals.f90	molecular + metals merged ionization system	105
lower_atmosphere/h3p_cooling.f90	H_3^+ infrared cooling (Miller et al. 2013)	104
lower_atmosphere/mol_rates.f90	molecular reaction rates (Koskinen et al. 2022, Table 1)	103
lower_atmosphere/lower_column.f90	analytic lower-atmosphere column (Koskinen et al. 2022)	90
radiation/lya_rt.f90	Ly α escape-probability RT (+ absorbing-bottom channel)	88
files_IO/opacity_input_read.f90	parse opacity.inp	85
flux/low_mach_dissipation.f90	gated fourth-difference contact-mode dissipation	79
nonlinear_system_solver/ion_cell_state.f90	named-field cell state of the ionization systems	77
radiation/hydrogen_n2_rates.f90	H($n=2$) rate coefficients (single source)	76
nonlinear_system_solver/System_HeH_TR_metals.f90	merged He 2^3S +metals system	68
time_step/steady_residual.f90	steady-residual assembly/norms	57
lower_atmosphere/lyman_werner.f90	H $_2$ Lyman–Werner photodissociation (Draine & Bertoldi 1996)	37
new-module total		6 324

So the 28 new core modules contribute 6 324 of the 13 457 added code lines; the remaining $\sim 7\,130$ added (and 699 removed) lines are edits spread across the 34 modified files — dominated by the cooling module (Cool_coeff.f90, 442 $\rightarrow$ 3 277 raw lines: CHIANTI metal-line fits, Fe tables, fine-structure statistical equilibrium, Penning and molec-

ular rate coefficients), followed by the ionization equilibrium and its diagnostics (ionization_equilibrium.f90, util_ion_eq.f90), the main driver, parameters/input parsing, IC construction, cross sections, and the post-processor, which wire the metals, molecules, opacity dispatcher, Roche potential, transonic/auto/Wind-AE IC, excited-H/Ly α physics, diffusion hooks, and the Brent/Newton solver upgrades into the existing solver. In total the EXHALE core is now $\approx 4.6\times$ the original code base (3 514 $\rightarrow$ 16 272 SLOC), touching most of the original files (34 of 47 modified, none deleted, plus 28 new); with the ported Wind-AE IC generator the full Fortran tree is 111 files / 21 018 SLOC ($\approx 6.0\times$). Beyond Fortran, the tree also carries the Python tooling (GUI, the EXHALE_transit.py transmission post-processor, plotting/analysis, the lower-atmosphere adapters and the elemental-flux closure driver) and $\approx 71\,000$ lines of documentation (94 md/tex files: the memos, the manual and this log) under docs/.

B Fortran source inherited unchanged from ATES

Of the 47 upstream Fortran files (ATES/ATES-Code-main), **13 are byte-for-byte identical** in EXHALE (re-measured 2026-08-30) — the numerics EXHALE reuses without touching: the included MINPACK nonlinear solver (the hybrid Powell hybrid method and its linear-algebra helpers, called by the ionization-equilibrium solver) and three small physics/setup pieces. Table 2 lists them with their role.

Five files left the byte-identical set since the 2026-07-02 edition of this table, each for a specific, documented reason rather than drift: time_step/RK_rhs.f90 gained the gated low-Mach dissipation hook (section 60); functions/UW_conversions.f90 and states/PLM_rec.f90 carry the state-array layout change (cells,3) $\rightarrow$ (3,cells) (2026-07 refactor); flux/Num_Fluxes.f90 gained a case default abort on an unknown flux name; and nonlinear_system_solver/System_implicit_adv_H.f90 was converted to the named-field cell state (ion_cell_state) with the rest of the System_* family. As before, the HLLC and ROE wave-speed estimators (flux/speed_estimate_HLLC.f90, speed_estimate_ROE.f90) are *effectively* untouched: their original logic is preserved verbatim and only a one-line defensive floor $p_{\min} = \max(p_{\min}, 1e-30 p_{\max})$ (plus its comment) was appended in the 2026-07-02 code review, so no original line was altered.

Table 2: The 13 Fortran files byte-identical between the original ATES (ATES-Code-main) and EXHALE (measured 2026-08-30), grouped by role.

File	Purpose
<i>MINPACK nonlinear-system solver</i> (Argonne; dir nonlinear_system_solver/)	
hybrd.f90	Powell hybrid method: find a zero of n nonlinear equations in n unknowns
hybrd1.f90	Simplified user-facing driver for hybrd
fdjac1.f90	Forward-difference approximation of the Jacobian
qrfac.f90	QR factorization via Householder transformations
qform.f90	Accumulate the orthogonal factor Q from the QR factorization
r1updt.f90	Rank-one update of the QR factorization
r1mpyq.f90	Apply the Givens-rotation sequence of the rank-one update
dogleg.f90	Dogleg (trust-region) step combining Gauss–Newton and steepest descent
enorm.f90	Euclidean norm of a vector (overflow/underflow safe)
dpmpar.f90	Machine floating-point precision / range constants
<i>Hydrodynamics, grid & radiation</i>	
states/Source.f90	Gravitational + geometric source terms
init/set_gravity_grid.f90	Pre-evaluate gravity at cell centers and interfaces
radiation/J_inc.f90	Energy-dependent incident stellar spectrum for photoionization